

The Puppeteer Papers

Alvaro Rivera

2026-08-27

Contents

Characterizing the Puppet	6
What this is	6
Single-volume reading	6
The papers	6
How they relate	14
What is Puppeteer, and why it appears here	14
How to read	15
How to cite	15
License	16
Anti-porous architecture	18
TL;DR	18
Claims this paper makes	19
When NOT to use this approach	21
1. Introduction: the invisible cost of porous designs	24
2. Three faces of porosity	28
3. The unified principle: anti-porosity	32
4. An instantiation: how Puppeteer realizes the three conditions	33
5. Empirical observations	38
6. Counter-arguments	43
7. Related work	44
8. Conclusion	49
Cross-references	50
Code provenance	51
Acknowledgments	51
References	51
Appendix A: Source-code verification	53
Appendix B: Diagnostic application to an independently-constructed schema	55
Program–value separability	60
TL;DR	60
Claims this paper makes	61

When NOT to use this approach	63
1. Introduction	64
2. Consequences of separability	69
3. Realization in a concrete runtime	71
4. Verb richness: surface vs depth	75
5. Empirical results	76
6. Counter-arguments	83
7. Related work	86
8. Conclusion	89
Code provenance	91
Acknowledgments	91
References	92
Appendix A — Code references	93
Reactions and the partition	98
TL;DR	98
Claims this paper makes	99
When NOT to use this approach	101
1. Introduction: a developer says things	102
2. The pragmatic partition: now versus deferred	104
3. The double purpose of Reactions	105
4. Friction as the architectural guardian	111
5. Three consequences (ranked)	112
6. Implementation in the Puppeteer framework	116
7. Comparison with Akka and Orleans	127
8. Counter-arguments	128
9. Related work	131
10. Conclusion: a developer says things — and now also at the endpoint	133
Appendix A. Code references	134
Code provenance	137
Acknowledgments	138
Appendix B. Bibliography	138
Preserving semantic continuity across actors	140
TL;DR	140
Claims this paper makes	141
1. Introduction	142
2. Genealogy	144
3. The assumption named	148
4. Contingency	151
5. Consequences	153
6. Why existing approaches do not dissolve the assumption	155
7. Reformulation	159
8. Instantiation	163
9. Relation to previous work in this paper series	176
10. Conclusion	177

Appendix A. Code references	178
Code provenance	179
Acknowledgments	180
Appendix B. Bibliography	180
Operations from the substrate	182
TL;DR	182
1. Introduction	183
2. Genealogy: four separate operational disciplines	184
3. The assumption named	189
4. The substrate theorem	190
5. Consequences: four equivalences as variants of replay	192
6. Why existing approaches duplicate work the substrate already does	207
7. Instantiation: realization in Puppeteer	212
8. Relation to previous work in this paper series	217
9. Counter-arguments	219
10. Conclusion	224
Appendix A. Source-code verification	225
Code provenance	227
Acknowledgments	228
Appendix B. Bibliography	228
Most infrastructure layers are symptoms of the persistence model:	
a construct for auditing production stacks	230
TL;DR	230
Claims this paper makes	230
1. Introduction	231
2. Related work	232
3. The construct: infrastructural symptom	233
4. Instantiation	236
5. Auditor workflow	242
6. The ladder of compensated layers	243
7. At the limit: from datacenter footprint to a deliverable appliance	245
8. What the lens makes visible and what it does not	246
9. Extension: from layers to roles	248
Code provenance	248
Acknowledgments	249
References	249
After the substrate: building software without a datacenter	251
TL;DR	251
Thesis	252
Claims this paper makes	252
1. Introduction	253
2. Cloud as library, not as habitat	255
3. Related work	257

4. The substrate that makes it possible	260
5. The demonstration	261
6. What it means in practice (a structural permission)	265
7. A diagnostic lens, briefly: when the server-role is accidental	268
8. Limitations and counter-arguments	269
9. Conclusion	272
Appendix A — Reproducibility	273
Appendix B — Role-level diagnostic table	276
Code provenance	278
Acknowledgments	278
References	278

Inference without Authority: the three authorities that govern an output	282
TL;DR	282
1. print	283
2. The Second Voice	284
3. The Symptom	285
4. The Three Authorities	286
5. The Observer	290
6. Limits	292
7. Related Work	294
8. Conclusion	296
Appendix A. Labs	297
Code provenance	299
Acknowledgments	299
Notes	300
References	300

Identity Precedes Staging: one play, many stages	302
TL;DR	302
1. The Domain and the Deployment	303
2. Two Experiments	305
3. What a Client Sees	307
4. Actor and Audience	310
5. Identity	311
6. Recognition	313
7. Where the Narrative Lives	316
8. Limits, and Two Observations Nobody Was Looking For	316
9. Related Work	321
10. Conclusion	333
Appendix A. Labs	335
Code provenance	349
Acknowledgments	350
References	351

The Assembled Verb: repertoire, assembly, and the authorship of a verb	357
TL;DR	357
1. The Ordinary Ground	358
2. Assembly Produces a Capability	361
3. The Criterion	367
4. The Subject	373
5. The Case That Makes the Plane Visible	377
6. Returning to the Corpus: Where Is V, and Where Did the Conti- nuity Go?	380
7. The Neighbours	382
8. Limits, and What Is Not Claimed	390
9. Conclusion	394
Disclosure	396
Appendix A — The Labs	396
Appendix B — The Search, Reported as Method	400
Appendix C — The Separation, Stated Over a Small Object	403
References	410

Characterizing the Puppet

A Unifying Explanation for the Architectural Observations Developed Throughout the Puppeteer Papers

Across these papers, independent architectural observations gradually accumulate: performative verbs, durable history, reactions, authority, staging, and repertoire. This series proposes that treating them as the faculties of a single computational entity provides a coherent explanation for those observations. Whether viewed as a design abstraction or as a computational characterization, the resulting software speaks more directly about its domain.

When software is modeled as if the Puppet were its fundamental modeling unit, the software speaks more directly about its domain.

What this is

Ten analytic theory papers (Gregor, 2006, Type I) on journaled-program substrates. Each names a structural property — porosity, separability, partition, continuity, substrate, symptom, accidental category, authority, identity, assembly — and traces its consequences.

Single-volume reading

For a continuous reading — preface plus the ten papers as numbered chapters, with table of contents — see the unified monograph:

Puppeteer: Journaled Programs and the Dissolution of Infrastructure (PDF)

Individual papers below are the canonical citable artifacts; the monograph is the recommended entry point for a first read.

The papers

#	Title	Version	DOI	One-line summary
1	Anti-porous architecture: a unified design principle for CQRS + Actor + Event-Sourcing systems	0.1-draft	10.5281/zenodo.20104863	<i>porosity</i> — the representational sparsity defect that database theory, domain-driven design, REST API design, and Event Sourcing each document under a different name — and the conditions under which it is rejected.
2	Program–value separability: the structural precondition for compilation, caching, and dense journaling in a DSL runtime	0.1-draft	10.5281/zenodo.20740697	The structural property of DSL programs that makes compilation, cross-invocation caching, and dense journaling possible at all; what is commonly read as three runtime optimizations is one consequence.

#	Title	Version	DOI	One-line summary
3	Reactions and the partition: opt-in eventual consistency in actor-native systems	0.1-draft	10.5281/zenodo.20792150	2021-10-16 The developer-controlled partition of an event's implied work into <i>now</i> and <i>deferred</i> , exercised through <i>Reactions</i> ; opt-in eventual consistency falls out as a downstream consequence, not a design choice.
4	Preserving semantic continuity across actors: a tell-based approach without orchestration	0.1-draft	10.5281/zenodo.21207063	2021-10-16 The assumption naturalized across fifty years of actor-systems literature — that causation between actors is operational, not programmatic — and shows that it is a contingent historical choice rather than a necessity of the actor model.

#	Title	Version	DOI	One-line summary
5	The Journal as Substrate: Unifying Deployment, Replication, Backup, and Offline Operation in Distributed Systems	0.1-draft	10.5281/zenodo.11340646	four operational disciplines documented separately in the distributed-systems literature as instances of a single structural condition, and names the property under which they dissolve into one mechanism.
6	Most infrastructure layers are symptoms of the persistence model: a construct for auditing production stacks	0.1-draft	10.5281/zenodo.20317450	the property by which a layer of infrastructure can be discriminated as compensating for a deficiency of the persistence model rather than serving a requirement of the problem domain. Redis is the canonical worked example.

#	Title	Version	DOI	One-line summary
7	After the substrate: building software without a datacenter	0.1-draft	10.5281/zenodo.10398908	Of course the joint operational consequence of the prior six papers: under a journaled-program substrate, cloud and microservice ecosystems become libraries the program invokes rather than habitats the program inhabits, and the datacenter ceases to be a structural requirement of running production software. The server-role becomes ephemeral as a side observation; the practical consequence is the headline.

#	Title	Version	DOI	One-line summary
8	Inference without Authority: the three authorities that govern an output	0.1-draft	10.5281/zenodo.2199637	The three authorities that govern an output — the domain (what exists), the actor (what becomes observable), and the assembler (where observation occurs) — and the <i>assertion beyond warrant</i> their collapse leaves at each end: a producer that stipulates a destination it cannot know, and an observer that infers a history it was never told.

#	Title	Version	DOI	One-line summary
9	Identity Precedes Staging: one play, many stages	0.1-draft	10.5281/zenodo.2189420	A one-line summary of the property that tells a domain apart from any <i>staging</i> of it — where it runs and who observes it — and measures one domain held fixed across five stagings and six clients at zero edits, against an orthodox ports-and-adapters version of the same domain built for the comparison. The baseline <i>decouples</i> where this arrangement <i>closes</i>; recognition of a routine in the record follows as a consequence rather than a subject.

#	Title	Version	DOI	One-line summary
0A	The Assembled Verb: repertoire, assembly, and the authorship of a verb	0.1-draft	10.5281/zenodo.21135421	2013-04-21 <i>assembly</i> as the plane at which the use of capabilities becomes a capability of a subject: one recorded definition whose content is the operations, exercised by reference, with com- mand/query modality — and thereby what enters the subject's history — attributed per performance at that plane rather than inherited. A criterion any record can be asked separates recorded constitution from coordination and description, measured on third-party engines and priced against the direct baseline.

#	Title	Version	DOI	One-line summary
---	-------	---------	-----	---------------------

How they relate

The papers form a chain of preconditions: each construct names a property that becomes available once the prior papers’ properties hold.

Paper 1 establishes the vocabulary of *porosity* and *anti-porosity*. Paper 2 names *program–value separability* as the structural property that compilation, caching, and dense journaling all depend on. Paper 3 names the *now/deferred partition* exercised through *Reactions*; the partition is exercisable because the journal is dense (Paper 1) and the Reactions are programs (Paper 2). Paper 4 extends semantic continuity across actor boundaries through *tell*, which sits inside the Reactions surface of Paper 3. Paper 5 reframes deployment, replication, backup, and offline operation as instances of a single substrate property. Paper 6 names *infrastructural symptom* as the property by which compensatory layers can be diagnosed in any architecture. Paper 7 observes the joint operational consequence of the prior six: under a journaled-program substrate the datacenter ceases to be a structural requirement of running production software, and software construction passes back into the hands of those who model the domain. Paper 8 turns from execution to output, naming the three authorities that govern what an output is and where it goes — the domain, the actor, and the assembler — and the unwarranted assertion their collapse leaves at either end of the voice. Paper 9 asks what holds still while both of those change, and answers with a property of the domain that a built system can be measured for: its identity precedes any staging of it, which is why a staging can be varied without touching it.

What is Puppeteer, and why it appears here

Puppeteer is a runtime that combines CQRS, the Actor Model, and Event Sourcing with a domain-specific language whose programs are journaled as the unit of persistence. Across the ten papers it appears as the instantiation — the existence proof that the conditions each construct defines can be realized in a working system. It is neither the subject of the papers nor required by them. Each construct could in principle be instantiated by other systems: an extension of Akka or Microsoft Orleans, a fresh runtime built around a journaled DSL, or a framework that has not yet been written. The role of the instantiation is to demonstrate realizability; the contribution lies in the constructs themselves and the conditions they name.

How to read

If your question is specific, an entry point may answer it without reading the whole series:

- “*Can a domain be held genuinely fixed while where it runs and who reads it both change — and how would I check that on a system I already have?*” → Paper 9.
- “*Which of the layers in my production stack are actually structural, and which are compensating for something?*” → Paper 6.
- “*What becomes operationally possible when production software no longer needs a datacenter?*” → Paper 7. (For the auxiliary question — “*is the server-role in my architecture a necessity, or a contingent artifact?*” — Paper 7 §7 offers a diagnostic lens.)
- “*How does a journaled program actually run — across actor boundaries, across machines, across failures, across deployments?*” → Papers 4 and 5.
- “*What are the structural preconditions on which the rest of the series rests — vocabulary, compilation, consistency model?*” → Papers 1 through 3.

For the full argument, read in order from Paper 1 through Paper 9. The papers are cumulative — each relies on vocabulary and conditions established by its predecessors. Paper 1 is the entry point.

All nine are working drafts (versions 0.1 through 0.4) and are open to feedback. Issues are welcome at <https://github.com/alvaroNCubo/puppeteer-papers/issues>.

How to cite

All papers in this series are self-deposited preprints on Zenodo and have not undergone peer review.

The following papers are archived in Zenodo with citable DOIs:

Paper 1 — *Anti-porous architecture*:

Rivera, A. (2026). *Anti-porous architecture: a unified design principle for CQRS + Actor + Event-Sourcing systems* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20404863>

Paper 2 — *Program-value separability*:

Rivera, A. (2026). *Program-value separability: the structural precondition for compilation, caching, and dense journaling in a DSL runtime* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20740697>

Paper 3 — *Reactions and the partition*:

Rivera, A. (2026). *Reactions and the partition: opt-in eventual consistency in actor-native systems* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20792156>

Paper 4 — *Preserving semantic continuity across actors:*

Rivera, A. (2026). *Preserving semantic continuity across actors: a tell-based approach without orchestration* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21207062>

Paper 5 — *The Journal as Substrate:*

Rivera, A. (2026). *The Journal as Substrate: Unifying Deployment, Replication, Backup, and Offline Operation in Distributed Systems* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21349146>

Paper 6 — *Most infrastructure layers are symptoms of the persistence model:*

Rivera, A. (2026). *Most infrastructure layers are symptoms of the persistence model: a construct for auditing production stacks* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20317450>

Paper 7 — *After the substrate:*

Rivera, A. (2026). *After the substrate: building software without a datacenter* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20398998>

Paper 8 — *Inference without Authority:*

Rivera, A. (2026). *Inference without Authority: the three authorities that govern an output* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21499637>

Paper 9 — *Identity Precedes Staging:*

Rivera, A. (2026). *Identity Precedes Staging: one play, many stages* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21894206>

Paper 0A — *The Assembled Verb:*

Rivera, A. (2026). *The Assembled Verb: repertoire, assembly, and the authorship of a verb* (v0.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.22135421>

License

This repository contains two kinds of content, licensed separately.

Papers and accompanying assets — the nine `0X-*.md` files, this README, and the asciinema recordings and GIFs under `paper7-assets/` — are licensed under Creative Commons Attribution 4.0 International (CC BY 4.0). You may copy, redistribute, adapt, and build upon them, including for commercial purposes, provided you give appropriate attribution to the author.

Code — the C# projects under `labs/`, used as reproducibility artifacts for the papers' empirical sections — is licensed under the Apache License 2.0.

The two licenses apply independently. Citing or quoting from the papers is governed by CC BY 4.0; using or modifying the code is governed by Apache 2.0.

Copyright © 2026 Alvaro Rivera.

Anti-porous architecture

TL;DR

State-oriented persistence and operation-oriented persistence are not merely different storage strategies; they encode fundamentally different representational primaries. The same encoding choice — *what the representation is shaped to record* — recurs at the domain layer (the in-memory type hierarchy) and at the endpoint contract (the wire DTO). When the choice is structure, the system pays the cost in sparsity; when the choice is operation, the system pays it nowhere because the form is sized to the instance that produced it.

Architectures that shape their domain models to fit a storage substrate produce *porous* designs under structural pressure: rich object graphs — recursive, heterogeneous, with different node and edge types — do not survive a relational schema, so domain classes either fragment across joined tables that return empty cells or collapse into wide tables of redundant primitives; endpoint DTOs accumulate optional fields to serve heterogeneous use groups behind a single contract; and persistence layers record *what state exists* without recording *how it came to be*.

This paper argues that these are surface manifestations of a single architectural decision — *structure-centric encoding*, the choice to make the representation primary the type-space the form must serve rather than the operation that determined the specific case. Four bodies of literature (database theory, domain-driven design, REST API design, Event Sourcing) have each diagnosed a face of this decision locally, in their own surface vocabulary, without naming the encoding decision behind it. We name the defect *porosity* and its principled inversion (operation-centric encoding cross-layer) *anti-porosity*, and we identify three simultaneous conditions a CQRS + Actor Model + Event Sourcing combination must satisfy to sustain it. A system satisfying these conditions exists; it is presented

in §4 as one realization, not as the foundation of the claim.

Claims this paper makes

1. **Porosity has three distinct surface manifestations, all instances of structure-centric encoding.**
 - *Domain*: the type hierarchy is recorded by its structural envelope — one table covering the universe of variants — rather than by the operation that produced the specific variant in each row. Two visible patterns: over-normalized schemas, where classes fragment across many tables joined by type and queries return rows of mostly-empty cells; and under-normalized schemas, where classes collapse into wide tables of optional fields conditioned on a **status** column.
 - *Endpoint*: a single API contract serving heterogeneous use groups records the union of fields any use case might require; each invocation expresses one case and populates only its slots. The contract is shaped by the type-space of use cases, not by the case the invocation actually expressed.
 - *Persistence*: representations capture *what state exists* (horizontal: the structural envelope of values) without capturing *how it came to be* (vertical: the operations that produced them). Three substrate-specific flavors: state-only recording in relational stores (vertical sparsity — operations discarded); fabricated edges in document stores when graph shape exceeds tree shape; and command-DTO-bloat in conventionally serialized event sourcing, where the event mirrors the input payload received rather than the script the execution actually ran.
2. **The three are not independent problems.** Each is a downstream consequence of the same upstream choice: structure as the representational primary at each layer. Once that choice is in place, every local fix recreates the pressure in a different layer.
3. **A combined CQRS + Actor Model + Event Sourcing implementation can reject all three porosities by making operation the representational primary at each layer.**
 - *Domain*: one class per role — one variant per operation outcome — with sum-type discrimination supported natively. The form a row takes is determined by the operation that produced it, not by the union of possible statuses.
 - *Endpoint*: DTOs with filterable optional parameters; the same contract serves heterogeneous use groups but each invocation selects only the slots its case expresses.

- *Persistence*: dense journals of *verbs* rather than state. Each entry records the script that executed and exactly the inputs that script consumed — not the type-space of inputs the command DTO could accept, not the structural envelope of state at a point in time. The state, however rich, is reconstructed deterministically by replay.
 - The three mechanisms project the same encoding decision (*operation*, not *envelope*) to three layers.
4. **The inverse encoding is realizable in production; existence is proof of construction, not of principle.** A system whose architecture jointly satisfies the three conditions of Claim 3 has run continuously in production since 2018. It descends from a 2005 wrapper through a sequence of internal projects, was ported from Java to .NET, and currently operates in the core of eCommerce and payments platforms — payment hubs, account-balance ledgers, KYC pipelines, customer-facing storefronts, and payment-processor integrations. The system is the Puppeteer framework; §4 describes how it realizes each condition (with §4.0 covering the historical genealogy). Because the construct *porosity* was named retrospectively, the system’s continued operation demonstrates that the three conditions are jointly satisfiable in production — not that the analytic frame explains the system’s success. Appendix B applies the diagnostic to a public schema constructed independently of the present authors. Code references throughout (Appendix A) point to the public repositories; comprehensive quantitative results are deferred to a companion case-study paper.
 5. **The contribution is identifying the unified encoding decision behind four locally-named defects.** Four bodies of work — relational database theory, domain-driven design, REST API contract design, and Event Sourcing — each documented surface manifestations of structure-centric encoding in their own vocabulary, without naming the encoding decision itself. CQRS, the Actor Model, and Event Sourcing are individually well-known constituent patterns; what we claim as novel is neither the patterns nor the detection of any single porosity, but the observation that the four traditions converge on a single architectural choice (structure as the representational primary), that *anti-porosity* names its principled inversion (operation as the representational primary), and that the inversion is implementable rather than merely conceivable. We note explicitly that other actor frameworks (Akka, Microsoft Orleans, Proto.Actor, EventStore) could *in principle* be extended to satisfy the three conditions, but none currently does (§7); the analytic frame’s independence from any specific realization is therefore *conjectural*, not established. Puppeteer is currently the only realization that jointly satisfies the three conditions, and the open invitation in §5 is the explicit mechanism by which the conjecture of broader realizability would be tested.
 6. **CQRS + Actor + Event Sourcing is interesting here for operation-centric encoding cross-layer, not for separation of con-**

cerns. These three patterns are conventionally introduced as separations: reads from writes (CQRS), concurrent units (Actor Model), events from projections (Event Sourcing). When applied together as a single discipline, they constitute a mechanism that projects operation-centric encoding to all three layers in which porosity manifests. This interpretation is available only once the encoding decision is named — porosity as the structure-centric default, anti-porosity as its inversion — rather than treating each layer’s symptom as a local problem.

When NOT to use this approach

The principle of anti-porosity has limits, and the implementation pattern that realizes it has narrower limits still. We name two regimes in which the implementation is not the right fit. In each case, anti-porosity may continue to apply as a diagnostic — porosity remains a real pathology — but the CQRS + Actor + Event Sourcing realization, as instantiated in Puppeteer, imposes costs that may not be repaid.

A. Domains where the actor cannot achieve state locality

The actor model with journal-based persistence assumes that the working set of events required to rehydrate the actor’s current state can be bounded. This holds when the actor’s automaton passes through cyclic states — created, modifiable, finalized, archived — such that, at some point in the cycle, prior events become safely evictable without altering the current state. Domains where the actor’s lifecycle has no such evictable point produce *unbounded rehydration*: every load replays a history that grows without ceiling.

Consider an airline flight. From the moment a schedule is published, an actor we may call `FlightOperation` accumulates events: the first booking, subsequent bookings, seat assignments, gate changes, boarding events, departure, landing. Throughout this sequence the state is genuinely active — answering “*is the flight bookable?*”, “*who is aboard?*”, “*has it departed?*” requires recent history. But once the flight has landed and reconciliation completes, the active state collapses: the flight no longer changes, and the remaining consumers are analytical (capacity-utilization reports, for instance) and can be served from periodic projections. At that point the entire history of bookings and gate changes can be marked evicted without compromising correctness. The actor has reached locality: its working set is bounded.

Compare this to an `Airline` actor that consolidates every flight, every passenger, every booking the airline has ever handled. Such an actor never reaches a “done” state. Its working set grows with operational history, and each rehydration replays material that has no bearing on any decision currently being made. The two designs share the same business domain; they differ in whether the actor’s responsibility has a natural end.

The structural analogy with virtual memory is exact. In an operating system, a process exhibits *locality of reference* when its accessed pages stay within a bounded working set; the OS evicts pages outside that set without affecting correctness. When locality is lost — when references span the full address space uniformly — the system enters *thrashing*: pages are reloaded as fast as they are evicted, and useful work approaches zero. The actor model with a journal exhibits the same structural property: events take the role of pages, the actor’s required history takes the role of the working set, and skips (event eviction) take the role of page replacement.

Virtual Memory	Actor + Journal
Working set bounded	Non-skipped events bounded
Locality of reference	Locality of automaton state
Page eviction	Skip / event eviction
Page size	Actor granularity
Thrashing	Unbounded rehydration

The locality property of an actor is, to a substantial degree, downstream of its naming. The word *actor* evokes *action*: a **Packer**, a **Dispatcher**, a **FlightOperation** — names that denote bounded responsibility with a natural end. Names that denote aggregations (*Inventory*, *Catalog*, *Customer*) imply unbounded responsibility, with no point at which prior events become evictable. A noun-aggregator actor is the structural equivalent of an OOP god class: state accumulates because nothing about the name suggests when it should stop. The same observation applies one level further down, to the partitioning of instances: an actor scoped per-country rather than per-shipment reproduces the problem even when the type is well-named. Granularity and partitioning are not separate decisions from naming; they are the same decision expressed at different levels.

Microsoft Orleans makes this granularity choice explicit in the very name of its activations — *grains* — emphasizing each as a small, self-contained unit. The naming heuristic developed here is consistent with that view and offers a verifiable test: if the proposed actor name is a noun-aggregator rather than a verb-doer, the design will likely fail at locality.

In typical production deployments, the cost of locality failure is partly amortized. Continuous-running datacenters confine rehydration to rare events — process restart, hardware failover — rather than steady-state operation; zero-downtime release patterns (developed in Paper 3) hide cold-start latency behind warm replicas during deployment; and journal storage is operationally cheap: it is economically negligible at relevant scales, and its access pattern — sequential append, with occasional sequential reads at rehydration — does not impose the random-I/O performance demands that drive RDBMS disk specifications. A

Puppeteer deployment is largely indifferent to disk speed, where a comparable RDBMS workload is bound by it. Locality is therefore best understood as a structural assumption of the model rather than a property whose violation is immediately catastrophic. The “when not to use” case bites when locality fails **and** the amortizing properties above do not hold — for instance, single-instance systems with frequent restarts, or workloads where actor cardinality is so high that eviction-and-reload occurs as part of normal operation.

The mechanism by which Puppeteer evicts journal events — the *skip* — presupposes locality: the framework can implement skips, but it cannot manufacture locality where the actor’s responsibility is unbounded. Locality is a property of the design, not of the framework.

B. Domains where the actor’s verbs cannot be kept fast

Each actor processes its mailbox serially: only one verb executes at a time on a given actor instance. The framework sustains thousands of requests per second per actor, at peak, on the assumption that every verb completes within a few milliseconds. A verb that blocks for hundreds of milliseconds — synchronous I/O to a slow service, an expensive query, a costly computation — does not merely slow itself; it blocks every subsequent verb in the mailbox, and ingest throughput collapses. The single-thread invariant of the actor turns from a correctness guarantee into a throughput cap.

Two patterns sustain the fast-verb invariant in practice. **I/O hoisting**: the caller resolves external dependencies — third-party calls, slow lookups, expensive joins — before invoking **perform**, and passes the resolved values as parameters. The actor receives data; it does not fetch. **Saga-phased I/O**: when the workflow genuinely requires interleaved external calls, the work is decomposed into a Saga. The actor advances one phase, releases, an external coordinator performs the I/O, and the response re-enters the actor as the next event. The actor never waits inside a single verb.

Where neither pattern applies — synchronous I/O is intrinsic to the verb, or the underlying algorithm is unavoidably slow — the actor’s serial processing is structurally incompatible with the workload. In those cases the framework’s claim of high per-actor throughput cannot be honored.

In both regimes, the principle of anti-porosity may continue to apply as a diagnostic; what fails is the realization, not the principle.

1. Introduction: the invisible cost of porous designs

This paper makes an analytic theory contribution. It identifies and names an encoding decision that prior literature has documented separately in surface vocabulary across four traditions (the construct, *porosity*); derives the design principle required to invert the decision across the three layers in which it manifests (*anti-porosity*); and presents an instantiation — a system whose underlying observations date to a 2005 prototype and whose current form has been in continuous production use since 2018 — as confirmation that the inversion is realizable. The contribution is the analytic frame; the instantiation is an existence proof of realizability. The genre is *theory for analyzing* (Gregor, 2006, Type I): analytic constructs that allow phenomena to be described, classified, and unified, with empirical evaluation supplementary. Hevner-style design science evaluation (Hevner, March, Park, & Ram, 2004) — which would require bundling artifact evaluation with construct introduction — is deferred to a companion case-study paper; the Gregor Type I frame separates the two cleanly. The paper is self-contained as a Gregor Type I analytic contribution: the frame stands on independent grounds and can be evaluated as such. Its position relative to companion papers in a larger series is additional context, not a precondition for its evaluation.

A defect recurs in mainstream software architectures: representations that carry more structure than the operations they describe ever populate. Relational schemas accumulate columns whose values depend on the row’s state — required for some, irrelevant for others. API contracts grow optional fields that serve different use groups under one route, with no contractual signal of which fields belong to which case. Event payloads serialize whatever the command DTO carried, regardless of which inputs the operation consumed. Each looks local, and each has a literature that diagnoses it locally: database theory names normalization anomalies and NULL semantics; domain-driven design names anemic and persistence-driven models; REST API design names overloaded contracts and field bloat; Event Sourcing names payload bloat and command-event coupling. Four traditions, four vocabularies, four sets of remedies. A partial cross-tradition unification has long existed under the name *object-relational impedance mismatch* (Ireland, Bowers, Newton, & Waugh, 2009), but it covers a single boundary — objects against relational stores — and treats the cost as a problem of mapping rather than as a consequence of an encoding decision. To our knowledge, no prior work names the encoding decision underneath the four surface manifestations at the level of generality this paper proposes; §7 reviews the prior unifications and positions the contribution as extension rather than discovery.

We call this defect *porosity*: the representation is shaped to record structure (the type-space the form must accept) rather than the operation that determined the specific case. Its visible symptom is widespread structural emptiness — fields

without values, columns whose meaning depends on the row, queries returning rows with cells that the application never reads, plus the absence of any record of how state came to be where state-only persistence is used. The defect emerges most visibly when complex or polymorphic abstractions are forced into relational tables (the historical default of structure-centric encoding), and intensifies as the abstractions become richer; it recurs at every layer where structure is treated as the representational primary, including endpoint contracts and event-sourcing payloads independent of any substrate. The vocabulary that names the encoding decision at the level of generality needed to see it across surfaces — beyond the object/relational boundary that impedance-mismatch literature already covers — has not, to our knowledge, been established in prior work. Developers do not say *this design is porous*; they say *this design is fine; the empty fields are normal*. Decades of structure-first modeling have made the trade-offs that produce porosity (over-normalization versus under-normalization, joins versus wide rows, NULLs versus duplicate tables, command-DTO-as-event versus per-call script) feel like inherent properties of software construction. They are not. They are the consequence of one architectural decision — structure as the representational primary — repeated millions of times across surfaces. Naming the decision is the first step toward seeing it.

This paper argues that porosity is not inevitable, that the four traditions cited above describe surface manifestations of a single encoding decision (structure as the representational primary), and that their independent local remedies are partial solutions to a problem that admits a unified one — the inverse decision applied across surfaces. We name the unified rejection of porosity *anti-porosity*, identify the conditions under which a CQRS + Actor Model + Event Sourcing combination sustains it across all three layers in which it manifests (domain, endpoint, persistence), and present a system that satisfies those conditions as confirmation that the principle is realizable. The historical genealogy of that system — how it came to satisfy the conditions before they were named — is presented in §4.0, alongside the mechanisms it uses to do so.

The paper is organized as follows. §2 enumerates three layers of porosity and argues that each is a face of the same defect. §3 introduces anti-porosity as a unified design principle and states the conditions a system must satisfy to realize it. §4 describes the origins and mechanisms by which one such system, Puppeteer, realizes those conditions; the section is descriptive of an instantiation, not constitutive of the principle. §5 reports operational observations drawn from production deployments. §6 addresses anticipated counter-arguments from the principle. §7 places this work in the context of related literature. §8 concludes.

1.1 A formal characterization

Porosity is not a database problem, nor an API problem, nor an event-sourcing problem. It is a defect of the representational primary — the architectural choice of what the representation is shaped to record. Four formal traditions, each diagnosing the defect from its own surface, converge on the same encoding

decision.

The defect, stated in one sentence:

*Porosity is the property of a representation that encodes **structure** (a schema, contract, payload shape, or state form) when encoding the **operation** that determined the specific case would preserve what that case actually contributed. Structure is **type-spaced**: the form covers the universe of cases the representation must serve. Operation is **instance-spaced**: the case that this invocation, this row, or this execution produced. Where the representation is shaped for the type-space, capacity that the instance does not use is paid in sparsity.*

The dual:

Anti-porosity, or density preservation, is the choice to make the operation the representational primary, encoded at each layer as an instance-shaped artifact — the script at the persistence layer, the sum-type variant at the domain layer, the selected use case at the endpoint layer — and to derive structural projections only where downstream consumption requires them.**

The defect has two visible forms. **Horizontal sparsity**: slots populated only by some instances of a shared form (NULL columns by status, optional DTO fields by use case, ES payload fields the execution did not consume). **Vertical sparsity**: the operational trace that the operation determined but the form did not capture (state schemas that record terminal values without the operations that produced them). Both are symptoms of the same encoding decision — structure as primary.

The four formalisms that follow each illustrate this defect in their own vocabulary; their convergence is on the encoding decision, not on a single surface.

Graph theory. Domain models can be formalized as *semantic graphs*: nodes typed by entity class, edges typed by relation, with both nodes and edges potentially heterogeneous in type (Diestel, 2017). The graph is itself an operational object — edges express verbs that relate nodes — but conventional relational projection captures only the nodes’ structural attributes, discarding the operational character of the edges. The graph-database and Semantic Web literatures (Lassila & Swick, 1999; Robinson et al., 2015) preserve graph shape directly; the relational literature does not.

Information theory. This paper’s terminology — *structural capacity* (bits available to encode the type-space of cases the form admits) and *informational content* (bits the specific instance actually contributes) — takes its inspiration from, but is not derived formally from, Shannon’s framework (Shannon, 1948). The lineage is explicit; the derivation is not. Shannon’s *entropy* and *channel capacity* are defined under assumptions (equiprobable states, channel noise as transmission perturbation) we do not establish here; the terms above are introduced to characterize representational shape, not channel behavior. Structure-

centric encoding fixes the form to cover the type-space; the instance’s contribution occupies a subset of that capacity, leaving the remainder as structural excess.

Storage engines. Each storage substrate exposes what we will call its *structural alphabet*: the set of shapes it can natively express. The term is introduced here; survey treatments (Hellerstein, Stonebraker, & Hamilton, 2007) catalogue substrates and their internals without isolating this property as an abstraction. Relational stores express rows $\times$ columns $\times$ types; document stores express nested trees of typed values; graph stores express labeled directed graphs. These alphabets are all *structural* — they capture forms that data can take, not operations that produce data. An operation-centric substrate exposes a different alphabet: the primary unit is verb-instance (an executable script), not row-shape. The journal substrate developed in §4.3 is one such case.

Database normalization theory. Codd’s relational model (Codd, 1970), refined through successive normal forms (Codd, 1972; Fagin, 1977), is the canonical historical formalism for representing data in the relational alphabet. Normalization is prescribed to eliminate update, insert, and delete anomalies. The anomalies are themselves diagnostic of structure-centric encoding: they arise because the relational model records state-as-form without the operations that produced it. The procedures that resolve them — splitting wide tables, introducing join columns, accepting NULLs in optional attributes — work within the structure-centric paradigm; they do not question the encoding decision.

Synthesis. Each tradition saw the defect from its own layer and named a local remedy. The unification proposed here is not over four symptoms but over one decision: across substrates and surfaces, the representational primary was chosen as structure rather than operation. Once the encoding decision is named, the four traditions become four windows onto the same defect.

The vocabulary of the definitions above is partly borrowed and partly introduced: *semantic graph* from graph theory; *sparsity* generalizes the information-theoretic notion; *structural alphabet* and *structural capacity / informational content* are this paper’s terms; *normal forms* and *anomalies* from database theory. The contribution is the encoding-decision frame — naming the architectural choice (structure-as-primary) that the four traditions each documented in their own surface vocabulary without naming as a single decision.

1.2 Terminology

This paper uses several near-synonyms to refer to aspects of the same conceptual unit. We distinguish them once and use them consistently.

- *Operation*: the conceptual unit — a verb applied to its consumed inputs. This is the artifact treated as the representational primary under operation-centric encoding.
- *Verb*: the action body defined once in the domain library (a method or

procedure with declared parameters). The same verb is invoked many times with different inputs.

- *Invocation*: the act of applying a verb to specific inputs at a specific time.
- *Script*: the textual representation of an invocation — the form persisted in the executable journal (§4.3). A script is the program-form of an operation.
- *Execution*: the runtime evaluation of the script against the current domain library, producing state.
- *Event*: in *conventional event sourcing*, the serialized DTO record produced after an invocation. This paper reserves *event* for the conventional ES usage and uses *script* or *operation* for the operation-centric form; the distinction is foundational (§2.3, §4.3) and not a naming preference.

The terms describe the same lifecycle phenomenon viewed from different vantage points: the verb is defined once, an invocation occurs, the script captures it, the execution produces state. The persisted artifact under operation-centric encoding is the script; the persisted artifact under conventional event sourcing is the event.

2. Three faces of porosity

2.1 Porosity in the domain layer

The domain layer’s primary structural artifact is the type hierarchy. Under structure-centric encoding, that hierarchy is projected onto a homogeneous form — a table — whose shape covers the type-space (all possible variants); each row occupies only one variant. The sparsity that follows is the cost of recording structure-of-the-class rather than the operation that produced the row.

In object-oriented modeling, the canonical pattern for representing entities that pass through distinct lifecycle states is inheritance: each state is a subtype, with its own fields, invariants, and methods. A purchase order, for example, naturally decomposes into **Draft**, **Requested**, **Paid**, **Dispatched** — each subtype carrying only the attributes and operations that its state admits. The transition from one state to another is the construction of a new subtype instance, not the mutation of fields on a single class.

The relational model does not support this pattern cleanly. A type hierarchy can be projected onto a relational schema via one of three known idioms — single-table, class-table, or concrete-table inheritance (Fowler, 2002) — each of which produces porosity in a different form. In practice, single-table inheritance dominates: the schema collapses the four subtypes into one wide table with a **status** column and optional fields populated only for some rows. Attributes that belong logically to **Paid** are nullable for rows still in **Draft**; attributes that belong to **Dispatched** are nullable for everything else. The hierarchy is recovered at read time by inspecting the **status** field and conditionally interpreting the rest. The semantic graph — a clean polymorphic structure — has been projected

onto a homogeneous vector space, and the gap between them surfaces as NULLs.

Formally, the relational projection encodes a sum type as a product type. The state hierarchy — **Draft**, **Requested**, **Paid**, **Dispatched**, each variant with its own fields and invariants — is naturally a tagged union: a sum type. The relational schema, lacking native sum-type support, encodes it as a product type — a single tuple of fields that must coexist regardless of variant, with the discriminator demoted to a column. The mismatch is structural. In the information-theoretic vocabulary of §1.1, sparsity — structural symbols present in the representation that carry no information for any specific instance — is the inevitable cost of forcing a sum-type structure through a product-type alphabet.

The choice is not driven by ignorance of the alternative. Developers familiar with object-oriented design recognize that inheritance is the structurally cleaner pattern; they adopt the property-field approach because the substrate makes it the path of least resistance. The other two idioms impose joins by type, schema duplication, and migration overhead severe enough to outweigh the modeling benefit. Porosity here is the signature of a forced compromise — not an absence of OOP wisdom, but a substrate that does not admit it.

The downstream effects extend beyond schema shape. State transitions on the porous table become updates against a row that other transactions may also read or modify; lock contention emerges as a structural cost of the design choice, not as an accident of high traffic. The broader observation — that a domain whose lifecycle was never logically concurrent now pays for transactional concurrency it did not request — is treated separately, as a distinct symptom of persistence-as-source-of-truth, in Paper 5.

2.2 Porosity in the endpoint layer

The endpoint layer’s primary contractual artifact is the DTO — a fixed-shape tuple of fields defining what flows between client and server. The same structure-vs-operation tension reappears one layer up: the DTO is shaped for the type-space of use cases the endpoint must serve, while each invocation expresses one specific case.

When an endpoint serves a single, homogeneous use case, the DTO is well-defined and dense: each field participates in the operation’s semantics. When an endpoint serves heterogeneous use groups behind a single route, the DTO accumulates optional fields covering the union of all cases — the contract is shaped by the universe, not by the case. This mirrors the wide table with **status**-conditioned columns of §2.1, with the encoding decision identical: structure as primary at the contract layer.

From the perspective of type theory, a DTO is typically modeled as a product type: a fixed collection of fields that must coexist. Heterogeneous use groups, by contrast, correspond naturally to a sum type — a tagged union — where each variant carries only the fields relevant to its case. When a product type is

forced to encode what is semantically a sum type, sparsity appears as optional fields and conditional validation.

Consider an endpoint `POST /orders` that serves both a customer self-checkout flow and an administrative batch-insertion flow. The customer flow specifies items, a delivery address, and a payment token; the administrative flow specifies override pricing, source attribution, and an internal correlation ID. A unified DTO carries the union of all fields, validates conditionally based on caller identity or feature flags, and provides no contract-level indication of which fields belong to which use group. Documentation shoulders the burden the schema cannot.

Formally, this is again a projection of heterogeneous semantic variants onto a homogeneous vector space. In terms of information theory, the DTO's structural capacity exceeds the information actually conveyed by any single call. The unused fields constitute representational sparsity: structural symbols present in the representation that carry no information for the specific operation being performed.

The defect manifests symmetrically on the response side. The same DTO that returns an order's full detail to a desktop client returns more than a mobile client requires; clients either over-fetch or fragment the operation into multiple round-trips, neither of which the original fixed-shape schema admits. The DTO, as a homogeneous projection of underlying semantic variants, reproduces at the wire contract level the porosity already observed in relational schemas of §2.1.

Costs accumulate at multiple levels. Validation logic becomes conditional. DTO families proliferate as variants are introduced to handle slightly different use groups. Client code absorbs out-of-band knowledge about which fields apply when. The contract — intended to be the disciplined surface between systems — becomes porous: its meaning depends on contextual knowledge that the contract itself does not encode.

2.3 Porosity in the persistence layer

The persistence layer's job is to durably capture what the system has done. Three regimes recur in practice — relational stores, document stores, and conventionally implemented event sourcing — and each manifests porosity differently while sharing a common cause: each chose structure as the representational primary. Relational and document stores record state (the structure of the world at a point in time), discarding the operations that produced it — vertical sparsity. Conventionally implemented event sourcing improves on state-only persistence by recording verbs, but persists the command DTO (the structure of what was received) rather than the script of what executed — horizontal sparsity, of a different shape from §2.1 but the same encoding decision.

Relational persistence captures state, not the operations that produced it. A row records that an order is in `paid` state and stores the amount and payment

method, but cannot answer how the order arrived there or what triggered the transition. The information loss is structural and asymmetric: state can be derived from operations by replay, but operations cannot be derived from state — the function from operations to state is many-to-one. Recording only state is therefore strictly less informative than recording operations; under evolved semantics (a bug fix, a richer derivation, a new derived field) or under audit (who triggered what, in what order), the missing direction is the one that matters and the one the schema cannot supply. The state-conditioned fields of §2.1 propagate to persistence: when invariants depend on state — for instance, **amount** is required only when **status = paid** — the schema cannot enforce them. The available alphabet (foreign keys, nullability, primitive **CHECK**) admits only primitive referential integrity, never the conditional, cross-field, cross-object invariants that emerge naturally from a sum-type domain. Invariant enforcement migrates to code by structural necessity, not by stylistic choice.

Document stores partially relieve the schema rigidity by allowing nested heterogeneous shapes. A purchase order can be a document with embedded items, addresses, and payment details, sparing the joins that relational schemas require. But the alphabet admits only a particular kind of structure: the labeled rooted tree. Domain semantic graphs are not constrained to trees — they admit cycles, references shared across documents, and edges of arbitrary type. When the same item is referenced from multiple orders, or when payments link to other documents, the missing edges must be fabricated at the application layer. The substrate handles trees natively and degrades the moment graph shape is required, in the strict graph-theoretic sense of §1.1.

Event sourcing, in its mainstream conception, improves on both by recording operations rather than state: replay reconstructs state by re-applying events. But in its conventional form — where events are serialized data structures mirroring the input commands — the entire input payload is captured regardless of which inputs the operation actually consumed. Consider an operation **RecordPayment** whose interface accepts twenty parameters; internally, the state-changing logic consumes four. The conventional implementation persists all twenty in the resulting event. The event payload is a fixed-shape tuple; the unconsumed fields constitute, in the informal sense established in §1.1, structural excess alongside the consumed content. Event sourcing inherits the porosity it was intended to escape: the journal records what the operation *received*, not what it *consumed*. The distinction is itself foundational and not merely a serialization preference. *Conventional event sourcing stores events (serialized DTOs); operation-centric persistence stores scripts (executable invocations)*. They are not interchangeable forms of the same data — they are different representational primaries, and the structure of each (received-DTO shape vs. consumed-input shape) is itself a face of the encoding decision §1.1 names.

The persistence layer thus exhibits porosity in three forms: state without operations in relational stores (vertical sparsity), fabricated edges where graph shape exceeds tree shape in document stores (horizontal sparsity of structural mis-

match), and command-DTO-as-event in conventionally serialized event sourcing (horizontal sparsity of received-not-consumed). The encoding decision is identical across the three: structure — state, tree, or command DTO — is treated as the representational primary, and the operation that produced it (or the script that executed) is either discarded or recreated as structural envelope. Three layers, one decision: §3 establishes the principle that addresses them as one.

3. The unified principle: anti-porosity

A system is anti-porous when its representations, at every layer, are *operation-centric*: shaped by the instance the operation determined rather than by the type-space the form might serve. In the vocabulary established in §1.1, this means choosing operation as the representational primary — encoding the script that executed (verb + consumed inputs), the variant the operation produced (sum-type instance), or the use case the invocation expressed (per-call wire selectivity) — rather than the structural envelope into which any instance of the type must fit. In the language of type theory, this means representing heterogeneity as sum types rather than product types; in the language of information theory, it means structural capacity sized to the instance’s contribution rather than to the type-space.

A criterion of demarcation. The structure-vs-operation encoding decision is observable wherever a representation is shaped for a universe of cases rather than the specific instance: function signatures with unused optional parameters, Protocol Buffer schemas with many **optional** fields, HTML forms serving heterogeneous use groups, generic-type parameters never specialized to specific types. The present paper draws a narrower line: it analyzes layers that persist artifacts whose lifetime exceeds the operation that wrote them — the in-memory aggregate (the domain), the wire DTO at the contract boundary (the endpoint), and the durable schema or journal (persistence). Representations whose lifetime is bounded by a single operation — a function call’s argument list, a generic type parameter at one instantiation, an HTML form rendered for one request — exhibit the same encoding decision but do not pay the cost in *stored* sparsity. The frame may apply to them with appropriate adjustment; the present paper does not claim it across them.

The principle manifests as three simultaneous conditions, each inverting the structure-centric default observed in §2. First, the domain layer admits sum-type structure natively: a state hierarchy is encoded as variants (one per operation outcome), not as a product type with a **status** discriminator and conditionally populated columns. Second, the endpoint layer admits per-call selectivity at the wire: contracts express sum-type discrimination across heterogeneous use groups (one variant per invocation), not a fixed-shape DTO that unions all of them. Third, the persistence layer records operations and only the inputs the operations actually consumed: the journal carries the script that executed (an instance-shaped artifact), not the structural envelope of what was received.

The three conditions are not independent fixes; they are the consequence of a single decision applied at three surfaces. As §2 made structurally evident, each manifestation of porosity arises from the same architectural choice — to treat structure as the representational primary, the type-space as the form, and the instance’s contribution as a subset that leaves the remainder sparse. Anti-porosity is the inverse choice — to treat operation as the representational primary — applied as a single discipline across the three layers. A patch applied only at the domain (aggressive denormalization, hand-rolled inheritance) leaves the API and journal demanding structural envelopes; a patch applied only at persistence (compact event encoding) leaves the journal’s density at the mercy of upstream layers’ choices. Anti-porosity is not the sum of three local choices; it is a single architectural decision projected to three surfaces.

Operationally, anti-porosity is *density preservation*: the representation’s structural capacity matches its informational content because the form was sized to the operation’s contribution rather than to the type-space — consumed content retained, structural excess discarded. The combination of CQRS, the Actor Model, and Event Sourcing — conventionally presented as three separations of concern — admits a different reading once anti-porosity is named as the principle they jointly satisfy: they constitute, when applied as a single discipline, a mechanism for density preservation across representations. The principle does not prescribe a particular framework. Akka, Microsoft Orleans, Proto.Actor, or any system that combines actors, CQRS, and event sourcing could in principle satisfy the three conditions; doing so requires a substrate decision none of those frameworks has currently made (§7). The next section describes one realization — the Puppeteer framework — selected as the existence proof for the principle, not as its definition. The realization rests on a single architectural choice: that operations themselves — verbs invoked by callers — be recorded as the durable representation. This extends the principle of *code-as-data* (familiar from Lisp at the language layer) to the persistence layer in a narrower form than Lisp’s strong homoiconicity. §4.3 develops the property formally as *script-form persistence*, and the resulting artifact as the *executable journal*.

4. An instantiation: how Puppeteer realizes the three conditions

This section describes how the Puppeteer framework realizes the three conditions stated in §3. It is the only section of this paper in which authorial voice (“we”) and the specifics of one implementation are concentrated. The realization is presented as an existence proof for the principle, not as its definition; the conditions could be realized differently in another framework that made comparable substrate decisions (§7).

4.0 Origins of the instantiation

A practical encounter with the underlying phenomenon predates the present analysis by nearly two decades. In 2005, work that eventually became the Puppeteer framework began as a wrapper around a DLL that exposed services in the then-emerging vocabulary of *web services*. While building this wrapper, an unexpected property of the construction emerged. By intercepting the parameters of every incoming call and writing them, in order, to a log — what would now be called a *journal* — the full state of the underlying automaton could be recovered without inspecting its code. Starting from any consistent prior state and replaying the log, the automaton arrived at exactly the same final state. Persistence, traditionally treated as a separately designed concern, had emerged as a side effect of the wrapper. The phenomenon was named *autopersistence*, after its observable effect rather than from any theoretical premise. What the wrapper-log approach made visible was something the relational-persistence approach had hidden: the log was *dense* — it contained no empty fields, recording only what each call had actually used. The relational schemas the same applications wrote to, by contrast, were full of NULLs, columns that changed meaning by row, and joins that returned mostly-empty cells. The journal was *dense*; the schema was *porous*. The label and the principle followed; the observation came first. The mechanisms described in §4.1–§4.3 are the present-day form of properties already implicit in the 2005 wrapper: a domain visible to the framework only by reflection, parameters captured exactly as the call carried them, and a journal that records operations rather than state. Between the 2005 prototype and the present design, the framework was used in a sequence of internal projects, ported from Java to .NET, and progressively refined; the form described in §4.1–§4.3 has been in continuous production use since 2018.

4.1 Roles, not concepts: the library

Anti-porosity in the domain layer admits a single thesis: **a sum type, not a table**. The framework partitions the domain by *roles* — operations that act — rather than by *entities* — nouns that exist. Each role is a subtype: a class with its own fields, invariants, and verbs. The purchase order example of §2.1 is encoded directly as a sum type, with **Draft**, **Requested**, **Paid**, and **Dispatched** realized as distinct subtypes of a common base, each carrying only the attributes its role admits. The system’s domain is the union of these roles, not a flattened table indexed by a **status** column.

The classes that realize these roles form what the framework calls the *library* — pure conceptual artifacts that do not know which “play” they participate in. A **Packer** is a *class puppet*: it knows how to pack; it does not know whether the packer instance is being persisted, replayed, or audited. The framework discovers domain types reflectively at runtime, scanning the configured library assemblies (**Actor.cs:15** declares the marker base class; **DomainLibraries.cs:117–128** performs the assembly scan; **ActorHandler.cs:61** invokes it at handler construction). The domain need not register itself; the framework inspects what

is there. The continuity is structural, not anecdotal: the 2005 wrapper (§4.0) intercepted parameters without inspecting the DLL it wrapped; reflection-based discovery preserves that property as a structural feature of the current design.

Polymorphism is a first-class concept of the DSL itself, distinct from how it is realized in the host language. Argument-parameter compatibility is computed at parse time via subtype checks (`AstExpression.cs:236–246`, `AreCompatible`); runtime substitution of subtypes is handled by the symbol table (`SymbolTable.cs:134`, `IsSubclassOf` check). Where the host language (C#) optimizes for verbose precision, the DSL optimizes for script readability and for boundary decoupling. Type promotion at the call site is permissive: polymorphic arrays promote to whatever collection type the library declares (`List`, `IEnumerable`, array, and so on); parameters accept any enumeration and promote at parse time to the specific enum the library expects. The DSL writes `Credit`, the DTO carries `Credit` as text, and the library’s `PaymentMethod` enum receives the matching constant — no lexical change at any boundary. The contract is structurally decoupled from the domain class: each side may evolve its enum vocabulary independently, with the DSL mediating. **This is not syntactic sugar but a design decision: the DSL is shaped to read as domain narrative, not as a transcription of a programming language.**

Anti-porosity in the domain layer follows from the alignment of three properties: role-oriented partitioning (sum-type variants in the type hierarchy), reflection-based discovery (the framework imposes no schema on the library), and DSL-level polymorphism (sum-type discrimination is honored end-to-end in the operation contract). The remaining two layers — endpoint and persistence — are realized by separate but compatible mechanisms, treated in §4.2 and §4.3.

4.2 Parameter modifiers: `?` and `Eval`

The framework makes the direction of every value flow explicit through four parameter modifiers: `In`, `Out`, `InOut`, and `Eval`. The first three mirror the in/out/ref convention familiar from systems languages, ensuring caller-puppet isolation: the puppet cannot modify the caller’s environment outside the declared direction. The fourth, `Eval`, is unique to this realization. The modifier system serves two complementary purposes: explicit isolation at the call contract, and honesty in the operation’s journal record.

For `Out` parameters, the journal writes the literal character `?` as a placeholder (`Parameters.cs:755–757`). The reason is structural: an `Out` parameter’s value is computed by the puppet, not supplied by the caller — there is no input value to record. At replay, the deserializer reads `?` and produces a default value of the parameter’s type (`Parameters.cs:780–782`); the puppet’s logic re-executes and fills the slot. The journal thus records exactly what the call carried at the input boundary — output slots appear as honest reservations, not fictitious inputs.

`Eval` parameters address a different problem: values the puppet uses that are not supplied by the caller and are not deterministic across replays. Gener-

ated identifiers and game-session state — a random number or a gameboard captured for a game’s session — are typical cases: values that pertain to the operation’s semantics but cannot be re-derived at replay. The `Eval` modifier directs the puppet to capture the value, on first execution, as a literal-assigning script (`name = (type)(value);`), persisted as part of the operation’s journal record (`Parameter.cs:33–36`, `163–224`, `228–247`). On replay, the script re-executes and assigns the same literal — determinism restored. V1’s interpreted DSL had an `eval` command for the same purpose; V2 replaced it with the parameter modifier because the command form could not always be compiled (`EvalStatement.cs:52–55`).

The modifier system ensures the journal carries exactly what the call was:

- Caller inputs recorded as values.
- Puppet outputs recorded as `?` reservations.
- Non-deterministic values recorded as literal-capturing scripts.

Anti-porosity at the parameter level emerges from two complementary disciplines:

- Explicit direction at the call boundary.
- Explicit capture of non-determinism.

? and Eval are dual mechanisms that operationalize anti-porosity at parameter granularity. Integration into the broader journal — and the script-form property that allows replay against an evolved domain — is treated in §4.3.

Boundary cases of Eval. Two cases of the `Eval` mechanism are worth naming explicitly, both of which the production deployments cited in §5 have had to address.

External dependencies. When an `Eval` expression resolves to a value derived from external state — a currency exchange rate at the moment of invocation, a foreign-key lookup against another service — the framework captures the resolved literal but not the dependency. The journal entry is `rate = (decimal)510.0`, not `rate = CurrentExchangeRate("USD", "CRC")`. On replay against a future library, the captured literal is reproduced exactly; the external state may have changed. This is correct under *snapshot semantics*: the journal records what the system actually saw at the time of the operation, not what the world is now, and replay reconstructs the historical execution. The framework’s contract is historical faithfulness, not present truth. Applications that require present-truth on replay (uncommon, since replay is typically used for state reconstruction, audit, or upgrades against evolved library logic) must re-derive the external value explicitly in the new library rather than relying on the captured literal. The boundary case is worth naming because the contract is sometimes assumed to be present truth and the framework’s contract is not that.

Sensitive values. `Eval` captures literal values into the journal, persisted as text.

For values that should not appear in plaintext at rest — cryptographic nonces, payment tokens, PII — the framework requires deployment-level mitigation. Three mitigations are available and have been used in the production payment deployments cited in §5: (a) at-rest encryption of the journal storage substrate (available in all four shipping backends); (b) application-level redaction policies applied at write time (specific to the application, e.g., card-number tokenization before `Eval` captures the value); (c) in-script transformation so that the `Eval`’d value is already cryptographically opaque before capture (e.g., the verb captures a hash or token, not the secret). The framework does not provide column-level encryption *inside* the script representation, because an encrypted literal in a script would not parse and execute on replay; opaque encrypted blobs must be decrypted before script evaluation, which is the application’s responsibility. The point worth naming here is that `Eval`’s literal-capture property does not exempt journal storage from the standard at-rest-encryption disciplines other persistence substrates require. Conventional ORM-style serialization shares this property and the same mitigation playbook; the framework inherits both the property and the obligation.

4.3 Operations, not state: the executable journal

In the framing developed here, the primary property of event sourcing is not temporal traceability but density preservation: it preserves semantic density that state-based persistence models discard.

By storing operations rather than state projections, the journal maintains an isomorphic representation of the semantic graph that tabular or document projections render sparse. The mechanism that achieves this we call *script-form persistence*: each entry in the journal is simultaneously data — storable, indexable bytes — and a program — a parsable, executable script in the framework’s DSL. The journal does not store a serialization of the actor’s state; it stores the script that produced it. We call the resulting artifact the *executable journal*: a durable log of executable DSL scripts, replayable against the current library to reconstruct state.

The relationship to *homoiconicity* in its strong sense, familiar from Lisp (McCarthy, 1960; the term itself coined by Mooers and Deutsch, 1965), is one of *lineage, not equivalence*. Lisp’s homoiconicity has consequences — macros that operate on S-expressions as primitive data structures of the language, eval-quote symmetry, code-modifying-code at the language layer — that the present realization does not claim. The DSL scripts in the executable journal are parsed by the framework parser and executed by the framework runtime; they are not consumed as primitive data structures by other DSL scripts, which is the property that would justify the strong term. The narrower property the present paper claims, and the only one its argument requires, is that persisted journal entries are textual representations of executable programs in the actor’s DSL — retrievable as text, parsable into ASTs, executable against the current library. Bash scripts, SQL stored procedures, and shell command logs share this nar-

rower property; we therefore do not claim homoiconicity in the strong sense. *Executable journal* and *script-form persistence* are the precise terms for what the argument requires; the lineage to Lisp is acknowledged without claim of identity.

Conventional event sourcing persists the command DTO; the present realization persists the execution script. The script encodes consumed inputs only — the bloat of received-but-unused parameters does not arise by construction.

The script-form representation enables a capability that state-storing substrates structurally cannot offer: re-execution against new semantics. When the domain library evolves — a logic fix, a richer derivation, a new computation depending on past inputs — the journal replays against the updated library, producing corrected projections without altering the historical record. Storage engines that persist state, by contrast, can only restore the state that was written; the operations that produced it have been discarded.

The density preservation property is structural, not accidental. **A script cannot contain parameters the execution did not use, because the script is generated after parameter direction and evaluation rules have been applied** (§4.2). The journal therefore inherits density from the modifier discipline: caller inputs recorded as values, puppet outputs as ? reservations, non-deterministic values as literal-capturing scripts. Two operational properties verify this in code. First, density is preserved through an asymmetry between inputs received and inputs consumed: only the latter are serialized (`Parameters.cs:755–757`, `:780–782`; `Lexer.cs:600`, where ? is tokenised as `TokenType.question`). Second, the journal admits exactly two primitive operations — append and read-forward (`JournalWriter.cs:65`; `JournalReader.cs:35`) — exposed through a unifying interface (`Diary.cs`); the vocabulary of relational stores — SELECT, UPDATE, DELETE, INSERT, and the transactional locking that surrounds them — is absent by construction, not by convention.

A journal of DSL scripts therefore combines three properties — script-form persistence (executable text retrievable and re-runnable), graph preservation, and density preservation — under one representational substrate.

5. Empirical observations

The empirical claims in this paper are minimal by design, consistent with the *theory for analyzing* genre declared in §1 (Gregor, 2006, Type I): the contribution is the analytic frame (the encoding decision underneath four locally-named manifestations), for which an existence proof of realizability is supplementary. Quantitative measurements of the instantiation are deferred to a companion case-study paper; the analytic frame stands without them. The principle predicts that representations satisfying the three conditions of §3 will exhibit higher information density per entry than relational projections of the same opera-

tions; the magnitude of the resulting storage-volume gap depends on the workload regime (broadcast vs pointwise vs analytical), characterized below. The worked example illustrates the broadcast regime where the gap is largest, and the deployed instantiation listed at the end of the section confirms the principle survives in production.

A worked example. Consider a generic domain in which buyers issue purchase orders against future scheduled events; an event is *confirmed* once it has occurred; and an event is *settled* with one of several outcomes (award, refund, restatement). The pattern recurs across ticketing, conditional pre-orders, and milestone escrow. Three primitive verbs realize the lifecycle:

Phase	Statement	Scope
Pre-event ($\times n$)	<code>order = Company.Purchase(buyer, items, events, amount, currency)</code>	one order per call; ranges over items $\times$ events
Event occurs	<code>event.Confirm(date, operator)</code>	one event; implicitly all order-items bound to it
Resolution	<code>event.Settle(outcome, date, operator)</code>	one event; implicitly all order-items bound to it

The receiver determines scope; no row enumeration is required.

The journal model. The framework’s V2 pattern stores each action’s verb body once in an action library; the journal records per invocation only the action identifier, the parameter values, and minimal administrative metadata (timestamp, operator, entry ID). The journal grows with parameters, not with action bodies. (V1’s raw-script representation, in which the verbatim script appears in each entry, is preserved for backward compatibility.) This split mirrors the separation between definition and invocation in any compiled or interpreted language; in this realization, that separation extends to persistence.

Storage backends. The conditions of §3 concern the representational shape of journal entries, not the substrate that physically stores them. Four backends ship with the framework — filesystem (UTF-8 script entries in append-only files), in-memory (for testing and ephemeral actors), and two relational engines (MySQL and SQL Server) — and the script-form, density, and append-only access patterns are preserved across all four. The relational backends expose the journal through standard SQL interfaces familiar to operations teams; entries remain executable scripts queryable as text columns. Choice of backend is a deployment decision; the structural properties the paper attributes to the journal are invariant across that choice. A commonly raised concern — that a journal is a black box for operations staff who must inspect it under pressure —

is therefore largely a deployment-configuration question rather than a property of the principle.

The structural gap, and the workloads under which it appears. A `Confirm` invocation in the journal carries the parameters of one verb plus its administrative metadata — typically tens of bytes. The same operation expressed as relational mutations must enumerate every order-item bound to the event, copying parent dimensions onto every child row, because SQL’s `JOIN` cannot quantify — it materializes Cartesian projections of existing rows, it does not abbreviate them. For an event with many bound items, the relational expansion exceeds the script representation by several orders of magnitude.

The gap is regime-dependent, and presenting only the favorable case would mislead. Three regimes are worth distinguishing.

Broadcast operations over large aggregates (the worked example above): one verb expressed as the framework’s script applies to N items via the receiver’s implicit quantification; the relational form must enumerate N rows. The write-amplification ratio is roughly N . The structural gap is large — several orders of magnitude when N is in the thousands — and grows with aggregate cardinality. Events bound to many order-items, role permissions broadcast to large user populations, and price changes applied across product catalogues all fall here.

Pointwise operations (one verb, one row affected): the journal grows by one script entry; the relational form mutates one row. The write-amplification ratio is roughly 1. The structural gap on storage volume is absent. The information density per entry is still higher under the operation-centric form (the entry records the script, not the received DTO), but the storage footprint is comparable. Pointwise updates — credit one wallet, increment one counter, mark one item shipped — fall here.

Analytical workload (read-heavy, snapshot queries): the journal does not address the workload directly; analytical queries run against projections, not the journal itself. The journal grows monotonically with operations regardless of read demand; the relational projection sized to the analytical workload occupies a stable footprint determined by current state. On *storage volume*, the journal accumulates without bound while the projection does not — the ratio inverts and favors relational storage. The trade is paid in journal storage (which §4.3 and Paper 5 argue is operationally cheap and write-only) and in projection-replay latency for cold projections; the gain is the ability to rebuild any projection against evolved semantics. Whether the trade is favorable depends on workload composition and on whether projection-rebuild is a valued operational capability.

The structural gap is therefore *strongest under broadcast regimes, neutral under pointwise regimes*, and *inverts in storage volume under purely analytical workloads*. Information density per entry remains higher under operation-centric encoding in all three regimes; the storage-volume gap is regime-specific.

Why the gap is structural. The compactness of the script representation

rests on three model properties:

- **Quantification.** The receiver implicitly ranges over the affected aggregate; the verb applies to all members at once.
- **Polymorphism.** A single `Settle(outcome, ...)` accepts heterogeneous outcomes as a sum type, sharing representation across branches.
- **Parameters as metadata.** Operator, date, and outcome travel as arguments of one statement; in the relational form they are copied into every affected row, because the row is the unit of identity.

The relational form has none of these. A `JOIN` cannot quantify universally; an `UPDATE ... WHERE ...` issues the directive, but its result is enumerated rows. The porosity of the relational layer denotes here, in concrete terms, the structural consequence of substituting “ $\forall \text{ item} \in \text{event}$ ” with “ $N \text{ copies of the antecedent}$.”

Forensic observation. The diagnostic is not specific to this example. Any mature relational schema can be read for porosity directly: counting a verb’s parameters against the columns of the rows its invocation affects, identifying NULL-allowed fields whose values the originating operation never had to supply, and noting cross-product tables that materialize what the verb expressed as quantification. Appendix B applies this procedure in detail to the *Ordering* bounded context of Microsoft’s `dotnet-eShop`, a reference architecture constructed independently of the present authors, and reports specific findings against its EF Core schema. Under the regime distinctions above, the eShop case exhibits horizontal sparsity in the domain layer (per-state populated columns) and vertical sparsity at the persistence layer (`Ignore(b => b.DomainEvents)`), both substrate-induced; the write-amplification ratio of the workload is left as an open question, since the analysis is structural rather than measured.

Existence proof in production. A system whose architecture jointly satisfies the three conditions of §3 has been in continuous production use since 2018, after more than a decade of preceding refinement (described in §4.0) traceable to a 2005 prototype. Puppeteer, the framework described in §4, currently runs structurally distinct subsystems within the core of eCommerce and payments platforms: payment hubs, account-balance ledgers, KYC pipelines, customer-facing storefronts and experiences, and payment-processor integrations. Code references throughout this paper (Appendix A) point to verifiable mechanism implementations. A separate case-study paper presents quantitative observations from these deployments — endpoint latency distributions, journal growth rates, replay performance, and developer-velocity comparisons — alongside the operational details (deployment, workload, infrastructure) that fall outside the scope of an analytic paper.

What the existence proof does and does not establish. Three limitations are worth naming. *First, retrospective interpretation.* The construct *porosity* was named after the system existed (§4.0); the system therefore demonstrates that the three conditions are *jointly satisfiable in production*, not that the frame *explains* its operational success. *Second, single-realization risk.* The inverse re-

alization is currently coextensive with Puppeteer. Appendix B reduces this risk by applying the porosity diagnostic to a public reference schema (Microsoft’s **dotnet-eShop**) constructed independently of the present authors — but Appendix B analyzes a *porous* system, not an *anti-porous* one constructed by an independent party. *Third, falsifiability of the principle as such.* The analytic frame predicts that operation-centric encoding admits density preservation; it does not currently predict outcomes the system would falsify. Specifically: a system that operates well while violating the principle would show the conditions are not necessary; a system that satisfies the conditions and fails would show they are not sufficient. We have not identified instances of either.

Open invitation. Three forms of external evidence would strengthen the analytic claim and are explicitly invited. *(a)* Analyses of independently-constructed systems under the porosity diagnostic (analogous to Appendix B), particularly systems built before 2024 that satisfy operation-centric encoding through different mechanisms — Akka Persistence, Microsoft Orleans + EventStore, or domain-specific event-sourced systems would be natural candidates. *(b)* Independently-constructed systems that satisfy the three conditions of §3 and fail under operational pressure, which would falsify sufficiency. *(c)* Independently-constructed systems that operate well at scale while explicitly violating the three conditions, which would falsify necessity. The author maintains a companion repository where such analyses are linked alongside the present paper.

Note on evaluation deferral. Deferring quantitative measurements to a companion case-study is a pattern recurrent across the paper series. Under Gregor Type I the deferral is supported — the analytic frame stands on independent grounds — but each analytic paper has so far transferred its empirical debt forward, with no paper yet redeeming the full debt. Three statements bound the position.

What is already published. Companion Paper 5 (*Substrate operations*) includes labs L1–L6 reporting append latency, density-per-event, cross-DC convergence parity, and buffer/RDBMS comparisons — bounded empirical observations scoped to Paper 5’s structural claims about substrate primitives, not a comprehensive case-study of Puppeteer’s production deployment.

What the present paper does not claim. Claims about Puppeteer’s holistic operational performance — endpoint latency distributions, developer-velocity comparisons, replay performance, longitudinal journal growth — are not made here. The present paper invokes the system only for the structural observations §4 lists and the existence claim of §5. Subsequent analytic papers may, like Paper 5, report bounded empirical observations scoped to their own structural claims; the comprehensive case-study is separate.

What is owed. The comprehensive case-study paper is owed; the analytic series will not continue indefinitely deferring it. A reader who finds the analytic series progressing while the case-study remains unwritten is observing the failure mode

the present authors name here and commit to avoid.

6. Counter-arguments

This section addresses the strongest objections to the argument advanced in §§1–5. Each objection is presented in its strongest form before its rebuttal; rebuttals draw on the principle articulated in §§1–3, not on the specific realization of §4.

“This is just CQRS done well.” CQRS already separates read and write; what does anti-porosity add? Two things, addressed in Claim 5 and §3. First, CQRS is one of four traditions that document a face of structure-centric encoding locally; anti-porosity is the unified diagnosis of structure-centric encoding as the common defect. CQRS implementations vary widely in this respect — some recreate porous DTOs and porous event payloads despite separating reads from writes. Second, the contribution of this paper is not any pattern (CQRS or otherwise) but the recognition that the four traditions converge on a single architectural choice — structure as the representational primary — which prior work has not named with sufficient generality to see across surfaces. **CQRS addresses separation of concerns; anti-porosity addresses the encoding decision (what the representation is shaped for). The two operate at different conceptual layers.**

“In small systems, porosity doesn’t hurt.” True. The operational costs of porosity scale with system size, longevity, and rate of change. The conceptual claim, however, does not depend on these. Porosity is a structural property visible at any scale; whether one chooses to address it depends on the regime in which the system operates. This paper does not advocate adopting any specific framework for small or short-lived systems; *When NOT to use this approach* explicitly enumerates regimes where the realization pattern is not the right fit, and the principle itself remains diagnostic in those regimes — porosity remains a real pathology, even where its correction would not repay its cost. **The claim of this paper is therefore diagnostic rather than prescriptive: it names a defect whose relevance depends on scale, not one whose correction is universally mandatory.**

“Joins are sometimes necessary.” Anti-porosity does not prohibit joins. As illustrated in §5, joins migrate to the read side, where enumeration serves analytical needs rather than polluting the representation of operations. The unified principle (§3) constrains only how state is *recorded* — not how it is queried for derivative purposes. Read projections may be denormalized, joined, indexed, and aggregated however a downstream consumer requires; the journal remains dense, but analytical and reporting use cases retain the full vocabulary of relational queries against derived views.

How read projections are constructed under anti-porosity. Two mechanisms are available, both operation-centric. First, *replay-on-demand*: a consumer requests a projection by replaying the journal against a projection-specific reducer. This

is correct, evolution-tolerant (the reducer can be updated and the projection rebuilt), and cheap for projections that aggregate small subsets of the journal. Second, *incremental projection via reactions*: a long-lived observer subscribes to the operation stream and updates a projection store as each verb commits. The observer is described as a *Reaction* in the companion Paper 3; it is itself operation-centric — the reaction processes the same operation that the journal records, not a relational diff of state — and therefore does not reintroduce porosity at the source-of-truth boundary. Projections may be relational (typed rows, indexed for query patterns) without polluting the journal, because the relational projection is *derived* from operations rather than *defining* the source of truth. The cost paid for relational reads is conventional (denormalization, indexing, eventual consistency lag), not structural; under anti-porosity the relational projection is one consumer of the source of truth, not its custodian. The full treatment of reactions, their consistency model, and projection rebuild strategies is the subject of Paper 3.

“Business intelligence and analytics require SQL.” Conventional BI assumes that the relational store *is* the source of truth. Under anti-porosity, the operation log is the source of truth, and analytical projections are downstream consumers. SQL access remains available — to the projection, not to the source of truth. **Because the operation log stores operations rather than state (§3, §4.3), projections for BI can be recomputed with evolving semantics — a capability unavailable when the relational store is itself the historical record.**

“Our data lake is the source of truth.” This is an organizational decision, not a structural property of the system. The data lake’s role can be reframed: it remains a shared projection consumed by downstream teams, while the operation log is the source of truth for the producer system. Anti-porosity is preserved internally even when external contracts demand denormalized projections at the system boundary. **The distinction is between organizational source of truth and architectural source of truth. Anti-porosity concerns the latter.**

7. Related work

This paper sits at the intersection of multiple bodies of work. Two prior streams attempted cross-tradition unification — *object-relational impedance mismatch* and the *data abstraction* tradition — and are reviewed first; four further traditions addressed faces of porosity locally and are reviewed after. Each is positioned for its relationship to the present paper, not for completeness; the literature review is not exhaustive (the methodological caveat is stated explicitly in §7.x below).

Prior unification attempts. The diagnosis that database-shaped representations and behavior-shaped objects do not align cleanly has a long history un-

der the name *object-relational impedance mismatch*. Ireland, Bowers, Newton, and Waugh (2009) provide the canonical peer-reviewed treatment: a systematic classification of the mismatch’s faces along structural, behavioral, and architectural axes, with explicit comparison of the mapping strategies (ORM, OODB, document-relational) that have been proposed against each face. That paper is the closest prior naming to the present paper’s *porosity*: both describe the cost of recording objects under a relational primary, both diagnose the cost (NULLs, joins, fragmented hierarchies) as structural rather than incidental, and both observe that the proposed remedies — work *within* the relational paradigm, do not question it.

The present paper differs from impedance-mismatch literature on three points. First, the impedance-mismatch frame is binary — object vs relational — and does not generalize to surfaces where neither side of the binary applies: an endpoint contract is not “object” or “relational”, an event-sourcing payload is neither either, yet both exhibit the same defect under the encoding-decision frame of §1.1. Second, impedance-mismatch literature names the *layer* at which the mismatch occurs (the object/relational boundary) without naming the encoding decision underneath it (structure as the representational primary); once that decision is named, the same defect becomes visible at endpoint and event-sourcing layers where no relational substrate is present. Third, impedance-mismatch remedies historically work *within* the relational paradigm (ORMs, document mappers, polyglot persistence, schema-on-read) rather than questioning structure as primary; the present paper proposes the inverse choice (operation as primary) projected across surfaces.

The *data abstraction* tradition (Parnas, 1972; Liskov & Zilles, 1974) and its successors in module-level encapsulation address a related but distinct concern: hiding implementation behind operations exposed by an interface, so that the representation can change without callers noticing. The principle “the representation should fit the operations that act on it” is older than the present paper’s diagnosis, and to that extent the present paper is descended from this tradition rather than novel against it. What this paper adds is the observation that the principle is honored at the module/class boundary but routinely violated at the persistence, contract, and event-payload boundaries — surfaces that the data-abstraction literature did not theorize as needing the same discipline. The contribution is the projection of the data-abstraction principle to surfaces where it had not been previously claimed, plus the unification with impedance mismatch under a single encoding-decision frame.

Algebraic data types and type theory. The sum-type vs product-type framing used throughout this paper (§2.1, §3) is borrowed directly from algebraic data type theory (Hindley, 1969; Milner, 1978), where the distinction has been formally established for decades. The contribution of the present paper is not this framing itself but its application to runtime representation: *which encoding the representational primary adopts at each layer* is the architectural decision the four traditions of §2 each documented in their own vocabulary. Type theory

provides the formal language; the paper applies it as an architectural diagnostic across surfaces that compile-time type systems do not reach.

Operation-based CRDTs. In distributed-systems literature, operation-based CRDTs (Shapiro, Preguiça, Baquero, & Zawirski, 2011) privilege operation over state in their replication semantics: the durable artifact replicated between nodes is the operation, not the resulting state. This is operation-centric encoding at the replication layer, established and well-defined. The present paper’s contribution is the cross-surface generalization: the same encoding decision can be made at the persistence primary, the domain primary, and the endpoint primary of a single-node system, not only at the replication boundary between nodes. The unification is across surfaces, not within the replication layer alone.

Append-only architectures. Append-only durable artifacts appear in Datomic, Kafka with log compaction, and the broader event-sourcing tradition. These systems share the structural property — additions only, no in-place mutation — but differ from the present realization in *what is appended*: append-only databases and logs typically append serialized state-deltas or command DTOs, not executable scripts of the operation that consumed specific inputs (§4.3). Append-only is therefore necessary but not sufficient for operation-centric encoding; the latter requires the additional decision that what is appended *is the script*, not the DTO.

Four applied traditions. *Domain-driven design* (Evans, 2003; Vernon, 2013) documents *anemic domain models* and *persistence-driven design* as pathologies and prescribes ubiquitous language and aggregate-rooted modeling as remedies. It diagnoses the symptom — a domain that lacks behavior — without naming the encoding decision behind it (structure-centric domain representation in service of structure-centric persistence). *REST API design* (Fielding, 2000; Richardson and Ruby, 2008) and the GraphQL specification (2015) document overloaded endpoints and fixed-shape DTOs, prescribing resource orientation and field selection respectively. GraphQL field selection moves toward operation-centric encoding at the wire — each invocation expresses what it consumes, not the union — but the discipline lacks corresponding operation-centric expression in the domain or persistence layers. *Event Sourcing* (Young, 2010; Fowler, 2005; Vernon, 2013) documents event-payload bloat and command-event coupling and prescribes command-event separation, upcasting, and versioning; it recognizes that persisting state is insufficient and persists operations instead, but typically represents those operations as fixed-shape command DTOs — the envelope of what was received rather than the script of what executed — a partial inversion that does not complete to operation-centric encoding. *Database theory* (Codd, 1970; Kent, 1983; Fagin, 1977) is the oldest and most formal of the four; it documents normalization anomalies and NULL semantics and prescribes successive normal forms. It names the effects (anomalies) without questioning the underlying decision: the relational model records state as structure, never operations as primary.

These four traditions have rarely engaged with each other on the

question of representational form, perhaps because their domains appear to address different problems. This paper argues that they are observing the same encoding decision through different lenses. Each tradition named a face of structure-centric encoding in its own layer and prescribed remedies that work within the structure-centric paradigm rather than questioning it. Among the cross-tradition unifications already reviewed, impedance-mismatch literature covers one surface (objects/relational) and the data-abstraction tradition addresses representation-fitting-operations at the module level; neither, to our knowledge, names the encoding decision (structure vs operation as the representational primary) at the cross-surface level proposed here.

Four formalisms. The grounding established in §1.1 invokes existing formal vocabulary, with two abstractions introduced here. Graph theory (Diestel, 2017), with extension to typed and labeled graphs in the Semantic Web (Lassila & Swick, 1999) and graph databases (Robinson, Webber, and Eifrem, 2015), provides the language of semantic graphs. Information theory (Shannon, 1948) inspires — without formally underwriting — the framing of representations in terms of *structural capacity* and *informational content*; both terms are this paper’s, not Shannon’s. Storage-engine literature (Hellerstein, Stonebraker, and Hamilton, 2007) catalogues the relational, document, and log substrates whose properties we abstract under the term *structural alphabet*; that abstraction is introduced here. Database normalization theory (Codd, 1970, 1972; Fagin, 1977) provides the canonical historical formalism for the relational case. The contribution is the observation that these four already-established formalisms — together with the two bridging terms — agree on a single defect that none names with sufficient generality across substrates. **The paper’s formal vocabulary is therefore mostly assembled and partly introduced — explicitly so where introduced.**

Aspect-oriented programming and ORM annotations. The aspect-oriented programming tradition (Kiczales et al., 1997; AspectJ; Spring AOP) identified persistence as one of the ugliest cross-cutting concerns and sought to externalize it from domain code. The most widespread practical instantiation is ORM-style annotations (the Java Persistence API, Hibernate, Entity Framework). These approaches share a goal with anti-porosity — keeping the domain class free of persistence concerns — but operate at a different level: they decorate the domain class with framework-specific metadata that describes how to map its fields to a relational substrate. The substrate’s structural alphabet is preserved; only the syntactic appearance of pollution is moved from the class body to its annotations. Anti-porosity, by contrast, addresses the substrate itself: by replacing relational projections with an executable journal of operations (§4.3), the question of how to map domain fields to relational columns disappears. **The domain class need not be annotated, because the persistence substrate no longer requires a projection of its fields.**

Other actor-model frameworks. Multiple frameworks implement combina-

tions of CQRS, the Actor Model, and Event Sourcing. Akka (Lightbend) on the JVM, Microsoft Orleans (Bernstein et al., 2014) with its grain abstraction, Proto.Actor across multiple languages, and the dedicated EventStore database all share substantial conceptual infrastructure with the realization presented in §4. They differ from it, and from one another, in a single consistent respect: **each persists events as serialized data structures defined at the command boundary, rather than as executable scripts derived after parameter evaluation.** The underlying storage organization typically retains relational patterns — category-indexed streams, table-backed projections — rather than per-actor append-only structures. None currently articulates anti-porosity as a unified principle, and none persists scripts in the script-form sense developed in §4.3. The differences are not failures of these frameworks against their stated goals — each does what it sets out to do — but they illustrate that anti-porosity requires a deliberate decision at the substrate layer that mainstream actor frameworks have not made. The decision is available to them: any of these frameworks could be extended to record operations as scripts rather than as DTOs, and to migrate the storage organization from category-indexed streams to per-actor append-only journals. **The distinction is not in the use of actors, CQRS, or event sourcing, but in what is chosen as the durable representational unit.**

Scope of the literature review. The review above is positioned rather than exhaustive. We did not conduct a systematic literature review in the PRISMA sense: we did not enumerate ACM DL, IEEE Xplore, or dblp queries with pre-declared keywords and inclusion criteria. The traditions reviewed were selected because they are the ones in which the present authors have prior fluency and whose vocabulary the paper engages with directly. Other prior naming attempts may exist that the present authors are unaware of; the falsifiable form of the paper’s novelty claim is therefore: *“we are aware of no prior work that names structure-vs-operation as the representational primary at the cross-surface level proposed here, and we engage explicitly with the two cross-tradition unifications we did find (impedance mismatch, data abstraction)”*. A reader who knows of a prior naming the paper missed is invited to write to the author; the contribution will be revised accordingly.

What this paper adds. The paper does not propose a new pattern. It identifies the encoding decision (structure vs operation as the representational primary) that the four applied traditions each described in surface vocabulary; positions the contribution as an *extension* of impedance-mismatch literature (covering surfaces beyond the object/relational boundary) and of the data-abstraction tradition (projecting its principle to persistence, contract, and event-payload surfaces); names the defect (*porosity*) and its principled inversion (*anti-porosity*); and demonstrates that the inversion is implementable as a unified discipline across surfaces. The Puppeteer framework serves as proof-of-existence. The analytic frame is *in principle* independent of any specific realization, but in practice the only realization currently known to the authors that jointly satisfies the three conditions is Puppeteer; the §7 review of “Other actor-model frameworks”

enumerates frameworks that could be extended to satisfy them but do not currently. Independence of the frame from any specific realization is therefore *conjectural*, not established; the open invitation in §5 is the explicit mechanism by which the conjecture would be tested. **The novelty lies in naming the encoding decision at a level of generality the prior unifications did not reach, and in showing that the inversion is operationally realizable in at least one production system.**

8. Conclusion

This paper has identified the encoding decision — *structure vs operation as the representational primary* — that four applied traditions each documented in surface vocabulary, and that prior cross-tradition unifications (impedance mismatch, data abstraction) named at narrower scope. We have named the structure-centric defect *porosity*, and the operation-centric inversion *anti-porosity*. Anti-porosity requires three simultaneous conditions: a domain whose hierarchy admits sum-type structure natively, an endpoint whose contract permits per-call selectivity, and a persistence layer that records the script that executed rather than the state that resulted or the command that was received. A system that satisfies all three through a single architectural decision — the executable journal — has been in continuous production use since 2018, after more than a decade of preceding refinement traceable to a 2005 prototype (§4.0); that system, Puppeteer, is presented in §4 as the existence proof for the principle, not as its definition. The analytic frame is *in principle* independent of any specific realization, but Puppeteer is currently the only realization we know of that jointly satisfies the three conditions; whether other systems could realize them through different mechanisms is open to test (§5 *Open invitation*). In the vocabulary established in §1.1, anti-porosity aligns structural capacity with informational content across the three representational layers a system exposes.

The implications extend beyond the paper. When porosity is removed at the substrate, much of the infrastructure traditionally introduced to compensate for substrate mismatch becomes optional rather than essential: caches that hid slow joins, ORMs that translated between domain and table, message queues that moved flat tuples between systems. The question of how systems built on porous substrates may adopt the new discipline progressively follows from the autopersistence property the framework inherits from its 2005 origin: a system whose initial state and event sequence are observable can be reconstructed independently of how it persists internally; a system whose only durable artifact is state cannot reconstruct the operations that produced it.

Several lines of work follow. A comprehensive case-study paper will report quantitative measurements from the framework’s production use; bounded empirical observations supporting specific structural claims have already appeared in Paper 5 (labs L1–L6) and are referenced in §5 of the present paper. Six companion papers extend the present argument across the series (Papers 2–7), enumerated

in the Cross-references section below. Independently, the forensic procedure introduced in §5 invites verification against public corpora — an exercise initiated in Appendix B against `dotnet-eShop` and open to readers for replication in their own domains.

The contribution of this paper is the identification of the encoding decision as the common cause behind four locally-named surface manifestations, the analytic frame that allows the decision to be diagnosed across surfaces (§§1–3), the demonstration that the inversion is realizable in production through one existence proof (§4) and that the diagnostic reproduces in a system the present authors did not build (Appendix B), and the explicit invitation for external falsification (§5 *Open invitation*). Whether the construct survives broader scrutiny depends on what readers find when they apply the diagnostic in their own domains and on whether independently-constructed systems satisfying or violating the three conditions of §3 come to light. Naming the encoding decision is a precondition for these tests, not a substitute for them.

Cross-references

This paper establishes vocabulary that subsequent papers in the series reuse:

- *Paper 2 — Dual compilation*: the anti-porosity principle as it manifests at the language layer (interpreted versus compiled execution paths).
- *Paper 3 — Reactions and the partition: opt-in eventual consistency in actor-native systems*: the consistency model that anti-porosity admits, with deferred derivative behaviors expressed as domain code rather than externalized pipelines; also the mechanism by which read projections are constructed under anti-porosity (§6).
- *Paper 4 — Preserving semantic continuity across actors: a tell-based approach without orchestration*: cross-actor coordination under operation-centric encoding; the journal’s density is prerequisite for the handoff mechanism.
- *Paper 5 — The Journal as Substrate: Unifying Deployment, Replication, Backup, and Offline Operation in Distributed Systems*: the substrate operations enabled by the executable journal; reports labs L1–L6 that empirically support its structural claims (also referenced from §5 of the present paper for the evaluation-deferral discussion).
- *Paper 6 — Most infrastructure layers are symptoms of the persistence model: a construct for auditing production stacks*: extends the anti-porosity argument outward to the broader infrastructure ecosystem, treating compensating layers (caches, ORMs, message queues) as artifacts whose necessity dissolves once the substrate is changed.
- *Paper 7 — After the substrate: building software without a datacenter*: extrapolates the argument forward, examining what is possible once journal-as-substrate is the operational primary rather than a substitute for relational persistence.

Code provenance

Source-code references in this paper resolve against the public Puppeteer repository at commit **b42d0f7** (2026-05-26), which ports the SimpleX/StageManager/Usher fixes from the active development branch on top of the prior public snapshot. The snapshot is archived in Software Heritage under the following persistent identifier:

```
swh:1:dir:177e2d61486bdbdfd2d5e774fcf392a45406e60d;  
  origin=https://github.com/alvaroNCubo/puppeteer;  
  anchor=swh:1:rev:b42d0f76b4278c36a45c221cc1453e3ee8ffb3de
```

A prior snapshot — commit **2f31f96** (2026-05-18) — is also archived in Software Heritage and remains resolvable for reference:

```
swh:1:dir:10e7e6bad7eb77b6c2e406762026177f95c3ae92;  
  origin=https://github.com/alvaroNCubo/puppeteer;  
  anchor=swh:1:rev:2f31f9674a5de816bdf1bf9d8360ff218a02e4da
```

The line numbers cited in the body have been verified against **b42d0f7** and remain stable across the two snapshots; the SimpleX integration did not move the cited mechanisms.

Inline references of the form `file.cs:NN` (e.g., `ActorHandler.cs:846`) resolve against the **b42d0f7** snapshot. A reader can construct per-file SWHIDs by adding the qualifiers `;path=<path>;lines=<NN>` to the **b42d0f7** directory SWHID above. Future commits to the repository may renumber lines; the SWHID preserves the cited state independently of any future change to the repository or its hosting.

Acknowledgments

The author used large language models (including Claude and ChatGPT) as editorial assistants for language refinement, structural feedback, and literature navigation. All original ideas, terminology, theoretical constructs, and technical content presented in this work are solely the author’s.

References

Bernstein, P., Bykov, S., Geller, A., Kliot, G., & Thelin, J. (2014). *Orleans: Distributed virtual actors for programmability and scalability* (Microsoft Research Technical Report MSR-TR-2014-41). Microsoft Research.

- Codd, E. F. (1970). A relational model of data for large shared data banks. *Communications of the ACM*, 13(6), 377–387.
- Codd, E. F. (1972). Further normalization of the data base relational model. In R. Rustin (Ed.), *Data base systems* (Courant Computer Science Symposia 6, pp. 33–64). Prentice-Hall.
- Diestel, R. (2017). *Graph theory* (5th ed.). Springer.
- Evans, E. (2003). *Domain-driven design: Tackling complexity in the heart of software*. Addison-Wesley.
- Fagin, R. (1977). Multivalued dependencies and a new normal form for relational databases. *ACM Transactions on Database Systems*, 2(3), 262–278.
- Fielding, R. T. (2000). *Architectural styles and the design of network-based software architectures* (Doctoral dissertation, University of California, Irvine).
- Fowler, M. (2002). *Patterns of enterprise application architecture*. Addison-Wesley.
- Fowler, M. (2005). *Event sourcing*. [martinfowler.com. https://martinfowler.com/eaDev/EventSourcing.html](https://martinfowler.com/eaDev/EventSourcing.html)
- GraphQL Foundation. (2015). *GraphQL specification*. <https://spec.graphql.org>
- Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.
- Hellerstein, J. M., Stonebraker, M., & Hamilton, J. (2007). Architecture of a database system. *Foundations and Trends in Databases*, 1(2), 141–259.
- Hevner, A. R., March, S. T., Park, J., & Ram, S. (2004). Design science in information systems research. *MIS Quarterly*, 28(1), 75–105.
- Hindley, R. (1969). The principal type-scheme of an object in combinatory logic. *Transactions of the American Mathematical Society*, 146, 29–60.
- Ireland, C., Bowers, D., Newton, M., & Waugh, K. (2009). A classification of object-relational impedance mismatch. In *2009 First International Conference on Advances in Databases, Knowledge, and Data Applications (DBKDA '09)* (pp. 36–43). IEEE.
- Kent, W. (1983). A simple guide to five normal forms in relational database theory. *Communications of the ACM*, 26(2), 120–125.
- Kiczales, G., Lamping, J., Mendhekar, A., Maeda, C., Lopes, C., Loingtier, J.-M., & Irwin, J. (1997). Aspect-oriented programming. In M. Akşit & S. Matsuoka (Eds.), *ECOOP'97 — Object-oriented programming* (pp. 220–242). Springer.
- Lassila, O., & Swick, R. R. (1999). *Resource description framework (RDF) model and syntax specification* (W3C Recommendation). World Wide Web Consortium. <https://www.w3.org/TR/1999/REC-rdf-syntax-19990222/>

- Liskov, B., & Zilles, S. (1974). Programming with abstract data types. *ACM SIGPLAN Notices*, 9(4), 50–59.
- McCarthy, J. (1960). Recursive functions of symbolic expressions and their computation by machine, Part I. *Communications of the ACM*, 3(4), 184–195.
- Milner, R. (1978). A theory of type polymorphism in programming. *Journal of Computer and System Sciences*, 17(3), 348–375.
- Mooers, C. N., & Deutsch, L. P. (1965). TRAC, a procedure-describing language for the reactive typewriter. In *Proceedings of the 20th National Conference of the ACM* (pp. 229–246).
- Parnas, D. L. (1972). On the criteria to be used in decomposing systems into modules. *Communications of the ACM*, 15(12), 1053–1058.
- Richardson, L., & Ruby, S. (2008). *RESTful web services*. O’Reilly Media.
- Robinson, I., Webber, J., & Eifrem, E. (2015). *Graph databases* (2nd ed.). O’Reilly Media.
- Shannon, C. E. (1948). A mathematical theory of communication. *Bell System Technical Journal*, 27(3), 379–423; 27(4), 623–656.
- Shapiro, M., Preguiça, N., Baquero, C., & Zawirski, M. (2011). Conflict-free replicated data types. In *Stabilization, Safety, and Security of Distributed Systems (SSS 2011)* (pp. 386–400). Springer.
- Vernon, V. (2013). *Implementing domain-driven design*. Addison-Wesley.
- Young, G. (2010). *CQRS documents*. https://cQRS.files.wordpress.com/2010/11/cQRS_documents.pdf

Appendix A: Source-code verification

This appendix consolidates the source-code references cited throughout the paper. All paths are relative to the public Puppeteer repository; line numbers reflect the source at commit `b42d0f7` (see *Code provenance* above). The framework’s source was migrated from Spanish to English identifiers during paper development; the line numbers and file names below reflect the English-language source as it currently stands.

A.1 Reflection-based discovery (§4.1)

Claim	File	Lines
Actor marker base class for domain entry points	Puppeteer/Actor.cs	15

Claim	File	Lines
Reflection-based discovery of domain types via assembly scan	<code>Puppeteer/EventSourcing/DomainLibraries.cs</code>	117–128
Domain library loaded at handler construction	<code>Puppeteer/EventSourcing/ActorHandler.cs</code>	61
Polymorphic argument-parameter compatibility at parse time (<code>AreCompatible</code>)	<code>Puppeteer/EventSourcing/Interpreter/Libraries/AstExpression.cs</code>	236–246
Runtime substitution of subtypes via symbol table (<code>IsSubclassOf</code> check)	<code>Puppeteer/EventSourcing/Interpreter/SymbolTable.cs</code>	134

A.2 Parameter modifiers (§4.2)

Claim	File	Lines
Four parameter modifiers defined (<code>In</code> , <code>Out</code> , <code>InOut</code> , <code>Eval</code>)	<code>Puppeteer/Parameter.cs</code>	33–36
<code>Eval</code> value setter generates literal-assigning script on first invocation	<code>Puppeteer/Parameter.cs</code>	163–224
<code>EvalScript</code> property (valid only for <code>Eval</code> parameters)	<code>Puppeteer/Parameter.cs</code>	228–247
Parser recognises <code>Eval</code> modifier in parameter declarations	<code>Puppeteer/Parameters.cs</code>	65, 110–111
<code>Eval</code> parameters' <code>EvalScript</code> serialized to journal	<code>Puppeteer/Parameters.cs</code>	334–367
<code>Out</code> parameters serialize as ? placeholder	<code>Puppeteer/Parameters.cs</code>	755–757
? for <code>Out</code> parameters deserializes to default value	<code>Puppeteer/Parameters.cs</code>	780–782

Claim	File	Lines
? tokenised as <code>TokenType.question</code> in the lexer	<code>Puppeteer/EventSourcing/Interpreter/Lexer.cs</code>	600
V1 <code>eval</code> statement preserved (interpreted only); V2 redirects to parameter modifier	<code>Puppeteer/EventSourcing/Interpreter/Libraries/EvalStatement.cs</code>	52–55

A.3 Executable journal (§4.3)

Claim	File	Lines
Journal high-level interface (append-style writes, read-forward only)	<code>Puppeteer/EventSourcing/DB/Diary.cs</code>	228–337
Append primitive (<code>AppendRecord</code>)	<code>Puppeteer/EventSourcing/DB/FileSystem/JournalWriter.cs</code>	65
Read-forward primitive (<code>ReadAll</code>)	<code>Puppeteer/EventSourcing/DB/FileSystem/JournalReader.cs</code>	35
Script entries persisted as UTF-8 bytes (<code>EncodeScriptEvent</code>)	<code>Puppeteer/EventSourcing/DB/FileSystem/BinaryEventCodec.cs</code>	41–101
<code>ReplayEvent</code> re-executes scripts at journal replay	<code>Puppeteer/EventSourcing/ActorHandler.cs</code>	846
On-demand parser invocation for type resolution at replay	<code>Puppeteer/EventSourcing/Follower/Reaction.cs</code>	1027–1029

Appendix B: Diagnostic application to an independently-constructed schema

B.1 Purpose

This appendix applies the diagnostic procedure of §5 (“Forensic observation”) to a schema constructed independently of the present authors. The purpose is to test, in the smallest available form, whether the construct *porosity* is identifiable in a system whose authors had no relationship to the construct’s naming. A

positive identification reduces — though does not eliminate — the risk named in §5 that the construct is coextensive with the present authors’ realization.

The system selected is the *Ordering* bounded context of Microsoft’s **dotnet-eShop** (<https://github.com/dotnet/eShop>), the official .NET reference implementation for cloud-native e-commerce. The Ordering context implements the canonical DDD + CQRS + event-sourcing patterns that §2 and §7 of this paper engage with. References below resolve against commit **9b4f943** (2026-current upstream **main**).

B.2 Method

The diagnostic of §5 asks three questions of any persisted representation:

1. **Horizontal sparsity (§2.1):** are columns of a table populated only by some rows, with the discriminator silently encoded in another column?
2. **Vertical sparsity (§2.3a):** are domain operations recorded as persistent artifacts, or only their terminal effects on state?
3. **Operation vs row (§5):** where the substrate forces enumeration of rows, would an operation-centric representation quantify (one verb over an aggregate) rather than enumerate?

We apply each question to the schema established by the initial EF Core migration plus its 2023–2024 refinements.

B.3 Observations

Schema under analysis. Migration **20230925222426_Initial.cs** (lines 122–162) creates the **ordering.orders** table with eleven columns:

Column	Nullability	Set by
Id	NOT NULL	sequence orderseq
Address_Street, _City, _State, _Country, _Zip	NULL	Address value object (owned entity flattened to 5 columns)
OrderStatusId	NOT NULL	enum discriminator
Description	NULL	varies by state transition (see below)
BuyerId	NULL	set by SetPaymentMethodVerified() , not at construction
OrderDate	NOT NULL	constructor
PaymentMethodId	NULL	set by SetPaymentMethodVerified()

Horizontal sparsity test (§2.1). The `OrderStatus` lifecycle (`Ordering.Domain/AggregatesModel/OrderAggregate/OrderStatus.cs:7–14`) defines six states: `Submitted`, `AwaitingValidation`, `StockConfirmed`, `Paid`, `Shipped`, `Cancelled`. The schema collapses all six into a single discriminator column (`OrderStatusId`). The `Description` field is the clearest horizontal-sparsity case: its meaning is conditioned on `OrderStatusId`, and its values are assigned literally as the lifecycle advances (`Order.cs:115,126,138,151`):

OrderStatusId	Description
Submitted	NULL
AwaitingValidation	NULL
StockConfirmed	"All the items were confirmed with available stock."
Paid	"The payment was performed at a simulated \"American Bank checking bank account ending on XX35071\""
Shipped	"The order was shipped."
Cancelled	variable per cancel path

The `Description` field is the operation's residue, persisted into a structural column whose meaning the schema does not encode and whose population is conditional on a sibling column. This is the §2.1 defect literally — a sum-type discriminated field encoded as a product-type column. `BuyerId` and `PaymentMethodId` are nullable for the same reason: they are determined later in the lifecycle by `SetPaymentMethodVerified()`; the schema accepts orders that have not yet been verified by populating `NULL`.

An anecdotal aside, not part of the diagnostic. The domain class `Order` (line 23) carries a private boolean field `_isDraft` accompanied by `#pragma warning disable CS0414 // The field 'Order._isDraft' is assigned but its value is never used`. This is in-memory dead code — plausibly a minor maintenance artifact rather than evidence of structural porosity, and we explicitly do not count it as diagnostic. We mention it because the analogous pattern at the persistence layer would be diagnostic — columns retained across migrations whose values no operation reads — but the present analysis does not enumerate such columns in the schema and the inference from one to the other is suggestive, not probative. Readers should weight the in-memory observation accordingly.

Vertical sparsity test (§2.3a). The Entity Framework configuration (`Ordering.Infrastructure/EntityConfigurations/OrderEntityTypeConfiguration.cs:9`) declares `orderConfiguration.Ignore(b => b.DomainEvents)`. The aggregate's domain operations — `OrderStartedDomainEvent`,

`OrderStatusChangedToAwaitingValidationDomainEvent`,
`OrderStatusChangedToStockConfirmedDomainEvent`,
`OrderStatusChangedToPaidDomainEvent`, `OrderShippedDomainEvent`,
`OrderCancelledDomainEvent` — are dispatched in memory and excluded from persistence by construction. The schema captures the terminal state (`OrderStatusId = Shipped`), not the operations that produced it. The transition trace exists only ephemerally, during the request that produces it, and is then dropped. This is the §2.3a defect literally: state-as-form persisted, operation-as-trace discarded.

Outbox as §2.3b instance. Migration `20231026091055_Outbox.cs:14–30` adds an `IntegrationEventLog` table:

Column	Type
<code>EventId</code>	uuid
<code>EventTypeName</code>	text
<code>State</code>	int
<code>TimesSent</code>	int
<code>CreationTime</code>	timestamp
<code>Content</code>	text (JSON-serialized payload)
<code>TransactionId</code>	uuid

The `Content` column holds the integration event serialized verbatim — the command DTO as received, not the script of what executed. This is the §2.3b defect at the outbox boundary: the event payload is shaped by the type-space of any integration event the system may emit, not by the specific operation that produced this event. Additionally, the `State` column itself reproduces the §2.1 defect — the outbox event’s own lifecycle (`Pending`, `InProgress`, `Published`, etc.) collapsed into a discriminator column.

Operation vs row test. The `SetStockConfirmedStatus()` operation (`Order.cs:108–117`) is, semantically, one verb applied to the aggregate. Under the relational projection, no row in the `orders` table records that this operation occurred; only `OrderStatusId = StockConfirmed` and `Description = "All the items..."` are observable, both attributes of the terminal state. To answer the question “*at what time did this order transition to StockConfirmed?*”, the schema has no answer — the operation’s timestamp is not in the schema (only `OrderDate`, the constructor time, is recorded). To recover the transition history one must join against the dispatched (but unpersisted) domain events, which is impossible after the request that produced them returns. The operation existed; the row does not record it. An operation-centric representation would quantify (one verb in the journal: `SetStockConfirmed(orderId, at: timestamp)`); the relational form has no quantifier and no place to put the verb.

B.4 What this demonstrates

The diagnostic reproduces. The Ordering bounded context of **dotnet-eShop** — constructed by Microsoft engineers without reference to the construct *porosity* — exhibits all three of the manifestations enumerated in §2: horizontal sparsity in **orders.Description** (§2.1), vertical sparsity by explicit **Ignore(b => b.DomainEvents)** (§2.3a), and contract-induced sparsity at the **IntegrationEventLog.Content** blob (§2.3b). The construct is therefore not strictly coextensive with the present realization: its diagnostic identifies the defect in a system the construct’s authors did not build.

B.5 What this does not establish

Three caveats are worth naming.

First, this is a *forensic* demonstration, not a *normative* one. The dotnet-eShop architects produced a competent reference implementation under the prevailing structure-centric default. The team’s choices (single-table inheritance for **Order**’s lifecycle, owned entities for **Address**, outbox for reliable event delivery, explicit **Ignore** for in-memory domain events) are within the standard playbook for **DDD + CQRS** on a relational substrate. The diagnostic identifies that the system *exhibits* porosity, not that it is *wrong* by some externally-imposed standard. eShop serves its purpose; the porosity is the cost of the encoding decision the system makes, not a defect of execution within that decision.

Second, this demonstration analyzes a *porous* system, not an *anti-porous* one constructed by an independent party. The reduction in single-realization risk is therefore partial: it shows the diagnostic reproduces externally, not that the inverse encoding has been independently constructed and successfully operated.

Third, the analysis is small in scope (one bounded context of one system, one migration sequence). A systematic survey of schemas — across multiple companies and frameworks — would strengthen the claim that porosity is a default rather than a local artifact. We do not undertake such a survey here; the present demonstration is sized to the falsifiability claim of §5 (that the construct is identifiable in independently-constructed systems), not to a comprehensive empirical mapping.

Program–value separability

TL;DR

Compilation, caching, and dense journaling are **not** features layered atop interpretation. They **are** downstream consequences of a single structural commitment: values live outside the script, not embedded within it.

The transition often described as “adding compilation” inverts the causal order. A DSL whose programs embed their values has no stable identifier, no template to cache, no separable sub-program to compile ahead of time — interpretation is not a stylistic choice but a structural necessity. Once values move outside the script, the program becomes a function awaiting its arguments — $F(x_1, x_2, \dots, x_n)$, as the runtime documents itself. And it is now the *same* program on every call, whatever arguments it receives: it has acquired an identity of its own, independent of its values. That identity is the pivot. The runtime decisions below are not consequences of separability directly so much as consequences of a program’s having a stable identity — and separability is what confers it.

From that single act of separation, four runtime decisions cohere. Scripts with externalized parameters are eligible for compilation to IL via Expression trees, cache as reusable programs with action identifiers, persist to the journal as compact action entries, and replicate with entropy bounded by argument vectors. Scripts with hardcoded values are not cached, run once, and persist as their literal text. The four are not parallel design choices; they are projections of a single precondition. This paper develops and measures the first three (compilation, caching, dense journaling); the fourth, replication-bounded entropy, follows from the same precondition and is developed in a companion paper — it is named here, not claimed as a result of this one.

This precondition operates on the substrate characterized in the previous paper of this series — Puppeteer’s journal as *code-as-data*:

programs persisted in their own executable text. The prior paper develops this as *script-form persistence* (the *executable journal*) and is careful to claim it in a sense narrower than Lisp’s strong homoiconicity, which it explicitly disclaims; we adopt the label *code-as-data* from this point on, with that formalism standing as established there. The compactness of these entries is not merely syntactic: each names an operation substantially richer than its surface signature, with the asymmetry scaling with the domain library’s depth (§5.5).

This is the principle of program–value separability: a DSL program becomes identifiable, compilable, cacheable, and persistable densely only when it is separable from its values. Externalized parameters are the structural precondition under which compilation, caching, and dense journaling become possible at all.

Claims this paper makes

1. **Causal inversion: program–value separation is a necessary condition for the stable-identity faces.** Compilation, caching, and dense journaling in the runtime described here are not optimizations layered atop interpretation. They are downstream consequences of *program–value separability* — the structural property that values live outside the script body, not embedded within it. A DSL program embedding its values has no stable identifier under which to reuse a specialization across invocations, no template to reference in the journal rather than copy, no bounded argument stream to replicate. The separation is necessary for these faces, not optional — though it need not be *declared* rather than recovered by analysis (§1.2), and it is not claimed that caching of every kind requires it (memoization of results does not). (*Verification: structural, within this class of runtimes — a script embedding its values produces a different cache key per invocation, so cross-invocation reuse, reference-based journaling, and bounded replication collapse together; the argument and its two standard objections are made precise in §1.2. The runtime documents the model itself: “Un script funciona como $F(x_1, x_2, \dots, x_n)$ ” (ActorHandler.cs:1085) — F is the program; the x_i are its externalized values.*)
2. **DSL programs admit a structural binary taxonomy.** *Parametric scripts* are separable from their values: cacheable as reusable programs with action identifiers, persistable as compact action entries in the journal, and economically compilable. *Literal scripts* are non-separable: never cached, ephemeral, and persisted as their literal text. Whether either class is compiled or interpreted at runtime is governed by a separate per-actor policy that, by default, follows the parametric/literal split. (*Verification:*

ActorHandler.cs:1091-1093 documents the rule verbatim. Implementation: PrepareCommandProgram (ActorHandler.cs:1097-1148) decides via parameters.HasUserParameter() and assigns one of three JournalEntry values — IsScript (literal), IsNewAction (parametric, first invocation), IsExistingAction (parametric, cached invocation). Compilation flag set per program by AdjustCompilationMode (Program.cs:134-150) under the CompilationModePolicy enum.)

3. **Interpretation and compilation are strategies over a shared AST substrate.** They are not separate execution paths over translated representations. Both walk the same AST: interpretation walks it directly; compilation specializes it into IL via Expression trees. Both paths pay the same parsing cost; only the compiled path pays the compilation cost. (*Verification: Program.Execute() (interpretation) and Program.ExecuteExpression() $\rightarrow$ ProgramExpression().Compile() (compilation) operate over the same Statement tree produced by Parser.Parse().*)
4. **Journal density emerges from separability, not from a compression scheme.** A stable $F(x_1, \dots, x_n)$ collapses repeated invocations into (actionId, values) tuples, with the script definition persisted once and referenced thereafter. The compactness is structural, not designed. (*Verification: Diary.WriteNewActionEntry (definition + first invocation values) vs Diary.WriteActionEntry (actionId + values only on subsequent invocations).*)
5. **Verb richness: surface vs depth.** A single DSL verb may invoke orchestrations richer than its surface signature suggests. This is not the result of translation between layers; the DSL is the invocation language for domain operations implemented directly in the host language. Journal density is therefore the asymmetry between small persisted tokens and the live-domain operations they represent. (*Verification: §4 develops the construct and §5.5 measures it on two open-source DDD aggregates — eShop’s Order purchase verb (.NET Foundation 2026) dispatches $\alpha = 9$ host-language invocations with a 24-method static closure ($\beta/\alpha \approx 2.7\times$); Grzybek’s SubscriptionPayment verb (Grzybek 2024) dispatches $\alpha = 2$ with a 73-method static closure ($\beta/\alpha \approx 36\times$). α is exact; β is a crude static lower bound on reachable host surface (§5.5), and at $n=2$ the $2.7\times$ -vs- $36\times$ gap is not attributable to any single cause. The supported claim is the qualitative one — the persisted token is far smaller than the behavior it names — carried by α and by journal density (§5.4), not by the β/α magnitude.)*
6. **Hot-loaded DSL programs over a stably-loaded domain.** The runtime supports hot-loaded DSL programs against a stably-loaded domain library: the assemblies configured as the actor’s domain libraries are reflectively cached at first use; new DSL scripts can combine the types they carry in new ways without assembly reload. This enables ad-hoc procedure invocations against a live, stateful actor — without redeploy-

ing endpoints or altering the domain library. (*Verification: DomainLibraries.GetOrLoad(params Assembly[]) caches public types statically per deduplicated assembly set; ActorV2.Using(scriptForChk, scriptForCmd) introduces new scripts per invocation against that cached surface area.*)

When NOT to use this approach

The principle of program–value separability has limits, and the implementation pattern that realizes it has narrower limits still. We name six regimes in which the analysis presented here either does not apply, applies only partially, or invites overreach.

A. This paper is not advocating for compiling all DSL programs unconditionally

The structural taxonomy described here distinguishes parametric from literal scripts; both are first-class. Literal scripts — ad-hoc oracle invocations against a live actor, exploratory queries against a stateful runtime, configuration-time experiments — pay zero amortization cost because they incur no compilation. Treating them as legacy or transient would discard a structurally valid mode of interaction.

B. This paper does not address consistency under concurrent actors

The runtime described operates with strict per-actor ordering and isolation. Multi-actor consistency, reaction propagation across actor boundaries, and the contract under which observable state remains coherent are treated formally in a companion paper of this series.

C. This paper does not claim program–value separability is sufficient for the four-fold coherence

Separability is necessary but not sufficient. SQL prepared statements achieve compilation and caching while exhibiting only two of the four faces — the persisted artifact is row data, not the parameterized statement itself, so the journal is not code-as-data (script-form persistence does not emerge). Other realizations may exhibit further subsets. Puppeteer demonstrates one realization where all four faces emerge by structural commitment; the converse is not entailed by separability alone.

D. This paper does not extend separability to general-purpose programming languages

The principle is formulated for DSL runtimes — domain-specific languages whose programs are short, structured, and amenable to identification by their textual form. General-purpose programming languages introduce variables captured from enclosing scope, side-effecting state in the host, and identity that exceeds the parameter contract. Whether separability admits a meaningful formulation in that broader setting is a separate question, and we do not take it up here.

E. This paper does not claim that hot-loaded DSL programs replace traditional deployment

Claim 6 articulates a runtime capability, not a deployment recommendation. Hot-loaded DSL programs against a stably-loaded domain enable ad-hoc invocation; redeployment of the domain library remains the appropriate move when the implementation itself changes. The two are complementary, not alternatives.

F. This paper does not address security, sandboxing, or resource bounds for hot-loaded scripts

Claim 6 grants a running actor the ability to accept and execute DSL programs formed at call time. The structural argument made here is silent on the operational concerns that capability raises in production: sandboxing of the DSL evaluator, per-script resource limits (CPU, memory, wall-clock), authentication and authorization over which callers may submit ad-hoc invocations, and the trust model under which the supplied script is admitted at all. These concerns require separate machinery — out of scope for the conceptual contribution offered here, but unavoidable for any deployment that exposes hot-loaded invocation beyond a closed boundary.

1. Introduction

This paper makes an analytic theory contribution. It identifies a structural property of DSL programs that prior literature has touched but not isolated as one — the construct: *program-value separability*; derives the runtime consequences that follow when the property is held — the principles: compilation, caching, and dense journaling, developed and measured here, together with a fourth, replication-bounded entropy, that follows from the same precondition and is developed in a companion paper; and presents an instantiation — a system in which those principles have been realized in production — as confirmation that the construct is realizable. The contribution is conceptual; the instantiation is

the existence proof of realizability, not the substance of the claim. The genre is the one Gregor (2006) names *theory for analyzing* (Type I): it introduces constructs and a taxonomy that let the phenomenon be described and classified, and states a necessary-condition claim about their consequences, with empirical evaluation supplementary. The Hevner-style design-science *evaluation* of the artifact (Hevner, March, Park, & Ram, 2004) — bundling artifact assessment with construct introduction — is deferred to a companion case-study paper; the Type I frame separates the two cleanly, and the instantiation here serves only as an existence proof of realizability.

The structural property at the center of this paper is not, in itself, new, and §7 documents its lineage. Partial evaluation, lambda lifting, closure conversion, and template instantiation all turn on a separation between a program and the values it will receive; each predates this work by decades. What is new is the dependent variable. In those traditions, separability is instrumental to *execution*: a binding-time or free-variable analysis recovers the static/dynamic split, a transformation consumes it to emit a faster residual program, and the split is discarded once the specialized code exists. This paper repositions the same property as the precondition for *persistence and identity*. The separated program is not consumed by a transformation but persisted as the durable, replayable, replicable unit of state — written once as a definition, referenced thereafter by a stable identifier and an argument vector, and reconstructed by deterministic replay. Compilation becomes one face among four, and not the load-bearing one. The claim this paper defends is that compilation, caching, dense journaling, and replication-bounded entropy are not four independent optimizations but four projections of this single precondition once the persisted artifact is the program rather than its effects — a connection the execution-specialization literature has no occasion to draw, because it never treats the program as state.

1.1 Calibrating the surface

You already understand this taxonomy. A SQL prepared statement is parameter-separable: the database parses it once, builds a query plan once, and reuses both for thousands of invocations with different bind values. An ad-hoc query offers no such surface — each is a fresh string, parsed and planned afresh. The same structural distinction governs DSL programs in any runtime that aspires to compilation, caching, and dense persistence. What changes between SQL and the runtime described here is not the principle, but the surface to which it applies. The principle, in other words, is neither new nor Puppeteer’s: it is one every database programmer already relies on locally, each time a prepared statement is reused. What is new is the recognition that the same condition, taken as general rather than local, is what underwrites compilation, caching, dense journaling, and replication alike — and that a runtime can be built to honor it everywhere, not only at the query boundary.

That surface tends to be unfamiliar in a specific way. The professional reader

of this paper has very likely spent a career persisting application state in relational tables. Where event sourcing has appeared in their work, it has typically appeared in tabular form: events as rows, schemas as projections, queries as joins against materialized views. The implicit assumption — that every representation of state ultimately resolves into something table-shaped — is not a failure of imagination but the horizon that industrial practice has produced over several decades. A runtime in which the persisted artifact is the program itself, in which density emerges from separability rather than from compression of rows, has had no canonical place in that landscape.

Calibration begins by separating two layers that the dominant paradigm tends to collapse into one. The first layer is the implementation: domain types, methods, business rules — written in the host language, indistinguishable in form from any well-designed object-oriented or functional library. The second layer is the invocation: the language by which a running actor is commanded and queried. In runtimes of the kind described here — of which Puppeteer, the framework introduced in the prior paper of this series, is one — that invocation language is a small DSL whose programs read as scripts, persist as journal entries, and execute as either compiled lambdas or interpreted walks over their own AST. The two layers are not translations of one another. The host language carries the implementation; the DSL carries the contract by which the implementation is reached. Conflating them — assuming the DSL must duplicate, transcribe, or shadow the host code — is where most readings of the architecture begin to diverge from its actual structure.

Three readings should be pre-empted at the outset: this is not double programming, not a translation layer, not a mapping specification. The domain implementation lives once, in the host language. The DSL is the language by which that implementation is invoked. The relationship is closer to that between SQL and a relational engine than to that between a model and a generated DTO: SQL does not duplicate the engine’s logic; it names operations that the engine performs. A short SQL statement can trigger joins, locks, and execution plans far heavier than its surface text. The DSL of an actor runtime carries the same kind of leverage — small surface, deep behavior — and the persisted journal records not the unfolding of that behavior but the invocations that produced it. With the surface calibrated, the formal question can now be asked: what does it take, structurally, for a DSL program to be compilable, cacheable, and amenable to dense persistence?

1.2 A formal characterization

The question can be sharpened. Compilation requires a stable artifact to specialize. Caching requires a stable identifier under which to file the result of that specialization. Dense journaling — persistence that does not balloon with the size of every invocation — requires a stable description of the operation that can be referenced rather than copied. Each of these three demands the same property of the program: that it admit, at the moment it is named, a separation

between the program itself and the values it will receive.

This property is **Program–Value Separability**. A DSL program is separable when its textual form names an operation parameterized by values supplied externally at invocation, rather than carrying those values within its body. Separability is structural, not stylistic: it is a property the program either has or does not have. It is decidable syntactically at preparation time, by inspection of the parameter declarations on the program’s surface. The principle that follows admits a compact statement: program–value separation is a necessary condition for the *stable-identity* faces — caching a specialization under a key reused across invocations, journaling the program by reference rather than by copy, and replicating it as a bounded argument stream. Two scopings keep this from collapsing into a definition. The necessity is of the separation itself, not of any particular way of obtaining it: a runtime may have the programmer declare parameters, as the one described here does, or recover the separation by analysis. And the necessity is of *these* faces, not of caching in general — caching of other objects, such as memoized results, requires no separation at all. Necessity is the strong claim; its scope and the two natural objections to it are made precise shortly.

What separability buys, in a word, is *identity*. A program whose values are supplied from outside has a textual form that is stable across invocations — it is the same program the second time and the thousandth — and so it can be named; a program that embeds its values is a different string on every call, and no two invocations are the same program at all. Identity is the pivot on which the rest turns. Once a program has an identity independent of its values, caching has something to file under, compilation has a stable artifact to specialize, the journal has a reference to record in place of a copy, and replication has a unit to ship. The faces traced in this paper are best read not as four properties of separability but as consequences of the *program identity* that separability confers — and it is the absence of that identity, not the presence of embedded values as such, that makes a value-embedding script uncacheable, uncompileable-for-reuse, and undense in the journal.

Separability admits a binary structural taxonomy of DSL programs. A **parametric script** is separable from its values: it is filed under a stable identifier in a runtime cache and persisted to the journal as a compact reference plus an argument vector. A **literal script** is non-separable: it is not cached, runs once as written, and persists as its literal text. Whether either class is compiled or interpreted at runtime is governed by a separate per-actor policy that, by default, follows the parametric/literal split — but admits override. Under that default policy, the runtime documents the joint regime in its own comments: scripts without user parameters are interpreted, uncached, and persisted as Script entries; scripts with user parameters are compiled, cached with an ActionId, and persisted as Action entries. The two strict properties — caching and journal entry format — are encoded as values of a single enum decided at the moment a program enters preparation. The taxonomy is not analytical but operational: the runtime acts on it.

Necessity is not sufficiency. A program that admits separability gains the structural possibility of being cached, persisted as a compact reference, and amortizing its compilation — but does not by that fact alone exhibit dense journaling. SQL prepared statements illustrate the gap: they are separable, they cache, they compile — yet the persisted artifact in the database is row data, not the parameterized statement itself. The third face — a journal in which the named operation is the persisted entry, not its downstream effects — does not emerge from separability alone. It requires the further commitment, code-as-data, that the runtime treat its programs as the persisted artifact: what is recorded in the journal is the program parameterized and the values it received, not the state changes those values produced. The sections that follow trace how separability, joined to code-as-data, yields the three faces in turn — and where a runtime that holds separability without the second commitment falls short.

Necessity, made precise. Two objections test the claim, and locating where each lands sharpens it rather than defeats it. The first is *memoization*: a runtime can cache the result of an invocation keyed on its inputs with no parameters declared anywhere. This is genuine caching, but of a different object — it caches outputs, not the program. A memoized value-embedding script still hits only on identical re-invocation, never across distinct argument vectors, and yields neither a reusable specialization, nor a reference-based journal entry, nor a bounded replication stream. Memoization caches effects; the faces at issue cache the operation. The second objection is stronger: a runtime can *recover* a separation the programmer never declared — canonicalizing literal scripts and abstracting their literals into a template after the fact, the inverse of constant folding. This does not refute the claim; it relocates it. What such a runtime recovers *is* program-value separation, obtained by analysis instead of by declaration. The separation remains the precondition; canonicalization is merely another route to it, one that pays an analysis cost and yields a separation that is provisional and runtime-derived rather than stable and surface-given. So the claim is not that a programmer must declare parameters — it is that the four-fold coherence requires the separation to exist, however reached, and that a *stable, syntactically decidable* separation is what makes the separated program serviceable as the durable identity under which journaling and replication operate, where an analysis-recovered one is not. Throughout, the definition of separability is syntactic — a property of program text — while the claim concerns that property’s consequences; the two are kept apart precisely so the claim is not true by definition.

The DSL program characterized above — the textual artifact with declared parameters — is the unit of separability. When the term *program* appears at the substrate level in subsequent papers of this series (the unit that is recorded, replicated between nodes, and replayed against state), it denotes the pair: a **domain library** (the compiled classes and verbs against which the DSL is interpreted, loaded statically and identical across replicas) and a **journal of invocations** (the ordered sequence of action references with parameter bindings that the runtime has applied). Replay is the deterministic re-execution of the journal

against the library; what persists, replicates, and reconstructs state is this pair. A DSL program in the sense of §1.2 is the unit; the substrate-level program is the cumulative state-producing artifact built from many DSL invocations executed against a stable library. The two readings are not in tension: separability of the unit is what makes the cumulative artifact dense; the cumulative artifact is what makes the unit’s separability operationally consequential.

The central claim can now be stated without hedging. It is not that separability improves compilation — compilation is merely its most visible consequence, and the paper would stand were it removed, on caching, dense journaling, and replication-bounded entropy alone. The claim is that compilation, caching, dense journaling, and replication-bounded entropy are four operational projections of one structural condition: a program’s having an identity independent of its values. Nor is the condition peculiar to this runtime or this DSL. No system can give a program a *reusable* identity — and therefore none of the consequences that depend on one, a cached specialization, a reference recorded in place of a copy, a replayable unit — while its values remain embedded in its representation. Separability is the general precondition; the runtime of §3 is one case study in what follows once it is met.

2. Consequences of separability

2.1 Compilation as specialization

The first consequence of separability is that the program admits ahead-of-time specialization. A separable program is, in form, a function awaiting its arguments — its body refers to values that will arrive at invocation, not to literals fixed at write time. That structural shape is precisely what a specializer requires. Given the program once, the runtime can resolve its references, lower its abstract syntax tree into a typed expression tree, and emit the executable form in advance of any specific call. The values are bound later; the work of preparing the program is done once.

The pipeline that produces this specialized form proceeds in three stages. The parser yields an abstract syntax tree in which user parameters appear not as values but as named slots — placeholders to be supplied at invocation. The runtime then traverses that tree and emits a typed expression: each named slot is bound to a parameter expression of the host runtime, and each operation of the script is bound to the corresponding host-language method or property access. Compilation of that expression yields an executable delegate over the parameter list — a function that takes the values supplied at invocation and runs the program against them.

Specialization, considered alone, is a one-time cost paid against an indefinite number of invocations. The parser runs once for the program; the expression tree is constructed once; the compiler emits the delegate once. Each subsequent invocation supplies new values into the existing delegate and runs through host-

language code that has already been resolved, typed, and emitted. The amortization is structural: a separable program admits compilation that pays for itself across repeated invocation, while a non-separable program offers nothing to amortize against. The work, however, is only economical if the specialized form persists between invocations.

2.2 Caching as identification

Persistence between invocations requires an identifier under which to file the work that has been done. Separability supplies that identifier directly: the script string of a separable program is stable across invocations because its values are external to it. The same parameterized program text, encountered a second time, is recognizable as the same program — and the cached specialization belongs to it. A non-separable program has no such anchor; every invocation produces a different string, and no two invocations are the same program at all. Caching is not an addition; it is what becomes possible once the program has a stable name.

The mechanism is direct. On first encounter, the runtime parses the script, prepares its specialized form, and stores both under an action identifier — a stable handle assigned to the script when it is first encountered. The script string serves as the lookup key; the action identifier serves as the persistent reference under which the program’s specialization lives. On subsequent encounters with the same script, the lookup succeeds, and the runtime retrieves the specialization without repeating the work that produced it. The cache is not a memoization of values; it is a registry of programs.

On a cache hit, the work that has been preserved is the program itself: the resolved tree, the typed expression, the compiled delegate. The values arriving at this invocation are bound into the delegate’s parameter slots and the program runs against them. From the program’s standpoint, nothing has changed between the first invocation and the thousandth — the script names the same operation, parameterized by the same slots, executed against the same domain. What differs across invocations is only the values supplied; what persists is the program.

2.3 Dense journaling as reference

The third consequence of separability is that persistence becomes reference rather than duplication. The program, having earned a stable identifier in the cache, can be written to the journal once — as a definition, with the script’s text and its parameter declarations recorded in full. Each subsequent invocation is recorded not by duplicating the program text but by referring back to that definition under its identifier, accompanied only by the values supplied at that invocation. Where a non-separable runtime must record every invocation as a complete script — body, values, and all — a separable runtime records the program once and its invocations as compact references against it.

The mechanism distinguishes three kinds of journal entries. When a parametric program is encountered for the first time, the runtime writes a definition entry: the script’s text, its parameter declarations, and the values supplied at this first invocation, all recorded together under the program’s newly assigned identifier. Subsequent invocations of the same program write only a reference entry — the identifier of the definition, plus the values for that call. A non-separable program, having no stable identifier, is written each time as a script entry: its text in full, since it cannot be referred back to anything that came before. Three entry kinds, decided mechanically at preparation time from the program’s separability, encode the program’s relation to its identity. The distinction is, at bottom, ontological: the entry type records whether the thing persisted *is* a program with an identity of its own — an Action, referable and replayable — or merely a script that ran once and has none. The journal is, in this sense, a ledger of which operations possess operational identity and which do not; separability is the test it applies, and the entry type is the verdict it records.

The density follows. For a parametric program invoked many times, the journal’s storage scales with the cumulative size of the argument vectors, not with the cumulative size of the script’s body — the body has been recorded once, and every later invocation costs only the bytes of its values. This is what makes the journal dense: nothing is copied that has been preserved by reference, and what survives is the operation that took place, not the state that resulted. The persisted artifact is the program — the script as text, parameterized, with its invocation values — rather than the downstream effects of running it. In code-as-data terms, the program lives in the journal as itself, not as a record of what it produced. Compilation, caching, and dense journaling are thus not independent optimizations but successive consequences of the same structural property: once a program is separable from its values, it becomes nameable; once nameable, it becomes cacheable; once cacheable, it becomes referable in persistence. The mechanism by which the runtime accomplishes this — the pipeline that compiles, the cache that retains, the journal that references — is the subject of §3.

3. Realization in a concrete runtime

3.0 Origins of the instantiation

The realization described in this section was not engineered to demonstrate the principle articulated above. The principle was identified by inspecting a runtime that had, over years of independent development for production use, come to satisfy it. The genealogy is structural rather than programmatic — first the artifact, then the principle that explains why the artifact cohered. What follows describes the runtime as it stands; the relationship between its mechanisms and the principle of §1.2 is a recognition, not a derivation.

The artifact is open to inspection, and the claims that follow are meant

to be checked rather than taken on trust. The runtime is public at github.com/alvaronCubo/puppeteer, archived in Software Heritage at the commit cited under *Code provenance*; every `file.cs:NN` reference in this section resolves against that public source. The benchmark harness of §5 is likewise public (this paper repository’s `labs/` and the runtime’s `tests-local/`), and the host aggregates it exercises are public MIT projects. The existence proof offered here is therefore verifiable in the sense that matters for the claim: a reader can read the cited code, build it, and re-run the measurements. The earlier, internal lineage of the runtime is not part of the proof — only the public snapshot is — and where the public snapshot differs from the internal development line, the public snapshot governs every reference and figure in this paper.

3.1 Compilation pipeline

The mechanisms described here are not independent engineering decisions but direct realizations of the separability principle established in the preceding sections. In the runtime described, the compilation pipeline runs from text to executable in three well-defined stages. A `Parser.Parse()` call lowers the script into an abstract syntax tree of `Statement` and `AstExpression` nodes. Each `Statement` carries an `ExecuteExpression` method that, when traversed, emits a `System.Linq.Expressions` node bound to the host parameters and the host-language operations the script names. The runtime composes these into a single `Expression.Lambda` and calls `.Compile()` to produce an executable delegate. Three stages — parsing, lowering, and compilation — produce the artifacts on which the runtime operates: the AST, the expression tree, and the compiled delegate.

The mechanism is straightforward to follow in the source. `PrepareCommandProgram` (`ActorHandler.cs:1097-1148`) orchestrates the first encounter: it rents a parser from a pool, parses the script, and decides — by parameter presence — whether to enter the parametric path. On the parametric path, the program’s `ProgramExpression` method (`Program.cs:182-244`) walks the AST: for each `Statement`, it calls the `Statement`’s `ExecuteExpression`, which returns an `Expression` node bound to the runtime’s parameter and output expressions. These nodes accumulate into a `BlockExpression`, which `ProgramExpression` wraps in `Expression.Lambda<Func<Parameters, Output, string>>` (`Program.cs:244`). The compilation itself is one line: `_executable = programExpression.Compile()` (`Program.cs:163`). The compiled delegate lives in the program’s `_executable` field; subsequent invocations skip the compile step and call the delegate directly with the values supplied at the call site.

The same mechanism applies one level deeper, to a feature the runtime calls *Eval parameters*. A parameter declared with the `Eval` modifier carries its own sub-script, which the runtime treats as a separable program in miniature: parsed once, compiled to its own delegate, and cached against the parameter’s name. The cache structure is a dictionary keyed by parameter name to a tuple of the

eval script and its compiled executable (`Program.cs:305`). On each invocation, the runtime compares the parameter’s current script to the cached one; if they differ, the entry is invalidated and the sub-program is re-parsed and re-compiled (`Program.cs:323-331`). The principle scales: any region of the program that has a stable textual identity admits the same treatment as the whole.

3.2 Interpretation retained

The runtime preserves interpretation as a first-class execution mode, not as a legacy fallback. When a script presents itself without user parameters, or when a per-actor policy directs the runtime to interpret regardless of parameter presence, the program executes by walking the AST directly: each Statement carries an `Execute` method that runs against the program’s current state, with no expression tree built and no delegate compiled. The interpreted path costs zero compilation but pays the cost of tree traversal at every invocation. Its place in the runtime is precisely the inverse of the compiled path’s: useful where invocation is one-shot or unanticipated, useless where the program will be reused.

The decision is governed by a per-actor `CompilationModePolicy` enum with three values — `Automatic`, `AlwaysCompiled`, and `AlwaysInterpreted` — declared on the actor and defaulted to `Automatic` (`Actor.cs:8-13`). On the parametric path, `PrepareCommandProgram` calls `AdjustCompilationMode` with `useInterpretedMode: false`; on the literal path, with `useInterpretedMode: true` (`ActorHandler.cs:1117-1131`). Under `Automatic`, the policy honors that hint: parametric programs compile, literal ones interpret. Under `AlwaysCompiled` or `AlwaysInterpreted`, the hint is overridden in favor of the policy’s name (`Program.cs:134-150`). Execution itself is dispatched by another switch on the same policy: `Perform` calls `ExecuteExpression` if the program is in compiled mode and `Execute` otherwise (`ActorHandler.cs:1955-1968`).

Caching applies asymmetrically to the two principal kinds of DSL invocation. For commands — programs that mutate actor state and persist to the journal — the runtime caches only when the script declares user parameters (`ActorHandler.cs:1117`). For queries, checks, and emit invocations — programs that read state without persistence — the runtime caches when the script declares user parameters or when no parameters are supplied at all (`ActorHandler.cs:1405`). The asymmetry is operational. Commands run under a write lock and serialize, so caching pays only when the same parametric program is invoked many times. Queries run under a read lock and parallelize, so a cache entry amortizes regardless of parametric reuse — a frequently-emitted query gains from caching even when its parameter set is fixed. One final observation: cosmetic variation in the script’s text across releases generates distinct cache identifiers, but the runtime’s reaction mechanism — treated in a companion paper — matches behavior against semantic patterns rather than identifier equality, so coherence is preserved across textual variants.

3.3 Hot-loaded DSL programs over a stably-loaded domain

The third element of realization is the runtime’s support for hot-loaded DSL programs against a stably-loaded domain library. The phrase deserves care, because it is easy to misread. The hot element is the DSL: at any time during the actor’s life, a new script — never seen before by this runtime instance — can be supplied, parsed, prepared, and executed without restart, redeployment, or reflection over a re-loaded assembly. The stable element is the domain library: the host-language types and methods that the DSL invokes are loaded once, by reflection over the assemblies configured as the actor’s domain libraries, and held in a cache for the lifetime of the process. Hot-loaded programs combine stable domain elements in new ways; they do not bring new domain elements into being. The runtime is hot at one layer and stable at the layer beneath it.

The mechanism is brief in the source. When an actor is constructed, the public types of its configured library assemblies are walked once by reflection and cached against a deduplicated key: `DomainLibraries.GetOrLoad(params Assembly[])` (`DomainLibraries.cs:77-115`, with both a single-assembly and a multi-assembly overload), invoked from `ActorHandler` (`ActorHandler.cs:61`) with the actor’s `LibraryAssemblies`. The cache survives the lifetime of the process; the same assembly set, loaded by a second actor, reuses the same domain dictionary. Against that fixed surface, new DSL scripts arrive through the actor’s fluent interface — `ActorV2.Using(scriptForChk, scriptForCmd)` (`ActorV2.cs:32`) — at any time the actor is running. There is no `AppDomain` reload, no `AssemblyLoadContext` unload, no file-watcher over the assembly: the surface area is fixed at first load and dynamics live entirely above it.

The operational consequence is direct: a running actor can be queried or commanded with a script formed at call time, against any combination of the domain operations the configured library assemblies carry. A configuration adjustment, a one-off audit query, a derivation that the published endpoints do not anticipate — each can be expressed as a fresh DSL program and executed against live actor state without intermediate deployment. The relationship resembles gRPC in its directness — a procedure call against the live system rather than a navigation of REST resources — but without a pre-declared interface contract: the script itself is the contract, formed at call time. What has been shown in this section is that separability is not merely a formal property of DSL programs but one that a concrete runtime can recognize, act upon, and encode into its execution, caching, and persistence behavior. The realization characterized here is one of an extensible family: other actor runtimes — Akka, Orleans, Erlang/OTP — admit, in principle, the same construction, since a parameterized DSL surface above a host-language domain library would yield the same four-fold coherence under the same separability commitment. The expressive reach of such on-the-fly programs depends on the richness of the verbs the domain library exposes, which the next section takes up.

4. Verb richness: surface vs depth

Separability makes the program nameable; verb richness determines how much domain meaning that name carries. The verbs that a DSL invokes are not method calls on data structures. A verb in this setting names an operation against a live, stateful actor — an orchestration that the host language has been written to perform on behalf of the domain. A single verb may invoke many internal methods, traverse domain relationships, evaluate derived properties, trigger downstream behaviors, and write to multiple regions of state, all in the course of executing one DSL statement. Where a setter assigns a field, an actor verb concludes a transaction — moving the actor from one consistent state to another. The asymmetry between the surface — the few characters that name the verb in the script — and the depth — the volume of domain computation that runs when it is called — is structural, not incidental.

Consider a verb that any reader from an e-commerce or supply-chain background will recognize: the `Confirm` of a purchase order. The script that names this verb at the DSL surface is short — only a handful of tokens. What runs when the script is invoked is substantially larger: a verification that the order’s reservations are still valid, a hold-to-allocation conversion across one or more inventory ledgers, a tax computation that depends on the order’s line items and their fiscal contexts, a posting to the financial ledger, and a set of reaction triggers that propagate the confirmation to subscribed views and downstream actors. Each of these steps is itself a tree of host-language method calls; the leaf operations include arithmetic, predicate checks, persistence writes, and outbound messages. The runtime gives the consistency boundary that separates this depth from its surface first-class support: a check script and a command script can be supplied together via `ActorV2.Using(scriptForChk, scriptForCmd)` and dispatched as a single invocation through `PerformCheckThenCommand (ActorV2Invocation.cs:69)`. The check runs against the actor’s current state under a read lock; if its assertions hold, the command runs under a write lock and is persisted to the journal. If the check fails, no state change is applied and no journal entry is written — the actor’s consistent state is preserved by construction.

`Confirm` is one example among many. The same asymmetry obtains for any verb that names a domain transition rather than a field assignment — `Settle`, `Allocate`, `Reconcile`, `Release`, and a hundred others that a well-designed domain library will expose. Verb richness is not an artifact of this particular runtime; it is a property of how domain operations are conceived and named. The implication folds back into the journal density argument of §2.3: when each persisted entry refers to a domain operation that may run dozens to hundreds of internal method calls, the journal’s compactness in bytes is matched by an enormous compactness in semantic information per byte. The journal does not record state changes; it records named transactions, each of which carries its full operational meaning by reference to the domain library that knows how to run it. The value of separability would be modest if the verbs it named were shallow.

It becomes profound when each name stands for a full domain transaction.

Under porous substrates, the asymmetry inverts entirely. There, every domain element must persist into a relational schema, and deep domain logic becomes a serialization burden — each derived state, each accumulator update, each new field propagates as schema complexity, ETL transforms, and projection overhead. Under a code-as-data substrate, the same depth is rewarded, not penalized: the domain library can be made arbitrarily expressive without expanding what the journal must record. Complex abstractions compose without friction: no DTOs, ORM annotations, or projection layers stand between the domain operation and its persistence. Porosity makes deep domains expensive to represent; anti-porosity makes them economical to compose. The two directions reinforce one another: the richer the domain, the more meaning each Action carries per persisted byte; and because the journal pushes no persistence concern back into the domain, the domain is free to grow richer still. Density and domain depth are not in tension — each makes room for the other, the plain structural consequence of recording operations instead of their effects. This asymmetry is why journal density matters at all. The journal is not compressing syntax — shaving bytes from a line of text; it is preserving a reference to an operation whose semantic volume exceeds its textual surface. An `ActionId` stands not for a line but for a domain transaction: were a verb a mere setter, density would be a trivial saving, but because a verb can name a whole transaction, the compact entry preserves something large by reference rather than copying it. This is the point the prior paper makes from the other side — the substrate stays dense because what it records are operations, not their unfolded effects. The empirical magnitude of these numbers — how many operations a typical verb dispatches, how many bytes a typical journal entry occupies, how the ratio behaves at scale — is the subject of §5.

5. Empirical results

5.1 Methodology

The measurements that follow span two complementary axes. The first axis is *program complexity*: a synthetic straight-line arithmetic kernel parametric on integer inputs (depth 5 to 100 statements); a DSL-rich kernel exercising control flow (for, if), arithmetic, and parameter binding (~500 dispatched operations per invocation); and a multi-item purchase command operating against an external rich-domain aggregate. The first two tiers isolate runtime overhead independent of host-language work; the third locates that overhead within a representative end-to-end transaction.

The second axis is *host codebase*. The production-verb measurements are replicated against two independent open-source MIT-licensed DDD aggregates that share a structural shape (rich behavioral methods on aggregate roots) but represent disjoint business domains: `dotnet/eShop`’s `Order` aggregate

(e-commerce ordering with multi-item cart and a four-step state machine), and `kgrzybek/modular-monolith-with-ddd`'s `SubscriptionPayment` aggregate (subscription billing with a payment-lifecycle verb). The dual-codebase design is deliberate: if the measured properties of compilation, caching, and journaling are structural consequences of the runtime rather than artifacts of a particular domain, the same property should appear on both codebases. The selection criterion was structural similarity (DDD aggregates with rich behavior methods rather than CRUD entities) over an unrelated business domain. This criterion selects on the dependent variable, and we are explicit about what that does and does not buy: choosing two rich-behavior aggregates tests whether the measured properties survive a change of domain *within* that class — it does not probe where the mechanism stops paying off. To probe the floor rather than the favourable case, §5.2 and §5.5 add an adversarial *flat-CRUD* verb (a single field setter), and §5.6 reads the result: the CRUD verb is itself separable, so all four faces still obtain. What varies with the domain is not whether they hold but the *magnitude* of the speedup (which persists even on a trivial verb) and the separate, non-face property of *verb richness* (which collapses to $\sim 1\times$ when the domain is shallow). Both aggregates are reachable as public source for reproduction; Appendix A lists the per-section datasets and the commit SHA.

Two hosts are not a statistical sample, and no claim of generality is rested on them. The structural claims of this paper are carried by the argument of §1–§4; the two codebases serve a narrower, falsificationist purpose — to test whether the predicted properties are artifacts of one domain. Were they domain-specific, they would not appear on a second, independently-authored aggregate from a disjoint business domain. That they appear on both, with the same structure and differing only in magnitude, removes domain-specificity as an explanation; it does not, and is not offered to, estimate a population parameter. The two hosts are a robustness replication, not a survey.

Measurements were produced against the public runtime commit `b42d0f7` (Appendix A), built in **Release**, on a 13th-generation Intel Core i9-13900 (24 physical / 32 logical cores), 64 GB RAM, Windows 11 (build 26200), .NET 9.0.14 (RyuJIT, AVX2). Two classes of measurement are reported and they differ in kind. *Deterministic* measurements — journal entry counts and payload bytes (§5.4), DSL dispatch counts and static call-graph closure (§5.5) — are exact, independent of build configuration and of any timing instrument, and reproduce bit-identically across runs; they carry the structural claims and are verified identical between Debug and Release. *Timing* measurements — the compiled-versus-interpreted speedup (§5.2), cold compile cost (§5.2), and eval-cache hit/miss (§5.3) — use the instrument each regime requires. Steady-state throughput (the speedup) uses **BenchmarkDotNet** (v0.14): 8 warm-up iterations and 15 measured iterations of 20,000 invocations each, with tiered compilation disabled (`DOTNET_TieredCompilation=0`) so the dynamically-emitted compiled delegate is fully optimized from the first measured invocation rather than averaged across a Tier-0→Tier-1 transition; means are reported with their 99.9% confidence half-interval. One-shot costs (a program's cold compile, an

eval sub-program’s recompile) are isolated at their exact call site by a runtime instrumentation hook and timed with Stopwatch — the appropriate instrument for a single hundreds-of-microseconds event, not a tight steady-state loop. Cache and pool hit rates use Interlocked counters; journal-density measurements parse the runtime’s binary journal format directly. The pool-hit-rate figures (§5.3) and the cache-footprint figures (§6.3) are such counter and GC measurements — insensitive to build configuration and to the exact commit — taken on the same public runtime line; they are reported as measured rather than separately re-pinned to **b42d0f7**, whose pool and cache code paths they exercise unchanged. The bench follows a bootstrap-then-measurement pattern: a non-parametric bootstrap script establishes the facade in the actor’s persistent symbol table once, and subsequent invocations bind per-iteration values via the runtime’s parameter API while the script string stays constant, so the compiled-program cache hits on every call — the regime the amortization argument names. Raw per-iteration data, the BenchmarkDotNet environment reports, and the commit SHA are in the companion **data/** directory (Appendix A).

5.2 Compilation amortization

Compiled execution shows a steady-state speedup over interpreted execution that varies with where the program’s work resides. For a synthetic arithmetic kernel of depth 100 — wholly DSL-bound — the speedup is $\approx 3.1\times$ (interpreted 2.99 μs , compiled 0.97 μs). For a DSL-rich kernel of ~ 500 dispatched operations exercising for-loops, conditionals, and parameter binding, it is $\approx 2.2\times$ (interpreted 16.0 μs , compiled 7.3 μs). For a multi-item purchase verb against the eShop **Order** aggregate — one DSL dispatch cascading through the **Order** constructor, three **AddOrderItem** calls, and a four-step state-machine walk to Shipped — it is $\approx 1.5\times$ (interpreted 1.97 μs , compiled 1.31 μs). These are BenchmarkDotNet means from a single run with tiered compilation disabled; the within-run 99.9% confidence half-interval is a few percent, but the ratios carry process-to-process variation of order ± 0.2 (the variation caveat in §5.6 applies), so they are read to about one significant figure. The pattern is monotonic: the more of the per-invocation work is DSL-bound, the larger the speedup; the more it is domain-bound, the smaller. The runtime amortizes only the AST-traversal overhead — the host-language code the verb dispatches runs identically in both modes, so a verb dominated by host work shows the compiler’s contribution as a smaller fraction. (Under Debug builds, with neither path optimized by the JIT, the same curve spans $1.49\times$ – $4.10\times$; the Release figures here are the defensible ones and the curve’s direction is unchanged.)

A deliberately adversarial datapoint sharpens what the curve measures. A flat-CRUD verb — a single field setter invoked as **o = c.SetValue(v)**, with $\alpha = 1$ host dispatch — does *not* collapse to a $1\times$ speedup: it measures $\approx 1.9\times$ (interpreted 0.63 μs , compiled 0.34 μs). The speedup does not track domain richness; it tracks the fraction of per-invocation time spent in the DSL dispatch that

compilation removes — reflection-based method resolution and parameter binding — and that cost is present even when the host method is trivial. What lowers the speedup is host work, not domain simplicity: the production verb, one dispatch into a heavy aggregate, sits lowest ($\approx 1.5\times$) precisely because its host work dilutes the fixed DSL overhead, and an I/O-bound verb would approach $1\times$ from below. The flat-CRUD case is genuinely adversarial only for a *separate* property — verb richness, which collapses to $\beta/\alpha = 1\times$ there (§5.5) and which is not one of the four faces — and for none of the faces themselves: the CRUD verb is separable, so it still caches, journals densely, replicates with bounded entropy, and compiles with amortization. The adversarial case maps the *magnitude* of the speedup and the *depth* a verb names; it does not turn any face off.

Specialization is paid once, at first encounter, not at every invocation. Compiling a single program scales about linearly with its statement count: for straight-line arithmetic kernels the cold compile cost, isolated at the `Compile()` call site, runs from 0.56 ms (p50, depth 5) to 4.30 ms (p50, depth 100) — a slope of roughly 39 μ s per statement. The eShop purchase verb is a single dispatch and therefore a small expression tree; it compiles cold in 380 μ s at p50 (516 μ s at p95; N=100 distinct script variants forcing cache misses). Against the 0.70 μ s per-invocation steady-state saving for that verb, the cold compile recovers itself after roughly 540 invocations — within the first moments of any actor that outlives its warm-up. These numbers are not performance claims but empirical confirmation of the structural amortization predicted by separability; the comparison is internal (compiled vs. interpreted on the same runtime), not a competitive benchmark against external systems (§5.6).

5.3 Cache amortization

The same amortization pattern applies recursively to sub-programs. A parameter declared with the Eval modifier carries its own sub-script that the runtime treats as a separable program in miniature; on a cache hit, the sub-program pays only the cost of invoking its cached delegate, while on a cache miss — when the sub-script’s text differs from the cached entry — it re-parses, rebuilds, and re-compiles. Across three sub-program complexities (a two-term arithmetic expression, a fifty-term arithmetic kernel, and a method call against a facade-bound counter on the eShop side), the cache-hit cost is 0.5–0.6 μ s at p50 while the cache-miss cost is 214–243 μ s at p50. The miss-to-hit ratio is roughly $400\times$ regardless of sub-program complexity — a separation of nearly three orders of magnitude that confirms the principle extends to any region of the program with a stable textual identity.

Allocation pressure is amortized at a finer granularity by parser and parameter pools. Three workloads measured under AlwaysCompiled policy: 1,000 single-thread invocations against a stable parametric script recorded a 100% parameter-pool hit rate; 1,000 single-thread invocations against distinct cold-cache scripts recorded a 100% hit rate on both pools; 5,000 invocations across 8

parallel threads against distinct scripts recorded 99.90% on parsers and 99.86% on parameters — twelve total misses across the parallel workload’s ~10,000 pool rents. Allocation of new Parser and Parameters instances is incurred at startup and under brief thread-fanout transients only; under steady state, both pools service the rent without allocation.

5.4 Journal density

In a parametric purchase workload against the eShop `Order` aggregate — a nine-line DSL script cascading through order construction, four `AddOrderItem` calls, and a four-step state-machine walk — 99.8% of journal entries land as compact action references: 1,001 invocations produce 1,001 action entries plus a single Define entry that carries the script body once. Each action entry’s payload averages 115 bytes — the argument vector for seventeen parameters (user identity plus four products’ details plus shared modifiers). The Define entry’s payload is 642 bytes — the script body itself, persisted once. Had each invocation stored the literal script text instead, the action payload would be $5.6\times$ larger. The runtime’s choice to persist named operations rather than their textual content is what makes this density possible. This density does not arise from compression techniques but from reference instead of duplication.

The same compaction structure appears on the second host. Against Grzybek’s `SubscriptionPayment` — a two-statement DSL surface (`NewWaitingPayment` followed by `MarkAsPaid`) over a payment-lifecycle aggregate — $N=1,000$ parametric invocations produce 1,001 Action entries (99.8% of the journal) plus a single Define entry of 207 bytes carrying the script body. Action payloads average 60 bytes. The literal-script storage would be $3.5\times$ larger. The ratio is more modest than eShop’s because both the script body and the argument vector are smaller on a narrower verb; the structural property — arguments scaling with invocations while the body is persisted once — is identical.

The magnitude of the ratio is a function of how the host’s domain API surfaces its data, not of the compaction mechanism. eShop’s `AddOrderItem` carries the full product specification per item — pid, name, price, discount, picUrl, units — so per-iteration argument vectors are dense (115 B). Grzybek’s `NewWaitingPayment` accepts five primitive parameters and `MarkAsPaid` takes none, so the argument vector is short (60 B). Hosts whose verbs identify catalog entries by short stable references compound the ratio further; hosts whose verbs string together longer parametric sequences also compound it. Either way the structural property holds: arguments scale with invocations; the script body does not.

5.5 Verb richness

For the eShop purchase verb the DSL script dispatches $\alpha = 9$ host-language invocations per DSL invocation — an exact, deterministic count, reproduced without variance across 1,000 runs. To gauge the host-language surface those

9 entry points reach, a syntactic Roslyn walker computes the *static forward closure* of the call graph, excluding trivial accessors and cutting off at the project boundary: $\beta = 24$ methods within `dotnet-eShop/src/Ordering.Domain`. For the Grzybek `SubscriptionPayment` verb, $\alpha = 2$ (its facade folds value-object construction into one host call) and $\beta = 73$ within `Payments.Domain` plus the shared `BuildingBlocks/Domain` (275 declarations indexed, 105 trivial accessors filtered). The β/α ratios are $2.7\times$ (eShop) and $\sim 36\times$ (Grzybek).

What β does and does not measure must be stated plainly. β is a *static lower bound on reachable declared surface*, not a count of operations executed: it over-counts (branches never taken at runtime remain in the closure) and under-counts (it does not follow virtual dispatch, delegates, or reflection). It is also sensitive to assembly partition — moving code across `.csproj` boundaries moves the cutoff and changes β — so it measures reachable surface *under a chosen boundary*, not an invariant of domain depth. The exact, faithful quantity is α : one DSL token dispatches several host operations. β corroborates only the coarse direction the surface-vs-depth claim of §4 needs — that the host neighborhood a single verb names is much larger than the verb’s textual surface — and nothing finer.

The $2.7\times$ -versus- $36\times$ gap, in particular, should not be read as the measured effect of a single cause. With two hosts, author coding style, aggregate granularity, and persistence-driven design pressure are fully confounded, and the data cannot apportion the difference among them. Both `Ordering.Domain` and `Payments.Domain` are RDBMS-anchored DDD aggregates — `private set` properties, EF-Core owned-entity annotations, `int?` foreign keys instead of object references, little polymorphism on the roots — and it is *plausible* that this representational pressure flattens a graph a richer domain would deepen. But that is a hypothesis these two points illustrate, not a relationship they establish; and the very features a richer domain would add — polymorphism, delegates — are exactly what the static walker cannot follow, so the conjecture that β would grow by orders of magnitude on such a domain is not one this instrument can confirm. What the measurement does support is the modest, exact claim of §4: a single persisted DSL token names a host-language operation substantially larger than itself, so the journal records far less than the behavior it commands. That observation rests on α and on the journal-density result (§5.4), both exact — not on the β/α magnitude or any causal reading of it.

An adversarial floor case makes the dependence explicit. A flat-CRUD verb — a single field setter that calls nothing — has $\alpha = 1$ and $\beta = 1$, so $\beta/\alpha = 1\times$ and the surface-vs-depth asymmetry vanishes entirely. Verb richness is therefore a property of the host domain’s depth, not of separability or the runtime, and it is not one of the four faces: where the domain is shallow there is simply no depth to name. Its collapse disables none of the faces — the flat-CRUD verb, being separable, still caches, journals densely, replicates with bounded entropy, and compiles with amortization; the journal entry is simply a compact reference to a verb that happens to name little.

5.6 Limitations and threats to validity

Several boundaries qualify the magnitudes above. *Sample of hosts.* Two aggregates are a robustness replication against domain-specificity (§5.1), not a sample from which a population magnitude can be estimated; the generality this paper claims is structural (§1–§4), and the empirical role of eShop and Grzybek is to show the predicted properties are not artifacts of one domain. *Single environment.* All timings come from one machine and one runtime build, so absolute microsecond figures are machine-specific. The portable quantities are the ratios — speedup, miss-to-hit, density — each a comparison of two paths under identical conditions, where machine-constant factors cancel. The speedup ratios additionally vary by order ± 0.2 across process launches (ordinary microbenchmark reality), so they are read to about one significant figure of confidence — their ordering and direction, not their third digit. The journal-density and dispatch-count figures, being deterministic counts, do not carry this variation. *Steady state versus production regime.* The speedup is measured with tiered compilation disabled, which isolates the fully-optimized steady state of both paths and removes the Tier-0→Tier-1 transition that otherwise makes the dynamically-emitted compiled delegate bimodal; under the default tiered regime a long-lived actor reaches the same steady state, with wider variance during warm-up. *Synthetic kernels.* The arithmetic and DSL-rich kernels isolate runtime overhead independent of host work; they are not stand-ins for whole applications, and the production verb is included precisely to locate that overhead inside a representative end-to-end transaction.

What none of these qualifications touch is the structural result, because it rests on the *deterministic* measurements rather than the timings. The journal-density and verb-richness figures (§5.4–§5.5) are exact counts and bytes, identical across builds and machines, and they carry the paper’s central empirical observation: that the journal’s expressiveness per byte rises as the host domain grows richer — each persisted action names a deeper orchestration — while the persisted footprint per invocation does not. The relationship runs both ways. A richer domain library makes each named action carry more meaning per stored byte; and the DSL, carrying no persistence concern of its own, lets the domain be modeled as richly as its designers wish without expanding what the journal must record. The timing results confirm only the *direction* of the amortization argument — that compilation pays for itself, within a few hundred invocations — not a universal constant; the structural claim does not depend on their magnitude.

No external baseline, by design. The speedup of §5.2 is an internal, controlled comparison — compiled versus interpreted execution of the same DSL program, against the same host domain code, under the same journaling, on the same runtime — so the only variable is the one separability is about: the AST-traversal overhead that compilation removes. It is deliberately not a competitive benchmark against Orleans, a hand-written compiled expression tree, or a SQL stored procedure. Those alternatives differ from the runtime here on several axes at

once — language, persistence model, domain encoding, transport — so a latency comparison would measure a confounded mixture rather than the separability effect, and this paper makes no claim that the runtime is faster than any of them. The external comparison the thesis does require is structural rather than chronometric: whether a system exhibits the four-fold coherence at all. On that axis SQL prepared statements (§6.4, §7.6) and Truffle/GraalVM (§7.2) are compared explicitly — they share separability and its execution-side consequences but, persisting rows or volatile machine code rather than the named operation, exhibit only a subset of the four faces. The speedup confirms that compilation amortizes as the principle predicts; it is not offered as evidence of competitive performance, and “compiling beats interpreting one’s own AST” is indeed its expected, modest content — the load-bearing empirical results are the deterministic density and dispatch measurements, not the timings.

Selection of hosts, and the adversarial case. The two production hosts were chosen for structural similarity (rich-behavior DDD aggregates), which is selection on the dependent variable: on its own it can only show the mechanism surviving a change of domain within that class, not where it stops paying off. The flat-CRUD verb of §5.2 and §5.5 is the corrective, and it locates the limits precisely — but the limits are not where a first reading might place them. The CRUD verb is itself separable, so all four faces still obtain: it caches, journals densely, replicates with bounded entropy, and compiles with amortization. What the adversarial case varies is not *whether* the faces hold but two quantities that are not faces at all: the *magnitude* of the compiled speedup (a gradient set by how much per-invocation time is DSL dispatch versus host work — $\approx 1.9\times$ even for the trivial verb, falling toward $1\times$ only as host or I/O work dominates) and *verb richness* ($\beta/\alpha \rightarrow 1\times$ when the verb is a bare setter), a property of the host domain’s depth that the journal-density argument uses but that separability does not entail. So this does not qualify the title’s claim. Separability remains the precondition for the four faces, and they obtain wherever a program is separable; the adversarial case maps how much each consequence delivers across domains, not a boundary where the consequences disappear. The boundary where separability holds but a face does *not* is a different case entirely — a separable runtime lacking code-as-data, exhibiting caching and compilation but not dense journaling — and it is SQL prepared statements (§6.4), the worked example of necessity without sufficiency.

6. Counter-arguments

6.1 “*This is just JIT compilation*”

A reviewer familiar with managed runtimes might object that the runtime described here is just-in-time compilation in disguise. The objection conflates two distinct compilation events. The .NET CLR’s JIT translates IL bytecode into native machine code at first invocation of any method; this happens identically whether the runtime executes its DSL by walking an AST or by invoking a

compiled delegate. What the runtime adds is a separate, prior compilation: from DSL script to typed Expression tree to IL — the work that produces the delegate the JIT will eventually translate. The speedup measured in §5.2 is a function of this DSL→IL stage, not of the IL→native stage. Both modes deliver IL to the JIT in the same way, so JIT effects cancel; what remains is the AST traversal overhead that the DSL→IL stage amortizes away. Calling this “just JIT” would erase the layer of specialization that separability makes possible.

6.2 “*Why retain interpretation? Just compile everything*”

A second objection: if compilation is so much faster, why retain interpretation at all? The answer is that compilation is economical only when the program will be reused. A script seen once and never again — an ad-hoc query against a live actor, a one-off configuration adjustment, an exploratory invocation formed at call time — provides no future invocations against which to amortize the compilation cost. Forcing such scripts through the compilation pipeline pays a cold compile cost (§5.2) — hundreds of microseconds to a few milliseconds, depending on program size — and discards the result; the compiled delegate is never invoked a second time, and the journal records the script literally rather than as a referenced action. The runtime’s `AlwaysCompiled` policy permits this on demand, but the default policy correctly recognizes that not every script is amortizable. Separability is necessary for compilation to make sense; reuse is the condition under which compilation is economical.

6.3 “*Dual paths add memory cost*”

A third objection: maintaining a compiled delegate per parametric program inflates memory at scale, and the dual-path arrangement compounds the cost. The measurement disagrees. In a single-actor cache holding 100, 1,000, 10,000, and 100,000 distinct parametric programs under `AlwaysCompiled` policy, the per-entry footprint stabilizes at approximately 6 KB across all four scales: 6,039 bytes per entry at 100 programs, 5,981 at 1,000, 5,977 at 10,000, and 5,930 at 100,000 — no super-linear overhead, no hash-table degradation. The marginal memory of an invocation against an already-cached program is effectively zero: 100,000 invocations against a single cached program retain 104 bytes total. A production actor caching dozens of distinct parametric programs and invoking them at scale incurs a memory cost on the order of hundreds of kilobytes. The cache scales linearly with distinct programs, not with invocations — well below the threshold at which the dual-path arrangement would be a structural concern.

6.4 “*This is just SQL prepared statements*”

A final objection: this is just SQL prepared statements rebranded — the same parametric-template-plus-bind-values pattern that relational databases have used for decades. The pattern is indeed the same; what differs is what the runtime persists. A SQL prepared statement is parameter-separable and

admits compilation and caching: the database parses it once, builds a query plan once, and reuses both for many invocations with different bind values. But what the database persists is row data — the projections produced by executing the statement against the underlying tables. The parameterized statement itself is not the persisted artifact; the rows it touches are. The runtime described here persists the operation rather than its row-level effects. Two of the four faces of separability — compilation and caching — are present in both systems; the journal density that follows from persisting the named operation is what SQL prepared statements do not achieve. Separability is necessary; pairing it with code-as-data is what produces the dense journal.

6.5 “*A verb dispatching dozens or hundreds of host methods produces an opaque runtime*”

A reviewer might object that a verb whose static call-graph closure reaches dozens or hundreds of host-language methods (§5.5) produces an opaque runtime, and that debugging across such depth is intractable. The objection misreads the architecture. The depth lives in the host language: domain classes, methods, business rules — debuggable with the standard tools the host already provides. The DSL surface only names what the host invokes; it does not introduce a separate execution layer that the host debugger fails to penetrate. Standard practice in the runtime described here proceeds by test-driven development with end-to-end test cases, with breakpoints and step-through performed in the host language; once a script is moved to a runtime endpoint, the same breakpoints continue to apply against the domain library it invokes. More consequentially, the journal makes post-mortem debugging tractable in a way porous substrates do not allow. When a production failure surfaces at journal entry id N, that entry refers — by name — to the exact script that produced the defect, parameterized by the exact values the script consumed. A breakpoint at the journal entry, a step-through of the named operation, and the defect reproduces; the failure is then replicated in a small end-to-end test case, fixed, and released. The dense journal is not the opacity it might appear to be on a superficial reading; it is among the runtime’s most powerful debugging artifacts, because the persisted artifact is the operation that ran, not a downstream trace of its effects.

Each objection treats compilation, caching, and journaling as independent engineering choices. The argument of this paper is that they are consequences of a single structural property; the objections dissolve when that property is made explicit.

7. Related work

7.1 Partial evaluation

The structural condition this paper formalizes has a close ancestor in **partial evaluation**, as developed by Jones, Gomard, and Sestoft (1993) and the Futamura projections (Futamura, 1971/1999). Partial evaluation specializes a program with respect to a subset of inputs known in advance, producing a residual program that depends only on the dynamic arguments — and it requires, structurally, that the program’s static and dynamic inputs be separable. Where the partial-evaluation literature applies the technique to general-purpose host languages, this paper applies the same separability principle to a domain-specific language whose programs are short, parametric by construction, and stored as journal entries; in this setting, separability is syntactic, decidable from the program’s surface rather than from static analysis. Two differences are load-bearing. First, partial evaluation *recovers* the static/dynamic split by binding-time analysis and *consumes* it: the residual program is the artifact of interest, and the separation is discarded once specialization is complete. Here the split is declared on the program’s surface and *persisted* — it survives into the journal as an (identifier, argument-vector) pair and into replication as the unit that replays to state. Second, the further consequences traced in §2 — caching as identity, dense journaling, replication-bounded entropy — fall outside the partial-evaluation tradition’s concern, which is execution speed rather than persistence; they follow from pairing separability with code-as-data persistence. Partial evaluation is the closest ancestor of the compilation face alone; on the other three it is silent.

7.2 Self-optimizing AST interpreters (Truffle/GraalVM)

The most direct comparison is not a classical technique but a production system: the Truffle framework and the GraalVM compiler (Würthinger et al., 2012, 2017). Truffle hosts a language as a self-optimizing abstract syntax tree whose interpreter nodes rewrite themselves under runtime profiling, and GraalVM applies partial evaluation to the interpreter specialized to a given AST — realizing the first Futamura projection in practice, from interpreter-plus-program to machine code, cached per call target and guarded by deoptimization. On the compilation face (§2.1, §3.1), Truffle is more sophisticated than the runtime characterized here: it profiles observed types and values, speculates aggressively, inlines across a polyglot boundary, and deoptimizes when a speculation fails. A reviewer is right to ask what program–value separability adds once Truffle is on the table. The first half of the answer is that Truffle already turns on it: its specialization holds the AST constant while frame arguments vary, and its compilation amortizes only when the same AST recurs across invocations — so it *presupposes* program–value separation. A program embedding its values would present a fresh AST per call and defeat Truffle’s reuse exactly as §2.2 describes it defeating the cache. Truffle is thus not a counterexample to the necessity

claim but an independent witness to it, operating on the same precondition.

The second half is that Truffle never leaves the execution axis. Its specialization is volatile execution machinery — machine code keyed to an in-process call target, provisional under deoptimization, discarded at process exit, bound to no stable name, replicated to no other node. There is no journal in which the specialized program is the persisted entry, no identifier under which it is referenced across time, no replication of the program as the unit of state. The separation Truffle exploits is operational and implicit, recovered and re-validated by the runtime; the separation this paper formalizes is declared on the program’s surface, decided once at preparation time without profiling or speculation, and elevated to a durable identity — a content-derived action reference that is at once the cache key, the journal entry, and the replication unit. This makes the necessity claim of §1.2 precise: declared separability is not necessary for fast execution — Truffle shows a profile-driven runtime can compile and cache without it — but it is necessary for the persistence faces, because a provisional, in-memory, profile-derived boundary cannot serve as a stable journal reference or a replication anchor. The mechanism follows the same divide: Truffle partially evaluates a host-language interpreter against the AST; the runtime here lowers the DSL AST directly into a typed expression tree and a delegate over the declared parameter vector, with no interpreter to specialize and no speculative guard to maintain.

7.3 Lambda lifting

Lambda lifting (Johnsson, 1985) is a program transformation that rewrites locally-defined functions whose free variables capture an enclosing scope into top-level functions that receive those variables as explicit parameters. The transformation makes the function’s dependencies syntactically visible — and, by doing so, makes the function amenable to ahead-of-time compilation, separate compilation, and uncluttered call-graph analysis. The structural insight is the one this paper generalizes: dependencies that remain implicit in the program’s body cannot be specialized over, but the same dependencies promoted to the surface can. Where lambda lifting transforms programs that were already written and externalizes their captured environment, the runtime described here requires programs to declare their values as externalized parameters from the start. Lambda lifting opens compilation; this paper extends the same principle to caching, dense journaling, and replication-bounded entropy by joining separability with code-as-data persistence.

7.4 Closure conversion

Closure conversion (Appel, 1992) is a compilation technique that translates a function with captured lexical environment into an explicit closure record — a pair of function pointer and environment data — that the runtime carries as a single value. Closure conversion makes implicit lexical dependencies explicit, but

preserves them as runtime data: the closure travels with its environment record, retained in memory, and re-bound on each invocation. Lambda lifting and the principle of this paper sit on the other side of the same choice: rather than carry the environment, eliminate it by promoting the variables that would have been captured into top-level parameters supplied at the call site. The runtime described here makes that choice mandatory and declarative — values must be supplied at invocation, not captured from a surrounding scope — which is what permits the journal to record the program by name without any captured-environment record traveling alongside it.

7.5 Template instantiation

Template instantiation in general-purpose languages — exemplified by C++ templates (Stroustrup, 1994; Vandevor, Josuttis, & Gregor, 2017) and the broader tradition of generic programming — applies the separability principle in another setting: a template body names operations parameterized by types or compile-time constants, and the compiler instantiates each template with concrete parameters at compile time, producing specialized code per instantiation. The structural pattern is the same as the one described here: a program that admits parametric reuse must be written so its parameters are separable from its body. Template instantiation operates over types and compile-time values; the runtime described here applies the same principle over runtime values supplied at invocation. The distinction is operational: instead of monomorphizing the program once per parameter tuple, the runtime compiles the program once and rebinds its arguments on every call. Both arrangements presuppose what this paper formalizes.

7.6 Prepared statements

Prepared statements in relational database systems (SQL standard ISO/IEC 9075:2023; Hellerstein, Stonebraker, & Hamilton, 2007) are the closest operational analog to the runtime described here. A prepared statement carries a parametric template (with placeholders for bind values), the database parses and plans it once, caches the plan, and reuses both for many invocations with different argument vectors — the same parametric-template-plus-bind-values pattern §1.1 used as pedagogical anchor for the principle this paper formalizes. The two systems share separability and its first two consequences: compilation (the query plan) and caching (the plan cache). They diverge on what the runtime persists. A relational database persists row data; the prepared statement itself is not the persisted artifact. The runtime described here persists the named operation rather than its row-level effects, completing the four-fold coherence with journal density and replication-bounded entropy. The principle this paper formalizes generalizes beyond any one runtime — it characterizes the structural condition any DSL runtime must satisfy to admit compilation, caching, and dense persistence at all.

These traditions arise in different domains — compilers, functional languages, generic programming, database engines — and they share one structural observation: programs whose dependencies are syntactically separable from their bodies admit forms of specialization that programs with embedded values cannot. What they do not share is what this paper adds. In every case above, separability is instrumental to execution and the separation is consumed by a transformation or a query planner; in none is the separated program the persisted, replayable, replicable unit of state. The contribution here is not to observe separability again but to relocate it — from an enabler of execution specialization to the precondition of a persistence regime — and to show that compilation, caching, dense journaling, and replication-bounded entropy are one precondition’s four projections rather than four independent techniques. Prepared statements make the gap concrete: they supply the closest operational analog and still exhibit only two of the four faces, because they persist rows rather than the named operation. Separability is necessary for the four-fold coherence; code-as-data persistence is what carries it past the two faces the prior art already reaches.

8. Conclusion

The argument of this paper can be stated compactly. Program-value separability — the syntactically decidable property that a DSL program declares user parameters rather than embedding values in its body — is the structural precondition under which compilation, caching, and dense journaling become economically meaningful. The four runtime faces traced here are not parallel design choices: they are downstream consequences of separability, paired with code-as-data persistence in the case of dense journaling. The runtime characterized in §3 realizes the principle concretely; the magnitudes reported in §5 confirm the predicted structural amortization across two open-source DDD aggregates — a compiled-versus-interpreted speedup that scales monotonically with DSL-bound work (roughly $1.5\times$ on the eShop purchase verb to $3.1\times$ on a synthetic arithmetic kernel), a cold compile cost ($\approx 380\ \mu\text{s}$ for the eShop purchase verb) that recovers itself within several hundred invocations, and a journal in which $\approx 99.8\%$ of parametric entries are compact action references producing a footprint several-fold denser than the equivalent literal-script storage. These magnitudes are drawn from two RDBMS-anchored DDD aggregates; whether host domains without that representational pressure would widen the surface-vs-depth gap further is a plausible conjecture, not a result that two data points — or a static call-graph walker blind to polymorphism and delegates — can establish.

The principle takes its place alongside the analysis of the prior paper of this series, *Anti-porous Architecture*, which characterized porosity as a representational sparsity problem and density preservation as the operational consequence of an event-sourced DSL journal. Where the prior paper named the defect that

domain representations on tabular substrates accumulate, this paper names the condition under which that defect can be avoided.

Prior paper	This paper
Problem: porosity	Condition for avoiding it
Density preserved	Program–value separability
Executable journal	Stable program identifier
Operations vs state	Functions vs instances

The two papers describe the same phenomenon from opposite sides: density preservation is what is achieved; separability is what makes it achievable. Together, they argue that density in domain representation is not an optimization but a structural property that follows from how programs relate to their values.

Stated as a design space, the two papers occupy a cell that little else does. Two properties are at work, and each is common on its own. *Separability* — values external to the program — is everywhere: every prepared statement, every parameterized query has it. *Persisting the operation* rather than its effects is also well known: event sourcing is built on it. What is rare is their conjunction. SQL is separable but persists rows, not the statement; event-sourced logs persist operations but their events are command-shaped *data*, not separable programs with an identity of their own; ordinary CRUD has neither. A runtime that is *both* separable *and* persists the operation produces what neither tradition has alone: a persistent operation carrying a stable identity independent of its values — an *Action*. Separability (this paper) supplies the identity; code-as-data persistence (Paper 1) makes the identified operation the durable artifact. The Action is precisely what falls out when both hold at once, and it is the corner of the design space — separable *and* operation-persisting — that this series occupies and that the surrounding traditions, each holding one property without the other, do not reach.

The principle is not bound to the runtime characterized in §3. In an Orleans- or Akka-style actor framework with a parameterized command DSL above a host-language domain library, the same precondition would predict the same four-fold coherence; in a service that uses prepared statements over an event-sourced log of named operations, three of the four faces are already available, and the fourth — replication-bounded entropy — follows once the persisted artifact is the operation rather than its row-level effects. The realization in §3 is one example of a construction the principle admits, not the only path to it.

Two extensions of the principle remain to be developed in subsequent work. The fourfold coherence described here does not exhaust the operational consequences of an externalized-parameter, locality-bound substrate: a persistent local write buffer with asynchronous remote replication, for instance, is structurally enabled by the same commitments — the actor’s isolation guarantees

that locally-buffered, not-yet-replicated entries are invisible to other contexts by construction, not by convention. The persistent buffer and its zero-downtime implications are treated in a companion paper. A second companion paper takes up the reaction mechanism whose semantic-pattern-matching has been mentioned in passing throughout this argument, and the consistency contract under which observable state remains coherent across actors. In each case the same structural observation continues to apply — what makes the journal compact, what makes deep domains compose without friction, is the separation of the program from its values. Porosity makes deep domains expensive to represent; anti-porosity makes them economical to compose.

Code provenance

Source-code references in this paper resolve against the public Puppeteer repository at commit `b42d0f7` (2026-05-26) — the same public snapshot cited by the prior paper of this series. The empirical measurements of §5 were produced against this commit, built in Release (Appendix A, §5.1). The snapshot is archived in Software Heritage under the following persistent identifier:

```
swh:1:dir:177e2d61486bdbdfd2d5e774fcf392a45406e60d;  
  origin=https://github.com/alvaronCubo/puppeteer;  
  anchor=swh:1:rev:b42d0f76b4278c36a45c221cc1453e3ee8ffb3de
```

The core files cited inline (`ActorHandler.cs`, `Program.cs`, `Actor.cs`, `ActorV2.cs`, `ActorV2Invocation.cs`, `DomainLibraries.cs`) are byte-identical between this commit and the prior public snapshot `2f31f96` (2026-05-18); the line numbers below resolve against either.

Inline references of the form `file.cs:NN` (e.g., `ActorHandler.cs:38`) resolve against this snapshot. A reader can construct a per-file SWHID by adding the qualifiers `;path=<path>;lines=<NN>` to the directory SWHID above. Future commits to the repository may renumber lines; the SWHID preserves the cited state independently of any future change to the repository or its hosting.

Acknowledgments

The author used large language models (including Claude and ChatGPT) as editorial assistants for language refinement, structural feedback, and literature navigation. All original ideas, terminology, theoretical constructs, and technical content presented in this work are solely the author's.

References

- Appel, A. W. (1992). *Compiling with continuations*. Cambridge University Press.
- Futamura, Y. (1999). Partial evaluation of computation process — an approach to a compiler-compiler. *Higher-Order and Symbolic Computation*, 12(4), 381–391. <https://doi.org/10.1023/A:1010095604496> (Original work published 1971 in *Systems, Computers, Controls*, 2(5), 45–50.)
- Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.
- Grzybek, K. (2024). *Modular monolith with DDD* [Computer software]. GitHub. <https://github.com/kgrzybek/modular-monolith-with-ddd> (commit 91c8ef2, accessed 17 August 2026).
- Hellerstein, J. M., Stonebraker, M., & Hamilton, J. (2007). Architecture of a database system. *Foundations and Trends in Databases*, 1(2), 141–259.
- Hevner, A. R., March, S. T., Park, J., & Ram, S. (2004). Design science in information systems research. *MIS Quarterly*, 28(1), 75–105.
- International Organization for Standardization. (2023). *ISO/IEC 9075:2023 Information technology — Database languages — SQL*.
- Johnsson, T. (1985). Lambda lifting: Transforming programs to recursive equations. In J.-P. Jouannaud (Ed.), *Functional programming languages and computer architecture* (pp. 190–203). Springer-Verlag. (Lecture Notes in Computer Science, Vol. 201)
- Jones, N. D., Gomard, C. K., & Sestoft, P. (1993). *Partial evaluation and automatic program generation*. Prentice Hall.
- .NET Foundation. (2026). *eShop: A reference .NET application* [Computer software]. GitHub. <https://github.com/dotnet/eShop> (commit 9b4f943, accessed 17 August 2026).
- Rivera, A. (2026a). Anti-porous architecture: a unified design principle for CQRS + Actor + Event-Sourcing systems. *Puppeteer Papers Series*, Paper 1. Zenodo. <https://doi.org/10.5281/zenodo.20404863>
- Stroustrup, B. (1994). *The design and evolution of C++*. Addison-Wesley.
- Vandevoorde, D., Josuttis, N. M., & Gregor, D. (2017). *C++ templates: The complete guide* (2nd ed.). Addison-Wesley.
- Würthinger, T., Wöß, A., Stadler, L., Duboscq, G., Simon, D., & Wimmer, C. (2012). Self-optimizing AST interpreters. In *Proceedings of the 8th Symposium on Dynamic Languages (DLS '12)* (pp. 73–82). ACM. <https://doi.org/10.1145/2384577.2384587>

Würthinger, T., Wimmer, C., Humer, C., Wöß, A., Stadler, L., Seaton, C., Duboscq, G., Simon, D., & Grimmer, M. (2017). Practical partial evaluation for high-performance dynamic language runtimes. In *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI 2017)* (pp. 662–676). ACM. <https://doi.org/10.1145/3062341.3062381>

Appendix A — Code references

The references below cite source locations in the Puppeteer codebase as `file:line` pairs against the public commit `b42d0f7` (see *Code provenance*); the cited core files are byte-identical at the prior public snapshot `2f31f96`. The benchmark harness that produced §5 is public: the MSTest cold-compile, eval, journal-density, and dispatch-count benches live in the runtime repository under `tests-local/` at commit `b42d0f7`, while the BenchmarkDotNet speedup harness, the synthetic compile/eval sweeps, and the Roslyn call-graph walkers live in this paper repository’s `labs/` directory (`lab01-bdn-speedup/`, `lab05-eshop-roslyn/`, `lab05-grzybek-roslyn/`). Datasets cited in §5 — raw per-iteration CSVs, the BenchmarkDotNet environment reports, and per-lab summaries — are stored in the companion `data/` directory of this paper repository, each stamped with the commit SHA. The measurement environment is recorded in §5.1.

§1.2 — Formal characterization

Reference	Location	What it shows
Runtime self-documentation as $F(x_1, \dots, x_n)$	<code>ActorHandler.cs</code> <code>:1085</code>	Comment articulating the parametric program model
Parametric/literal regime documented in code	<code>ActorHandler.cs</code> <code>:1091-1093</code>	Comment encoding the binary taxonomy: scripts without user parameters are interpreted, uncached, persisted as Script entries; scripts with user parameters are compiled, cached with an ActionId, persisted as Action entries

§3.1 — Compilation pipeline

Reference	Location	What it shows
Orchestration on first encounter	<code>ActorHandler.cs:1097-1148</code>	<code>PrepareCommandProgram</code> rents a parser, parses the script, decides parametric vs literal path
AST traversal	<code>Program.cs:182-244</code>	<code>ProgramExpression</code> walks <code>Statements</code> , accumulates <code>Expression</code> nodes into a <code>BlockExpression</code>
Lambda composition	<code>Program.cs:244</code>	<code>Expression.Lambda<Func<Parameters, Output, string>></code> wraps the <code>BlockExpression</code>
Compilation step	<code>Program.cs:163</code>	<code>_executable = programExpression.Compile()</code>
Eval parameter cache structure	<code>Program.cs:305</code>	Dictionary keyed by parameter name to <code>(EvalScript, Compiled)</code> tuple
Eval recompile on script change	<code>Program.cs:323-331</code>	Cache lookup, invalidation, re-parse, re-compile

§3.2 — Interpretation retained

Reference	Location	What it shows
<code>CompilationMode</code> Policy enum	<code>Actor.cs:8-13</code>	Three policy values (<code>Automatic</code> , <code>AlwaysCompiled</code> , <code>AlwaysInterpreted</code>); default is <code>Automatic</code>
Policy hint passed at call site	<code>ActorHandler.cs:1117-1131</code>	<code>AdjustCompilationMode(useInterpretedMode, policy)</code> — <code>useInterpretedMode</code> : <code>true</code> for literal path, <code>false</code> for parametric
Policy resolution switch	<code>Program.cs:134-150</code>	Sets <code>IsCompiledMode</code> from policy + hint
Execution dispatch	<code>ActorHandler.cs:1955-1968</code>	<code>Perform</code> calls <code>ExecuteExpression</code> if compiled, <code>Execute</code> otherwise
Commands caching rule	<code>ActorHandler.cs:1117</code>	Caches only when <code>parameters.HasUserParameter()</code> is <code>true</code>
Queries caching rule	<code>ActorHandler.cs:1405</code>	Caches when <code>EMPTY_PARAMETERS</code> or <code>HasUserParameter()</code> is <code>true</code> (asymmetry vs commands)

§3.3 — Hot-loaded DSL programs

Reference	Location	What it shows
Domain library loading	<code>DomainLibraries.cs:77-115</code>	<code>GetOrLoad(params Assembly[])</code> reflects public types and caches them once per deduplicated assembly set; a single-assembly overload remains for the back-compat path
Library binding to actor	<code>ActorHandler.cs:61</code>	<code>libraries = DomainLibraries.GetOrLoad(LibraryAssemblies)</code>
Fluent script invocation	<code>ActorV2.cs:32</code>	<code>Using(scriptForChk, scriptForCmd)</code> introduces new scripts at runtime against the cached domain

§4 — Verb richness

Reference	Location	What it shows
Check-then-command fluent dispatch	<code>ActorV2Invocation.cs:69</code>	<code>PerformCheckThenCommand()</code> — read-lock check + write-lock command + journal write, with rollback on check failure

§5 — Empirical datasets

The datasets that produced the magnitudes reported in §5 are stored in the companion `data/` directory of the paper repository. The production-verb measurements are produced against two independent open-source aggregates:

- **eShop** — `dotnet/eShop` (MIT), `src/Ordering.Domain/AggregatesModel/OrderAggregate/Order.cs`. Multi-item purchase via 10-arg constructor plus four `AddOrderItem` calls plus a four-step state-machine walk to `Shipped`.
- **Grzybek** — `kgrybek/modular-monolith-with-ddd` (MIT), `src/Modules/Payments/Domain/SubscriptionPayments/SubscriptionPayment.cs`. Subscription payment lifecycle via `Buy` factory plus `MarkAsPaid`.

The synthetic kernels in §5.2–§5.3 (arithmetic, DSL-rich) carry no external-codebase dependency; their datasets are in `data/paper2-synthetic/`.

Section	Dataset directory	Lab	Host	What it contains
§5.2 (compilation speedup)	data/lab01-esshop/	Lab 1	eShop + synthetic	BenchmarkDotNet compiled-vs-interpreted (Release, <code>DOTNET_TieredCompilation=0</code>); GitHub-markdown + CSV reports incl. the BDN environment block
§5.2 (compile cold cost, production verb)	data/lab02-esshop/	Lab 2 Tier 3	eShop	Cold compile cost per distinct script variant, N=100, isolated at the <code>Compile()</code> call site
§5.2 (compile cost by depth)	data/paper2-synthetic/	Lab 2 synthetic	synthetic	Cold compile cost vs statement count (depths 5–100), 50 variants per depth
§5.3 (eval cache hit vs miss)	data/lab03-esshop/	Lab 3 Tier C	eShop	Stable vs mutating eval text; per-iteration end-to-end and isolated eval-compile ticks
§5.3 (eval cache by complexity)	data/paper2-synthetic/	Lab 3 synthetic	synthetic	Hit vs miss for 2-term and 50-term eval sub-programs

Section	Dataset directory	Lab	Host	What it contains
§5.4 (journal density)	<code>data/lab04-eshop/</code>	Lab 4	eShop	Action / Script / Define entry counts and payload bytes parsed directly from BinaryEventCodec journal_*.bin
§5.4 (journal density, replication)	<code>data/lab04-grzybek/</code>	Lab 4	Grzybek	Same entry-type analysis applied to the SubscriptionPayment lifecycle verb
§5.5 (verb richness)	<code>data/lab05-eshop/</code>	Lab 5	eShop	α DSL dispatch counter (runtime) plus β Roslyn forward closure over <code>Ordering.Domain</code>
§5.5 (verb richness, replication)	<code>data/lab05-grzybek/</code>	Lab 5	Grzybek	$\alpha + \beta$ over <code>Payments.Domain</code> plus shared <code>BuildingBlocks.Domain</code>

Each dataset directory contains a `headline.md` summary, raw CSVs, and a Git SHA stamping the runtime version against which the lab was run. The Roslyn walker source for the β half of §5.5 lives at `labs/lab05-eshop-roslyn/` and `labs/lab05-grzybek-roslyn/` — both are self-contained .NET 9 console projects that can be re-executed against the public source trees of the respective aggregates.

Reactions and the partition

TL;DR

When CQRS introduces eventual consistency, the standard treatment is as a cost: a contract between separated read and write stores that the application must absorb. This paper argues that for actor-native systems the contract is upside-down. The primitive is not eventual consistency at all; it is the developer's pragmatic partition of the work an event implies into *now* (what the verb's contract promises before responding) and *deferred* (placed wherever the developer chose — same thread, another node, the cluster).

Reactions are the artifact through which the partition is exercised — and they are not async events under another name. Each Reaction matches a *trajectory* through the actor's journal, a sequence of one or more stages naming a saga or lifecycle, statically compiled against the domain's libraries with typed identifiers propagated across stages. Reactions serve two purposes simultaneously. The first is **pragmatic deferral** — for latency budget, parallelization, scaling. The second is **categorical segregation** — keeping operational concerns (email, BI, third-party messaging, telemetry) from propagating into domain methods. Reactions are *not* technology-free; they are precisely where that technology lives, **encapsulated** so that other domain verbs do not inherit the dependency, and the encapsulation is structural — the language itself blocks the operational concern from leaking back. The library, in turn, stays free of that propagation, built under its legitimate programming paradigm (currently object orientation). Both purposes share one guardian: **friction**. If Reactions cost the developer more than inline code, the partition is not exercised; verbs degrade and the domain library is contaminated by precipitation. The friction is an architectural responsibility of the framework, not a discipline question for the developer.

Three consequences fall out of the primitive when it is low-friction. The primary one is a **domain library that closes** — that can be

declared *task done* because it is structurally free of present and future operational tools. The other two are **fast verbs by construction** and **opt-in eventual consistency** (the developer names each deferred effect by hand, inverting the CQRS-classical assumption). **Eventual consistency is the shadow of the partition, not the design choice.**

Claims this paper makes

This paper makes four core claims. The detailed assertions are folded into them as sub-points; two statements the argument relies on but does not originate are listed afterward as supporting observations.

1. **The now/deferred partition is the primitive — drawn, not deduced, and placed.** Within an event’s frame, formal correctness places *everything* in *now*: every domain effect entailed by the verb’s contract should, formally, resolve before the event closes. The *now/deferred* split is therefore not a decomposition deduced from correctness but a pragmatic boundary the developer **draws** inside the event’s implied work — feature by feature, governed by latency budget, parallelization opportunity, and the categorical-purity argument of claim 2(b), not by the formal contract. The partition also carries a **placement**: at the moment of deferral the developer decides where the deferred work runs, along a proximity/urgency axis (close-to-action $\leftrightarrow$ remote-and-deferred). Drawing the cut and placing it are one act, not two.
2. **The artifact that realizes the partition serves a double purpose, structurally enforced, and is declarative trajectory matching — not an event handler.** In the instantiation, the artifact is the Reaction.
 - **(a) Pragmatic deferral** — work the verb cannot afford within its latency budget, or that is better run in parallel, is deferred.
 - **(b) Categorical segregation** — extrinsic operational concerns (email, BI, telemetry, webhooks) land in the artifact, not in domain methods. The segregation is of *propagation*, not of presence: the artifact invokes the technology, encapsulated so other domain verbs do not inherit the dependency. This extends the anti-porosity principle of Paper 1 to a fourth layer — the operational edge — and the encapsulation is *structurally* enforced, not conventional: a Reaction’s emit path runs as a query (`PerformEmit, isQuery:true`, blocking `expose` and `journal writes`), so the language refuses to let the operational concern leak back into the actor’s state (`ActorHandler.cs:2408-2479`). The legitimate paradigm inside the library (classes, inheritance, invariants, domain methods) is preserved precisely because the encapsulation is non-leaky (§3.1).

- **(c) Mechanism: trajectory matching.** The artifact is a rule the engine applies over the actor’s journal — a multi-event *trajectory* of **Seek** stages, statically compiled against the domain’s **Libraries** at build time (a stage naming a class the domain does not contain fails to compile), with typed identifiers propagated across stages. The matcher is CEP-style and not claimed as novel (§9.3); what is distinctive is the substrate, the actor’s own journaled invocations (`Pattern.cs`, `Reaction.SolveActionReferences()` at `Reaction.cs:1444`, `MatchTree.TryMatchAtLevel()`).
 - **(d) Two placement modes, not a slider.** The placement of claim 1 is realized as a choice between **Cue** (push loop on a separate thread, against the actor’s live in-memory state) and **Job** (on-demand against the replicated journal, freely placeable), further narrowed by activation scope (`DirectorOnly/CastOnly/Company`) (`Reaction.cs:20-24`, `Reaction.cs:68-82`).
 - **(e) Elision closes the journal-density loop.** When the pattern correlates a trajectory, its constituent events can be marked elided (`Metadata.Elide`, §6.5), so the *skips* of Paper 1 become the natural outcome of a declarative correlation rather than a post-hoc optimization.
3. **Friction is the architectural guardian of the partition.** The artifact must make declaring deferred work no costlier than inlining it, or both purposes collapse together: if deferring costs more lines or cognitive load than inline code, the partition is not exercised — verbs degrade *and* the domain library is contaminated by precipitation (the worse failure, being hard to revert and propagating couplings). Friction is an architectural responsibility of the framework, not a question of developer discipline.
 4. **Three consequences follow when the partition is low-friction, ranked.** The primary one is **a domain library that closes** — declarable *task done* by being free of present and future operational tools (§5.1); its operational form is that the domain becomes a single real artifact maintained over years, rather than classes scattered and drifting across services (§5.1). Secondary are **fast verbs by construction** (§5.2) and **opt-in eventual consistency**, in which the developer names each deferred effect by hand, inverting the CQRS opt-out assumption (§5.3). The three are not features or design choices; they fall out of the primitive (claims 1 + 2 + 3) when its exercise is low-friction.

Supporting observations (relied on, not originated here): - *Logical, not physical, CQRS.* Within an actor, reads and writes operate on a single in-memory state, strongly consistent at the actor boundary; the eventual-consistency contract of classical CQRS does not apply there. This is the semantics of a stateful actor (Hewitt, Bishop, & Steiger, 1973; Agha, 1986), recorded as the ground on which §5.3 reframes the consistency vocabulary — not a contribution. (*Verification: PerformCmd/PerformQry over a shared actor state in §6.*) - *Con-*

tinuation across the deployment switch. Deferred work in flight when a node hands off is not lost: a Reaction’s cursor is durable and replicated with the journal (`DiaryStorage.GetReactionCheckpoint`), so the successor resumes from the last *confirmed* position — at-least-once, as within a single node (§3.6). The zero-downtime coordination that makes the handoff *seamless* (`Theater.aliveGate`) is a separate companion-paper topic; it governs *when* the successor goes live, not *whether* work survives. (The one common operational exception is a coordinated cut-over of a Kafka producer/consumer pair on a message-format change.)

When NOT to use this approach

The primitive of pragmatic partition is general; the implementation pattern that realizes it has narrower limits. We name three regimes in which the analysis presented here either does not apply, applies only partially, or invites overreach.

A. Domains where every effect must be transactionally atomic with the response

Some domains — core ledgers, certain regulated transaction systems — require that every observable effect of a verb commit in lockstep with the response. This bounds the *deferral*, not the model. The developer still exercises the partition by drawing it at zero deferral (everything in *now*, §2.2): the verb’s effects stay on the command path and there is no deferred branch to admit eventual consistency. This is, in fact, where the construct’s correctness is *most* verifiable — the operation runs against a single in-memory state and is recorded as one journal entry, so it is atomic and exactly-once by the journal’s uniqueness (a journal-internal effect, not the at-least-once external case of §3.6). What the domain forbids is *deferring* an effect to a Reaction whose delivery is at-least-once; such an effect is kept in *now*, or made idempotent at its sink (§3.6). The opt-in framing of claim 4 is then exercised as opt-out-of-nothing.

The standing objection to a journal substrate in these domains is not correctness but *rehydration cost*: reconstructing an aggregate such as a running balance by replaying every operation is $O(N)$. The discipline that answers it is an **in-band state checkpoint** — opening a command with a journaled *observation* of the precondition it depends on. The command reads its own current state into a parameter (an `Eval` over the present value) and records that observation alongside the operation, so the resulting entry is self-contained: it carries both the prior observable state and the effect. Concretely, a counter: a debit command on an account evaluates `account.currentBalance()` as a parameter at the head of the operation — the observation is taken before the debit applies and journaled in evaluated form as part of the same entry — so the entry records both the balance the operation observed and the debit it applied, and no earlier entry is needed to interpret it. (The observation rides the `Eval` parameter path — the precondition the command reads before it operates; it is distinct from `expose`, a script statement the developer places to externalise a computed value into the

entry’s `ExposeData`, §6.7.) The journaled observation is not state smuggled past the journal — it is a historical fact, as legitimate as the operation that follows, which reads it as context. Because one entry now reconstructs both, predecessor entries can be marked elided (`Metadata.Elide`, §6.5) and rehydration drops toward $O(1)$ per actor — the journal-density mechanism of Paper 1 applied to a high-volume aggregate, trading disk for rehydration time at the extreme of one checkpoint per operation. With this discipline the journal substrate is viable in precisely the atomic, high-throughput domains a bare reading of the partition would exclude — making this the domain where the construct is at once applicable and, by the journal’s uniqueness, its correctness verifiable. That answers the question the exclusion would otherwise raise: where correctness matters most, the construct is not absent but at its most checkable.

B. Teams without appetite to think in terms of partition decisions

The primitive shifts a design responsibility onto the developer that some teams do not want. A team that treats Reactions as “fire and forget” — without modeling the contract that follows — loses the very property the primitive was supposed to confer. The framework can lower the friction; it cannot supply the design judgment.

C. Workflows where the deferred branch’s failure is structurally inadmissible

A pragmatic partition presupposes that the deferred branch can fail or retry without compromising the verb’s already-committed contract. Workflows in which a downstream failure of the deferred work must abort the verb retroactively — sagas with compensating actions are one such pattern, but the right place to discuss them at depth is the companion paper on Choreography. Some such workflows are admissible under the partition; others are not. The honest position is that the partition is not the right primitive for every long-running workflow.

1. Introduction: a developer says things

This paper makes an analytic theory contribution. It identifies a structural primitive that frameworks for actor-native systems have left implicit — the developer-controlled partition of an event’s implied work into a “now” branch and a “deferred” branch with a placement decision attached — and derives the design principles required to exercise that primitive at scale. The instantiation that demonstrates the construct is realizable is the Reaction primitive of the Puppeteer framework, treated in §6; it is the existence proof of realizability, not

the substance of the claim. The genre is the one Gregor (2006) names *theory for analyzing* (Type I): it introduces a construct that lets the phenomenon be described and classified, with empirical evaluation supplementary. The Hevner-style design-science *evaluation* of the artifact (Hevner, March, Park, & Ram, 2004) is deferred to a companion case-study paper; the Type I frame separates construct introduction from artifact assessment cleanly.

A developer is modeling a purchase-order endpoint. They are not thinking about transactions, consistency contracts, or CQRS sides. They are thinking about roles and what each role does next: “*the cashier confirms the purchase and locks the requested items. Then a confirmation goes out, and the rewarder evaluates whether any promotional campaign applies.*” They say it at a whiteboard, and a few hours later it is in code — not as a set of architectural decisions translated through a vocabulary, but as the same sentence the developer just spoke. In C# a developer says things — defines classes, writes methods, encodes invariants. What an actor-native runtime adds is that the same act of saying things now extends to the endpoint level: who handles the request, what the response promises, what is handled afterwards.

That extension carries a question the developer has not been asked to answer in industrial practice for some time, and most of this paper is about it. When the developer said “*first this, then that, then the other*”, they drew a partition — between what the verb’s response promises *now* and what is *deferred and placed elsewhere*. They drew it without naming it, on pragmatic grounds: latency, parallelization, and the sense that some concerns simply do not belong inside a domain method. The literature has names for the consequences of that partition — domain libraries that close, fast verbs by construction, opt-in eventual consistency — but the developer never reached for those names. They reached for a sentence about who does what, and the runtime was kind enough to let them write the sentence as the program.

The classical actor model — Akka, Orleans, Proto.Actor and their lineage — provides actor isolation and message passing, but it does not provide a declarative, journal-respected mechanism for what the verb commits to as part of its contract. What happens after the verb is, in those frameworks, side effect: a method call into another grain, a `tell` to another actor, a registered timer, a write to a queue. None of those are part of the actor’s program in the sense the journal preserves (substrate-level *program* throughout this paper refers to the construct of Paper 2 §1.2: the pair of domain library and journal of invocations). This paper argues that Reactions are that missing mechanism — and that without it, the journal-density property established in Paper 1 and Paper 2 cannot survive contact with operational requirements that span the verb’s response.

This paper takes that sentence apart in turn. §2 names the partition the developer drew. §3 examines the artifacts the developer composed to write it: a verb on the actor — a command, a query, or a check-then-command — and the Reactions that handle what the verb deferred. §4 looks at what makes writing a Reaction cheap or expensive, and why the answer matters more than it appears

at first. §5 follows what falls out when the writing is cheap: three consequences, ranked. §6 grounds the discussion in code from the public Puppeteer codebase. §7 places the design against Akka and Orleans, §8 addresses the predictable objections, §9 sets the work in the broader literature, and §10 returns to the developer’s sentence and observes what changed.

2. The pragmatic partition: now versus deferred

2.1 The first cut is *who*

The first cut the developer makes when modeling is not “what feature is this” — it is *who*. Cashier. Packer. Deliver. Rewarder. Once the actors are named, the work distributes itself: each role naturally absorbs the verbs that belong to it, and the partition between *what I do now* and *what I see go by and react to afterwards* emerges from the modeling, not from a separate engineering decision. The developer never says “this is opt-in eventual consistency”; they say what each role does. That sentence carries both role and feature — the cashier role *and* the priority order, what is settled before responding versus what is handled after — and it is not feature-by-feature; it is role-by-role-with-feature.

2.2 The partition is drawn, not deduced

Formally, the situation is sharper. Every domain effect implied by the event sits, by formal correctness alone, in *now*: nothing in the verb’s contract entitles the developer to defer anything. The license to defer is therefore not deduced — it is **drawn**. The developer draws it while modeling, on pragmatic grounds, and the grounds are three. First, a latency budget the verb cannot exceed without breaking the actor’s serial-mailbox invariant — the topic of §5.2. Second, an opportunity to parallelize work across roles or across nodes — the placement question of §2.3. Third, a categorical sense that certain concerns (a confirmation going out, an analytics ping, a third-party update) do not belong inside a domain method at all, regardless of timing — the case developed in §3.3. The three grounds are independent of one another, and a single deferral may rest on more than one.

2.3 Placement is part of the same act

At the same act of drawing the now/deferred partition, the developer also decides where the deferred work runs. The axis is not continuous: it has two canonical points named explicitly by the framework. *Cue* runs immediately on a separate thread, against the actor’s live in-memory state; *Job* runs on demand against the replicated journal, freely placeable on a follower in another pod or another machine. Section 6.1 develops the two modes in detail. The point relevant here is that partition and placement are one decision, not two — choosing what gets deferred and choosing which of the two structurally distinct modes carries it are the same act, not separable concerns.

This partition — its drawing and its placement, taken as one act — is the primitive of the rest of the paper. Everything that follows is downstream of it. §3 examines the artifacts through which the developer writes the partition (a verb on the actor and the Reactions that handle what the verb deferred), naming the two purposes those Reactions serve. §4 looks at what makes writing them cheap or expensive. §5 follows what falls out when the writing is cheap.

3. The double purpose of Reactions

Before naming what Reactions exclude from the domain library, the paper must name what they do **not** exclude. The domain library is not tool-free. It is built under the legitimate programming paradigm — currently object orientation — that sustains the language of the domain itself. Once that legitimate tool is named (§3.1), the two reasons to defer can be developed without inviting the misreading that “pure” means “tool-free”: pragmatic deferral (§3.2) and categorical segregation of operational pollution (§3.3).

3.1 The legitimate tool inside the domain

The rich library is written in a host language under a programming paradigm — currently object orientation. Classes, inheritance, polymorphism, invariants, domain methods, and the composition of conceptual primitives are the *legitimate tool* of the domain. They live inside the library; they are what the library is made of. The categorical segregation developed in §3.3 does not exclude this paradigm — on the contrary, it is what protects it.

The DSL of an actor runtime does not compete with this paradigm. It is not a second implementation of the domain; it is the language by which the domain is invoked. The relation is structurally analogous to that between SQL and a relational engine: SQL does not duplicate the engine’s logic; it names operations the engine performs. A short SQL statement can trigger joins, locks, and execution plans far heavier than its surface text. The DSL of an actor runtime carries the same kind of leverage — a few lines invoke domain methods whose implementation lives once in the library and whose behavior may be deep. Two endpoints that share invocation lines are not duplicating knowledge; they are composing distinct calls over the same library, in the sense Hunt and Thomas (1999) named when they defined DRY as the prohibition on duplicate *representation of knowledge*, not on the appearance of similar text. The polymorphism of the domain is a first-class citizen of the DSL (Paper 1 covers the parser and symbol-table treatment of inheritance and interfaces).

A second observation belongs here, anticipating §5.1: the closability argument that follows is about the boundary of the library, not about its exhaustiveness. The library can keep growing in domain concepts under its legitimate paradigm — new classes, new invariants, new relations — and that growth is healthy and expected. Practical incompleteness is not a defect of the closability argument.

What closes is the absorption of *operational* tools that change with the ecosystem; what remains open is the domain’s own conceptual development.

3.2 Pragmatic deferral

The first reason a developer defers work is operational. The verb of the actor must complete promptly: an actor sustains its target throughput when its mailbox advances quickly, and a verb that takes too long delays every subsequent verb behind it. The constraint is not that I/O must be absent from the verb — formally permitted, an I/O may sit inside a **PerformCommand** — but that I/O accumulates badly. Many verbs each making a network call, or one verb whose latency is unbounded, both collapse the throughput the actor model is supposed to sustain. The fast verb is structural rather than aspirational, and the developer protects it by deferring.

Three patterns, each at a different scope, sustain the fast verb. *I/O hoisting at the caller* — the caller resolves expensive lookups and passes the resolved values into the **PerformCommand**, so the actor receives data and never fetches. *Saga-phasing between events* — when a workflow genuinely requires interleaved external calls, the actor advances one phase, releases, an external coordinator does the I/O, and the response re-enters as the next event. *Reactions* — work that the verb commits to as part of the contract but does not block on, deferred to a **Cue** or **Job** (§6.1). The three are not alternatives; they compose. What §3.2 names is the third — the case where the deferral is owned by the actor itself, declared as a **Reaction**.

A practical recommendation belongs to the third pattern. A **Reaction**’s body is fast-verb code, with the same throughput constraint as any verb on the actor, and it is the right home for a side effect (sending a Kafka message, calling an external service): were the same side effect to live inside a **PerformCommand**, it would violate the actor principle that commands compute against in-memory state without external I/O; lifting it into a **Reaction** keeps the actor’s command path pure while keeping the side effect’s logic in the domain layer where it belongs. The delivery guarantee that side effect inherits is *at-least-once* — the **Reaction**’s action is confirmed only after it runs (§3.6) — so an external side effect declared here must be idempotent at its sink; a journal-internal effect (**Metadata.Elide**) is the one kind that is exactly-once, by atomic commit (§3.6).

The same discipline governs multi-datacenter replication, but with a distinction the single-datacenter case lets one skip. Under journal replication, what travels between datacenters is the action — its **actionId** and parameters — and each datacenter runs the same domain code, and the same **Reaction**, against its own locally-replicated journal. So *every datacenter fires the Reaction*. Whether that is symmetry or duplication depends entirely on the effect’s target.

When the target is itself **per-datacenter** — a *local* broker whose *local* consumer writes a *local* RDBMS or BI store — N independent firings are exactly the intent: each datacenter maintains its own copy, the topology is symmetric,

and a synchronous write to a *shared* external RDBMS (which would make the datacenters contend on every match) is the anti-pattern to avoid. But when the effect has a **single, globally-visible target** — one confirmation email to the customer, one call to a payment processor — N datacenters produce N effects: N emails, N charges. That is not symmetry; it is a duplication bug. Replication does not make such an effect idempotent, and §3.6 does not claim it would. A globally-singular non-idempotent effect must therefore be either made idempotent or deduplicated at its shared target — the per-match idempotency key of §3.6 lets the target collapse the N firings into one — or confined to a single firing, which is what **DirectorOnly** (§6.6) is for: it restricts a Reaction to the primary replica so the effect is not duplicated across the topology (for a globally-singular effect the topology must make that primary unique across datacenters, not per-datacenter). The symmetric-topology argument licenses cross-datacenter replication of an effect only when its target is per-datacenter; it does not license duplicating a user-visible singular effect.

The recommended pattern is therefore to keep slow I/O *outside* the Reaction body altogether: the Reaction emits a message to a broker — Kafka in the canonical case — and a consumer downstream (a BI sink, an analytics writer, a third-party adapter) performs whatever heavy I/O the requirement demands. A typical BI consumer reduces to one insert and several selects against its own store; the actor never sees that load. The framework permits direct RDBMS calls inside a Reaction; the recommendation is not to use them.

3.3 Categorical segregation: encapsulation, not exclusion

Operational concerns demanded by requirements but extrinsic to the domain — email, BI, third-party messaging, telemetry, webhooks — land in Reactions, not in domain methods. The crucial point: the segregation is of *propagation*, not of *presence*. A Reaction can live inside the same library as the domain methods it observes — declared in the same project, next to the classes it cares about. What the framework requires is that the Reaction be its own artifact, with its own scope of effects, rather than collapsed into the body of a **PerformCommand**. The structure is therefore one of co-location without co-mingling. A Reaction whose purpose is to send a confirmation email will, of course, contain a call to an email service; the technology lives there, encapsulated, near the domain methods it serves but distinct from them. What the segregation prevents is that the email concern leak into a domain method that ought to know nothing about email — and that the next domain verb the actor processes inherit a dependency on the email API by virtue of being in the same class. *Reactions are not technology-free; they are the encapsulation that keeps technology from propagating into the domain.*

The function is structurally analogous to an anti-corruption layer inDDD, but architectural rather than stylistic: any complete instantiation treats the deferral artifact as a first-class citizen of the language, and the developer’s defaults route operational concerns there. In Puppeteer this artifact is the Reaction.

The legitimate paradigmatic tool inside the library (§3.1) is preserved precisely because the encapsulation is non-leaky.

The non-leak is structural, not stylistic. The Reaction’s emit path runs as a query: `PerformEmit` is invoked with `isQuery:true`, which blocks the script from using `expose` or declaring globals, takes a read lock that runs in parallel with queries and other emits, and leaves the journal untouched (§6.5; verification at `ActorHandler.cs:2408-2479`). The technology of the Reaction (a Kafka producer, an HTTP webhook, a BI sink) executes; the journal does not record it as a domain event; the actor’s state does not absorb the call. The language refuses to let the operational concern leak back into the actor in the parse phase, before the developer ever has the chance to make a mistake about it.

Section 3.3 connects directly to Paper 1: the anti-porosity transversal, treated there as domain + endpoint + persistence, extends here to a fourth layer — the operational edge of the domain. *Reactions are the anti-porous boundary at the operational edge.*

3.4 Orthogonality of purpose and placement

The two purposes named in §3.2 and §3.3 — pragmatic deferral and categorical segregation — and the two placement modes named in §6.1 — `Cue` and `Job` — are independent decisions. A pragmatic deferral may place close to the action as a `Cue` near-continuation, or far from it as a `Job` running on a journal-only follower; a categorical segregation may do the same. The cells of the resulting two-by-two are all populated in practice: a confirmation email may be a `Cue` (purpose categorical, placement local) or a `Job` on a follower (purpose categorical, placement remote); a heavy projection update may be a `Cue` against the live actor (purpose pragmatic, placement local) or a `Job` distributed across the cluster (purpose pragmatic, placement remote). The developer chooses both dimensions at the moment of declaring the artifact; the two dimensions remain independent in any complete instantiation — there is no “default placement for categorical concerns” or “obligatory mode for pragmatic deferrals.” Each declaration states its own combination.

A second observation about the partition belongs here. The partition is not only a *pragmatic* decision (latency, parallelization, categorical separation); it is also a *preservation* mechanism. When the actor replicates across datacenters (§3.2), the journal travels with it and so does the code: each datacenter executes the same Reactions on its own actor against its own locally-replicated journal. Were the verb’s consequences side effects bolted onto the actor at runtime — `tells`, registered timers, callbacks — the replication would either lose them or reproduce them incoherently. Because Reactions are part of the program, declared in the DSL and replicated as code, the actor behaves as its program says, in every datacenter, every time it rehydrates.

3.5 Pattern matching against the actor's trajectory

A Reaction is not an event handler. The framework's primitive is more general — and more difficult to find an analogue for in adjacent runtimes. A Reaction is a rule the engine applies over the actor's journal: a sequence of one or more **Seek** stages that, taken together, describe a *trajectory* through the actor's history. The pattern matches not when one event arrives, but when the actor has, over its lifetime, gone through the stages the rule names. Identifiers captured in one stage propagate to the next, correlating the events: an **order** named in the first **Seek** is the same **order** that must appear in the next, and in the next. The Reaction is part of the actor's program — declared in the DSL, parsed at build time against the domain's libraries, replicated as code together with the journal — not a side effect attached to the actor at runtime.

Three properties are jointly characteristic. *Multi-event*: the smallest match is a saga, not a single message. *Statically compiled*: the pattern is parsed against the domain's **Libraries** in build phase — `[_:Customer].InitiateOrder(order:Order)` resolves to the class **Customer** and to the signature of **InitiateOrder**; if either does not exist in the domain, the Reaction fails to construct. *Trajectorial, not payload-based*: what the matcher inspects is the sequence of *invocations* the actor recorded — its conduct — not the body of any one message. These three properties separate the mechanism from event handlers in Akka, grain reminders in Orleans, and topic consumers in Kafka — all of which fire on a single message rather than a trajectory. They do *not* separate it from complex event processing: correlated multi-event matching with bound, propagated variables is CEP's defining capability, and typed, host-integrated CEP engines already compile patterns against application types (§9.3). The Reaction matcher is best read as CEP over a particular substrate, not as a new matching primitive. The substrate is what differs: the matched events are the actor's own journaled *invocations* — the executable journal of Paper 1 and Paper 2 — so the pattern, the journal, and the domain share one type universe and one persistence substrate rather than a stream fed to an external engine. The contribution this paper claims is the now/deferred partition (§1), not the matcher that serves it.

This is the property that licenses everything else. Without trajectorial matching, “*the cashier's verb is fast and a Reaction handles the rewarder afterwards*” reduces to a callback. With it, what the developer has written is “*when, in the cashier's history, an order was initiated, then confirmed, then paid, react*” — a declaration about how the actor's lifecycle unfolds, not about what to do when a single event arrives. The matching machinery is examined in §6.2; a worked example with the order lifecycle is in §6.3.

3.6 Reactions are journal-indexed, not event-driven

A Reaction is not an event handler. It is a journal-indexed program cursor, and its guarantees are two semantics that must not be conflated: the cursor's

advance over the journal is exactly-once, while the external *effect* an action performs is at-least-once. Both follow from how the cursor is committed.

Failure model. Processes are crash-stop. The journal and each Reaction’s cursor are durable; the cursor — one position per **Seek** of the pattern — is part of the actor’s persistent state, advances only forward, and is validated on load so that a non-monotonic checkpoint is refused. Recovery is by replay from the last *confirmed* cursor position. Nothing below assumes a distributed transaction across the journal and any external system; the runtime does not provide one.

The cursor’s advance — exactly-once, never silently dropped. The cursor is monotonic and validated on load (a non-monotonic checkpoint is refused, *Failure model* above), so each confirmed position is reached once: a replay re-confirms the same position rather than minting a new one. The cursor advances to *confirmed* only after the Reaction’s action has run; if the action throws, or the process fails before the confirmed checkpoint is written, the position is not confirmed and the entry is replayed on recovery (**MatchTree.cs:1565-1575**). The construct therefore never silently skips a match — a failure to execute surfaces as a gap that triggers retry, not a lost effect. The result is *at-least-once execution riding on an exactly-once cursor*: the cursor never double-counts, but the action it gates may run more than once. That at-least-once execution rests on the cursor’s monotonicity, not on deduplication or idempotency written inside the Reaction body; those address the *external* edge, below.

What the cursor does not by itself guarantee — no duplicate external effects. Because the action runs before its checkpoint is confirmed, a crash between an external effect (a Kafka send, an HTTP call) and the confirmed-checkpoint write re-runs the action on recovery — the classic dual-write exposure. The runtime binds no external effect to the cursor advance, so end-to-end exactly-once for such an effect is a property of the *sink*, not of the Reaction: it holds when the downstream consumer deduplicates (an idempotency key, an upsert keyed by the entry id or **actionId**) or when the effect is naturally idempotent. The broker-mediated, locally-symmetric topology of §3.2 narrows the exposure — duplicates stay local and dedup is a local concern — but does not remove the requirement.

Where exactly-once does hold structurally. One class of effect is exactly-once by construction: a journal-internal mutation committed atomically with the cursor. **Metadata.Elide** (§6.5) marks the matched entries elided and advances the checkpoint in a single store commit (**MarkEventsAsElidedWithCheckpoint, MatchTree.cs:1512**) — effect and cursor are one atomic write, so neither a duplicate nor a gap is possible. Exactly-once is available precisely where the effect and the cursor share one transactional store; for effects outside it, the honest guarantee is at-least-once plus an idempotent sink. A measurement bears out the steady state — a **Cue** activated across several hundred matches fired once per match with no duplicate (§6.9) — but that run induced no crash, so it exercises the common path, not the crash-window duplicate this section names.

The same discipline underwrites cross-actor causation. A **Causation.Continue** (§6.5) records the send in the sender’s journal under this same cursor, so the *recording* of the causation inherits the at-least-once-never-dropped guarantee above; the delivery semantics of the send itself are the transport’s, and reach exactly-once on the same terms as any external effect — an idempotent receiver. The journal is the boundary of causation because the causation is recorded as program, not because the transport is transactional.

4. Friction as the architectural guardian

The two purposes of §3 share a guardian. If Reactions cost the developer more lines or more cognitive load than inline code, the partition is not exercised. Both purposes collapse simultaneously: verbs degrade *and* the domain library is contaminated by precipitation.

Reactions must be at least as ergonomic as inline code. If deferring costs the developer more than not deferring, the partition isn’t exercised — verbs degrade AND the domain library is contaminated by precipitation. The contamination is the worse failure: it is hard to revert and propagates couplings. The friction is an architectural responsibility of the framework, not a failure of the developer’s discipline.

The contamination is the worse of the two failures, for two reasons. First, a slow verb is observable and gets fixed: monitoring catches it, profiling diagnoses it, the team rewrites the slow verb on a known schedule. A polluted domain method looks correct and does not — its tests pass, its code reads cleanly, the email it sends is the right email. The dependency on the email API is invisible from the outside, hidden behind the method’s signature, until the day the developer needs to revert it. Second, the contamination propagates. Once `sendEmail(...)` lives inside a domain method, every caller of that method inherits the coupling — directly, by transitive call; indirectly, by sharing a class with that method and inheriting its dependencies in test setups, in mocking frameworks, in the surface area of integration. Reverting the design later requires touching every call site, every test, every place the class has appeared as an object of composition.

The architectural responsibility is therefore not to teach the developer discipline. It is to make Reactions, by construction, no more expensive than inline code. The framework’s contribution is not the partition itself — the developer always *could* have deferred, in any actor framework. The contribution is friction reduction along two dimensions. The first is language-level cost: declaring a Reaction is a few lines of fluent DSL, comparable to or shorter than an inline implementation of the same effect (§6.1, §6.5). The second is textual proximity: Reactions can be written next to the endpoint they observe — in the same file, a few lines below the **PerformCommand** they continue — so that the developer reading the endpoint sees, on the same screen, what else happens around it:

this confirmation goes out, that loyalty rule fires, this analytics counter is updated. The Reaction’s location is itself a documentation aid. When the cost of doing the right thing exceeds the cost of doing the wrong thing, no amount of design review will hold; when the costs are equal or inverted — and when the right thing is also the more legible thing — the partition exercises itself. The framework that takes this seriously is the framework that gets the partition for free.

5. Three consequences (ranked)

The three consequences below fall out of the primitive (§3) when its exercise is low-friction (§4). They are ranked: closability is primary (§5.1); fast verbs by construction (§5.2) and opt-in eventual consistency (§5.3) are secondary. The last of the three is where the paper’s central inversion becomes concrete: eventual consistency is not an architectural commitment the system makes up front but the shadow the partition casts — a consequence of a modeling decision, named one deferral at a time (§5.3).

5.1 A domain that closes

The primary consequence of the partition, when its exercise is low-friction, is a domain library that closes. The library is not just kept clean of operational pollution at any one moment; it is structurally protected against the absorption of operational tools that arrive in the future. Email today, Slack tomorrow, a webhook protocol the year after — none of these find their way into a domain method, because the path of entry has been removed at the language level (§3.3). The library that the team is writing now will still be the library the team is reading in three years, and the conceptual primitives it contains will still mean what they meant when they were declared.

At some point a domain must be declarable as “task done”. A library whose definition must absorb every present and future operational tool — email today, Slack tomorrow, a webhook protocol the year after — is a library that never closes. The only way it closes is if it stays pure: free of operational concerns by construction. Reactions are the artifact that makes that closure possible.

Two clarifications matter. First, what closes is the boundary, not the conceptual growth of the domain itself. The library can keep absorbing new domain concepts under its legitimate paradigm — new classes, new invariants, new relations — and that growth is healthy and expected. Practical incompleteness is not a defect of the closability argument:

If Puppeteer is not exhaustive, that is a practical limitation, not a claim that the domain must be seen as complete.

Second, “task done” is an organisational sentence, not a software-engineering

one. A team that can declare its domain library complete with respect to operational tools can build everything else — refactor, onboarding, upgrade, extension — on a stable foundation. In most enterprises today, “the domain” does not exist as a library at all. It exists as classes scattered across each csproj, copied between services, drifting independently over years until “the customer model” or “the order model” has become a dozen incompatible structures across the codebase. The team speaks of these as if they were one thing; in operational reality they are not. A library that closes its boundary to operational tools while remaining open to conceptual growth is the structural condition under which a single domain artifact, maintained as one over years, becomes economically possible at all.

The closability argument extends the anti-porosity transversal of Paper 1 to a fourth layer. Domain libraries shaped for the operations they describe and the conceptual primitives they admit, not for the operational tools requirements happen to demand this year, are libraries that the team can put down and pick back up without forgetting what they meant.

5.2 Fast verbs by construction

Fast verbs are a construction, not a postulate, and the construction is layered. The actor’s own command-path overhead is small enough to vanish into the domain work it carries: a minimal in-memory verb — one field write, no host work — dispatches and persists in $\approx 0.4 \mu\text{s}$ on commodity hardware (BenchmarkDotNet, Release, tiered compilation disabled; environment and method in §6.9), so a verb’s latency is governed by the domain methods it invokes, not by the runtime around them. The actor’s verbs are fast first because the actor’s state lives in memory: **PerformCmd** and **PerformQry** operate on it without touching disk or external stores. This is the logical-CQRS observation of this paper — and the broader argument of Paper 1. They are fast also because the domain methods invoked from a verb are themselves rich but local: a single verb can dispatch many calls into the host language and across many levels of the domain library, each operation cheap because each operates on the same in-memory state (Paper 2, in the empirical measurements of the purchase example). The Reaction’s contribution is third in this layering, not first: it preserves the speed of the verb against the requirements that would otherwise force I/O into it.

The structural mechanism is, in any instantiation that exercises the partition, that the deferral artifact runs outside the verb’s write semaphore. In Puppeteer, a Reaction does not run inside the write semaphore of **PerformCommand**; it runs after the command has persisted, on a separate thread (in the case of **Cue**) or on demand against the journal (in the case of **Job**). The verb returns to the caller without waiting for any deferred step to begin. The actor’s serial-mailbox invariant — which would otherwise impose a throughput cap whenever a verb does anything slow — becomes a throughput floor: peak throughput is what verbs are *able* to achieve once the slow work has been hoisted out by I/O hoisting at the caller, by saga-phasing between events, or by Reactions on the deferred

branch.

What §5.2 names, then, is the third layer of the explanation. The first two — in-memory state, rich local methods — are already present in any actor framework that respects them. What this paper adds is the condition under which they survive contact with operational requirements: a `Reaction` primitive cheap enough to absorb every effect that would otherwise be tempted to live inside the verb. Section 6 develops the mechanism in code references; the structural commitment is summarised in `Reaction.cs:1113-1255` and `ActorHandler.cs:2408-2479`.

5.3 Opt-in eventual consistency

The third consequence honors the forward-reference from Paper 1 §6: the eventual consistency that classical CQRS treats as a defining commitment becomes, in actor-native systems, a shadow of the partition rather than its design choice.

Two honesties frame what follows. First, the in-actor half is not new: that a stateful actor reads and writes a single in-memory state, strongly consistent within its own boundary, is the semantics of the actor model itself (Hewitt, Bishop, and Steiger, 1973; Agha, 1986) — the logical-CQRS observation *records* it, it does not discover it. Second, what this section offers is therefore not a mechanism but a *reframing* — an analytic lens, in the Gregor (2006) Type I sense of §1, on where eventual consistency comes from and how it is bounded. The claim is not that opt-in eventual consistency is a new capability (any stateful actor has consistent local reads); it is that, once the deferred branch is named per-`Reaction`, the eventual-consistency contract stops being a system-wide property the application must absorb and becomes a per-verb list the developer wrote by hand. That is a change in how the consistency surface is *described and reasoned about*, not in what the runtime mechanically does — which is exactly the kind of contribution a theory-for-analyzing makes.

The classical model is opt-out. Read and write stores are physically separated; their synchronisation is asynchronous by default; the application accepts the contract, including its observability gap, as the price of the architecture. Greg Young’s *CQRS Documents* make the trade-off explicit: stronger contracts are recoverable, but only by application-level effort layered over the framework’s defaults.

The actor-native model is opt-in. The actor itself is the read model — `PerformCmd` and `PerformQry` operate on a single in-memory state (the logical-CQRS observation) — so within the actor’s boundary, reads and writes are strongly consistent without any ceremony. The eventual-consistency contract appears only at the deferred branch: when the developer declares a `Reaction`, what happens after the verb’s response is named explicitly. That act of naming is what the contract is. Each deferred effect is documented at the point of deferral, with a name and a placement, and the system makes no further claim about its observability than what the developer has stated.

This inversion is *structural*, and the structural claim is stronger than legibility. Within the actor, reads and writes are strongly consistent (the logical-CQRS observation); eventual consistency arises *only* at a deferred branch, and every deferred branch is a declared Reaction. So the application-level eventual-consistency surface is **co-extensive with the declared deferral set** — it is exactly the Reactions, no more and no less — and the runtime introduces no application-level eventual consistency outside it. That is a closure property of the program, not a documentation convenience: a reader does not merely *find* the windows more easily; the windows *are* the deferrals, enumerable from the source by construction. (The cross-replica dimension — the journal’s asynchronous replication across nodes, §3.2 — is a separate consistency parameter, equally explicit as a topology choice, not eventual consistency smuggled in elsewhere.) A classical CQRS system has no such closure: its eventual consistency is a global property of the read/write separation, located in topology, synchronisation policy, and lag — not bounded by any enumerable set in the program. What this paper does *not* claim is a human-factors result (that enumerability lowers defect or comprehension rates — an empirical question for the deferred case-study, §6.10) nor a mechanised proof of the closure (a formalisation in operational semantics is future work); the claim is the structural invariant itself, which is what a theory-for-analyzing contributes.

It is worth saying plainly what naming the partition as the primitive buys, given that this paper concedes the matcher is CEP (§9.3) and logical CQRS the classical pattern read at a different level (§9.1). The partition is not the *outbox pattern*, which is a delivery mechanism for a single deferred effect (record the intent locally, relay it) — a tool one uses *inside* a deferred branch, not the decision of what to defer. It is not *sagas*, which coordinate a multi-step workflow with compensation across events — a different concern, bounded in *When NOT to use C*. And it is not the folklore “defer side-effects”, which is advice with no consequences attached to it. The analytic capacity the construct adds is that *one* design act — the developer’s now/deferred cut, with placement — is the common upstream cause of three properties the literature treats separately: a domain library that closes (§5.1), fast verbs (§5.2), and opt-in eventual consistency (this section). Naming it as the primitive is what licenses two moves nothing on that list licenses: reasoning that a single lever, friction (§4), governs all three at once; and re-attributing eventual consistency as *derived from the cut* rather than *imposed by the architecture* — the cause-and-effect order classical CQRS assumes, inverted. For a theory-for-analyzing (Gregor Type I) that unification and re-attribution is the contribution. A programming-languages or systems treatment would press further — an operational semantics in which the closure property above is a theorem and the three consequences are derived formally — and that is work a companion paper, not this one, would carry.

6. Implementation in the Puppeteer framework

This section grounds the discussion in code from the public Puppeteer codebase; the path-and-line citations throughout resolve against the commit recorded under Code provenance.

6.1 The two modes: Cue and Job

The `ReactionMode` enum exposes two values, `Cue` and `Job`, declared in `Reaction.cs:20-24`. The two are not points on a continuum chosen by the developer at runtime; they are structurally distinct execution modes, picked when the Reaction is defined and binding the rest of its semantics.

`Cue()` is the live-actor mode. The framework registers the Reaction via `actorHandler.AddRecordWrittenCallback(...)` after an initial catch-up batch (`Reaction.cs:863`); from that point on, every record written to the journal wakes the Reaction’s push loop with the new entry, and the catch-up poll itself waits on that same signal rather than on a fixed sleep. Because the Reaction shares the actor’s in-memory state with the director — the same rehydration, no separate automaton — its access is inexpensive while the actor is hot. End to end, from a command being issued to the Reaction’s action firing, latency was measured at a median of ≈ 1.3 ms (p99 ≈ 2 ms) — the signal-driven path, not a polling backoff (§6.9). `ReadForward()` is mandatory: a `Cue` cannot read backward, since its purpose is to react to the actor’s ongoing trajectory, not to its past.

`Job()` is the journal-driven mode. There is no push loop. The Reaction runs only when an external scheduler or a call to `Reactions.Execute()` triggers it. It reads the journal directly rather than the actor’s memory, which makes the Reaction *portable*: any process with read access to the replicated journal can host the work. The case enabling specialised workers — *“this pod runs the followers for campaigns; that pod runs the heavier analytics followers; another pod activates only periodically to handle older events”* — is built on this primitive. Both `ReadForward()` and `ReadBackward()` are admissible.

`Job` does not rehydrate the actor. When the Reaction needs domain-computed data, the canonical mechanism is `expose` (§6.7) — the command pre-writes the data to the journal alongside the action, and the Reaction reads it during its own rehydration without ever waking the actor. This keeps `Job` cheap to host on a journal-only follower; reaching into the live actor would defeat the cost model that justifies the mode in the first place. The non-rehydration boundary is therefore not a missing feature but a structural commitment that preserves the distinction between the two modes.

The two modes correspond to two structurally distinct points on the proximity/urgency axis introduced in claim 2(d). A comparative table:

Aspect	Cue	Job
State source	Actor’s live memory (same process)	Rehydration from journal
Lifts a separate automaton	No	Yes (follower hydrates)
Trigger	<code>OnRecordWritten</code> push callback	Explicit invocation
Latency	Sub-second	Batch / on-demand
Distributable	No	Yes
Direction	<code>ReadForward</code> only	Forward or backward

The choice between them is an architectural decision, not a tuning knob. A Reaction modeled as **Cue** cannot transparently become a **Job**: the two assume different things about what is in memory, who hosts the work, and whether the actor must be live for the Reaction to make progress. This is the §2.3 axis materialised — partition and placement are one act because, at the moment of partitioning, the developer is also choosing which of these two modes will carry the deferred work. Both modes operate as journal-indexed cursors per §3.6: the choice between them governs *who walks the journal and when*, not whether the cursor exists.

6.2 Static pattern matching against the domain libraries

The Reaction’s binding to the domain is established at build time, not at run-time. Three layers of validation make this concrete.

The first layer is static parsing in `Pattern.cs`. Every reference of the form `[_:Class]`, `[var:Class]`, or `Class(param:Type, ...)` is resolved against the actor’s **Libraries**, case-insensitively. If the class or the constructor signature does not exist in the domain library, the Reaction fails to construct — there is no runtime path that can recover from a missing type. The Libraries here are the same domain types that Paper 1 named as the legitimate tool inside the library and that Paper 2 treated as the substrate against which DSL programs name their operations; the Reaction is one more language layer that compiles against that substrate.

The second layer runs per event during rehydration. `Reaction.SolveActionReferences()` (`Reaction.cs:1444`) parses the script of each `ActionEventData` exactly once and caches it under an LRU of 100 entries; identifiers that the pattern captured are marked `IsParameter=true` so subsequent stages can reuse them without re-parsing. The cache amortises the cost of parsing the same script many times across many events; the LRU bound prevents memory growth in workloads with high script diversity.

The third layer is the runtime walk. `MatchTree.TryMatchAtLevel()` traverses

the tree of stages by BFS (`WithSharedHydration`, where captures from earlier stages remain visible to later ones) or by DFS (`WithIndependentHydration`, where each branch carries its own hydration context). The choice between the two is part of the `Reaction`'s definition and shapes how identifier propagation behaves across `Seek`, `ThenSeek`, and `ThenFinalSeek` stages.

The parser, importantly, does not execute the script's expressions. It only knows the types. Matches over non-literal values are therefore by structure and type, not by value equality — with two exceptions: literals (a constant in the pattern) and `Id` parameters with `IsParameter==true` (the captured identifiers of §6.3). The composition rules round out the layer: a `Reaction` has at least one `Seek` and any number of `ThenSeek` stages, optionally terminated by a `ThenFinalSeek`; each `Seek` may carry several `OnMatch` clauses, all of which must match against the same underlying script. The build-time validation stops a `Reaction` from being defined against a domain it does not fit; the runtime walk only ever evaluates patterns the Libraries have already approved.

6.3 Multi-event patterns: trajectory, not event

The simplest non-trivial `Reaction` names three stages of an order's life and correlates them by both the order instance and a primitive identifier:

```
actor.Reactions
  .DefineReaction("OrderCompleted")
  .Job().DirectorOnly().ReadForward().WithSharedHydration()
  .Seek("Initiate")
    .OnMatch("[_:Customer].InitiateOrder(order:Order,
      ↪  number:int)")
  .ThenSeek("Confirm")
    .OnMatch("order.Confirm(number)")
  .ThenFinalSeek("Pay")
    .OnMatch("order.Pay(number, [_:PaymentMethod])")
  .Program.Emit("...");
```

Read aloud, the pattern says: *“when, in the actor's history, a customer initiated an order with a given number N, and that same order with that same number N was later confirmed and paid — react.”*

Five details of this small example carry the weight of the primitive.

Identifier propagation. The captured `order` named at `Seek("Initiate")` is the same `order` *instance* that must appear at `ThenSeek("Confirm")` and at `ThenFinalSeek("Pay")`; the captured `number` is the same primitive *value*. The matcher correlates by typed identity for object captures and by literal equality for primitive captures — both kinds of identifiers, once captured, are treated as fixed for the remainder of the trajectory. A `Confirm` invoked with a different number, or for a different order instance, would not advance the trajectory; the trajectory is locked to *this* order with *this* number once the first stage matches.

Non-contiguity. Between **Confirm** and **Pay** the actor may have processed any number of unrelated commands; the matcher steps over them. The trajectory is a path through the journal, not a contiguous window.

One match per stage. Each **Seek** and **ThenSeek** matches exactly once by default. For one-to-many correlations — an order with several items being added before payment — the stage is marked **Many()**, valid at any position in the pattern. The surrounding language and its limits are documented in §6.4.

Pattern over conduct. What the matcher inspects is the sequence of invocations the actor recorded — its conduct as a sequence of **PerformCommand** calls — not the body of any one message. The Reaction sees what the actor *did*, not what arrived in its mailbox.

Static binding. **Customer**, **Order**, **PaymentMethod** are types from the actor's domain library, validated at build time against **Libraries**; the symbol table of the actor is available to the matcher, so polymorphism and type relations are respected without further configuration.

Three properties together — *trajectory*, *static binding*, *typed propagation* — set this apart from event handlers and grain reminders, which fire on a single message. Against complex event processing they do not set it apart but align with it (§3.5, §9.3): correlated multi-event matching is CEP's, and what differs is the substrate, not the matching power. The third property is what makes the example above readable as a saga rather than as three loosely-coupled callbacks.

6.4 Filters, time, quantifiers

The matcher accepts three families of refinement on top of the base pattern.

Where clauses. **Where(expression)** (**ReactionEngine.cs:71-83**) chains a boolean condition onto a **Seek**. Inside the expression the developer has access to **@Now** (the event's timestamp, non-nullable), **@EntryId** (the journal entry's identifier, non-nullable), the identifiers captured by the pattern, and the global **time** instance of the framework's **Temporal** type. **time.Days(n)**, **time.Hours(n)**, **time.Minutes(n)** and so on (**Reaction.cs:810**) construct **TimeSpan** values for relative windows. The framework rejects always-false comparisons at build time: comparing **@Now** or **@EntryId** against **null** raises a **LanguageException**, since their types do not admit a null sentinel.

Quantification. The base pattern matches each **Seek** exactly once. A stage marked **Many()** (**ReactionEngine.cs:146**) is the existential modifier for one-to-many correlations: it collapses multiplicity — firing on the first match and tolerating the rest on the same trajectory — and, unlike a positional repeat, it is valid at any stage, not only the first. Two refinements compose with it: **.AtLeast(n)** (**ReactionEngine.cs:166**), a structural quantifier requiring at least *n* matches before the stage is satisfied, and **.Accumulate(name)** (**ReactionEngine.cs:157**), which hands the raw accumulated set to the next stage under a name — a capture for the domain to compute over, not a computation the matcher performs.

Windows bound a stage relative to its parent match: `.Within(entries)` and `.Within(TimeSpan)` (`ReactionEngine.cs:194-220`) close when an event arrives outside the range — no timer, no eviction; the journal closes the window naturally — and `.Aged(TimeSpan)` (`ReactionEngine.cs:225`) is a firing gate that withholds a match until its closing event has settled that span behind the front of the journal. The exact-family quantifiers `None()`, `One()`, and `Exactly(n)` (`ReactionEngine.cs:248-265`) fire when a window closes with the expected count and therefore require a `.Within(...)` — in an open journal there is no other closing point. What the matcher deliberately does *not* admit is any operator that transforms or aggregates the matched data before delivering it: no `GroupBy`, no `Distinct`, no `Sum` or average. That is domain computation and lives in the domain library, not the matcher (§3.3, §6.8) — the matcher correlates a trajectory and hands the raw set over via `Accumulate`. `Within` and `Aged` are not valid on the first `Seek`, which has no parent match to anchor to.

Lifecycle. `SetExpirationDate(DateTime)` (`Reaction.cs:944`) shuts down the Reaction wholesale after a date; `IsExpired` (`Reaction.cs:952`) returns `true` once `DateTime.UtcNow` exceeds the expiration. `SetActive(bool)` (`Reaction.cs:938`) is the manual override. These two differ in scope from `Within(timeout)`: `Within` closes a single window inside a `Many()` stage; `SetExpirationDate` and `SetActive` toggle the entire Reaction. The two scopes are useful for different problems — `Within` lives in the pattern; `SetExpirationDate` lives outside it.

6.5 The three action planes

A Reaction takes exactly one action when its pattern matches. That action is expressed through one of three named *planes*, each describing a different surface of the system the matched trajectory is permitted to affect.

Plane	What the verb touches	What it cannot touch
Program	the actor's domain libraries, read-only	the journal, the actor's state, any global
Causation	the actor's own journal (a journaled cross-actor send)	the domain libraries or the metadata plane
Metadata	the journal's elision register	the domain state or any external effect

A Reaction is therefore not an unbounded scripting surface attached to a pattern. It is a constrained declaration of which surface of the system the matched trajectory is permitted to affect. The builder enforces that only one plane can be configured per Reaction; configuring a second plane after one has already been set is a construction error. The *exactly-one-action* rule is structural, not

conventional — the developer cannot accidentally route a single match into two different planes because the surface refuses the second configuration call at build time.

Program — read-only execution against the domain libraries

The **Program** plane executes a script under the same regime as a query. The parser runs in query mode, the script is forbidden from declaring globals, the **expose** keyword is rejected, the journal is not written, and the runtime takes a read lock that runs in parallel with other queries and emits.

The script may invoke arbitrary external technology — an HTTP client, a Kafka producer, a SignalR hub, a BI sink. What the language denies is any path by which those calls could feed back into the actor's state. This is not a discipline rule kept by convention; it is rejected at parse time, before the developer has the chance to write the leak. A Reaction may observe the actor and affect the outside world, but it cannot feed that effect back into the actor.

The optional **when**: guard performs a second read-only check at firing time:

- the pattern asserts that *this trajectory occurred in history*,
- the guard asserts that *it still makes sense to react now*.

The distinction matters under replay, catch-up after downtime, and replica rehydration — the three situations in which the temporal distance between history and the present can be arbitrary. The pattern is bound to the trajectory; the guard is bound to the clock.

A worked example. Consider a Reaction that, after a confirmed purchase, pushes a real-time toast notification to the customer's UI:

```
seller.Reactions.DefineReaction("NotifyPurchaseToast")
    .Cue().Company().ReadForward()
    .Seek("Purchase")
    .OnMatch("[s:Seller].purchase($orderId, $date, $amount,
↪ $customer)")
    .Program.Emit(
        "notifier.Push(@customer, @orderId);",
        when: "check (notifier.IsFresh(@date)) WARNING 'toast
↪ stale, drop';"
    );
```

The domain verb `Seller.purchase(...)` carries no reference to the notification hub, to UI presence, to toast timing, or to any operational concern of the notification layer. All of that lives inside the Reaction: the call to `notifier.Push(...)` invokes the extrinsic technology, and the guard decides at firing time whether the toast still has an audience. Under the steady path the push fires within seconds of the journal entry; under catch-up or replay, the same match arrives at the matcher hours or days later, the guard returns false, and the emit is sup-

pressed. The freshness policy is plain code in a library class — `ToastNotifier` — that the actor sees through its `Libraries` but that the domain verb never references. The segregation is guaranteed because there is no syntactic route by which this Reaction could import the notifier back into `Seller.purchase`.

Causation — journaled cross-actor continuation

The **Causation** plane is the journaled cross-actor send. A Reaction on this plane writes exactly one sentence into the sender’s own journal instructing another actor; the sender records the causation, the receiver processes the message independently through its own endpoint. Cross-actor coordination therefore appears in the program as a domain fact rather than as infrastructure — neither a saga step nor a tracing span nor a choreography event, but a sentence in the actor’s history.

Its full development — the structural conditions under which it preserves cross-actor program continuity, and the comparison against sagas, choreography, and distributed tracing — is out of scope here. In the present paper it is named only as the second of the three surfaces a Reaction may touch.

Metadata — declaring a trajectory closed

The **Metadata** plane marks the events covered by the pattern as elided. After a trajectory has been correlated, its component events no longer participate in any future decision; subsequent rehydrations walk past them.

This is the conceptual bridge to the *skips* introduced in Paper 1. What appeared there as a journal-density optimisation is here the natural outcome of declarative correlation. The developer is not cleaning the journal after the fact. The developer is declaring “*this trajectory is closed; its component events no longer decide anything.*” The skip is the consequence of that declaration, not its cause.

Parameter injection

`WithParameters(...)` allows captures from the pattern matching to be modified or augmented before the action runs, regardless of which plane is configured. The parameters are pre-populated with the pattern’s captures, so the idiomatic use is to enrich them rather than to construct them from scratch.

Why the planes matter

Without named planes, a Reaction would be an unbounded script runner attached to history — a feature whose semantics would depend on which calls the developer happened to write. With named planes, a Reaction becomes a one-sentence statement about which layer of the system is allowed to change when a historical pattern is recognised. The same primitive that detects the trajectory also declares the surface of effect, and the surface of effect is selected

from a closed set of three rather than from the open set of all calls a script could make.

That restriction is what makes the segregation between

- domain verbs,
- cross-actor causation,
- operational side-effects, and
- journal maintenance

reliable across teams and across time. The reliability is not a property of the developer's discipline. It is a property of the surface the language permits the developer to address.

6.6 Activation: where the Reaction runs

After `Cue()` or `Job()`, the developer chooses an activation scope (`Reaction.cs:68-82`). `DirectorOnly()` runs the Reaction only on the primary replica — appropriate when the work must be unique across the topology: orchestrations that should not be duplicated, and the **globally-singular non-idempotent effects of §3.2** (a confirmation email, a charge to a payment processor) that would otherwise fire once per datacenter. Within a cluster, `DirectorOnly` collapses those N firings to one; across datacenters it does so only if the topology elects a *single* primary rather than one per datacenter (§3.2). It is the activation-scope answer to the duplication hazard §3.2 names — complementary to deduplicating at the sink by the per-match idempotency key (§3.6). `CastOnly()` runs only on followers (secondary replicas), enabling the specialised-workers case named in §6.1: followers that hydrate against the journal alone, hosting `Job` Reactions that the director does not need to evaluate. `Company()` runs on the director and on all followers; appropriate when each replica should hold its own evaluation of the trajectory, perhaps for local cache invalidation or for redundant emit paths.

Activation composes with the `StageManager` — the topology director of the Choreography module. The `StageManager` decides who is director and who is follower at any moment; the Reaction inherits that identity to decide whether to fire. A Reaction defined `CastOnly` becomes silent on the director and active on whatever process the `StageManager` has elected as a follower; a director-to-follower handoff therefore changes the population of running Reactions without any further action by the developer. The deployment-time mechanism that keeps such handoffs safe is `Theater.aliveGate` (the continuation observation); its treatment is out of scope for this paper.

6.7 expose versus `Program.Emit`

The two mechanisms are easy to confuse and are not the same thing. They are addressed in the same direction — feeding downstream systems data the actor

has computed, without forcing those systems to rehydrate the actor — but from opposite sides of the journal.

`expose` is a keyword of the actor’s language, valid only inside a `PerformCommand`. When the verb runs, the script can mark a value as exposed, and the framework persists it into the `ExposeData` of the resulting journal entry alongside the `Action`. It is the *writer’s* externalisation point: the cashier confirming an order can expose the order number, the calculated total, the resolved campaign — and downstream systems can consume those values directly from the journal without having to walk the actor’s state.

`Program.Emit` is the Reaction’s *reader’s* externalisation point. A Reaction running over the journal — whether `Cue` or `Job` — has access to the `ExposeData` of the events it traverses (`Reaction.cs:870`, `includeExposeData=true`), and it can use those values in its `Where` clauses, in its captures, and in the parameters of its emit. What it cannot do is *produce* an `expose`: `PerformEmit` blocks the keyword by parsing with `isQuery:true` (§6.5). The asymmetry is structural and intentional. `expose` is the writer’s hook into the journal; `Program.Emit` is the reader’s hook out of it.

6.8 What does not exist — honest limits

A technical paper that reports what a primitive supports without reporting what it does not is incomplete — but two kinds of absence must be told apart: **deliberate boundaries** that follow from the construct (and would be wrong to add) and **pending gaps** that may simply arrive. The Reaction primitive does not, today, provide:

- **Computation of domain values over the matched events (`Sum`, `Count`, `Average`, arithmetic over payloads) — a deliberate boundary, not a gap.** The matcher’s quantifiers shape the *set* of matches — mark a stage `Many`, require `AtLeast(n)`, bound it with a `Within` window, hand the raw set on with `Accumulate` (§6.4) — which is correlation, the matcher’s own work. They stop deliberately short not only of computing domain values but even of grouping or de-duplicating the set (no `GroupBy`, no `Distinct`): that shaping, too, is domain work over the accumulated capture. This is not a shortfall against CEP (§9.3): it is the anti-porosity boundary of §3.3 and Paper 1 applied to the matcher — value computation is *domain* work and lives in the domain library under its legitimate paradigm, not in the Reaction DSL. The Reaction recognizes a trajectory of historical facts (the journaled invocations, Paper 2) and defers; a decision that needs an aggregate invokes a domain method or decomposes into two Reactions. The cost is real and owned: you cannot fold a running `Sum` into the pattern itself.
- **Negative or absence patterns.** No `Without`, `NotMatch`, `Negate`, `Unless`, `Missing`. “*X happened but Y did not happen afterwards*” is not directly expressible. The conventional workaround is a periodic `Job` that

matches `X` and consults the actor's state via `Program.Emit` with a `when:` check to confirm `Y` has not occurred.

- **Some regex-style quantifiers — partly filled.** The matcher provides `Many()`, `AtLeast(n)`, and the exact family `None()` / `One()` / `Exactly(n)` (the exact family gated on a `.Within(...)` window, §6.4). Still absent as dedicated forms are `AtMost(n)`, `OneOrMore`, and `Optional`.
- **Absolute time windows** as a dedicated method. The relative `.Within(entries)` / `.Within(TimeSpan)` anchors a window to the parent match (§6.4); an absolute calendar window (`start–end`) is expressed instead via `Where("@Now >= ... && @Now <= ...")`.
- **Optional ThenSeek.** A `ThenSeek` cannot be marked optional. Optional branches are written as separate `Reactions`.
- **Job-mode rehydration of the actor.** A `Job` `Reaction` reads only the journal; it does not wake or rehydrate the actor itself. When domain-computed state is needed by the `Reaction`, the canonical mechanism is `expose` (§6.7), which pre-writes the data alongside the command. Silent rehydration inside `Job` is intentionally absent — collapsing this boundary would erase the structural distinction between `Cue` and `Job` (§6.1) and reintroduce the cost the design was avoiding.
- **Cross-actor dispatch with journal-recorded causation — once a gap, now resolved.** An earlier draft of this paper named this as a limitation: cross-actor causation was mediated by `ITransport` infrastructure outside the sender's journal, leaving the broker as the seam. Since then, the framework has gained the third action plane — `Causation.Continue` (the `Tell` primitive) — that records the cross-actor dispatch as a sentence in the sender's own journal while leaving the transport choice to the developer's `ITransport` and the target actor's processing entirely to its own endpoint. The present paper names the plane (§6.5) but does not develop the primitive; its full treatment — design rationale and comparison against sagas, choreography, and distributed tracing — is out of scope here. The bullet is kept as a record of how the gap was named before being filled.

These absences fall into the two kinds named at the top. The value-computation boundary and `Job`-mode non-rehydration are **deliberate**: each follows from a principle the construct rests on — the domain barrier (§3.3) and the `Cue/Job` cost distinction (§6.1) — and adding it would erode that principle. The rest — negative patterns, the remaining regex-style quantifiers (`AtMost`, `OneOrMore`, `Optional`), absolute-time sugar, optional `ThenSeek` — are **pending gaps**: matcher features consistent with the construct that simply are not built yet. A reader who expects a CEP engine's aggregation and windowing operators is importing a different boundary: there the engine computes; here the matcher recognizes a trajectory and the domain computes. The framework's scope is what it supports; what it does not support is either a boundary to respect or a gap to model around.

6.9 Measurement

The argument of this paper is structural, and the instantiation is an existence proof of realizability rather than the design-science *evaluation* of an artifact (§6.10). Two measurements nonetheless calibrate the magnitudes the body asserts. Both are taken against the public commit recorded under Code provenance (6a529c4), with the harness in `labs/lab03-reactions`. Environment: Intel Core i9-13900, .NET 9.0.14 (X64 RyuJIT), Windows 11, Workstation GC; Release builds; timings via BenchmarkDotNet with tiered compilation disabled (`DOTNET_TieredCompilation=0`). Figures are read to about one significant figure.

Verb overhead. A minimal in-memory verb — one field write, no host work — dispatches and persists in $\approx 0.4\ \mu\text{s}$. This is the actor’s own command-path overhead, the floor beneath every verb; a verb’s observed latency is therefore set by the domain methods it invokes, not by the runtime around them (§5.2). The figure is sub-microsecond, orders of magnitude under any millisecond budget, which is why the fast verb is structural rather than tuned.

Cue reaction latency. End to end — from a command being issued to a Cue Reaction’s action firing — latency was a median of $\approx 1.3\ \text{ms}$ (p99 $\approx 2\ \text{ms}$) over 2000 activations, the Reaction firing once per match with no duplicate (exactly-once held, zero double-fires). The run was crash-free, so it exercises the common path, not the at-least-once crash window of §3.6; what it shows is steady-state delivery and latency, not a refutation of that window. The figure is well under the sub-second bar §6.1 claims, and it is governed by the signal-driven push path: a newly journaled entry preempts the catch-up poll because that poll waits on `pushSignal (ActorReactions)`, the same signal the `RecordWritten` callback sets, rather than waiting out a fixed `Thread.Sleep` backoff. That the latency is a property of scheduling and never a floor of the construct is visible in the contrast: a naive fixed-backoff poll ($50\ \text{ms} \rightarrow 1\ \text{s}$) measures $\approx 0.13\ \text{s}$ on the identical path — a $\approx 100\times$ difference from a change that touches scheduling alone, leaving delivery and exactly-once untouched. The construct never bounded the latency.

Threats to validity. The timings are single-environment and machine-specific; only their order of magnitude is portable, and they are reported to one significant figure. No external-system baseline is offered: a latency comparison against another actor runtime would confound language, persistence model, and transport at once, and the comparison this paper rests on is structural, not chronometric. The load-bearing observations are the deterministic ones — the verb-overhead floor and exactly-once-per-match (the monotonic cursor of §3.6, `ActorReactions.cs:186`), build-independent by construction — while the timings only confirm the direction the structural argument predicts.

Adversarial case. The partition does not always pay. When the deferred effect is cheaper than the machinery that defers it — the thread handoff, the pattern match, the separate emit under a read lock — moving it into a Reaction lowers the verb’s latency but raises the total work the system does per event. The

construct buys a fast verb and a closed library; for a trivial effect on a low-traffic verb it can cost more than it saves, and inlining is then the honest choice. The partition earns its keep where the deferred branch carries real operational weight — the case it was designed for, and the case the *When NOT to use* regimes (A–C) bound from the other side.

6.10 Evaluation deferral

The measurements of §6.9 are bounded observations scoped to the structural claims of §5 and §6, not a comprehensive evaluation of the instantiation. Holistic operational performance — endpoint-latency distributions under production load, developer-velocity comparisons, behaviour across a cluster — is not measured here. Under the analytic frame (§1) the deferral is supported: the construct stands on the structural argument, with empirical evaluation supplementary. The honest qualification is that the series has so far carried this empirical debt forward rather than redeeming it; the comprehensive case-study is owed, and naming the debt here is part of not letting the analytic papers accumulate it indefinitely.

7. Comparison with Akka and Orleans

7.1 Akka actors and Become/Receive

Akka’s actors model behavior change with **Become** and **Unbecome**, swapping the receive partial function as the actor moves between modes. A booking actor might **Become(opened)**, then **Become(confirmed)**, then **Become(paid)**; each receive defines a different set of messages the actor accepts and how it handles them. The partition this expresses is across *actor states* — what the actor will accept next changes with the state — and the developer’s mental model is a state machine the actor walks through.

Reactions express a different partition. The actor’s **PerformCommand** is one verb, not many — it does not **Become** between calls — and the partition that Reactions add is across the *event boundary*: when a command has been processed and persisted, what happens afterwards is the Reaction’s domain. Akka’s mechanism is suited to actors whose interface changes as state evolves; Puppeteer’s mechanism is suited to actors whose interface is stable but whose effects after each verb are rich. Neither subsumes the other; they answer different questions.

7.2 Orleans grains, timers, and reminders

Orleans’ grains support timers (in-process, not durable across grain reactivation) and reminders (durable, scheduled). Both share the placement intent of **Job**: deferred work, possibly on another silo, scheduled rather than co-located. The structural differentiator is what gets *persisted*. An Orleans timer or reminder

is a runtime registration — scheduled in the grain’s activation lifecycle and re-registered on reactivation — that is not itself part of the grain’s durable state. A Reaction is part of the actor’s journaled program: declared in the DSL, parsed against the domain library (§6.2), and reconstructed from the journal on every rehydration (§3.5). Whether the deferral mechanism is recorded in the persisted program or attached to the runtime is an inspectable difference, and it is the one this section rests on. A separate, *ergonomic* observation — that declaring a Reaction is a few lines of fluent DSL beside the verb (§4) where the Orleans equivalent is procedural — is a design-rationale remark about friction, not a measured result: this paper runs no usability study and makes no quantitative “fewer lines” claim.

7.3 Where the design spaces converge and diverge

Akka, Orleans, and Puppeteer address the same structural pressures: the actor’s serial mailbox, the cost of state hydration, the cost of distributing work across nodes, the difficulty of keeping fault isolation honest. The three answer those pressures with different design commitments. Akka commits to a flexible behavior model in which the actor’s interface itself can evolve; Orleans commits to virtual actors that materialise on demand, with placement and scheduling as runtime concerns. Puppeteer commits to an externalised DSL with a homoiconic journal (Paper 2) and an anti-porous domain (Paper 1) — and Reactions are the language layer of that commitment: declarative pattern matching over the actor’s trajectory (§3.5), with a strict separation between the actor’s command path and what runs afterwards (§5.2). The Puppeteer position is one point in a non-trivial design space, not its peak.

Nothing in the present paper claims Reactions are the only realization of the construct. An Akka or Orleans extension that exposed a declarative DSL for trajectory matching with build-time binding to the host language’s types would satisfy the same construct; this paper presents one realization, not the realization.

8. Counter-arguments

8.1 “Eventual consistency is a known anti-pattern in business systems”

The objection treats eventual consistency as a uniform property of the system, taken or rejected wholesale, and reads the warning literature on it as applying everywhere it appears. The framing is correct for *opt-out* eventual consistency, the kind classical CQRS introduces by separating read and write stores: every read is potentially stale, every write is potentially unobserved, and the application has to recover stronger contracts wherever the business demands them. Under that framing, the warning is well-earned.

The model developed in this paper inverts the framing (§5.3). Reads and writes inside the actor are strongly consistent — **PerformCmd** and **PerformQry** operate on the same in-memory state. Eventual consistency appears only at the deferred branch, named by the developer one deferral artifact at a time, with a placement attached. There is no system-wide consistency contract that the application must compensate for; there is, instead, a list of named windows the developer has chosen to admit. The literature warning applies to the framework that imposes eventual consistency as a default; a reader of an instantiation that exercises this construct can answer “what is eventually consistent here?” by reading the deferral artifacts declared near the verb and nothing else.

The honest concession (the 20%): some business domains require that every observable effect of a verb commit in lockstep with the response — banking core ledgers, certain regulated transaction systems. The opt-in framing does not soften this; in those domains, the developer cannot extend a license that does not exist, and the partition’s primitive is not the right primitive (§When NOT, A). The objection lands where it lands; it does not generalise to every business system that adopts the framework.

8.2 “Reactions are async events with a different name”

The objection holds only if Reactions are read at their surface — as something that runs after a verb, on a different thread, with a callback shape. At that surface they look like async events, and the objection follows.

Read at the structural level of the construct, the artifact realizing the partition is something else. An async event handler responds to a single message; the construct’s matching primitive matches a *trajectory* — a sequence of one or more stages naming a saga, a lifecycle, or an anomaly through the actor’s history (claim 2(c), §6.3). The pattern is parsed against the domain’s libraries at build time, not against text at runtime: a reference to a class the domain does not contain fails to construct (§6.2). Identifiers captured in one stage are typed and propagated to the next, which is what correlates the events into a story — an async event handler has no analogue. The placement axis with two canonical points (in Puppeteer’s instantiation, **Cue** and **Job**) carries activation scopes (**DirectorOnly** / **CastOnly** / **Company**) that compose with the StageManager’s topology (§6.6); an async event has none of that surface. The action plane choices (**Metadata.Elide**, read-only **Program.Emit** with optional **when:** check, and the journaled cross-actor **Causation.Continue**) are constrained primitives — particularly the read-only **Program.Emit** whose underlying **PerformEmit** structurally blocks **expose** and journal writes (§6.5) — designed to enforce the categorical segregation of §3.3 in the parse phase, not at runtime.

The honest concession (the 20%): the surface mechanism — work happens after the verb returns, on a separate thread or on demand — is shared with async event systems. A reader who looks only at the trigger sees something familiar. The argument of this paper is that the trigger is the least interesting part of the

primitive.

8.3 “Even with low-friction Reactions, the developer can pollute the domain”

The objection is correct in spirit: friction reduction is not a substitute for discipline. A determined developer can put `sendEmail(...)` inside a `PerformCommand` even when a `Reaction` one line below would do the job; the framework cannot prevent it. Were the argument that friction reduction *replaces* developer judgment, the objection would land.

The argument is different. Friction reduction is the lever the framework can pull; developer judgment is the lever it cannot. When friction is high — when the right thing costs more than the wrong thing — no amount of training, code review, or architectural pep-talk holds at scale. When friction is low, training and review have something to lean on; the developer who chooses inline anyway is doing so against an obvious cheaper alternative, and the choice becomes legible to reviewers as a deliberate departure rather than a path of least resistance.

The honest concession is at the seam where friction reduction reaches its limit: a careless team will write polluted code under any framework, and a disciplined team will write clean code under any framework. The framework’s contribution is the *gradient* between these — making clean code the default, pollution the exception, and the cost of getting it wrong recoverable rather than irreversible. The objection becomes an argument for friction reduction, not against it: in a world where developer discipline is unreliable, the only thing the architecture can do is shape the cost surface so that the path of least resistance is also the path of least damage.

8.4 “DSL scripts repeat lines across endpoints — this violates DRY”

The objection assumes DRY as a textual prohibition. In its original formulation (Hunt & Thomas, 1999), DRY prohibits the duplicate representation of *knowledge*, not the appearance of similar lines: “*every piece of knowledge must have a single, unambiguous, authoritative representation within a system.*” The domain library lives once, in OOP; DSL scripts are the language in which an actor is invoked, structurally analogous to SQL. Two SQL queries that share JOIN patterns are not in violation of DRY — they are compositions over a shared schema. The same is true of two DSL endpoints that share invocation lines: the knowledge is in the library that the lines invoke, not in the lines themselves.

Concession (the 20%): if two DSL endpoints share *domain logic* — not invocation patterns but actual computation — that logic should refactor into a method of the library. The line is between invocation and implementation, not between unique and repeated text.

8.5 “Reactions should be able to rehydrate the actor for richer state queries”

The choice between modes is the structural distinction itself — full state queries belong to the live-actor mode, not to the journal-driven one (§6.1). When the data is anticipated by the developer, the journal pre-write mechanism (**expose** in Puppeteer, §6.7) puts the data where the deferral artifact will consume it during its own rehydration, without ever waking the actor. The journal-only-without-rehydration boundary is what makes the journal-driven mode cheap to host on a remote follower; lifting it would collapse the two modes into one fuzzy mode and reintroduce the cost the original design was trying to avoid.

Concession: there are rare cases of data that no one anticipated would be needed by a future Reaction, and that already exist in many past events without an **expose**. The workaround is operational — add **expose** now, run a migration that replays old events to write the **expose** retroactively, or build an auxiliary **Job** that materialises a side projection. Costly, but solvable, and rare enough that it does not justify reopening the primitive.

9. Related work

9.1 Actor Model and CQRS foundations

The actor model originates with Hewitt, Bishop, and Steiger (1973), which introduces the actor as the universal unit of computation: an entity with private state, a mailbox for incoming messages, and a single thread of behavior responding to them. The model has accumulated decades of refinement — Agha (1986), Hewitt (2010), the *Reactive Manifesto* (Bonér, Farley, Kuhn, & Thompson, 2014) — but the core remains: state isolated, computation message-driven, concurrency through the multiplication of actors rather than the sharing of state. Puppeteer’s actor sits in this lineage, with the addition that the journal of received messages is itself the persisted artifact rather than a derived log.

Command Query Responsibility Segregation (CQRS) is named in Young (2010), with intellectual antecedents in Meyer’s command-query separation (1988) and Fowler’s reflections on enterprise architecture (2002). The classical formulation separates write and read stores at the level of the system architecture, with eventual consistency between them as the design contract. This paper engages CQRS not as the primary frame but as the literature that names the consequences of the partition this paper develops; the “logical CQRS” supporting observation is the classical pattern read at a different level — one in-memory state, two verbs over it, the eventual-consistency contract appearing only at the deferred branch (§5.3, §8.1).

9.2 DDD and the anti-pattern literature

Domain-driven design originates with Evans (2003), where the *anemic domain* and *smart UI* antipatterns are diagnosed: domain types reduced to data containers with behavior emigrated to service layers, and business rules pulled into the presentation/transport edge. Vernon (2013) catalogs the implementational pitfalls of DDD-by-vocabulary-only — bounded contexts and aggregates named but never structurally enforced. Fowler (2003, in *AnemicDomainModel*) gives the antipattern its most cited articulation. These three sources name the industrial reality §5.1 of this paper describes: in most enterprises, “the domain” does not exist as a library — it exists as scattered classes whose names suggest cohesion and whose substance does not.

Fowler’s *Patterns of Enterprise Application Architecture* (2002) is the relevant catalog for the DTO antipattern: a Data Transfer Object intended for transport that becomes load-bearing in the domain, contaminating both. The same volume names the *Distributed Object* antipattern — the chatter problem this paper alludes to in §3.1 when it argues for the SQL-analogous use of a DSL as invocation rather than transcription.

Hunt and Thomas (1999) define DRY in *The Pragmatic Programmer* as the prohibition of duplicate *representation of knowledge*, not of repeated text — the formulation §3.1 and §8.4 recover against a more recent industrial reading. Sandi Metz’s *The Wrong Abstraction* (2016) provides the natural complement: the cure for repeated text is sometimes more text, not the wrong abstraction.

9.3 CEP and adjacent runtimes

Complex event processing has a substantial literature — Luckham’s *The Power of Events* (2002), Etzion and Niblett’s *Event Processing in Action* (2010) — concerned with detecting patterns over streams of events, often from heterogeneous sources, often without the type system of the consuming application available. Correlated multi-event pattern matching with bound, propagated variables is CEP’s defining capability, and typed, host-integrated engines — such as Esper, Flink CEP, and Drools Fusion — already compile or validate patterns against application types. This paper claims no novelty for the matching mechanism: the Reaction matcher is CEP-style matching, and these engines are its closest prior art. What differs is the substrate. CEP matches over external event streams; a Reaction matches over the actor’s own *journalled invocations* (the executable journal of Papers 1 and 2), so the matched units are domain-method calls resolved against the domain library at build time (§6.2) and the matcher, the journal, and the domain share one type universe. That is a difference of setting and integration, not of matching power — and the paper’s contribution is the partition construct (§1–§5) the mechanism serves, not the mechanism itself. CEP engines also *compute* over what they match — aggregations, windows; a Reaction’s matcher does not. It shapes the match set (group, de-duplicate, bound) but leaves value computation to the domain (§3.3, §6.8) — again a

placement of the boundary, not a weaker engine.

Akka and Orleans are the most relevant adjacent runtimes (§7). Their documentation is the canonical reference for Akka’s **Become/Receive** and Orleans’ grain timers, reminders, and stream APIs — the comparators against which §7 contrasts the Reaction primitive without claiming superiority.

9.4 Prior papers in this series

This paper builds on two prior contributions. *Anti-porous architecture: a unified design principle for CQRS + Actor + Event-Sourcing systems* (Paper 1) names the structural defect — porosity — that the present paper extends to a fourth layer (the operational edge of the domain, §3.3) and the legitimate tool — the domain library shaped by its operations — that §3.1 carries forward. *Program-value separability: the structural precondition for compilation, caching, and dense journaling in a DSL runtime* (Paper 2) names the structural property of DSL programs that makes their compilation, caching, and dense journaling possible at all; the static binding of Reaction patterns against the domain’s **Libraries** (§6.2) is one more layer of language built atop that substrate. Both prior papers are self-deposited preprints on Zenodo (Rivera, 2026a, 2026b) and have not undergone peer review, as is the present paper.

10. Conclusion: a developer says things — and now also at the endpoint

In §1 a developer was modeling a purchase-order endpoint by saying things. “*The cashier confirms the purchase and locks the requested items. Then a confirmation goes out, and the rewarder evaluates whether any promotional campaign applies.*” That sentence was the program before it was code. Each phrase in it carries one of the technical names this paper has used.

- “*The cashier confirms the purchase and locks the requested items*” — a verb on a single in-memory state shared by reads and writes; the literature calls this *logical CQRS* (a supporting observation of this paper).
- “*Then a confirmation goes out, and the rewarder evaluates...*” — the partition between *now* and *deferred*, with placement chosen at the moment of partitioning (claim 1). Each phrase after the first is itself a Reaction declared next to the verb.
- “*A confirmation*” lands in a Reaction rather than in a domain method because the email concern does not belong inside the verb the cashier exposes — categorical segregation (claim 2(b)), encapsulating the technology where it can be exercised without propagating into the rest of the domain.
- “*The cashier, the order, the rewarder*” — these live inside a domain library that closes its operational boundary while continuing to grow conceptually under its legitimate paradigm (claim 4).

The developer reached none of these technical names while modeling. They reached for a sentence about who does what. The names are downstream — useful for reasoning about the system, for defending the design, for situating it against Akka and Orleans and the broader CQRS literature. The sentence is upstream — useful for getting the work done.

At some point a domain must be declarable as “task done”. A library whose definition must absorb every present and future operational tool — email today, Slack tomorrow, a webhook protocol the year after — is a library that never closes. The only way it closes is if it stays pure: free of operational concerns by construction. Reactions are the artifact that makes that closure possible.

One of those downstream names carries more weight than the others. The CQRS literature treats eventual consistency as the architecture’s defining commitment — taken or refused wholesale, a property the application must then absorb wherever the business demands stronger contracts. Read through the partition, it is none of that: it is the shadow the partition casts, present exactly where the developer drew a deferral and nowhere else. That is the inversion this paper turns on — eventual consistency is a consequence of a modeling decision, not a property of the architecture.

In C# a developer says things. In the DSL of an actor-native runtime, a developer says things too — about endpoints, about who does what, about what gets handled afterwards. The act is the same; the surface is bigger. What is new is not how the developer thinks. It is what the language lets them say.

Appendix A. Code references

All references are to the Puppeteer codebase; the files below live under `Puppeteer/EventSourcing/` (the matcher files under `Puppeteer/EventSourcing/Follower/`). Path:line citations were verified against the anchored commit (2026-06-21). Releases after this date may shift line numbers; the symbolic references (method names, class names) are stable and locate the same artifact regardless of line drift.

Reaction primitive — `Reaction.cs`

Range	What it shows	Cited in
20-24	<code>ReactionMode</code> enum (<code>Cue</code> , <code>Job</code>)	§6.1
68-82	Activation builders (<code>DirectorOnly</code> , <code>CastOnly</code> , <code>Company</code>)	§6.6

Range	What it shows	Cited in
810	<code>time</code> global registered as <code>Temporal</code> instance	§6.4
863	<code>Cue</code> mode push-callback registration via <code>AddRecordWrittenCallback</code>	§6.1
870	<code>includeExposeData=true</code> during rehydration	§6.7
938	<code>SetActive(bool)</code> manual override	§6.4
944	<code>SetExpirationDate(DateTime)</code> shutdown	§6.4
952	<code>IsExpired</code> predicate	§6.4
981-983	The three action planes exposed as properties (<code>Program</code> , <code>Causation</code> , <code>Metadata</code>)	§6.5
1039	<code>SetMetadataAction()</code> internal hook for <code>Metadata.Elide()</code>	§6.5
1076	<code>EnsureNoActionConfigured()</code> build-time guard (exactly-one-action rule)	§6.5
1106	<code>WithParameters</code> parameter modifier	§6.5
1113-1255	Action handlers (<code>ExecuteAction</code> , <code>ExecuteProgram</code> , <code>ExecuteCausation</code>)	§5.2, §6.5
1256-1267	<code>ExecuteMetadata()</code> with the <code>MarkEventsAsElided</code> invocation (Metadata plane)	§6.5
1174	<code>PerformEmit</code> invocation from <code>ExecuteProgram</code>	§6.5
1444	<code>SolveActionReferences()</code> per-event resolution	§6.2

Reaction action planes — `Planes.cs`

Range	What it shows	Cited in
5-13	Header comment naming the three planes (Program , Causation , Me tadata) and their verbs	§6.5
44	Program.Emit(string script) — read-only emit	§6.5
53	Program.Emit(string script, string when) — read-only emit with check	§6.5
76	Causation.Continue(s tring script) — journalled cross-actor verb (Tell)	§6.5
142	Metadata.Elide() — mark matched entries as elided	§6.5

Reaction language layer — **ReactionEngine.cs**

Range	What it shows	Cited in
71-83	Where(expression) clause registration with build-time validation	§6.4
146-265	Quantifiers and windows (Many, Accumul ate, AtLeast, Within, A ged, None, One, Exactl y)	§6.4

Read-only emit path — **ActorHandler.cs**

Range	What it shows	Cited in
2408-2479	PerformEmit method (read-only, isQuery:tru e , read lock)	§3.3, §5.2, §6.5, §6.7

Range	What it shows	Cited in
2405	In-source comment marking the query-mode parse: <code>isQuery: true</code> blocks <code>expose</code> (which would persist to the journal) and global-variable declaration	§6.5
2423	<code>parser.Parse(isQuery: true, isCheck: false)</code>	§6.5
2464	<code>rwLock.EnterReadLock()</code> taken before <code>Perform</code>	§6.5

Other modules (referenced without line numbers)

File	What it shows	Cited in
Puppeteer/EventSourcing/Follower/Pattern.cs	Static parsing of pattern descriptions against <code>Libraries</code>	§6.2
Puppeteer/EventSourcing/Follower/MatchTree.cs	<code>TryMatchAtLevel</code> , BFS/DFS hydration walk	§6.2
Puppeteer/EventSourcing/Follower/MatchTree.cs:1565-1575	confirmed checkpoint saved only after <code>ExecuteAction</code> succeeds (at-least-once)	§3.6
Puppeteer/EventSourcing/Follower/MatchTree.cs:1512	<code>MarkEventsAsElidedWithCheckpoint</code> — elision + cursor in one atomic commit	§3.6, §6.5
Puppeteer/EventSourcing/DB/DiaryStorage.EventElisionStorage	<code>MarkEventsAsElided</code> storage interface	§6.5

Code provenance

Source-code references in this paper resolve against the public Puppeteer repository at commit [6a529c4](#) (2026-06-21). The snapshot is archived in Software Heritage under the following persistent identifier:

```
swh:1:dir:463f291ec5e4fa8abef24dc1fe3e46e5a5975739;  
  origin=https://github.com/alvaroNCubo/puppeteer;  
  anchor=swh:1:rev:6a529c45367ca2c32a891a067efb5be3fed589ec
```

Inline references of the form `file.cs:NN` (e.g., `ActorHandler.cs:38`) resolve against this snapshot. A reader can construct a per-file SWHID by adding the qualifiers `;path=<path>;lines=<NN>` to the directory SWHID above. Future commits to the repository may renumber lines; the SWHID preserves the cited state independently of any future change to the repository or its hosting.

Acknowledgments

The author used large language models (including Claude and ChatGPT) as editorial assistants for language refinement, structural feedback, and literature navigation. All original ideas, terminology, theoretical constructs, and technical content presented in this work are solely the author's.

Appendix B. Bibliography

Agha, G. (1986). *Actors: A model of concurrent computation in distributed systems*. MIT Press.

Akka documentation. (n.d.). <https://akka.io/docs/>

Apache Flink CEP documentation. (n.d.). <https://nightlies.apache.org/flink/flink-docs-stable/docs/libs/cep/>

Bonér, J., Farley, D., Kuhn, R., & Thompson, M. (2014). *The reactive manifesto* (v2.0). <https://www.reactivemanifesto.org/>

Drools Fusion documentation (complex event processing). (n.d.). <https://www.drools.org/>

Esper reference documentation. (n.d.). <https://www.espertech.com/esper/>

Etzion, O., & Niblett, P. (2010). *Event processing in action*. Manning Publications.

Evans, E. (2003). *Domain-driven design: Tackling complexity in the heart of software*. Addison-Wesley.

Fowler, M. (2002). *Patterns of enterprise application architecture*. Addison-Wesley.

Fowler, M. (2003). *Anemic domain model*. [martinfowler.com. https://martinfowler.com/bliki/AnemicDomainModel.html](https://martinfowler.com/bliki/AnemicDomainModel.html)

- Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.
- Hevner, A. R., March, S. T., Park, J., & Ram, S. (2004). Design science in information systems research. *MIS Quarterly*, 28(1), 75–105.
- Hewitt, C., Bishop, P., & Steiger, R. (1973). A universal modular actor formalism for artificial intelligence. In *Proceedings of the 3rd International Joint Conference on Artificial Intelligence (IJCAI-73)* (pp. 235–245).
- Hewitt, C. (2010). *Actor model of computation: Scalable robust information systems*. arXiv. <https://arxiv.org/abs/1008.1459>
- Hunt, A., & Thomas, D. (1999). *The pragmatic programmer: From journeyman to master*. Addison-Wesley.
- Luckham, D. (2002). *The power of events: An introduction to complex event processing in distributed enterprise systems*. Addison-Wesley.
- Metz, S. (2016). *The wrong abstraction*. sandimetz.com. <https://sandimetz.com/blog/2016/1/20/the-wrong-abstraction>
- Meyer, B. (1988). *Object-oriented software construction*. Prentice Hall.
- Microsoft Orleans documentation*. (n.d.). <https://learn.microsoft.com/en-us/dotnet/orleans/>
- Rivera, A. (2026a). Anti-porous architecture: a unified design principle for CQRS + Actor + Event-Sourcing systems. *Puppeteer Papers Series*, Paper 1 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20404863>
- Rivera, A. (2026b). Program–value separability: the structural precondition for compilation, caching, and dense journaling in a DSL runtime. *Puppeteer Papers Series*, Paper 2 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20740697>
- Vernon, V. (2013). *Implementing domain-driven design*. Addison-Wesley.
- Young, G. (2010). *CQRS documents*. https://cqrs.files.wordpress.com/2010/11/cqrs_documents.pdf

Preserving semantic continuity across actors

TL;DR

The actor model treats cross-actor causation as operational — message dispatch by a runtime, not a statement in any actor’s program. The treatment is not entailed by the model; it is a contingent design decision, adopted as the default frame of the canonical actor lineage and seldom examined as a choice. The cost is structural: an entire ecosystem of compensating patterns — sagas, choreography, distributed tracing, workflow engines — exists to reconstruct cross-actor causal chains that no actor’s program records. The treatment even erases a grammatical distinction — one dispatch verb collapsing an assertion of fact (past tense) and a command (present imperative) into a single send — which the programmatic alternative restores.

This paper observes that the assumption can be removed without altering any structural commitment of the actor model. Three conditions describe a system in which the cross-actor send is recorded as a sentence in the sender’s program: locality of writes (each journal records only its own actor’s activity), causation as program statement (the cross-actor send is a sentence in the sender’s journal), and no external coordinator (no party outside the participating actors decides what happens next). Any primitive satisfying these three conditions dissolves the assumption — on the dense journal-as-program substrate the conditions presuppose (the send recorded as a sentence, not a payload). *Tell* — a primitive in the Puppeteer framework — is one such realization.

Three implementations of the same domain — saga, choreography, tell — are exhibited side by side. The reader sees, in journals, that the saga places the joint history in a coordinator’s program, that the choreography places it in an external bus log, and that the tell instantiation places it in the sender’s own program. Four tests probe

these claims under tell. Three show that the journal alone supports replay reconstruction, cross-datacenter replication, and audit query — operations that require external apparatus under the assumption. The fourth shows that after a crash in the dispatch window the journal still records a stranded tell’s fate honestly, recovered from the transport’s testimony rather than fabricated. Each cross-actor edge is recorded as program in its sender’s journal; a multi-hop chain composes these per-edge records — programmatic at every hop, where the alternatives reconstruct the chain from artifacts outside any program (§8.5). **Fifty years of convention are not fifty years of necessity.**

Claims this paper makes

1. **The assumption named.** Cross-actor causation has been treated as operational rather than programmatic across fifty years of canonical actor-systems literature — from Hewitt et al. (1973) through Akka and Orleans. (*Verification: §2 + the genealogy table in §2.4.*)
2. **The assumption is contingent.** No theorem of the actor model entails it. The three structural commitments of the model — autonomy, message-based communication, isolation — define the ontology of actors, not the mechanics of their runtimes. (*Verification: §4.*)
3. **Three conditions resolve the assumption.** Locality of writes (C1), causation as program statement (C2), and no external coordinator (C3) describe a system in which semantic continuity is preserved across actors without violating any of the actor model’s commitments. The conditions are not design choices; they are the consequences of reconciling the model’s commitments with the construct. (*Verification: §7.1.*)
4. **Any primitive satisfying C1+C2+C3 dissolves the assumption — given a journal-as-program substrate.** The realization is one of many possible, but the conditions are not free-floating: C2 requires the send to be recorded as a *statement of the program* — a dense operation, not a serialized event — which presupposes the anti-porosity of Paper 1 and the externalized parameters of Paper 2. The claim is therefore *tell on such a substrate*, not *tell in any actor framework*; a framework that records sends as opaque payloads would satisfy C2 only by first adopting that substrate, which is the subject of Papers 1–2 and beyond this paper’s scope. (*Verification: §7.5.*)
5. **The actor model’s commitments survive the reformulation intact.** Autonomy, message-based communication, and isolation each remain unchanged; only the historical interpretation of what each actor’s

program is *permitted to say* is removed. (*Verification: §7.3.*)

6. **The compensating ecosystem exists because of the assumption.** Saga orchestrators, event-driven choreography, distributed tracing, and workflow engines each compensate, in their own way, for the absence of cross-actor causation in any actor’s program. Their sophistication, maturity, and widespread adoption are evidence of the cost of the assumption, not of its correctness. The scope is the *record* of cross-actor causation: tell relocates that record into the program; it does not provide the transactional recovery — compensation, rollback, timeouts — that sagas and workflow engines also carry, and does not claim to (§6.1). (*Verification: §6 + the side-by-side labs in §8.3.*)
7. **Tell exhibits the conditions in a runnable instantiation.** A Reaction whose `.Causation.Continue(...)` body issues a `tell` records the cross-actor assertion in the sender’s journal — as a typed message-action — and the receiver’s acknowledgment closes the round-trip. The sender’s journal becomes a self-contained record of that cross-actor edge; a multi-hop chain composes these per-edge records — programmatic at every hop, unlike the alternatives’ reconstruction from outside any program (§8.5). (*Verification: §8.2 + §8.3 Style 3 + §8.5.*)
8. **Auditing the cross-actor narrative becomes reading the program.** The journal shows what was sent, to whom, with what content, at what time, and with what acknowledgment — without correlation IDs, distributed tracing, or external aggregation. (*Verification: §8.5 G3.*)
9. **Replay reconstructs cross-actor state from the journal alone — and recovers each tell’s fate.** A fresh actor with no shared transport and no live receiver, replaying the journal, reconstructs the in-flight tell state; for a tell stranded by a crash in the dispatch window, the rehydrated actor reconstructs the pending tell and the transport testifies its fate, which the journal records as a verdict — acked, not-delivered, or honestly pending — so recovery never leaves the journal asserting a send that did not land. (*Verification: §8.5 G1, G4.*)
10. **Cross-datacenter replication preserves the cross-actor causal chain.** The journal carries the cross-actor causation across data centers because the causation was always recorded in a place replication can carry. (*Verification: §8.5 G2.*)

1. Introduction

When one actor causes another to act, where does that causal sentence live? Consider the ordinary case. An actor performs an operation and, as a result, sends a message to another actor. The first actor’s operation, state mutation,

and emitted events appear in its journal or trace; the second receives the message, and its own operation is likewise recorded. Yet the act of sending — the causal step that connects the two — appears in neither actor’s program. It is mediated by the runtime, the dispatcher, or the message broker: the sender’s journal does not record that it spoke, and the receiver’s journal does not record, in program terms, who spoke to it. Ask later why the second actor did what it did, and the answer has to be assembled from outside every program — partly from the first actor’s log, partly from the second’s, partly from the broker’s offsets, partly from a distributed trace. The causal chain exists. No program contains it.

This is so familiar that it rarely appears noteworthy. But it should. If actors are programs, and a sentence in one actor’s program causes an effect in another, then that causal sentence has every reason to appear as part of the first actor’s program. That it does not is no logical consequence of the actor model; it is a structural feature of how actor systems have historically been constructed.

This paper names that gap. The construct introduced is *semantic continuity*: the property of a program (in the substrate-level sense established in Paper 2 §1.2: the pair of domain library and journal of invocations) whose causal structure remains recorded as part of the program itself, even when its effects cross boundaries. The defect identified is the absence of semantic continuity at actor boundaries. The principles derived are the conditions under which it can be preserved without violating actor isolation and without introducing orchestration. The instantiation presented is *tell*, a primitive in which the cross-actor send is recorded as a sentence of the sender’s journal — a dense program operation, not a serialized payload, which presupposes the anti-porous journal established in Paper 1.

Methodologically, this is an analytic theory contribution: it names a structural assumption that the canonical literature on actor-based systems has documented in different forms without recognizing as a single construct, derives the principles under which the assumption can be rejected, and presents a running system in which those principles are realized as an existence proof — not as the substance of the claim. The weight of the argument is conceptual: the naming of the assumption (§3) and the account of why an entire ecosystem of patterns answers it (§6) carry the claim, while the instantiation (§8) is its existence proof. The genre is the one Gregor (2006) names *theory for analyzing* (Type I): it introduces a construct that lets the phenomenon be described and classified, with empirical evaluation supplementary; the Hevner-style design-science *evaluation* of the artifact (Hevner, March, Park, & Ram, 2004) is a separate undertaking this paper does not attempt. The conditions C1–C3 derived in §7.1 are the necessary conditions the construct entails, not a prescriptive method offered for adoption: deriving what a system must satisfy to preserve the construct is Type I analysis, where prescribing and evaluating an artifact would be a separate design-science (Type V) contribution this paper does not undertake. This is not a systems paper — it presents no performance benchmarks, fault-injection

metrics, or latency comparisons against existing actor frameworks; analytic theory measures contribution by the precision of the construct, the validity of the principles, and the realizability of the instantiation. Readers expecting quantitative comparisons against alternatives will find structural comparisons — what each pattern records, where the joint history lives — in §6 and §8.4.

§2 traces the genealogy of the assumption that produces this gap, showing that across five decades and multiple canonical generations of literature the assumption has remained continuous in substance though varied in form. §3 names the assumption explicitly: *causation between actors is treated as operational rather than programmatic*. §4 demonstrates that this assumption is contingent rather than necessary — no theorem of the actor model entails it. §5 examines the architectural and operational consequences of the assumption. §6 shows why existing responses to those consequences — sagas, choreography, distributed tracing, and workflow engines — cannot dissolve the assumption, because they reconstruct cross-actor flow after the fact rather than preserving it as program. §7 reformulates the model: under three explicit conditions, semantic continuity can be preserved across actors. §8 presents the instantiation, *tell*, through an illustrative case study against saga and choreography implementations of the same domain. §9 relates the construct to prior work in this paper series. §10 concludes.

2. Genealogy

2.1 The founding decision (1970s)

The actor model was introduced in 1973 by Hewitt, Bishop, and Steiger as a unifying account of concurrent computation in which “all of the modes of behavior can be defined in terms of one kind of behavior: sending messages to actors” (Hewitt et al., 1973, p. 235). Sending messages is positioned as a universal primitive, but the framing goes further: it is positioned as a *machine-level* primitive. “The basic unit of execution on an actor machine is sending a message in much the same way that the basic unit of execution on present day machines is an instruction” (Hewitt et al., 1973, pp. 236–237). The act of sending is therefore parallel to a hardware instruction — it is what the abstract machine performs *between* actors, not content of any actor’s program.

A different concurrency tradition emerged in 1978 with Hoare’s *Communicating Sequential Processes* (Hoare, 1978). It is worth distinguishing rather than assimilating: CSP’s communication is a *symmetric, synchronous* handshake at named ports — sender and receiver rendezvous, each blocking until the other is ready — which is the opposite of the actor model’s *asynchronous, fire-and-forget* send. These are different ontologies, and this paper does not fold them together. CSP earns one narrow point of contrast: there too, cross-process communication is a construct of the language, not a statement in either process’s program. But the

genealogy below tracks the actor lineage specifically; CSP is a foil, not a second witness to the same assumption.

The actor model was formalized in 1986 in Agha’s MIT thesis (Agha, 1986), which defined actors as behavior functions over messages and treated cross-actor effects as emitted output messages of the function. The formal account ratified what Hewitt et al. (1973) had stated informally: an actor’s program describes how it responds to the messages it receives; what happens *between* actors is the operational semantics of the model, not the content of any actor’s program.

By the close of the 1980s, the actor lineage — Hewitt’s informal formulation, Agha’s formalization — had stabilized a single architectural decision: the boundary between an actor and the message-passing layer was drawn deliberately, and the layer between actors was characterized as operational rather than programmatic. (CSP reached a structurally similar boundary by a different route and under a different ontology; the resemblance is worth a note of contrast, not a claim that one assumption spans both traditions.)

2.2 The pragmatic maturation (1990s–2000s)

The operational frame became productive in Erlang. Armstrong’s (2003) PhD thesis argued that fault tolerance becomes tractable precisely because processes do not share programmatic flow: failure is handled across process boundaries by links and exit signals. In Armstrong’s own terms, “a process can supervise the existence of another process by setting up a link to it. When a process terminates, it automatically sends exit signals to the process to which it is linked” (Armstrong, 2003, p. 217). Supervision is a structural pattern in which a supervisor observes the failure of a child process and reacts. The supervisor’s program is local; the failed process’s program does not extend into the supervisor; the cross-process relationship is mediated by VM-level links and exit signals. Cross-process causation is an infrastructure concern. The frame stabilized further by becoming useful: the operational-rather-than-programmatic separation was now what made fault-tolerant systems possible.

2.3 The modern instantiations (2010s–)

Two implementations carried the actor frame into modern production systems. Akka, the canonical JVM actor framework (Lightbend, n.d.), characterizes its fundamental message-send pattern directly:

“Tell is asynchronous which means that the method returns right away. After the statement is executed there is no guarantee that the message has been processed by the recipient yet. It also means there is no way to know if the message was received, the processing succeeded or failed.”

The verb is named — `tell` — and so is its absence of program-level effect. The dispatch occurs; the sender’s program does not record that it occurred, nor

whether the recipient acted on it.

Microsoft’s Orleans (Bernstein et al., 2014) introduced *virtual actors*: location-transparent grains whose dispatch is mediated by the Orleans runtime. Each grain has a key; the runtime decides where the grain lives and routes calls; the grain’s code observes only local state and method invocation. The runtime maintains the cross-grain dispatch logic as infrastructure; the grain’s program never sees nor records cross-grain causation as part of its narrative.

Decades apart, one assumption. In Hewitt et al. (1973): “*Sending a message to an actor makes no presupposition that the actor sent the message will ever send back a message to the continuation*” (p. 241). In Akka 2014–present: “*there is no way to know if the message was received, the processing succeeded or failed.*” The two formulations are decades apart and identical in substance. The verb **tell** has been part of actor-systems vocabulary throughout. What this paper observes is that the verb names the dispatch, not the program: a **tell** from actor A to actor B leaves no trace in either actor’s program. The verb has existed; the assumption that the verb’s effect was extra-programmatic has survived intact alongside it.

2.4 The naturalized assumption

Across fifty years and four canonical generations of literature — Hewitt’s foundational paper, Agha’s formalization, Armstrong’s pragmatic maturation, and the modern instantiations in Akka and Orleans — the question of whether cross-actor flow is a *statement* of any actor’s program is not the one the lineage asks. Where cross-entity causation is recorded at all — the message logs and causation identifiers of §6.4 — it is recorded operationally, beneath or beside the program, never as the program’s own sentence. The assumption does not appear as a defended thesis; it is the default frame within which the lineage operates.

The forms in which the assumption surfaces vary by generation; the substance is continuous.

Year	Work	Anchor	Form of the assumption
1973	Hewitt et al. — <i>A Universal Modular ACTOR Formalism</i> (IJCAI)	“The basic unit of execution on an actor machine is sending a message in much the same way that the basic unit of execution on present day machines is an instruction.” (pp. 236–237)	Sending is a machine-level primitive, parallel to a hardware instruction; therefore not content of any actor’s program.
1986	Agha — <i>Actors: A Model of Concurrent Computation</i> (MIT)	Actors are behavior functions; cross-actor effects are emitted output messages of the function.	The behavior function is local; the cross-actor effect is the <i>output</i> of the function, separate from the function’s body.
2003	Armstrong — <i>Making reliable distributed systems...</i> (KTH)	“A process can supervise the existence of another process by setting up a link to it [...] it automatically sends exit signals to the process to which it is linked.” (p. 217)	Cross-process flow is infrastructure-level (VM links and exit signals), not part of any process’s program.
2010s	Lightbend — <i>Akka Interaction Patterns</i>	“There is no way to know if the message was received, the processing succeeded or failed.”	The dispatch verb (tell) exists; its effect is explicitly outside the sender’s program — no acknowledgment record.

Year	Work	Anchor	Form of the assumption
2014	Bernstein et al. — <i>Orleans: Distributed Virtual Actors</i>	The Orleans runtime places grains and mediates cross-grain calls; grain code sees only local state.	Cross-grain dispatch is hidden by the runtime; the grain’s program never sees nor records cross-grain causation.

The continuity of this assumption across five decades is not the result of oversight, and still less of error. Each canonical contribution was optimizing a different property: Hewitt, Bishop, and Steiger, a uniform account of concurrency; Agha, a formal semantics of actor behavior; Armstrong, fault tolerance through process isolation; Akka and Orleans, distribution and location transparency. None was optimizing the *programmatic preservation of cross-actor causation*, because that was not the property under design. The assumption is the shadow of a different objective function, not a mistake the field failed to notice — which is why questioning it is not a charge against the lineage but a change in what one chooses to optimize. What follows does not dispute the effectiveness of the separation; it questions whether the separation is necessary.

§3 names this assumption explicitly. §4 demonstrates that it is contingent rather than necessary.

3. The assumption named

What §2 has shown to be continuous across fifty years admits a single one-line formulation. The structural claim that links the five entries of §2.4 — Hewitt’s “machine-level instruction”, Agha’s emitted output messages, Armstrong’s VM-supported links, Akka’s `tell`, Orleans’ runtime-mediated dispatch — is:

Causation between actors is operational, not programmatic.

The two terms in the formulation are the working categories of the rest of this paper. They also answer, for now, the question that opened it: *where does the causal sentence live?* Under the assumption, it lives in the operational layer — between actors, mediated by the runtime — and not in any actor’s program.

Operational designates effects that the system performs around or between actors but that no actor’s program records. Message dispatch by a runtime, su-

pervision links between processes, virtual-actor placement decisions, message-broker routing, and distributed-tracing correlation IDs are all operational in this sense: their existence is mediated by infrastructure, and their effect on the system is real, but no program in the system contains the act of mediation as a statement.


Programmatic designates an effect that an actor’s program contains as a *statement of the program* — an operation the program executes, observable within the program’s own narrative and re-executed when the program is replayed. A method invocation an actor performs on itself is programmatic in this sense; so is a state mutation it writes. The program records what it did, as the doing, and replay re-runs it.

The distinction has a linguistic signature. A *command* is a directive — it asks for what is not yet done, named in the imperative present (**ApplyReward**, **Confirm**); an *assertion* reports a fact already lived, named in the past (**PurchaseConfirmed**). Treating the cross-actor send as operational collapses the two into one runtime dispatch — a single send verb carrying both, as in the canonical frameworks (§2.4); treating it as programmatic keeps them distinct and lets the sender record the one it can truthfully make — the assertion of what it did. The verb tense is the surface mark of which act the sentence performs; the choice this paper argues for is the one under which that mark survives into the program (developed at §8.2).


The boundary is not whether the send is recorded in the actor’s own log. A serialized event can be appended there without being a statement of the program. Conventional event sourcing already does this: an aggregate raises a **MessageSent** or integration event when it dispatches to another context, and the event lands in the aggregate’s own stream. But that event is *data the program emitted*, not a *sentence the program runs* — on replay it is folded back into state by a reducer, never re-executed as a statement of the program, and the dispatch is relayed by an outbox or broker the program does not contain. By the test used here it is operational-in-effect despite its location. A recorded cross-actor send is *programmatic* only when the send is a dense operation in the program — an executable sentence naming its recipient and message by reference, replayed as program rather than re-applied as state. *Replayed as program* means the recorded statement is re-executed by the interpreter — the program re-runs, reconstructing the state the send produced — not that the external dispatch is repeated: no journaled operation repeats its external effects on replay (replaying a debit reconstructs the balance without re-contacting the bank), and the tell is no exception. The programmatic property is that the send is a re-executable statement the program re-runs, where a serialized event is data a reducer folds without the program ever re-executing a send. That density criterion — a recorded operation, not a serialized payload — is the *anti-porosity* established in Paper 1; the present paper takes it as a precondition and asks what becomes possible at the actor boundary once it holds.

The contrast can be visualized:

Under the assumption (§3): causation operational, not programmatic



Under the alternative (§7): causation as program statement



Both views describe the same flow; they differ in what is recorded as program. Under the assumption, the cross-actor send lives in infrastructure between actors and is invisible from any actor's program. Under the alternative, the cross-actor send is itself a journal entry on the sender's side, with a corresponding receipt entry on the receiver's side.

The naturalized assumption is the claim that, in the ontology of an actor system, the act of one actor causing an effect in another belongs to the *operational* category and not the *programmatic* one. The 1973 framing of message-sending as a hardware-instruction analogue is this claim. The 2014–present framing of **tell** as a fire-and-forget dispatch with no acknowledgment record is this claim. Across formulations, the assumption holds.

§4 demonstrates that the assumption, while continuous in the literature, is contingent rather than necessary: no theorem of the actor model entails it. §5 examines the consequences of the assumption — what is structurally lost when cross-actor causation lives outside any program. §6 shows why the existing repertoire of patterns built atop the actor model — sagas, choreography, distributed tracing, workflow engines — cannot dissolve the assumption, because they reconstruct cross-actor flow *after* the fact rather than preserving it as program. §7 reformulates: the assumption is contingent, the alternative is constructible, and the cross-actor flow can be a sentence of the program.

4. Contingency

The assumption named in §3 — that causation between actors is treated as operational rather than programmatic — would be necessary if it followed from a theorem of the actor model, from a foundational property the model formally entails, or from the structural commitments that distinguish actor systems from other concurrency paradigms. It does not. This section traces what the actor model formally requires and what it does not, and shows that the assumption is a contingent feature of how actor systems have historically been built rather than a consequence of what they are.

4.1 What the actor model entails

The actor model, in its canonical formulations from Hewitt et al. (1973) through Agha (1986), formally entails three structural commitments:

1. **Autonomy.** Each actor has its own state and processes messages serially. No actor’s behavior is contingent on the simultaneous execution of another’s.
2. **Message-based communication.** Actors communicate by passing messages. No actor reads or writes another’s state directly.
3. **Isolation.** An actor’s state is private to it. The address space of an actor is not shared.

These three commitments are what distinguish actor systems from shared-memory concurrency models. They define the ontology of actors, not the mechanics of their runtimes. They are the substance of “what the actor model is” in any reading of the canonical sources.

4.2 What it does not entail

The three commitments above do not entail that the *act of sending* be invisible to the sender’s own program. They do not entail that cross-actor dispatch be mediated by a runtime layer that owns the send. They do not entail that a record of “actor A spoke to actor B at time T” must reside outside actor A’s narrative.

In Hewitt et al. (1973), the act of sending is an action the actor performs. The framing of message-sending as a machine-level instruction analogous to a hardware instruction (Hewitt et al., 1973, pp. 236–237) is an interpretive choice that supports the architecture proposed in that paper; it is not a formal consequence of what an actor is. In Agha (1986), the actor’s behavior is modelled as a function from (state, message) to (next state, outgoing messages) (Agha, 1986, ch. 4). The outgoing messages are output of the function. The formalization neither requires nor prohibits the function from producing, in addition, a record of the send observable in the actor’s own program — the question is not raised. The formalization can be extended, without modifying any of the three commit-

ments above, to include as output not only the outgoing messages but also a record of the send observable within the actor's own program.

What the canonical formulations establish is the *boundary* between actors — autonomy, isolation, message passing as the exclusive cross-actor mechanism. They do not establish that the boundary be invisible from within. The opacity of the cross-actor send to the sender's program is a separate decision that the canonical sources adopted but did not derive.

4.3 The objection from isolation

The most common objection to making cross-actor sends programmatic is that doing so would violate actor isolation. The objection conflates two distinct properties.

Isolation requires that no actor read or write another's state. Recording a cross-actor send in the sender's journal does not require either. The sender writes to its own journal — its own state. The journal entry describes what the sender did: “*I sent a message of type M to actor B at time T .*” The receiver, when it processes the message, writes a corresponding entry to its own journal: “*I received a message of type M from actor A at time T .*” Both records are local; neither requires shared state. Isolation is preserved.

What changes is not who can read whose state, but what each actor's program is permitted to say about its own activity. That change is local to the actor's own program; it is invisible across the boundary.

4.4 The objection from orchestration

The second common objection is that recording cross-actor causation as program would introduce orchestration. The objection conflates a record of causation with a director of behavior.

Orchestration is the architectural pattern in which an external coordinator decides what each participating actor does next. The coordinator dispatches commands; the participants execute them. Causation flows from the coordinator outward.

A journal entry that records what an actor did is not a coordinator. It does not direct anything. Each actor still processes its own messages autonomously, decides what to do, and emits its own outgoing messages. The journal preserves what happened; an orchestrator would prescribe what should happen. Recording causation does not introduce a director. It introduces narration, not control.

4.5 Closing

The assumption stated in §3 is therefore not entailed by the actor model. It does not follow from Hewitt's formulation, from Agha's formalization, from the

requirement of isolation, or from the absence of orchestration. The opacity of the cross-actor send to the sender’s program is contingent in exactly this sense — not a theorem of the model — but it is not arbitrary: it is the shadow of a different objective function (§2.4), well suited to location transparency, decoupling, and fault isolation, and so adopted by the canonical sources and left in place. It is not, in any formal sense, a property the model requires.

The consequence is bounded but decisive: the absence of semantic continuity at actor boundaries is not a necessity of the actor model. It persists because it served other objectives, not because the model demands it — so removing it is a change of objective, available without abandoning a single commitment, rather than the correction of an oversight.

5. Consequences

When the assumption stated in §3 is in force, cross-actor causation is held outside any actor’s program. The consequences of this displacement are not operational inconveniences; they are structural. Four are examined here. Each is a distinct symptom; each is the same absence — the absence of semantic continuity at actor boundaries — observed from a different angle.

5.1 Auditability requires external reconstruction

The question “*why did this happen?*”, asked of an effect that crossed an actor boundary, cannot be answered by reading any single actor’s program. The sender’s journal records what the sender did; the receiver’s journal records what the receiver did; neither records the link. To answer the question, an auditor must consult tools that live outside the programs: correlation IDs threaded through messages by the runtime, distributed tracing systems that observe the runtime from outside, log aggregation pipelines that join records across actor boundaries.

These tools work, but their necessity is evidence that the program does not contain what they reconstruct. They do not extend the program; they substitute for it. Auditability therefore depends on infrastructure that compensates for what the program does not record.

5.2 Replay does not reconstruct cross-actor history

Replaying an actor’s journal reconstructs that actor’s local history: which messages it received, how it responded, what state it produced. Replay of two actors’ journals reconstructs two local histories. It does not reconstruct the joint history. The joint history never existed as a program artifact to be replayed.

The reason is direct. The act of one actor sending a message to another is not

in either journal as a program statement. The sender’s journal contains the operations leading up to the send; the receiver’s journal contains the operations following its receipt. The send itself — the event that joins them — has no representation in either record. Replay can therefore reconstruct each actor in isolation, but not the causal chain that connects them.

This consequence matters wherever event sourcing, time-travel debugging, or post-hoc analysis is intended to span more than one actor. The model that promises replayability delivers it within actors and not across them.

5.3 Cross-datacenter replication loses program-level causation

When journals are replicated across data centers — a standard pattern for disaster recovery and geographic distribution — the per-actor histories travel with them. The cross-actor causation does not, because it was never written anywhere that replication can carry. It lived in the dispatcher, the message broker, the runtime — components that are typically per-cluster or per-DC.

Recovery from a DC failure can therefore reconstruct each actor’s local state but cannot reconstruct, from the replicated material alone, the causal chain that led to the system’s pre-failure state. The chain is reconstructed by replaying the broker’s queue, by re-issuing messages that were in flight, or by reconciliation processes that assume the state and reconstruct backward. None of these steps is part of any actor’s program; all of them depend on infrastructure-level apparatus that may or may not have been replicated.

5.4 Debugging across actors is log archaeology

A developer asked to explain why an actor produced a particular state, when that state was caused by a chain of events crossing several actors, cannot read a single program to find the answer. The developer assembles the answer from logs, traces, broker dumps, and timestamps — pieces of evidence whose joining is itself an act of reconstruction. The discipline that emerges is log archaeology: the careful inference of a causal narrative from artifacts that are not narrative themselves.

The discipline is real and, in some organizations, mature. Its maturity is proportional to the absence of programmatic causation — the cost is not that the discipline fails but that it must exist at all, reconstructing from infrastructure what no program records. But its existence is the symptom. A program whose causation is recorded as program admits debugging by reading; a program whose causation is dispersed across infrastructure admits debugging only by inference.

5.5 Closing

The four consequences are not separate problems with separate fixes. They are four manifestations of a single structural absence: causation between actors is held outside any actor’s program, so any operation that requires reading the cross-actor causal chain — auditing, replaying, replicating, debugging — must reconstruct it from artifacts that lie outside the program.

Symptom	What is lost	What external apparatus reconstructs it
Auditability	the causal chain as part of any program	distributed tracing, correlation IDs, log aggregation
Replay coherence	the joint history reconstructible from per-actor journals	custom event-sourcing layers that span actors
Cross-DC replication	program-level causation when journals cross DCs	broker-level log replication, reconciliation processes
Debugging	the ability to read the cross-actor flow as a single program	log archaeology — manual correlation of timestamps and IDs

The cost of the assumption is not a list of operational difficulties to be addressed by better tools. It is the systematic displacement of causation from program to infrastructure. The tools that arise — tracing, correlation IDs, reconciliation processes, log archaeology — are not extensions of the program. They are evidence of what the program does not contain.

6. Why existing approaches do not dissolve the assumption

The consequences of the assumption named in §3 — auditability through external reconstruction, replay limited to single actors, cross-DC fragility, debugging by log archaeology — are not unrecognized in the actor-systems literature. Each has accumulated a corresponding repertoire of patterns intended to address it. This section examines the principal patterns and shows that they are responses to the consequences, not dissolutions of the assumption that produces them. Each pattern reconstructs cross-actor flow somewhere; none records it as a sentence of any participating actor’s program.

This section carries as much of the paper’s weight as the instantiation that follows. Once the assumption of §3 is named, sagas, choreography, distributed tracing, and workflow engines stop reading as unrelated tools and resolve into responses to one question — *where does the causal sentence live?* — each answering it the same way: by placing the sentence somewhere outside the participating actors. That so many mature patterns are answers to a single absence is the paper’s central claim; §8 is its existence proof.

6.1 Sagas

The saga pattern, in both its orchestrated and choreographed forms, addresses the problem of multi-actor business transactions: an operation that requires several actors to act in sequence, with compensating actions if any step fails.

In the orchestrated form, a saga orchestrator holds a state machine that represents the cross-actor flow. The orchestrator dispatches commands to the participating actors and listens for their responses; if a step fails, the orchestrator emits compensating commands that undo or correct earlier steps. The cross-actor flow is therefore programmatic, but it is the orchestrator’s program — not the program of any participating actor. From the perspective of an actor receiving the orchestrator’s command, the message is indistinguishable from any other; the actor’s program does not know it is in a saga.

In the choreographed form, the orchestrator is dissolved into a protocol. Each participating actor publishes events when it completes a step; the next actor in the flow subscribes to those events and triggers its own step. The cross-actor flow is encoded across the actors as a distributed protocol — but no single actor’s program contains the flow. Each actor’s program contains its local response to events; the flow itself is the implicit composition of those responses, observable only from outside.

Both forms treat the cross-actor flow as a separate concern from any participating actor’s program. The orchestrated form externalizes it to a coordinator; the choreographed form distributes it across event subscriptions. Neither form makes the cross-actor send a sentence of the sender’s program. The assumption stated in §3 is preserved by construction. The saga makes cross-actor flow programmatic only by relocating it outside the actors whose behavior it coordinates. A saga also does more than relocate the record: it implements transactional recovery — compensating actions, retries, step boundaries — a concern that is genuinely hard independently of where the causal record lives. *tell* does not provide that recovery and is not a substitute for it; what *tell* addresses is narrower and prior — *where the record of the cross-actor send lives*. The two are different instruments, compared here only on that axis. They compose rather than compete: *tell*’s relocation of the record is the substrate a receiver-side saga can sequence over — each asserted fact a step the saga recognizes, a compensating action itself another asserted fact (**SaleReversed**) rather than a coordinator’s command.

6.2 Distributed tracing

Distributed tracing — OpenTelemetry, Zipkin, Jaeger, and similar frameworks — addresses the problem of observability across actor boundaries. A trace context is propagated through messages by the runtime; spans are emitted by participating components; a trace storage backend reconstructs the cross-actor causal chain from the spans, presenting it as a tree or timeline.

The reconstruction is real and useful. It also confirms the structural absence it compensates for. The trace is built from artifacts that exist outside the participating actors' programs: the trace context is metadata threaded through messages by middleware; the spans are emitted by instrumentation that observes the runtime; the trace storage is a separate service that joins the spans. None of these artifacts is part of any actor's program. They are the apparatus that makes the absence navigable.

A trace is also lossy in a way that program is not. A trace records that a span with given attributes occurred; a program records what was said and why. The two are not equivalent. The trace is sufficient for observability; it is not the program of the cross-actor flow. It exists precisely because no such program exists.

6.3 Workflow engines

Workflow engines — Temporal, Cadence, AWS Step Functions, Camunda, Argo Workflows — externalize the cross-actor flow into a dedicated programming environment. The workflow is written as a separate program in the engine's own model; the engine dispatches activities to participating services; the engine's program records the flow as it executes.

Workflow engines are unusual among the patterns considered here because they do produce a programmatic record of the cross-actor flow. The flow is a program — an explicit, replayable, durable program. But the program belongs to the workflow engine, not to any participating actor. The actors are activities invoked by the workflow; their programs do not contain the workflow.

The asymmetry matters. From the perspective of a participant, an activity invocation is a remote procedure call from an external system; the participant's program records that it received an invocation and produced a result. The relationship between activities — the flow that connects them — lives in the workflow engine, in a separate codebase, in a separate execution model. Reading the participant's program does not reveal the workflow; reading the workflow does not reveal the participant's local program. Two complementary records that do not compose into one. The workflow and the actors form a bipartite description of what, in a semantically continuous system, would be a single program.

6.4 Causal and message logging

A sharper objection comes not from the business-flow patterns above but from distributed-systems fault tolerance, where recording what one process sent another is decades old. Message logging — in its sender-based, receiver-based, and causal variants, surveyed by Elnozahy, Alvisi, Wang, and Johnson (2002) — records the messages a process sends or receives so that a crashed process can be reconstructed by replaying them (Johnson & Zwaenepoel, 1987). Lamport’s happened-before relation (Lamport, 1978) and the vector clocks that refine it (Fidge, 1988) record cross-process causal order directly. These mechanisms pre-date this paper by decades and unambiguously record cross-entity causation. A distributed-systems reader is entitled to ask: how is a *tell* recorded in the sender’s journal different from a sender-based message log, which has existed since the 1980s?

The difference is the operational/programmatic distinction of §3, and message logging falls cleanly on the operational side. A message log is maintained by the recovery protocol beneath the application, for the recovery protocol’s purposes: it captures message payloads and delivery metadata so the fault-tolerance layer can resend or replay them after a crash, and it is read by that layer, not by any actor’s program. The application’s program does not contain the act of sending as a sentence; the log is apparatus the runtime keeps *about* the program — the same category as the distributed trace of §6.2, differing in purpose (recovery rather than observability) but not in kind. A vector clock is narrower still: it records the *order* of causation, not its *content* — that B’s state causally followed A’s, not what A said to B or why.

Tell records the send as a sentence in the actor’s program: a dense DSL statement in the journal that *is* the program — the anti-porosity of Paper 1, the externalized parameters of Paper 2 — read as narrative and replayed as program. The contrast is exact. A sender-based message log answers “*what bytes must I resend to recover this process?*”; the journal answers “*what did this actor say, to whom, and why?*” The first is a recovery artifact, payload-oriented and type-erased; the second is the program. The same holds for event-sourcing causation and correlation identifiers (Young, 2010; Vernon, 2013) and for the process-manager pattern (Hohpe & Woolf, 2003): each threads an identifier or a coordinating state machine through messages so the cross-entity chain can be reassembled downstream — operational reconstruction by the same logic as §6.1–6.3. That cross-entity causation has been captured for thirty years is not in dispute. What no prior mechanism does is record it as a sentence of the sending actor’s own program. Message logging is, if anything, the strongest witness for §3’s distinction: it is the most mature, most studied apparatus for capturing cross-actor sends, and it lives entirely on the operational side of the line.

6.5 Closing

The patterns examined fall into two groups. Orchestrated sagas, choreographed sagas, distributed tracing, and workflow engines (§6.1–6.3) are responses to the same consequence: cross-actor flow is not in any participant’s program, and something must compensate for that absence — a coordinator’s state machine, an event-subscription protocol, a trace storage backend, a workflow engine’s separate program. Causal and message logging (§6.4) arises for a different purpose — fault-tolerant recovery — but shares the decisive property: the cross-actor send it records lives in a log beneath the program, not as a sentence within it.

Pattern	Where the flow is encoded	Does it preserve the flow as program in the actors?
Saga (orchestrated)	In the orchestrator’s state machine	No — externalized to a coordinator
Saga (choreographed)	Across event handlers in each actor	No — only local responses appear
Distributed tracing	In the trace storage, observed by an external system	No — the program is observed, not extended
Workflow engine	In the workflow engine’s separate program	No — the workflow is a separate artifact
Causal / message logging	In the recovery layer’s message log, beneath the program	No — apparatus for crash replay, not a program statement

The five “No”s share a structure. In every case, cross-actor flow is made programmatic only by placing the program somewhere other than in the actors whose behavior constitutes the flow. The compensation always occurs outside the participants.

These patterns are therefore not solutions to the assumption stated in §3; they are architectural responses that accept the assumption and build systems around it. Their sophistication, maturity, and widespread adoption are not evidence that the assumption is correct. They are evidence that the absence of semantic continuity is costly enough to justify entire subsystems dedicated to reconstructing what the program does not contain.

7. Reformulation

The diagnosis is complete. The assumption stated in §3 — that causation between actors is treated as operational rather than programmatic — is not

entailed by the actor model (§4); produces a structural absence of semantic continuity at every actor boundary (§5); and survives in every existing pattern that responds to the resulting consequences (§6). The remainder of the paper takes the diagnosis as established and asks what would have to be the case for the assumption to be dissolved. This section answers that question by deriving the conditions under which semantic continuity is preserved across actors, without violating any of the actor model’s structural commitments and without introducing orchestration.

7.1 Three conditions

The conditions are derived directly from §4. The actor model entails autonomy, message-based communication, and isolation (§4.1). The construct introduced in §1 requires that the cross-actor causal chain be recorded as part of some actor’s program (semantic continuity). Reconciling the two yields three conditions, each of which addresses one structural commitment of the model and one element of the construct. The conditions are not design choices; they are the consequences of reconciling the model’s commitments with the construct.

Condition C1 — Locality of writes. When an actor sends a message to another actor, the sender records the act of sending in its own journal, and only there. The receiver, on processing the message — including an acknowledgment returned for an earlier send — records the receipt in its own journal, and only there. No actor writes to another’s state. This condition preserves the model’s isolation commitment: each journal is local property; recording cross-actor causation does not require shared state.

Condition C2 — Causation as program statement. The cross-actor send appears as a sentence in the sender’s program — a journal entry whose content names the recipient and the message. Reading the sender’s journal reveals not only what the actor did locally but also the act of speaking that connected its program to another actor’s. The cross-actor edge becomes part of the program’s narrative rather than an artifact reconstructed from outside it. The construct introduced in §1 — semantic continuity — is realized at the sender’s boundary by this condition; symmetrically, the receiver’s program records the receipt as a sentence about who spoke to it.

Condition C3 — No external coordinator. No party outside the participating actors authors what each actor does next — what it does, which message it sends, how it responds. The message-passing layer may carry those messages, with whatever reliability it provides; what C3 forbids is an external *author* of the flow — a coordinator whose state machine chooses each actor’s next step — not an external *carrier* of it. Each actor processes its own messages autonomously and decides its own response. The journal records what happened; no orchestrator prescribes what should happen, and no participant outside the actors is required to interpret the record. This condition preserves the absence of orchestration that §4.4 identified as a non-negotiable property of

the alternative.

The three conditions are independent. C1 alone preserves isolation but does not establish continuity; C2 alone establishes continuity but, without C1, would risk shared-state writes; C3 alone preserves autonomy but does not address the recording question. Together, the three conditions yield a system in which cross-actor causation is recorded as program in each participant's local journal, dispatched through the existing message-passing layer, and coordinated by no external party.

7.2 *tell* as primitive

A primitive that satisfies the three conditions by construction can be defined.

Define *tell* as a sentence in an actor's program that, in a single act, (a) records an entry in the sender's own journal naming the recipient and the message, (b) dispatches the message through the existing message-passing layer, and (c) requires no coordination with any party outside the sender and the recipient.

The three components correspond to the three conditions. Component (a) satisfies C1 — the journal entry is written only to the sender's journal — and C2 — the entry is the program statement that names the cross-actor causation. Component (b) inherits the existing message-passing semantics of the actor model; nothing about the dispatch mechanism is new. Component (c) satisfies C3 — the act involves only the sender's local write and the dispatch; no orchestrator participates.

The primitive is conceptual. Its naming follows actor-systems convention: in Akka and elsewhere, *tell* designates a fire-and-forget cross-actor send (§2.3). The collision with Akka's verb is not accidental but argumentative: the verb has been part of actor-systems vocabulary for over a decade alongside the assumption that the verb's effect was extra-programmatic. The primitive defined here keeps the verb and changes what the verb records. Where Akka's *tell* leaves no trace in either actor's program (§2.3), this primitive's effect is to make the trace the program. The difference is not in the dispatch semantics but in what the program is allowed to say about the dispatch.

The canonical effect can be stated in one sentence: *tell* turns the journal into the boundary of causation. The boundary between actors does not disappear; it remains the locus of message passing. What changes is that the boundary becomes inscribed in each participating actor's program.

7.3 What changes, what does not

A reader familiar with the actor model may ask whether the conditions above amount to a different model. They do not. Each of the actor model's three commitments survives unchanged.

Autonomy is preserved. Each actor still processes its own messages serially. *tell* introduces no requirement of synchronization with other actors; the sender does not wait for the recipient. The conditions add a write to the sender's local journal; they do not couple the sender's progression to the recipient's.

Message-based communication is preserved. Actors still communicate by passing messages, exclusively. *tell* uses the same message-passing layer that the actor model already specifies; no shared memory, no remote procedure calls, no synchronous read of another actor's state. The dispatch mechanism is unchanged; only the recording at each end is added.

Isolation is preserved. No actor reads or writes another's state. The sender's journal entry is written by the sender. The receiver's journal entry, if any, is written by the receiver. The two records are local; they reference each other by content, not by shared address.

What changes is what each actor's program is permitted to say about its own activity (echoing §4.3). What does not change is what each actor can do, observe, or share. The opacity of the cross-actor send to the sender's program — which §3 named as the assumption and §4 demonstrated to be contingent — is the only thing that the conditions remove. The actor model remains intact; only the historical interpretation of what must remain invisible to the program is removed.

7.4 The shape of the act

The structural shape of *tell* is rendered below. The diagram is included not to introduce new content but to make the relations among the components of §7.2 visible at a glance. The diagram is deliberately mundane: nothing new is introduced except the presence of the journal entries.

sequenceDiagram

```
participant ProgA as Actor A's program
participant JA as A's journal
participant MPL as Message-passing layer
participant ProgB as Actor B's program
participant JB as B's journal

ProgA->>JA: tell msg to B – record the assertion
ProgA->>MPL: dispatch msg
MPL->>ProgB: deliver msg
ProgB->>JB: process and record receipt entry
```

Two journal writes (in JA and JB), one dispatch through the existing message-passing layer, no participant outside the sender and the recipient. The diagram reproduces the three conditions visually: writes are local (C1), the cross-actor flow is recorded as program at both ends (C2), and no coordinator is present (C3).

7.5 Closing

The conditions are not extensions of the actor model; they are the conditions under which the assumption named in §3 can be removed without altering the model’s structural commitments. *tell* is one realization of the conditions; any realization that satisfies C1, C2, and C3 would dissolve the assumption identified in §3. But the conditions presuppose a substrate: C2 requires the cross-actor send to be a *statement of the program* — a dense operation, not a serialized event (§3) — which is the anti-porosity of Paper 1 and the externalized parameters of Paper 2. The generality is therefore over primitives *on a journal-as-program substrate*, not over actor frameworks in general; a framework that records sends as opaque payloads would satisfy C2 only by first adopting that substrate, which is the subject of Papers 1–2 and beyond this paper’s scope.

The reformulation is conceptual. The remainder of the paper presents the empirical question: can the conditions be realized in a running system? §8 answers by exhibiting an instantiation and demonstrating it through an illustrative case study.

8. Instantiation

8.0 Origins of the instantiation

The instantiation discussed below is provided by Puppeteer, an actor-based framework whose journal-as-program substrate predates the analysis presented in this paper. The framework’s lineage runs from a 2005 autopersistence prototype — domain classes that persisted themselves through reflection, with no schema decisions in the domain code — through a DSL that emerged as the persistence substrate, to event sourcing as the runtime’s primary discipline. By the time the conditions of §7 were articulated, the framework already satisfied them; the present section reads as observation of that alignment, not as derivation of the framework from the conditions. This invites a fair question of circularity — were C1–C3 read off Puppeteer and then presented as model-derived? They are not: §7.1 derives them from the actor model’s three commitments and the construct alone, citing no framework feature. That a framework built for other reasons, years earlier, already satisfies them is therefore independent evidence that the model-derived conditions are satisfiable — not a sign they were retrofitted.

8.1 The case study domain

The case study is a simplified loyalty domain modelled on a production scenario from a deployment of the framework. A Seller actor confirms a purchase order via a domain command. A RewardEngine actor holds a registry of campaigns; for each purchase event, the RewardEngine evaluates which campaigns qualify

(by date and amount thresholds) and applies the corresponding rewards to the customer.

The flow is intentionally minimal: one actor produces a domain event, another reacts to it, and the relationship between the two needs to cross an actor boundary. The simplicity is deliberate — the case study illustrates a structural property of the cross-actor mechanism, not the complexity of any particular business workflow.

8.2 The instantiation: *tell* in Puppeteer

Puppeteer’s Reaction surface exposes three planes — *Program* for in-actor read-only effects, *Metadata* for journal-metadata changes, and *Causation* for cross-actor causation. The three planes name what the verb touches; *tell* lives on the third.

The Seller’s Reaction is defined as follows:

```
seller.Reactions.DefineReaction("PurchaseFunnelToRewards")
  .Job().Company()
  .WithSharedHydration()
  .Seek("Purchase")
    .OnMatch("[s: Seller].purchase($order, $date, $amount, $customer)")
  .Causation.Continue(@"
    tell PurchaseConfirmed
      with @order, @date, @amount, @customer
      to RewardEngine('rewards-1')
      once @order;
  ");
```

The *tell* statement is a sentence in the actor’s DSL, and the kind of sentence matters: the Seller *asserts a fact it has lived* — *PurchaseConfirmed*, named in the Seller’s own vocabulary — addressed to the *RewardEngine* (to *RewardEngine('rewards-1')*), carrying the values that fact involved (with *@order, @date, @amount, @customer*) under a per-utterance identity taken from the order it carries (*once @order*), so one compiled action issues a distinctly-identified assertion for each purchase. It does not invoke a method on the *RewardEngine*, and it does not name how the message travels. This follows the single discipline the journal obeys throughout this series: an actor’s program may record only what that actor could itself have said. The Seller can say *that a purchase was confirmed*; it cannot say *how the RewardEngine applies rewards* — that is the *RewardEngine*’s verb, recorded in the *RewardEngine*’s own journal — nor *which broker carries the message*, which is deployment, not program. The Reaction is established once and thereafter watches the Seller’s journal; when the domain command *s.purchase(...)* lands as an entry, the standing Reaction matches it and its *.Causation.Continue(...)* body fires, journaling the assertion on the Seller’s side.

The assertion is journaled as a typed *message-action* — defined once (its signature deduced from the values it carries) and then invoked. This is the same define-then-invoke shape *every* parameterized operation takes on the dense journal-as-program substrate this series builds on (Papers 1–2): the domain purchase, issued with typed parameters (**date**, **amount**), is journaled the same way, so the shape below appears twice — once for the purchase, once for the assertion. After the bridge delivers the assertion to the RewardEngine and the receiver acknowledges, the Seller’s journal contains six entries:

```
[0] s = Seller();
[1] (define of the message-action s.purchase –
its typed signature)
[2] s.purchase('ord-100', 5/9/2026, 250, 'cust-42');
[3] (define of the message-action PurchaseConfirmed –
its typed signature)
[4] tell PurchaseConfirmed
      with 'ord-100', 5/9/2026, 250, 'cust-42'
      to RewardEngine('rewards-1')
      once 'ord-100';
[5] tell ack 'ord-100' from RewardEngine('rewards-1');
```

The values are passed as typed parameters, not interpolated into the script text — the definition records the signature `s.purchase(order, date, amount, customer)` and the invocation records the bound arguments, so the journal stores the operation and its values separately, exactly as the tell does. The tell’s identity is the order it carries (**once @order**), which the invocation binds to **'ord-100'** — the ack in entry [5] returns under that same key. The tell’s signature is inferred once — at first use, on the live path — and recorded in entry [3]; replay replays that recorded definition rather than re-inferring it, so the message-action is identical and order-independent across replays. This is the define-and-invocation discipline of Paper 2, under which a name is defined once and thereafter invoked; a reused name resolves as any operation does there. The inference is a write-time convenience, never a replay-time computation.

The same separation fixes what crosses the boundary. What the RewardEngine receives is the asserted fact and its values — the bound arguments of entry [4] — and only those; the parameter names, the message definition, and any operational envelope the Seller keeps about the act (the identity or request context under which the purchase was recorded) stay in the Seller’s own program and do not travel with the tell. A tell is a domain speech act: a value the receiver’s domain needs must be *in the assertion*, carried as one of its values, because the receiver maps the asserted message to a command it owns and can read nothing the assertion did not say. This is the cross-actor face of a discipline the journal obeys throughout — domain values move as arguments; the metadata a program keeps about its own operations is a separate channel that is not part of what it asserts (the *operational-vs-program* distinction of §3).

The three conditions of §7 are realized by these entries.

- **C1 (Locality of writes).** The six entries are in the Seller’s journal only. The RewardEngine’s journal records its receiving operation independently — the two journals never share storage.
- **C2 (Causation as program statement).** Entries [3]–[4] are the cross-actor assertion rendered as a DSL sentence — defined once as a typed message-action, then invoked. Reading the Seller’s journal alone reveals that it asserted `PurchaseConfirmed` to RewardEngine, with what values, at what time. Entry [5] closes the round-trip with the receiver’s acknowledgment.
- **C3 (No external coordinator).** C3 forbids an external party that *authors* the flow — one that decides, from its own state, what each actor does next. It does not forbid a *carrier*. The Seller’s own program makes the assertion — it records, as a sentence of its program, that it told the RewardEngine and what — and the bridge only carries that assertion onward. The journal names no transport at all: which broker or protocol carries the message is a deployment binding resolved outside the program, not a sentence in it (see below). A saga coordinator is the opposite: its state machine authors the flow, which then lives as a quasi-domain artifact outside every participant’s program, and folding it into a participant would dirty that domain with coordination it should not carry. Giving delivery a clean home in the carrier is what lets each actor’s journal hold only the causal record — the fact that it spoke. (The reproducibility lab exhibits this edge two ways: a single-process didactic run in which the bridge also stands in for the receiver’s own processing, and a separated-receiver run in which a pure in-process broker carries the envelope while the RewardEngine runs its own consumer — mapping the asserted message to a command it owns, journaling that command in its own journal, and acknowledging autonomously. The second instantiates the carrier/receiver split this clause turns on, with no party standing in for the receiver; in a deployment the delivered message is processed by the recipient’s own program in the same way.)

A note on entry [5]. The acknowledgment is not the RewardEngine writing into the Seller’s journal. The RewardEngine emits an ordinary result in its own journal; the transport routes an acknowledgment envelope back to the Seller; the Seller’s own handler processes that inbound envelope and records `tell ack ...` in the Seller’s journal. For the ack message the Seller is the *receiver*, so C1’s receiver clause applies unchanged — an actor records its own receipt in its own journal. The cross-actor edge is two messages, the assertion outbound and the ack returning, each narrated by the actor that processed it; neither actor writes to the other’s journal.

Delivery and correlation are separated by design. The journal is the durable record of what was asserted and what was acknowledged — correlation; redelivery, timeout, dead-lettering, and backoff belong to the transport. Routing, too, lives outside the program: the sentence names the addressee by its logical role (`to RewardEngine('rewards-1')`), and a deployment-level binding — not

a clause in the journal — resolves that role to a physical route (which transport, which topic). The journal therefore reads identically no matter which transport that binding selects; the wire never enters the actor’s voice. The delivery guarantee is whatever the chosen transport provides; the journal’s correlation role is invariant across all of them. If an acknowledgment is lost, the transport’s retry redelivers and the assertion’s identity (**once @order** — an author-chosen key, here the order the assertion carries — or a content hash when **once** is omitted) is the receiver-side deduplication key, so a redelivered tell does not duplicate the effect; a replayed journal re-executes the assertion — rebuilding the in-flight record — but its dispatch to the transport runs only on live execution, so replay does not re-send (the transport never sees a message from rehydration). A crash in the narrow window between journaling the assertion and dispatching it strands the envelope — journaled as issued, yet never handed to the transport; on recovery the rehydrated sender reconstructs that pending tell from the journal alone and the transport, the sole authority on delivery, testifies its fate, which the sender records as a verdict in its own voice (§8.5 G4) — so a crash no longer leaves the journal asserting a send that never landed: each issued tell is resolved (acknowledged, or marked unacknowledged by its addressee) or honestly left pending for the transport to settle, bounded by what that transport can still testify. The delivery model is thus at-least-once with identity-based deduplication; the consistency machinery it rests on — the now/deferred partition and its deduplication discipline — is developed in Paper 3, on which this construction builds.

The framework rejects **tell** outside **.Causation.Continue(...)**. A direct **tell** from a top-level command throws a runtime exception, with the error message pointing the developer at the correct surface. The cross-actor send is always a sentence in some Reaction’s Causation body, never a free-floating call.

One feature of the sentence repays attention: the message is named in the past — **PurchaseConfirmed**, a fact the Seller has already lived — never in the imperative. This is not a style convention but a consequence of the criterion: an actor can assert only what has already happened in its own history, and by the time the assertion is journaled the purchase is a settled entry. A *command* — the directive at a request edge, which may still be refused — is named in the present imperative (**ApplyReward**, **Confirm**); to name a command in the past would be a category error, an order to do what is already done. The tense is the surface mark of which kind of act the sentence is.¹ This is where the construct parts from the canonical actor frameworks: Akka exposes a single send verb (**!**, *tell*) for every message and accordingly advises naming them all in the past tense — folding the report of a fact and the issuing of a command into one form, which is right for the events and wrong for the commands, because one

¹The surface mark is English verb tense because the DSL is written in English; the distinction it carries — asserting a settled fact versus issuing a directive — is not itself tied to tense, and a language that grammaticalizes aspect rather than tense would mark the same distinction by other means. The claim is about the two kinds of act, not about tense as a linguistic universal.

verb cannot tell them apart. Here the two are kept distinct and the tense marks the difference — the assertion the sender makes about its own past, and the command the receiver runs in its own present, landing in two journals, in two voices.

8.3 Three implementations side by side

The same domain — the Seller confirms a purchase, the RewardEngine applies qualifying campaigns — is exercised in three styles. The code that distinguishes each style and the journals each produces are reproduced below without commentary. The interpretation follows in §8.4.

The saga and choreography implementations are the author’s own, written for this illustration and deliberately kept minimal (Appendix A); they are not independent or performance-tuned baselines. The illustration is accordingly structural, not competitive: it shows *where* each style records the joint cross-actor history, not that tell is faster, more resilient, or operationally simpler. Those are measurable dimensions this paper does not test (§1).

Its weight is definitional, not empirical. The alternatives are written in their canonical shape — an orchestrated saga whose participants are commanded by a coordinator that owns the flow, and a choreography whose actors publish to a shared bus — so the reader can see in journals what each pattern’s defining feature entails for where the joint history lands. The obvious objection — that a saga could have each participant record its own role — names exactly the tell move: a participant that records its cross-actor participation as a statement of its own program has, by that act, satisfied C2. The contrast is therefore definitional: the alternatives place the joint history outside any participant’s program because their defining feature *is* externalized coordination. The journals below illustrate that truth; they are not measured evidence that one pattern outperforms another, and a differently-built saga would not change where externalized coordination, by definition, leaves the record.

Style 1 — Saga (orchestrated)

A SagaCoordinator actor drives the workflow via direct commands to participants:

```
saga.Using("step = 'PurchaseRequested';").PerformCommand();

seller.Using("s = Seller();").PerformCommand();
seller.Using("s.purchase(@order, @date, @amount, @customer);")
    .WithParameters(p => {
        p["order",    typeof(string)]    = "ord-100";
        p["date",     typeof(DateTime)]  = new DateTime(2026, 9, 5);
        p["amount",   typeof(decimal)]   = 250m;
        p["customer", typeof(string)]    = "cust-42";
    })
```

```

        .PerformCommand();

saga.Using("step = 'PurchaseConfirmed';").PerformCommand();
rewards.Using(@"
    foreach (c in loyalty.Campaigns()) {
        if (c.Applies(@date, @amount) == true) { c.Reward(@order, @customer); };
    };")
    .WithParameters(p => {
        p["order",    typeof(string)]    = "ord-100";
        p["customer", typeof(string)]    = "cust-42";
        p["date",     typeof(DateTime)]  = new DateTime(2026, 9, 5);
        p["amount",   typeof(decimal)]   = 250m;
    })
    .PerformCommand();

```

```
saga.Using("step = 'RewardsApplied';").PerformCommand();
```

Every value crosses the DSL boundary as a typed parameter — `p["date", typeof(DateTime)]`, `p["amount", typeof(decimal)]`, the string ids too — bound through `.WithParameters(...)`, never string-interpolated into the script text. Each command is issued separately, so each lands as its own journal entry; a parameterized command is journaled as a message-action (define-then-invoke), and two campaigns added by the *same* parameterized command share one definition. The `RewardEngine` holds two campaigns — `C-newcomer` (min 10, qualifies for the 250 purchase) and `C-highroller` (min 1000, does not) — so the receiver's loop applies one of the two. After the run, the three journals contain:

SagaCoordinator's journal (3 entries):

```

[0] step = 'PurchaseRequested';
[1] step = 'PurchaseConfirmed';
[2] step = 'RewardsApplied';

```

Seller's journal (3 entries):

```

[0] s = Seller();
[1] (define of the message-action s.purchase)
[2] s.purchase('ord-100', 5/9/2026, 250, 'cust-42');

```

RewardEngine's journal (6 entries):

```

[0] loyalty = RewardEngine();
[1] (define of loyalty.AddCampaign)
[2] loyalty.AddCampaign('C-newcomer', 1/1/2020, 10);
[3] loyalty.AddCampaign('C-highroller', 1/1/2020, 1000);
[4] (define of the reward loop)
[5] foreach (c in loyalty.Campaigns()) { ... c.Reward('ord-100', 'cust-42'); };

```

Style 2 — Choreography (event-driven, no coordinator)

An event bus mediates the cross-actor handoff:

```
bus.Subscribe(ev => {
  if (ev.StartsWith("PurchaseConfirmed:")) {
    rewards.Using(@"
      foreach (c in loyalty.Campaigns()) {
        if (c.Applies(@date, @amount) == true) { c.Reward(@order, @customer); };
      };")
    .WithParameters(p => {
      p["order",    typeof(string)]    = "ord-100";
      p["customer", typeof(string)]    = "cust-42";
      p["date",     typeof(DateTime)]  = new DateTime(2026, 9, 5);
      p["amount",   typeof(decimal)]   = 250m;
    })
    .PerformCommand();
  }
});

seller.Using("s = Seller();").PerformCommand();
seller.Using("s.purchase(@order, @date, @amount, @customer);")
  .WithParameters(p => {
    p["order",    typeof(string)]    = "ord-100";
    p["date",     typeof(DateTime)]  = new DateTime(2026, 9, 5);
    p["amount",   typeof(decimal)]   = 250m;
    p["customer", typeof(string)]    = "cust-42";
  })
  .PerformCommand();
bus.Publish("PurchaseConfirmed:ord-100");
```

After the run:

Seller's journal (3 entries):

```
[0] s = Seller();
[1] (define of the message-action s.purchase)
[2] s.purchase('ord-100', 5/9/2026, 250, 'cust-42');
```

RewardEngine's journal (6 entries):

```
[0] loyalty = RewardEngine();
[1] (define of loyalty.AddCampaign)
[2] loyalty.AddCampaign('C-newcomer', 1/1/2020, 10);
[3] loyalty.AddCampaign('C-highroller', 1/1/2020, 1000);
[4] (define of the reward loop)
[5] foreach (c in loyalty.Campaigns()) { ... c.Reward('ord-100', 'cust-42'); };
```

Bus log (1 entry):

```
published: PurchaseConfirmed:ord-100
```

Style 3 — Tell (Puppeteer)

A Reaction on the Seller observes its own purchase and issues a **tell** from its `.Causation.Continue(...)` body:

```
seller.Reactions.DefineReaction("PurchaseFunnelToRewards")
  .Job().Company()
  .WithSharedHydration()
  .Seek("Purchase")
    .OnMatch("[s: Seller].purchase($order, $date, $amount, $customer)")
  .Causation.Continue(@"
    tell PurchaseConfirmed
      with @order, @date, @amount, @customer
      to RewardEngine('rewards-1')
      once @order;
  ");
```

```
seller.Reactions.Execute(); // once defined, the Reaction stands over the journal and applies
                             // to every matching entry –
it is not invoked per command. A Job is
                             // swept on demand (driven here for a deterministic test); a Cue
                             // Reaction runs continuously on its own thread.
```

```
seller.Using("s = Seller();").PerformCommand();
seller.Using("s.purchase(@order, @date, @amount, @customer);")
  .WithParameters(p => {
    p["order",    typeof(string)] = "ord-100";
    p["date",     typeof(DateTime)] = new DateTime(2026, 9, 5);
    p["amount",   typeof(decimal)] = 250m;
    p["customer", typeof(string)] = "cust-42";
  })
  .PerformCommand();
```

```
// the standing Reaction observes the purchase and asserts PurchaseConfirmed to RewardEngine;
// the bridge delivers and acks back
```

The tell's identity is `once @order` — the per-utterance key *is* the order id, so one compiled action issues a distinctly-identified assertion per purchase, and the ack returns under that same id.

After the run:

Seller's journal (6 entries):

```
[0] s = Seller();
[1] (define of the message-action s.purchase)
[2] s.purchase('ord-100', 5/9/2026, 250, 'cust-42');
[3] (define of the message-action PurchaseConfirmed)
[4] tell PurchaseConfirmed
```

```

        with 'ord-100', 5/9/2026, 250, 'cust-42'
        to RewardEngine('rewards-1')
        once 'ord-100';
[5] tell ack 'ord-100' from RewardEngine('rewards-1');

```

RewardEngine's journal (6 entries):

```

[0] loyalty = RewardEngine();
[1] (define of loyalty.AddCampaign)
[2] loyalty.AddCampaign('C-newcomer', 1/1/2020, 10);
[3] loyalty.AddCampaign('C-highroller', 1/1/2020, 1000);
[4] (define of the reward loop)
[5] foreach (c in loyalty.Campaigns()) { ... c.Reward('ord-100', 'cust-42'); };

```

8.4 Structural reading

The three implementations in §8.3 exercise the same logical flow but record it differently. Three pairs of journals were exhibited.

Saga: the SagaCoordinator's journal contains the workflow narrative — **PurchaseRequested** → **PurchaseConfirmed** → **RewardsApplied**. The Seller's journal records its local purchase only; it does not know it is part of a saga. The RewardEngine's journal records its local reward only; equally unaware. Three journals, three local stories. Only the coordinator's journal contains the joint history.

Choreography: no coordinator exists. The Seller's journal records the local purchase; the publish to the bus is invisible to the actor. The RewardEngine's journal records the local reward; the receipt from the bus is invisible to the actor. The bus's own log records the publish. No actor's journal contains the joint history; the bus log, an external infrastructure artifact, is the only place the cross-actor handoff is recorded.

Tell: the Seller's journal contains the purchase, the assertion (journaled as a typed message-action — defined once, then invoked), and the ack — six entries in all (each parameterized operation is a define plus an invocation) that constitute the joint history as a sequence of DSL sentences. The RewardEngine's journal records its local reward, as in the other styles. Under tell, the sender's journal alone reconstructs this cross-actor edge — the single hop from Seller to RewardEngine; §8.5 examines what happens when the chain is longer.

The three styles produce equivalent business outcomes. They differ structurally in where the cross-actor causal chain is recorded. The difference is not cosmetic. The saga coordinator, the event bus log, the distributed trace, and the workflow engine all exist to compensate for the absence identified in §3. Under tell, that program-level absence is gone: the sender's journal already contains the cross-actor narrative those patterns reconstruct elsewhere, so the apparatus built to recover it has nothing to recover.

The claim is about the *record*, not the wire. Tell does not eliminate the message-passing infrastructure: delivery remains operational, carried by whatever transport a deployment-level binding resolves the addressee to (§8.2), and the bridge that routes envelopes is part of it. What moves into the program is the causal record — the assertion the sender makes, and the acknowledgment of receipt — not the act of delivery. The contribution is the relocation of the narrative into the program, not the removal of transport.

Style	Joint history location	Audit path
Saga (orchestrated)	In the saga coordinator’s journal exclusively	Read the coordinator’s journal; participants’ journals are half-stories
Choreography (event-driven)	In no actor’s journal; the bus log is the only joint artifact	Merge participants’ journals via correlation id, plus the bus’s log
Tell (program-level cross-actor primitive)	In the sender’s journal as a sequence of DSL sentences	Read the sender’s journal entries [tell] and [ack] verbatim

In the first two styles, an additional architectural element is required to make the cross-actor flow observable as a narrative: a coordinator, a bus log, a trace backend, or a workflow program. In the tell style, no additional element is introduced. The narrative exists where the actor model already records program: the journal.

8.5 Property validation

Four property tests probe the claims of §5 and §6, all reproduced by the harness in `labs/lab04-tell` run against the public commit recorded under Code provenance. The first three (G1–G3) demonstrate that consequences claimed in §5 — auditability through external reconstruction, replay limited to single actors, cross-DC fragility — are reversed under tell, exhibiting properties that would be inaccessible if the assumption named in §3 were in force; they are intentionally chosen to mirror the compensating patterns of §6, each exhibiting a property that, under the assumption of §3, requires an external architectural pattern to achieve. A fourth (G4) turns to a sharper, adversarial question: whether the record stays honest when a tell never crosses.

G1 — Replay coherence (closes §5.2). The first test stages an in-flight tell: the envelope leaves the Seller but the bridge does not deliver it before the test asserts. A fresh actor instance with the same name, no shared transport, no live receiver, and no in-memory state replays the journal alone. Replaying the journal re-executes the tell statement; under replay it rebuilds the in-flight

record rather than re-dispatching, so the replayed actor reconstructs the cross-actor state from the journal alone — the program re-ran, with no duplicate send. The joint history exists as a program artifact and replay reaches it.

G2 — Cross-DC replication (closes §5.3). The second test replicates the Seller’s journal entry-by-entry to a fresh actor in an independent storage tier — the analogue of moving across data centers. The replicated actor, with no transport connectivity to the original receiver, reconstructs the dedup state from the replicated bytes alone. The cross-actor causal chain travels with the replication because it was always recorded in a place replication can carry.

G3 — Audit query (closes §5.1). The third test asks the audit question — *why did this happen?* — by reading the Seller’s journal directly. The cross-actor assertion is entry [4] (the message-action invocation); the acknowledgment is entry [5]. The cause-effect chain is reconstructed without distributed tracing, correlation IDs, or log aggregation.

The audit this closes is the *causal* one — what was said, to whom, when, and with what acknowledgment — because that is what the program records. It is not the *identity* audit — under which user, tenant, or request context an actor acted — which is operational metadata in the sense of §3 and lives, where a deployment keeps it, in a separate envelope keyed to the entries, not in the cross-actor causal record. Claim 8 is a claim about causal auditability; the two audits answer different questions and travel on different channels (Paper 5, §5.2).

Each of these properties is a consequence of cross-actor causation being recorded as program. Under saga, choreography, tracing, or workflow approaches — where it is not — the same properties are reachable only by consulting artifacts outside the participating actors.

G4 — Tell-fate recovery, the honest record under crash. G1–G3 stage an in-flight tell and show the *record* survives a crash, a move across data centers, an audit. A sharper question is what the record *says* when the tell never crosses. The send is journaled before the post-commit dispatch hands its envelope to the transport; a crash in that window strands the envelope — journaled as issued, never delivered, never acknowledged. Left unaddressed this is the one place the “sender’s journal alone” property could lie: the journal would assert a send that did not happen. The fourth test stages exactly this crash and rehydrates the sender. Replay reconstructs the set of *pending* tells — issued, neither acknowledged nor settled — from the journal alone; for each, the transport, the sole authority on delivery, testifies its fate and the sender records the verdict *in its own voice* — `tell 'ord-100' unacknowledged by RewardEngine` when the transport reports the envelope failed, the ordinary `tell ack ...` when it reports the envelope was delivered and only the acknowledgment was lost, and nothing while the fate is still in flight (the transport keeps ownership). The verdict names the addressee the Seller did not hear back from, not the broker that carried the message: A may say “*RewardEngine never acknowledged*” — a fact within its own universe — but not “*per the broker that carried it,*” which

is infrastructure it has no standing to assert. This is the dual of §6.4’s causal and message logging: there, the record of what crossed lives in a recovery layer beneath the program, consulted only by crash-replay machinery; under tell it is a sentence in the sender’s program, so recovering it is reading the journal, not excavating infrastructure. After recovery the sender’s journal no longer asserts a send that never landed: each issued tell is either resolved — acknowledged, or marked unacknowledged by its addressee — or honestly carried as still pending, for the transport to settle when it can. The guarantee is honesty about the outcome, not omniscience: a transport that cannot testify — one whose own record of a tell’s fate did not survive the failure — answers **InFlight**, and the tell stays pending rather than acquiring a fabricated verdict. Delivery stays the transport’s; the verdict is the journal’s.

The multi-hop limit — an adversarial case. G1–G3 are duals of the consequences named in §5, and so are chosen to show what tell does well. The honest counter-case is a longer chain. The case study is a single hop — the Seller tells the RewardEngine. Suppose instead a chain assembled the way every hop is: a Reaction on A tells B; B carries its own Reaction that observes the entry its receipt produces and, from that Reaction’s **.Causation.Continue** body, tells C. Nothing propagates the chain automatically — each actor opts in with its own Reaction (the envelope’s causal identifier is per-hop-local, not a threaded chain id), so the hops remain autonomous, as C3 requires. Each hop is recorded as program in the journal of the actor that originated it — A’s journal holds the A→B tell and its ack, B’s journal holds both its receipt of A and the B→C tell it issued, C’s journal holds its receipt of B. No single journal holds the whole A→B→C chain; reconstructing it end to end means composing A’s and B’s journals (linkable by envelope identifier across them). In this, tell inherits a distributed-history property of the kind §6.1 identified in choreography — but the difference is in kind, not in absence. Under choreography no participant’s journal records the cross-actor edge as program at all; the chain lives only in the bus log and is reconstructed from non-programmatic artifacts. Under tell every edge is a programmatic record in its sender’s journal, so the multi-hop chain is a composition of programs — each hop locally complete, auditable, and replayable on its own. Tell makes each edge programmatic and local; it does not centralize a multi-hop chain into one journal. The “sender’s journal alone” property (§8.4) is therefore a per-edge guarantee: it holds for the hops an actor originates, which is the whole chain only when the chain is a single hop.

8.6 Closing

The case study and the defensive tests together constitute the existence proof. The conditions of §7 are realizable: a system in which the cross-actor send is recorded as a sentence in the sender’s program, dispatched through the existing message-passing layer, and coordinated by no external party can be built and exercised. The instantiation in Puppeteer is one such system. The framework’s own production deployment, which the case-study domain is modelled on (§8.1),

is proprietary and is not exhibited here; the existence proof this paper offers is the runnable instantiation above. Other realizations of the conditions are possible (§7.5); the present section establishes only that at least one is.

9. Relation to previous work in this paper series

The journal exhibited in §8.3 — a sequence of DSL sentences in the sender’s program, recording the cross-actor send and its acknowledgment — required several structural preconditions to be a viable substrate. The journal had to be dense rather than porous: filled with operations, not with type-erased payloads. It had to record the operations with their parameter references intact, not with values inlined as literals. It had to maintain a discipline that separates immediate from deferred work, with a guardian for the boundary between them. Each precondition has been the subject of prior structural analysis in the present series.

Paper 1 introduces *porosity* — the representational sparsity that arises when domain state is recorded as serialized data structures rather than as programmatic operations. Anti-porosity is the design principle that the journal records what was said, not what was stored. Without that property, the entries the reader saw in §8.3 — `tell PurchaseConfirmed with ... to RewardEngine('rewards-1')` — could not be programmatic at all; they would be opaque payloads.

Paper 2 introduces *externalized parameters* as the structural precondition under which compilation, caching, and dense journaling become possible at all. Without externalized parameters, the journal could not record what was said with parameter references intact — values would be inlined as literals, or the script would lose its connection to the actor’s symbol table. The tell sentence in §8.3 carries `@order`, `@date`, `@amount`, `@customer` as references rather than literals; this paper is what makes that representation possible.

Paper 3 introduces the *partition* between immediate and deferred work, with Reactions as the guardian of the boundary. Paper 3 names the *Causation* plane — `.Causation.Continue(...)` — as the third surface a Reaction may touch (Paper 3 §6.5), and records (Paper 3 §6.8) that the cross-actor case an earlier draft had listed as a limitation is resolved by that plane, while leaving the primitive’s full treatment out of scope. The Reaction surface that Paper 3 establishes is the surface on which `tell` lives: the `.Causation.Continue(...)` body the reader saw in §8.3 is where the cross-actor send is permitted to appear. The present paper is the treatment Paper 3 deferred.

The series, taken together, defends a single architectural property under different framings:

Puppeteer preserves semantic continuity inside an actor. Tell preserves semantic continuity across actors.

The first sentence is the joint contribution of Papers 1 through 3: the structural conditions under which the journal can serve as a program-level substrate for what an actor does. The second sentence is the contribution of the present paper: the cross-actor extension of that substrate that the reader saw exhibited in §8.3.

Read as a sequence, the series traces a single thread — the operation. Paper 1 establishes that operations precede state; Paper 2, that they carry their own parameters and can author themselves; Paper 3, that they can generate further operations as reactions; and the present paper, that an operation can be addressed, as an assertion, to another actor. *Tell* is the cross-actor extension of the journal-as-program substrate the earlier papers build, not messaging added to it from outside.

Papers 1–3 are self-deposited preprints on Zenodo (Rivera, 2026a, 2026b, 2026c) and have not undergone peer review, as is the present paper. This paper rests structurally on them — C2 presupposes the substrate Papers 1 and 2 establish (§3, §7), and the Reaction surface on which `tell` lives is Paper 3’s (§8.2) — so it should be read as the latest in a preprint chain, not as resting on peer-reviewed foundations.

10. Conclusion

The actor model has, for fifty years, treated cross-actor causation as an operational concern of the runtime rather than as a construct of the program. The treatment was productive: actor systems became fault-tolerant, scalable, and reliable precisely because the separation between an actor’s program and the message-passing layer was operationally effective. Out of that productivity grew the ecosystem of patterns analyzed in §6 and exhibited in §8.3 — saga orchestrators, choreography buses, distributed tracing, workflow engines. Each compensates, in its own way, for the program-level absence the reader saw in the journals.

The present paper observes that the absence is not entailed by the actor model. It is a contingent design decision adopted by the canonical sources and propagated through the lineage as its default frame — productive enough that the question of whether the send could be a statement of the program was seldom raised. The conditions under which the absence can be removed — locality of writes, causation as program statement, no external coordinator — preserve every structural commitment of the actor model. A primitive that satisfies the three conditions, whether named *tell* or otherwise, makes the cross-actor send a sentence in the sender’s program; the apparatus that exists to reconstruct that

narrative after the fact has nothing left to reconstruct. The message-passing layer itself remains — delivery stays operational — but the causal record no longer lives outside every actor’s program.

The contribution of this paper is conceptual. The instantiation is the existence proof; the existence proof is the journal the reader saw in §8.3.

The question that opened the paper has, under tell, a different answer than the one the field gave by default: *where does the causal sentence live?* — in the sender’s own program, as a sentence the journal records and replay re-runs.

Fifty years of convention are not fifty years of necessity. The actor model does not need to be replaced. Its commitments — autonomy, message-based communication, isolation — survive the reformulation intact. What changes is the historical interpretation of what must remain invisible to the program. The interpretation, named explicitly in §3, examined for contingency in §4, traced through its consequences in §5, shown to require an entire ecosystem of compensating patterns in §6, and contrasted with the alternative in §8.3, is the only thing the present paper asks the field to revise.

Appendix A. Code references

The labs cited in §8 are publicly available in the Puppeteer codebase. The references below are to the test suite of the framework’s repository.

Cross-actor Puppeteer/EventSourcing/Follower/	primitive	surface	—
File	What it shows	Cited in	
Reaction.cs	Reaction class, Action terminator dispatch, replay-safe action execution	§7.2, §8.2	
Planes.cs	The three plane types (<code>ProgramPlane</code> , <code>CausationPlane</code> , <code>MetadataPlane</code>) and their property accessors	§8.2	
ReactionEngine.cs	Pattern matching surface, <code>.OnMatch(...)</code> , plane passthroughs	§8.2	

Reproducibility lab — **labs/lab04-tell/**

The runnable harness that reproduces every journal exhibited in §8 ships with this paper (and in **paper04-data.zip**), built against the public runtime commit recorded under Code provenance. The framework’s own end-to-end tests for the same scenarios live in the private fork and are not part of the public clone; this lab is the public, self-contained equivalent.

File	What it shows	Cited in
Program.cs	The harness, one method per scenario: the three-style side-by-side case study (orchestrated saga, event-driven choreography, tell — same domain, three journal locations of the joint history); G1 replay coherence; G2 cross-DC replication; G3 audit query; G4 tell-fate recovery across the crash window; and the negative gate for tell outside .Causation.Continue(...) . Prints each actor’s journal and a PASS/FAIL line per assertion.	§8.2, §8.3, §8.4, §8.5
LoyaltyDomainStubs.cs	Domain-side stubs for Seller , RewardEngine , Campaign — kept minimal so the focus remains on the cross-actor mechanism	§8.1

Code provenance

Source-code references in this paper resolve against the public Puppeteer repository at commit **37ad9cf** (2026-07-05). The snapshot is archived in Software Heritage under the following persistent identifier:

```
swh:1:dir:2b3ef9f68cf6e3a7a8e32ec73121122cc2ac74c5;  
origin=https://github.com/alvaroNCubo/puppeteer;  
anchor=swh:1:rev:37ad9cf44e7749f8fa0d2da8cc0c1093ad5a7e83
```

Inline references of the form `file.cs:NN` (e.g., `ActorHandler.cs:38`) resolve against this snapshot. A reader can construct a per-file SWHID by adding the qualifiers `;path=<path>;lines=<NN>` to the directory SWHID above. Future commits to the repository may renumber lines; the SWHID preserves the cited state independently of any future change to the repository or its hosting.

Acknowledgments

The author used large language models (including Claude and ChatGPT) as editorial assistants for language refinement, structural feedback, and literature navigation. All original ideas, terminology, theoretical constructs, and technical content presented in this work are solely the author's.

Appendix B. Bibliography

- Agha, G. (1986). *Actors: A model of concurrent computation in distributed systems*. MIT Press.
- Armstrong, J. (2003). *Making reliable distributed systems in the presence of software errors* [Doctoral dissertation, Royal Institute of Technology (KTH)]. https://erlang.org/download/armstrong_thesis_2003.pdf
- Bernstein, P. A., Bykov, S., Geller, A., Kliot, G., & Thelin, J. (2014). *Orleans: Distributed virtual actors for programmability and scalability* (Technical Report MSR-TR-2014-41). Microsoft Research.
- Elnozahy, E. N., Alvisi, L., Wang, Y.-M., & Johnson, D. B. (2002). A survey of rollback-recovery protocols in message-passing systems. *ACM Computing Surveys*, 34(3), 375–408.
- Fidge, C. J. (1988). Timestamps in message-passing systems that preserve the partial ordering. In *Proceedings of the 11th Australian Computer Science Conference* (pp. 56–66).
- Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.
- Hevner, A. R., March, S. T., Park, J., & Ram, S. (2004). Design science in information systems research. *MIS Quarterly*, 28(1), 75–105.
- Hewitt, C., Bishop, P., & Steiger, R. (1973). A universal modular actor formalism for artificial intelligence. In *Proceedings of the 3rd International Joint Conference on Artificial Intelligence (IJCAI-73)* (pp. 235–245).

- Hoare, C. A. R. (1978). Communicating sequential processes. *Communications of the ACM*, 21(8), 666–677.
- Hohpe, G., & Woolf, B. (2003). *Enterprise integration patterns: Designing, building, and deploying messaging solutions*. Addison-Wesley.
- Johnson, D. B., & Zwaenepoel, W. (1987). Sender-based message logging. In *Proceedings of the 17th International Symposium on Fault-Tolerant Computing (FTCS-17)* (pp. 14–19).
- Lamport, L. (1978). Time, clocks, and the ordering of events in a distributed system. *Communications of the ACM*, 21(7), 558–565.
- Lightbend. (n.d.). *Akka core: Interaction patterns* [Akka documentation]. <https://doc.akka.io/libraries/akka-core/current/typed/interaction-patterns.html> (Internet Archive snapshot, 2026-02-16: <https://web.archive.org/web/20260216052600/https://doc.akka.io/libraries/akka-core/current/typed/interaction-patterns.html>)
- Rivera, A. (2026a). Anti-porous architecture: a unified design principle for CQRS + Actor + Event-Sourcing systems. *Puppeteer Papers Series*, Paper 1 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20404863>
- Rivera, A. (2026b). Program–value separability: the structural precondition for compilation, caching, and dense journaling in a DSL runtime. *Puppeteer Papers Series*, Paper 2 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20740697>
- Rivera, A. (2026c). Reactions and the partition: opt-in eventual consistency in actor-native systems. *Puppeteer Papers Series*, Paper 3 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20792156>
- Vernon, V. (2013). *Implementing domain-driven design*. Addison-Wesley.
- Young, G. (2010). *CQRS documents*. https://cQRS.files.wordpress.com/2010/11/cQRS_documents.pdf

Operations from the substrate

TL;DR

The canonical distributed-systems literature treats four operations — deployment without downtime, cross-datacenter replication, backup, and offline catch-up — as separate disciplines. Each has developed its own mature apparatus: blue-green and rolling updates for deployment; change-data-capture and log shipping for replication; snapshots and point-in-time recovery for backup; message queues and convergence protocols for offline operation. Systems that require all four must implement and operate all four independently.

This paper shows that the separation is contingent. Under the journal-as-substrate condition established in prior papers of this series — where the journal records operations rather than states, is written in the same language as the program, and consists of compact named entries — the four disciplines reduce to a single primitive.

Deployment is replay to the present head. Replication is the same replay at another site. Backup is the corpus held without replay. Offline operation is replay after a delay. The canonical apparatuses become structurally subsumed: each reconstructs, from outside the program, work the substrate already performs by construction.

Six laboratories measure the substrate under conditions where these equivalences hold. A version handover reaches the present head at 532k entries/s; cross-site transfer carries 33–96.7 bytes per event against payloads 3–9× larger; passive consumers converge across heterogeneous backends (FileSystem, MySQL, SQL Server) to identical in-memory state in 18/18 cells; local append latency runs $\approx 3\times$ below co-located RDBMS engines under matched durability. §7 shows how four runtime primitives — each present for independent structural reasons — compose into the operations the canonical literature

implements as separate subsystems.

For a journaled program, deployment is replay, replication is sharing history, backup is copying the program, and offline operation is delayed replay.

1. Introduction

This paper makes an analytic theory contribution. It identifies a structural property whose absence the canonical distributed-systems literature has assumed without naming it as a single construct; derives the equivalences that follow when the property is present; and presents an instantiation showing it is realisable in production. The contribution is conceptual; the instantiation serves as an existence proof rather than as the substance of the claim. The genre is the one Gregor (2006) names *theory for analyzing* (Type I): it introduces a construct that lets the phenomenon be described and classified, with the measurements reported in §5 exhibiting the régime under which the construct holds rather than evaluating an artefact. The Hevner-style design-science *evaluation* of the artefact (Hevner, March, Park, & Ram, 2004) — a comparative systems study of throughput, availability, or fault tolerance against existing stacks — is a separate undertaking this paper does not attempt.

This is neither a systems paper proposing a new mechanism nor a survey of existing apparatuses. It does not introduce a deployment tool, a replication protocol, a backup format, or an offline-operation pattern. It examines a structural property that prior literature has implicitly assumed absent and shows that, under its presence, four operational disciplines of distributed systems reduce to a single primitive. In this genre, contribution is measured by the precision of the construct, the validity of the equivalences derived from it, and the realisability of the instantiation. Sections 5 and 7 exhibit the instantiation and report measurements establishing the régime in which the equivalences are observable.

Distributed systems must address four operational dimensions: release without downtime, cross-datacenter replication, recovery after data loss, and periods of disconnectivity (offline operation). The canonical approach treats these as separate disciplines, each with distinct apparatus and cost semantics: blue-green and rolling updates; change-data-capture and log shipping; snapshots and point-in-time recovery; message-queue buffers and convergence protocols. These bodies of practice are mature and widely adopted, yet they evolved largely independently. Systems that require all four must implement and operate four distinct stacks, each with its own engineering and cost.

The fragmentation is rarely questioned because it is familiar. But the four disciplines can be restated more simply: they vary when and where a program is read — at the present head and in place (deployment), at the present head

at another site (replication), held at another site without reading (backup), and read after a delay (offline catch-up). They differ along two axes — temporal position and spatial position — over a single artefact. If that artefact were the program itself, recorded as operations rather than states, the four disciplines would be four readings of the same primitive: reading the program in order.

This paper names that artefact and the equivalences it admits. The construct introduced is the journal as the substrate of the program (in the sense established in Paper 2 §1.2: the pair of domain library and journal of invocations). Under the anti-porosity condition of Paper 1 and the separability condition of Paper 2, the journal records operations rather than states, is written in the same language as the program, and reduces each entry to a compact reference to a parametric verb and its arguments. The journal, so constituted, is not a record kept by the program; it is the program written out over time.

Four equivalences follow. Deployment is replay to the present head. Replication is the same replay at another site. Backup is the corpus held without replay. Offline operation is replay after a delay. The canonical apparatuses above are responses to the absence of the substrate condition; under its presence they become structurally subsumed by the substrate.

§2 traces the genealogy of the four disciplines as parallel bodies of literature. §3 names the assumption that produced the fragmentation. §4 states the substrate theorem and the four equivalences. §5 develops each equivalence and reports measurements exhibiting the régime under which it holds. §6 examines four canonical apparatuses — blue-green deployment, change-data-capture, snapshot backup, and message-queue buffer — and shows that each reconstructs, from outside the program, work the substrate already performs. §7 presents an instantiation in production. §8 relates the construct to prior papers in this series. §9 addresses counter-arguments. §10 concludes.

2. Genealogy: four separate operational disciplines

For decades, four operational concerns — deployment downtime, cross-datacenter replication, backup, and offline operation — have evolved as separate bodies of literature and practice. Each developed its own vocabulary, its own apparatus, and its own cost model. The genealogies traced below are not exhaustive; they are sufficient to establish that the canon treats these four as distinct disciplines, each with a dedicated mechanism, and that the resulting fragmentation of the operational stack has not been challenged within any of the four traditions.

2.1 Deployment downtime

The earliest deployment practice in production systems was the cutover: stop the running service, swap binaries, start it again. Downtime was accepted as the cost of release. The first patterned response was articulated by Fowler in *BlueGreenDeployment* (Fowler, 2010), which named a procedure in which “you ensure that you have two production environments, as identical as possible” and traffic is switched from one to the other at release time. The mechanism is environmental: two parallel deployments, one apparatus for routing, one for binary distribution.

The continuous-delivery literature ratified the pattern. Humble and Farley (Humble & Farley, 2010) generalized blue-green into a broader release-engineering discipline in which deployment is a build-pipeline output rather than an operator action, and downtime is a regression to be detected rather than a default to be accepted.

The modern instantiations are canary release (Sato, 2014) and the orchestration-mediated rolling update (Burns et al., 2016; Kubernetes, n.d.). Canary release introduces partial traffic exposure before full cutover, treating the deployed binary as a hypothesis tested against a fraction of production load. The Kubernetes rolling update replaces pods incrementally under a readiness-probe gate, making the deployment apparatus a property of the orchestrator rather than of the application.

A parallel strand of the deployment canon addresses schema and data migration. Ambler and Sadalage (Ambler & Sadalage, 2006) formalized *database refactoring* as a discipline in which every change to the data store is recorded as a forward-and-backward pair of scripts versioned alongside the application code. The pattern is instantiated by migration frameworks such as Rails Migrations, Liquibase (Liquibase, n.d.), and Flyway (Flyway, n.d.), which apply versioned scripts to the data store at release time. The migration apparatus is separate from the traffic-switching apparatus: it runs adjacent to — or before — the cutover, mutates the store in place, and produces a schema or data shape that the next binary version expects.

Across these mechanisms the shape is constant: a separate apparatus is built whose function is to absorb the change — of binary, of schema, or both — without interrupting the service. The apparatus is external to the program; the program is unaware that it is being replaced or migrated.

2.2 Cross-datacenter replication

Replication across datacenters originates in database engines. MySQL replication (MySQL, n.d.) established the canonical primary/replica topology in the 1990s: the primary writes a binary log of statements or row-level events, and replicas tail that log. The replication mechanism lives inside the database engine; the application program is unaware that replication is occurring.

The pragmatic maturation came with change-data-capture (CDC). Debezium (Debezium, n.d.) and Maxwell (Maxwell, n.d.) read the database’s binary log and re-emit each change as an event on an external transport, typically Kafka. CDC inserts an apparatus between the database and downstream consumers: a process that watches state changes in one system and reconstructs them as event streams for others. The reconstruction step is necessary precisely because the originating system does not emit events — it emits state changes that CDC observes and translates.

The modern instantiations extend the apparatus topologically. Cassandra’s multi-datacenter replication (Lakshman & Malik, 2010) embeds the replication topology in the storage engine itself, allowing writes accepted in any datacenter to propagate to peers under a configurable consistency level. Kafka MirrorMaker (Kreps, Narkhede, & Rao, 2011) generalizes the pattern for log-structured systems: a dedicated process consumes from one cluster and re-produces to another, with offsets tracking progress per topic.

Across these mechanisms the shape is constant: cross-datacenter replication requires a separate apparatus — engine-internal, engine-adjacent, or fully external — whose function is to observe the changes happening in one location and reproduce them elsewhere.

2.3 Backup

The backup canon predates the others. Agent-based backup (Bacula, n.d.; Veeam, n.d.) developed in the 1990s and 2000s as a discipline in which a dedicated process periodically reads the production data store and writes a copy to an archival medium. The agent is external to the application; backup occurs at the storage layer.

The pragmatic maturation introduced snapshot semantics. The WAFL filesystem (Hitz, Lau, & Malcolm, 1994) established the modern primitive: a point-in-time copy taken at the filesystem level, exploiting copy-on-write so that the snapshot is cheap and consistent with respect to the moment of capture. Database engines adopted the same pattern under the name *point-in-time recovery* (PostgreSQL, n.d.), in which continuous archival of the write-ahead log permits reconstruction of any committed state within the archived range.

The modern instantiations move the apparatus to object storage. Cloud-era backup (AWS, n.d.; Azure, n.d.) treats archival as a property of the storage tier rather than of a dedicated agent: lifecycle policies move data to colder tiers automatically, and recovery is a matter of retrieval from the archive.

Across these mechanisms the shape is constant: backup is a separate apparatus that produces a copy of state and is exercised at recovery time to reconstruct that state. The backup procedure observes the data store from the outside; the program being backed up does not participate.

2.4 Offline operation

The offline-operation canon is the most heterogeneous of the four, because it spans network-layer, transport-layer, and application-layer treatments of disconnection. The early precedent is store-and-forward messaging: UUCP and, later, SMTP queue undeliverable messages locally and retry until the remote endpoint is reachable. Disconnection is absorbed by a queue at the boundary of the system.

The pragmatic maturation generalized the boundary queue into a programmable apparatus. RabbitMQ (RabbitMQ, n.d.) introduced dead-letter exchanges to capture messages whose delivery is impeded, deferring them for retry or human inspection. Kafka (Kreps, Narkhede, & Rao, 2011) generalized the pattern further: consumer offsets are explicit, durable, and per-partition, so a consumer that goes offline resumes from its last committed offset when it returns. The queue is no longer a remediation mechanism for failed delivery; it is the substrate over which the consumer’s progress is tracked.

The modern instantiations carry the pattern into application-level state. Eventually-consistent stores (DeCandia et al., 2007) treat each replica as a local source of truth that converges with peers when connectivity permits. The Dynamo paper established the design vocabulary — vector clocks, hinted handoff, read-repair — that subsequent systems — Cassandra, Riak, and Redis replication (Redis, n.d.) — instantiate in varying degrees.

Once again, the shape is constant: offline operation requires a separate apparatus — a queue, an offset, a convergence protocol — whose function is to absorb the period during which the system cannot reach its counterpart.

2.5 Where the operational complexity lives

The four disciplines surveyed above can be compared along a single matrix. Each row names a discipline; each column names a dimension of the apparatus the canon developed to handle that discipline.

Discipline	Where the complexity lives	Apparatus	Cost
Deployment	Outside the program, in the release pipeline	Two-environment switch (blue-green); incremental pod replacement under a readiness gate (rolling update); fractional traffic exposure (canary); versioned schema and data migration scripts (Liquibase, Flyway, Rails Migrations)	Duplicate production environment; orchestration layer; manual or automated cutover procedure; forward-and-backward migration pairs
Cross-datacenter replication	Inside the storage engine, or in an adjacent CDC layer	Binary-log shipping (engine-internal); change-data-capture (engine-adjacent); multi-DC topology (engine-aware)	Replication lag; reconstruction overhead in CDC; consistency-versus-availability configuration surface
Backup	Outside the program, in a dedicated archival pipeline	Agent-based copy; filesystem or engine snapshots; continuous log archival (PITR); object-store lifecycle	Recovery-time objective and recovery-point objective; storage-tier cost; periodic exercise of the recovery procedure
Offline operation	At the boundary between system and counterpart	Persistent queue with retry semantics; durable consumer offsets; convergence protocol over replicated stores	Buffer sizing; offset management; eventual-consistency reasoning at the application layer

Three observations follow from the table. The location of the complexity differs across rows: the deployment apparatus lives in the release pipeline, the replication apparatus lives inside or beside the engine, the backup apparatus lives in a separate archival pipeline, and the offline-operation apparatus lives at the system’s boundary. The mechanism in each row was developed without reference to the others: the blue-green literature does not cite the CDC literature; the snapshot literature does not cite the consumer-offset literature. The four literatures evolved in parallel, each solving what it believed to be a different problem. And the cost in each row is paid separately: a system that wishes to be zero-downtime, cross-datacenter, recoverable, and tolerant of disconnection budgets independently for each of the four columns.

The canon, taken as a whole, presents these four as four different problems requiring four different apparatuses. §3 names the assumption that made this fragmentation appear necessary.

3. The assumption named

What §2 shows to be continuous across four separate bodies of literature can be stated as a single structural assumption:

Operational concerns of a system are addressed by apparatus around the program rather than by the program itself.

Across the four disciplines, the response to each operational concern is the introduction of a mechanism external to the program. Blue-green deployment builds a parallel environment to absorb the cut-over. Change-data-capture inserts a translation layer between stores to carry replication. Snapshot agents observe the data store from outside and write captures to archival media. Message-queue brokers hold messages at the system boundary to absorb periods of disconnection. In each case, the mechanism is not a sentence of the program it serves. It observes the program from outside, maintains a parallel artefact, and is operated independently of the program’s own execution.

The assumption embedded in these practices is that operational concerns require such external apparatuses, each developed independently to address a specific condition.

The alternative this assumption leaves unexplored is that the program might address these concerns by varying how it is read — by changing when or where the program is replayed, without leaving the program’s own representation. Under this alternative, the four disciplines become four readings of a single primitive. The question becomes structural: under what condition would the program admit such readings?

§4 names that condition and the four readings it permits. The apparatuses

surveyed in §2 remain appropriate responses in their original context; what this paper shows is that they arise from the absence of the substrate condition, which, if present, subsumes the need for each apparatus.

4. The substrate theorem

4.1 The substrate

The first two papers in this series established two structural properties of the journal that this paper takes as preconditions. Paper 1, §3 named *anti-porosity*: the journal records the operations of the program rather than the states they produce, and admits no representational sparsity at the boundary between domain code and persisted form. The journal is therefore *code-as-data* — entries are executable sentences in the same language the program is written in, in the narrow *script-form* sense Paper 1 establishes rather than Lisp’s strong homoiconicity — and dense — every effect that crosses the boundary is named in that language, not summarised as state delta. Paper 2, §1.2 and §3 named *separability*: programs that are parametric in their arguments admit a compiled form in which the operation’s body is cached under an identifier and the journal entry reduces to that identifier plus the call’s arguments. Together the two conditions yield a journal whose entries are compact, named operations rather than serialised states.

Under these two conditions the journal is not a record kept by the program; it is the program written out over time. Each entry is a sentence the program has uttered; the journal is the corpus of those sentences in the order they were uttered. The naming move of this paper is to call this artifact *the substrate of the program*: the journal does not store the program’s outputs, it constitutes the program’s existence over time. The program and the journal are the same artifact under two readings — one synchronous, one diachronic. Under the synchronous reading, the runtime is the program. Under the diachronic reading, the journal is the program.

One further property, established in Paper 1 §4.2, is what makes the substrate *replayable without divergence*, and every equivalence of §4.2 rests on it: the arguments in an entry are *values captured at execution time, not expressions re-evaluated on read*. A nondeterministic or state-derived argument — a generated identifier, a clock reading — is resolved once, on the writing path, and the resolved value is frozen into the entry; a later read reuses that value rather than recomputing it. (*Verification: the capture-at-execution of the Eval modifier, Paper 1 §4.2 — on first execution the value is captured as a literal-assigning script* `name = (type)(value)`; *persisted in the entry, and on replay that literal is re-assigned rather than recomputed*; *Parameter.cs:33–36, 163–224, 228–247.*) Were an argument re-evaluated on read, replaying it would diverge from the run that wrote it, and deployment, replication, backup, and offline catch-up

would each reconstruct a different program than the one recorded.

4.2 The theorem statement

The theorem of this paper can be stated as a single proposition.

Given a program whose execution is recorded as a journal satisfying the anti-porosity condition of Paper 1 and the separability condition of Paper 2 — that is, a journal that is code-as-data, dense, and made of compact named operations rather than serialised states — the following four equivalences hold.

E1 — Deployment is replay. A new version of the runtime, started from the substrate, reconstructs the program by reading the journal sequentially and applying each entry. Nothing about the new version’s instantiation differs in kind from what the previous version did at every entry up to that point; the act of bringing a process to a state in which it can serve requests is the act of replaying the substrate up to its present head. Deployment therefore differs from steady-state operation only in the position of the read cursor over the same corpus.

E2 — Replication is sharing history. A second instance of the program, running on a different machine or in a different site, reconstructs the program by reading the same sequence of entries the first instance wrote. The act of replicating is the act of transmitting the substrate from one site to another and replaying it at the destination; no state extraction, no schema mapping, and no change-data inference enter the act. Replication therefore differs from deployment only in the site at which the same corpus is replayed.

E3 — Backup is copying the program. A consumer that reads entries from the substrate and persists them locally — without instantiating the runtime that would replay them — holds a copy of the program in the form in which it was written. The act of backing up is the act of duplicating the substrate; recovery is the act of replaying the duplicate. Backup therefore differs from replication only in whether the destination chooses to replay the same corpus or merely to hold it.

E4 — Offline operation is delayed replay. A consumer that is unreachable when entries are written, and that catches up by reading them later from a persistent buffer or directly from the source, applies the same replay to the same sequence as a consumer that received the entries in real time. The act of operating offline is the act of replaying the substrate after a delay; the delay does not enter the program, only the cursor of the consumer. Offline operation therefore differs from replication only in the elapsed time between the write and the read of the same corpus.

The four equivalences are not parallel restatements of a single mechanism. Each addresses one classical operation of distributed systems (§2) and shows that the operation, as the canon presents it, dissolves into a single primitive — replay

— once the substrate’s properties are taken as given. E1 addresses the time axis (a new version replacing an old one in place); E2 addresses the space axis (a new site running the same program); E3 addresses the storage axis (a consumer holding the corpus without running it); E4 addresses the connectivity axis (a consumer reading later than the writer wrote). The four together exhaust the operational variations that the canonical literature in §2 treats as separate disciplines. The operational specifics of how each replay is realised — handoff, transmission, passive ingestion, catch-up — are taken up in §5; §4 confines itself to the structural claim.

4.3 Apparatus: the substrate and its four replay arrows

The shape of the theorem can be rendered as a simple spatial arrangement. No new content is introduced; the arrangement makes the four equivalences of §4.2 visible at a glance, with the substrate at the centre and the four conditions arrayed around it.

The substrate — the journal as program — sits at the centre. From it radiate four equivalences, each corresponding to one classical operational discipline:

- **E1 — Deployment:** replay to the present head.
- **E2 — Replication:** replay of the same corpus at another site.
- **E3 — Backup:** the corpus held without replay.
- **E4 — Offline operation:** replay after a delay.

The vertical pair (E1, E4) varies the temporal position of the replay — at the present head, or after a delay. The horizontal pair (E2, E3) varies the spatial position of the corpus — replayed at a different site, or merely held at one. The canonical literature in §2 names four disciplines by looking at the destinations; viewed from the substrate, only replay occurs, under four conditions on when and where.

5. Consequences: four equivalences as variants of replay

Each of the four sub-sections that follow develops one equivalence of §4.2. The sub-section opens with the operational realisation of the equivalence — the structural intervals through which the program is read under that variant — and follows with a measurement that exhibits the régime in which the equivalence holds. Each lab is designed to isolate the structural content of its equivalence, not to optimise throughput; performance figures are reported as evidence of the régime, not as benchmark claims.

5.1 E1 — Deployment is replay

§4.2 stated the equivalence: a new version of the runtime, started from the substrate, reconstructs the program by reading the journal sequentially and applying each entry. §5.1 develops the operational realization of this equivalence — the handover window during which an instance about to become authoritative reads the substrate up to the present head, and the precise sense in which that window is replay and nothing else.

The handover window

A new version of the runtime begins as a *follower*: a process that reads from the same substrate the live instance writes but does not yet accept requests. The follower’s task is to reach the present head of the substrate. Once it has — and once a brief synchronization tail has cleared — the follower is admitted as authoritative and the previous version is released. Three structural intervals comprise the window:

1. **Bulk replay.** The follower opens the journal from entry zero (or from a checkpoint, if the substrate is partitioned with one) and applies each entry to its in-memory state in the order written. The state reconstructed at the end of bulk replay equals, entry-for-entry, the state the previous version reached by writing those same entries.
2. **Lock and synchronization tail.** The previous version is paused briefly so that no further entries are written. The follower reads the small residue of entries that arrived between the start of bulk replay and the pause — the *handover tail* — and applies them. The state at the end of the tail equals the state the previous version held at the moment it was paused.
3. **Gate flip.** The follower is admitted as the authoritative instance. From this point on the substrate’s writer is the new version; the previous version is released.

The operational specifics of each interval are realized by a single primitive in the instantiation (§7.1 develops the gate mechanism); what §5.1 establishes is that the *content* of the window is replay. No state extraction, no schema migration, and no inference step enters the operation. The new version becomes authoritative by reading what the previous version wrote, in the order it was written, until the read cursor reaches the position the previous version had left it at.

What the substrate measures

L1 (lab paper05-lab1-redblack-replay-time @ 99e8202) times the handover window over four journal sizes ($N \in \{1k, 10k, 100k, 1M\}$ entries) under the compact-action régime named in Paper 2. The measurement brackets the inner bulk-replay interval and the handover tail separately with the public `Stopwatch` API — around `follower.Start(asFollower: true)` and the lock/synchronisation tail — so that the structural content of the window —

the act of replay — is reported in its own right rather than aggregated with orchestration overhead.

Under the compact-action régime with `AlwaysCompiled` compilation, the substrate replays N entries at a sustained rate of **532 thousand entries per second** (p50, $N = 100\,000$), completing the bulk replay in **187.8 ms** (p95 = 190.6 ms). The rate is reached as the runtime amortises between $N = 10\,000$ and $N = 100\,000$ entries and **continues to climb through the $N = 1\,000\,000$ anchor to 1.76 million entries per second** — fixed startup cost spread over more events. The handover tail is below 0.3 ms at every N measured. **Bulk replay accounts for more than 99.8 % of the deploy window in every cell.** The replay rate is sustained — indeed it improves — across a tenfold journal range; were it linear in expanded state rather than in compact-event count, the rate would fall as state grew. It does not.

Three properties of the measurement deserve naming. First, the substrate cost of deployment is bounded by compact-event count, not by expanded state — replaying ten times the journal takes no more than ten times the time, and here less: the per-event rate climbs as fixed startup cost amortises over more entries. Second, the orchestration tail is empirically negligible at every measured size; the deploy window *is* the replay window to within a fraction of one percent. Third, the régime under which this holds is the régime Paper 2 already named the canonical one: parametric verbs encoded as `ActionId` plus arguments, not literal scripts. A journal of literal scripts would replay slower in proportion to script length; this paper concerns the compact-action substrate named in Paper 2.

Apparatus

The handover window can be rendered as a sequence between four roles. The leader is the live instance accepting writes; the journal is the substrate; the follower is the new instance approaching the present head; the gate is the predicate that decides which instance is authoritative at any moment.

```
sequenceDiagram
    autonumber
    participant L as Leader (authoritative)
    participant J as Journal (substrate)
    participant F as Follower (new version)
    participant G as Gate

    Note over F,G: Follower starts; gate closed for follower
    F->>J: open from entry 0 (or checkpoint)
    loop bulk replay
        J-->>F: next entry
        F-->>F: apply entry to in-memory state
    end
```

```

Note over F: follower at near-head
F->>L: request pause
L->>L: stop accepting writes
loop handover tail
    J-->>F: residual entry written before pause
    F->>F: apply
end
Note over F: follower at present head
F->>G: open for follower
G->>L: close for leader
Note over L: leader released; follower authoritative

```

The diagram contains no operation other than reading from the journal and applying. The leader writes nothing between pause and release; the follower reads from the same artifact at the same offsets the leader was about to read had it continued. The gate’s flip is bookkeeping: it nominates which instance is authoritative; it neither produces nor consumes substrate content.

One canonical exception

The equivalence developed here covers the in-process deployment of a single program. One canonical case lies outside its scope: a coordinated cut-over of a producer/consumer pair on an external message broker (typically Kafka) when the message format itself changes between versions. There the two endpoints must release together, because the substrate they share is the broker’s topic — a contract between independent programs — rather than the journal of either. The construct of this paper does not deny this case; it narrows its scope. Within a journaled program, deployment is replay; across two journaled programs that share a broker topic with a changing schema, the cut-over remains a separate concern. Paper 3, which develops the partition primitive this paper’s Reactions build on, named this exception in its claim 8 on in-actor continuity; it is preserved here without amendment.

To the next equivalence

§5.1 has shown that bringing a new version of a program to a state in which it can serve is replay of the substrate up to the present head, and that the orchestration around that replay is residual at every measured size. The equivalence varies the position of the read cursor in time — at the present head of the same substrate. §5.2 varies the position of the substrate in space: a second instance, at another site, reconstructs the same program by reading the same sequence of entries the first instance wrote.

5.2 E2 — Replication is sharing history

§4.2 stated the equivalence: a second instance of the program, running on a different machine or in a different site, reconstructs the program by reading

the same sequence of entries the first instance wrote. §5.2 develops the operational realization in two parts. First, that the bytes streamed between sites are structurally dense — the substrate’s wire form carries named operations, not serialized states, by construction. Second, that a consumer reading those bytes reaches the same program as the writer by replay alone, without recourse to a coordinator. The transport across which the bytes travel is treated as a witness — an off-the-shelf service that satisfies the substrate’s requirement on consumer-pull delivery.

The compact wire

Paper 2, §2.3 (*dense journaling as reference*), established that under separability the journal entry of an invocation reduces to a reference to the parametric verb’s definition plus the arguments of the call. The verb’s body is not present in the entry; it lives once, named, in the corpus’s vocabulary. Replication across datacenters is the transmission of those entries, not of states; the wire encodes the program as it was uttered, not as it was applied.

The entries so transmitted are the domain program: operations and their argument values. Operational metadata *about* the program — the identity, origin, or request context under which an operation was recorded — is, where a deployment maintains it, a distinct channel keyed to the same entries, not part of the entries themselves. It rides the same substrate mechanisms (replicated alongside the corpus, copied with the store), but it is neither a domain value nor an operation, and the equivalences of this paper concern the program the journal constitutes, not the audit envelope a deployment may keep around it. The distinction matters at the cross-actor boundary, where the two channels part ways (Paper 4, §8.2): a value the receiver’s domain needs crosses as an argument; the sender’s envelope does not cross at all.

L2 (lab `paper05-lab2-bytes-per-event @ 99e8202`) instruments the codec on the substrate’s wire surface, comparing the compact encoding (`ActionId` + parameters) against the literal encoding (script DSL + parameter assignments) at three tiers of verb richness, with and without an opportunistic `gzip` on the literal form.

Bytes-per-event in compact form is 33 at the trivial-arithmetic tier (68-byte script body), 33 at the branching-arithmetic tier (99-byte body), and 96.7 at the production-shaped tier (681-byte body). The literal form at the same three tiers is 118, 150, and 910.7 bytes. The ratio between literal and compact rises from **3.6× at tier 1** to **9.4× at tier 3**. Applying `gzip` to the literal form closes the gap only to **3.2× / 3.9× / 4.0× at tier 3**. The structural saving cannot be recovered by an opportunistic compressor, because the script body is not present in the compact form at all — it lives once in the definition record (909 bytes at tier 3, amortized to ≈ 0.91 bytes per invocation over 1 000 calls) and is referenced by a 4-byte `ActionId` thereafter.

The cross-validation against Paper 2 Lab 4 sharpens the claim. Tier 3’s compact

payload is 67.7 bytes here; Paper 2 Lab 4 measured 67 bytes for the production verb that the synthetic tier-3 stand-in calibrates against. The two numbers agree within rounding. At full production scale, Paper 2 Lab 4 reports the literal-to-compact ratio reaches $20\times$ over a thousand invocations of a real ticket-purchase verb. The $9.4\times$ headline here is a conservative lower bound under the lab’s literal-projection; the production figure is reported in Paper 2 Lab 4.

Three properties of the measurement deserve naming. First, density is structural rather than algorithmic — compression operating after the fact cannot recover what naming achieved by construction. Second, the saving grows with verb richness; a domain whose operations are named transitions (Paper 2 §3) replicates per event at a width that does not scale with the operation’s internal complexity. Third, replication across datacenters, in this régime, is event streaming over a substrate that does not carry state by construction — the wire is the program as uttered, not as expanded.

The push-to-pull inversion

The substrate is written by the producer and read by the consumer; the role of the transport between them is to hold the corpus durably and to deliver entries at the consumer’s pace. The producer does not know the consumer’s rate; the consumer’s progress is measured against the substrate’s own offsets, not against any signal from the producer. A consumer that lags absorbs the lag in its own cursor; the producer’s write path is unaffected.

This inversion is not a property of any particular delivery service. It is a property of the substrate. The producer’s journal is the source of record; any consumer downstream is a reader of that record. An off-the-shelf webhook delivery service satisfies the consumer-pull requirement without adding any structure the substrate did not already license: it accepts an inbound POST from the producer’s local journal callback, holds it durably, and offers it to subscribers at the rate they fetch. The fact that an off-the-shelf service suffices is the relevant observation; the service is a witness to the substrate’s structural sufficiency, not an extension of it.

This is the *passive-consumer* arrangement: the consumer reads from a durable buffer at its own pace, and no party coordinates the two sides. It is distinct from a synchronous replicated-actor topology — one instance authoritative, others following under an elected role — which is a different mechanism with different coupling and is not the arrangement measured here. E2 is the claim that the *passive* path needs no coordinator, not that every replication topology dispenses with one.

The shape can be rendered between five roles: the primary instance at site A, the primary’s journal, the delivery service, the consumer at site B, and the consumer’s journal.

sequenceDiagram

```

autonumber
participant P as Primary (site A)
participant Ja as Journal A (substrate)
participant T as Delivery (off-the-shelf, e.g. SVIX)
participant C as Consumer (site B)
participant Jb as Journal B (substrate)

P->>Ja: append entry n
Ja->>P: callback (entry n written)
P->>T: post entry n
Note over T: held durably; offered at subscriber's pace
C->>T: fetch next
T-->>C: entry n
C->>Jb: append entry n
C->>C: apply entry to in-memory state

```

The diagram contains no synchronous coupling between producer and consumer. The producer's append-and-post is local; the consumer's fetch-and-apply is local; the durable buffer between them is the delivery service. The substrate at A and the substrate at B contain the same entry at the same offset — the two journals are bit-equal under the measurement examined next.

The symmetric consumer

§4.2 stated that replication “differs from deployment only in the site at which the same corpus is replayed.” The structural content of this claim is that the consumer at site B, running the same binary as the producer at site A, derives the same program from the corpus it reads as the producer derived from the corpus it wrote. No coordinator is required to synchronise the two sites; the substrate is the coordinator. This property removes the need for any external coordinator between sites: coordination is implicit in the shared corpus rather than enacted by a privileged runtime.

L3 (lab `paper05-lab3-inproc-symmetric @ 99e8202`) measures this property at the level of Reactions — the partition primitive of Paper 3 — under two terminators, `Emit` and `MarkAsSkip` (Elide). Two instances of the same actor binary consume the same event stream; the lab compares, at the byte level, both the callback tuples each instance's Reactions produce and the journal segments each instance writes. Across 6 cells ($N \in \{100, 500, 1\,000\} \times K = 2$ repetitions), **all 6 cells produce 0 callback byte diffs and 0 journal-segment byte diffs**. The instance at site B reconstructs the Reaction output entirely from the journal prefix it reads.

The discipline that makes this work without a coordinator is a property of the program, not of the runtime. A Reaction's guard refers only to data resident in the corpus — in L3, the arguments the program journaled when it invoked each operation, frozen at write time by the anti-porosity of Paper 1. Whatever

a Reaction at B needs to evaluate a guard travels in the corpus the program writes; whatever the program did not write there, no remote evaluator can see. The substrate carries the program, and the program carries the contract on which its consumers depend.

Paper 4, §5.3 and claim 10, named the closely related property for cross-actor causation: the journal record of the act of speaking from one actor to another is what survives cross-datacenter replication, where canonical apparatuses (CDC, log shipping, message-queue replication) reconstruct state and lose the causal chain that produced it. §5.2 of the present paper completes the picture from the substrate side: the journal record is what the consumer reads, and the consumer’s Reactions derive from it without needing the producer’s runtime to be reachable.

Two caveats are worth naming, in the spirit of full disclosure. L3’s transport is in-process — a channel between two instances in the same process — and its measurement caps at $N = 1\,000$ entries (beyond which the harness becomes memory-bound). The symmetric property is structurally insensitive to the transport’s distance and to N : a Reaction’s match path is per-entry, and the journal-segment diff at $N = 1\,000$ is identically zero across all measured cells. But the lab demonstrates the property in vitro; full cross-datacenter delivery is the operational mode realised in §7.2, and the higher- N regime is bounded by the consumer’s append throughput rather than by anything the matcher does.

What §5.2 establishes, and what §5.3 develops next

§5.2 has shown that replication, in the substrate’s terms, is the streaming of named operations over a wire whose density is by construction; that the transport between sites can be an off-the-shelf consumer-pull service; and that a consumer running the same binary reaches the same program as the producer by reading the corpus it shares. The equivalence varies the spatial position of the corpus while preserving replay as the means of reconstruction.

§5.3 varies the spatial position of the corpus without reconstructing the program at the destination. A consumer that reads the substrate and persists it locally — without instantiating the runtime that would replay it — holds the program in the form it was written. Backup, in this view, is replication arrested at the moment of holding rather than carried through to replay.

5.3 E3 — Backup is copying the program

§4.2 stated the equivalence: a consumer that reads entries from the substrate and persists them locally — without instantiating the runtime that would replay them — holds a copy of the program in the form in which it was written. The canonical literature on backup (§2.3) treats backup as a captured state, archived periodically by an external apparatus and replayed only on recovery. The equivalence stated here re-reads the operation: for a journaled program,

backup is the program materializing in another substrate, and the destination's choice is whether to hold the corpus or replay it.

The structural argument

The equivalence follows in four steps from the substrate condition of §4.1:

1. If the journal is the program (Papers 1 and 2), the artifact whose duplication preserves the program is the journal.
2. Duplication of the journal is the streaming of named entries to a destination substrate — the program materializing there, sentence by sentence, in the order it was uttered.
3. The destination need not instantiate the runtime to receive the corpus; it need only persist what arrives. The consumer is a substrate, not a process.
4. When and if the destination chooses to replay, the corpus reconstructs the program. The choice between *holding* and *replaying* is independent of the act of materializing.

Three symmetries

Backup, in this view, is positioned by its relation to the three other equivalences of §4.2. The position is fixed by what the destination commits to, not by any new operation.

- **Backup and replication.** Replication (§5.2) is the destination streaming the corpus and replaying it; backup is the destination streaming the corpus and holding it. The same protocol governs both. Replication is backup carried through to replay; backup is replication arrested before replay.
- **Backup and deployment.** Deployment (§5.1) is replay at a different time of the same site, ending with the new instance admitted as authoritative; backup is replay-eligible material held at a different site, with no instance ever admitted as authoritative there. The same protocol reads the corpus forward; what differs is the gate flip at the end — present in deployment, absent in backup.
- **Backup and offline operation.** Offline operation (§5.4) is the consumer reading the corpus with a delay between write and read but eventually catching up and serving; backup is the consumer reading the corpus and never serving. The same protocol delivers entries; what differs is whether the consumer ever takes a turn at serving requests.

The three symmetries are not corollaries discovered after the fact. They follow from the substrate condition directly: each of the four equivalences is read off the substrate by varying one of two axes — when the read happens, and what the consumer commits to once it has read — and the four exhaust the combinations. §4.3 named this arrangement.

What the substrate measures

L4 (lab `paper05-lab4-passive-consumer @ 99e8202`) instruments the materialize cycle on three storage backends under the compact-action régime of L1. The dataset spans 3 backends $\times$ 3 journal sizes $\times$ 2 protocol layers $\times$ $K = 2$ measurement repetitions = 36 cells, plus 3 catch-up cells.

Under régime $N = 100\,000$ compact entries, the destination journal converges at $\approx$ **2 150 events/sec on FileSystem**, $\approx$ **952 events/sec on MySQL**, and $\approx$ **766 events/sec on SQL Server (SQL Edge stand-in)**. The wire round-trip itself accounts for **less than 0.6 % of the cycle** at every cell; the substrate cost is the destination’s append rate, not the operation. A replica actor instantiated over the destination journal — a verification step, not part of the backup operation — reaches the same in-memory state as the primary in **18 of 18 measured cells across the three backends and three journal sizes**.

The independence from backend is the empirically significant property here. The equivalence is not a feature of any storage technology; it holds wherever the journal can be appended to.

The destination as cursor

The destination’s three reading regimes — near the head of the substrate, paused after an interruption, or newly instantiated after the primary has run for some time — share a single protocol. Each is a cursor on the substrate, read forward from its present position to the present head. Catch-up after a retention loss is not a separate mode; it is the same read, started further back. L4’s catch-up cells confirm this: the catch-up throughput matches the steady-state throughput within stochastic variation on each backend. The substrate does not contain a distinction between *current* and *recovering*; it contains a sequence of entries, and each consumer holds a cursor over it.

Caveats

Two caveats are worth naming, in the spirit of full disclosure. First, L4’s transport between primary and destination is in-process — a local proxy, not a real cross-datacenter wrapper; the network latency $\times$ number of round-trips a cross-datacenter deployment adds is bounded but not measured here. Second, L4 measures the *passive* half of the equivalence — the destination consuming the corpus — not the act of promoting a destination to authoritative status; that promotion is the E1 case §5.1 already developed, and is structurally available to any destination that holds the corpus to a chosen `EntryId`.

The operational realisation of E3 — the modes under which a destination starts, how the program declares which entries materialize where, the transport that carries them, and the dimensions along which the construct scales (heteroge-

neous backends, multi-destination fanout, per-datacenter topology) — is developed in §7.3.

To the next equivalence

§5.3 has shown that backup is the program materializing in another substrate; that backup is positioned by its relation to the three other equivalences along two structural axes — when the read happens, and what the consumer commits to once it has read; and that the corpus convergence rate on heterogeneous backends is independent of the construct. The equivalence varies the consumer’s commitments while preserving replay as the means by which the program is reconstructed if and when needed.

§5.4 varies a different axis: the delay between when the substrate is written and when it is read. Offline operation, in this view, is replay after a gap — the consumer reads the same corpus the producer wrote, but the read trails the write by an arbitrary interval.

5.4 E4 — Offline operation is delayed replay

§4.2 stated the equivalence: a consumer that is unreachable when entries are written, and that catches up by reading them later from a persistent buffer or directly from the source, applies the same replay to the same sequence as a consumer that received the entries in real time. The canonical literature on offline operation (§2.4) treats the buffer as an apparatus added at the system’s boundary — a message queue, a durable offset, a convergence protocol — whose function is to absorb the period during which the system cannot reach its counterpart. In the regime described here, no such apparatus exists: the substrate itself is the durable buffer, and offline operation reduces mechanically to cursor advancement over the same corpus.

If the producer writes to a substrate that is local — local to the producer’s process, locally durable — then there is no point at which the producer needs to know whether any consumer is reachable. The producer’s write is to its own substrate. A consumer reading from that substrate, whether it reads in real time or after an interruption, holds a cursor over the same corpus; when the cursor advances, the consumer applies the entries between the old position and the new one. The buffer of the canonical literature is therefore not added; it is the substrate itself.

Offline as the limit of replication

§5.2 developed replication as the consumer reading the corpus the producer wrote, at a different site, with the transport between them holding entries durably until the consumer fetches them. §5.4 is the limit case: the consumer’s read trails the write not by transport latency but by an arbitrary interval — minutes, hours, or longer. The protocol is identical. The act of catching up

is the act of advancing the cursor over the entries that accumulated in the meantime.

The same mechanism applies in the reverse direction. A producer whose canonical remote store is unreachable continues writing to its local substrate; the remote store, when reachable again, is a consumer that fell behind. The role of *producer* and *downstream store* is asymmetric only in steady-state framing; structurally, both are cursors over the same corpus.

What the substrate measures

L5 (lab `paper05-lab5-offline @ 99e8202`) measures the offline equivalence by configuring an actor’s substrate as a local persistent journal with an asynchronous flush to a remote canonical backend (MySQL and Azure SQL Edge). The remote backend is then partitioned for a measured interval (~5 s of write volume); the partition is resolved; the backlog drains.

Three properties of the measurement deserve naming.

1. **Decoupling of write latency from the remote.** Under online operation, the primary appends at **516 μ s p50 against MySQL** and **429 μ s p50 against SQL Edge** — **6.0 $\times$ and 2.7 $\times$ faster** than the unbuffered direct path. The producer’s lock release time is bounded by the local substrate’s append, not by the remote’s commit latency.
2. **Partition latency is not worse than online latency.** During the partition, the buffered primary’s per-append p50 is 333 μ s on MySQL and 339 μ s on SQL Edge — the same régime as online; the write lock never traverses the network. The substrate’s local persistence is what makes this hold; the absence or presence of the remote does not enter the write path.
3. **Catch-up is linear in the backlog.** On reconnect, the accumulated entries drain at the remote backend’s steady-state ingest rate — ≈ 296 events/sec on MySQL, ≈ 500 events/sec on SQL Edge — with **zero events lost** across the measured cells. Catch-up is not a separate mode; it is the remote’s cursor advancing from its last position to the present head.

Three caveats are worth naming, in the spirit of full disclosure. First, L5 partitions the remote by stopping its container — a realistic partition mode but not the only one; half-open connections without container death would exercise the connection-timeout path differently and are not measured here. Second, the partition-phase latency being at or below the online phase is partly an artifact of single-host scheduling — the asynchronous flush thread is idle during partition, freeing CPU for the writer; in production with separate threads on separate cores, the artifact shrinks, but the structural property — the write path is independent of the remote’s reachability — holds either way. Third, *what* drains on reconnect is journal entries to the canonical backend — the remote’s cursor advancing to the present head; the measured path does not re-emit cross-actor

tells. Tell re-delivery on a replay-style catch-up is a separate path belonging to the tell primitive (Paper 4), with its own delivery semantics and not exercised here; the *zero-events-lost* result is a property of journal drain to the backend, not a claim about tell re-delivery on replay.

The operational realisation of E4 — the configuration that enables the local substrate as the primary’s source of truth, the asynchronous component that flushes to the canonical backend, and the catch-up path on reconnect — is developed in §7 alongside the other operational details.

To the next equivalence

§5.4 has shown that offline operation is replay after a gap, and that the gap is absorbed by the substrate’s own local persistence rather than by a separate apparatus. The four equivalences of §4.2 are now developed.

§5.5 turns to the cost that underlies all four: replay against a local substrate is cheap because append against a local substrate is cheap. The latency budget that supports the four equivalences is examined as a single property of the substrate, not as a property of any one operation.

5.5 The latency budget

The four equivalences of §4.2 reduce four classical operational disciplines to a single primitive: replay. Replay is the act of reading the substrate forward from a cursor and applying each entry. Its cost is the cost of reading and applying entries. For a substrate that is read often and written often, the cost of *writing* entries — appending to the journal — is the cost basis that supports the family of replay-based operations. §5.5 examines that basis.

The structural claim that underlies E1–E4 is this: every replay-based operation reduces to advancing a cursor over N entries, and therefore inherits the substrate’s per-entry append latency as its dominant cost term. The actor’s apply cost is identical across all regimes because it operates on the same in-memory state. The only term that varies between regimes is the cost of appending entries to the substrate.

The relevant comparison is between two regimes for the substrate’s persistence. Under the regime this paper develops, the substrate is a local journal whose persistence is one *fsync*-bounded append per entry. Under the canonical regime (§2), the substrate’s role is played by an RDBMS whose persistence is one commit-bounded *INSERT* per entry. The latency gap between the two regimes propagates linearly to the cost of deployment, replication, backup, and offline catch-up. The four operational disciplines therefore share a single cost basis: the substrate’s append latency.

What the substrate measures

L6 (lab `paper05-lab6-latency-budget @ 99e8202`) measures per-entry append latency on three local backends under a one-event = one-durable-commit régime: a local FileSystem journal, a co-located SQL Server (Azure SQL Edge stand-in), and a co-located MySQL. The dataset spans 3 backends $\times$ K = 5 runs $\times$ N = 100 000 entries per backend = 1.5 million samples, with the first 1 000 appends per cell discarded as warm-up.

Backend	Append p50 (μ s)	Append p95 (μ s)	Append p99 (μ s)	Ratio to local FS (p50)
Local FileSystem journal	328	540	3 434	1.0 $\times$
Co-located SQL Server	1 170	1 916	2 479	3.56$\times$
Co-located MySQL	1 005	1 536	1 861	3.06$\times$

Durability is parameterised identically across the three backends: the FileSystem backend invokes a per-entry flush-to-disk; SQL Server runs with default per-commit log flush; MySQL runs with `innodb_flush_log_at_trx_commit=1` and `sync_binlog=1`. The measurement therefore compares the same physical durability guarantee across all three — one entry, one durable commit.

The journal-append regime favours the substrate by $\approx 3\times$ over both RDBMS engines at the median, and the ordering holds at p95 (the FileSystem journal remains the fastest). The advantage is *structural*: both RDBMS engines land in the same $\approx 3\times$ band, so the gap is a property of the substrate’s append being one *fsync* per record versus the RDBMS being one *fsync* per redo-log entry plus the protocol overhead of a transaction. At p99 the ordering inverts — the FileSystem journal’s tail (3 434 μ s) exceeds both engines’ — because a direct `fsync` is exposed to occasional host-level flush and GC stalls that the RDBMS engines’ buffered log-writers absorb; the median and p95 advantage is the operative figure, and the p99 tail is a durability-flush artifact, not a throughput property.

The corollary to E1–E4 follows. Every replay-based operation costs $\approx 3\times$ less per entry against the local-journal substrate than against a co-located RDBMS in the durable-commit regime. Deployment at N entries (E1, §5.1), replication at N entries (E2, §5.2), backup ingestion at N entries (E3, §5.3), and catch-up after a partition at N entries (E4, §5.4) inherit the same factor.

Caveats

Two caveats are worth naming, in the spirit of full disclosure.

First, the comparison is *conservative against the substrate*. The RDBMS engines in L6 run on ephemeral container storage — MySQL on a tmpfs mount, SQL Edge on a Docker named volume — whereas the local journal goes through the host’s NVMe controller. On a production RDBMS deployment with real disk, the RDBMS side would slow further; the $3\times$ ratio reported here is therefore a lower bound, not an upper one. Adding network latency for a remote RDBMS — a typical production topology — adds a flat additive term per commit, pushing the ratio further still.

Second, the comparison can be narrowed by relaxing durability on the RDBMS side. Group commit, batched inserts, and asynchronous-durability modes — SQL Server’s `DELAYED_DURABILITY` and MySQL’s `innodb_flush_log_at_trx_commit=2` — all close part of the gap by trading durability for throughput. These are honest engineering choices; L6 measures the durable-commit baseline because the substrate’s local-journal regime keeps full durability and the apples-to-apples comparison demands the same of the RDBMS side. A reader who deploys an RDBMS-backed substrate with relaxed durability will see a smaller gap; the structural claim — that replay cost is bounded by append latency — does not change, only the constant.

To the next sub-section

§5.5 has shown that the cost basis of the four equivalences is one quantity — the substrate’s per-entry append latency — and that against a co-located RDBMS in the durable-commit regime, the local-journal substrate runs at $\approx 3\times$ lower latency per entry. The result is structural rather than engine-specific: two RDBMS engines, two implementations, the same ratio band. §5.6 closes §5 by naming two further consequences of the substrate — observability and debug — that fall out of the same property without requiring separate apparatuses.

5.6 Forensic operations as substrate consequences

Two further operations follow from the substrate without requiring separate apparatuses. The substrate is a total, ordered record of every named operation performed by the program. Conventional observability stacks reconstruct this record indirectly by correlating logs, metrics, and traces emitted outside the program. Under the substrate condition, the record already exists in primary form; any observability layer becomes a reader of the substrate rather than a reconstructor of program history. Sagas do not appear as patterns to be inferred post-hoc from logs; they are named at write time by the program’s partition primitive (Paper 3) and therefore exist in the substrate as first-class entries.

Replay is deterministic — by the argument-capture precondition of §4.1, the same corpus, applied to the same binary, produces the same in-memory state at the same cursor position, because every argument in the corpus is a value frozen at write time rather than an expression re-evaluated on read. A runtime issue traceable to an entry is therefore reproducible by replaying the substrate

to that entry’s position; a conditional breakpoint over the entry’s identifier is trivial because the identifier is a position on the cursor, and the cursor advances one entry at a time. Debugging reduces from forensic analysis over logs to advancing the cursor to a chosen position in the substrate. Both observations exist on the same axis as §5.1–§5.5: they require no additional mechanism beyond the substrate itself. §6 turns to the cost in operational complexity that the canonical disciplines pay to achieve, by separate apparatuses, what the substrate already provides.

6. Why existing approaches duplicate work the substrate already does

The four disciplines surveyed in §2 are not unrecognised by their respective traditions. Each has developed, over decades, a refined apparatus to address the operational concern it names — blue-green and rolling-update for deployment, change-data-capture and log shipping for replication, snapshot agents and point-in-time recovery for backup, message-queue offsets and convergence protocols for offline operation. The maturity of these apparatuses is not in question. What this section examines is whether any of them dissolves the operational concern it addresses, or whether each reconstructs, from outside the program, work that the substrate condition of §4.1 already does.

This section is not a critique of existing practices but a structural comparison between what those practices reconstruct and what the substrate condition provides by construction. The four sub-sections below examine one canonical apparatus per discipline. The pattern of each examination is the same: name what the apparatus does, identify the structural property it relies on, and observe that the substrate condition provides that property by construction. The point is not that the apparatuses are wrong but that they are responses to the absence of the substrate condition. Once the condition holds, the work each apparatus performs is already done by the program’s own act of writing the journal.

6.1 Blue-green vs substrate

Blue-green deployment (Fowler, 2010) provisions two parallel production environments and switches traffic between them at release time. The apparatus has three components: a duplicate environment, a router that decides which environment is authoritative at any moment, and a procedure for bringing the standby environment to the state from which it can accept traffic. The first two are operational infrastructure; the third — bringing the standby to a serving state — is what the apparatus does that the substrate condition recasts.

Under the substrate condition (§4.1), the act of bringing a process to a state in

which it can serve is the act of replaying the journal up to the present head (E1, §4.2). The new version’s instantiation differs from the old version’s continued operation only in the position of its read cursor. The duplicate environment is therefore not a parallel system to which state must be propagated; it is a process whose cursor on the same corpus has not yet reached the present head. The replay that the apparatus performs implicitly — by re-running migration scripts, by exporting and re-importing state, by re-priming caches — is performed explicitly by the substrate’s reading discipline.

The structural redundancy is not that blue-green is wrong; it is that blue-green’s hardest step, the act of synchronising the standby with the live state, is replay reconstructed manually. The reconstruction is manual because the live state, in the canonical regime, is not the journal of the program; it is a database whose entries record what is, not what was said. Under the substrate condition the manual reconstruction has nothing to do, because the corpus from which the standby reads is the same corpus the live instance writes. The router and the duplicate environment remain; what dissolves is the apparatus that bridges the two.

6.2 Change-data-capture vs event streaming

Change-data-capture (Debezium, n.d.; Maxwell, n.d.) reads the binary log of a relational store and re-emits each row-level change as an event on an external transport. The apparatus exists because the originating system does not emit events; it emits state changes, and CDC observes those changes and translates them back into the event form that downstream consumers require. The translation step is the apparatus.

The translation is informationally lossy in one direction and inferential in the other. State changes do not carry the operation that produced them; CDC infers, from a sequence of column-level deltas, what the application program intended. The inference is necessarily heuristic — two distinct programmatic operations can produce identical column-level deltas, and the CDC consumer cannot recover the distinction. The apparatus is therefore not only a transport; it is a reconstruction layer that approximates events from state, with the loss inherent in that approximation.

Under the substrate condition the reconstruction has no work to do. The originating system emits events — they are the entries of the journal, written in the language of the program (Paper 1, anti-porosity) and reduced to compact named operations (Paper 2, separability). The wire across which replication travels carries those entries directly (§5.2); no inference from state, no column-level delta extraction, no schema mapping enters the act. The CDC apparatus is structurally duplicative: the substrate emits exactly the artifact CDC reconstructs — minus the reconstruction loss, and minus the apparatus that produces it.

6.3 Snapshot backup vs passive consumer

Snapshot backup (Hitz, Lau, & Malcolm, 1994; PostgreSQL, n.d.; Veeam, n.d.) captures the state of a data store at a chosen moment and writes that capture to an archival medium. Recovery instantiates the captured state and resumes operation from it. The apparatus has two components: the snapshot mechanism that produces the capture, and the recovery procedure that consumes it. Both operate on state — the captured artifact is a frozen image of what the program had stored, not a record of what the program had done.

The structural property the apparatus relies on is that the program’s state, if captured atomically at a moment, suffices to characterise the program at that moment. The property holds for programs that are reducible to their state. It does not hold for programs that are reducible to their history: a snapshot of the state at moment t does not record the operations that produced it, only their cumulative effect, and recovery from the snapshot loses the operational record.

Under the substrate condition the program is not reducible to its state; it is constituted by the sequence of operations in its journal (§4.1). A consumer that holds the corpus — without instantiating the runtime that would replay it — holds the program in the form in which it was written (E3, §4.2; §5.3). The choice between holding and replaying is a property of the destination, not of the captured artifact. The snapshot apparatus produces a state image and discards the operational record; the substrate emits the operational record continuously and lets the destination decide whether and when to replay.

The redundancy is structural in two directions. First, the snapshot apparatus is unnecessary: the substrate is already the program, copied continuously. Second, the snapshot apparatus is also insufficient: it captures state, which is less than the program — recovery from a snapshot reconstructs *where* the program was, not *what it did to get there*. The passive consumer of §5.3 holds the program, and any moment in the program’s history is replayable from the corpus the consumer holds.

6.4 Message queue buffer vs persistent local journal

The message-queue pattern places a durable buffer between a producer and a consumer that may not be reachable in real time. RabbitMQ’s dead-letter exchanges, Kafka’s consumer offsets, and the eventually-consistent stores derived from Dynamo (DeCandia et al., 2007) each place an apparatus at the boundary between the system and its counterpart. The apparatus is external to the producer; the producer writes to its store, and a separate mechanism — the queue, the broker, the convergence protocol — handles the case in which the consumer is unreachable.

The structural property the apparatus relies on is that the producer’s store and the consumer’s store are different artifacts, with a transport between them that must be made durable. Under that structural premise, the queue is the right

apparatus: it absorbs the period during which the consumer cannot reach the producer, and it persists the messages that would otherwise be lost.

Under the substrate condition the producer’s store and the consumer’s store are not different artifacts at the level of representation. The producer writes to its substrate; the consumer reads from a copy of the same substrate, or from the substrate directly. The substrate is itself the durable artifact, and offline operation reduces to the consumer’s cursor trailing the producer’s write head by an arbitrary interval (E4, §4.2; §5.4). No queue is added between the two parties because the substrate is already the queue — it persists every operation in the order it was uttered, and a consumer that has been unreachable for an arbitrary period catches up by reading the entries that accumulated.

The redundancy here is the most structural of the four. Where the queue, the broker, and the convergence protocol exist as boundary infrastructure, the substrate is itself the boundary — it is what holds the program between writer and reader, between connected and disconnected, between now and later. The infrastructure that the canonical literature places between two stores is, under the substrate condition, the property of the single store both parties share.

6.5 The substrate match table

The four examinations of §6.1–§6.4 share a structure. Each canonical pattern relies on a property the substrate condition supplies by construction; each apparatus performs work that the journal, written as the program, already does. The pattern is summarised in the table below, in the form of Paper 4 §6.4: a single question — *does the apparatus address a property the substrate already provides?* — answered for each pattern.

Pattern	Property the apparatus relies on	Does the substrate condition already provide it?
Blue-green deployment	The standby must be brought to the same state as the live instance before traffic switches	Yes — replay of the same corpus to the present head (E1, §5.1) brings the standby to that state by construction; no external state-propagation step is required

Pattern	Property the apparatus relies on	Does the substrate condition already provide it?
Change-data-capture	The originating store emits state changes; events must be reconstructed from them	Yes — the substrate emits the operations themselves (Papers 1 and 2; §5.2); no reconstruction step is required and no inferential loss is incurred
Snapshot backup	The program is characterised by its state at a moment; recovery instantiates that state	Yes — the program is constituted by the sequence of operations in the journal (§4.1, §5.3); the corpus held by a passive consumer is the program, not a state capture of it
Message queue buffer	A durable apparatus is required between the producer’s store and the consumer’s store	Yes — the substrate is the durable artifact both parties share (§5.4); a consumer reading later than the producer wrote reduces to cursor advancement over the same corpus

The four “Yes”s share a structure. In every case, an apparatus that the canon developed to bridge an absence of the substrate condition becomes structurally subsumed once the condition holds. The pattern remains usable — none of the four is wrong as engineering — but the work it performs duplicates work the substrate already does. The maturity, sophistication, and widespread adoption of the four patterns are not evidence that the substrate condition is unnecessary. They are evidence that the absence of the substrate condition is costly enough to justify four separate apparatuses to compensate for it, each developed independently by a different tradition.

The diagnosis of §2–§6 is now complete. §7 turns from the construct to the instantiation: a system in production whose runtime already satisfies the substrate condition, whose four operational disciplines fall out of the substrate without dedicated apparatus, and whose existence is the evidence that the construct names a realisable property and not only a conceptual one.

7. Instantiation: realization in Puppeteer

The construct of §§3–6 makes claims about a class of systems characterised by the substrate condition. The present section exhibits one such system. The voice from this point forward is concentrated and authorial: the runtime described below is *Puppeteer*, the framework the present author maintains, and the operational decisions named are decisions taken in deployments that run today. The framework is presented here as an existence proof for the construct rather than as the subject of the paper; the four equivalences hold in any runtime that meets the substrate condition, of which Puppeteer is one.

7.0 Origins

Paper 3, §6 introduces the runtime decomposition this section relies on: *Performance* is the local hosting of a single actor over a configured storage, *Theater* is the host for one or many Performances on a node, *Ensemble* is a pool of actors with routing, and *StageManager* is the peer-to-peer replicated state machine that coordinates topology across nodes. The full vocabulary is presented in Paper 3 §6 and is not re-introduced here. §7 names only the runtime surfaces on which the four equivalences of §4.2 land: the *aliveGate* (E1, §7.1), the per-event push callback wired to a webhook delivery service (E2, §7.2), and the *Materialize* construct that declares which entries materialise in which destinations (E3, §7.3). The offline equivalence (E4) lands on the same write path the local-journal régime of §5.4 already exhibited, and is referenced rather than re-developed here.

7.1 Theater.aliveGate: the red-black mechanism

The handover window of §5.1 is realised in the runtime by a single boolean predicate: the *aliveGate*, a manual-reset event held on every **Performance** instance. The gate is the locus of the deployment apparatus, and its life-cycle is the operational realisation of the structural intervals named in §5.1.

A **Performance** is started in one of two modes. A primary starts with `Start()` (default `asFollower: false`): the runtime hydrates the actor from its substrate, raises the `OnFirstHydration` and `OnHydrated` callbacks, and sets the gate. A follower starts with `Start(asFollower: true)`: the runtime hydrates the actor from the same substrate, but the gate remains reset and the public `IsAlive` predicate reports `false`. External callers — request dispatchers, write coordinators — block on `WaitUntilAlive` until the gate is set.

The handover is a single sequence over the gate. The follower invokes `LockWhileNotSynchronized()`, which pauses writes on the live primary and returns the `EntryId` at which the pause occurred. The follower then catches up

over the residual tail with `CatchUpFromJournal(targetEntryId)`, which replays pending events under the actor’s write lock until the follower’s cursor reaches the targeted `EntryId`. `UnlockAndRunAlive()` flips the follower’s gate and releases the primary. From the perspective of any external caller, the system transitions from “primary alive” to “follower alive” in a single atomic moment — the gate flip — preceded by replay of the substrate and nothing else.

The structural content of the apparatus is therefore minimal. The gate is book-keeping over the substrate’s read cursor; the runtime contains no separate deployment subsystem, no state-extraction layer, and no schema-mediated handoff. The orchestration tail measured in §5.1 — below 0.3 ms at every measured N — is the cost of the gate flip and the residual-tail replay, not of any apparatus. Appendix A lists the exact source locations of each named surface.

7.2 Event streaming: push-to-pull inversion

The substrate’s push-to-pull inversion (§5.2) is realised in Puppeteer by composing three primitives that already live in the runtime. None of them was introduced for §7.2; each existed for its own structural reason, and together they satisfy the consumer-pull requirement.

First, the journal exposes a per-record callback. `Diary.OnRecordWritten` is invoked by the storage layer for every entry written, with the entry’s `entryId` and its raw byte representation. The setter accepts a single subscriber; the companion `AddRecordWrittenCallback` chains additional subscribers without disturbing the existing one, so that multiple downstream consumers can attach to the same primary without the primary knowing how many.

Second, the partition primitive of Paper 3 exposes a *Cue* mode for Reactions. A `Cue()` Reaction runs continuously: it batches over the substrate from the current cursor to the present head and then enters *push mode*, subscribing to the same `OnRecordWritten` callback above. The push loop applies each entry as it arrives, matches it against the Reaction’s pattern, and dispatches the Reaction’s terminator if the pattern matches.

Third, the substrate’s externalisation point is the terminator of a `Cue` Reaction. The terminator’s body — DSL the program writes — is the place at which the runtime delegates to whatever external delivery service the deployment has elected. We observe that an off-the-shelf webhook delivery service (SVIX, in the deployments referenced here; equivalent products would suffice) satisfies the substrate’s consumer-pull requirement: the Reaction’s terminator posts the entry to the service from a per-entry callback; the service holds the entry durably; downstream subscribers fetch from the service at their own pace.

No code in `Puppeteer/` or `Choreography/` references SVIX. The runtime exposes the callback and the partition primitive; the wiring to a particular webhook service is a property of the deployment, not of the framework. The substrate’s claim is that any service satisfying the consumer-pull contract is sufficient; the

off-the-shelf adoption is the empirical witness to that sufficiency.

7.3 Passive consumer: replicate without an actor running

The passive consumer of §5.3 is the operational realisation that, in earlier drafts of this paper, was named as an honest gap — a property derivable from the substrate condition but not yet exhibited as a system in production. The gap closed on the day this section was drafted, by composition of patterns the runtime already had. The composition has since been formalised under the construct *Materialize*. The remainder of this sub-section walks the composition brick by brick, in the spirit of the andragogical commitment of Paper 2: a reader who has followed §§4–6 is at risk of dismissing E3 as “stated, not exhibited”; the present sub-section makes the assembly visible so that the dismissal has nothing to dismiss.

The instantiation has four bricks. The reader is asked to hold them separately and observe how they compose.

Brick 1 — Same binary, alternate startup mode. A destination is a process running the same actor binary as the primary, started in an alternate mode. The mode admits no external requests: the `aliveGate` of §7.1 is not opened for the destination’s lifetime. Two such modes exist. `Performance.Start(asFollower: true)` starts the actor as a *follower* — a process whose `Job()` Reactions run against a shared substrate; the follower is a worker, not a backup, and its Reactions contribute markers back to shared auxiliary tables. The second mode, `/materialize`, runs no Reactions at all — it consumes pre-computed markers from the primary via the wire protocol described under Brick 3. The mode is selected by the host process at startup; the runtime is the same in either case.

Brick 2 — Per-event marker in the DSL. The program declares which entries materialise in which destinations through a marker on the Reaction’s metadata plane:

```
actor.Reactions.DefineReaction("...")
    .Job().Company()
    .Seek("...")
    .OnMatch("...")
    .Metadata.Materialize("DC-B");
```

`Metadata.Materialize("DC-B")` records, for every entry that matches the pattern, a row in the `EventMaterialization` table naming the destination. The row is the program’s declaration that this entry is part of the corpus the named destination is to receive. Reading the program reveals which entries are intended for which destinations, in the same DSL the program is written in — the construct is, in the sense of Paper 1, an instance of the code-as-data recording the substrate condition requires.

Brick 3 — The wire protocol between primary and destination.

The destination, when it wakes up, asks the primary for the entries it has not yet seen. The exchange is realised by four wire verbs on the primary's `actor.Materialization` sub-namespace. `ReadRecordsAfter(destination, fromEntryId)` returns the raw substrate records the destination is missing — Layer 1, the program as the primary wrote it. `ConfirmUntil(destination, entryId)` records the destination's acknowledgment on the primary side as a Max-monotonic watermark — the primary now knows the destination has the corpus up to that `EntryId`. `ReadReactions(destination)` and `ReadElidedRange(destination, from, to)` return Layer 2 derived state — the Reaction registry's atomic snapshot and the elision markers the primary computed in the same range — so that a destination electing to replay can reach the same in-memory state the primary reached without re-running the Reactions itself.

The destination side is the fluent client `MaterializeMirror`. The contract is two-layered:

```
mirror.Sync();                                // Layer 1 –
orchestrates (a) + (b).
mirror.AsProgramMirror().Sync();              // Layer 2 –
orchestrates (a) + (c) + (d) + (b).
```

`Sync()` fetches records and confirms the watermark; `AsProgramMirror().Sync()` additionally fetches the Reactions snapshot and elision markers, so that the destination's program state is reconstructible without instantiating the Reaction matcher. The destination decides what to do with the result — a passive backup persists the records and discards Layer 2; a replicated read-replica persists both and applies the elision; a snapshot-time mirror runs Layer 1 only.

Brick 4 — Recovery primitive. When a destination has been offline for an arbitrary period — or when an intermediate delivery service has exhausted its retention — the destination resumes by issuing `Sync()` against its current watermark. The primary serves the missing entries from its substrate; the destination's watermark advances; the round-trip is the same as steady-state. There is no separate recovery mode in the runtime, because there is no separate steady-state mode either: every fetch is `Sync()` against the watermark, and every catch-up is `Sync()` over a wider range. The recovery primitive is the steady-state primitive.

The four bricks compose into a single operation. The program declares which entries materialise in which destinations (Brick 2). A destination process — running the same binary in `/materialize` mode (Brick 1) — fetches records from the primary's Materialization surface (Brick 3) at its own pace, confirms what it has received, and catches up after any interruption by the same call (Brick 4). The runtime contains no replicator, no separate backup subsystem, and no synchronisation layer between primary and destination beyond the four wire verbs and the per-Reaction marker. The operation falls out of the substrate;

the construct names the operation that the program already declares.

Three dimensions of the construct deserve naming, because each is structurally present without requiring a separate apparatus.

Heterogeneous backends. `Diary` selects its storage implementation polymorphically from `(DatabaseType, connectionString)` at construction time. A primary running on the local `FileSystem` can materialise to a destination running on MySQL or SQL Server, and the wire protocol carries the same raw records across the heterogeneity — `ReadRecordsAfter` returns Layer 1 records regardless of which backend the primary stores them in. The four wire verbs are implemented on each of the four backends with the same contract; L4 (§5.3) exhibits this independence with **18 of 18 cells reaching equal in-memory state across three backends and three journal sizes**.

Multi-destination fanout. `AddRecordWrittenCallback` chains additional subscribers without displacing the existing one. Multiple destinations may register against the same primary; the primary’s write path is unaware of how many. Each destination holds its own watermark on the primary side; each progresses at its own rate; the primary serves them independently. Multi-backup of one actor — and the per-destination retention policies that go with it — falls out of the chain primitive.

Per-datacenter topology. A datacenter is an operational placement of `Performance` instances; it is not a construct the actor knows about. A primary at datacenter A and a destination at datacenter B is a topology decision realised by the operator: a local SVIX at each site for the Cue feed of §7.2, and `/materialize` processes at each site whose mirror clients fetch from the local SVIX or directly from the primary. The substrate’s locality property (the journal is the actor’s, not the system’s; cross-ref §6.2 of Paper 3) extends laterally: each datacenter’s destinations are wired to the local instance of the delivery service, and cross-site bandwidth is consumed only by the SVIX-to-SVIX replication the off-the-shelf service performs.

A gap is worth naming here, in the same spirit of full disclosure. The `Diary.OnRecordWritten` setter is wired today for primary `FileSystem` and for the buffered régime; for primary MySQL or SQL Server the setter is a silent no-op. The canonical case of the substrate (a `FileSystem`-primary deployment) is unaffected, and the §7.2 push-to-pull path runs against the `FileSystem` callback as designed. A future deployment with primary MySQL would require the callback to be wired in the MySQL storage path — a localised patch, not a structural change.

7.4 Numbers

The operational régime of the four bricks above runs under the latency budget §5.5 already developed. The substrate cost of materialisation is the cost of appending the program at the primary (which §5.5 measures) plus the cost of

appending at the destination (which §5.3's L4 measures). No new numbers are introduced here; the relevant ones live in §5.1 (deployment window), §5.2 (wire density), §5.3 (destination convergence), §5.4 (offline catch-up), and §5.5 (per-entry append latency). Together they characterise the régime in which the four bricks above operate; together they materialise the construct in production-grade numbers.

8. Relation to previous work in this paper series

The substrate condition stated in §4.1 — that the journal is code-as-data, dense, and made of compact named operations — is not asserted *ex nihilo* by the present paper. Each of its components is the conclusion of prior structural analysis in this series, and each entered §4.1 as a precondition rather than as a claim of this paper. Without those preconditions the four equivalences of §4.2 would not follow; with them, the four equivalences are corollaries of the substrate's properties under variations of when and where the read happens. §8 makes the chain of dependence explicit.

Paper 1 introduces *anti-porosity* as a design principle for the boundary between domain code and persisted form. The journal records what the program said, not the data structures the program manipulated; entries are sentences in the same language the program is written in, and every effect that crosses the boundary is named in that language rather than summarised as a state delta. Anti-porosity is the precondition under which the journal can be the program at all: a porous journal records states whose operational origin is lost, and the substrate condition reduces to recording a derivative of the program rather than the program itself. §4.1 takes anti-porosity as given; without it, E2 (replication as sharing history) and E3 (backup as copying the program) collapse into the canonical regimes of §2.

Paper 2 introduces *separability* as the structural property that makes the journal entry of a parametric verb a compact reference plus its arguments. The verb's body is not present in the entry; it lives once, named, in the corpus's vocabulary. Separability is the precondition under which the wire of §5.2 carries operations at a density that does not scale with the operation's internal complexity, and it is the precondition under which the latency budget of §5.5 holds — replay is linear in compact-entry count, not in expanded state, only because separability has reduced the entry to its reference form. §5.5's three-times advantage against co-located RDBMS engines is, in part, the operational expression of the structural property Paper 2 names.

Paper 3 introduces *the partition* between immediate work, done at the actor's writer-lock cursor, and deferred work, done by Reactions that read the journal forward and produce their own derived entries. Paper 3 §6 establishes the runtime decomposition — *Performance, Theater, Ensemble* — that frames the

operational realisation of §7. The mechanism that admits a new instance as authoritative (§5.1’s gate flip) is named in Paper 3 claim 8 as a property derivable from the partition; the same paper names the cross-actor extension as a forward reference closed by Paper 4. §5.2’s symmetric-consumer property — that two instances of the same binary, reading the same corpus, produce identical Reaction output (L3) — depends on a Reaction reading only corpus-resident data: in L3, the arguments the program journaled when it invoked each operation — frozen at write time by Paper 1’s anti-porosity — so whatever the remote consumer needs to evaluate the guard travels in the corpus the program writes.

Paper 4 introduces *tell* as the primitive under which cross-actor causation is recorded as program in each participant’s local journal. Paper 4 claim 10 names cross-datacenter replication as the canonical case in which *tell*’s program-level recording survives where the conventional apparatus (CDC, log shipping) reconstructs state and loses the causal chain. §5.2 of the present paper completes the picture from the substrate side: the journal record is what the consumer reads, and the consumer reconstructs not only the local actor’s state but, when *tell* entries are present, the cross-actor causal chain those entries record. The substrate condition makes Paper 4’s cross-DC property hold without an external mechanism that observes the producer’s runtime.

The dependence is summarised in the table below.

Paper	What it establishes	What §4.1 takes as given	Where it is load-bearing in this paper
Paper 1	Anti-porosity: the journal records operations, not state	The journal is code-as-data and dense	E2 wire density (§5.2); E3 the corpus is the program (§5.3); §6.3 snapshot is informationally less
Paper 2	Separability: the entry of a parametric verb is ActionId + arguments	The journal entry is compact, named, parameter-referenced	Wire density at production scale (§5.2; L2); replay rate linear in compact count, not state (§5.1; L1); latency budget (§5.5; L6)

Paper	What it establishes	What §4.1 takes as given	Where it is load-bearing in this paper
Paper 3	Partition: immediate vs deferred; Reactions; runtime decomposition (<i>Performance</i> , <i>Theater</i> , <i>Ensemble</i>)	Reactions and the partition primitive; the runtime apparatus referenced in §7	Gate flip apparatus (§5.1; §7.1); symmetric consumer via corpus-resident reads (§5.2; L3); apparatus vocabulary (§7.0)
Paper 4	<i>tell</i> primitive: cross-actor causation as program in sender's journal	The cross-actor recording the substrate carries across sites	Cross-DC causal chain travels with replication (§5.2; cross-ref claim 10)

The four prior papers can be read as the structural derivation that the present paper takes as substrate. Each named one property that the journal had to possess for the program to live in it; the present paper takes the four properties as given and reads off the four operational equivalences that follow. Where each prior paper went from one structural commitment to its consequences, the present paper goes from four structural commitments to a single primitive — replay — under which the four canonical operational disciplines collapse. The relation is therefore not parallel: §4.1's substrate is what the prior four established, and §4.2's theorem is what follows once it is in hand.

9. Counter-arguments

The substrate condition and its four equivalences invert vocabulary the canonical literature has used productively for decades. A reader familiar with that literature is likely to raise objections that are not addressed by §§2–8. Four are taken up here. The treatment of each follows the same shape: name the valid aspect of the objection, then identify the structural premise on which the invalid aspect depends and observe that the substrate condition removes that premise. The refutations are made from the construct, not from any framework that instantiates it.

9.1 “Loading a large journal takes time”

Canonical assumption. Bringing a system to a serving state requires reconstructing its persisted history end-to-end, and the cost of that reconstruction grows without bound with the system’s operational age.

The objection. A journal that accumulates entries over the operational life of a program grows without bound. Bringing a new instance to a serving state by replaying the journal therefore takes longer the older the program becomes. A canonical anecdote names a relational system whose initial load — from a binary log accumulated over a year — required roughly an hour before the system could accept traffic. If deployment is replay (§5.1), then deployment is bounded below by that load time, and the substrate’s account of deployment makes the deploy window worse, not better.

What is valid. Replay is linear in the entries it reads. A journal twice as long takes twice as long to replay. The objection correctly observes that the absolute deploy time grows with the program’s history under any substrate-based account of deployment.

What does not follow. The hour-long load named in the anecdote is not replay over the substrate condition; it is reconstruction of state from a log whose entries are state-change deltas. The two régimes have different cost basis. Under the substrate condition, the entry is a compact named operation — **ActionId** plus arguments — and replay is the application of those operations in order (Papers 1 and 2). The per-entry cost is the in-memory apply of one operation, not the disk-bound write or schema-mediated commit that state-change replay incurs.

§5.1’s measurement materialises the structural distinction. L1 reports a sustained replay rate of **532 thousand entries per second** at the median ($N = 100\,000$), rising to **1.76 million entries per second** at $N = 1\,000\,000$ — the rate is bounded by neither N nor the per-entry apply cost, and the orchestration tail around bulk replay is below 0.3 ms at every measured N . A journal of ten million compact entries replays in the order of ten seconds at this rate, not an hour. The objection generalises a measurement made in one régime — state replay over a binary log — to a régime in which the entries are operations, and the generalisation does not hold. What grows linearly is the number of operations the program has uttered; what does not grow is the cost of uttering each one again in memory.

One canonical case lies outside this argument and is preserved as an honest limit. Where the cost-dominant operation at a single entry is itself expensive — a network round-trip to an external service, a complex domain calculation — replay carries that cost per entry just as live operation did. The substrate does not make any one operation cheaper; it makes the bookkeeping around the operation negligible.

9.2 “Change-data-capture already solves replication”

Canonical assumption. Cross-datacenter replication is properly addressed by reconstructing events from the originating store’s state-change deltas, mediated by a CDC pipeline between primary and replica.

The objection. Cross-datacenter replication is a mature engineering concern. Change-data-capture pipelines, log shipping between database engines, and managed replication services route writes from primary to replica with bounded lag and well-understood failure modes. The substrate’s account of replication (§5.2) appears to re-derive a problem that the canon has solved repeatedly. What additional structural claim does the substrate condition support that the existing apparatus does not?

What is valid. CDC pipelines work. Production systems replicate across datacenters with CDC every day; the apparatus is reliable, vendor-supported, and well-tooled. As an engineering solution, it solves the operational concern of moving state between sites.

What does not follow. The structural premise of CDC is that the originating store records state changes, not operations, and that events must be reconstructed from those state changes by inferential reading of the binary log (§2.2, §6.2). The reconstruction is lossy: two distinct programmatic operations that produce identical state changes are indistinguishable at the CDC layer. The wire across which CDC traffic travels carries a derived artifact — column-level deltas — at a width that scales with the size of the affected rows rather than with the operation that caused them.

The substrate condition replaces the reconstruction with direct emission. The journal entry is the program’s operation, in the language of the program (Paper 1), reduced to a reference to a parametric verb plus its arguments (Paper 2). The wire carries the operation as uttered, not the state changes it produced. §5.2’s measurement (L2) reports a compact wire of **33 to 96.7 bytes per event** depending on verb tier, against literal-form payloads of **118 to 910.7 bytes**; the literal-to-compact ratio reaches **9.4× at the production-shaped tier and 20× at full production scale** (Paper 2 Lab 4). Applying **gzip** to the literal form closes the gap only to **4.0× at tier 3** — the structural saving is not algorithmic but representational, and a compressor operating after the fact cannot recover what naming achieved by construction.

The objection is therefore precise in scope. CDC solves replication as state propagation. The substrate condition does not deny that solution; it shows that under the construct’s conditions the problem CDC solves does not arise, and the wire carries a different artifact. The two are not in competition for the same engineering decision; they are at different layers of the design.

9.3 “Snapshot backup is simpler”

Canonical assumption. Backup is a captured image of state at a chosen moment, archived by an external apparatus and re-instantiated at recovery time.

The objection. Backup as a periodic state snapshot is a forty-year-old discipline with well-understood operational properties: recovery time objective (RTO), recovery point objective (RPO), storage tiering, lifecycle policies. The substrate’s account of backup (§5.3) requires a continuously consuming destination — a process that holds the corpus as it grows — which appears operationally more complex than periodic state capture. Why prefer continuous to periodic?

What is valid. Periodic snapshots have an operational vocabulary the field knows. The discipline is mature, vendor-supported, and integrates with object-store lifecycle infrastructure. As an engineering choice for systems whose program is reducible to state, the snapshot apparatus is a sound default.

What does not follow. The structural premise of snapshot backup is that the program is reducible to its state — that a captured image at moment t suffices to characterise the program at that moment (§6.3). The premise holds for programs whose journal is, in effect, a derivative of state changes. It does not hold for programs whose journal is the program itself.

Under the substrate condition, the artifact that characterises the program is the journal, not any state derived from it. A snapshot at moment t captures *where* the program had arrived; the corpus to t captures *how it got there*. The two are not equivalent informational artifacts: from the snapshot, the operational history is lost; from the corpus, the snapshot at any past moment is recoverable by replay. The substrate-based account of backup is therefore not simply a different operational pattern with the same content; it is the strictly more informative choice when the construct’s conditions hold.

The cost is honest. A continuously consuming destination, even one whose role is to hold the corpus and never replay it, must be reachable, must persist what arrives, and must catch up when it falls behind. §5.3’s measurement (L4) reports steady-state convergence of **$\approx 2\,150$ events/sec on FileSystem**, **≈ 952 events/sec on MySQL**, and **≈ 766 events/sec on SQL Server**, with per-instance verification reaching equal in-memory state in **18 of 18 cells across three backends and three journal sizes**. The numbers establish that the operational obligation is bounded — a destination that ingests at the rates above is not a degenerate case — and that the corpus arrives identically across heterogeneous backends. The objection’s premise that simplicity favours snapshot is honest engineering; the construct’s reply is that the simpler artifact carries less information, and the choice between holding state and holding program is the choice the substrate condition makes available.

9.4 “Append does not scale beyond N writes/sec”

Canonical assumption. A journaled append is bounded by the per-store write throughput of the underlying medium, and that ceiling constrains any substrate-based account to per-actor write rates the medium can absorb.

The objection. Every persistent store has a write-throughput ceiling. A substrate that records every operation as a journal append inherits whatever ceiling the underlying store imposes; production systems that exceed that ceiling are constrained either to batched-write modes or to relaxed-durability commits. The substrate condition therefore does not generalise to high-throughput workloads, and the four equivalences hold only for systems that fit within the per-actor write rate the journal can absorb.

What is valid. A local-journal substrate, like any persistent store, has a per-entry append latency that bounds its write throughput. §5.5’s measurement (L6) reports **328 μ s p50 on the local FileSystem backend** under one-entry one-durable-commit semantics, corresponding to roughly three thousand entries per second of single-writer append rate. A workload that exceeds that rate on a single actor cannot be absorbed by single-writer append at the same durability guarantee.

What does not follow. The single-actor ceiling is not a ceiling on the construct; it is a property of the configuration in which one actor handles one event stream. The substrate condition is preserved under sharding (one actor per shard, each with its own journal — the per-actor invariant of Paper 3, §6.2 and §6.8) and under local buffering with asynchronous replication to a canonical backend (§5.4’s offline equivalence applies to the steady-state régime, not only to outage absorption). Neither sharding nor buffering is foreign to the construct; both fall out of the substrate’s locality property — the journal is the actor’s, not the system’s.

The honest scope is therefore narrower than the objection suggests. The construct does not claim that a single actor with one journal absorbs unbounded write rates; it claims that within the rate the journal can absorb at a given durability guarantee, the four equivalences hold. §5.5’s table records that the local-journal régime favours the substrate by $\approx 3\times$ **over both co-located RDBMS engines at p50 and at the long tail**, and that the advantage is structural (both RDBMS engines land in the same band). The trade-off is not journal-vs-no-journal; it is local-append-vs-network-commit, and the substrate’s ceiling is the higher of the two before any sharding is considered.

Relaxing durability changes the constant. Group commit, batched inserts, and asynchronous-durability modes on the RDBMS side close part of the gap by trading durability for throughput; the substrate’s local-append rate can be raised similarly by group-flush semantics. §5.5’s caveats name both choices as honest engineering. What the comparison preserves under all durability régimes is the structural claim: replay-based operations inherit append latency as their dominant cost term, and a local append against a substrate is, in every measured

régime, at least as fast as a comparable commit against a co-located transactional store.

10. Conclusion

The canonical literature on distributed systems has long treated deployment, cross-datacenter replication, backup, and offline operation as four operational disciplines, each supported by its own apparatus. This treatment proved productive: blue-green environments, change-data-capture pipelines, snapshot agents, and message-queue buffers matured into reliable engineering practice. From that productivity grew the fragmentation analysed in §2 and §6 — four parallel literatures, four parallel apparatuses, and four parallel cost models, borne independently by systems that must be zero-downtime, cross-datacenter, recoverable, and tolerant of disconnection.

This paper observes that the fragmentation is not entailed by the operations themselves. It follows from the assumption named in §3: that operational concerns are addressed by apparatus around the program rather than by the program itself. Under the substrate condition established by prior papers in this series, that assumption no longer holds. The journal is the program written out over time, and four equivalences follow from reading it under variations of when and where (§4.2): deployment is replay at the present head; replication is the same replay at another site; backup is the corpus held without replay; offline operation is replay after a delay.

Viewed from this condition, the apparatuses surveyed in §6 appear as mechanisms that reconstruct, from outside the program, work that the substrate already performs by construction. Blue-green deployment reconstructs replay manually; change-data-capture reconstructs operations from state; snapshot backup captures a derivative of the program’s history; message-queue buffers introduce a durable artefact to absorb disconnection that the substrate’s own persistence already absorbs.

The contribution of this paper is conceptual. The instantiation in production serves as an existence proof; the measurements in §5 exhibit the régime in which the four equivalences are observable. One artefact — the substrate — supports four operational readings without requiring separate mechanisms.

For a journaled program, deployment is replay, replication is sharing history, backup is copying the program, and offline operation is delayed replay. The four operational disciplines named in §2 remain meaningful; what changes is the structural locus of operational concerns: from externally-added apparatus to internal readings of the program’s substrate.

Appendix A. Source-code verification

The constructs named in §7 are publicly available in the Puppeteer codebase. The line references below are pinned against the master branch as of the drafting of this section. Three tables organise the surface by sub-section: the *aliveGate* mechanism (§7.1), the push-to-pull primitives (§7.2), and the *Materialize* construct (§7.3).

§7.1 — Theater.aliveGate

Surface	File	Line(s)	What it shows
<code>aliveGate</code> <code>manual-reset</code> <code>event</code>	<code>Choreography/Theater/Performance.cs</code>	34	Reset while follower or during handover; Set when alive
<code>Start(bool asFollower = false)</code>	<code>Choreography/Theater/Performance.cs</code>	100	Two-mode start; sets gate iff primary
<code>WaitUntilAlive(CancellationToken)</code>	<code>Choreography/Theater/Performance.cs</code>	136	External callers block here
<code>IsAlive</code> predicate	<code>Choreography/Theater/Performance.cs</code>	167	False while follower; true after handover
<code>LockWhileNotSynchronized()</code>	<code>Choreography/Theater/Performance.cs</code>	177	Pause writes on primary; return <code>EntryId</code>
<code>UnlockAndRunAlive()</code>	<code>Choreography/Theater/Performance.cs</code>	183	Flip gate; release primary
<code>CatchUpFromJournal(long targetEntryId)</code>	<code>Puppeteer/EventSourcing/ActorHandler.cs</code>	2159	Replay residual tail under write lock
<code>RehydrateFromEvent</code>	<code>Puppeteer/EventSourcing/DB/Diary.cs</code>	228	Diary-side bulk replay primitive

§7.2 — Push-to-pull primitives

Surface	File	Line(s)	What it shows
Diary.OnRecordWritten setter	Puppeteer/EventSourcing/DB/Diary.cs	145	Per-record callback (single subscriber)
AddRecordWrittenCallback	Puppeteer/EventSourcing/DB/Diary.cs	160	Chain N subscribers without displacing existing
ReactionMode.Cue	Puppeteer/EventSourcing/Follower/Reaction.cs	19	Continuous push-mode Reaction
Cue batch → push loop wiring	Puppeteer/EventSourcing/Follower/Reaction.cs	563-579	Catch-up batch then subscribe to callback
SVIX integration	—	—	Not present in Puppeteer/ or Choreography/; wired at the per-API host process

§7.3 — Materialize construct

Surface	File	Line(s)	What it shows
Performance.Start(asFollower: true)	Choreography/Th eater/Performance.cs	100, 118	Alternate startup mode; SuppressReactionJournaling flag
Reaction.MaterializeFromMetadata(destination)	Puppeteer/EventSourcing/Follower/Reaction.cs	669	DSL marker plumbing: records destination on action
actor.Materialization sub-namespace	Puppeteer/Materialization.cs	(whole file)	Register / Deregister / List / ReadRecordsAfter / ConfirmUntil / ReadReactions / ReadElidedRange

Surface	File	Line(s)	What it shows
MaterializeMirror fluent client	Puppeteer/MaterializeMirror.cs	(whole file)	<code>Sync()</code> (Layer 1); <code>AsProgramMirror().Sync()</code> (Layer 2)
MirrorSyncResult immutable struct	Puppeteer/MaterializeMirror.cs	170-202	Records / ReactionsSnapshot / ElisionMarkers + watermark advance
Diary backend polymorphism	Puppeteer/EventSourcing/DB/Diary.cs	53-78	(<code>DatabaseType, connectionString</code>) $\rightarrow$ InMemory / FileSystem / MySQL / SQLServer
Local-buffer (offline régime)	Puppeteer/EventSourcing/DB/Diary.cs	28, 34, 80-135	<code>IsBuffered</code> ; <code>ReplicationAgent</code> async drain to canonical storage

Code provenance

Source-code references in this paper resolve against the public Puppeteer repository at commit **99e8202** (2026-07-13). The snapshot is archived in Software Heritage under the following persistent identifier:

```
swh:1:dir:c55af9b6466302ed402cb579440ceec870e449c8;
  origin=https://github.com/alvaroNCubo/puppeteer;
  anchor=swh:1:rev:99e8202c808c3d79ca7b1dfbdae0ffb7872c940e
```

Inline references of the form `file.cs:NN` (e.g., `ActorHandler.cs:38`) resolve against this snapshot. A reader can construct a per-file SWHID by adding the qualifiers `;path=<path>;lines=<NN>` to the directory SWHID above. Future commits to the repository may renumber lines; the SWHID preserves the cited state independently of any future change to the repository or its hosting.

The six laboratories' source and datasets accompany this paper as **paper05-data.zip**. The aggregated results each lab reports — the headline and summary tables this paper cites — are included in full. To keep the archive light, Lab 6's raw per-event sample log (`samples.csv`, $\approx$ 60 MB / $\sim$ 500 000 rows) is omitted and replaced by a note giving the lab source and the runtime commit needed to regenerate it.

Acknowledgments

The author used large language models (including Claude and ChatGPT) as editorial assistants for language refinement, structural feedback, and literature navigation. All original ideas, terminology, theoretical constructs, and technical content presented in this work are solely the author's.

Appendix B. Bibliography

Ambler, S. W., & Sadalage, P. J. (2006). *Refactoring databases: Evolutionary database design*. Addison-Wesley.

AWS. (n.d.). Managing your storage lifecycle. Retrieved from <https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lifecycle-mgmt.html>

Azure. (n.d.). Overview of operational backup for Azure Blobs. Retrieved from <https://learn.microsoft.com/en-us/azure/backup/blob-backup-overview>

Bacula. (n.d.). Bacula backup software. Retrieved from <https://www.bacula.org/>

Burns, B., Grant, B., Oppenheimer, D., Brewer, E., & Wilkes, J. (2016). Borg, Omega, and Kubernetes. *Communications of the ACM*, 59(5), 50–57.

Debezium. (n.d.). Debezium. Retrieved from <https://debezium.io/>

DeCandia, G., Hastorun, D., Jampani, M., Kakulapati, G., Lakshman, A., Pilchin, A., Sivasubramanian, S., Voshall, P., & Vogels, W. (2007). Dynamo: Amazon's highly available key-value store. In *Proceedings of the 21st ACM SIGOPS Symposium on Operating Systems Principles (SOSP '07)* (pp. 205–220).

Flyway. (n.d.). Flyway database migrations. Retrieved from <https://flywaydb.org/>

Fowler, M. (2010). BlueGreenDeployment. *martinfowler.com*. Retrieved from <https://martinfowler.com/bliki/BlueGreenDeployment.html>

Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.

Hevner, A. R., March, S. T., Park, J., & Ram, S. (2004). Design science in information systems research. *MIS Quarterly*, 28(1), 75–105.

Hitz, D., Lau, J., & Malcolm, M. (1994). File system design for an NFS file server appliance. In *Proceedings of the USENIX Winter Technical Conference*.

Humble, J., & Farley, D. (2010). *Continuous delivery: Reliable software releases through build, test, and deployment automation*. Addison-Wesley.

- Kreps, J., Narkhede, N., & Rao, J. (2011). Kafka: A distributed messaging system for log processing. In *Proceedings of the NetDB Workshop, 6th International Workshop on Networking Meets Databases* (co-located with SIGMOD 2011).
- Kubernetes. (n.d.). Deployments: Rolling update deployment strategy. Retrieved from <https://kubernetes.io/docs/concepts/workloads/controllers/deployment/#rolling-update-deployment>
- Lakshman, A., & Malik, P. (2010). Cassandra: A decentralized structured storage system. *ACM SIGOPS Operating Systems Review*, 44(2), 35–40.
- Liquibase. (n.d.). Liquibase. Retrieved from <https://www.liquibase.org/>
- Maxwell. (n.d.). Maxwell’s daemon. Retrieved from <https://maxwells-daemon.io/>
- MySQL. (n.d.). MySQL 8.0 reference manual: Replication (Chapter 17). Retrieved from <https://dev.mysql.com/doc/refman/8.0/en/replication.html>
- PostgreSQL. (n.d.). Continuous archiving and point-in-time recovery (PITR). Retrieved from <https://www.postgresql.org/docs/current/continuous-archiving.html>
- RabbitMQ. (n.d.). Dead letter exchanges. Retrieved from <https://www.rabbitmq.com/dlx.html>
- Redis. (n.d.). Redis replication. Retrieved from <https://redis.io/docs/management/replication/>
- Rivera, A. (2026a). Anti-porous architecture: a unified design principle for CQRS + Actor + Event-Sourcing systems. *Puppeteer Papers Series*, Paper 1 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20404863>
- Rivera, A. (2026b). Program–value separability: the structural precondition for compilation, caching, and dense journaling in a DSL runtime. *Puppeteer Papers Series*, Paper 2 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20740697>
- Rivera, A. (2026c). Reactions and the partition: opt-in eventual consistency in actor-native systems. *Puppeteer Papers Series*, Paper 3 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20792156>
- Rivera, A. (2026d). Preserving semantic continuity across actors: a tell-based approach without orchestration. *Puppeteer Papers Series*, Paper 4 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21207062>
- Sato, D. (2014). CanaryRelease. *martinfowler.com*. Retrieved from <https://martinfowler.com/bliki/CanaryRelease.html>
- Veeam. (n.d.). Veeam backup & replication user guide. Retrieved from <https://helpcenter.veeam.com/docs/backup/vsphere/overview.html>

Most infrastructure layers are symptoms of the persistence model: a construct for auditing production stacks

TL;DR

Many layers taken for granted in production systems — caches, ORMs, message queues, distributed locks, application servers, orchestration clusters — do not exist because the problem requires them, but because the underlying persistence model does.

This paper names that property *infrastructural symptom*, defines the two diagnostic indicators that distinguish it from a structural requirement, provides an auditor workflow to read any layer under those indicators, and shows a journaled actor-native system in which the canonical symptoms disappear.

The construct is diagnostic, not prescriptive: it makes visible which parts of a stack compensate for the model rather than serve the problem.

Claims this paper makes

1. Some infrastructural layers in production deployments exist to compensate for structural deficiencies of the underlying persistence model rather than for structural requirements of the problem being solved (§3).
2. The property naming this distinction — *infrastructural symptom* — is read through two diagnostic indicators: compensation (the layer compensates for an identifiable model defect) and dissolution (the layer loses justification when the defect is removed) (§3).

3. An auditor workflow guides the application of the two indicators and produces a reading — model-dependent, structural, or unclassified — per (layer, use-case) pair (§5).
4. A journaled actor-native model instantiates a system in which four mechanisms — locality, journal density, journal-as-causal-substrate, self-contained substrate — dissolve the seven canonical categories of compensating infrastructure (§4.1).
5. The Redis case admits the workflow with per-use-case reading: six use-cases dissolve as symptoms; two remain structural (§4.2).
6. The ladder of compensated layers is open-ended; the workflow extends to backup tooling, service mesh, distributed-tracing/APM, and future emergent layers (§6).
7. The model permits appliance-class deployment as a limit case in which categories of business software become physically shippable — an architectural permission, not an operational proof (§7).

1. Introduction

This paper makes an analytic contribution: a diagnostic typology for reading existing stacks. It identifies and names a structural pattern that prior literature has documented case by case without articulating as one (the typology); derives the indicators under which the pattern reads as model-dependent (the criteria); and presents an instantiation in which those indicators read positively and the pattern is absent (the existence proof). The genre is closer to *theory for analyzing* in the sense of Gregor (2006) than to design theory: evidence is presented in the form of a working artifact, but the artifact illustrates the lens’s reading rather than validating it. The contribution is conceptual; the instantiation confirms that the pattern is contingent rather than structurally necessary, not that any specific system is the only way to dissolve it.

Production software systems are routinely deployed alongside layers of supporting infrastructure: in-memory caches, object-relational mappers, message queues, distributed locks, application server frameworks, orchestration clusters, full operating systems. These layers are treated as standard components of any non-trivial deployment. Architectural discussions assume their presence by default; tutorials and reference architectures position them as foundational; engineering teams budget for them, staff them, and monitor them. Their absence is treated as a property of toy systems.

This paper names a property that some — not all — of these layers share. We call this property *infrastructural symptom*: a layer is an infrastructural symptom when it exists to compensate for a structural deficiency in the underlying persistence model rather than for a structural requirement of the problem being solved. The distinction is routinely elided; the two cases are treated alike in architecture, in tooling, and in engineering culture. Once the property is named, each layer in a given stack becomes auditable under the question: is

this compensating for something, and if so, for what?

Section 2 situates the paper among prior work. Section 3 defines the construct and derives the two diagnostic indicators that distinguish infrastructural symptom from structural requirement. Section 4 presents the instantiation: a journaled actor-native system in which the canonical layers — caches, glue queues, locks, application server scaffolding — are absent not by omission but by structural irrelevance. Section 5 specifies the auditor workflow for reading a stack under the typology. Section 6 generalizes from the canonical case to a chain of compensated layers, from in-memory cache up through full operating system. Section 7 examines the limit case, where the chain collapses to a domain artifact deliverable as an appliance. Section 8 examines what the lens makes visible and what it does not.

2. Related work

The genre of this paper is closer to *theory for analyzing* (Gregor, 2006) than to design theory. An analytic paper of this kind identifies a typology, derives the indicators under which the typology reads, and presents an instantiation that demonstrates the reading produces non-trivial discrimination. The contribution is conceptual; the instantiation is the existence proof, not the substance of the claim.

Critique of DBMS-centric architecture as the dominant paradigm for production systems has accumulated in the literature without naming the construct presented here. Stonebraker and Çetintemel (2005) observe that the relational database has become a one-size-fits-all hammer applied to use-cases for which it is structurally ill-suited, and predict the proliferation of specialized stores. Kleppmann (2017) catalogs the patterns by which contemporary deployments combine caches, queues, ORMs, and search indices around a relational core, documenting their pairwise interactions and operational costs. Helland’s *Life beyond Distributed Transactions* (2007) makes a parallel observation about the proliferation of coordination machinery under distributed-transaction semantics. His *Immutability Changes Everything* (2016) extends a related argument to the substrate itself: append-only representation dissolves much of that machinery by construction. The present paper names what these critiques document piecewise — the unifying property of compensation — and renders the discrimination available as a diagnostic lens.

The instantiation presented in §4 is constructed from three lines of prior work. The actor model originates with Hewitt (1973), articulating computation as message passing between isolated entities. Event sourcing as a persistence pattern is associated with Fowler (2005), which characterizes the program’s history as the durable artifact rather than a snapshot of current state. Command Query Responsibility Segregation, articulated by Young (2010), separates the read and write paths so they no longer compete for the same model. Vernon (2015) brings

these three together as practical patterns in *Reactive Messaging Patterns with the Actor Model*; that synthesis is the immediate ancestor of the present instantiation. The journaled actor-native system used as the existence proof here combines these threads; it does not invent them.

This paper is the sixth in a series. The series develops the construct’s foundations across five preceding papers. Paper 1 introduces porosity as the representational defect that this paper’s construct identifies in its compensation form. Paper 2 develops program-value separability and the externalized-parameters precondition that supports journal density as a §4.1 mechanism. Paper 3 introduces Reactions and the pragmatic partition between work-due-before-responding and deferred work, which §4.1 invokes as the journal-as-causal-substrate mechanism. Paper 4 introduces the Tell primitive and cross-actor causal continuity, used in §4.1 for the same mechanism. Paper 5 develops the journal as substrate for deployment, replication, backup, and offline operation, which §6 references as the substrate-mediated lifecycle that dissolves application-server scaffolding. Paper 7 carries the same discrimination upward from infrastructural layers to architectural roles; §9 introduces this extension.

The contribution unique to this paper is conceptual. The construct *infrastructural symptom* names a property that prior work documents in pieces but does not unify. The two diagnostic indicators — compensation and dissolution — give the property its shape. The auditor workflow of §5 guides its application without requiring adoption of any particular alternative. The ladder of §6 and the limit case of §7 extend the typology’s reach. The instantiation of §4 demonstrates that the typology reads cleanly under the indicators; the demonstration is not the contribution.

3. The construct: infrastructural symptom

We state the construct before naming its two diagnostic indicators and the workflow that guides their application.

An infrastructural symptom is a layer of infrastructure that exists to compensate for a structural deficiency in the underlying persistence model rather than for a structural requirement of the problem being solved.

The construct operates at the level of purpose, not at the level of product. A given physical layer in a deployment may serve multiple purposes; the construct discriminates among those purposes one at a time. The unit of analysis is the use-case, not the layer.

A layer reads as an infrastructural symptom for a given use-case when both of two diagnostic indicators are positive.

The compensation indicator. The layer’s purpose for that use-case is traceable to a specific deficiency of the underlying persistence model. The deficiency must be identifiable: not “we use this layer because it is fast,” but “we use this

layer because the persistence model exhibits property P, and this layer compensates for P.” When the compensation indicator is positive, asking what would happen if P were not present? yields a non-trivial answer.

The dissolution indicator. When that deficiency is removed — by adopting a persistence model in which P does not arise — the layer loses justification for that use-case. The deficiency’s absence is sufficient to make the layer redundant for that purpose; nothing else about the problem domain, workload, or operational context needs to change.

The dissolution indicator requires that M’ — the counterfactual model in which P does not arise — be *constructible*: a persistence model with operationally defined semantics, viable under the workload, hardware, and programming-model constraints of the use-case under examination. A counterfactual whose viability is unspecified — an imagined absence of P without a model that exhibits the absence in operational form — does not count as a model in which D does not arise; it only restates the wish that D were absent. The lens does not admit such counterfactuals. The instantiation of §4 is offered precisely as one such constructible M’, sufficient to make the dissolution indicator non-vacuous for the seven canonical categories.

Both indicators must be positive. Compensation without dissolution describes a genuine patch whose value survives model evolution. Dissolution without compensation is not constructible: a layer cannot dissolve under a model change unless its purpose was tied to the model in the first place.

The construct is not a universal accusation against infrastructure. Many layers in production systems exist as structural requirements — they serve purposes rooted in the problem domain, independent of any persistence model. Three categories are typical:

- **External-system integration.** Layers mediating communication with systems outside the boundary — payment gateways, partner APIs, sensor networks — compensate for no model deficiency. The integration problem remains under any persistence model.
- **Cryptographic and protocol enforcement.** TLS termination, request signing, OAuth verification address security or protocol requirements rooted in the problem domain. No persistence-model change makes encryption optional.
- **Genuinely expensive computation.** Analytics engines, machine-learning inference services, video encoders exist because the computation itself is intrinsically expensive. That cost does not change under a different persistence model.

The boundary between symptom and structural requirement is not a property of the layer as a product but of the purpose for which it is used. The same caching system, deployed to absorb read-vs-write contention at the persistence layer, instantiates an infrastructural symptom. Deployed instead to memoize

a deliberately expensive computation, it instantiates a structural requirement. The reading is produced per use-case.

Seven categories of layer are encountered regularly in production deployments and admit reading under the two indicators. The third column describes, in model-property terms rather than system terms, what a persistence model under which both indicators read positively would offer instead.

These seven are not a privileged set but an entrenched one: the layers most consistently assumed by default in conventional production stacks, taught as foundational, and budgeted as standing operational concerns. Other commonly-deployed layers — logging frameworks, schedulers, service discovery, secret managers, feature-flag servers — admit reading under the same workflow of §5 and produce more varied results. Some dissolve (in-process logging, internal scheduling and discovery); some are structural (secret managers, whose purpose is rooted in access control rather than persistence; the A/B-testing facet of feature-flag servers, rooted in product methodology); some split per use-case (logging at internal vs. external boundaries; service discovery at internal vs. external). The lens discriminates; Table 1 catalogs the architectural baseline against which the discrimination is most consequential.

Table 1. Diagnostic reading: seven canonical categories of compensating infrastructure under the two indicators of §3.

Layer	Underlying defect compensated	What a model with both indicators positive offers instead
In-memory cache	Reads contend with writes at the persistence layer; reconstructing state from storage is expensive	State local to the unit of computation; no contention between read and write paths
Object-relational mapper	Domain objects do not align with relational rows; mapping is opaque and lossy	Domain objects are the unit of persistence; no translation layer
Message queue used as glue	Components cannot directly invoke deferred work without losing causal record	Deferred work is journal-recorded and replayable; queues are not the substrate of causality
Distributed lock	Multiple writers contend for the same logical entity across processes	A single point of authority per logical entity; serialization is intrinsic, not externalized

Layer	Underlying defect compensated	What a model with both indicators positive offers instead
Application server framework	The persistence model leaks into application code; lifecycle plumbing must be added	The domain is the application; lifecycle is a property of the substrate, not added scaffolding
Orchestration cluster as coordination substrate	Stateful coordination across instances requires external choreography	State has location; coordination is between actors, not between containers
Full operating system	The runtime, the framework, and the application require disparate substrates	A minimal substrate suffices to run the domain artifact

The seven rows are not exhaustive. Backup tooling, service mesh, and distributed-tracing frameworks admit similar analysis and are discussed in Section 6. Whether such a persistence model exists, and at what cost, is the subject of Section 4.

The construct is diagnostic, not prescriptive. It enables auditing existing stacks under the two indicators; it does not specify how to build a model in which infrastructural symptoms do not arise. Construction is the work of Section 4, which exhibits one such model, and of Section 5, which specifies the auditor workflow that produced the diagnostic table.

The value of the typology is diagnostic. A practitioner who does not adopt the instantiation of Section 4 can nevertheless use it to read which layers in their stack are tied to their persistence model and which are tied to their problem domain. That visibility is independent of the choice to act on it.

4. Instantiation

4.0 Genealogy

Puppeteer’s earliest layer dates to 2005, when a persistence library internally called *autopersistencia* adopted a single design rule: the classes representing the domain would carry no persistence concerns. Persistence would be discovered by reflection over the class hierarchy; storage details lived elsewhere. The rule was modest in scope and pragmatic in motivation. Its consequence was structural: the domain became a clean substrate, unobstructed by the technical aspects that conventional architectures embed inside it.

At no point in this evolution was there an intention to eliminate caches, queues, locks, or application scaffolding; those layers were assumed to belong in any serious system. Their later absence was not a design objective but an empirical outcome. This distinction matters for the argument of this paper: the construct is not dissolved here by architectural ideology but as a byproduct of unrelated design constraints that happened to satisfy the indicators derived in §3.

The natural follow-up question was: if domain classes carry no persistence concerns, where does state live? The answer that emerged across the following years became the framework’s first generation. State lives in a journal, and the journal records the program itself — not serialized events, not row deltas, but the DSL script that produced the state, replayable in order. The narrative of execution, not the photograph of memory, became the unit of persistence — a property whose absence, in the terms of §3, is the deficiency that caches, ORMs, and queue-glue exist to compensate for in DBMS-centric architectures.

The framework’s second generation added four properties, each developed in preceding papers of this series. Parameters were externalized, allowing programs to be compiled, cached, and recorded as references rather than literal scripts (Paper 2). Reactions were introduced as the developer-controlled partition between work due before responding and work deferred to placement (Paper 3). The Tell primitive established cross-actor causality recorded in the journal rather than orchestrated externally (Paper 4). The journal itself became the catalog of the actor’s declarative vocabulary, dissolving the lateral storage of action definitions that the framework had carried since V1.

The framework was not designed to dissolve the construct named in this paper. The construct was identified retrospectively, after the chain of architectural decisions described above had converged on a system in which most of the canonical layers had become unnecessary. The work of this paper is, in that sense, a forensic reading of an existing artifact, not an engineering plan derived from a principle. The principle was a property of the artifact before it was a property of the literature.

4.1 The instantiation under the two indicators

Four mechanisms in the journaled actor-native model account for the dissolution of the seven canonical categories of §3. A single mechanism may dissolve more than one category, and the in-memory cache dissolves under two mechanisms because §3 records two distinct defects compensated by that layer. Each mechanism decomposes into several sub-mechanisms — properties of the model that, taken together, exhibit the umbrella behavior. Table 2 correlates each mechanism with its sub-mechanisms, the §3 defects it dissolves, and the canonical layers that become symptoms under it.

Table 2. Mechanisms in the journaled actor-native model mapped to the §3 defects they dissolve and the canonical layers that become symptoms.

Mechanism	§3 defect compensated	Canonical layer that becomes symptom
Locality • privacy of state • actor-as-serializer • placement	Reads contend with writes at the persistence layerMultiple writers contend for the same logical entity	In-memory cache (contention defect)Distributed lock
Journal density • journal-as-program (homoiconic) • replay-as-reconstruction • domain-class-as-unit-of-persistence	Reconstructing state from storage is expensiveDomain objects do not align with relational rows	In-memory cache (reconstruction defect)Object-relational mapper
Journal as causal substrate • Reactions recorded in journal • Tell recorded in journal • causality boundary is the journal	Components cannot directly invoke deferred work without losing causal recordStateful coordination across instances requires external choreography	Message queue used as glueOrchestration cluster
Self-contained substrate • domain artifact is the application • substrate carries lifecycle, persistence, dispatch • journal-mediated release handoff • minimal OS surface	The persistence model leaks into application codeThe runtime, framework, and application require disparate substrates	Application server frameworkFull operating system

Locality. State is private to the actor that owns it; the actor processes commands and queries serially against its own memory. There is no shared store from which concurrent reads must be isolated from concurrent writes. Two of the §3 defects vanish mechanically under this property: the read-versus-write contention that motivates in-memory caches, and the inter-process contention that motivates distributed locks. The compensation indicator is vacuous for both. Placement is the developer-controlled partition treated in Paper 3. This mechanism reads positively under the actor-native preconditions established across the preceding papers in this series: the actor’s state fits in memory and is processed serially over its private storage (Papers 1 and 3), and cross-actor continuity is mediated through the journal by Reactions and Tell rather than through shared queries (Papers 3 and 4). Aggregates that exceed the in-memory precondition — extremely hot single entities, datasets larger than a single host’s working set

— fall outside this reading and require actor splitting or partitioning, which is a design technique distinct from the lens itself.

Journal density. The journal does not record serialized snapshots or row-by-row deltas; it records the program — the DSL script — that produced the state. Reconstruction is replay, not deserialization; replay of a compact script is cheap. The domain class is the unit of persistence, with no intermediate row to map. Both §3 defects compensated by caches and ORMs — reconstruction expense and object-row impedance — are dissolved mechanically. The mechanism is treated as compilation in Paper 2 and as homoiconic representation in Paper 1. Implementation: `Puppeteer/EventSourcing/DB/Diary.cs`. Replay cost for long-lived actors is bounded by complementary substrate primitives: Distill compacts the journal in place (`Puppeteer/EventSourcing/DB/Diary.cs:264`); elision markers and Materialize watermarks enable offload-and-trim through a materialization destination (Paper 5, §7.3). Together these primitives fill the role that Rolling Snapshots (Young, 2010) play in conventional event-sourced systems, without re-introducing snapshot semantics.

Journal as causal substrate. Pragmatic deferral, developed in Paper 3 as Reactions, records work the verb does not perform before responding as a journal entry, with the causal link to the originating event preserved. Cross-actor continuity, developed in Paper 4 as Tell, records dispatch from one actor to another in the journal as well. The journal becomes the boundary of causality. In DBMS-centric systems, causality crosses the persistence boundary and must be reconstructed with message queues used as glue and orchestration clusters used as coordination substrates; here, causality never leaves the journal. Implementation: the `reactions` field in `Puppeteer/EventSourcing/ActorHandler.cs:38`.

Self-contained substrate. The runtime is a small loader; the application is the domain artifact. Lifecycle, persistence, and dispatch are properties of the substrate, so no external process is required to host or coordinate the application artifact. Release handoff and rolling deployment, handled in DBMS-centric deployments by application server orchestration, are journal-mediated here and are the subject of Paper 5. The application server framework as a category has no equivalent role; the operating system surface is reduced to the minimum required, a limit case treated in §7.

Together, these four mechanisms account for the dissolution of the seven canonical categories of §3 mechanically, not by intent. None was added to the model with the construct of this paper in mind; each was developed for the reasons enumerated in §4.0.

4.2 The Redis case in detail

Redis occupies a distinctive position in the contemporary infrastructure landscape: a single product whose presence is taken for granted across deployments that are otherwise architecturally dissimilar. The diversity of its typical use-cases — cache, session store, pub/sub, queue, distributed lock, leaderboard,

rate limiter — does not reflect the diversity of a single problem domain. It reflects the diversity of the pressures that DBMS-centric architectures present to application designers, each of which Redis conveniently absorbs in a familiar form. This makes Redis the canonical case for inspection under the construct of §3. The analysis that follows is not, however, about Redis as a product; it is about the defects of the underlying persistence model that Redis exists to absorb — defects that other compensating layers also address in variant forms, treated briefly after the canonical table.

Table 3 applies the two indicators of §3 to the most common Redis use-cases, tracing each to the specific defect it compensates and the mechanism in §4.1 that dissolves it.

Table 3. Application of the diagnostic typology to canonical Redis use-cases.

Redis use-case	§3 defect compensated	Native response in journaled actor-native
In-process cache for database results	Reads contend with writes at the persistence layer; reconstructing state from storage is expensive	Actor state is the cache (§4.1 Locality + Journal density)
Session storage	Stateless application servers cannot retain user state across requests	The user actor retains session naturally as part of its state (§4.1 Locality + Self-contained substrate)
Pub/sub among internal components	No causal record of cross-component events; coordination must be reconstructed	Reactions record deferred work in the journal with causal links preserved (§4.1 Journal as causal substrate)
Message queue for deferred work	Components cannot directly invoke deferred work without losing causal record	Reactions and Tell record causation in the journal (§4.1 Journal as causal substrate)
Distributed lock for shared entity	Multiple processes write the same logical entity across boundaries	The actor owns its entity; serialization is intrinsic, not externalized (§4.1 Locality)

Redis use-case	§3 defect compensated	Native response in journaled actor-native
Leaderboard / sorted set	Real-time aggregation against the persistence layer is expensive	Query against actor state or a follower; replay is cheap (§4.1 Locality + Journal density)

Each row reads positively under both the compensation and dissolution indicators of §3. The pub/sub row covers internal-component coordination; the external-systems variant is structural and is treated below.

Two Redis use-cases do not read positively under the dissolution indicator and are therefore structural under the typology. The first is pub/sub directed at external systems: webhooks to third-party consumers, event distribution to partners, broadcast to subscribers outside the operator’s boundary. The external system does not dissolve under any persistence-model change; the distribution problem persists. The second is rate limiting applied to external API consumption: per-customer quotas, third-party API throttling, abuse mitigation against external traffic. The rate limit is a property of the external relationship, not of the internal model. Redis, or any comparable substrate, remains the appropriate tool for both.

Other compensating layers in the DBMS-centric ecosystem exhibit variants of the same pattern. Memcached shares the cache analysis: locality and journal density make in-process caches mechanically redundant. Distributed coordinators such as ZooKeeper share the lock and coordination analysis: the actor model makes the underlying contention vacuous. Message brokers such as RabbitMQ share the queue analysis: Reactions and Tell record causation in the journal.

Event-storage products such as EventStore DB and Kafka deserve a separate observation. Both adopt a journaled posture — they record events rather than serialized snapshots, and they offer replay. But they instantiate the journal as an external infrastructure layer that the application consumes through a client API, not as the substrate of the application itself. Causality is recorded but remains external to the unit of computation; coordination between event production, event consumption, and application state still requires application-side glue. The product captures the right idea and externalizes it, leaving the application code in the compensating position. Under the construct, these products do not dissolve the symptom; they relocate it.

A second observation about these compensating layers is that they do not absorb their compensation freely. Each requires the application to carry the boilerplate of its use: cache invalidation logic, lock acquisition and release with

retry and fencing semantics, serialization between domain objects and external representations, key naming and TTL discipline, consumer offset management and rebalancing handlers, schema registry coordination, idempotency keys for at-least-once delivery. The application code carries the porosity imposed by RDBMS representations (Paper 1) as its baseline; the compensating layer adds a second porosity, this time imposed by the layer itself: its own API surface, lifecycle, and consistency contracts. Domain logic interleaves with both. Under the instantiation, the compensating layer and the boilerplate that the application carries to use it are both absent. The actor's code addresses the problem domain; the substrate addresses the rest.

Redis as a product is not what the typology identifies; nor is any single compensating layer treated above. What the typology reads is the architectural pattern in which a persistence model exposes defects and successive layers — caches, brokers, coordinators, even externalized event stores — accumulate to absorb them. Replacing one such layer while retaining the model that produces the others does not change how the typology reads. The unit of change is the persistence model itself, not any individual layer.

5. Auditor workflow

The typology of §3 enables reading per use-case, but practitioners face stacks containing dozens of compensating candidates. Without an explicit workflow for application, the typology risks remaining theoretical: invoked at the level of architecture review, but absent from the operating discipline that distinguishes one layer from the next. The workflow that follows guides the application of the two indicators of §3, producing a reading in a finite number of steps for any (layer, use-case) pair under examination. It can be applied without adopting the instantiation of §4.

Input: layer L deployed for use-case U.

Output: reading in {model-dependent, structural, unclassified};
defect D identified when model-dependent.

Step 1 – Compensation indicator

Identify the defect D of the underlying persistence model for which L compensates in use-case U.

If a defect D is identifiable, proceed to Step 2.

If no defect is identifiable AND L's purpose is rooted in the problem domain (external-system integration, cryptographic or protocol enforcement, intrinsically expensive computation), return structural.

If no defect is identifiable AND L's purpose is opaque to the auditor (organizational, regulatory, historical, team-skill), return unclassified.

Step 2 – Dissolution indicator

Consider a counterfactual model M' in which D does not arise.

M' must be constructible: a persistence model with operationally defined semantics, viable under the workload, hardware, and programming-model constraints of U .

If no such M' is constructible, return unclassified.

If L retains a role in M' for U , return structural.

Otherwise return model-dependent with D .

Four notes on application. First, the unit of analysis is the pair (layer, use-case), not the layer alone. The same product may produce different readings for different uses; the workflow must be run once per use-case. Second, identifying the defect in Step 1 requires explicit attribution to a property of the persistence model. “*We use this layer because it is faster*” does not name a defect; “*we use this layer because the model serializes writes against a shared store*” does. The workflow depends on the discipline of this attribution. Third, the *unclassified* reading exists precisely so that the workflow does not silently absorb opaque or non-model-rooted layers as structural by default; an honest auditor records what the lens cannot see rather than forcing a binary outcome. Fourth, the diagnostic table of §3 and the Redis case of §4.2 are worked examples of the workflow applied to canonical and product-specific layers respectively. Each row in those tables is the reading produced by the workflow for one (layer, use-case) pair.

The workflow renders readings; it does not prescribe action. A model-dependent reading does not require that the layer be removed in the current system. It indicates that the layer’s continued presence is structurally tied to retaining the persistence model that produces the defect. Practitioners may rationally retain a model-dependent layer — for legacy preservation, contractual constraint, team familiarity, migration cost — but they do so knowing that the layer carries a debt that another model would not impose. The workflow’s product is visibility, not directive.

6. The ladder of compensated layers

The seven canonical categories enumerated in §3 are not a list. Read in order, they form a ladder of increasing scale: in-memory cache (per-process), object-relational mapper (within-application), message queue and distributed lock (between processes), application server framework (the application stack), orchestration cluster (across machines), and full operating system (the hosting substrate itself). The order is not arbitrary; it follows the expanding radius of the defect the layer compensates for. Each rung compensates for a deeper limitation of the model; each rung dissolves under the alternative for the same reason.

Three additional categories that §3 forward-referenced extend the ladder.

Backup and disaster-recovery tooling — point-in-time restore, write-ahead log archiving, snapshot orchestration — exists because state in DBMS-centric architectures lives in mutable rows whose history is not the primary artifact. Under the model, the journal is the backup, and rehydration is the recovery operation; both reduce to a single substrate property treated more fully in Paper 5.

The service mesh — inter-service routing, discovery, traffic shaping, mTLS, sidecar instrumentation — exists because stateless application servers, by design, do not retain coordination state internally. Under the model, much of this coordination is internal: Reactions and Tell handle deferred work and cross-actor dispatch as journal entries. The parts that cross the operator’s boundary (TLS to external partners, traffic shaping for external clients) remain structural and continue to belong in a mesh or comparable substrate.

The distributed-tracing and APM stack — span propagation, trace ingestion, request-flow visualization, latency aggregation — requires discrimination per use-case, treated in Table 4 below.

Table 4. Use-case-level reading for the distributed-tracing and APM stack.

Observability use-case	Reading	Reason
Distributed tracing to reconstruct what happened inside the system	Model-dependent	The journal already records the trace; replay is deterministic
Distributed tracing across boundaries with external systems	Structural	The external system does not dissolve under any persistence-model change
APM as <i>“find the bug in this incident”</i>	Model-dependent (predominantly)	Replay localizes the bug without requiring external instrumentation
APM as continuous capacity, latency, and error-rate telemetry in production	Structural	Resource consumption is a real-world property, not a model artifact

The first and third rows read as model-dependent; the second and fourth read as structural requirements. Debug as a one-shot operation dissolves more cleanly than continuous telemetry because the journal substitutes directly for the debugger. Continuous telemetry is mixed: the internal-trace component dissolves, but the capacity-and-cross-system component remains. §8 treats this as a worked example of the typology’s reading at the boundary of its applicability.

The typology’s readings do not eliminate the products from a deployment. The framework does not dissolve APM and OTel; it dissolves the internal-tracing use-case while leaving APM and OTel plugged in at the points where they instantiate a structural requirement. The same observation applies to mesh products at external boundaries, and to backup tooling for legacy systems retained alongside a journaled core.

The ladder is not closed. New layers will emerge to compensate for new defects of DBMS-centric architectures — feature-flag servers, schema registries, secret stores, sidecars for auth and identity, configuration-as-a-service products. Each new layer can be audited under the workflow of §5 and read under the typology of §3. The ladder is open upward, sideways, and into the future because the underlying persistence model continues to generate compensations; the workflow walks each new rung as it appears.

7. At the limit: from datacenter footprint to a deliverable appliance

The workflow of §5 walks each rung of the ladder as it appears; §6 showed the ladder extends. The natural question is what remains when the workflow of §5 marks every rung as model-dependent. §7 names that limit.

What remains is the domain artifact, the journal that records its execution, and the minimal substrate that runs them. The substrate retains what is irreducibly needed: a runtime, journal storage, network I/O, and a scheduling loop. Not required are the cluster orchestrator, the application server framework, the object-relational mapper, the external cache, the distributed coordinator, and the general-purpose operating system in its hosting form. The domain artifact is the application; the substrate is what is required to run it; everything else has been audited and removed under the typology.

The artifact’s architectural requirements are independent of scale. The same architecture sits behind a multi-datacenter cluster with cross-region replication (where the journal and Tell handle topology), behind a Kubernetes red-black rolling deployment (where the journal mediates the release handoff), behind a single server (where the artifact runs against local storage), and behind an on-premises appliance hosting an accounting system, a clinic management system, a hardware-store inventory and point-of-sale system, or a professional-services back office. The substrate’s storage requirement reduces to two operations — appending entries and advancing a cursor over them — which together compose the substrate’s single primitive of replay (Paper 5); both are executed by commodity flash at marginal cost. The RAM requirement scales with the domain’s working state, not with the deployment topology. The difference between these shapes is operational variation over the same architecture. What bounds the appliance case is the customer’s domain — its data volume, user count, workload heaviness — not anything the architecture introduces.

This is an engineering-compatibility claim, not an operational demonstration. The deployments of the framework’s home site run in conventional datacenter infrastructure across eight domain installations; no appliance-class on-premises shipment is currently among them. What §7 argues is that the path to that shape is bounded by ordinary engineering parameters — runtime footprint, journal latency on flash storage, fleet management for updates — and not by an architectural obstacle that the lens has not already named. The barriers that conventional architectures encounter at the appliance limit (a DBMS that does not fit the box and its administration overhead, an application server that demands a separate tier, an orchestrator with no peer to coordinate with, a cache for a model that does not need it) are precisely the layers the workflow of §5 reads as symptoms. The model’s permission for the appliance shape is the structural counterpart of those symptoms’ dissolution.

This bears restating because of what it implies. Whole categories of small-and-medium-business software — accounting, clinic management, retail inventory, professional services — are today delivered on infrastructure scaled to a fraction of their actual workload, because the conventional stack obliges a datacenter footprint for a database, an application server, and the orchestration to keep them running. The barrier is not only to delivery but to development: building software for a single hardware store or a single clinic requires the infrastructural expertise of running a datacenter, which is the structural reason such products often do not exist at all. The lens reads each of those layers as a symptom. What §7 names is the limit case made available by that reading: software that is shippable as hardware to the customer, hosted in a single appliance, operated by the customer or a local service provider.

At the structural limit, the architecture permits software to be delivered as a physical artifact: an appliance the customer owns, not a datacenter the vendor leases on their behalf. The barrier to that shape is not in the application or in the customer’s needs — it is in the infrastructure that conventional architectures require. The barrier dissolves with the symptoms.

8. What the lens makes visible and what it does not

Sections 3 through 7 showed the typology in action: canonical layers in §3, the Redis case in §4, the workflow in §5, the extended ladder in §6, and the limit case in §7. Each of those sections assumed conditions — that the persistence model is identifiable, that use-cases are decomposable, that a counterfactual model is constructible. §8 treats what the lens makes visible and what it does not. The typology still applies; the readings become coarser and the translation to action grows more complicated.

The canonical boundary case is observability. Table 4 (§6) shows that the APM stack contains four use-cases with mixed readings: two model-dependent

(internal distributed tracing, one-shot debug) and two structural (cross-system tracing, continuous capacity telemetry). Each reading is correctly assigned per use-case. But a single OTel deployment in a production system serves all four use-cases simultaneously without surgical separation. Removing the instrumentation that serves the internal-tracing use-case affects the instrumentation that serves cross-system tracing; the two share span propagation, collector pipelines, and storage backends. The typology reads the use-case, not the deployment. Operational reality does not admit the same separation, and the practitioner retains the deployment intact and carries the model-dependent use-cases with it.

Four boundaries describe what the lens makes visible and what it does not.

Quantification. The workflow returns a reading per use-case; it does not predict cost saved, performance gained, or operational complexity reduced. Quantification is the work of instrumentation, not of a diagnostic typology — distinct from a limitation, this is a property of the genre.

Stack prescription. A stack populated by model-dependent layers can be identified as suboptimal, but the typology does not prescribe an alternative. The instantiation of §4 demonstrates that one exists; it does not specify the unique alternative for any given system.

Adoption prediction. A model-dependent reading does not inform the difficulty of migration: organizational costs, contractual constraints, vendor lock-in, and engineering skills carry independent weight.

Opaque models. The workflow requires identifying the defect of the underlying persistence model in Step 1 of §5. Systems whose persistence model is opaque — third-party SaaS exposing only an API, proprietary platforms without published internals — fall outside its applicability and are recorded as *unclassified* under the workflow’s third reading.

Section 5’s closing observation extends here: the visibility that the lens renders is independent of the action the practitioner takes. Model-dependent layers may remain in deployments for legitimate reasons — legacy preservation, contractual obligation, organizational constraint, migration cost. The typology documents the debt; it does not resolve it.

The lens reads cleanly where the persistence model is explicit and characterizable, where use-cases are decomposable, and where a counterfactual model is constructible. Where one of these conditions fails, the typology still applies but its readings grow coarser, or the workflow falls back to *unclassified*. Where it reads cleanly, what the lens surfaces is not a smaller stack but a domain that the model permits to be modeled, librated, and evolved without being interleaved with compensating concerns. The utility of the typology is proportional to the decomposability of the system under examination.

Persistence is one axis along which representational pressure deforms a domain — the axis this paper has traced. There is at least one more. When interaction

is modeled before the domain — when aggregates take the shape of the screen rather than the screen rendering the actor’s output — the same compensation pattern reappears in a different register: UI components, view-models, request DTOs, and step-state objects accumulate around a domain that no longer fits its own substrate.

The lens of accidental category reads there with the same discrimination it produced for infrastructural layers and for the server-role. The actor’s existing output boundary — the primitives by which it emits to its invoker, agnostic of who consumes — already factors transport out of the domain. This axis is not developed here.

9. Extension: from layers to roles

The typology of §3 reads among *layers* of infrastructure — components that occupy positions in a stack and compensate for the persistence model from those positions. Architectural *roles* — the entity that accepts authoritative writes, the master in replication, the bootstrap issuer, the orchestrator — admit the same reading under the same two indicators. A role is an accidental category when its necessity is traceable to a specific representational choice (what nodes replicate to share state) and the role loses justification when that choice is replaced.

The same lens therefore applies beyond layers. Under a substrate in which nodes replicate programs rather than data, state, or operations, the server-role does not survive as a structural requirement. Its dissolution is not an architectural objective; it is a structural consequence of the same representational choice that dissolves the layers described in §6.

The two analyses operate at different levels of abstraction — first on layers, then on the roles that those layers exist to serve — without modifying the typology itself. The reading scales upward unchanged.

Code provenance

Source-code references in this paper resolve against the public Puppeteer repository at commit **2f31f96** (2026-05-18). The snapshot is archived in Software Heritage under the following persistent identifier:

```
swh:1:dir:10e7e6bad7eb77b6c2e406762026177f95c3ae92;  
  origin=https://github.com/alvaronCubo/puppeteer;  
  anchor=swh:1:rev:2f31f9674a5de816bdf1bf9d8360ff218a02e4da
```

Inline references of the form `file.cs:NN` (e.g., `ActorHandler.cs:38`) resolve against this snapshot. A reader can construct a per-file SWHID by adding the qualifiers `;path=<path>;lines=<NN>` to the directory SWHID above. Future commits to the repository may renumber lines; the SWHID preserves the cited state independently of any future change to the repository or its hosting.

Acknowledgments

The author used large language models (including Claude and ChatGPT) as editorial assistants for language refinement, structural feedback, and literature navigation. All original ideas, terminology, theoretical constructs, and technical content presented in this work are solely the author's.

References

- Fowler, M. (2005). Event sourcing. Retrieved from <https://martinfowler.com/eaDev/EventSourcing.html>
- Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.
- Helland, P. (2007). Life beyond distributed transactions: An apostate's opinion. *Proceedings of the 3rd Biennial Conference on Innovative Data Systems Research (CIDR)*, 132–141.
- Helland, P. (2016). Immutability changes everything. *Communications of the ACM*, 59(1), 64–70.
- Hewitt, C., Bishop, P., & Steiger, R. (1973). A universal modular ACTOR formalism for artificial intelligence. *Proceedings of the 3rd International Joint Conference on Artificial Intelligence (IJCAI)*, 235–245.
- Kleppmann, M. (2017). *Designing data-intensive applications: The big ideas behind reliable, scalable, and maintainable systems*. O'Reilly Media.
- Rivera, A. (2026a). Anti-porous architecture: a unified design principle for CQRS + Actor + Event-Sourcing systems. *Puppeteer Papers Series*, Paper 1. Zenodo. <https://doi.org/10.5281/zenodo.20404863>
- Rivera, A. (2026b). Program–value separability: the structural precondition for compilation, caching, and dense journaling in a DSL runtime. *Puppeteer Papers Series*, Paper 2. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/02-program-value-separability.md>
- Rivera, A. (2026c). Reactions and the partition: opt-in eventual consistency in actor-native systems. *Puppeteer Papers Series*, Paper 3. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/03-reactions-and-partition.md>
- Rivera, A. (2026d). Preserving semantic continuity across actors: a tell-based approach without orchestration. *Puppeteer Papers Series*, Paper 4. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/04-cross-actor-continuity.md>

- Rivera, A. (2026e). The journal as substrate: unifying deployment, replication, backup, and offline operation in distributed systems. *Puppeteer Papers Series*, Paper 5. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/05-substrate-operations.md>
- Rivera, A. (2026g). The server is not a structural requirement: identifying accidental architectural roles under journaled programs. *Puppeteer Papers Series*, Paper 7. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/07-server-as-accidental-category.md>
- Stonebraker, M., & Çetintemel, U. (2005). “One size fits all”: An idea whose time has come and gone. *Proceedings of the 21st International Conference on Data Engineering (ICDE)*, 2–11.
- Vernon, V. (2015). *Reactive messaging patterns with the actor model: Applications and integration in Scala and Akka*. Addison-Wesley.
- Young, G. (2010). CQRS documents. Retrieved from https://cQRS.files.wordpress.com/2010/11/cQRS_documents.pdf

After the substrate: building software without a datacenter

TL;DR

Building software for a small clinic, a hardware-store inventory system, or a regional logistics back office requires the operational expertise of running a datacenter — not because the domain demands it, but because the conventional cloud-microservice architecture does. The barrier is not the problem being solved; it is the stack that the problem is conventionally solved on.

This paper observes that under a journaled-program substrate — developed piece by piece across the prior six papers of this series — the datacenter ceases to be a structural requirement of running production software. Cloud and microservice ecosystems, when accessed from a journaled program, become libraries the program invokes rather than habitats the program must inhabit. Software can be shipped in shapes the conventional stack does not permit: a single on-premises appliance the customer owns, a peer mesh of devices that cooperate without a central server, a hybrid that invokes cloud services as libraries. The paper establishes the structural permission for these deployment shapes; their end-to-end operational demonstrations are engineering work the paper does not undertake.

A working instantiation makes the observation operationally inspectable: a port of the Ordering bounded context of Microsoft’s `dotnet/eShop` reference application to a journaled-program substrate, replicated across three nodes that bootstrap each other through paper QR codes. After bootstrap, the cluster operates against its own journals; none of the services that originally constituted the Ordering microservice are reachable from it; the cluster continues to function under each one’s absence.

A side observation, stated honestly: the role that prior literature names *server* becomes ephemeral under this substrate — discharged once per peer during the

analog bootstrap, then absent from the running cluster. The construct *accidental category*, applied here at the role layer, is offered as a diagnostic lens the reader may apply to their own architecture; it is not the contribution. The contribution is the practical observation that journaled-program substrates make datacenter-less software construction operationally inspectable, in one working case.

Thesis

Under a journaled-program substrate, the datacenter ceases to be a structural requirement of running production software. The components that the conventional cloud-microservice stack treats as habitats — database, application server, orchestration cluster, observability platform — become libraries the program invokes from its own journaled context, or fall away entirely. Software can be deployed in shapes the conventional stack does not permit: a single on-premises appliance the customer owns, a peer mesh of devices (workstations, phones) that cooperate without a central server, a hybrid that invokes cloud services as libraries. *Structurally compatible* and *operationally shipped in production* are distinct: the paper establishes the former through one working case; the latter is engineering work the paper does not undertake.

The role that prior literature names *server* — the entity that accepts authoritative writes, owns coordination, mediates between clients, and provides the address against which other parties direct their traffic — becomes ephemeral under this substrate. It is discharged once per peer during a bootstrap that can be carried by analog means and decided by humans, then absent from the running topology. This side observation connects the paper to Paper 6’s lens: the server-role reads as an *accidental category* — the role-level analog of *infrastructural symptom* — once we change what nodes replicate.

Evidence. Three operating nodes derived from a port of the Ordering bounded context of Microsoft’s `dotnet/eShop` reference application to a journaled-program substrate, bootstrapped by paper QR codes and authorized by human decisions. After bootstrap, the cluster operates against its own journals; none of the services that originally constituted the Ordering microservice are reachable from it; the cluster continues to function under each one’s absence. The original deployment topology required a datacenter; the ported topology runs on three modest hosts and one paper QR code per peer.

Claims this paper makes

1. Under a journaled-program substrate, cloud and microservice ecosystems become libraries the program invokes rather than habitats the program inhabits; this relation is achieved through the substrate’s first-class deferred-work mechanism (Reactions, Tells) plus the actor-lock’s structural cost on

raw I/O in the verb body, rather than through application-level resilience discipline (§2).

2. The substrate’s three properties — locality, journal as substrate of the program, and self-contained runtime — combined with the actor’s write-lock semantics, make the library relation the structurally cheaper path for deferred external work rather than a per-call developer discipline (§4).
3. A working instantiation, constructed by porting the Ordering bounded context of `dotnet/eShop` to a journaled-program substrate and deploying it across three nodes that bootstrap each other through paper QR codes, makes the observation operationally inspectable: after bootstrap, the cluster operates against its own journals, none of the services that originally constituted the Ordering microservice are reachable from it, and the cluster continues to function under each one’s absence (§5).
4. The bootstrap of a new node — the operation that prior literature treats as the residual case requiring a persistent server — admits a presentation in which the information transferred crosses by analog means (F2) and the authorization is exercised by humans (F4), and the software portion that remains (F1, F3, F5) lasts only as long as the handshake, after which the issuer process exits (§5.3).
5. The substrate of §4 permits whole categories of small-and-medium-business software — accounting, clinic management, retail inventory, professional services — to run on hardware sized to the customer’s domain, rather than on infrastructure scaled to the operator’s stack. This is a structural permission, not an operational demonstration; the present paper does not exhibit an end-to-end customer-delivered appliance (§6).
6. The role that prior literature names *server* becomes ephemeral under the substrate — discharged once per peer during bootstrap, absent from the running topology thereafter — as the structural counterpart of claims 1–5. The construct *accidental category*, applied at the role layer, extends Paper 6’s *infrastructural symptom* to roles; the extension is a diagnostic lens, offered as a side benefit rather than as the headline (§7).
7. The contribution is the practical observation that journaled-program substrates make datacenter-less software construction operationally inspectable in one working case; the accidental-category construct is the analytic tool through which the observation may be applied to other architectures (§7, §9).

1. Introduction

This paper extends the *infrastructural symptom* construct of Paper 6 from infrastructural layers to architectural roles, an extension that Paper 6 §9 anticipates. The genre is closer to *theory for analyzing* in the sense of Gregor (2006) than to design theory: the construct itself is Paper 6’s, applied at a different level of abstraction. The novel material of this paper is the role-level demonstration — a working instantiation in which the server-role dissolves — and the operational

observation that, under the same conditions, cloud ecosystems pass from architectural dependencies to libraries invoked by the program. The contribution is conceptual at the extension level and operational at the demonstration level.

Software for a small clinic, a hardware-store inventory system, a regional logistics back office: these are familiar domains. They are not technically deep — the data models are bounded, the workflows are tractable, the user counts are modest. Yet building any of them today as production-grade software requires assembling a stack — databases, application servers, orchestration clusters, observability platforms, deploy pipelines — that the domain does not demand. The barrier is not the domain. It is the operational expertise required to run the conventional stack.

This paper observes what happens to that stack under a journaled-program substrate. Across the prior six papers of this series the substrate has been developed piece by piece: anti-porous representation (Paper 1), program-value separability (Paper 2), the pragmatic partition for deferred work (Paper 3), cross-actor causal continuity (Paper 4), the journal as substrate of the program (Paper 5), and the diagnostic lens for reading which infrastructural layers compensate for the persistence model rather than serve the problem (Paper 6). Each was argued on its own terms. This paper observes their joint operational consequence.

The consequence is that the datacenter ceases to be a structural requirement of running production software. Cloud and microservice ecosystems, when accessed from a journaled program (in the sense established in Paper 2 §1.2: the pair of domain library and journal of invocations), become libraries the program invokes rather than habitats the program must inhabit. Software can be deployed in shapes the conventional stack does not permit: a single on-premises appliance the customer owns; a peer mesh of devices (workstations, phones) that cooperate without a central server; a hybrid that invokes cloud services as libraries. The paper establishes this consequence as structural permission demonstrated through one working case (§5); end-to-end production deployment in any of these shapes is engineering work it does not undertake (§6).

The paper’s argument runs in three movements. §2 names the operational relation: when external services are reached from a journaled program, what was a dependency (whose absence is an outage) becomes a library invocation (whose absence is a journaled failed call). §3 situates the paper among prior work on local-first, peer-to-peer, agent-centric, federated, and actor-model distributed systems. §4 sketches the substrate properties that make the library relation default rather than disciplined. §5 makes the observation operationally inspectable: a port of the Ordering bounded context of Microsoft’s `dotnet/eShop` reference application — a widely-used example of cloud-microservice architecture in the .NET ecosystem — to a journaled-program substrate, deployed across three nodes that bootstrap each other through paper QR codes. The cluster operates against its own journals; none of the services that originally constituted the Ordering microservice are reachable from it; the cluster continues under each one’s absence.

§6 examines what the observation implies for software construction in practice as a *structural permission*: software becomes deployable in shapes the conventional stack does not permit — a single appliance the customer owns, a peer mesh of devices that cooperates without a central server, a hybrid that invokes cloud services as libraries — and the barrier of operational expertise dissolves with the architectural requirement that produced it. §6 also marks the boundary between structural permission and end-to-end operational demonstration; the present paper establishes the former, not the latter.

§7 introduces, briefly, a diagnostic lens that the reader may apply to their own architecture: the *accidental category* — a role whose necessity is traceable to a representational choice and that loses justification when the choice is replaced. The server-role of the conventional cloud-microservice stack reads as accidental under the substrate; what was a permanent fixture of the topology becomes ephemeral, discharged once per peer during analog bootstrap and absent from the running cluster thereafter. The construct extends to the role layer what Paper 6 introduced for the layer of infrastructure; it is a tool the reader may use, not the contribution of the paper.

§8 examines limitations honestly, including what the present demonstration does and does not establish. §9 concludes.

The contribution is the observation about datacenter optionality, made operationally inspectable in one working case. The construct is offered as a side benefit — the lens through which the observation can be applied to other architectures — not as the central claim. The choice of `dotnet/eShop` for the instantiation is rhetorical: it is a widely-used reference application for cloud-microservice patterns in the .NET ecosystem, whose original deployment topology makes the contrast visible. The framework underlying the instantiation was not designed to support this paper’s claim; the property is identified retrospectively, after a chain of unrelated architectural decisions converged on a system in which the conventional infrastructure requirements had quietly stopped being requirements. The work is, in that sense, a forensic reading of a class of artifacts in which the property happens to hold — consistent with the analytic genre Gregor (2006) describes as *theory for analyzing*.

2. Cloud as library, not as habitat

A cluster of journaled-program nodes accesses cloud and microservice ecosystems as libraries the program invokes, not as habitats the program must inhabit. The structural difference is not in what happens when an external service is unreachable — resilience patterns developed for microservices already achieve graceful degradation under that condition — but in the substrate’s relation to external invocation by default. Resilience patterns (Hystrix at Netflix; Polly in .NET; sidecar-based circuit breakers in Istio; the catalog in Nygard, 2018) achieve library-like consumption at the application level: each external call is

wrapped in an opt-in policy, each fallback handler is explicitly written, and the discipline of doing this consistently is a property of the development team, not of the runtime. Under the journaled substrate, library-like consumption is the substrate’s idiomatic path: external invocations made through Reactions (Paper 3) or Tells (Paper 4) are recorded as journal entries, the outcome (including failure) is part of the program’s durable state, and the journal is replayable as part of the program’s history. The substrate makes this path structurally favored rather than purely disciplinary: raw I/O inside the verb body is permitted as domain code but carries a lock-cost penalty as the substrate’s structural disincentive; external invocations made through Reactions and Tells are journaled by mediation through those constructs. The boundaries of this property — the lock-cost penalty on raw I/O in the verb body, the developer convention of placing slow reads in the caller’s code before the verb is invoked, and the deferred-work mediation through Reactions and Tells — are stated explicitly as falsifiability conditions in §8. The journal-recorded failure is not merely observability metadata but program state available to subsequent Reactions and Tells.

The demonstration in §5 makes this concrete. After the port described in §5.1 and the bootstrap of §5.3, none of the services that previously had to be running for the Ordering microservice to exist are reachable from the cluster. The domain code runs entirely inside each node; what previously required an API host, a database, and an event bus is now a library invoked by journaled programs. The cluster operates against its own journals; the cloud topology of the original microservice has no point of contact with the cluster after the bootstrap completes.

The same relation generalizes to external services that the cluster invokes during normal operation. If the identity provider is unreachable at the moment of invocation, the journal records the failed invocation and the cluster continues. If a payment gateway is unreachable, the journal records the attempt and its outcome. The external service’s absence is recorded as data, not experienced as outage. The cluster does not inhabit any of these services; it invokes them at chosen moments through journaled Reactions (Paper 3) and proceeds with their result, including the result of their absence.

This operational relation inverts the formulation that has organized the local-first software literature:

Table 1. The inversion of Kleppmann’s formulation, captured as a six-word pair.

Kleppmann (Local-first)	The present observation
<i>Own your data in spite of the cloud.</i>	<i>Record the cloud; don’t inhabit it.</i>

Local-first software treats the cloud as something the software must survive

without. A journaled cluster neither survives without the cloud nor depends on it; it invokes cloud services when useful and proceeds without them when not. The two relations differ in what the absence of a cloud service means for the cluster’s continued existence.

The observation shows that cloud services, when used from a journaled program, are consumed as invocable capabilities rather than as the substrate in which the program must reside.

3. Related work

Local-first software (Kleppmann et al., 2019) articulates seven ideals for software that operates primarily on devices owned by their users rather than on remote services. Replication is mediated by Conflict-free Replicated Data Types (CRDTs), which allow concurrent edits to merge without coordination. The position with respect to cloud infrastructure is precise: cloud services are admissible, but the software must continue to function in their absence. The title of the paper — *Local-First Software: You Own Your Data, in spite of the Cloud* — captures the stance: the cloud is neither rejected nor required. The server-role is reduced to one of several optional services, and its elimination is the explicit objective of the architectural pattern.

Peer-to-peer append-only logs constitute a second line. Secure Scuttlebutt (Tarr et al., 2019) replicates per-identity signed logs across a gossip network without a central coordinator; the log, not the data it carries, is the unit of synchronization. Hypercore (Holepunch, n.d.) and the Dat protocol (Ogden et al., 2017) generalize the same pattern, treating the append-only feed as a primitive replicated by a distributed hash table. In each case the unit of replication is a log of data — events, document edits, file fragments — and the application logic that interprets the log lives in the consuming application. The server’s elimination is the design target of the network layer.

Holochain (Harris-Braun et al., 2018) makes the agent the source of truth. Each agent maintains a local source chain — an immutable, signed sequence of its own transactions — and contributes to a validating distributed hash table. The architecture inverts the data-centric assumption of blockchain rather than addressing the cloud directly; the elimination of the server-role is a consequence of the agent-centric inversion. The replicated artifacts are the agent’s chain and the validated DHT; application logic, written as zomes, is interpreted locally by each agent.

Federated systems form a fourth line. Matrix federates an event graph across cooperating homeservers; ActivityPub federates streams of activities; the AT Protocol of Bluesky splits hosting across a Personal Data Server, a Relay, and an AppView, with account portability as the displacement mechanism. The server-role is preserved but pluralized: every participating host fulfills the role for some subset of users. The cloud is partitioned across cooperating providers

rather than centralized or rejected.

Adjacent work admits brief mention. Solid (Sambra et al., 2016) proposes Personal Online Datastores that decouple identity from application providers; the server-role moves from application owner to data host without being eliminated. Earthstar (Earthstar Project, n.d.) pursues local-first synchronization without a distributed hash table, sitting between Kleppmann’s CRDT model and Hypercore’s gossip topology. Neither treats the server-role as a category whose referent is contingent.

Each of these lines treats the server-role as a *design target* — an architectural feature to redesign, redistribute, replace, or pluralize. None of them names the property that the present paper identifies: that the role is contingent on a choice about what nodes replicate, and that under a different choice the role has no referent at all. Table 2 summarizes the four families and the position of the present construct.

Table 2. Five families of prior work and how each treats the server-role. The bottom row identifies the construct of the present paper; the columns describe what each family replicates between nodes, where its application logic resides, why a privileged server is absent (or pluralized), and how each family situates itself with respect to cloud infrastructure.

Family	What is replicated	Where the logic lives	How the server-role is treated	Relation to the cloud
Local-first (Kleppmann)	CRDTs / data	In the application	Eliminated as a design goal	Survival without it
Hypercore / Dat / SSB	Append-only logs	In the application	Eliminated as a network goal	Replaced by peers
Holochain	Source chain + DHT	In the agent	Eliminated by agent-centric inversion	Treated as an alternative architecture
Federated (Matrix / ActivityPub / AT Protocol)	Events / activities	On federated servers	Pluralized across hosts	Fragmented across providers
The present construct	Programs	In the journal of each node	Not a pre-requisite (accidental category)	Subsumed as a library

The discriminator is the second column. What each family replicates between nodes — CRDTs and data deltas, append-only logs of data, agent source-chains, federated events and activities — determines what residual role survives. The present construct replicates neither data, deltas, logs, nor events: it replicates *programs*, in the substrate-level sense of Paper 2 §1.2 (the pair of domain library and journal of invocations). The fourth column is downstream of this choice. Under a representational choice in which what travels between nodes is the program itself, the server-role’s necessity does not survive. The fifth column is downstream of the fourth: under the present construct the cloud is neither evaded, replaced, nor pluralized, but reused as a library within the program.

Beyond these four families that treat the server-role as a design target, a fifth body of work shares direct lineage with the substrate of §4 without making the server-role its design objective. Erlang/OTP (Armstrong, 2003) is the longest-running production actor system; its supervision trees and `gen_server` abstraction shape the ecosystem from which the substrate descends, though OTP nodes carry explicit names and `gen_server` processes carry explicit server-role at the level of the diagnostic lens (§7). Microsoft Orleans (Bernstein et al., 2014) introduces virtual actors with location-abstracted activation: an actor’s host is determined by silo-cluster membership, not by application code. Orleans dissolves the location-of-the-actor but retains the silo cluster’s coordination layer — the server-role is internalized in the membership protocol rather than externalized to the application, but it is not dissolved at the lens’s reading. Akka Persistence (Lightbend, n.d.) and EventStoreDB (Event Store Ltd., n.d.) are the most direct comparators on the actor-and-event-sourcing axis: both persist and replicate the *events* an actor emits, with the actor’s behavior recovered by replaying those events through code that lives outside the journal — the unit of replication is the event stream, and the program that interprets the stream is deployed alongside the runtime as application code. The present substrate inverts this relation: the unit of replication is the program itself, and what was previously called the event is recovered as the trace of the program’s execution. Under the lens of §7 this difference is not a feature comparison but a representational one: events are what the program produces, not what the program *is*. Amazon Dynamo (DeCandia et al., 2007) operates as eventually-consistent without master, with conflict resolution per-key; it is the canonical AP-without-master system within the CAP positioning (Brewer, 2000; Gilbert & Lynch, 2002). Datomic (Cognitect, n.d.), a commercial immutable database, treats its journal as the source of truth but routes all writes through a single transactor — under the lens of §7, the transactor itself is an accidental category that the present substrate dissolves.

Two theoretical results bracket the structural claim. The CALM theorem (Hellerstein, 2010; Ameloot et al., 2013) establishes a formal correspondence between monotonic computation and coordination-freeness — a program admits coordination-free distributed execution iff its computation is monotonic. The substrate of §4 does not appeal to monotonicity directly, but the dissolution it documents at the role layer is consistent with CALM’s prediction: when state

propagation is by replay of a deterministic program over a journal rather than by reconciliation of competing writes, the conditions under which coordination becomes structurally necessary are restricted. Conflict-free Replicated Data Types (Shapiro et al., 2011), the mechanism underlying Kleppmann’s local-first model, achieve coordination-free convergence at the data layer; the present substrate achieves the same property at the program layer, replicating script invocations rather than data deltas.

The substrate of §4 combines the actor model (Hewitt, 1973), event sourcing (Fowler, 2005), and CQRS (Young, 2010), synthesized as practical patterns by Vernon (2015). An additional move — domain classes carrying no persistence concerns, discovered by reflection — antedates the synthesis by roughly a decade in the framework underlying the instantiation; that genealogy is described in Paper 6 §4.0.

This paper is the seventh in a series. Papers 1–6 develop the substrate-level foundations on which the present observation rests, culminating in Paper 6’s *infrastructural symptom*, of which the diagnostic lens of §7 applies the same two indicators to architectural roles — a move Paper 6 §9 anticipates explicitly.

The contribution unique to this paper is the observation that under a journaled-program substrate, cloud and microservice ecosystems pass from architectural dependencies to library invocations (§2), with the practical consequence of making the datacenter optional for production software construction (§6). The instantiation of §5 makes the observation operationally inspectable in one working case. The diagnostic lens of §7 — *accidental category*, applied at the role layer as an extension of Paper 6’s *infrastructural symptom* — is offered as a side benefit: the analytic tool through which the observation can be applied to architectures beyond the one demonstrated here.

4. The substrate that makes it possible

The substrate underlying this paper’s observation is the journaled actor-native framework developed across the prior six papers of this series. Its architectural lineage is documented in Paper 6 §4.0 and is not retraced here. The substrate is described in the abstract throughout this section; its concrete instantiation, which the demonstration in §5 builds upon, is the framework whose source is publicly available (see Appendix A).

The term *program* throughout this section is used in the substrate-level sense of Paper 2 §1.2: the pair (domain library, journal of invocations) that nodes replicate to share state. The substrate’s representational choice — replicating programs in this sense — is what makes it distinct from prior actor-and-event-sourcing systems. CRDT systems replicate state deltas; pub-sub and message-queue systems replicate messages; event-sourced systems (Akka Persistence, EventStoreDB, conventional CQRS+ES stacks) replicate events. The substrate of this paper replicates programs. The three properties of the sub-

strate that follow are entailments of this choice, not three independent design decisions; together they make the cloud-as-library relation of §2 default rather than disciplined.

First, each node owns a local journal of invocations — DSL scripts (in the unit sense of Paper 2) that, on replay, reconstruct the node’s state deterministically. The journal is the cumulative state-producing artifact; the state is its replay, not a captured snapshot. No external store mediates between the node and its state.

Second, the domain library — the classes and verbs that the journal scripts invoke — is loaded by reflection at startup; no node hosts the domain as a service for others. Every node carries the same library; the substrate is symmetric across the cluster.

Third, replication is journal-to-journal: nodes exchange script references and parameters, not state diffs, and reach convergence by replaying the same sequence of scripts against the same library. The substrate’s primitive is *replay* — the single primitive to which Paper 5 reduces deployment, replication, backup, and offline operation.

None of these properties requires a node to fulfill a privileged role. Each is intrinsic to the substrate, not externalized. When the cluster of §5 reaches external services, it does so against this substrate’s storage: the journal records the invocation and its outcome as part of the program’s history, and the cluster continues by replaying the journal. That is what makes cloud-as-library the default — the library relation is structural, not a property of how the developer wrote the call site.

The implementation source is publicly available at the commit pinned in Code Provenance; reproduction details are consolidated in Appendix A.

5. The demonstration

The observation of §2 — that under a journaled-program substrate the cloud becomes a library and the datacenter becomes optional — is made operationally inspectable by the demonstration in this section: a port of the Ordering bounded context of Microsoft’s **dotnet/eShop** reference application to the substrate of §4, deployed across three nodes that bootstrap each other through paper QR codes.

5.1 The port

The instantiation ports the *Ordering* bounded context of Microsoft’s **dotnet/eShop** reference application — a widely-used example of microservice-oriented cloud architecture in the .NET ecosystem — onto the substrate of §4.1. The diagnostic value of the port is what it makes visible: the lines of the original

codebase that the substrate cannot host without modification are a small minority, concentrated in persistence-and-deployment glue. The aggregate, the value objects, the lifecycle verbs, and the invariants move into the journaled substrate without modification. What was previously named “the Ordering microservice” turns out to be a clean domain library wrapped in infrastructure compensating for a different substrate; the port separates the two.

5.2 The three-node deployment

Three identical processes — each running the same substrate with the ported domain library loaded by reflection — deploy across three Docker containers. The coordination mesh is fully connected: three pairwise channels forming the complete $N(N-1)/2$ graph at $N = 3$. Three is the minimum peer count at which the dissolution claim is unambiguous: with two peers, any rotation can be read as one serving the other; with three, no subset of two dominates the third.




Figure 1. The three-node coordination mesh. Each node is a peer; none is privileged.

No node carries the write-acceptance privilege permanently; the role rotates symmetrically across peers. After the rotation completes, the three journals are byte-coincident.

Director rotation is the specific peer-coordination mechanism chosen for the demonstration. It is part of the StageManager pattern — a peer-to-peer coordination layer built on top of the actor substrate, located in `Choreography/StageManager/Stage.cs:470` — and is not a property of the actor model itself. The rotation is not consensus-based in the Raft (Ongaro & Ousterhout, 2014) or Paxos sense; the substrate does not elect Directors internally, run voting, maintain term or epoch numbers, or provide fencing tokens. Promotion is an external decision; the new Director’s identity propagates via best-effort announce. The substrate’s contribution is journal-determinism between peers (Paper 5), composable with various coordination patterns and transports; StageManager with Director rotation, over Kestrel/HTTPS, is the combination chosen for §5’s Docker-container demonstration. What it does not remove, and does not claim to remove, is the broader consensus problem — a different question with different machinery.

5.3 The analog bootstrap

The bootstrap of a new peer proceeds in five phases. **F1** opens an invitation on the issuer side. **F2** transmits the invitation to the joining peer by any means, including analog ones — printed paper, a verbal dictation, an email. **F3** opens an authenticated connection back to the issuer. **F4** pauses for a human decision to approve or reject. **F5** seals a credential to the joining peer’s public key and closes.

Two of the five phases are sub-software: F2 (the information transfer) and F4 (the authorization decision). Neither requires a service to be running anywhere; the carrier of F2 is whatever the operator chooses, and the decision-maker of F4 is human. The other three phases (F1, F3, F5) are themselves server-role-like discharge: the issuer listens on an address (F1), accepts a connection (F3), and seals a credential (F5). The discharge is genuine, not concealed. What §5 documents is that this discharge is ephemeral, time-bounded, and terminating: the three software-mediated phases are local to the issuer process, which exits when the bootstrap is complete. The transport stack that carries F3 (HTTPS/TLS via Kestrel; see Appendix A.4) is a property of the underlying transport, not of the journaled substrate — the substrate’s contract is that whatever the transport delivers, the journal records and replays. **The issuer process does not act as a server, coordinator, authority, registry, or rendezvous point after bootstrap; its only role is to transfer a credential during the brief lifetime of the handshake, after which it ceases to exist.**

After all three containers complete the handshake — each with its own QR code carried by analog means and approved by a human — the issuer process exits. The containers continue to operate.

```
ordering-a  Up 11 seconds    0.0.0.0:6443->6443/tcp
ordering-b  Up 11 seconds    0.0.0.0:6444->6443/tcp
ordering-c  Up 11 seconds    0.0.0.0:6445->6443/tcp
```

Output of `docker compose ps` after the issuer exits. No surviving network path connects the containers to the issuer; the bootstrap papers are in a drawer.

The information that constituted the bootstrap traveled by paper. The decision that authorized it was made by a human. The cryptographic protection of the credential happened during the brief lifetime of the issuer process. The artifact that survives is the three running containers.

5.4 The empirical matrix

The demo runs across a two-by-two matrix: *in-process* versus *cross-container*, and *inline* versus *parametric*. Each cell is a complete execution of the workload across the three nodes; the in-process row reduces the bootstrap to direct invocation, while the cross-container row runs the full §5.3 protocol.

Table 3. Empirical matrix of the existence proof. Each cell reports the final

journal entry count, identical on all three peers and byte-coincident on disk.

Workload regime	In-process (3 Stages, single OS process)	Cross-container (3 Dockers, real TLS + analog bootstrap)
Inline (literal scripts)	22 entries	22 entries
Parametric (Define + bind)	7 entries	7 entries

The four cells produce identical final entry counts per workload regime, and the three peers in each cell are byte-coincident on disk. The cross-container row confirms that the operational properties of the substrate survive the addition of the analog bootstrap and the real cryptographic stack; the in-process row is the structural rehearsal that isolates the substrate’s intrinsic properties from the network. The compaction from inline to parametric is real and measured; its precise ratio and the conditions under which it varies are treated in Paper 2.

Byte-coincidence is the structural witness of the headline observation: it shows that no peer is privileged — the cluster operates symmetrically, with no node fulfilling the role of authoritative writer or central coordinator. Operational metrics — write throughput, append latency, cross-DC convergence — live elsewhere in the series: Paper 5 lab L1 reports 451k events/sec on the substrate, lab L6 reports 347µs p50 append on the FS backend, and lab L4 demonstrates bit-exact cross-DC convergence under sustained workload. The present experiment is not a benchmark; it is the existence proof for the observation of §2, alongside the operational floor that Paper 5 establishes.

In the parametric regime, the journal records seven entries rather than five: each rotation Director emits its own Define + Invocation pair, because the Define cache is local to the Stage that emits it, not replicated. The journal is the catalog of each peer’s declarative vocabulary; the cache’s compactness lives on the Stage axis, not on the journal axis.

5.5 Partition tolerance and re-join

The same artifact additionally exhibits partition tolerance and re-join from disk. A peer can exit the cluster while the others continue operating; a fresh process can mount the absent peer’s data directory and rehydrate its journal-relative state from disk alone, without contacting the issuer or the surviving peers; and the rehydrated peer integrates into the still-running cluster via peer-to-peer catch-up.

Five phases describe the sequence:

- **R1.** The three peers converge to a shared journal state.

- **R2.** One peer exits; the others continue. The cluster retains a Director and accepts new work.
- **R3.** The surviving Director issues additional work; the two surviving journals advance; the absent peer's on-disk journal is unchanged.
- **R4.** A fresh process — distinct in-memory identity, same data directory — reads the disk and reconstructs the journal-relative state of the absent peer, without contacting the issuer and without contacting the surviving peers.
- **R5.** The rehydrated peer integrates into the still-running cluster via peer-to-peer catch-up; the gap is replayed from a surviving peer; subsequent work brings all three peers back to a shared state.

The properties exercised by this scenario are not separable. The cluster's continued operation in the absent peer's absence is only possible if no node is the cluster's source of authority. The rehydration without consulting the issuer is only possible if the disk is a self-contained witness of the peer's prior participation. The catch-up between peers is only possible if any surviving peer is capable of replaying the gap; no specialized recovery node is invoked.

Operationally, the scenario shows that a peer can leave the cluster, the cluster can continue accepting work in its absence, and the peer can re-join by reading the journal that survived its disappearance on its own disk — no central authority is consulted at any point in this loop. The role of the issuer was discharged once, at the original onboarding, and is not invoked again.

The issuer is invoked exactly once per peer. The credential sealed during the bootstrap of §5.3 is preserved across restarts, and the peer's continued participation is verified by the journal's cryptographic integrity, not by reference to a running service. If a peer loses its disk, it re-onboards as a new peer; if it preserves its disk, it returns as itself, regardless of how long it was absent.

6. What it means in practice (a structural permission)

The observation of §2 — cloud as library, not as habitat — and the demonstration of §5 together state that under a journaled-program substrate, the datacenter ceases to be a structural requirement of running production software. This section examines what that implies for software construction in practice, and what it does not.

The existing landscape. Software delivery without a datacenter is not a new category. Appliance-delivered business software (Odoo Community, Sage 50, QuickBooks Desktop, SAP Business One, Splunk Enterprise) and self-hosted applications (Nextcloud, Plex, Home Assistant, Synology DSM) deliver workloads without a datacenter, in some cases for decades. In-memory databases (VoltDB Community, SQLite, DuckDB) run application data outside cloud in-

infrastructure. Edge-computing platforms (Cloudflare Workers, AWS IoT Greengrass) distribute logic outside the datacenter. Embedded business systems — PLCs, point-of-sale terminals, ATMs — have run software on modest hardware for forty years. The substrate of §4 does not invent any of these deployment shapes; it does not compete with the systems above; it does not claim that the categories they occupy do not exist.

What the substrate adds. The reframing is narrower than the categorical “datacenter optional” headline would suggest in isolation. The systems above each deliver software outside the datacenter, but each pays a price for that delivery: bundled database administration (Nextcloud, Sage, Odoo); no native event sourcing in most self-hosted applications; in-memory databases as storage layers without application-level semantics for replay; no journal as a unified substrate for deployment, replication, backup, and offline operation (per Paper 5). The journaled-program substrate brings the joint operational profile of cloud-native event-sourced microservices — event sourcing native (Paper 1), program-value separability for compiled cached invocations (Paper 2), the pragmatic partition between work-due-before-responding and deferred work (Paper 3), cross-actor causal continuity through Tell (Paper 4), the journal as substrate of the four operational disciplines (Paper 5), and the diagnostic lens for reading which infrastructural layers compensate for the persistence model (Paper 6) — to hardware modest enough to fit any of the deployment shapes the systems above already inhabit. Operating an on-premises business system in the conventional stack composes several layers of operational expertise — database administration, application hosting, backup, updates, coordination between components — at the customer site or via the reseller; the journaled substrate collapses these into the substrate itself, so storage, backup, updates, and dispatch become properties of the runtime rather than concerns the operator composes. The substrate is not necessarily faster or cheaper than the existing systems — that is a separate empirical question — but its *operational surface* is structurally smaller. The unique contribution is this bundle on modest hardware; the categories of datacenter-less delivery themselves predate the substrate.

The contribution can be considered illustratively (not as evidence) through a concrete scenario. A developer or small team encounters a customer too small to be addressed economically by existing mid-market or enterprise systems and not adequately served by generic spreadsheets, single-purpose SaaS, or bespoke development. Under the conventional stack, the developer faces a ladder of operational concerns before they reach the domain: *I need a server — at the customer or in a cloud — and either way I need to provision a database, configure an application tier, set up backup, arrange monitoring, and handle updates before I can start writing the domain.* The toy alternative — a spreadsheet, a script, a one-off utility — does not survive contact with the customer’s actual usage. The ladder does not begin at the domain; it begins below it. The journaled-program substrate flattens the ladder: the developer can write the domain as their first concern (a library of classes and verbs in the language they already know), while the runtime substrate handles persistence, backup,

updates, and dispatch. The operational surface that the conventional stack ten-steps to is, under the substrate, one or two — and those one or two do not demand operational specializations the developer does not have. This scenario is offered for illustration; the paper makes no controlled-study claim about which developer populations are or are not currently served by existing systems, nor about which barriers (operational, marketing, support, sales, regulatory) are binding in particular cases.

New deployment shapes. The substrate is compatible with deployment shapes the conventional cloud-microservice stack does not permit. Beyond the single-appliance shape — software shippable as a deliverable artifact the customer owns, sized to the domain — the same substrate enables a peer mesh: devices that cooperate without a central server, where each peer participates as a node in the §5 sense and the cluster operates from the peers’ joint awokeness rather than from a server’s uptime.

The §5 demonstration is precisely this mesh at $N=3$ on Docker containers, composing three layers: the substrate (the actor model with journaled programs re-played against a shared domain library), a peer-coordination pattern (the StageManager pattern with its Director rotation), and a transport (Kestrel/HTTPS). The substrate’s core does not prescribe a particular coordination pattern or transport; StageManager and Kestrel/HTTPS are demo-specific choices appropriate to Docker. The transport-layer pluggability is exercised in the published artifact: alongside the Kestrel/HTTPS used in §5, a second transport implementation based on the SimpleX messaging protocol has been implemented and verified to run on Android, and — when substituted for Kestrel/HTTPS as the peer transport in the §5 three-node Docker scenario (an `smp-server` container providing the SMP relay; the three `ordering-*` containers connecting to it via `ConfigureTransport(TransportType.SimpleX, ...)` instead of HTTPS) — produces the same convergence outcome with bit-equivalent final journal entries to the HTTPS variant (`Choreography/Transport/SimpleX/` in the public repository at commit `b42d0f7`). The layered framing — substrate / coordination pattern / transport — is therefore not hypothetical but exercised across at least two transport implementations. Other deployment shapes compose the substrate with different layer combinations. Workstations on a corporate LAN can reuse the demo’s Kestrel-based transport directly — a clinic with three workstations is the closest direct lift of the demonstration. Mobile peer meshes compose the substrate with the SimpleX-based transport (its Android viability verified at the transport layer) and a coordination pattern appropriate to opportunistic synchronization rather than continuous Kestrel sessions; the operational engineering of a full mobile peer-mesh deployment — battery management, NAT traversal between home networks, OS suspension handling, cellular reconnection patterns — lives at the operational layer downstream of transport choice and is not undertaken by the present paper. The substrate’s contribution is the actor-journal-replay model; transport and coordination are downstream choices the deployment makes.

The contrast with specialized peer-to-peer systems — blockchain consensus, distributed hash tables, mesh-routing protocols — is that the substrate’s mechanism (journalled program, replayed deterministically against a shared domain library) is not specialized to a particular peer-mesh problem; the same substrate composes with whatever transport-and-coordination layer the deployment requires, serving a delivery mesh of riders coordinating inventory, a small clinic where nurses share a patient roster during their shifts, a classroom of students cooperating on a shared exercise, a community game where players’ devices form the world’s persistent state — each with the transport and coordination layers chosen appropriately for that deployment substrate.

What remains. §2 and §5 demonstrate operationally that a journalled cluster operates without a datacenter footprint in one case — the Ordering bounded context of `dotnet/eShop` — at one scale — $N=3$ — under one transport — HTTPS/TLS, on Docker containers. From that demonstration, this section derives by structural argument that appliance delivery, peer-mesh deployment, and accessibility-of-development are architecturally permitted. The paper does not demonstrate an end-to-end appliance delivery to a customer in production, a peer mesh of mobile devices in field use, a developer-to-customer flow with operational expertise lower than the conventional stack requires, or the long-tail viability of journalled software for any of the application categories listed. The structural permission is what the paper establishes; the end-to-end demonstration of each deployment shape is engineering work the paper acknowledges as outside its scope. Beyond the architectural permission, the substrate does not address sales channels (how the customer learns this software exists), distribution (how it reaches the customer), support models (who answers the phone when something fails), regulation (where data resides and whose privacy laws apply), or pricing (license vs subscription vs free). These are downstream of architecture and depend on choices the architect does not control. The shapes the developer might choose (single appliance at the customer site, peer mesh across devices, cluster across machines, hybrid with cloud invocations) are not new categories; they are deployment topologies for which polished products often do not exist or are poorly fitted at the small end of the market, and the obstacle to building those products has multiple components. This paper addresses one of those components — the architectural one. The others remain. The conditions under which empirical evidence would refute the structural permission are stated in §8.

7. A diagnostic lens, briefly: when the server-role is accidental

The substrate’s three properties of §4 and the demonstration of §5 make a side observation possible: the role that prior literature names *server* — the entity that accepts authoritative writes, owns coordination, mediates between clients, provides the address against which other parties direct their traffic — becomes

ephemeral rather than permanent. In the demonstration of §5, no node carries the four functions of the server-role permanently. Acceptance of authoritative writes and ownership of coordination redistribute under Director rotation; mediation and addressing distribute across the peer-to-peer mesh; partition tolerance and re-join distribute across surviving peers via catch-up (§5.5); the bootstrap issuer (§5.3) participates once per peer and exits. **The server is not a permanent identity. It is a transferable coordinator role any peer can assume in turn — and any peer can leave the cluster and re-join without reference to a coordinator outside it.**

At the level of construct, this is Paper 6’s *infrastructural symptom* lens applied at the role layer: the server-role reads as an *accidental category* — the role-level expression of infrastructural symptom — when its necessity is traceable to a specific representational choice (Paper 6’s *contingency* indicator at the role layer) and disappears when that choice is replaced (Paper 6’s *dissolution* indicator). The reading operates per role-in-use-case; many roles remain structural under any substrate — identity providers backed by cryptographic root authority, gateways mediating external systems, computationally expensive services — and no substrate change makes them otherwise. The full apparatus of the lens — formal indicators, the constructibility bound on counterfactual substrates, per-use-case discrimination — is in Paper 6 §3 and is not retraced here. Appendix B catalogs seven canonical architectural roles under the two indicators, as the role-level analog of Paper 6’s Table 1, for readers who want to apply the lens to systems beyond the one demonstrated in §5.

The lens is offered as the analytic tool through which the consequence of §2 and §6 can be applied to architectures beyond the demonstration of §5 — not as the contribution of the paper.

8. Limitations and counter-arguments

The demonstration of §5 carries specific limitations that follow from the choices made for the demo. They are reported here honestly; none of them affects the observation of §2 or its practical consequence in §6, but each affects what the present experiment alone establishes.

Identity in the port is local. The original `dotnet/eShop` deployment uses an external identity provider; the port substitutes a local string-typed identifier. The lens of §7 does not read identity providers as dissolving — identity providers backed by cryptographic root authority are structural requirements, as Appendix B notes — and the demo does not exhibit one. A second experiment that integrates an external identity provider as a Reaction-invoked library would extend the empirical surface without changing the lens’s reading.

Cross-context dependencies are not portrayed. The Ordering bounded context of `dotnet/eShop` carries references resolved by the Catalog and Basket bounded contexts; the port treats those references as opaque identifiers. The

same port mechanism applies to Catalog and Basket; the present experiment exhibits Ordering only.

The two limitations above warrant clarification, because they are easy to read as selection bias when they are in fact bounded-context isolation in the Domain-Driven Design sense (Evans, 2003; Vernon, 2013). The Ordering bounded context, in any architecture — microservices, monolith, or journaled-program substrate — references Catalog and Basket entities by identifier rather than by full object, because cross-context object resolution is itself an integration boundary, not a property of Ordering’s domain. Identity, similarly, lives in its own bounded context (the Identity microservice in the original eShop); Ordering does not authenticate users, it accepts a `UserId`. Removing those from the port is not eliminating components that “need a server”; it is respecting the same boundaries DDD prescribes. Appendix B itself lists identity providers backed by cryptographic root authority as structural requirements — the lens never reads identity as dissolving, so its exclusion is not adverse evidence.

The empirical matrix is at $N = 3$. The mesh in §5.2 connects three peers — the minimum non-trivial fully-connected mesh — using a specific peer-coordination mechanism (Director rotation) chosen for the demonstration. The substrate’s structural claim is independent of node count and of the specific coordination mechanism: it is that journaled programs replicate without a privileged coordinator at any scale at which the journal-determinism property holds between peers. Whether the specific Director rotation mechanism demonstrated at $N=3$ scales operationally to larger N is engineering work the paper does not undertake; other coordination mechanisms (gossip, hierarchical overlays, sharding) are well-understood in distributed systems and could be used in conjunction with the substrate’s journal-determinism at larger scales, but the present paper does not demonstrate any such combination. The substrate’s structural permission holds at any N where journal-determinism between peers is preserved; the operational scaling characteristics at large N are not demonstrated and are explicitly out of scope.

The cross-container partition-tolerance scenario is captured in a parallel artefact, not in the matrix of §5.4. The natural cross-container equivalent of the in-process scenario in §5.5 — stopping a container, allowing the running cluster to continue accepting work in its absence, restarting the container, and watching it rehydrate from its mounted volume and rejoin via catch-up against a surviving peer — is implemented in the same orchestrator that produced the matrix (the `--rehydrate-demo` flag of the demo script) and depends on no substrate feature that is not already exercised by the in-process test in §5.5. The decision to keep §5.4 to the four cells of the matrix, rather than expanding to a fifth cell, was rhetorical: the in-process exercise establishes the property structurally. The cross-container capture documents the same property under a different process boundary; it does not add a new claim.

The compaction ratio between inline and parametric workloads depends on Director rotation. Each new Director journals a fresh Define +

Invocation pair, so the Define cache compacts on the Stage axis rather than on the journal axis. The compaction is real and measured at three rounds; its precise value at larger workloads is treated in Paper 2.

What would falsify the structural permission. The structural permission framing of §6 absorbs many objections by deferring to engineering work downstream of architecture — operational difficulties at scale, transport-specific failure modes, deployment-environment quirks. None of these would refute the structural permission; they would calibrate the cost of exercising it. The structural permission is refuted only by failure modes that show the substrate cannot host the deployment shape even in principle. This subsection names four such conditions.

- **A runtime central component is required for intra-cluster operation at $N > 3$.** If extending the §5 demonstration to $N=4$ or higher required introducing a process that must remain alive — a coordinator, a directory server, a bootstrap re-issuer — beyond the peers themselves, the cluster’s peer-symmetric character would not generalize from $N=3$, and the appliance and peer-mesh shapes would not be structurally compatible with cooperation beyond the demonstration’s scale.
- **Journal-determinism breaks under transport change.** The substrate’s claim of transport pluggability requires that the journal records (and replays) the same operations regardless of which transport carried them. If swapping Kestrel/HTTPS for an alternative transport — SimpleX, analog transfer, or any other — produced journals with different content for the same operations, or state not recoverable from replay, the transport-pluggability claim of §6 would be refuted.
- **Raw I/O in verb bodies systematically breaks journal replay.** §2 describes cloud-as-library as the substrate’s idiomatic path, with raw I/O inside the verb body permitted as domain code under a known lock-cost penalty as the substrate’s structural disincentive. If such I/O could systematically produce journal entries that the substrate could not replay deterministically — beyond the lock-cost penalty and the convention boundary of using Reactions or Tells for external invocation — the substrate’s contribution to cloud-as-library would reduce from substrate-supported to purely convention-mediated, and Paper 2’s “journal is the program” claim would be similarly narrowed.
- **Peer-mesh requires global identity for cluster-to-cluster operation.** If two journaled clusters established independently — each through analog bootstrap (§5.3) — could not be reconciled without consulting a shared identity authority, the peer-mesh’s independence from central authority would not extend to inter-cluster operation, and the deployment-shape claims of §6 would narrow.

Each condition is testable. The paper does not undertake these tests; it states the structural permission and the conditions under which empirical evidence

would refute it.

The lens of §7 extends to other architectural roles that admit the same analysis. Service discovery, queue-as-glue between bounded contexts, configuration-management roles, and distributed-tracing collectors are auditable under the two indicators; the present experiment does not exercise them, and the substrate’s relation to each is left as a forward question. The lens admits this open extension by design — its diagnostic value is in reading which categories in any given system are contingent, not in enumerating them exhaustively.

9. Conclusion

This paper observed what happens to the conventional cloud-microservice stack under a journaled-program substrate developed across the prior six papers of this series: the datacenter ceases to be a structural requirement of running production software. Cloud and microservice ecosystems, when accessed from a journaled program, become libraries the program invokes rather than habitats the program must inhabit. Software becomes shippable as a physical artifact — an appliance the customer owns, sized to the customer’s domain, not a datacenter the vendor leases on the customer’s behalf.

The observation was made operationally inspectable by porting the Ordering bounded context of `dotnet/eShop` to the substrate of §4, deploying it across three nodes that bootstrap each other through paper QR codes, and observing that after bootstrap the cluster operates against its own journals — none of the services that originally constituted the Ordering microservice are reachable from it, and the cluster continues under each one’s absence. The demonstration does not argue uniqueness of the substrate; it establishes that the observation is inspectable in one working case.

The practical consequence (§6) is, as structural permission, an opening in the space of deployment shapes the conventional stack does not permit. Software becomes structurally compatible with delivery as an on-premises appliance, with cooperation across a peer mesh of devices (workstations, phones) without a central server, and with hybrid topologies that invoke cloud services as libraries. The shapes are not new categories — existing systems (Odoo Community, Sage 50, QuickBooks Desktop, Splunk Enterprise, Nextcloud) already deliver business software without a datacenter — but the journaled substrate offers a different operational profile that does not require the operator to compose specialized expertise (database administration, application hosting, backup, updates, orchestration) on top of the architecture. The paper establishes the structural permission; whether the permission becomes operational practice in any of these shapes is engineering work the paper does not undertake.

A side observation, stated honestly, is that the role prior literature names *server* becomes ephemeral under this substrate: discharged once per peer during analog bootstrap, then absent from the running topology. The construct *accidental*

category (§7), applied at the role layer as an extension of Paper 6’s *infrastructural symptom*, is offered as a diagnostic lens the reader may apply to their own architecture. The construct is the tool; the practical observation is the contribution.

The lens does not promise universal dissolution. Identity authorities, payment networks, sensor integrations, and computationally intensive services remain structural under any substrate. The reading is per role-in-use-case, not per system. What the lens offers is examinability: any production architecture can be examined role-by-role, asking whether each role is tied to the problem being solved or to the persistence model chosen.

Paper 6 applied the lens at the layer level, identifying many infrastructural layers — caches, queues, locks, application servers, orchestration clusters — as compensations for the underlying persistence model. The present paper observes the joint operational consequence: with those layers absent, and the architectural roles they serve revealed as accidental, the datacenter itself becomes optional, and the construction of production software passes back into the hands of those who model the domain. From this point on, treating the datacenter as a structural requirement is no longer an assumption but a choice that can be made explicit and evaluated.

Appendix A — Reproducibility

The artifact described in §5 is publicly available. The references in this appendix to file paths, test method names, and orchestrator scripts resolve against the released code; the artifact comprises the substrate framework, the runtime, the ported **Ordering** domain library and its test suite, the demo orchestrator, the supporting reports, and the asciinema recordings, hosted across two repositories described in §A.1.

A.1 Source and suite

The artifact is hosted in two public repositories under <https://github.com/alvaroNCubo>. The substrate framework (**Choreography**/), the runtime (**Puppeteer**/), the ported **Ordering** domain library and its test suite (**tests-local/UnitTestPaper7EShop**/), the demo orchestrator (**docker**/), and the supporting reports referenced in §A.7 (**notes/paper7_phase1_results.md**, **notes/paper7_phase2_results.md**) are in <https://github.com/alvaroNCubo/puppeteer>. The asciinema recordings and GIFs of §A.8 are in the **paper7-assets/** directory of <https://github.com/alvaroNCubo/puppeteer-papers>. Inline code citations of the form **file.cs:NN** resolve against the commit pinned in Code Provenance.

A.2 Demo orchestrator

The orchestrator script `docker/run-demo.sh` provisions three Docker containers, runs the analog bootstrap handshake against each, and exercises one cell of the matrix per invocation. Three flags select the scenario:

- `bash docker/run-demo.sh` — cross-container inline (22 entries final).
- `bash docker/run-demo.sh --parametric` — cross-container parametric (7 entries final).
- `bash docker/run-demo.sh --rehydrate-demo` — captures the cross-container leave-and-rejoin scenario referenced in §8: stops one container, allows the cluster to advance, restarts the container, and observes its rehydration and catch-up.

Each invocation runs in approximately 60–90 seconds on Docker Desktop.

A.3 Tests cited

The empirical claims in the body are exercised by specific tests in the suite:

Body section	Property	Test method	Project
§5.3	F1–F5 bootstrap handshake under real cryptography	EndToEnd_RealCryptoOverRealTls_RoundsTripIdentity	UnitTestChoreography/UsherOnboardingTests.cs
§5 (single peer operates alone)	Force-promoted Stage operates without peers	SingleStage_OperatesAlone_WhenPromotedForce	UnitTestPaper7EShop/ThreeNodeOrderingTests.cs
§5.4 (in-process, inline)	Three-stage Director rotation, inline regime	ThreeStages_DirectorRotation_AllConverge_HappyPath	UnitTestPaper7EShop/ThreeNodeOrderingTests.cs
§5.4 (in-process, parametric)	Three-stage Director rotation, parametric regime	ThreeStages_DirectorRotation_AllConverge_HappyPath_Parametric	UnitTestPaper7EShop/ThreeNodeOrderingTests.cs
§5 (rotation under cancellation)	Cancellation branch replicates across three stages	CancellationBranch_Replicates_AcrossThreeStages	UnitTestPaper7EShop/ThreeNodeOrderingTests.cs
§5.4 (instrumentation)	DLL load + per-phase convergence + journal bytes snapshot	Metrics_F1_Snapshot	UnitTestPaper7EShop/ThreeNodeOrderingTests.cs

Body section	Property	Test method	Project
§5.5 (partition tolerance)	Leave-and-rejoin from disk, peer-to-peer catch-up	<code>OneStage_Offline_OthersAdvance_Rehydrate_CatchUp_Parametric</code>	<code>UnitTestPaper7EShop/ThreeNode0rderingTests.cs</code>
§5.5 (rehydration without onboarding)	Rehydrated Stage restores identity from disk alone	<code>RehydratedStage_RestoresIdentityAndJournal_WithoutOnboarding</code>	<code>UnitTestPaper7EShop/ThreeNode0rderingTests.cs</code>

A.4 Cryptographic stack

The cryptographic primitives underlying the F1–F5 handshake of §5.3 are documented in `Choreography/Usher/Crypto/*.cs`. The implementations use BouncyCastle 2.6.2:

- Ed25519 keypair generation, signing, verification: `Ed25519StageKeyGenerator.cs`, `Ed25519Signer.cs`, `Ed25519SignatureVerifier.cs`.
- Ed25519-to-X25519 derivation (Edwards-to-Montgomery, $u = (1 + y) / (1 - y) \bmod p$): `Ed25519ToX25519.cs`.
- Sealed-box AEAD using X25519 ECDH and ChaCha20-Poly1305, with the participating public keys bound as additional authenticated data: `SealedBoxPayloadSealer.cs`.
- Kestrel TLS listener and self-signed certificate generation with SAN: `HttpsTransportListener.cs`, `SelfSignedCert.cs`.
- TOFU certificate fingerprint pin via `SocketsHttpHandler.SslOptions.RemoteCertificateValidationCallback`: `HttpsClientFactory.cs`.

A.5 Share-link URI

The share-link emitted by the issuer takes the form `puppeteer-usher://v1/{base64url(json)}`. The JSON payload carries the transport address of the issuer endpoint and a SHA-256 fingerprint (`fp` field) of its TLS certificate. The encoder/decoder is defined by `IUsherShareLinkEncoder` in `Choreography/Usher/`. A representative share-link under load is 449 characters in length, within the alphanumeric capacity of a QR code at version 14 with error correction level M (approximately 535 characters).

A.6 Wire format

The HTTPS transport between containers carries a `ChannelPurpose` byte after the sender identifier in each frame. This wire format (v2) prevents Define and

Invocation channels from collapsing into the same channel-dictionary key when Director rotation runs the parametric workload across multiple containers; the in-process regime is unaffected. The wire format is implemented in the transport handlers under `Choreography/Transport/Https/`.

A.7 Supporting reports

Two markdown reports in the `notes/` directory of the same branch document the experiment in detail:

- `notes/paper7_phase1_results.md` — the in-process row of the matrix (eleven sections plus three addenda).
- `notes/paper7_phase2_results.md` — the cross-container row, the hardening cycle that converged on the demo, the three-Docker mesh, the rotation protocol, and the parametric cell.

A.8 Recordings

Eleven cast recordings (asciinema format) and eleven GIF renders capture the four cells of the matrix of §5.4 and the rehydration property of §5.5, one recording per Docker container per cell. The recordings are published alongside this paper as supplementary material, hosted in the `paper7-assets/` directory of the `puppeteer-papers` repository.

Appendix B — Role-level diagnostic table

The lens of §7 — Paper 6’s *infrastructural symptom* applied at the role layer — produces a reading per role-in-use-case. The same node, identified in a deployment as “API server,” reads as an accidental category when deployed to accept authoritative writes against a shared database, and as a structural requirement when deployed to enforce cryptographic identity claims. The discriminator is the function, not the name. Seven canonical categories of architectural role admit reading under the two indicators of the lens; their substrate-property counterparts are summarized below as the role-level analog of Paper 6’s Table 1.

Table B.1. Diagnostic reading: seven canonical categories of architectural role under the two indicators of Paper 6’s lens, re-expressed at the role layer.

Architectural role	Representational choice that necessitates it	What a substrate satisfying both indicators offers instead
Server as write authority	Nodes replicate data; consistency requires a single point of acceptance	Each node accepts writes locally; the journal records the program, which replays deterministically on every node
Master / primary in replication	Nodes replicate state diffs; a primary must order them	Each node owns its journal; cross-node causality is recorded by reference, not orchestrated
Coordinator in distributed transactions	Multiple writers contend for a logical entity across processes	Each logical entity has a single point of authority intrinsic to the substrate; serialization is not externalized
Stateful gateway	Clients cannot directly invoke deferred work without losing causal record	Deferred work is journal-recorded; the substrate carries causality, not the gateway
Application server / runtime container	The runtime, the framework, and the application require disparate substrates	A minimal substrate suffices to run the domain artifact
Bootstrap service	New nodes require an issuer of credentials and addresses	Bootstrap is information transfer that crosses by any means, including analog artifacts; the issuer process is ephemeral
Orchestrator as coordination substrate	Stateful coordination across instances requires external choreography	Coordination is between actors; the substrate provides location and continuity intrinsically

The seven rows are not exhaustive; service discovery, queue-as-glue between bounded contexts, and configuration management admit similar analysis. The lens is diagnostic, not prescriptive: it enables auditing existing architectures under the two indicators of Paper 6, independently of any decision to adopt the

substrate of §4.

Code provenance

Source-code references in this paper resolve against the public Puppeteer repository at commit **b42d0f7** (2026-05-26). This commit includes the substrate code described in §4-§5, the reproducibility artifacts of Appendix A (**docker/**, **notes/**), and the SimpleX-based transport implementation referenced in §6 (**Choreography/Transport/SimpleX/**). Earlier commits in the repository preserve the substrate’s prior states: **91118fc** (2026-05-22) added the reproducibility artifacts; **2f31f96** (2026-05-18) is the earlier substrate snapshot before the SimpleX transport fixes (C2S envelope decryption, batch parsing, catch-up piggyback routing, StageManager Director-rotation race, Usher F5 invitation-consumed race) of commit **b42d0f7**. The earlier substrate snapshot is archived in Software Heritage under the following persistent identifier:

```
swh:1:dir:10e7e6bad7eb77b6c2e406762026177f95c3ae92;  
  origin=https://github.com/alvaroNCubo/puppeteer;  
  anchor=swh:1:rev:2f31f9674a5de816bdf1bf9d8360ff218a02e4da
```

Inline references of the form **file.cs:NN** (e.g., **ActorHandler.cs:38**) resolve against **b42d0f7**. A reader can construct a per-file SWHID by adding the qualifiers **;path=<path>;lines=<NN>** to the directory SWHID above for the **2f31f96** snapshot; for the post-**2f31f96** state (which includes the SimpleX transport fixes of **b42d0f7**), the GitHub **b42d0f7** reference is canonical. Future commits to the repository may renumber lines; the SWHID preserves the cited **2f31f96** state independently of any future change to the repository or its hosting.

Acknowledgments

The author used large language models (including Claude and ChatGPT) as editorial assistants for language refinement, structural feedback, and literature navigation. All original ideas, terminology, theoretical constructs, and technical content presented in this work are solely the author’s.

References

- Ameloot, T. J., Neven, F., & Van den Bussche, J. (2013). Relational transducers for declarative networking. *Journal of the ACM*, 60(2), 1–38.
- Armstrong, J. (2003). *Making reliable distributed systems in the presence of software errors* [Doctoral thesis]. KTH Royal Institute of Technology.
- Bernstein, P. A., Bykov, S., Geller, A., Kliot, G., & Thelin, J. (2014). *Orleans: Distributed virtual actors for programmability and scalability* (Tech. Rep. MSR-TR-2014-41). Microsoft Research.

- Brewer, E. A. (2000). Towards robust distributed systems [Keynote address]. *Proceedings of the Nineteenth Annual ACM Symposium on Principles of Distributed Computing (PODC '00)*.
- Cognitect. (n.d.). *Datomic: A transactional database with a flexible data model, queryable as a log of changes* [Software project]. Retrieved May 22, 2026, from <https://www.datomic.com/>
- DeCandia, G., Hastorun, D., Jampani, M., Kakulapati, G., Lakshman, A., Pilchin, A., Sivasubramanian, S., Voshall, P., & Vogels, W. (2007). Dynamo: Amazon’s highly available key-value store. *Proceedings of the 21st ACM SIGOPS Symposium on Operating Systems Principles (SOSP '07)*, 205–220.
- Earthstar Project. (n.d.). *Earthstar: simple personal data sync* [Software project]. Retrieved May 22, 2026, from <https://earthstar-project.org/>
- Evans, E. (2003). *Domain-driven design: Tackling complexity in the heart of software*. Addison-Wesley.
- Event Store Ltd. (n.d.). *EventStoreDB: The database for event sourcing* [Software project]. Retrieved May 23, 2026, from <https://www.eventstore.com/eventstoredb>
- Fowler, M. (2005). Event sourcing. Retrieved from <https://martinfowler.com/eaDev/EventSourcing.html>
- Gilbert, S., & Lynch, N. (2002). Brewer’s conjecture and the feasibility of consistent, available, partition-tolerant web services. *ACM SIGACT News*, 33(2), 51–59.
- Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.
- Harris-Braun, E., Luck, N., & Brock, A. (2018). *Holochain: Scalable agent-centric distributed computing* (Alpha 1 white paper). Holo. <https://www.holochain.org/documents/holochain-white-paper-alpha.pdf>
- Hellerstein, J. M. (2010). The declarative imperative: Experiences and conjectures in distributed logic. *ACM SIGMOD Record*, 39(1), 5–19.
- Hewitt, C., Bishop, P., & Steiger, R. (1973). A universal modular ACTOR formalism for artificial intelligence. *Proceedings of the 3rd International Joint Conference on Artificial Intelligence (IJCAI)*, 235–245.
- Holepunch. (n.d.). *Hypercore: append-only logs for the peer-to-peer web* [Software project]. Retrieved May 22, 2026, from <https://docs.pears.com/building-blocks/hypercore>
- Kleppmann, M., Frazee, P., Gold, J., Graber, J., Holmgren, D., Ivy, D., Johnson, J., Newbold, B., & Volpert, J. (2024). Bluesky and the AT Protocol: Usable decentralized social media. In *Proceedings of the ACM*

CoNEXT-2024 Workshop on the Decentralization of the Internet. ACM. <https://doi.org/10.1145/3694809.3700740>

Kleppmann, M., Wiggins, A., van Hardenberg, P., & McGranaghan, M. (2019). Local-first software: You own your data, in spite of the cloud. In *Onward! 2019: Proceedings of the 2019 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software* (pp. 154–178). ACM. <https://doi.org/10.1145/3359591.3359737>

Lemmer-Webber, C., Tallon, J., Shepherd, E., Guy, A., & Prodromou, E. (2018). *ActivityPub* (W3C Recommendation, 23 January 2018). World Wide Web Consortium. <https://www.w3.org/TR/activitypub/>

Lightbend. (n.d.). *Akka Persistence* [Software documentation]. Retrieved May 23, 2026, from <https://doc.akka.io/libraries/akka-core/current/typed/persistence.html>

Matrix.org Foundation. (n.d.). *Matrix specification* [Online document]. Retrieved May 22, 2026, from <https://spec.matrix.org/>

Microsoft. (2026). *dotnet/eShop: A reference .NET application implementing an eCommerce site* [Software repository, commit 9b4f943, 2026-04-21]. Retrieved May 22, 2026, from <https://github.com/dotnet/eShop>

Nygard, M. T. (2018). *Release it! Design and deploy production-ready software* (2nd ed.). Pragmatic Bookshelf.

Ogden, M., McKelvey, K., & Madsen, M. B. (2017). *Dat: Distributed dataset synchronization and versioning* [White paper]. Code for Science. <https://github.com/datprotocol/whitepaper>

Ongaro, D., & Ousterhout, J. (2014). In search of an understandable consensus algorithm. *Proceedings of the 2014 USENIX Annual Technical Conference (USENIX ATC '14)*, 305–319.

Rivera, A. (2026a). Anti-porous architecture: a unified design principle for CQRS + Actor + Event-Sourcing systems. *Puppeteer Papers Series*, Paper 1. Zenodo. <https://doi.org/10.5281/zenodo.20404863>

Rivera, A. (2026b). Program–value separability: the structural precondition for compilation, caching, and dense journaling in a DSL runtime. *Puppeteer Papers Series*, Paper 2. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/02-program-value-separability.md>

Rivera, A. (2026c). Reactions and the partition: opt-in eventual consistency in actor-native systems. *Puppeteer Papers Series*, Paper 3. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/03-reactions-and-partition.md>

Rivera, A. (2026d). Preserving semantic continuity across actors: a tell-based approach without orchestration. *Puppeteer Papers Series*, Paper

4. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/04-cross-actor-continuity.md>
- Rivera, A. (2026e). The journal as substrate: unifying deployment, replication, backup, and offline operation in distributed systems. *Puppeteer Papers Series*, Paper 5. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/05-substrate-operations.md>
- Rivera, A. (2026f). Most infrastructure layers are symptoms of the persistence model: a construct for auditing production stacks. *Puppeteer Papers Series*, Paper 6. <https://github.com/alvaroNCubo/puppeteer-papers/blob/main/06-infrastructure-symptom.md>
- Sambra, A. V., Mansour, E., Hawke, S., Zereba, M., Greco, N., Ghanem, A., Zagidulin, D., Abounaga, A., & Berners-Lee, T. (2016). *Solid: A platform for decentralized social applications based on linked data* (Technical Report). MIT CSAIL & Qatar Computing Research Institute. http://emansour.com/research/lusail/solid_protocols.pdf
- Shapiro, M., Preguiça, N., Baquero, C., & Zawirski, M. (2011). *Conflict-free replicated data types* (Research Report RR-7687). INRIA.
- Tarr, D., Lavoie, E., Meyer, A., & Tschudin, C. (2019). Secure Scuttlebutt: An identity-centric protocol for subjective and decentralized applications. In *Proceedings of the 6th ACM Conference on Information-Centric Networking (ICN '19)* (pp. 1–11). ACM. <https://doi.org/10.1145/3357150.3357396>
- Vernon, V. (2013). *Implementing domain-driven design*. Addison-Wesley.
- Vernon, V. (2015). *Reactive messaging patterns with the actor model: Applications and integration in Scala and Akka*. Addison-Wesley.
- Young, G. (2010). *CQRS documents*. Retrieved from https://cQRS.files.wordpress.com/2010/11/cQRS_documents.pdf

Inference without Authority: the three authorities that govern an output

TL;DR

Every program has a `print`, and it is taken for granted that the code deciding *what* to emit also decides *where* it goes. These are distinct authorities. Pulled apart, they reveal a third: what an output *is* — its projection — is the actor’s to author, not the domain’s. An output is governed by three competences: the domain (what exists), the actor (what becomes observable), and the assembler (where observation occurs). Collapsing any two forces one authority to pronounce a decision it cannot know — an *assertion beyond warrant*, the paper’s construct, which sits at both ends of an output: a producer stipulates a *where* it cannot know; an observer given only snapshots *infers* a *how* it was never told (the genuine inference the title names). Restoring the authorities dissolves both. The contribution is analytic: an architecture is sound not by how it divides code, but by whether the authorities within it match the decisions each part has the standing to make.

Dependencies. This paper is part of the *Puppeteer Papers*, a series of self-deposited preprints, and rests on two of them: the actor’s speech and `tell` (Paper 4) and the journal as substrate (Paper 5), constructions established there rather than re-argued. Its symptom criterion is stated and justified here (§3), echoing — but not depending on — the one Paper 6 applied to infrastructure. Methodologically, it is an analytic theory contribution in the sense of Gregor’s (2006) theory for analyzing (Type I): the three authorities and the *assertion beyond warrant* are constructs by which an architecture is described and judged, while the labs of Appendix A are an existence proof of realizability, not a design-science evaluation of cost or benefit.

1. `print`

Every programming language has a `print` statement. It is so familiar that it has become invisible. A beginner learns that `print` displays text; a professional, that it writes to a log; a distributed-systems engineer, that it publishes to a topic, updates a read model, refreshes a cache, or emits telemetry. On what `print` does, everyone agrees. On a prior question — who is entitled to decide *where* its output goes — almost no one pauses.

This is not a paper about printing. It is a paper about authority over the destination of an output, and about what the familiar shape of `print` has quietly kept out of view.

Consider an accountant. She can specify, in complete detail, how the total of a customer's invoice is computed from an order: which lines are taxable, how discounts compose, how rounding is applied, when a line is excluded. Every rule of that projection is hers to state. Put to her a different question — should this invoice be written to MySQL or to PostgreSQL? delivered over a message topic, held in a cache, pushed to a dashboard? — and her answer is immediate, and correct: she has no idea. The destination is simply not something the computation of an invoice has an opinion about.

Yet business software is built, routinely, to answer exactly that question in the same place it computes the invoice. In a typical web service the whole of it is a single statement — `return Ok(invoice)` — and that one line names the total to show *and*, at the same stroke, fixes that the result is JSON, that it travels over HTTP, that it carries a 200, that it is synchronous. The module that knows how to total an order has also named the format, the protocol, the status, the timing. Two decisions — *what* becomes observable and *where* observation occurs — are folded into one body of code, and the fold is so customary that it goes unseen, in precisely the way `print` goes unseen.

The accountant sharpens the point because she is the producer of the observable and still cannot name its destination. This is not a limitation peculiar to accountants. It is a property of the relation between a computation and its output: the party that authors what becomes observable is not, in general, the party positioned to decide where observation occurs. The producer of an output can decide *what becomes observable*; it cannot decide *where observation occurs*.

If that is so, two questions follow, and the rest of this paper is their pursuit. First (§2–§4): if the producer does not decide the destination, who does — and what does that reveal about what an output actually *is*, once it is no longer mistaken for a write? Second (§5 onward): if the destination is bound elsewhere, how does an *observer* come to know what was made observable — and why does the conventional answer force the observer into an inference it should never have had to make?

2. The Second Voice

Section 1 ended with a separation. One party knows what should become observable — the invoice, its line items, its total. Another must decide where observation occurs — the database, the topic, the cache. The question is no longer whether these are different decisions; §1 has shown that they are. It is who is entitled to pronounce the second one.

The producer cannot simply acquire that standing. Nothing about how an invoice total is computed teaches where invoices are kept; the rules of the projection are silent on the choice of sink, and no refinement of them will make one speak it. The knowledge is not there to be found. The decision must therefore be pronounced somewhere else.

It is worth being exact about “somewhere else,” because the usual answers are too weak. The destination is not merely computed in another layer, delegated to another module, or hidden behind another class. Those are the same speaker — one program, one author, arranging its own internals. What §1 forces is not a rearrangement of code but a difference of *standing* — and standing, here and throughout, is epistemic. It follows the ordinary norm of assertion: one may state as fact only what one is in a position to know. The producer’s want of standing to name the destination is therefore not a rule laid upon it; it is the plain consequence of a fact about its knowledge — the destination is not in it. The destination belongs to a different authority, and must be spoken in another voice; the authority is the ground, the voice only how it makes itself heard.

This series has met the notion of a voice before. Paper 4 established that an actor has one: it records its deeds in its own vocabulary, and may say only what it could genuinely have said. That discipline turned on the boundary of a single voice. The present separation is the discovery that, for output, a single voice is not enough — because a single voice carries a single competence, and the output answers to two. The actor can pronounce the projection: it lived the order, it knows the total. It cannot pronounce the destination: nothing it lived tells it where the total is kept. It has no standing to pronounce it.

This is the paper’s claim in one line. The actor speaks the projection; another authority speaks the destination. Where Paper 4 showed that the actor speaks, this paper shows that, for its output to be described correctly, the actor does not speak alone.

Because it turns on standing and not on mechanism, the separation is easy to counterfeit. A system can be split into layers, modules, classes, even two languages, and still speak with a single voice: splitting code does not split authority. A second module is not necessarily a second voice. What tells them apart is not whether there is another class, but who could author the second decision without modifying the first — for a second voice exists only when the second decision can be authored, replaced, and evolved independently, because it belongs to a different authority.

The principle outruns output. An architecture is not measured by how many layers it holds, but by whether the entitlements within it match the decisions each part has the standing to make. Measured so, most software has a single voice and pronounces the destination with it — what becomes observable and where, spoken as though they were one competence when they are two. That folding is so ordinary as to be invisible; §3 takes it up as what it is: a symptom.

3. The Symptom

Return to the collapse §2 ended on. When one voice is made to speak both what becomes observable and where, something has to fill the place where the second authority should have spoken. The voice has no knowledge of the destination — that was the whole point; nothing it lived tells it where its output is kept. Yet a destination is produced all the same. It has to be: the code must compile, the write must resolve, the line must end somewhere. So the voice supplies what it cannot know. It *asserts* the destination — stipulating it, or taking a framework’s default — and lets that stand as fact.

That is the phenomenon, and it is worth being exact about it. The producer does not reason its way to MySQL; it weighs no evidence and concludes nothing. It settles on a destination and lets it stand — writes it, or inherits a default. The hardcoded sink is not a fact the producer possessed, nor one it inferred; it is a value it was forced to fabricate, standing in for an authority that was never given a voice.

This does not indict the producer’s other work, and the distinction is the whole of it. The projection *is* an inference, and a warranted one: the invoice total is derived, the subtotal computed, and the producer has the standing to derive them — it lived the order. What sits beside that warranted inference is not a second inference but a bare *assertion*: the destination, committed to as fact without the standing to know it. The symptom is not that a voice infers — inference, where warranted, is its proper work. It is that a voice asserts beyond its warrant.

None of this was derived from a prior construct. The assertion was forced by the argument of §1 and §2 — a fused authority, a decision it cannot know — before any name was reached for. What remains is to name it, and to say by what test the name is earned. Call a feature of a system a *symptom* when it bears two marks. It *compensates*: it stands in for something absent, doing the work of a capability the system does not have. And it *dissolves*: supply the missing capability and the feature vanishes, with nothing left for it to do. Together the marks separate what a problem genuinely demands from what a deficiency elsewhere induced — a feature the problem itself requires would persist even in an otherwise ideal system, while one that disappears the moment some other lack is repaired was never answering the problem, only the lack. Dissolution-on-repair is the diagnostic: the counterfactual that tells an induced feature from

an intrinsic one.

The unwarranted assertion bears both marks exactly, now about a decision rather than a component. It compensates: the producer cannot know the destination, so the assertion stands in for the competence that could — it is what the single voice says because a second voice was not there to say it. And it dissolves: give the destination its own authority, let a second voice pronounce it, and the assertion simply disappears; there is nothing left for it to compensate for. What looked like a decision the producer had to make turns out to be an artifact of who was allowed to speak. That is what earns the name. The symptom is an **assertion beyond warrant**: a value a voice commits to as fact only because it was made to pronounce a decision it had no standing to know. (The same test recurs across the series: Paper 6 applied these two marks to infrastructure layers — a cache, an object-relational mapper, a queue — asking whether each existed because the problem demanded it or because the persistence model was deficient. Here the object is a decision, one level in, rather than a component; the criterion is the one just stated, not one borrowed to be believed on another paper’s word.) Its most literal form — a genuine inference — waits at the other end of the output, in §5; that is where the word this paper is named for fits without strain.

This is also what §2’s criterion looks like from the inside. An architecture is sound when its authorities match its decisions; the symptom is the moment they do not — when a decision is pronounced by a voice with no claim to it, and the gap between the two is papered over by an assertion nothing warrants. Read this way, a hardcoded destination is not a small infelicity of style. It is the visible trace of a missing authority.

Which raises the question a reader has been holding since §2. If the separation is forced by the argument, and its absence leaves a mark this legible — does any system actually honor it? §4 answers.

4. The Three Authorities

By §3 the argument has separated two authorities: the actor, which pronounces the projection, and the assembler, which pronounces the destination. But the projection has a boundary on its other side as well, and naming it completes the structure.

The actor authors the projection — but not out of nothing. It authors it over a material it did not invent: the objects, the state, the operations that make an order an order. That material is the domain’s. The domain knows what exists; the actor knows only what, of what exists, becomes observable. These are not one competence, and the distance between them shows itself the moment a domain object renders itself: in doing so the domain has reached past *what exists* into *what becomes observable* and taken the projection, which was never its to take. The symptom there is not that the object prints. It is that the

material spoke as though it held the projection’s authority.

Three competences, then, not two — and the series has been uncovering them in turn. Paper 4 found that the actor may not name the *transport* that carries its voice to a peer. This paper found that it may not name the *sink* where its voice lands for an observer. And now a third, pointing the other way: the domain may not author the *projection* — what becomes observable belongs to the actor, not to the material. The actor stands in the middle, its competence bounded on both sides:

Authority	Knows	Governs
Domain	what exists	the material
Actor	what becomes observable	the projection
Assembler	where observation occurs	the destination

Table 1. The three authorities that govern an output.




Figure 1. The actor bounded on both sides — a soft, negotiable edge to the domain and a hard, fixed edge to the assembler — with each symptom (§3–§4) marked as a reach across an edge into a decision the reaching party has no standing to make.

Read against this, the two symptoms are one kind of thing: a competence reaching across a boundary. A domain object that renders itself is the domain reaching up into the projection; a hardcoded sink is the actor reaching up into the destination. In each, one authority pronounces a decision that belonged to another and papers the gap with an assertion it has no standing to make. This is not a stack of layers — layers stack the same speaker — but a division of what each party is in a position to know.

A realization of the structure follows directly from it: give each competence its own place to speak, and the symptoms have nowhere to form. A reader who has come this far is right to ask whether anything actually builds it. One system does. Puppeteer is one such realization, and — because it was not built to make this paper’s point — the separation it exhibits was already in place before the point was named: found, not staged. That is the whole of the claim, and it is a demonstration of buildability, not a measurement. Three labs on the running framework (Appendix A) make it concrete: the same actor and the same projection reach genuinely different destinations — real SQL Server and MySQL among them — with no edit to the producer; one journaled fact is observed by three distinct views without a single new domain method; and an observer reaches the same knowledge whether it is told or polls. Each is reported against a fused baseline, which pays one edit — one producer change, one domain method — for what the separation gives for nothing. What the labs do not do is measure: no cost against that baseline at scale and no benefit in production, which are left to other work. The three authorities are the argument’s, not the framework’s.

In it, a program is written across two languages. The domain — an `Order` and its `OrderItems`, the operations each genuinely performs — is written in a general-purpose host language, and it holds *what exists*. An actor is a journaled unit: its state is not stored and mutated in place but reconstructed by replaying the operations it has recorded (Papers 1–2; the journal of Paper 5), so it has a *history* — the journal of those operations — and, at any moment, a *state* — the object graph that history reconstructs. What the actor does when a command reaches it is written in a small DSL — it declares, branches with `if`, walks with `foreach`, and emits with `print` and its durable form `expose` — and what it authors is *what becomes observable*: it walks its state (the order’s items, each item’s product name and its subtotal — unit price times units), selects the primitive leaves the projection needs, and emits them, naming a logical output and no destination at all.

In the DSL the projection is a walk that narrows and emits, and where it narrows is a boundary of its own. Over Microsoft’s `dotnet/eShop Order` aggregate — the same port the systems papers use — an invoice over discounted lines reads:

```
foreach (item in order.OrderItems) {
    if (item.Discount > 0) {
        print item.ProductName product,
              item.UnitPrice * item.Units subtotal;
    }
}
```

Every name is the domain’s own — `OrderItems`, `ProductName`, `UnitPrice`, `Units`, `Discount` — and the subtotal is computed on the way out (`UnitPrice * Units`), a derivation, not a stored field. The narrowing sits in the actor’s `if`; it could instead be pushed into the material, the domain exposing the narrowed collection as a verb of its own (`order.DiscountedItems()`), which is legitimate

only when *discounted items* is a notion the domain could genuinely own — worth shaping its write representation for. That shaping is not hypothetical here: the eShop item already copies its product’s name onto itself rather than pointing at a catalog, so a line can be read without leaving it. Pushing a selection down is the same trade; making it for a view the domain does not own is the symptom. Either way, **print** emits values under names and names no place. The assembler, outside the actor, binds that logical output to *where observation occurs*: a table in MySQL, a topic in Kafka, a hub, a file, the console, nothing — and whether it is pulled or pushed. That the script is unchanged either way is not asserted here only rhetorically: the framework’s own output interface documents it — *pull versus push is a property of the destination, never of **print*** — with the sink configured outside the actor, at the hosting handler, identically across every topology.¹ The shape is Unix’s **>**, with the modern range of sinks in place of a file.

And the shell’s other operator has a counterpart too. Where **>** redirects a program’s output to a sink, **|** pipes it to another program; Puppeteer’s **|** is the **tell** of Paper 4 — the same projection sent not to a store but to a peer actor, received as a told message. The two are distinct planes, not one stream routed a single way, so an actor’s output can want both operators at once — **program params > target | peer params**, redirected to a sink *and* piped to a peer. No shell grants that in one stream; it makes you choose **> file** or **| program**. Puppeteer does not, because the destination is the assembler’s to bind and the pipe is Paper 4’s **tell** to carry — and the projection in the middle, as before, learns of neither.

Put §2’s test to it. Who could send the invoice to Kafka instead of MySQL without modifying the actor? The assembler alone; the actor is not touched, not recompiled, not reread. The same actor that writes to a warehouse in production will, run from a command line with its output bound to the console, print the same invoice to the terminal, where it can be piped into the ordinary tools. Two destinations, one projection, and the projection learned of neither. That independence is the operational face of the separation: each competence can be authored and revised without leave from the others.

These two edges are not equally hard, and the structure is more honest for saying so. The actor/assembler boundary does not move: no refinement of how a total is computed will yield where totals are kept, so the destination is never the producer’s to name — the knowledge is simply absent from it. The domain/actor boundary is not fixed in that way. Whether *discounted items* is *what exists* — a notion the domain genuinely owns, worth a place in its model — or only *what becomes observable* — a narrowing the actor imposes for a view — is a modeling judgment, and for a given concept it can fall either way without anyone overreaching; that is why the narrowing above could be pushed into an **order.DiscountedItems()** as a legitimate alternative rather than a mistake. The line between the material and the projection is drawn per concept; the line between the projection and the destination is not drawn but discovered. What

does not move is the rule that governs both once a line is set: neither side may reach across it — a domain object that renders itself has reached past whatever it models as *what exists* into a projection it never chose to own, and that is the symptom wherever the boundary sits. Table 1 names the three authorities and Figure 1 draws their two boundaries — but the two are not edges of one hardness: one is a fact about absent knowledge, the other a judgment about owned knowledge. The table lists the competences; the figure marks that their boundaries are not cut equally deep.

There is a way to see the hard edge without arguing it: test for it. If the destination were the producer's, no test of the projection could run without also standing up the sink — to exercise what the actor computes you would first have to say where its output goes. The opposite is what holds. The projection can be exercised end to end — the order placed, the invoice totalled, the result read back — with no destination bound at all, on the pull path alone: the answer returns to the caller and nothing is sent anywhere. That a test of the domain needs no sink is not a convenience the framework grants; it is the operational proof that the sink was never in the domain to supply. This is the epistemic claim made verifiable, and falsifiable in the plainest sense the word allows: were the destination truly the producer's, no such test could pass without it. Here the analysis of §4 meets something checkable — the labs of Appendix A run exactly this test, exercising the projection while the destination stays a separate, unbound choice.

One primitive has been named but not yet used. Where **print** makes a projection observable in the moment, **expose** persists it into the actor's history — the projection enters the record the actor keeps, durable and replayable. That difference is the hinge of §5, because it is the difference between an observation that merely passes and one that enters a narrative an observer can later be told.

Puppeteer is offered as *a* realization, not *the* realization. The paper does not claim this DSL is necessary; §2 was explicit that each authority may be housed in any language, or in none. The claim is that the three authorities are real, that they can be kept apart, and that where they are, the material, the projection, and the destination each leave the others' voice. That one working case shows the separation is not only arguable but built — and what it opens for the party on the receiving end, the observer, is §5.

5. The Observer

Every section so far has stood beside the producer. Turn now to the party on the other end — the one that receives what was made observable. Call it the observer: a dashboard, a report, a downstream service, a screen. There is a question about the observer this paper does not take up — how it *interprets* what it receives, reading a sequence of acts for the routines they compose. The question here is narrower, and prior: how does the observer come to *know* the

state at all, and does the symptom of §3 reach it too? It does — and its situation turns out to mirror the producer’s exactly.

An observer comes to know the state by one of two means. It can read a *snapshot* — the configuration as it stands right now: this invoice, this total, this set of open orders. Or it can receive a *narration* — the sequence of operations that brought the state about: the order placed, the line added, the discount applied, the invoice settled. The two are not interchangeable. Both yield the same configuration; only the narration preserves how the configuration came to be. A snapshot is an answer with its reasons erased.

Now give the observer only snapshots, and ask of it anything the snapshot does not carry — how the total was reached, what has changed since it last looked, why an order stands where it does. It cannot read the answer off; the snapshot does not hold it. Yet it must answer all the same, and so it does what the producer did in §3: it reconstructs what it was not given. It diffs successive snapshots to guess what changed; it re-derives the path from the endpoint; it polls, and polls again, rebuilding a history from a series of stills. The observer infers the history it was not given.

And this is the same defect, at the far end of the voice — but here, at last, it takes the form of a genuine inference. The history was never the observer’s to know; it belonged to the actor, which lived it. Reconstructing it from snapshots is an assertion reached beyond the observer’s standing, exactly as the stipulated sink was — except that the observer truly *infers* it, diffing and re-deriving, where the producer merely stipulated. **This is the inference the paper is named for:** not the producer’s quiet stipulation, but the observer’s laborious reconstruction of what it was never told. What §3 found at the point of emission recurs at the point of reception: the producer asserted a *where* it could not know; the observer infers a *how* it was never given. One defect — assertion beyond warrant — book-ending the output, met once by stipulation and once by inference.

The polling is not the observer’s failing; it is forced, in the same way the producer’s stipulation was forced. An observer handed a snapshot and left to rebuild the rest has been given the configuration but not the account behind it — a boundary that shows the answer and withholds its reasons. Its polling is the mechanical signature of a missing narration.

The dissolution is the one already seen, read from the other side. The earlier papers built a substrate in which the record is not a snapshot but an account: the journal (Paper 5) is the sequence of operations the actor lived, and each entry is something the actor could genuinely have said (Paper 4). Deliver that to the observer — the record the actor keeps, into which **expose** enters the observations meant to last — and the observer has nothing left to reconstruct. It is not diffing stills; it is being told.

What the observer has then is *testimony*: it comes to know the state the way anyone comes to know a fact they did not witness — on the word of the party that did. And this names the shape of the whole paper. The two — the

producer's stipulated *where* and the observer's inferred *how* — were born in one place: absent information. The producer was not given the destination, and asserted one; the observer was not given the account, and reconstructed it. Neither gap is closed by a mechanism. Each is closed by a kind of knowledge delivered to the party that lacked it — the destination, owned by a second authority; the account, carried by the actor's own record. The observer stops being an inference engine; it still computes the view it needs, but over an account it was given, not a reconstruction it was forced to fabricate. Spoken in full and received as testimony, the voice leaves nothing at either end for an unwarranted assertion to fill.

This resolution has a boundary, and in distributed practice not a small one. Testimony is knowledge on another's word, and only as strong as the standing of the one who gives it. Within a trust boundary — producer and observer under one authority, no adversary between them — the account can be taken as given, and the observer is spared its reconstruction. Across a trust boundary, which for anything leaving its own trust domain is the ordinary condition rather than the exception, the observer cannot simply believe what it is told: it must verify, and verification is a need this paper names but does not meet (§6). The dissolution just reached is therefore real but scoped — it frees the observer from rebuilding a *withheld* history, not from checking an *untrusted* one.

Paper 4 gave the actor speech; this paper lets the observer receive it as testimony. But a narration received is not yet a narrative recognized — what the observer makes of the account it is told, reading a sequence of acts as the routines the actor performed, is a further question, beyond the scope of this paper.

6. Limits

An argument is only as trustworthy as the boundary it draws around itself, the more so when it ends on a dissolution — the point at which it is easiest to claim more than was shown.

The hardest limit follows from the paper's own root. The symptom was traced to absent information: a party forced to act on knowledge it was not given fills the gap with an assertion nothing warrants, and restoring the authority that holds the knowledge makes it vanish. But that cure works only where the absence is *architectural* — where the knowledge exists, held by some party, and was merely withheld by a collapsed boundary. It does nothing for an absence imposed by the world. An observer that must not only read but *write* while cut off from the account is denied the state by physics, not by architecture; its speculation about what it cannot see is genuine, and no rearrangement of authorities abolishes it. §5 frees the observer from reconstructing a history that was withheld; it does not free a writer from guessing at a history that has not reached it. The three authorities dissolve absent-information-by-design; they leave untouched the real difficulty of absent-information-by-disconnection.

A second limit lies in the word the paper ends on. Testimony is knowledge on the word of another, and it is only as good as the standing of the one who gives it. The observer is spared inference precisely because it trusts the source — the account reaches it as something the actor could genuinely have said. And across a trust boundary that assumption usually fails — between organizations, over a public network, anywhere the producer is not under the observer’s own authority — so the untrusted producer is the default case in distributed systems, not an edge of it. There an adversarial or unreliable producer makes testimony insufficient, and the observer needs verification, a different question this paper does not take up. This is the boundary of the result’s scope, not a caveat within it: the argument resolves absent-information inside a trust domain and stops where trust does.

Nor does the paper solve delivery. It has said that the assembler binds the destination and carries the account to the observer; it has not said how the account survives a partition, or reaches a thousand observers rather than one. Those are questions of transport, answered by the substrate the earlier papers built, not by the division of authorities argued here. Fan-out scales with the transport, not with the journal — a slow or lossy channel is a real constraint, and an orthogonal one.

And the symptom is conditional on the knowledge being absent — worth stating plainly, because it bounds the claim. Where a single hand genuinely holds both knowledges — a sink fixed and known at authoring time, in a program that will only ever deploy one way — naming the destination in the producer asserts nothing it cannot back, because nothing was unknown; there is no symptom, only a coupling. What separation buys there is not correctness but evolvability: the freedom to change the sink without editing the producer. The epistemic claim is the stronger one, and it bites wherever the two knowledges genuinely part — most of production software, but not all of it.

Finally, the structure earns its keep only where there is a second end for it to have. In a domain where the party that produces the output and the party that observes it are one — a single-player program with no downstream reader — the symmetry has nothing to bite on, and the apparatus is overhead. The claim is not that every output has three separate authorities in practice, but that where the roles are genuinely distinct, so are the competences, and collapsing them leaves a mark.

A final scope, and it is the widest. This paper argues a single axis — the output — and demonstrates the three authorities only there. The choice is not arbitrary: the output is the axis whose boundary is **hard**. The destination is simply not in the producer, and no refinement puts it there, so the separation shows without noise. Other axes submit to the same authority *question* — what the world speaks in, the clock a computation reads, the draw it makes — but their boundaries are subtler, because a capture on those axes must decide *what* to take and *when*, where **print** decides nothing of the sort. This paper does not settle them; it claims the output alone, as the first and cleanest axis of a

principle, not the whole of one. To read it as a promise of total portability would be to credit it with work it has not done — and to invite exactly the question it does not answer.

7. Related Work

The idea that a program should not name where its output goes is not new; what is new is reading it as a matter of authority. The mechanism has a long lineage. Unix bound the destination outside the program — a process writes to a file descriptor and the shell decides, through `>` and `|`, whether that descriptor is a terminal, a file, or another program (Ritchie & Thompson, 1974); Plan 9 pushed the same indirection into per-process name spaces (Pike et al., 1992); flow-based programming (Morrison, 1994) and Kahn process networks (Kahn, 1974) gave components ports and let an external topology decide what connects to what. All of these late-bind the destination. None asks whether the program was *entitled* to name it; the binding is offered as composition, not as competence.

Three lines come closer, and each holds one edge of the argument without its reframe. Coordination languages (Gelernter & Carriero, 1992) split a program into a computation model and a coordination model — what it computes apart from how it is connected — which is the material/assembler distinction in all but name, drawn as modularity, with no epistemology attached. The object-capability model (Miller, 2006) is built entirely on authority — no ambient authority, least privilege, the right to act conferred rather than assumed — but turns that authority on access and effect, not on the projection or destination of an output. One line holds the split, the other the authority; neither turns either on output as an epistemic act.

The third line is the sharpest test, and the paper’s novelty stands or falls on it. Algebraic effects and their handlers (Plotkin & Pretnar, 2009) give the inversion its most rigorous form: a computation *performs* an operation — **print** among them — without fixing its interpretation, and a handler installed by an enclosing context decides what becomes of it, covering both directions of the boundary with a semantics rather than a convention. Were the contribution only *that the destination is bound elsewhere*, effects would already hold it whole, and this paper would be late binding with a story attached. The delta is that the effect discipline is a mechanism and the claim here is a criterion that judges the mechanism’s use. An effect system relocates *where* an interpretation is written; it says nothing of whether the party writing it had the standing to know the value it supplies. Nothing in it stops a handler from hardcoding MySQL — the stipulation of §3, now lodged inside the handler — and the program stays perfectly well-typed, because effect typing tracks *which operations are handled*, a matter of control and safety, never whether a handler was entitled to the value it names. Effect-safety and authority-soundness are therefore independent: a computation can be fully effect-safe and still carry the symptom, which lives on a different axis — not who *intercepts* an operation, but who is in a position to

know its value. This is the epistemic-versus-compositional distinction the paper turns on, and it is load-bearing, not cosmetic: withdraw it and nothing lets one call an effect-handled `print` with a fixed sink defective, for by the calculus's own lights it is not. The epistemic reading is not a gloss on late binding; it is the axis along which a program that already uses effects is graded.

Two structural gaps confirm the boundary is not the same one. Effects split an operation from its interpretation — a binary — where this paper finds three authorities; and the finest of the three, the domain's inability to author the projection (the self-rendering object of §4), has no counterpart in an effect, whose handled value the producer computes with no sense that its *shape* answers to a different competence than the *material* it is drawn from. And effects are producer-side throughout: they say nothing of how an *observer* comes to know a state, and so leave untouched the second half of the argument — the snapshot-versus-narration symptom of §5, where the same defect returns with no operation to handle. The split, the inversion, the authority: each neighbouring line holds one, none holds all three, and none reads output as an epistemic act.

Closest, and worth separating out, is the revised Single Responsibility Principle, which asks that a module answer to a single *actor* — a stakeholder with the authority to request change (Martin, 2017, ch. 7). This is separation by authority, the very criterion argued here — but its authority is organizational: who may *ask* for a change. The authority of this paper is epistemic: who is in a position to *know*. The producer does not withhold the destination out of good manners toward another team; it cannot supply it, because the knowledge is not in the projection. That distinction is what separates the present claim from the ordinary counsel of layered architecture. Controllers and views, services and DTOs, records — these divide technologies and structure, not authorities. When a controller returns a view, a DTO, a file, the same code that chose what to show has also fixed that it is HTTP, and JSON, and synchronous; the DTO does not defer the destination but freezes the projection into a value, and changing the sink recompiles the producer.

The objection nearest to a practitioner is that this is dependency injection, or hexagonal architecture — ports and adapters — renamed. It is best met on its own ground, which is the test suite, not philosophy. Hexagonal inverts the dependency: the domain declares a *port*, an output interface, and an adapter chosen from outside implements it. That is genuine late binding — but a port is still something the domain names and depends upon, and to test the domain one must supply a double for it, a fake sink standing in for the real. This separation leaves no port. `print` depends on no output interface, not even an abstract one; it emits values under a logical name and knows of no destination to invert. The difference is observable, not interpretive: count the doubles a domain test must stand up for its output — under hexagonal, at least one; here, none, because there is nothing there to double (§4). Inversion relocates a dependency; this removes it, and that is the line injection cannot cross, because

injection presupposes a thing to inject.

A contemporary preprint approaches the same vocabulary from the opposite direction, arguing that tool boundaries do not confer epistemic warrant on what an agent takes *in* as observation (Romanchuk & Bondar, 2026). Where it asks what an agent is entitled to believe from its inputs, this paper asks what a producer is entitled to decide about its outputs; the shared words — warrant, epistemic — mark an adjacency, not an overlap. The intuition beneath that adjacency is not itself new: information-flow control has long held that a value’s classification is not laundered by crossing a boundary — a label follows the data through whatever module it passes (Denning, 1976; Sabelfeld & Myers, 2003) — and the same instinct animates data provenance. What the neighbouring preprint adds is the agent-era statement of it; what steadies the positioning here is the older observation, so the neighbourhood does not rest on a single recent source.

Two long-standing models sit across from the claim rather than beside it. Communicating sequential processes (Hoare, 1978) and the actor model (Hewitt et al., 1973) both name the recipient of a message explicitly — a channel, a mail address. That explicitness is precisely what the present argument identifies as the reintroduced coupling: to name the party at the other end is to take, in the sender’s voice, a decision the sender has no standing to make.

What remains, after all of these, is an unclaimed intersection. The mechanism of late binding is old; the authority framing exists for access; the coordination split exists for modularity; the effect inversion exists for semantics. But output read as an epistemic boundary — a producer entitled to what becomes observable and not to where it lands, the naming of a destination as an inference reached beyond warrant — is a description the literature has left open. The neighbouring terms are already taken: *observability* by telemetry, *output as a boundary* by information-flow security, **print**-versus-logging by engineering hygiene. The reading offered here fills a gap it did not have to invent.

8. Conclusion

This paper began at the most familiar instruction in programming — **print** — and found beneath it, where everyone thought the matter settled, a question no one was asking: who is entitled to decide where an output goes. Pulling on it opened not one hidden decision but three. What exists, what of it becomes observable, and where observation occurs are separate competences, held by separate authorities — the domain, the actor, and the assembler of Table 1.

The construct followed from the division. Wherever one authority is made to pronounce another’s decision, it must reach a conclusion it has no standing to reach — an assertion beyond its warrant — and such an assertion is a symptom in the sense §3 defined: it compensates for an authority that is absent, and dissolves the moment that authority is restored. The symptom proved to sit

at both ends of the output. The producer, denied a voice for the destination, stipulates a *where* it cannot know; the observer, denied the account, infers a *how* it was never told. One defect — absent information — met once by stipulation and once by inference.

What dissolves both is not a mechanism but a restoration of authority. Give the destination its own voice, and the producer stops asserting a where it cannot know. Deliver the record as what it is — the sequence the actor lived, an account it could genuinely have given — and the observer stops reconstructing how the state arose. It returns from being an engine of reconstruction to being, simply, a reader of what it was told.

None of this is a fact about **print**, or about any one system. The three authorities are a way of reading any architecture. The usual measure — how many layers, how clean the separation of code — misses what matters. The measure that holds is the one this paper has applied throughout: an architecture is sound not by its divisions of code but by whether the authorities within it match the decisions each part has the standing to make. That is a test a reader can carry back to a system already built and run against it — as Paper 6 offered a test for infrastructure, this offers one for the authorship of output.

Set against the series, the result is a mirror, and it has a spine of three words, one to a paper. Paper 4 showed that the actor *speaks*. This paper shows that it does not speak alone, and that what it speaks reaches the far end as *testimony* — received, where that word can be trusted, on the word of the one who lived it, not reconstructed by the one who did not. What testimony becomes when it is finally read — its *narrative* — the series has still to take up. Each of these makes one fundamental artifact explicit, and together they are the material of a synthesis this series is working toward, not the business of any single paper. Here it is enough to have shown that an output has authorities, that they can be kept apart, and that when they are, no one at either end of the voice is left to infer what someone else was in a position to say — within the trust domain where that voice can be believed, the observer, at the last, is simply told. Where it cannot, verification remains, and that frontier is the paper's edge, not its undoing.

Appendix A. Labs

The three authorities of §4 are exhibited here on the running framework. These labs are illustration, not measurement: each demonstrates that a separation the argument reached analytically is *buildable*, and that it holds mechanically against a *fused baseline* — the ordinary arrangement in which the two decisions are written together. None measures cost against that baseline at scale, or benefit in production; those are the questions §6 leaves open. In the manner of the substrate paper's labs, each result is a count the separation drives to zero — the zeros are the claim.

Lab A — the destination (the assembler’s authority). One actor records an order and projects it with a single `print` script; the projection is then bound, from outside the actor, to one destination after another — a real SQL Server table, then a real MySQL table — changing only the writer. Both backends receive identical rows, and the actor is neither recompiled nor reread. The wire format (TOON or JSON) and the direction (pull or push) are bound in the same outside place. A fused baseline — the projection and the sink written together — pays one producer edit for each new destination; the separated actor pays none. And the pull test makes the epistemic point empirically: the projected total is read back through a query with nothing pushed — a domain result exercised while no destination takes part, which is the operational proof that the sink was never in the domain (§4).

Lab B — the projection (the observer’s authority: what). From one journaled fact, three distinct observers are added — a fulfilment view, a finance view whose figure is *derived in the projection* (unit price $\times$ units), and a catalog view — each a projection reaction over the same fact, each authored without adding a single method to the domain. The domain’s surface is unchanged, confirmed by reflection; a fused baseline pays one domain method per view.

Lab C — the direction (the observer’s authority: how). The same fact reaches an observer by two routes: it is *told* — a reaction carries it across a `tell` to the observing role — or the observer *polls* it with a query. Both deliver the identical observation, and the producer’s domain holds no method for either direction: being told is a reaction, polling is a query, and neither is the domain’s to name.

Lab D — testability (the hard boundary is observable). The separation of §4 is not only argued; it can be tested for. A domain projection is exercised end to end — an order recorded, its total read back — with no destination bound at all: no sink, no port, no test double for output. That the test needs none is the operational proof that the sink was never in the domain; were it the producer’s, the test could not run without it. Against a hexagonal (ports-and-adapters) baseline the difference is a count: an output test of the ported domain must stand up at least one double — the port’s — while the separated domain stands up zero, because there is no port to double.

Lab	Separation	Separated cost	Fused baseline
A	destination bound outside the actor (real SQL Server, MySQL)	0 producer edits per destination	1 producer edit per destination
B	N views over one fact	0 new domain methods for 3 views	3 domain methods

Lab	Separation	Separated cost	Fused baseline
C	told vs polls	0 domain methods for the direction	no fused analog
D	domain output tested with no destination	0 test doubles for output	≥ 1 (hexagonal port)

The labs live in the framework’s test suite; the real-backend case runs against containers and self-excludes from the default run. They show the separation is not only arguable but built, and that its mechanical benefit against a fused baseline is a set of edits not made — leaving the economic question, whether those unmade edits are worth the apparatus, to §6 and to future work.

Code provenance

Source-code references in this paper resolve against the public Puppeteer repository at commit **0bf947b** (2026-07-22). The snapshot is archived in Software Heritage under the following persistent identifier:

```
swh:1:dir:4cce51c877c0836f5a561ce2d25bca9f7800ee71;
  origin=https://github.com/alvaroNCubo/puppeteer;
  anchor=swh:1:rev:0bf947bd6563e34cb141e3b5ba6cd13b4a811023
```

The three source files this paper names — `Puppeteer/IOutputSink.cs`, `Puppeteer/EventSourcing/Follower/Planes.cs`, and `PuppeteerCli/AttachCommand.cs` — resolve against this snapshot. A reader can construct a per-file SWHID by adding the qualifier `;path=<path>` to the directory SWHID above. Future commits to the repository may move these files; the SWHID preserves the cited state independently of any future change to the repository or its hosting.

The four laboratories’ source and datasets accompany this paper as **paper08-data.zip**. Each lab’s result table — the counts this paper cites in Appendix A — is included in full; the labs are count-based, so no large raw sample log is produced or omitted.

Acknowledgments

The author used large language models (including Claude and ChatGPT) as editorial assistants for language refinement, structural feedback, and literature navigation. All original ideas, terminology, theoretical constructs, and technical content presented in this work are solely the author’s.

Notes

¹ `Puppeteer/IOutputSink.cs` (interface `IOutputSink`, struct `PushDocument`). Its documentation comment states the separation directly: “*Pull vs push is a property of the destination, never of `print`. The DSL script is identical either way... the sink is assembly-agnostic: it is set via `OutputTarget(sink)`... because the mechanism lives at the `ActorHandler`, not in any one topology.*” The four output planes an actor’s reaction can address — `Program.Emit` (ephemeral projection), `Causation.Continue` (`tell`), `Outbox` (durable, exactly-once-recorded), `Metadata` (elide / materialize) — are defined in `Puppeteer/EventSourcing/Follower/Planes.cs`. The console binding referenced in §4 is `puppeteer attach`, whose REPL runs the actor with its output on `Console.Out` (`PuppeteerCli/AttachCommand.cs`), pipeable into ordinary shell tools.

References

- Denning, D. E. (1976). A lattice model of secure information flow. *Communications of the ACM*, 19(5), 236–243. <https://doi.org/10.1145/360051.360056>
- Gelernter, D., & Carriero, N. (1992). Coordination languages and their significance. *Communications of the ACM*, 35(2), 96–107. <https://doi.org/10.1145/129630.129635>
- Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.
- Hewitt, C., Bishop, P., & Steiger, R. (1973). A universal modular ACTOR formalism for artificial intelligence. *Proceedings of the 3rd International Joint Conference on Artificial Intelligence (IJCAI)*, 235–245.
- Hoare, C. A. R. (1978). Communicating sequential processes. *Communications of the ACM*, 21(8), 666–677. <https://doi.org/10.1145/359576.359585>
- Kahn, G. (1974). The semantics of a simple language for parallel programming. In J. L. Rosenfeld (Ed.), *Information Processing 74: Proceedings of the IFIP Congress* (pp. 471–475). North-Holland.
- Martin, R. C. (2017). *Clean architecture: A craftsman’s guide to software structure and design*. Prentice Hall.
- Miller, M. S. (2006). *Robust composition: Towards a unified approach to access control and concurrency control* [Doctoral dissertation, Johns Hopkins University].
- Morrison, J. P. (1994). *Flow-based programming: A new approach to application development*. Van Nostrand Reinhold.
- Pike, R., Presotto, D., Thompson, K., Trickey, H., & Winterbottom, P. (1992).

The use of name spaces in Plan 9. *Proceedings of the 5th ACM SIGOPS European Workshop*, 72–76.

Plotkin, G., & Pretnar, M. (2009). Handlers of algebraic effects. In G. Castagna (Ed.), *Programming Languages and Systems (ESOP 2009)*, Lecture Notes in Computer Science (Vol. 5502, pp. 80–94). Springer. https://doi.org/10.1007/978-3-642-00590-9_7

Ritchie, D. M., & Thompson, K. (1974). The UNIX time-sharing system. *Communications of the ACM*, 17(7), 365–375. <https://doi.org/10.1145/361011.361061>

Rivera, A. (2026d). Preserving semantic continuity across actors: a tell-based approach without orchestration. *Puppeteer Papers Series*, Paper 4 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21207062>

Rivera, A. (2026e). The journal as substrate: unifying deployment, replication, backup, and offline operation in distributed systems. *Puppeteer Papers Series*, Paper 5 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21349146>

Rivera, A. (2026f). Most infrastructure layers are symptoms of the persistence model: a construct for auditing production stacks. *Puppeteer Papers Series*, Paper 6 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20317450>

Romanchuk, O., & Bondar, R. (2026). Semantic laundering in AI agent architectures: Why tool boundaries do not confer epistemic warrant. *arXiv preprint arXiv:2601.08333*.

Sabelfeld, A., & Myers, A. C. (2003). Language-based information-flow security. *IEEE Journal on Selected Areas in Communications*, 21(1), 5–19. <https://doi.org/10.1109/JSAC.2002.806121>

Identity Precedes Staging: one play, many stages

TL;DR

A domain and its deployment are usually written together, so a new place to run means editing the domain. They can be separated completely, and this paper measures one that is. A Tetris board was held fixed while five stagings and six clients changed around it: zero domain edits, zero test doubles. The domain declares no interface for what drives it or for what it emits, so the contract that would be a port is provided by the substrate and bound in the staging. None of that geometry is new. What the paper adds is a measurement of it, against a ports-and-adapters version of the same game built for the comparison — which matches the zero per staging. What it does not match is of another kind: ports and adapters *decouples* a domain, while this arrangement *closes* one. Two measurements: a ported domain declares obligations outward — three driven ports, and with them three stand-ins without which twenty of sixty-four tests do not run, an operation per port on its public surface, and a state it will accept from outside — and it needs a reconstitution surface of its own, costing 56 lines. This one declares none and needs none, replay re-entering by the operations the acts were performed through. Calling those two *closure* is a reading of them rather than a third count. Two findings arrived unsought, neither of them this paper's result and both bounding how it may be read: an incomplete record answers plausibly rather than failing loudly, so carry the recorded state and assert a replay against it; and a fact a domain emits must be derivable inside one actor, which bounds how a domain may be cut.

Dependencies. This paper is part of the *Puppeteer Papers*, a series of self-deposited preprints, and rests on four of them: the actor's speech and **tell** (Paper 4), **Reaction** read as the recognition of a routine (Paper 3), the server read as an accidental category, ephemeral after bootstrap (Paper 7), and testimony (Paper 8), whose observer receives an account it is told. Paper 8 noted, without pursuing it, that a narration received is not yet a narrative recognized; this paper takes up the recognition, and reaches it not as its subject but as a

consequence of the identity it argues. Methodologically it is an analytic theory contribution in the sense of Gregor's (2006) theory for analyzing (Type I): identity and staging are constructs by which an architecture may be described and compared, and the paper offers no prescription for building one — its instrument is a question that can be asked of a system already built (§10), not a rule for making one. The labs of Appendix A are accordingly an existence proof of realizability, not a design-science evaluation of cost or benefit, and where the text has slipped into evaluative language the slip is corrected rather than defended. Stated in the vocabulary of software-engineering research method rather than of information systems: the contribution is a descriptive model whose validation is by example — a system built and measured — in Shaw's (2003) classification of contribution and validation types. And in the terms of Stol and Fitzgerald's (2018) framework, it trades generalizability over actors and realism of context away entirely, being one domain on one framework, in exchange for a precise and reproducible measurement of a single quantity.

1. The Domain and the Deployment

In most systems, the code that defines what a system does is not cleanly separable from the code that defines where it runs and which clients it serves. Moving a domain from a console to a web server means editing it; adding a second kind of client means adding to the thing the client observes; splitting one process into several is treated, uncontroversially, as building a different system. Where a system runs, and whom it serves, are treated as properties of the system itself — so that a monolith and its distributed successor are counted as two systems rather than one system deployed two ways.

There is a field in which this separation is routine rather than aspirational: theatre. A play is written once and staged in many venues, for many audiences, without being rewritten, and no one counts the Broadway production and the school-gymnasium production as different plays. The term is taken from there: throughout this paper a *staging* is a deployment — a particular place a domain runs and a particular client it serves — and the domain is what is staged. Two units are counted in what follows and they are not the same: a **staging** is a deployment configuration, of which the example has five, and a **host** is an executable project, of which it has twelve, eleven reaching the domain through the actor and one writing bytes to a pipe and reaching nothing. Counts of hosts are counts of projects that had to compile and run; counts of stagings are counts of arrangements the domain was placed in. It is a naming convenience, not an argument, and earns its keep only if the separation it names can be made real and measured — which is the rest of the paper.

The claim is not that the domain is *decoupled* from its deployment, which is ordinary advice. It is that the domain is **closed**: it declares no obligation for a deployment to meet, so the dependency between the two runs one way only, and that direction is a checkable property of the built system. The domain compiles

and runs with no staging bound to it. Every staging refers to the domain; the domain refers to no staging. In that precise sense the domain comes first: a staging is constructed against a domain that already exists and does not know of it. This is all the title means by *precedes* — a direction of dependence in the build, not a claim about time or metaphysics — and §2 measures it directly.

This is not dependency inversion under another name, though the difference takes some care to state and §9 states it. Ports-and-adapters architecture keeps adapters outside the domain, and for whatever a domain needs *from* the world it has the domain declare an interface and depend on it. In the model examined here the domain declares no interface at all — not for what drives it, nor for what it emits. It is entered by calls on its own operations, and it emits facts under logical names, while where those facts go, to whom, and what drives them are bound entirely outside it. Whether that is *cheaper* than the ported alternative is a separate question, and mostly it is not: a ported version of this game was built and counted, and a new transport costs its domain nothing either (§9, Appendix A). What the ported design still has, and this one does not, is an open obligation — an interface it declared and depends on someone to supply. That difference, not a cost, is what §9 measures.

Placed in the series, this is a question of a specific kind. Paper 4 asked who may speak; Paper 8 asked who may decide what an output is; the series puts its questions in that form. This paper asks a prior one — not who is entitled, but what an entitlement is attributed to, since a subject that does not hold still across its stagings cannot bear one at all. It is a question about the identity of a domain, not about infrastructure or topology. A later paper in the series will need that question answered before it can treat a set of domains as a reusable repertoire; that is not this paper's concern, which is the narrower, checkable claim that a domain's identity is prior to and independent of its staging. One consequence is developed in §6 and §7: when a domain's acts are recorded as they are performed, the sequence of what it did is available whole from its journal — read from that record, wherever a copy of it is kept, rather than assembled out of fragments held separately.

One question follows, and a second that the first makes unavoidable. The first (§2, §5): does anything actually hold a domain fixed while its staging changes — across both where it runs and which client it serves — can that be measured rather than asserted, and what does the invariance reveal? The second (§3–§4, §6) is not a parallel contribution but what the first needs in order to be legible at all: since a client never reads a domain directly, through what does it read, and does it read the domain's present or its past? An invariance nobody can observe is not a result, so the reading has to be accounted for; that is the whole of the second question's standing here, and §6 says where it stops.

One thing about the order of the work belongs here too, because it decides how the rest is to be read. **The arrangement preceded its account; the experiments were built to expose and bound it.** The domain, the actor and most of the stagings of Experiment A existed before *identity*, *staging* and

closure were words for anything here, and none of them was built to satisfy a construct. The labs are the other half: the ported baseline, the three containers, and the growth of §8.2 were made for this argument, and §8.3 says plainly where that weakens what they show. So the constructs are instruments for reading a construction rather than a design derived from a theory, and the question they are answerable to is whether they describe faithfully something already built — which is the sense in which the labs are validation by example and not a test of a hypothesis.

2. Two Experiments

The claim of §1 is checkable, and this section checks it against a running example: a Tetris game, whose board and rules are a small but complete domain, written as a **Well** and a set of operations over it. The domain is held fixed, and two things are varied around it in turn — where it runs, and which client it serves. The measurement in each case is the same: the number of changes the variation forces on the domain source. The evidence is the git history of the example repository, and the anchors are collected in Appendix A.

The domain’s independence has a concrete form before anything runs — the shape of the build. Every executable host in the example reaches the domain through a single actor, and the domain reaches back to nothing but the language’s own library:

```
11 hosts          -> TetrisActor
    TetrisAi, TetrisConsole, TetrisObserver, TetrisServer, TetrisStage,
    TetrisStageDuo, TetrisStageDuoTls, TetrisStageServer, TetrisWatch,
    TetrisWeb, TetrisWebRest

TetrisActor       -> TetrisDomain, Puppeteer, Choreography
TetrisDomain.Tests -> TetrisDomain

TetrisDomain      -> (none)
```

Figure 1. Every declared dependency between the example’s projects, read out of the project files on its main branch rather than drawn, and grouped so that the direction is visible: an arrow means *this project is built against that one*. Every staging depends on the same domain; the domain depends on none. Fourteen of the fifteen projects appear — **TetrisSend**, the pipe writer, is omitted because it references nothing and is referenced by nothing, being a separate executable that only writes bytes. Read downward, every arrow leads toward **TetrisDomain** and none leads away from it.

The figure answers one question, *who depends on whom*, and it is worth saying what it does not answer. It does not show that the domain leaves nothing open — no port declared, no contract, no persistence or transport requested. That is a different property, it is the stronger of the two, and it is measured separately

against a built alternative in §9. A reader unfamiliar with this language’s build files needs only the arrow: a project that names another cannot be compiled without it, and a project that names none can be compiled alone.

Declaring no reference is not the same as naming nothing, and one qualification belongs with the figure rather than after it. Three hosts do name a domain type, reaching it through the actor’s reference, which in this language is transitive: each writes `typeof(TetrisDomain).Assembly` to hand the assembly to the framework (`sm-duo-tls/Program.cs:33`, and the same line in `sm-server` and `sm-duo`). What they name is an empty public anchor and nothing else, so a staging holds a name and not a handle. Lab E reads that fence at the level of members and reports its two authored exceptions.

What the domain does instead is *satisfy conventions*, which is a different relation from declaring a dependency and the difference is load-bearing. The framework finds the domain by reflection, so its operations must be ordinary methods with coercible parameters, and the empty public anchor is there so a host can name the assembly in its own source — a convenience, not a requirement, as Lab F reports. What separates conformance from dependency is not interpretive: the domain compiles and its own tests run with the framework absent from the build graph entirely (Lab E). A dependency cannot be absent from the build and still be met; a convention can.

That direction is what §1 called *precede*, and one note on vocabulary avoids confusion later. With arrows running from a dependent to what it depends on, the domain is a *sink*: every staging’s references lead to it and none leads out. Reverse the arrows and it is the *root*, every staging reachable from it and none from another. Where later sections call the domain the stagings’ common ancestor, that reversed orientation is what is meant; §5 says what the property does and does not establish.

Experiment A — change the stage. The same `Well` is run on a console; in a browser, with input and output carried over a real WebSocket connection, and in a second browser variant over HTTP with server-sent events; across two co-hosted actors coordinated by a `StageManager`, once in memory and once over a real TLS channel; and across three separate Docker containers, each a peer joined to the others over real container-to-container TLS. These are genuinely different deployments — different processes, machines, transports, and wire formats. The domain source is identical across all of them: a diff of the domain directory between the first staging and the last comes back empty (Appendix A). What changes from one staging to the next is host code — the shell that binds an input source and an output sink to the actor. The example’s own name for that shell is exact: the host is “an accidental shell” (`actor/IGameHost.cs:14`), and the same `Well` runs on `PerformanceV2` or on `StageV2` without knowing which. The example builds clean and its domain tests pass, unchanged, in every staging.

A distributed demonstration invites more credit than it earns, so its boundaries

belong in plain sight. Three containers run the same `Well` as three `StageManager` peers over genuine container-to-container TLS; a scripted driver on the Director plays the game, and the two casts converge to a byte-identical board. Adding that staging left the diff of both the domain and the actor directories empty. What is *not* claimed is any test of resilience: no peer was killed and no partition induced, so partial failure is untouched (§9). The setup, the bootstrap, and what each node is handed are in Lab B.

Experiment B — change the client. With the stage held at one host, the client is varied instead. The same board is played by a human at a keyboard reading an on-screen grid; by an automated player that sends moves over a pipe and reads the board through a view it computes for itself; by two passive observers, one that pulls the board and one that receives it pushed; and by a browser, in which a player and a spectator are both written in JavaScript and reach the game over the network. Six clients over one `Well` — some in C#, one a PowerShell script, two in a browser — and again the domain directory does not change across any of them (Appendix A). Six is the number of additions; counted by how differently the board is actually represented, the labs hold two, a grid and a vector of column heights, and Appendix A reports both numbers rather than the flattering one alone. What each client adds is an input adapter, an output adapter, or both; none of them is a change to the game.

The two experiments vary independent things. Experiment A changes where the domain runs; Experiment B changes who observes it and how. Neither touches the domain, and that independence is the result: its constancy does not depend on where it is staged or on who is watching. The identity §1 argued to be prior to the staging appears here as an invariant the two variations cannot move — and the build graph says why it cannot be moved: nothing the domain refers to changed, because it refers to no staging and to no framework in the first place.

How much these counts separate this arrangement from a competent ports-and-adapters one is a question §9 answers with a baseline that was built and counted, and the answer is less than the counts alone suggest: a hexagon’s rule model takes a new transport at no cost either. What the counts establish here is the invariance itself, not a saving over an alternative.

3. What a Client Sees

Experiment B added clients by adding adapters, and the word carries the second half of the argument, so it is worth making precise. A client never reads the domain directly. It reads through an adapter, and the domain is the same behind every one of them.

What the domain emits is not a rendering — but calling it a *raw* fact would overstate the case, and the difference needs a criterion rather than an adjective. What is emitted is the set of occupied interior cells, and that set is computed rather than stored: the domain unions the settled pile’s cells with the falling

piece's and clips the result to the frame's interior (`domain/Well.cs:322-331`). The pile itself holds a set of positions (`domain/Pile.cs:26`). So what leaves the domain is a derivation the domain performs, not a field read off storage — and the domain project, as §2 established, declares no dependency on the framework or on any staging: it has no notion of a screen, a renderer, or a client. The substrate pushes that set outward on each move (`actor/TetrisActor.cs:43-47`), forwarding it rather than projecting it further.

Derivation, then, cannot be the line, since the automated player's column profile is a derivation too. The criterion this paper uses is the one the series applies elsewhere, and it is a single question a reader can put to any domain's source: **does a rule of the domain depend on the notion?** If some operation is stated in those terms, the notion is the domain's own, and what is emitted in them is a fact it holds; if no operation mentions it and only a consumer needs it, it is a view computed for a reader. Applied here the two separate cleanly. Occupancy by position is the vocabulary the rules are written in — collision is a membership test over cells, and what counts as a complete row is counted in them — so the emitted set is in the domain's own terms. Column heights, the gaps between them, their bumpiness appear in no operation of the `Well`; nothing in the game is decided by any of them, and only the automated player wants them. That is why the one is emitted and the other computed outside.

What the criterion buys is that the line is inspectable rather than asserted: a third party can read the domain and check which notions its rules turn on. What it does not buy is a line that is sharp or permanent. Which notions a domain owns is a modelling judgment — Paper 8 called this the negotiable boundary of the two it drew — and a `Well` that grew an operation deciding something by column height would move it, legitimately. The claim here is not that the boundary is objective, only that it is auditable, and that in this domain the audit comes out as this section describes.

Each client turns the fact into something it can use, and the turning is the client's work, not the domain's. The human's on-screen grid is produced by a renderer that lays the bitmap out as squares (`actor/BoardRenderer.cs:19-48`). The browser produces the same grid a second time, independently, in JavaScript — its code is commented as mirroring that renderer (`web/Program.cs:169-178`) — so the terminal grid and the browser grid are two adapters over one frame. The automated player needs something else: it reasons poorly over a bitmap, so it lifts the same frame into a column-height profile — a skyline, the gaps and wells between columns, a few aggregate figures (`tools/pile-scan.ps1:54-83`) — a form suited to deciding where a piece should go. None of these is the board's own appearance; nothing in the domain decides anything by a picture. It emits a set of cells in its own terms, and each client projects that through an adapter chosen by the host that runs it.

One fact about how that adapter came to exist belongs in the argument rather than only in the disclosure at the end of this paper, because it is the sharpest evidence available for the claim of this section. The automated player is an

instance of a large language model, and the column-height view it reads the board through was written *by that instance* during the lab, with the author's permission, to suit the form in which it reasons (`tools/pile-scan.ps1`). Nobody translated the board for it. It was handed the same emitted fact every other client gets, and it authored its own projection of that fact into the shape it could use. A projection belonging to the party that reads it is, in this one case, not an interpretation of the architecture but an observed event: the reader was a different agent from the domain's author, and it built the adapter itself. How that came about, and what it means for weighing the evidence, is disclosed in the acknowledgments.

This is where the on-screen grid loses its privilege. It is natural to treat the pixels a human sees as the real board and the automated player's vector as a derived abstraction of it. The system does not: both are adapters over the same emitted fact, and neither is closer to the domain than the other. The grid feels primary only because it is the one a person happens to read; the domain has no such preference, because it produces neither — it produces the fact both are made from. The automated player makes this hard to miss precisely because its adapter looks nothing like a human's. Faced with a client that reads the board as a vector of column heights, one notices that a human, too, only ever reads it through a rendering, and had simply stopped noticing.

This arrangement has a much older statement in the database literature, and the paper is better for saying so. Logical data independence — that an application should be insulated from the structure in which data is held, and served instead by views derived from it — was part of the relational model from the beginning (Codd, 1970), and the three-schema architectures that followed made the separation of a stored schema from per-application views explicit. “The domain emits facts and each client projects them” is that separation, reached from the side of the domain. What differs is who owns the projection. Under data independence the mediating system owns the schema and derives the views: the projection is the database's work, performed on the application's behalf. Here there is no mediating schema authority at all — the domain emits its own facts in its own vocabulary, and each client's projection is authored outside both the domain and any central schema, by the staging that binds it. The insulation is the same; the party doing the insulating is not.

This connects directly to the boundary Paper 8 drew. There, what an output *is* — its projection — was shown to belong to the actor that authors it, not to the domain that holds the material; a domain object that renders itself reaches past what it knows into a decision it does not own. Here the same boundary is seen from the client's side: the projection is authored per client, outside the domain, and the domain that would have rendered itself instead emits a fact and lets each client author its own view. The adapters are the input sources and output sinks of the earlier papers — a keyboard or a pipe on the input side; a terminal grid, a numeric vector, or an HTML canvas on the output side — bound to the actor from outside and exchanged without the domain's knowledge. Whether a

client is served by pulling or by being pushed to is a choice of the same kind, made in the same place, and the two modes have long been catalogued in the publish/subscribe literature (Eugster, Felber, Guerraoui, & Kermarrec, 2003). That the domain does not know which adapter is attached is the same fact, seen once more, that let the stage and the client vary in §2 without moving it.

4. Actor and Audience

Section 3 answered *through what* a client reads the domain: an adapter, never the raw domain, and no adapter privileged. A second question sits beside it — what, in time, a client reads — and the answer is narrower than the temporal reading of it invites. Acting and watching are distinct roles and what each may have differs; that is a separation of roles, not a law about time. The distinction matters, because the temporal version of the claim is either trivial or false.

Take the obvious shortcut first, since it is the one that has to be dismissed with care rather than in passing. Each mutating command could return the board it produced, and the client that issued the move would have its picture at once. There is nothing incoherent in that. It is *read-your-writes*, one of the named session guarantees (Terry et al., 1994); it is usually what a commanding client wants; and it is often chosen for good reasons of latency. What such a response cannot do is serve a client that issued no command, and that limit is close to definitional — no request, no response to be served from. So what follows is a design corollary rather than a logical necessity: a system whose only path to a view is the command’s response has no path at all for an observer that does not command, and every client in §2 and §3 that was not driving the game is exactly such an observer.

This is not a new mechanism, and — the point deserves to be conceded rather than discovered — the temporal reading of it is not new either. The field already separates the two operations: command-query separation, and command-query responsibility segregation after it, hold that a view is not served from the write (Meyer, 1988; Young, 2010), and event sourcing builds a view as a projection over recorded acts rather than a snapshot of current state (Fowler, 2005). And that what an external consumer holds is the past was argued two decades ago in the vocabulary of data: data outside a service lies beyond the reach of that service’s transactions and is *temporally disconnected* from it, so what has left is what was and not what is (Helland, 2005). Nor is the structure beneath that new: in distributed systems it is the causal past. Lamport (1978) showed that the events available to an observer form a partial order — what can have reached it is what happened before it — so an observer’s knowledge is constitutively of the past, and a “present” shared across separated parties is not a thing that can be observed at all. This section therefore does not claim the past-tense view as a finding; it sits inside a structure distributed systems have had since 1978, and the data literature stated for consumers two decades ago.

Nor is the present unavailable to an audience, and this paper’s own labs say so: one of the six clients polls the board and is handed the current state (Lab C). An audience can see the present. What differs by role is narrower — a live read taken *in order to act* belongs to the party acting. The game’s host queries the current board, is a piece falling, so as to choose its next move, and that reading is part of the act; an observer polling the same board is not deciding anything, and is being shown a configuration. The difference is one of standing, not of timing, and in that form the claim is not this paper’s to make: Paper 4 gave the actor its voice, and Paper 8 established who is entitled to say what an output is. This section applies their result to the seats in a theatre rather than extending it.

Naming the two apart is where this paper needs the one before it. An observer given only the present is given a snapshot — the configuration as it stands — from which, to learn how that configuration arose, it must reconstruct the history; that reconstruction is the defect Paper 8 traced. An observer given the past is given the account itself, the sequence of acts the domain performed, which Paper 8 named *testimony*: knowledge held on the word of the one that lived it, not rebuilt from stills. The record-reading view of this paper is therefore not merely consistent with Paper 8 — it is Paper 8’s resolution made operational. To see a domain by reading its record is to receive its narrative as testimony, and it is the only way an audience comes to know *how* the state arose. What polling supplies is the configuration; what the record supplies is the account.

This returns to what the paper is about. Because what an audience is *told* is the recorded narrative rather than a state read off the moment, the thing told is a standalone object: it exists in the record, the same for every audience, independent of who reads it and when. And it is the same for every audience for the same reason the board was — one domain, referenced by every staging, produced it. The invariant §5 names — the domain as the common ancestor of its stagings — is what makes the narrative one narrative, readable the same from every seat. What remains the same when the staging changes is, in the end, what there was to see.

5. Identity

The two experiments leave a single structural property to state, and it is statable without interpretation: a domain was run on five stagings and read by six clients — between them, in two genuinely different representations — and through all of it the domain referred to no staging, every staging referred to it, and its content did not change. In the dependency graph the domain is a sink; in the reversed graph it is the root from which every staging is reachable, which is what §2 meant in calling it their common ancestor.

What that property licenses one to say is a second question, and the word this paper uses for it invites more metaphysics than the evidence supports. *Identity*

is offered as a reading of the property above and not as a further finding: a thing that holds still under that much variation is the sort of thing an entitlement can be attributed to, which is what the rest of the series needs of it.

Two cautions follow, and they mark exactly where measurement stops. The first concerns the word *position*. A position in a graph individuates nothing — a different domain, or this one rewritten, would occupy the same one — so the property measured here cannot be what makes this domain the domain it is; nor is it even the graph's only sink, since the language's own library is another. The second is that the property is evidence about a *relation* and not a criterion of identity: it says that what the stagings were built against is one thing, and that its content owes them nothing. Calling that an identity is a reading of the evidence rather than a part of it, and the two are better kept apart than run together, because the measurement is checkable and the reading is not. What this paper measures is independence and priority. *Identity* is the word for what independence and priority license one to speak of, and the title should be read in that light.

One echo in that title is worth disowning before a reader supplies it. *Identity precedes staging* has the shape of Sartre's *existence precedes essence*, and inverts it — an essence prior to any concrete appearance is close to what that formula was written to deny. The resemblance is accidental and the claim is not the philosophical one. Nothing here is an essence: what precedes a staging is a position in a dependency graph and a source directory that did not change, both of which can be read off a built artifact by someone who has never heard of the phrase. A reader who hears the echo should hear it as coincidence of form, not as an argument being smuggled.

Read that way, *precede* is not a claim about time or ontology but the direction of the dependence. A staging is built against a domain that already exists and does not know of it — the way a library is compiled before, and in ignorance of, the programs that will call it. The domain does not *survive* its stagings; surviving would make the stagings the primary thing and grant the domain the modest virtue of enduring them. The relation runs the other way: the stagings were built because the domain was there to be built against. Nothing the domain refers to changed across the experiments because it refers to no staging at all — so its content was never at risk from a change downstream of it.

This has a second face, less formal than the graph and worth stating because it is what makes the identity legible to a person rather than only to a build tool. A domain with an identity is one whose parts can be understood as stable roles. To know what the board of the game *is* — that it holds a pile of settled cells, spawns and moves a falling piece, clears full rows, ends when a new piece cannot enter — is to know, structurally, what it does and does not do, what may be asked of it, and what it will never answer. This is not a psychological claim about how people understand stories; it is the observation that a role with a fixed identity fixes its own boundary, and that boundary is what domain-driven design reached for with its aggregates and bounded contexts — arrived at here

by asking what stays the same, rather than by drawing a diagram of layers. The invariance measured in §2 and the legibility described here are the same thing seen twice: what holds still across every staging is what there was to understand in the first place.

6. Recognition

If a domain keeps its identity across every staging, then whatever it did keeps its identity too. The sequence of a domain's own acts — a piece spawned, moved, dropped; a row cleared — is recorded in its journal as the domain performed them, and because the domain is the same on every stage, that record reads the same from every one. This is the first, plain consequence of §5: the account of what happened is available directly from the record the acts were written to, whole wherever that record is kept, and does not have to be reassembled from fragments held elsewhere.

That availability is what lets a client do more than watch. Given the sequence of acts, a reader can recognize in it the routine the acts compose — that these three moves and a drop were the placement of one piece, that this run of placements filled a row. The term needs pinning down, since the rest of the section leans on it, and it is not this paper's to define: a *routine*, in the sense Paper 3 gave it, is a pattern over journal entries together with the correlation that binds them into one trajectory — a reaction seeks an opening entry and then a closing one, and what it matches, between them, is the routine. So recognition is an operation with a definition and an implementation, not a figure of speech, and it reads the record as it stands rather than inferring a history from snapshots.

Recognition in this paper is that operation applied to a domain whose identity holds: because the acts are the same across stagings, the routine they compose is recognizable across stagings too. That step has been measured (Appendix A, Lab H). One reaction, seeking a spawn and then the drop that ends it, recognizes the same two placements in the same order over three different stagings: the journal of a single process, the journal a warm server keeps while driven over a named pipe, and the journals of the three containers — read inside each node, against its own store. Within that third staging the three nodes' records are byte-identical. *Across* the three stagings what matches is the sequence of acts — same count, same order, same shape, verb for verb — while the entry identifiers differ by a constant, the cluster writing three idempotent seeding acts where one process writes one. The domain's diff for the whole exercise is empty. So the claim is a result rather than an inference, on a short trajectory: two placements, not two hundred.

That pair of results is worth reading together, because what moved and what did not are not the same kind of thing. The identifiers moved: the record's bookkeeping is a fact about where the domain ran. The acts did not: a spawn on a console and a spawn inside a container are one operation with the same

parameters, and each leaves behind an entry of the same kind. So what held still across the three stagings was not the record's representation, which shifted, but what the act was and what performing it left behind. A staging changes who is present when a domain acts and changes nothing about what the acting does. That is the sense in which §5's identity reaches past the source text — the invariance is not only in what the domain is written as but in what its verbs accomplish when performed — and it is Paper 4's account of an actor's speech that makes this a claim rather than another way of stating the empty diff.

Three things the lab found are more interesting than the confirmation, and §6 is narrower for them. The first is that this record offers no handle to correlate on: a spawn names a piece *type* and a drop names nothing at all, so the trajectory between them is tied by order alone and not by any identity the acts carry. The second is worse, and it answers part of the question this section leaves standing — whether a reading can be wrong. It can, and silently. Landing is not an act here: the board lands a piece inside a tick or a drop, privately, so a piece that comes to rest under gravity leaves its opening unclosed, and the *next* piece's drop closes it. The count comes out right and the correlation comes out wrong, with nothing to signal the difference. The third is that entry identifiers are not invariant across stagings — the three-node staging writes three seeding acts where one process writes one, so everything after shifts — and the closing identifier is the only handle a match hands its reader. The routine recognized is the same; the name by which a reader could refer to *this instance* of it is not.

The second of those generalizes beyond this lab, and it is the same shape as the finding of Lab G. A routine is recognizable only in the moments of it that were *acts*. Where a moment was a private consequence — a landing folded into a tick, a pile absorbing cells — the record holds no entry, and no pattern can seek what was never recorded. Recognition is bounded not by the reader's ingenuity but by what the domain performed out loud.

That bound exposes a gap in this paper, and it is better named here than left for a reader to find. Section 3 offered a criterion for the *vocabulary* of a record — does a rule of the domain depend on the notion? — which is auditable against the domain's own operations. There is no companion criterion for its *granularity*: nothing in this paper says which moments must be acts. Yet granularity is what recognition rests on, so the fidelity of a record as an *account* depends on a decision the paper leaves unguided while the fidelity of its *vocabulary* has a test. And the failure mode is the worst kind, the one Lab H found: an absent act yields not an error but a smaller coherent story, matched to the wrong closing entry, with the count coming out right.

No criterion is offered, and the honest statement is that this paper does not have one. What its labs supply is three cases any candidate would have to accommodate, reported as three cases and not as a rule. Absorbing a landed piece had to become an act when the board was cut in two, because two parties had to speak and the pile's own record needed an act of its own (§8). A landing had to be an act for a reader that *reads* the record, since Lab H's reaction could

not seek what was never written. And a score and a difficulty level did not have to be acts at all, being folds over acts already recorded (Lab I). Together they point at what a *reader* of a record needs to identify, as against what a *re-performer* of it needs — but three cases are not a criterion, and offering the hint as one would be the overreach §5 declines elsewhere.

It is worth locating the gap, since it is not peculiar to this arrangement. Choosing which moments a model records is a modelling decision of the same kind as drawing an aggregate’s boundary, for which the domain-driven design literature offers heuristics and worked judgment rather than a test (Evans, 2003). What this arrangement changes is only when the mistake surfaces: a boundary drawn badly resists the code being written against it, whereas a moment that should have been an act surfaces much later, in a reader that recognizes the wrong thing and reports it to no one. That is an open problem this paper raises and does not close, and it is worth being exact about what kind of question it is. Which moments must be acts is not a question about representation — about how faithfully a record renders what happened — but about what the domain is *doing* when it performs one. This paper has no account of that.

One consequence of it belongs beside §8.2 rather than here alone, because it is the same cost seen from the other end. Growth in this arrangement is additive when the new thing is a value or a rule over acts already recorded — that is what Lab I measured. Discovering that a *moment* should have been an act is the other case, and it is the expensive one, because every record already written lacks that act and no projection over those records can supply it. What the domain would have to do is not add an operation but reinterpret a history that does not contain the distinction, which is exactly the schema-evolution cost §8.2 concedes to the event-sourcing literature (Overeem, Spoor, & Jansen, 2017; Overeem, Spoor, Jansen, & Brinkkemper, 2021). So the missing criterion is not only an intellectual gap: it is the decision whose mistakes are dearest to correct, and that is the strongest reason to want one.

This is where the paper meets the one before it. Paper 8 followed an output to the party that receives it and showed that, within a trust boundary, the receiver comes to know the state as *testimony* — an account it is told rather than one it reconstructs. But an account received is not yet a routine recognized: being told the sequence of acts is prior to seeing, in that sequence, what was done. This paper takes the next step, and reaches it as a consequence rather than a new claim — an account whose subject keeps its identity can be recognized as the same routine wherever it is told.

And the recognition is the client’s, performed through the same adapter §3 described. The board does not announce “a row was cleared”; it records the acts, and each client recognizes the routine through the view it holds. The automated player reads its column-height vector and recognizes a well being filled; a person reads the grid and recognizes the same thing; neither recognition is the domain’s, and neither is privileged over the other. What both recognize is one routine, because beneath both adapters is one domain with one identity

— which is where the paper began. That no recognition is privileged does not settle whether recognition is itself governed: who is entitled to read a routine, and whether a reading of one can be wrong, is a question this paper reaches and leaves standing.

7. Where the Narrative Lives

Section 6’s availability has a counterpart, and it wants naming carefully. When what a system does is spread across autonomous services, each with its own store, the sequence of what happened is in none of them. It has to be assembled afterwards, by correlating records that were produced separately — which is part of what distributed tracing and correlation identifiers are for. That work has a developed literature: Google’s Dapper (Sigelman et al., 2010, a technical report rather than a refereed paper), and the refereed line that followed it (Mace, Roelke, & Fonseca, 2015; Sambasivan et al., 2016; Kaldor et al., 2017). None of that is treated here as a deficiency, and nothing in this paper says a system should do without it. The point concerns *completeness at a reader*: whether some single record holds the whole account, or whether the whole exists nowhere and must be joined together out of fragments.

Put that way, the section survives its own Lab B, which would otherwise contradict it — three machines keep three journals, converging by replication and by a catch-up path, so the narrative is in three places and reading it does depend on a mechanism. What makes that unlike correlation is not that there is one copy but that the three are copies *of one record*: each holds the same acts, each alone holds the whole game, and a reader at any node joins nothing to anything. So the honest form of the claim is *one logical narrative, replicated, with convergence the framework guarantees* — weaker than “one place”, and what the labs support. It is scoped, too, to one actor’s own first-person account (Paper 4): a system of many actors has many accounts, and relating them is a question this paper does not take up (§8). The model does not remove the need to know what a system did; it changes where that knowledge already sits.

8. Limits, and Two Observations Nobody Was Looking For

The claim is bounded, and the boundaries are worth drawing, the more so because the result is easy to overstate into a promise it does not make.

8.1 What the claim covers

It holds only where there is something for it to hold of. A domain with a genuine identity, staged more than one way, is where the invariance has teeth; a program too small to have a domain distinct from its deployment, or one that will only ever run one way, has no second staging to be invariant across, and

the apparatus is overhead. The claim is not that every program’s domain is separable from its deployment in practice, but that where the two are genuinely distinct, the separation is real, buildable, and measurable — and that collapsing them, where they were distinct, is what the ordinary practice does.

A second limit, and the more consequential one, follows from what the demonstrated domain is. The **Well** is a pure computational core: it holds rules and state, and its outputs are data. It keeps no external store, provokes no effect in the world, reads no clock, authorizes nothing, and depends on no other party — it never has to *cause* anything. That is exactly the case in which “the domain emits facts and each client projects them” suffices, because every consumer really is a client: it wants to know, and nothing turns on whether it ever does.

What the example does not exercise is a domain that must bring about an external effect and then reason about the result — charge a card, notify a partner, settle a payment. There the consumer is not a client but an obligatory collaborator, and “the domain emits facts and each client projects them” is not a description of that relation. Nothing here was built that way, so nothing here measures it. Paper 4 (Rivera, 2026d) takes up cross-actor continuity and its **tell**, which is where a reader with such a domain would look; whether it settles this case is not a question these labs asked.

8.2 Invariance under staging, not under evolution

Identity here is invariant under staging, not under evolution. The claim is that a *change of staging* does not change the domain — not that the domain is frozen. Requirements grow, boundaries are redrawn, and verbs are retired; a domain that does any of those is still a domain no staging can call for a change to. Two measurements bound that, one in each direction.

The domain grew. A score and a difficulty level were added to the board — the authority case, a value any client could compute and which therefore has to be single-valued, and a rule over acts already recorded (Appendix A, Lab I). Nothing was deleted but doc comments, no signature changed and no verb removed, shown by diffing a sorted inventory of the board’s members. What the lab supplies that this paper had not measured is the converse of §2: when the domain grew, all twelve host projects kept running with no edit, checked by running them rather than compiling them. The cost of *adopting* the new capability then fell only where it was wanted — nothing for four of them, one line for five, four lines for each of the two browser hosts, and thirty-two for the input host, the one place where the level actually changes what happens.

The domain was re-cut. One role was modelled too large and became two, by authoring the two and then reading the original’s recorded acts — the same read rehydration performs — and driving each new role to perform its own, so that each ends holding a record in its own voice (Lab G). The original’s journal is not cut or rewritten; it is read as the account of what happened, and kept.

Eleven of the twelve host projects were untouched while the domain divided beneath them, and the twelfth needed one line.

Where such a boundary may be cut has a criterion, and §8.4 relies on it: **a boundary is admissible where no decision on either side reads state the other party mutates concurrently.** The cut of Lab G passes it because the board’s phases are disjoint — the pile is immutable while a piece is falling, and the piece is gone once it has been absorbed — so neither role ever reads the other’s state mid-change. The criterion is about concurrency and not about size, and it is what makes a tight-looking split sound.

That second one could look like a threat: if the boundaries *within* a domain can be redrawn, is the identity being measured merely that of the current cut? The distinction the paper turns on answers it. A re-decomposition is a change to the domain, made by whoever authors it for reasons of modelling; it is not a change of staging, and no staging can call for one. What carries across is the record — the acts were performed and remain performed — so the append-only character of every record involved survives the redrawing untouched.

The shape of growth is additive, and one measurement supports that rather than establishing it. An operation added is a new act in the record rather than a reshaping of the acts already there, so growth accretes; Lab I is that shape observed once, on the authority half of the change. The half that adds an *act* rather than a value — a piece held and swapped, a move undone, a second player — is unclimbed here, and so is the retirement of a verb, which in a production system is marked obsolete and kept for as long as any recorded act still invokes it. Both are named as future work rather than claimed. Nor should additive growth be read as a claim that evolving a recorded history is free: empirical study of event-sourced systems finds converting recorded events as their schemas change to be a real and recurring cost (Overeem, Spoor, & Jansen, 2017; Overeem, Spoor, Jansen, & Brinkkemper, 2021). Nothing measured here bears on it. The invariance reported is across stagings of a domain, not across revisions of it.

8.3 Threats to validity

One threat to the result’s validity cannot be argued away, and it belongs in the open. Every staging and every client counted here was written by the same author as the domain, and a metric of the form *the domain did not change* is gameable by construction: any code can be held constant if all variation is pushed outside it, and an empty diff is equally consistent with a domain designed backwards from a known set of stagings. What would discriminate between those is a staging nobody anticipated, introduced by a party with no part in the design.

The record supplies some of that and not all of it, and the part it supplies is checkable in the example’s own history rather than asserted here. The domain’s last commit is dated 29 June 2026 and it has not been modified since. Three

stagings and clients were added on 23 July, twenty-four days later — a gesture client, a constrained-device client, and the three-machine deployment — in working sessions separate from the one that authored the domain, and the two clients appear in no earlier artifact. The domain’s diff across all three is empty. Two of the three are not labs of Appendix A, having been withdrawn there for want of measured evidence about the clients themselves; what is cited here is only their effect on the domain, which was measured, and which is all this particular argument needs.

It remains weaker than an independent test. The author commissioned the work, and cross-machine deployment had been named as a next step in a host comment the day after the freeze. One structural fact bounds part of the worry without dispelling it: the fence is kept by the compiler and not by anyone’s discipline, since the domain declares no dependency and its verb-bearing types are internal (§2), so staging concerns could not have been threaded quietly into the domain even had someone wished to. What that does not bound is the other form of the objection — that a domain’s *surface* may have been chosen with the known stagings in mind — and no property of a build graph rules that out. Nor would another client on this same domain settle it, which is worth saying because it is the obvious next move and the wrong one: the objection is about who authored *this* domain, so what answers it is the arrangement applied to a domain its author did not design for it. That is a second worked domain, not a seventh client, and this paper does not have one.

A further limit is now measured rather than suspected, and it narrows the result: invariance across *transports* is not distinctive of this arrangement. A ports-and-adapters version of the same game, built and counted (Appendix A, Lab F), also takes a new transport without touching its domain — twice over. So the zeros of Experiment A, taken as counts per staging, do not separate this arrangement from a competent ported one; what separated them in measurement was a new *capability*, a game outliving its process, which cost the ported domain a change and cost this one nothing. The claim to carry forward is about capabilities a record supplies, not about stagings as such. And the distribution stagings — the two coordinated actors and the three machines — were measured against no baseline at all, since a ported design has no distribution story to port; those zeros stand uncontested rather than confirmed.

That correction has a limit of its own, and it runs the other way. Every other limit in this section narrows what the paper claims; this one narrows the concession just made. The baseline is the author’s own. It was written after the fact, by someone who already knew which property it had to exhibit, and written deliberately clean — its eleven rule files differ from the journaled domain’s by one line each, the namespace. What it establishes is therefore what ports and adapters *permit*: a rule model with no project references, no packages and no framework named, unmoved by a new transport. It is not evidence that such domains are what the pattern yields in practice, and nothing of that kind is claimed here in either direction. So the comparison does its job — it shows

this paper’s per-staging counts are not distinctive, and that correction stands — while establishing nothing about the population of ported systems, which would take a study of a kind not attempted here. A reader inclined to read Lab F as showing that the arrangement argued for is already common should note that it shows only that a competent instance of the alternative can be built, by the same author, for the purpose; which is the same standing this paper’s own artifact has, and no more.

Two narrower limits belong here as well. The demonstration is one domain on one framework, and it is an existence proof, not a measurement: it shows the separation is buildable and that its mechanical cost is a set of edits not made — a domain diff that stays empty — not that the separation is worth its apparatus at scale, or cheaper in production, which are questions for other work. And the distributed staging of §2 is bounded as that section stated: its peers meet through an out-of-band bootstrap, its coordinator role does not rotate, its transport trusts a certificate on first use, and its replication reaches agreement through a catch-up path rather than an uninterrupted stream. Each bounds the deployment; none touches the domain’s invariance, which is the only thing the staging was there to test.

One of those two observations bears on this paper’s own claim and so belongs here rather than with them. A fact a domain emits must be derivable within a single actor’s state, and that is a property of the substrate constraining a modelling decision *inside* the domain — influence running against the direction this paper measures. It declares no dependency, imports nothing, and its diff stays empty, so §2’s measured claim is untouched. What it does mean is that **declaring no dependency on the framework is not the same as being unshaped by it**. The influence is narrow, and it falls on how a domain may be divided rather than on what it says.

One axis is deliberately left aside, and it closes the list of limits. This paper varies where a domain runs and who observes it, and holds the domain invariant across both; it does not treat which actor talks to which — the addressing and topology of a system of many domains — a separate question with its own answer, not folded into this one.

8.4 Two observations that bound the interpretation

During the experiments two observations emerged repeatedly. Neither is part of the claim advanced here; they are reported because they bound its interpretation and motivate later work.

The first is a hazard, and the evidence for it is the strongest in the paper: **an incomplete record answers plausibly**. Three labs met it independently, by three routes, each taking it for a local incident — a routine correlated to the wrong closing entry with its count still right, a truncated replay returning a board internally consistent and in one case exactly correct, a rehydration silently dropping every zero-argument act (Labs H, I, G). A gap in a record yields not an

error but a smaller coherent story. That is unlike *gray failure*, whose defining feature is differential observability, the application’s view diverging from the observer’s (Huang et al., 2017): here there is none, since every reader reads the same consistent record and no second view exists to disagree. What detects it is therefore not another observer but a ground truth carried forward, which gives a rule cheap enough to state for anyone measuring a journaled system, not this one alone — **a replay measurement must carry the state recorded at play time and assert against it**, or it will pass while being wrong. Two qualifications: it is a rule about measuring rather than running, and the exposure is not peculiar to this arrangement, since any reader reconstructing state from an append-only record has it. It also meets the one row of §9’s table of Waldo’s four differences marked as not addressed — partial failure of a *peer* is exercised nowhere here, while partial failure of a *record* is what these labs found, and it is quiet in a way a dead peer is not.

The second is a constraint, and it is narrower than the first in a way worth stating before it is read: it is a property of *this substrate’s* emitting plane rather than of journaled domains in general, and unlike the first it rests on an argument with checked premises rather than on a measurement — nothing about the loss was measured, because the split actor of that lab exposes no sink, no formatter and no output wiring, so there was no push channel there to lose (Lab G carries both qualifications). On this substrate, then: **a fact a domain emits must be derivable within a single actor’s state**. A complete frame was a join over the pile’s cells and the falling piece’s, computed inside the domain; after the board is cut in two those live in two actors, and a projection on the emitting plane reaches only its own actor’s state, so neither role can push a whole frame (Lab G carries the premises and their anchors). Emitting a joint fact therefore bounds the admissible decompositions of a domain — a boundary may be cut wherever §8.2’s concurrency test permits, except across a fact the domain must emit whole. What this does *not* touch is §3 or Lab D, whose counts stand as taken, nor observation of a divided domain: a party reading both records and joining them is an ordinary adapter, which is what §3 argues for. That reader was not built and nothing is claimed for it. What closes is the push path.

9. Related Work

The separation of a domain from the machinery that runs it is not a new aspiration; reading it as a *measured invariance* rather than a design prescription is where this paper sits among its neighbours.

9.1 What kind of claim the concessions are against

What follows concedes each of this arrangement’s mechanisms to prior work, and every concession is meant literally. But prior art on the parts of a configuration is not prior art on the configuration, and that should be clear before the concessions begin. An architecture in the foundational account is three things

and not one — its elements, the *form* that constrains them, and the rationale for that form — where form is the properties of and relationships among the elements (Perry & Wolf, 1992). Same elements under a different form, different architecture. A form is what is claimed here: a domain that declares nothing on either side, whose acts are recorded by a substrate it does not name, staged by parties that name it. Each element of that is old and this section says where each comes from.

The form has four constituents, and because the contribution is the form and not the parts, they are set out together here rather than left for a reader to gather from the places each is mentioned.

	Constituent	Credited to	In the baseline?	What it accounts for
i	a rule model that declares no dependency	information hiding; functional core, imperative shell	yes — measured tie	nothing distinctive
ii	entered by calls on its own operations, calling nothing outward	inversion of control	yes	nothing distinctive
iii	declares nothing about what becomes of what it emits	implicit invocation	no	the empty publicly callable surface
iv	its acts are recorded as it performs them	system prevalence; event sourcing	no	history and durability at no cost to the domain

Table 1. The four constituents of the arrangement, what each is owed to, whether the ported baseline of Lab F has it, and which measured difference each accounts for. Every constituent is old; the claim is the set.

Whether the *form* is old was settled by measurement rather than by assertion in a related-work section. The baseline holds the first two and lacks the last two, and each surviving difference traces to one of the two it lacks. That the domain

has no publicly callable surface follows from (iii): a port must be visible outside the domain that declares it, and where nothing is declared there is nothing to expose. That history and durability cost the domain nothing follows from (iv) and not from (iii) — a substrate could carry a domain’s emissions to a sink without recording any of them, and then a restart would find nothing to replay. The record is what makes the second difference, and it is named as a constituent for that reason.

9.2 Work that isolates a domain

Most of that work can be compared on one property rather than mechanism by mechanism, and comparing it that way is what the construct above buys. **Existing work isolates a domain. None of it asks whether the domain is finished.** Isolation is a relation between a domain and its surroundings, and every line below achieves some version of it; completeness is a property of the domain alone, and no line below measures it.

Prior work	How it isolates a domain	Asks whether the domain is complete
information hiding (Parnas, 1972)	encapsulate the decisions most likely to change	no — a principle to follow, not a property to check
domain-driven design (Evans, 2003)	a domain layer kept clear of infrastructure	no — and its repository is an obligation the domain declares
functional core, imperative shell (Bernhardt, 2012a, 2012b; Peyton Jones & Wadler, 1993)	purity: effects confined to a shell	no — its subject is testability, and it is conceded below
model-driven architecture (Kleppe, Warmer, & Bast, 2003; Mellor & Balcer, 2002)	platform independence, by transformation	no — what runs is a derivative, not the model
twelve-factor (Wiggins, 2011; Kratzke & Quint, 2017)	configuration separated from code	no — its subject is the deployment, not the domain
software product lines (Clements & Northrop, 2001; Apel et al., 2013)	one codebase, a variant per configuration	no — it manages intended difference
ports and adapters (Cockburn, 2005; Martin, 2017)	declared ports, adapters outside	no — the ports <i>are</i> the incompleteness

Prior work	How it isolates a domain	Asks whether the domain is complete
location transparency (Hewitt et al., 1973; Agha, 1986; Armstrong, 2003)	an actor reachable without regard to where it runs	no — its subject is the address of a peer
command-query segregation and event sourcing (Meyer, 1988; Young, 2010; Fowler, 2005)	a view not served from the write	no — the mechanism §4 stands on, not a claim about the domain

Table 2. Prior work on isolating a domain, against the question this paper asks of one. Three rows are also concessions of mechanism and are developed below: ports and adapters, because it is the comparison the result must survive; functional core, because it reaches the same zero test doubles; and the two that own this arrangement’s input and output edges. Two borrowings of vocabulary are acknowledged rather than compared — *accidental* is Brooks’s (1987), used here of the staging while his distinction partitions difficulty rather than identity, so nothing reported here is a silver bullet; and the theatre is Laurel’s (1991), whose subject is the audience’s experience where this paper takes only the distinction between a play and its staging. Brooks’s cut was later applied to software’s own materials by Moseley and Marks (2006), separating essential logic from the accidental complexity that state and control flow impose on it — made within a program where this one falls between a domain and its stagings, and the reason the concession to functional core below is owed at all: keeping an essential part free of accidental machinery is an old goal with a literature.

That tradition also states the isolation goal this paper measures, and it is worth naming the one point where its own answer stops short, since a reader who practises it will feel that point before any other. Domain-driven design asks that the domain layer be kept clear of infrastructure, and for persistence it offers the repository: the interface declared *in the domain layer*, the implementation outside it (Evans, 2003). That is a compromise rather than an oversight, and Lab F shows why it is forced — reconstituting an aggregate requires a way in, so a design of that shape must have the domain offer something. In the vocabulary above, the repository is an obligation the domain leaves open, and a domain-driven design domain is decoupled rather than closed. This one offers no repository interface, no factory for a mapper, and no restore constructor, because reconstitution is the recorded acts being performed again through the domain’s own operations: the way in is the way it was always entered.

The scope of that difference should be stated narrowly, because it is easy to inflate into a claim about the whole tradition. It concerns *isolation from in-*

frastructure and nothing else. Domain-driven design is also a discipline of modelling — a language shared with domain experts, contexts bounded deliberately, a model refined against the business — and a Tetris board demonstrates none of that. Nor does this paper reach the domains that tradition is usually brought in for, which act on the world and then reason about what happened; §8.1 concedes them. What is claimed is one axis: on the isolation goal both share, the arrangement here pays no repository, and the ported baseline built for comparison shows what paying one costs.

9.3 Ports and adapters, and what closure adds to it

Closest, and the comparison the result must survive, is hexagonal architecture — ports and adapters (Cockburn, 2005) — and its restatement as clean architecture, the domain at the centre with every dependency pointing inward toward it (Martin, 2017). The comparison is usually put as a question about staging invariance, and put that way it is settled and uninteresting: a ported version of this same game was built, given the same stagings one at a time, and counted, and two of its transports cost its rule model nothing at all (Appendix A, Lab F). Per staging the two arrangements are indistinguishable, and this paper concedes it without reservation.

The difference that does separate them is not a count of edits, and it is easier shown than argued. Set what each arrangement still requires of someone else side by side, taking every row from Lab F or Lab E:

	What is measured	Ported baseline	Here
I	outward obligations the domain declares	3 driven ports — board output, piece selection, state	0
↳	stand-ins its own tests require	3, without which 20 of 64 tests do not run	0
↳	publicly callable domain operations	at least one per port	0 — one type with no members
↳	a state admissible from outside	GameState , arriving through the state port	none; only acts
II	a reconstitution surface of its own	a restore constructor and a pile factory: 56 lines added, 5 removed	none — replay re-enters by the operations the acts were performed through

What is measured	Ported baseline	Here
---------------------	-----------------	------

Table 3. What each arrangement still requires of someone else — **two independent measurements, not five**. The rows marked $\downarrow$ are consequences of the fact above them and are shown because each is separately checkable, not because each is separate evidence: a stand-in is a driven port seen from the test side; a port must be visible outside the domain that declares it, so where nothing is declared there is nothing to expose; and the state the baseline will accept arrives through one of those same three ports, the one its own source calls *a third driven port, added by staging 4*. The quantity is unfulfilled obligations, not lines of code.

The conjunction of those two measurements is what this paper calls **closure**, and the word is doing no work the counts do not: *the domain declares no obligation outward, and its reconstitution requires no surface of its own*. A port is an obligation the domain itself declared and left for someone else, so the first fact is what makes the three rows under it follow; and reproducing by the operations the acts were performed through is what makes the second one hold, since there is no event distinct from the act for a separate replay surface to take. A domain with neither obligation is finished. Ports and adapters *decouples* a domain; this arrangement *closes* one.

Two established senses of the word meet here, and both are better dealt with than left for a reader to supply. One is favourable and is claimed deliberately: in the language of modules and linking a component with no unresolved external reference is *closed*, and the closure of a set of modules is what a linker computes when nothing outward remains to be found. That is very nearly the property measured above, arrived at by asking a different question, and the resemblance is not a coincidence to be excused. The other is unfavourable and is denied. In the open/closed principle, *closed* means closed to modification (Meyer, 1988) — and this domain is not that. It grows: a score and a difficulty level were added to it, additively and without a signature changing (§8.2, Lab I). What is closed here is the set of obligations the domain leaves to others, not the set of changes its author may make to it.

Three things bound that, and the first is the most likely to be misread. Closure is not visible in the domain *considered in isolation*: a hexagon’s driving port is the domain’s own API, which nobody supplies and nobody doubles, and a domain needing nothing outward declares no driven port either — the **Well** as bare rules over a board is such a domain (§8.1). All five rows appear only once the domain is *deployed*, when its state must outlive a process and arrive on a second machine. In the ported arrangement those became declared needs; here they were never needs, the acts having been recorded as they were performed.

The second bound is on the word. Calling the sum of those counts *closure* is a reading of them, as §5's *identity* is a reading of the structural property, and the paper keeps reading and measurement apart here as it does there. The third is that history and durability are one row of the table and not the claim; purity, conceded below to functional core and imperative shell, buys freedom from effects and closes none of these rows. The obvious objection — that orthodox practice persists an aggregate through an adapter-side mapper without touching the rule model at all — is answered by measurement rather than by argument: the baseline attempted exactly that design and still needed a seam authored inside its own rule model, for a reason its source states and Lab F sets out.

One measured result runs the other way and belongs in this comparison rather than in its appendix. The build graph does not separate the arrangements: a ported rule model is equally free of project references, packages, and framework names. So the testability argument above must be narrowed to what it can bear. What cannot be exercised without a stand-in is the ported *application service* and its three driven ports, without which twenty of its sixty-four tests do not run; the ported *rule model*, taken alone, is as testable and as dependency-free as the journaled domain. The driven and driving sides are not symmetric.

It is fair to press the point: if the substrate provides the mechanism, has the port not simply moved into the framework — universal, unnamed, but a port all the same? Two answers, and the second is the one that can be checked. A term stretched that far stops separating anything: a domain declares neither the garbage collector, nor the scheduler, nor the collections library it computes with, and calling each of them a port empties the word. What *port* earns its keep by naming is an interface a component **declares and depends upon**, because that is the thing with the consequences this paper counts — an implementation per staging, a stand-in per test. And measured there the difference is not interpretive: the domain of §2 compiles and passes its tests with the framework absent from its build graph, which a domain that declares a driven port cannot do, since a port with neither implementation nor double is not something a test can run against at all.

Ports and adapters:

```
Adapter --calls--> [driving port = the domain's OWN api] --> Domain
Domain --declares--> [driven port] <--implements-- Adapter
(only the driven port is a dependency of the domain, and a domain
that needs nothing outward declares none)
```

Here:

```
Domain          Staging --binds--> input source / sink / client
(no declared dependency either way; the state that must outlive the
process and reach a peer is the substrate's doing, not a port the
domain asks for)
```

The contract has not been inverted; it has moved off the domain's boundary, and two places are involved rather than one. The *contract* is the sub-

strate's: it defines what an output is, what may command the domain, and how either is carried. The *binding* is the staging's — which sink, which client, which transport — and it happens later than a reader is likely to assume, a destination being attached to a live actor rather than fixed at build or deployment (`Choreography/Theater/PerformanceV2.cs:415`). So the formulation used throughout is that the contract is provided by the substrate and bound in the staging, never that it moved into the staging, which would credit the staging with a mechanism it only wires up. The same shape holds for the transport carrying an actor's messages to another, the domain naming no wire there either; that belongs to the axis this paper sets aside (§8) and to Paper 4, and is mentioned so the output side is not read as a special case. What the move costs is measured rather than reasoned (Appendix A, Lab F): a ported arrangement pays no domain edit per staging either, and its driving side needs no double, a hexagon implementing its driving port rather than depending on it. Where it pays is on the driven side — three ports without stand-ins for which twenty of its sixty-four tests do not run — and in its publicly callable surface, at least one operation per port against none here. That is the observable trace of the move and not the point of it. The point is that the boundary the contract sits on is no longer inside the domain, which is what lets the domain be the common ancestor of §5: a port is a thread from the domain back to the shape of its clients, a fragment of a staging lodged inside it, and while it is there the identity is not clean.

9.4 Who owns the two edges

The geometry itself has a prior claimant, and it is the nearest neighbour this paper has. Implicit invocation — a component announces an event, components registered elsewhere are invoked, and the announcer names no recipient and declares no interface for one — was formalized as a design space three decades ago (Garlan & Notkin, 1991) and studied as the means by which independently built tools are integrated and independently evolved (Sullivan & Notkin, 1992); its later forms are surveyed as publish/subscribe (Eugster et al., 2003). A domain that emits facts under logical names and says nothing about who receives them is doing that, and the coupling that remains — agreement on the names — is what that literature calls nominal coupling. The geometry of §2 is therefore implicit invocation's geometry, and the honest statement is that this paper did not discover it.

Three things distinguish what is claimed here, and all three are narrower than the geometry. The first is the direction the vocabulary faces. In implicit invocation the event vocabulary is an *integration* vocabulary: announcer and listeners agree on a set of events in order to be integrated, and the names exist because something will listen. The names a domain emits here are its own — the operations it performs and the facts they leave behind — and they would stand unchanged if nothing listened at all; the staging adapts to them, never the reverse. The second is that implicit invocation is offered as a mechanism, with

a design space to explore and trade-offs to weigh, whereas the claim here is an invariance measured on a built system: the domain's diff across five stagings and six clients. The mechanism makes that zero possible; it neither asserts nor measures it. The third is the conclusion drawn. That literature concludes that components may be integrated and evolved independently. This paper concludes that the domain stands as the common ancestor of every staging of it — a claim about what a domain is, not about how the parts of a system are connected.

The input edge has a prior claimant of the same kind. That a framework calls the application's code rather than the reverse was named *inversion of control* when frameworks were first described as a form of reuse (Johnson & Foote, 1988), and it is the mechanism by which a domain can be driven without declaring an interface for whatever drives it: the substrate invokes the domain's operations, and the domain calls nothing outward. Section 2's input-side zero rests on that mechanism and the credit belongs there. What the term later came to name in practice — a container that wires declared dependencies together (Fowler, 2004) — is a different thing, and the distinction matters for the reason Paper 8 gave: injection presupposes something to inject, and where a domain declares no port there is nothing to wire. The mechanism is inversion of control; the measurement, and the identity it reveals, are what is added here.

One practitioner pattern has arrived at that same geometry from a different motive, and it deserves an explicit concession. Functional core, imperative shell places a dependency-free core inside a shell that owns all input and output, and its best-known consequence is precisely the count reported above: the core can be tested with no test doubles, because there is nothing there to double (Bernhardt, 2012a, 2012b). That zero is therefore not this paper's discovery, and it should not be presented as one. The name belongs to practitioner literature, but the discipline beneath it does not: confining effects to a boundary so that the remainder can be reasoned about and exercised without them is the monadic treatment of input and output, argued in the refereed record three decades ago (Peyton Jones & Wadler, 1993). What differs is the claim the zero supports. The pattern buys its core's independence with *purity* — the core is functions without effects and the shell holds the mutation — and it is a discipline about the internal structure of a single program, whose shell is one imperative wrapper. The domain here is not pure: it mutates. What is recorded of that mutation is recorded by the substrate, and the domain never learns a journal exists — which is the measured claim of §2 and not a courtesy of this comparison. So purity is neither available as the mechanism nor needed as one: the independence is bought not by declining to mutate but by keeping the account of the mutation someone else's work. And its shell is not one wrapper but many stagings, one of them spread across three machines. Same zero, different result: the pattern explains why a *test* needs no double; this explains why a *staging* needs no domain edit, and what that says about the domain's identity.

9.5 The local/remote objection

Which raises the strongest objection available to this paper, and it is an old one. Waldo, Wyant, Wollrath, and Kendall (1994) argued that the distinction between local and remote cannot be papered over: latency, memory access, concurrency, and partial failure are irreducible, and systems that hide them fail precisely where reliability is demanded. Experiment A — a domain moved from one process to three machines with no edit — has the shape of the thing they warned against. The objection is accepted here rather than rebutted. Nothing in this paper claims those differences vanish, or that distribution is free. Nor does it help to gesture at the incidents a distributed staging happens to produce — an out-of-band rendezvous, a connect-readiness race, a certificate trusted on first use — as though each answered one of Waldo’s four. It is more useful to set the four out and say, for each, where it falls here and whether this paper does anything about it:

Waldo’s difference	Where it falls here	Status
Latency	the transport, and the replication that crosses it	present; nothing about it was measured, no timing was taken
Memory access	nothing is shared by reference at all: each node keeps its own copy of the record and reconstructs its own state from it	answered by the arrangement rather than borne by it
Concurrency	the connect-readiness race of Lab B; one writer per record	present, and the race is reported
Partial failure	not exercised — no peer was killed, no partition induced (Lab B)	not addressed

Table 4. Waldo et al.’s four differences against what this paper does with each. Two incidents a reader might expect to find here belong to neither column: an out-of-band rendezvous is a bootstrap question and trust-on-first-use a security one, and neither is among the four.

Read that way the honest statement is narrower and more interesting than the one it replaces. *Memory access* is the case where it hurts most, and it is not evaded: there are three copies of the board’s state in Lab B and no shared reference anywhere, which is not a workaround for the absence of shared memory but an arrangement that never assumed it. *Partial failure* is where this paper is weakest, and Lab B says so already. What holds is that the differences

which do appear are borne by the staging, and that none of them reached the **Well** — a claim narrower than location transparency’s, and deliberately so: the local/remote distinction is real and irreducible, and it is not a fact about the domain.

9.6 How the claim is checked

The method by which the claim is checked has a literature of its own. Software reflexion models compare a high-level model of a system against the dependency graph extracted from its source, reporting where the two converge, diverge, and fall silent (Murphy, Notkin, & Sullivan, 1995, 2001). The conformance-checking work that followed made such comparisons routine: languages in which permitted and forbidden dependencies are declared and then checked against the code (Terra & Valente, 2009), comparative accounts of the available static techniques (Passos et al., 2010), and tools — ArchUnit, NDepend — in which such rules are asserted as ordinary unit tests. What motivates that whole literature is architectural erosion: the drift of a built system away from the architecture intended for it (Perry & Wolf, 1992). Section 2’s reading of the build graph is a measurement of the same kind, and the resemblance is what lets this paper say *checkable* rather than *arguable* — with erosion as the instructive contrast, since it is what a dependency direction does when nobody is enforcing it, and the empty diff is what the same direction looks like when it is a property of the arrangement rather than a rule someone must police. The difference lies in what is checked. Conformance checking takes an intended architecture as its input and counts departures from it; there, the direction of dependence is a rule imposed and then policed. Here the direction is not a rule but the substance of the claim, and the quantity measured is not a violation count but an invariance under deliberate variation: change the staging, and see whether the domain moves. The first is an audit; the second is an experiment.

9.7 What §4 owes the data literature

The nearest neighbour of §4 is a data-management literature, and it must be credited carefully, because two distinct papers bear on it. *Immutability changes everything* argues for append-only computing: observed facts recorded and kept, results derived from them on demand, on the model of an accountant who corrects by adding an entry rather than by erasing one (Helland, 2015); Kleppmann and Kreps (2015) turn that into a way of composing systems around a log. The earlier paper is the one that reaches §4’s temporal claim. *Data on the outside versus data on the inside* holds that data outside a service is beyond the scope of that service’s transactions and temporally disconnected from it (Helland, 2005) — which is to say that what an external consumer holds is always the past. Section 4 therefore does not discover that; it arrives at it twenty years later, in the vocabulary of roles rather than of services, and the honest statement is that the data literature drew the conclusion first.

What survives the concession is narrower, and worth stating exactly. Helland's argument concerns where data sits relative to a transaction boundary and what may be assumed of it there. Section 4's concerns *offices*: that watching and acting are distinct roles; that the present is available only in the acting one; and that an audience is therefore constituted by the record rather than merely served from it. Its operative consequence is a design corollary the data framing has no occasion to state — a command's response cannot serve a client that issued no command, so a system whose only view path is that response has none for an observer — which follows not from a property of data but from whose the response is. The temporality is Helland's, and the structure beneath it Lamport's; the attribution of roles, and the corollary it yields, are what this section adds. It does not claim more than that, and §4 says why: returning state in a command's response is a recognized session guarantee (Terry et al., 1994), not an error, and the present is observable by an audience that polls for it.

9.8 The series, and what remains unclaimed

Two of this series' own papers hold pieces of the present result. Paper 7 read the server as an *accidental category* — ephemeral under this substrate, discharged at bootstrap and absent from the running topology after — and offered that as a lens rather than a result. The consequence drawn here is that where a domain runs is then not constitutive of it, which is one axis of the staging this paper varies. Paper 8 located, from the side of output, the boundary between what a domain holds and what an actor projects for a client, which is the adapter boundary of §3 seen from within a single output. This paper generalizes both: it varies the whole staging — where a domain runs and who observes it — and names what stays fixed across all of it as the domain's identity.

What remains unclaimed, after all of these, is the intersection — which is to say the form, in the sense this section opened with. Late binding of a destination is old; platform independence has a name and a methodology; location transparency has a production tradition; a dependency-free core has a pattern; being driven without declaring an interface for the driver has a name three decades old; and checking a dependency graph against an intended model has a literature. Each of those is conceded above, individually and without reservation. What none of them states is that the contract may leave the domain on *both* sides at once, so that a domain declares nothing about what drives it or about what it emits — and that the property this leaves behind, the domain standing as the common ancestor of every staging of it, is measurable after the fact on a system already built.

A reader is entitled to ask why a list of conceded parts should leave anything behind, and the answer is the one the section opened with rather than a rhetorical one. The claim is about a form, and a form is tested by building the alternative. The alternative was built, from parts every one of which is credited above, and it produced neither of the two differences that survive. Nothing here rests on

the parts being new. What it rests on is that assembling them the ordinary way does not arrive here — a claim about the arrangement, and one this paper publishes a falsification test for: build a ported design that obtains durability and a second machine without touching its rule model.

One such design is available and needs naming, since it is the first a reader who has built one will reach for, and it is nearer than the mapper of Lab F. An event-sourced hexagon declares an event-store port and reconstitutes by replay rather than by mapping, so it needs neither a restore constructor nor a factory for a pile, and the 56 lines this paper counts would not be spent. What it does need is a *reproduction surface* distinct from its command surface: an orthodox event-sourced aggregate exposes something of the shape of `Apply(event)` and `LoadFromHistory(events)`, so the rule model still grows a seam, only a different one. The residual claim stated exactly is therefore not that a ported design must pay a mapper’s seam, but this: reproducing by the domain’s *own* operations requires no reconstitution surface at all, because the substrate re-invokes the very verbs the acts were performed through, and there is no event distinct from the act for an `Apply` to take. That is *predicted* rather than measured — no event-sourced baseline was built, and building one is the lab that would close the test this paper publishes.

10. Conclusion

This paper began where practice is settled. A program’s domain and its deployment are written together, so changing where a system runs, or which client it serves, means changing the system. Set beside theatre, where one script is staged in many venues without being rewritten, that entanglement stops looking like a property of software and starts looking like a confusion of a play with its staging.

The confusion has a measurable opposite. One domain — a Tetris board — was run on a console, in a browser over a WebSocket, in a browser over server-sent events, across two co-hosted actors over real TLS, and across three separate machines in containers; and it was read by a person at a keyboard, an automated player over a pipe, two passive observers, and a browser drawing it in JavaScript. The number of changes that variation forced on the domain is zero, and the number of test doubles the domain’s own tests require — for what drives it or for what it emits — is also zero.

Behind those counts is a property of the domain rather than a discipline, and the counts are not where it lies: a ported rule model is equally free of references and packages, equally testable without stand-ins, and equally unmoved by a new transport (§9), so those counts distinguish a clean domain from an entangled one and not this arrangement from that one. **Ports and adapters decouples a domain; this arrangement closes one.** Section 9 puts that on two measurements — whether the domain declares any obligation outward,

and whether reconstituting it needs a surface of its own — and the 61 changed lines a ported rule model spent to survive a restart are one row of the second, not a separate result. The domain here declares no interface for what drives it or for what it emits, so the contract fixing what an output means, where it goes, and what may command it does not sit on its boundary at all: the substrate provides it and the staging binds it. That is not dependency inversion carried a step further; it is the contract in a different place, and it is what leaves every staging's references leading to the domain and none leading out — a sink in the dependency graph, the root of its reverse, which is all the paper means in calling the domain their common ancestor. What is checkable against a built artifact is that structural property together with the empty diffs; reading it as an *identity* is what §5 does, on the terms §5 sets. Read so, *precedes* names the direction of the dependence. A staging is built against a domain that does not know of it; the domain is not the thing that survives its stagings but the thing that makes them possible.

Two consequences followed with no new machinery. Watching and acting are distinct roles, and what each may have differs: a live read taken *in order to act* is part of the act and belongs to the party acting, while a view served from a command's response can serve no client that issued none. So an audience is constituted by the record rather than served a state read off the moment — not because the present is closed to it, since one of the six clients polls the board and is handed it, but because what an audience is told is an account and not a configuration. And because the domain keeps its identity, the account of what it did keeps its identity too, so the routine those acts compose is recognizable from any stage, through any adapter. Neither is a second claim; both are what the first one entails.

What can be carried away is an instrument rather than a prescription, and the difference is worth being exact about, since this paper declares itself the analytic kind of contribution. Nothing here establishes that an architecture is better for holding the property, which would be a judgment about what makes an architecture good. What the paper supplies is a question that can be put to a system already built and answered by measurement rather than opinion: does the domain inside it hold an identity independent of where it is staged and to whom? Vary the staging, and count what the domain had to change. The labs of Appendix A are that question asked once, of one domain, with the answer zero.

Two things came out of those labs that nobody set out to find. Neither is the result of this paper, and both are reported because they bound how its result may be read — the first by bounding what a reader of a record can be told, the second by bounding how a domain that keeps one may be cut. The first is a hazard, and the evidence for it is the strongest in the paper: **an incomplete record answers plausibly**. Three labs met it independently, by three routes, none investigating it — a routine correlated to the wrong closing entry with the count still right, a truncated replay returning a board that was internally

consistent and in one case exactly correct, a rehydration silently dropping every zero-argument act. A gap in a record yields not an error but a smaller coherent story, and because every reader reads the same record and it is consistent, no second view exists to disagree. The rule that follows is cheap and is offered to anyone measuring a journaled system, not just this one: carry the state recorded at play time and assert a replay against it. A replay checked only for internal consistency will pass while being wrong. The second is a constraint, and it is stated in §8: a fact a domain emits must be derivable within a single actor's state, so emitting a joint fact bounds how that domain may be divided — the one respect in which the substrate shapes the domain rather than the other way round, and the more interesting for being the exception. Whether an architecture *ought* to have the property — whether the arrangement pays for itself, and against what — is exactly what §8 and Appendix A decline to claim, and a reader who wants that will have to measure it.

Four questions the argument opened are left standing, each larger than this paper. If an entitlement must be attributed to something that holds still, then identity is prior to authority, and the series' earlier questions about who may speak rest on this one. Who is *entitled* to read a routine, §6 reaches without settling — though whether a reading of one can be wrong it does settle, and the answer is that it can, and silently, when a moment that mattered was never an act. Which is the third question, and the one this paper is most exposed on: what must be an act at all. A criterion is offered here for a record's vocabulary and none for its granularity, and it is granularity that an account's fidelity rests on (§6). And because a domain that declares no dependency cannot reach for another, whatever composes two domains is neither of them — the question that anything treating a set of domains as a repertoire would have to answer first. Closure is what this paper hands to that question. A domain that still declares obligations is a *part* of a system, completed by whatever supplies them; one that declares none is a whole, and a whole is the only kind of thing a repertoire can hold. Whether closure is sufficient for that, or merely necessary, is not settled here.

The audience changes. The stage changes. The play asks nothing of either.

Appendix A. Labs

The claim of §2 is a count, and this appendix reports the counts and says where each can be checked. These labs are illustration rather than evaluation: each shows that a separation the paper argues for is *built*, and what it costs the domain, which is nothing. In the manner of the earlier systems papers, every result here is a count the separation drives to zero, and the zeros are the claim. Two things are not measured and are marked as such rather than mixed in: whether the arrangement is cheaper or faster at scale, or better in production, which §8 leaves open; and what the distribution zeros are zeros *against*, since the one alternative arrangement that was built — the ported baseline of Lab F

— has no distribution counterpart to set beside them.

All of them run against one domain — the `Well` and its operations (`domain/`) — reached by every host through a single actor (`actor/TetrisActor.csproj:11`), which the example’s own comment calls an accidental shell (`actor/IGameHost.cs:14`). The domain project declares no reference of its own (Lab E).

Lab A — the stage. The same `Well` runs on a console; in a browser with input and output carried over a WebSocket; in a browser over HTTP with server-sent events; and across two `StageManager` actors, joined once in memory and once over a real Kestrel TLS channel. Five stagings, differing in process, transport, and wire format. A diff of the domain directory from the first staging to the last is empty, and the example builds with no errors. The domain’s own test suite also passes unchanged throughout, which is worth stating for what it is and no more: a regression check showing the stagings did not break the domain. Its coverage was not measured, so neither the suite’s size nor its passing is offered here as evidence that the domain is adequately exercised — a test count is not a coverage measurement, and nothing in this appendix should be read as one.

Lab B — three machines. Three Docker containers run the same `Well` as three `StageManager` peers on a private bridge network: one Director, two casts, joined over Kestrel TLS on a port never exposed to the host, with coordination and replication crossing genuine container-to-container TLS. A scripted driver on the Director plays the game to a non-trivial board; the three nodes converge on the same journal entry and the frame each pushes is byte-identical. Adding this staging left both the domain and the actor untouched — the diff of each across the whole increment is empty, and the files added are a new host, its Docker definitions, and notes.

Two details of that run belong here rather than in the body. The game was played by a script rather than by a person, but it was played: the `Well` genuinely mutated on one machine and the same state reached the others. And what each node is handed is the domain’s *assembly*, named through the anchor type — the `Well` itself is passed nowhere and constructed by no host, but built inside each node by a recorded act, the seeding command the actor issues on startup, so that a board exists on three machines without any of them ever having held one.

The weight this lab can carry is small, and saying so is more useful than dressing it up. It is a *single* run, reproducible from one script but not repeated, with no fault injection: no peer was killed, no partition induced, no message dropped. “The three converge” is therefore an existence result and not a claim about convergence under adversity, which repetitions and induced failures would be needed to support. Its other boundaries are the ones §2 states — an out-of-band rendezvous, a fixed Director, trust-on-first-use certificates, convergence completed by the framework’s catch-up path. And by this paper’s own §8 the lab is measured against no baseline at all, since a ported design has no

distribution story to port. What it establishes is that the domain’s diff stays empty when the staging is genuinely cross-machine, which is one more staging on the same axis as Lab A rather than a second kind of evidence.

Lab C — the client. With the stage held fixed, the client varies: a person at a keyboard reading an on-screen grid; an automated player sending moves over a pipe and reading a view it computes for itself; two passive observers, one pulling the board and one receiving it pushed; and a browser in which a player and a spectator are both written in JavaScript and reach the game over the network. Six clients over one **Well** — some in C#, one a PowerShell script, two in a browser — and the domain directory does not change for any of them. Each adds an input adapter, an output adapter, or both. One of those adapters was written by the client that reads through it rather than by the author of the domain: the automated player is an instance of a large language model and authored its own column-height view during the lab (§3; disclosed in the acknowledgments).

Two counts belong here rather than one, because they answer different questions and the larger flatters the smaller. Six is the number of clients added, and it is the number that bears on the domain’s diff: each was a separate addition, and none of the six cost the domain an edit. But six overstates how *differently* the board was read. Three of them — the person at the keyboard, the observer that pulls, the observer that receives a push — read the same grid over three different transports, and the browser’s grid is that grid once more, computed independently in JavaScript. Counted by representation instead of by client, these labs hold **two**: a grid of cells, and a vector of column heights. Counted by implementation, three: a C# renderer, a JavaScript one that mirrors it, and a PowerShell height profile. The six is what supports the invariance claim; the two is the honest measure of representational distance, and §3’s argument that no view is privileged rests on the two and not on the six. Both are readable off one table, so a reader need not take either count on trust:

Client	Reaches the domain by	Representation it reads
console, driven by a person	keyboard, in process	grid of cells (C# renderer)
automated player	a pipe to a warm server	vector of column heights (PowerShell)
observer	a query it issues, pulling	grid of cells (C# renderer)
watch	a pushed frame	grid of cells (C# renderer)
browser player	WebSocket	grid of cells (JavaScript)
browser spectator	server-sent events	grid of cells (JavaScript)

Table 5. The six clients of Lab C. Reading down the last column gives the count

that matters: two representations, one of them implemented twice in different languages. Reading down the middle column gives six distinct ways in, which is what the domain's empty diff was measured against.

Lab D — the adapter. One emitted frame reaches three different projections, none of them the domain's. The board emits its occupied cells — a set it computes, in its own vocabulary (`domain/Well.cs:322`, and §3 on why that is a fact it holds rather than a view); a renderer lays that out as a terminal grid (`actor/BoardRenderer.cs:19-48`); the browser produces the same grid independently in JavaScript, its code commented as mirroring that renderer (`web/Program.cs:169-178`); and the automated player lifts the same frame into a column-height profile with its gaps, wells, and aggregate figures (`tools/pile-scan.ps1:54-83`). Three implementations over one fact, and no domain method for any of them — but two representations, not three, since the browser's grid mirrors the terminal's; the height profile is the one that differs in kind.

Lab E — both sides testable, with the framework absent. The domain's own tests run with no staging bound at all: no host, no transport, no sink, and no double — not for what drives the domain and not for what it emits. Stronger, and this is the form in which the claim of §2 is checkable rather than interpretive: the framework is absent from the domain's build graph altogether. The domain project is four properties and no references —

The fence at the level of members. Every verb-bearing type in the domain is *internal*: the board, the pile, the shapes, the piece types, the seven pieces, the rule exception (`domain/Well.cs:34`, `domain/Pile.cs:21`, `domain/Frame.cs:29`, `domain/PieceType.cs:9`, `domain/Pieces.cs`, `domain/TetrisRuleException.cs:17`). The library's whole public surface is one empty type, so not even the actor's own project can invoke a domain operation in the host language — it obtains the assembly and passes it on. The compiler refuses where reflection admits.

Where in-language access was genuinely wanted it had to be granted by name, from inside the domain's own assembly, and it was: to the domain's test project and to one console host (`domain/AssemblyInfo.cs:8-9`). That is an authored exception, enumerable, and extended to no other staging — the four remaining stagings of Experiment A cannot name a domain type at all. Two qualifications belong with it. One grant is to a *staging*, so the claim is not that no staging may call the domain but that calling it takes an act of authorship inside the domain, by name. And that grant is no longer exercised — the console drives the actor and names no domain type, as its own project comment says — so the fence is in practice tighter than its source admits and the grant could be withdrawn without breaking a build. It is reported rather than tidied away to purchase a cleaner claim, and it is left standing: withdrawing it would tidy the fence at the price of changing what this lab measured, which is a poor trade for a count the paper leans on.

It would be wrong to say the domain references *nothing*: it targets a language runtime and uses that runtime’s collections, as every C# project does. The precise claim, and the only one this paper uses, is narrower — **the domain declares no dependency on the framework that runs it, or on any staging of it.**

```
<Project Sdk="Microsoft.NET.Sdk">
  <PropertyGroup>
    <TargetFramework>net9.0</TargetFramework>
    <Nullable>enable</Nullable>
    <RootNamespace>Tetris</RootNamespace>
    <AssemblyName>TetrisDomain</AssemblyName>
  </PropertyGroup>
</Project>
```

— with `dotnet list package` reporting no packages for it, no framework namespace imported by any of its sources, no framework attribute, and no framework interface implemented; the test project references a test runner and the domain and nothing further. The suite passes on that graph — a regression check, as Lab A says, not a coverage claim.

What this count does and does not separate is settled by the ported baseline of Lab F, which has the same clean graph: zero project references, zero packages, no framework named. So *this* count ties, and “zero references, zero packages” separates a clean domain from an entangled one rather than one arrangement from the other. What does not tie is on the driven side only: the ported application service cannot be constructed without stand-ins for its three driven ports, and twenty of its sixty-four tests cannot run without them, while the journaled domain’s tests need none.

Lab	What varies	Measured
A	5 stagings (console, WebSocket, server-sent events, StageManager in memory, StageManager over TLS)	0 domain edits; the domain suite passes unchanged (a regression check, not coverage)
B	3 machines (containers, container-to-container TLS)	0 domain edits and 0 actor edits; three nodes converge, byte-identical frame
C	6 clients — some in C#, one a PowerShell script, two in a browser; 2 distinct representations among them	0 domain edits
D	3 projections of one emitted frame	0 domain methods for any view

Lab	What varies	Measured
E	domain built and tested with the framework absent from its build graph	0 references, 0 packages, 0 doubles; the suite passes
G	the domain's own internal boundary — one role cut into two, stagings held still	11 of 12 hosts unchanged (the 12th, one line); 0 divergences over 47,783 steps; a whole game replayed into both roles, game over included; record $2.32\times$ (316 against 136), 6 extra entries per landing
H	one reaction over 3 stagings' journals (single process, warm server over a pipe, 3 containers read in place)	same 2 placements in the same order in all three; the 3 containers' journals byte-identical, and the acts matching verb for verb across stagings with entry ids offset by a constant; 0 domain edits — with no correlation handle, no landing act, and entry ids not staging-invariant
I	the domain itself grows — a score and a difficulty level added	+30 lines of code (98 added lines in all, 62 doc comments and 6 blank), 0 code lines deleted (3 doc-comment lines removed), no signature or verb changed; build graph byte-identical; 12/12 hosts keep running at 0 edits; adoption costs 0 for four, 1 for five, 4 for each browser host, 32 for the input host — only where wanted

Table 6. What each lab measured. Every entry is a count taken from the example’s own history or build, and none of these rows is a comparison — the one comparison, against a ported arrangement built for the purpose, is Lab F and is reported separately below.

Lab F — the ported baseline, built and counted. The comparison §9 rests on is a measurement rather than an estimate, and this is where it was taken. An orthodox ports-and-adapters arrangement of the same game was written — its eleven rule files differ from the journaled domain’s by one line each, the namespace, so the comparison is not rigged by writing a worse domain — and four stagings were added to it one commit at a time so that each could be counted: a console, a WebSocket host, a REST-and-server-sent-events host, and an automated player. All four run. Adding the baseline changed nothing on the journaled side: the diff of the domain, the actor, the solution file, and all fifteen existing hosts is empty, and the journaled side’s tests still pass.

A ported arrangement invites four estimates about what it costs a domain, and measurement contradicts two of them, ties a third, and leaves a difference that is not the one the estimates predict. Table 7 reports each against what was counted.

Estimate	Measured against the built baseline
an implementation per staging	confirmed — 2, 1, 1, 2 across the four stagings
a stand-in per driven port	confirmed — three; the application service cannot be constructed without them, and 20 of 64 tests cannot run at all
a domain edit per staging	refuted — the WebSocket and the server-sent-events stagings each cost the ported rule model nothing, identical to the journaled side, twice over
a stand-in “per side”	refuted — the driving side is zero, because a hexagon <i>implements</i> its driving port rather than depending on it (§9)
a distinguishing build graph	tie — both rule models: zero project references, zero packages, no framework named

Estimate	Measured against the built baseline
cost of a new <i>capability</i>	cross-process persistence moved the ported arrangement by 204 lines added and 9 removed, of which 61 changed lines — 56 added, 5 removed — fall inside the rule model itself; the journaled side gained the same capability, for the same client, at nothing

Table 7. The ported estimates, measured. Two are refuted, one ties, and the difference that survives is not the one the estimates predict.

What the baseline actually shows: the unit of account was wrong. A port costs a domain nothing per *staging* and everything per *capability its ports do not already carry*. Two new transports moved the hexagon not at all — the same zero this paper reports, reached by a different arrangement — so invariance across transports is not distinctive, and §2’s counts should not be read as if it were. What moved the hexagon was persistence: to let a game outlive its process, the ported rule model had to grow a restore constructor and a way to rebuild a pile from cells — 61 changed lines, added and removed, not added alone. The journaled arrangement acquired the same capability, for the same client, at no cost to the domain, because a record of the acts already existed. The measured difference is therefore not *per staging* but *per capability of the kind a record supplies for free*.

Read as a test of an arrangement rather than as a scoreboard, that is the more useful thing the baseline does, and §9 leans on it there. This lab holds two of the three constituents of the arrangement the paper argues for — a rule model free of dependencies, and a domain entered by calls on its own operations rather than by calling outward — and lacks the third, declaring nothing about what becomes of what it emits. Holding two of three, it produces neither surviving difference. Which is what makes the paper’s claim a claim about the assembly and not about any part of it, since every part is credited to prior work.

One number in that table needs a caveat stated plainly, since it is the most attackable figure in the baseline. Where a ported arrangement’s *domain* ends is a matter of layout, and this baseline places its rule model, its ports, and its application service in a single project — the ordinary arrangement at this size, and the one comparable to the single project the journaled domain occupies. Read that way, persistence cost 204 lines added and 9 removed. A reader who counts only the pure rule model should read it as 56 added and 5 removed, and that is the figure worth arguing with, being the part no choice of layout can move outside the domain. Both are reported, either reading is available without recomputing anything, and the choice changes no other measurement

in the table — the stand-ins are required by whichever project owns the port. On both readings the comparison is against nothing.

The rival design that would dissolve that number, and what it costs instead. There is a standard move in orthodox practice that would make the 56 lines an artifact of one implementation choice rather than a property of the arrangement: persist through a separate persistence model and a mapper on the adapter side, so the rule model never learns that anything is stored. If that were available here, the paper’s residual claim would fall with it. The reason it is not available lies in the model rather than in the choice, and it can be read off the code. `Pile` is `internal sealed` with a `private` constructor (`domain/Pile.cs:21, :28`); the only construction paths it offers are an empty pile and the landing transition (`:35, :56`); and it enforces one invariant by construction, that a pile never retains a complete row (`:16`).

A mapper in another assembly therefore has three ways in, and each lands somewhere. It can be handed a factory the domain authors, which is a change to the rule model. It can be granted `InternalsVisibleTo` — which reaches `internal` members but not `private` ones, so it still needs a domain-authored factory, and the grant is itself a line inside the domain naming a persistence assembly. Or it can use reflection, which reaches a private constructor and is what mainstream object-relational mappers in fact do; that route asks nothing of the domain, and it bypasses every check the constructors perform.

The baseline took the first road, and its own source records why. Its rule model grew a factory, `internal static Pile Of(int width, ImmutableHashSet<Position> cells)`, whose comment calls it the seam a state port needs to put a saved pile back, added for the fourth staging because “the only way in was through the model” (`baseline-hex/domain/model/Pile.cs:51`). It grew a restore constructor whose comment names the cost outright — the point at which persistence reached the model (`baseline-hex/domain/model/Well.cs:75`). That constructor re-asserts the invariants on the way in, the complete-row check among them (`baseline-hex/domain/model/Well.cs:305, :314`). Its assembly grants `InternalsVisibleTo` to its test suite and nothing else (`baseline-hex/domain/AssemblyInfo.cs`), and both new members are called from inside the domain project (`baseline-hex/domain/model/Well.cs:86`, `baseline-hex/domain/application/GameService.cs:197`), so the mapper never touches the model. The separate-persistence-model design is what this baseline does; it still required the seam.

Stated no more strongly than it should be: the claim is not that 56 lines is the necessary price. A model with public constructors and no invariants would pay less, and would have less to protect. The claim is that where a rule-bearing type is closed and enforces something, putting a saved instance back is a domain-side act — pay for it in the rule model, as the baseline did, or reach past the constructors and forgo the check.

The journaled side does neither, and the reason is qualitative rather than a count. Reconstitution there is re-performing the recorded acts through the domain's own operations, so **Integrate** runs again on the way back in and the invariant holds on the reconstituted pile exactly as it held on the original — no seam, and nothing bypassed. A mapper reconstitutes *around* the invariant or is let in through a door the domain opens; a replay reconstitutes *through* it. What is not measured is the third road: no reflection-based mapper variant was built, so that comparison is an argument with checked premises rather than a count, and building it is the next measurement this baseline wants.

One concession runs the other way, and belongs here rather than in a footnote. The build graph does not distinguish the arrangements: a ported rule model is also free of project references, packages, and framework names, so “zero references, zero packages” should stop carrying comparative weight — it distinguishes a clean domain from a dirty one, not this arrangement from that one.

A second concession was drafted beside it and did not survive checking. It read that the ported arrangement is cleaner on one point, needing no counterpart to §2's empty public anchor type. But the framework's seam takes an *assembly*, not a type (`Choreography/Theater/PerformanceV2.cs:60`), and reads one with `GetTypes()` under a filter that admits internal types (`Puppeteer/EventSourcing/DomainLibraries.cs:136,151`); the framework's own CLI loads domain libraries by full path (`PuppeteerCli/AttachCommand.cs:309`). Rebuilt with the anchor deleted, the domain exposes zero public types, and a host holding no reference to it, loading it by path, seeds a **Well** and spawns a piece to the same reported state. The anchor is this example's ergonomics and not the arrangement's cost, so the asymmetry claimed for it is withdrawn.

What remains unmeasured is distribution. The two `StageManager` stagings of Lab A and the three machines of Lab B have no ported counterpart here, deliberately: a hexagon has no distribution story to port, so building one would mean inventing an architecture rather than measuring one. Lab B's zero therefore stands uncontested by measurement rather than confirmed against a rival.

Lab G — re-decomposition. Every other lab holds the domain fixed and varies what surrounds it. This one does the opposite: it cuts the domain's own internal boundary while the stagings stay where they are. The **Well** was split into two roles, one owning the settled pile and one owning the falling piece; the original's record was then read — by the same introspection read that rehydration performs, of each action and its parameters — and the new roles were driven to perform their own. No journal was cut, transformed, or rewritten: the original's acts are identical before and after, with nothing added or removed, and both new records are append-only with contiguous entry identifiers, each holding only its own role's verbs.

The exercise runs over a whole game rather than a prefix of one: 129 acts to game over, 29 landings, 136 entries, every emitted act present in the record. Both roles reach the **Well**'s exact final state, game over included — which is the interesting part, because the re-cut moved the authority for *game over* to the pile role, and replaying the record into the two reproduces it with both agreeing. Two things make that a fair test rather than a private one. The framework's own rehydration of the same journal reaches the same state, so the record is a record in the ordinary sense and the lab's reading took nothing special out of it. And the full length is available at the substrate commit this paper cites, an earlier one having stopped the readable record short of the game's end.

What was measured. Eleven of the twelve hosts were untouched while the domain divided beneath them; the twelfth needed one line. The invariant held over 60 games and 47,783 steps compared against the single-**Well** version — 9,265 landings, 3,336 rows cleared, 39 games ended, no divergence — and 47,783 overlap checks that crossed the new boundary, asking each role separately, found no overlap; the roles never disagreed about a game being over. Those figures were re-run on the corrected example at the substrate commit this paper cites, and reproduced to the digit, so the equivalence result was never sensitive to which state the example sat in; only the entry counts were. A random player cleared no rows at all in twenty games, so two further playing policies had to be written before the collapse was exercised. Nothing became speculative, and two checks came out stronger than they had been.

What it cost. This is the one lab in which the domain does change, by design: **Well.cs** was untouched, **Pile.cs** took two lines, and 292 lines of new domain were written. Movement costs no cross-role message; a landing costs two tells and two acknowledgements. The record costs 2.32 times as much — 316 entries across the two roles against 136 for the same game — and almost all of the difference sits exactly where the two roles have to speak. Of the 180 extra entries, 174 are the 29 landings at six entries each: a landing writes the piece's handover, the pile's absorption, and a message and acknowledgement each way, against the one act the undivided board wrote. One more is the second seed, since two roles open two aggregates. The remaining five are one-time declarations — eleven templates across the two roles against six for the one, each written once however many times its act recurs. Entry counts depend on how verbs are recorded, and these were re-taken after every command in the example became a parametrized Action: a template plus a compact argument per act, where a literal script wrote one full sentence per call. That encoding is why the original comes to 136 rather than the 130 an earlier run reported, and the ratio has moved with it, from $2.38\times$ through $2.29\times$ to $2.32\times$. What the encoding does not touch is the shape of the finding: 129 acts, 29 landings, six entries per landing, and no divergence. A fourth state appeared, because landing is now two acts by two parties rather than one; a guard had to be restated; and a cell-set serialization moved inside the domain, a tell carrying ordered scalars. One further consequence is larger than a cost, and §8.4 states what it amounts to; the premises are here, each of them checked. The undivided board built a complete frame by unioning the pile's

cells with the falling piece's and clipping to the interior, a derivation performed inside the domain (`domain/Well.cs:322`). A projection on the emitting plane is produced by a reaction belonging to one actor, running on that reaction's thread once that actor's read lock is released (`Puppeteer/IOutputSink.cs:106`), so it reaches that actor's state and no other's. After the cut those two halves sit in two actors, and neither role can push a whole frame. A third premise is a fact about this lab's own artifact rather than about the substrate, and is why nothing about the loss was measured: the split actor exposes no sink, no formatter and no output wiring, so there was no push channel here to lose in the first place. The reader that would join the two records was not built and is not claimed.

The finding nobody anticipated. The record is not neutral between the two new roles. Every act in the **Well**'s journal turned out to belong to the piece role, and the pile role's own act — absorbing landed cells — appears nowhere in the original, because under the **Well** absorbing was never an act at all but a private consequence of a tick. The pile role's record therefore does not translate; it is generated, by the piece role performing and telling. So a re-decomposition does not divide an existing record between two roles: one inherits the original's voice, and the other's record is born as a consequence of it. Whether that holds in general, a single case cannot say.

Lab H — the same routine, recognized on three stagings. One reaction is defined outside the domain: it seeks a spawn, then the drop that ends the piece's descent, and what it matches between them is a placement. It is then run against three stagings — the journal of a single process; the journal a warm server keeps while a client drives it over a named pipe; and the journals of the three containers, read inside each node against its own store. It recognizes the same two placements, in the same order, in all three. The three containers' journals are byte-identical, and the domain's diff for the exercise is empty: a reaction is defined outside the domain and the domain learns nothing of it.

Three measured qualifications matter more than the confirmation. There is **no correlation handle** in this record: a spawn names a piece type and a drop names nothing, so the trajectory is tied by order alone. There is **no landing act**: the board lands a piece privately, inside a tick or a drop, so a piece that comes to rest under gravity leaves its opening unclosed and the next piece's drop closes it — two placements matched at one closing entry, with the count right and the correlation wrong, and nothing to signal it. And **entry identifiers are not staging-invariant**: the three-node staging writes three seeding acts where one process writes one, shifting everything after, while the closing identifier is the only handle a match hands its reader. The trajectory is short — two placements — so this is an existence result about recognizability, not a measurement of it at scale.

Lab I — the domain grows. A score and a difficulty level were added to the board: the first computable by any client and therefore required to be single-valued, the second a rule over lines already recorded. The growth was additive in the strict sense. Over the whole domain directory the change is 98

lines added and 3 removed, across three files rather than the two new ones alone: `Scoring.cs` at 39 lines and `Difficulty.cs` at 38, plus 21 added lines in `Well.cs`. Of the 98 added, 30 are code; the other 68 are the domain's own documentation convention — 62 doc-comment lines and 6 blank — and a reader should not take them for filler or for logic. All 3 removed lines are doc-comment prose in `Well.cs`. No code was deleted, no signature changed and no verb removed, established by diffing a sorted inventory of the board's members rather than by inspection. The build graph came out byte-identical: no packages, no references, and the two new files declare no imports whatsoever. The domain's suite went from 44 tests to 60, the original 44 untouched.

The measurement this paper had never taken is the converse ripple. When the domain grew, all twelve host projects kept running with no edit — verified by running them before and after from a throwaway worktree at the pre-change commit, not merely by compiling them. The cost of *adopting* the new capability then fell only where it was wanted, and the twelve divide exactly: nothing for four of them, one line for five, four lines for each of the two browser hosts, and thirty-two for the input host, the only place where the level changes what happens. Those figures are the sum over both growth commits and both adoption commits; taking only the first two understates the input host, which acquires its thirty-two lines in the adoption commit rather than when the level was added. That last is where the boundary this paper predicted becomes visible: the domain owns what level the game is at, the staging owns how fast gravity ticks, and the domain still holds no clock — which a reflection test now enforces.

One finding is a caution rather than a result, and it applies to every replay measurement including this paper's. Three of the eight replay fixtures did not *fail*: they answered. A game recorded as having cleared twenty rows replayed as having cleared two; one recorded at six replayed at four; and the third, recorded at zero, replayed at zero — correct, because the acts it lost happened to clear nothing, so the truncation was invisible in exactly the scalars a reader checks. The cause was a defect in the substrate's index rather than anything about the model, and it has since been fixed upstream, this lab being one of two independent reports that found it. The lesson survives the fix, and §8.4 states it as a rule: a replay checked only for internal consistency will pass while being wrong, so this lab now carries the state recorded at play time and asserts against it. That follows from what fusion means and is close to a tautology; no fused version was built. Until that exists, the comparison in this paragraph is an argument the paper makes, not a result it reports.

The nine laboratories' source and datasets accompany this paper as `paper09-data.zip`, laid out as the earlier papers in this series lay theirs out: one directory per lab under `labs/paper09-labX-...` holding its source and a README that states what it measures, what to run and which figure of this paper it produces, and a directory of the same name under `data/paper09-labX-...` holding its outputs and its write-up. A suite-level `labs/paper09-README.md` answers what the per-lab files cannot: that each lab stands alone and none de-

depends on another, which one to run if only one is run, the order by cost, and the one interaction that can mislead — Lab A measures that the domain did *not* change while Labs G and I change it deliberately, so running Lab A’s diff on their branches reports a difference that is a different measurement rather than a failed one. Two of the nine have no source of their own and say so in their README: Lab A, whose five stagings *are* hosts of the example and whose evidence is the empty diff across them, and Lab H, whose reaction is quoted in its write-up. Lab F is the largest, carrying the ported baseline in full. The labs are count-based, so no large raw sample log is produced or omitted, and every count cited above appears in full in the corresponding write-up.

Assembling that archive was itself a correction. The labs had been built on separate branches of the examples repository and never merged, so the material this section promises had been scattered across seven branches and reachable only by knowing which; a reader could not have obtained it from the repository named above. It is now consolidated, which is what makes the promise in the previous paragraph one a reader can act on.

One correction made during this work belongs in plain sight, because it touched the example, and an earlier draft of this appendix reported it as done when the artifact this paper deposits did not yet carry it. The frame each host pushes reaches its sink through a reaction over the journal, and a reaction over the domain observes recorded *operations*, never literal script text — a substrate rule, and a deliberate one. The example’s actor had been emitting its verbs as literal text. So the push channel never engaged: no frame was written at all, which cost three labs their evidence rather than merely changing its form. The player hosts were unaffected, because each renders from its own query of the board; the *viewers* were the loss — the client-authored column-height view of §3 had no frame to read, the live viewer of Lab D had nothing to repaint, and each node of Lab B wrote no frame of its own.

The correction is now in the deposited artifact, and it is larger than the one verb an earlier draft described. An operation is recorded as an operation only when it carries an argument, so the seed carries the board’s dimensions and a spawn carries its piece; the five verbs that move a falling piece carry nothing that varies — there is no value in *move left* — and each therefore carries the name of the act it performs. The seam the actor speaks to no longer offers an unparametrized form at all, which makes the omission unrepresentable rather than merely discouraged, and the same conversion was applied to the two-role actor of Lab G. Two consequences are reported where they fall: entry counts move, because a template written once and a compact argument per act is a different encoding from one full sentence per call (§8.2), and Lab G’s harness had been recovering the board’s dimensions by reading them out of the seed’s text, so it now reads them from the seed’s arguments as it already did for a spawn.

The correction was to the example’s actor, not to the domain: the domain’s diff across the entire period, including this fix, is empty — which is the invariance

this appendix reports, holding through a change to the machinery around it.

Code provenance

Source-code references in this paper are paths in **this repository**, and that is the one thing about provenance worth stating plainly: the example the labs measure is vendored here, at `labs/paper09-example/`, so the paper and the artifact it cites are one git history and there is a single commit to mark rather than two to reconcile. Example paths are cited relative to that directory — `domain/Well.cs:322`, `actor/IGameHost.cs:14` — and the ported baseline of Lab F, being a second and different arrangement, keeps its own prefix: `baseline-hex/...` under `labs/paper09-labF-porting-baseline/`. Line anchors resolve at the commit this paper is deposited from, and every one of them was re-resolved there before deposit rather than carried forward from the draft — nineteen in the example and the ported baseline, five in the framework.

One thing a vendored copy cannot carry is a history, and §8 rests on one: the dates that separate the domain’s last commit from the stagings added twenty-four days later. That chronology is checkable only where it happened, in the repository the example was written in (github.com/alvaroNCubo/puppeteer-examples), on `main` at `ec9a2d4`: the domain’s last commit is `fd8d94b`, the three additions of 23 July are `33f151c`, `8d7536e` and `974f62a`, and `git diff fd8d94b..main -- Tetris/domain/` is empty across the fifteen commits between. The copy here has one commit, the one that placed it, which is why Lab A of Appendix A is the single lab a reader cannot run against this repository and why its output is captured in `data/` for a reader who would rather not clone anything.

Framework source is cited in five files and nowhere else, and only where what is asserted is a property of the substrate that the example cannot establish: the scope of the emitting plane, which §8’s constraint on admissible decompositions turns on (`Puppeteer/IOutputSink.cs`, `Puppeteer/StageHook.cs`); the attachment of a destination to a live actor, which §9 uses to say when a destination is bound (`Choreography/Theater/PerformanceV2.cs`); and what a domain assembly must expose for the framework to read it, which Lab F’s withdrawn concession turns on and which an example carrying an anchor cannot settle by itself (`Puppeteer/EventSourcing/DomainLibraries.cs`, `PuppeteerCli/AttachCommand.cs`). Those five anchors resolve against the **public** Puppeteer repository (github.com/alvaroNCubo/puppeteer) at commit `2673d57`, which is also the minimum the labs’ engine requirement names. An earlier draft cited them against `dd67047`, a commit in the private Azure repository the runtime is developed in — a reference no reader could resolve, and the substantive provenance defect this pass fixed. Every other anchor in this paper is in the example, and every measurement reported is taken on the example rather than on framework internals. The example also contains two clients that no lab here counts — a webcam gesture input and an emulated constrained de-

vice — left out because their hardware legs were never run, so there is nothing measured to report about them.

One thing about the labs’ provenance belongs here rather than beside each count, and it is more specific than a general disclaimer. This draft’s measurements were taken while the example was still moving, and the labs do not all sit at one state of it: some were taken after the actor correction reported at the end of Appendix A and some before it. That correction changes how verbs are recorded, so it changes any figure computed from entry counts. One such figure existed — the record ratio of Lab G — and it has since been re-taken on the corrected example at the substrate commit this paper cites, which moved it and left everything else in that lab standing. No figure now reported comes from an uncorrected state. The correction never touched the figures the paper’s argument rests on: a diff over a source directory, a count of project references, a count of hosts that needed editing, and a comparison of final board states are all indifferent to how an act was written into the journal. Every lab is nonetheless to be re-verified at the single published commit and reported at that commit as part of the deposit, and where a count is sensitive to the example’s state the appendix says so rather than leaving a reader to reconcile dates.

A word on what is and is not public, since the two are easy to run together. The commit pinned above is in that public repository: `2673d57` is its `main`, so a reader who visits it today finds the source the five framework anchors resolve against. That snapshot is archived in Software Heritage as well, as the preceding paper in this series is, so the anchors outlive the repository that hosts them:

```
swh:1:dir:5f0a48469f6e23a6bbdf8cc63938099b632785c5;  
  origin=https://github.com/alvaronCubo/puppeteer;  
  anchor=swh:1:rev:2673d57445fad8e85cb950b1de62d7c57c14c3fc
```

One cleanup was weighed for the deposit and deliberately not made: the domain’s assembly carries an `InternalVisibleTo` grant to the console host which no host now exercises (Lab E reports it). Withdrawing it would rest the fence on the test suite’s grant alone — tidier, and a change to what Lab E measured — so it stays as reported. The example’s own anchors need no repository visit at all: they are in the accompanying archive, `paper09-data.zip`, whose contents are listed above. The three-machine lab reproduces from a single script, `labs/paper09-labB-three-machines/docker/run-demo.sh`, against a local container runtime. The labs are count-based, so no large raw sample log is produced or omitted.

Acknowledgments

The author used large language models (including Claude and ChatGPT) as editorial assistants for language refinement, structural feedback, and literature navigation. All original ideas, terminology, theoretical constructs, and technical content presented in this work are solely the author’s.

One disclosure belongs with the labs rather than with the writing, because a file it produced is cited above as evidence. The automated player of Lab C is an instance of a large language model, driving the game over a pipe; the view it reads the board through — the column-height projection of §3 (`tools/pile-scan.ps1`) — was written by that instance during the lab, with the author’s permission, to suit the form in which it reasons. That the adapter was authored by the client rather than by the author of the domain is not incidental to the argument. It is the sharpest case available of a projection belonging to the party that reads it, and it is reported here so that a reader can weigh the evidence knowing how it came about.

References

Entries without a refereed venue — practitioner articles, talks, screencasts, and technical reports — are given with their access date and a permanent archival snapshot, and are marked as such where the argument leans on them; where such a source sits behind a paywall, a freely available presentation of the same idea is listed beside it.

The letter suffixes on this series’ own entries denote the paper’s number in the series, not the order of appearance here — **2026c** is Paper 3, **2026d** Paper 4, **2026g** Paper 7, **2026h** Paper 8 — so that a given paper carries the same label wherever in the series it is cited. Letters absent below belong to papers this one does not cite.

Agha, G. (1986). *Actors: A model of concurrent computation in distributed systems*. MIT Press.

Apel, S., Batory, D., Kästner, C., & Saake, G. (2013). *Feature-oriented software product lines: Concepts and implementation*. Springer. <https://doi.org/10.1007/978-3-642-37521-7>

Armstrong, J. (2003). *Making reliable distributed systems in the presence of software errors* [Doctoral dissertation, Royal Institute of Technology (KTH), Stockholm].

Bernhardt, G. (2012a). *Boundaries* [Conference talk, Software Craftsmanship North America; practitioner source, no refereed venue, freely available]. Destroy All Software. <https://www.destroyallsoftware.com/talks/boundaries> (accessed 26 July 2026; archived at <https://web.archive.org/web/20260703082615/https://www.destroyallsoftware.com/talks/boundaries>)

Bernhardt, G. (2012b). *Functional core, imperative shell* [Screencast; practitioner source, no refereed venue, behind a paywall — see 2012a for the freely available presentation of the same idea]. Destroy All Software. <https://www.destroyallsoftware.com/screencasts/catalog/functional-core-imperative-shell> (accessed 25 July 2026; archived at <https://web.archive.org/web/20260703082615/https://www.destroyallsoftware.com/screencasts/catalog/functional-core-imperative-shell>)

<https://web.archive.org/web/20260723235945/https://www.destroyallsoftware.com/screencasts/catalog/function-core-imperative-shell>)

Brooks, F. P., Jr. (1987). No silver bullet: Essence and accidents of software engineering. *Computer*, 20(4), 10–19. <https://doi.org/10.1109/MC.1987.1663532>

Clements, P., & Northrop, L. (2001). *Software product lines: Practices and patterns*. Addison-Wesley.

Cockburn, A. (2005). *Hexagonal architecture* [Practitioner article, no refereed venue]. <https://alistair.cockburn.us/hexagonal-architecture/> (accessed 25 July 2026; archived at <https://web.archive.org/web/20260708163140/https://alistair.cockburn.us/hexagonal-architecture>)

Codd, E. F. (1970). A relational model of data for large shared data banks. *Communications of the ACM*, 13(6), 377–387. <https://doi.org/10.1145/362384.362685>

Eugster, P. T., Felber, P. A., Guerraoui, R., & Kermarrec, A.-M. (2003). The many faces of publish/subscribe. *ACM Computing Surveys*, 35(2), 114–131. <https://doi.org/10.1145/857076.857078>

Evans, E. (2003). *Domain-driven design: Tackling complexity in the heart of software*. Addison-Wesley.

Fowler, M. (2004). *Inversion of control containers and the dependency injection pattern* [Practitioner article, no refereed venue]. <https://martinfowler.com/articles/injection.html> (accessed 25 July 2026; archived at <https://web.archive.org/web/20260722213553/https://martinfowler.com/articles/injection.html>)

Fowler, M. (2005). *Event sourcing* [Practitioner article, no refereed venue]. <https://martinfowler.com/eaDev/EventSourcing.html> (accessed 25 July 2026; archived at <https://web.archive.org/web/20260714065500/https://martinfowler.com/eaDev/EventSourcing.html>)

Garlan, D., & Notkin, D. (1991). Formalizing design spaces: Implicit invocation mechanisms. In *VDM '91: Formal Software Development Methods*, Lecture Notes in Computer Science (Vol. 551, pp. 31–44). Springer. https://doi.org/10.1007/3-540-54834-3_5

Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.

Helland, P. (2005). Data on the outside versus data on the inside. *Proceedings of the 2nd Biennial Conference on Innovative Data Systems Research (CIDR '05)*, 144–153. (Reprinted in *Communications of the ACM*, 63(11), 111–118, 2020. <https://doi.org/10.1145/3410623>)

Helland, P. (2015). Immutability changes everything. *Proceedings of the 7th Biennial Conference on Innovative Data Systems Research (CIDR '15)*. (Also published in *ACM Queue*, 13(9), 2015,

- <https://doi.org/10.1145/2857274.2884038>, and in *Communications of the ACM*, 59(1), 64–70, 2016, <https://doi.org/10.1145/2844112>)
- Hewitt, C., Bishop, P., & Steiger, R. (1973). A universal modular ACTOR formalism for artificial intelligence. *Proceedings of the 3rd International Joint Conference on Artificial Intelligence (IJCAI)*, 235–245.
- Huang, P., Guo, C., Zhou, L., Lorch, J. R., Dang, Y., Chintalapati, M., & Yao, R. (2017). Gray failure: The Achilles’ heel of cloud-scale systems. *Proceedings of the 16th Workshop on Hot Topics in Operating Systems (HotOS ’17)*, Whistler, BC, Canada. <https://doi.org/10.1145/3102980.3103005>
- Johnson, R. E., & Foote, B. (1988). Designing reusable classes. *Journal of Object-Oriented Programming*, 1(2), 22–35.
- Kaldor, J., Mace, J., Bejda, M., Gao, E., Kuropatwa, W., O’Neill, J., Ong, K. W., Schaller, B., Shan, P., Viscomi, B., Venkataraman, V., Veeraraghavan, K., & Song, Y. J. (2017). Canopy: An end-to-end performance tracing and analysis system. *Proceedings of the 26th Symposium on Operating Systems Principles (SOSP ’17)*, 34–50. <https://doi.org/10.1145/3132747.3132749>
- Kleppe, A., Warmer, J., & Bast, W. (2003). *MDA explained: The model driven architecture — Practice and promise*. Addison-Wesley.
- Kleppmann, M., & Kreps, J. (2015). Kafka, Samza and the Unix philosophy of distributed data. *IEEE Data Engineering Bulletin*, 38(4), 4–14.
- Kratzke, N., & Quint, P.-C. (2017). Understanding cloud-native applications after 10 years of cloud computing — A systematic mapping study. *Journal of Systems and Software*, 126, 1–16.
- Lamport, L. (1978). Time, clocks, and the ordering of events in a distributed system. *Communications of the ACM*, 21(7), 558–565. <https://doi.org/10.1145/359545.359563>
- Laurel, B. (1991). *Computers as theatre*. Addison-Wesley. (2nd ed., Addison-Wesley Professional, 2013.)
- Mace, J., Roelke, R., & Fonseca, R. (2015). Pivot tracing: Dynamic causal monitoring for distributed systems. *Proceedings of the 25th Symposium on Operating Systems Principles (SOSP ’15)*, 378–393. <https://doi.org/10.1145/2815400.2815415>
- Martin, R. C. (2017). *Clean architecture: A craftsman’s guide to software structure and design*. Prentice Hall.
- Mellor, S. J., & Balcer, M. J. (2002). *Executable UML: A foundation for model-driven architecture*. Addison-Wesley.
- Meyer, B. (1988). *Object-oriented software construction*. Prentice Hall.
- Moseley, B., & Marks, P. (2006). *Out of the tar pit* [Presented at Software Practice Advancement (SPA 2006); no formal proceedings].

- <https://curtclifton.net/papers/MoseleyMarks06a.pdf> (accessed 25 July 2026; archived at <https://web.archive.org/web/20260718134750/https://curtclifton.net/papers/MoseleyMarks06a.pdf>).
- Murphy, G. C., Notkin, D., & Sullivan, K. (1995). Software reflexion models: Bridging the gap between source and high-level models. *Proceedings of the 3rd ACM SIGSOFT Symposium on the Foundations of Software Engineering (SIGSOFT '95)*; *ACM SIGSOFT Software Engineering Notes*, 20(4), 18–28. <https://doi.org/10.1145/222132.222136>
- Murphy, G. C., Notkin, D., & Sullivan, K. J. (2001). Software reflexion models: Bridging the gap between design and implementation. *IEEE Transactions on Software Engineering*, 27(4), 364–380. (The extended journal version of the 1995 paper, and the canonical citable form.)
- Overeem, M., Spoor, M., & Jansen, S. (2017). The dark side of event sourcing: Managing data conversion. *Proceedings of the 24th IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER '17)*, 193–204. <https://doi.org/10.1109/SANER.2017.7884621>
- Overeem, M., Spoor, M., Jansen, S., & Brinkkemper, S. (2021). An empirical characterization of event sourced systems and their schema evolution — Lessons from industry. *Journal of Systems and Software*, 178, 110970.
- Parnas, D. L. (1972). On the criteria to be used in decomposing systems into modules. *Communications of the ACM*, 15(12), 1053–1058. <https://doi.org/10.1145/361598.361623>
- Passos, L., Terra, R., Valente, M. T., Diniz, R., & Mendonça, N. (2010). Static architecture-conformance checking: An illustrative overview. *IEEE Software*, 27(5), 82–89.
- Perry, D. E., & Wolf, A. L. (1992). Foundations for the study of software architecture. *ACM SIGSOFT Software Engineering Notes*, 17(4), 40–52. <https://doi.org/10.1145/141874.141884>
- Peyton Jones, S., & Wadler, P. (1993). Imperative functional programming. *Proceedings of the 20th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '93)*, 71–84. <https://doi.org/10.1145/158511.158524>
- Rivera, A. (2026c). Reactions and the partition: opt-in eventual consistency in actor-native systems. *Puppeteer Papers Series*, Paper 3 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20792156>
- Rivera, A. (2026d). Preserving semantic continuity across actors: a tell-based approach without orchestration. *Puppeteer Papers Series*, Paper 4 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21207062>
- Rivera, A. (2026g). After the substrate: building software without a datacenter. *Puppeteer Papers Series*, Paper 7 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20398998>

- Rivera, A. (2026h). Inference without Authority: the three authorities that govern an output. *Puppeteer Papers Series*, Paper 8 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21499637>
- Sambasivan, R. R., Shafer, I., Mace, J., Sigelman, B. H., Fonseca, R., & Ganger, G. R. (2016). Principled workflow-centric tracing of distributed systems. *Proceedings of the 7th ACM Symposium on Cloud Computing (SoCC '16)*, 401–414. <https://doi.org/10.1145/2987550.2987568>
- Shaw, M. (2003). Writing good software engineering research papers. *Proceedings of the 25th International Conference on Software Engineering (ICSE '03)*, 726–736.
- Sigelman, B. H., Barroso, L. A., Burrows, M., Stephenson, P., Plakal, M., Beaver, D., Jaspan, S., & Shanbhag, C. (2010). *Dapper, a large-scale distributed systems tracing infrastructure* (Technical Report dapper-2010-1) [Technical report, no refereed venue]. Google Inc. <https://research.google/pubs/dapper-a-large-scale-distributed-systems-tracing-infrastructure/> (accessed 25 July 2026; archived at <https://web.archive.org/web/20260705194849/https://research.google/pubs/dapper-a-large-scale-distributed-systems-tracing-infrastructure/>)
- Stol, K.-J., & Fitzgerald, B. (2018). The ABC of software engineering research. *ACM Transactions on Software Engineering and Methodology*, 27(3), Article 11. <https://doi.org/10.1145/3241743>
- Sullivan, K. J., & Notkin, D. (1992). Reconciling environment integration and software evolution. *ACM Transactions on Software Engineering and Methodology*, 1(3), 229–268.
- Terra, R., & Valente, M. T. (2009). A dependency constraint language to manage object-oriented software architectures. *Software: Practice and Experience*, 39(12), 1073–1094. <https://doi.org/10.1002/spe.931>
- Terry, D. B., Demers, A. J., Petersen, K., Spreitzer, M. J., Theimer, M. M., & Welch, B. B. (1994). Session guarantees for weakly consistent replicated data. *Proceedings of the Third International Conference on Parallel and Distributed Information Systems (PDIS '94)*, 140–149.
- Waldo, J., Wyant, G., Wollrath, A., & Kendall, S. (1994). *A note on distributed computing* (Technical Report SMLI TR-94-29). Sun Microsystems Laboratories. (Reprinted in J. Vitek & C. Tschudin, Eds., *Mobile object systems: Towards the programmable internet*, Lecture Notes in Computer Science, Vol. 1222, pp. 49–64. Springer, 1997. https://doi.org/10.1007/3-540-62852-5_6)
- Wiggins, A. (2011). *The twelve-factor app* [Practitioner source, no refereed venue]. <https://12factor.net> (accessed 25 July 2026; archived at <https://web.archive.org/web/20260724222316/https://12factor.net/>)
- Young, G. (2010). *CQRS documents* [Practitioner source, no refereed venue]. https://cQRS.files.wordpress.com/2010/11/cQRS_documents.pdf (accessed 25

July 2026; archived at <https://web.archive.org/web/20240317190715/https://cqrs.files.wordpress.com/2010/11>

The Assembled Verb: repertoire, assembly, and the authorship of a verb

TL;DR

A(); B(); C(); consumes three capabilities and leaves three. This paper measures an arrangement where assembly *produces* one: the composition becomes a verb of a subject — one definition in the subject’s record, exercised many times by reference — and the paper’s instrument is a question askable of any record, which separates that from a coordinator’s record, a durable-execution history and a trace. Where such a unit is missing, the continuity still has to go somewhere, and shipped software authors it into the pieces: §6 measures those compensations, and a census of 72 named artifacts finds no name for a trajectory spanning subjects. So the consequence the series takes forward is not that composition became first-class, but where continuity may live: once the composition has a place of its own, it no longer has to be authored into its constituents — and the constituents need not know the future wholes in which they will participate.

Dependencies. This paper is part of the Puppeteer Papers, a series of self-deposited preprints, and rests on four of them: the actor’s speech and tell (Paper 4), whose two directions supply the assembly plane’s own statements in §5.1; the output authority (Paper 8), whose producer-is-not-the-authority geometry §2.4 meets again on the plane of historicity; Reaction (Paper 3), whose per-verb consistency contract §2.4’s movement parallels and whose mechanism §8’s corollary draws on; and identity across stagings (Paper 9), whose corpus-selection and anchor-verification discipline this paper follows. The methodological position — what kind of contribution this is, and the boundary against the nearest alternative — is stated in §1.2, beside the evidence it classifies.

1. The Ordinary Ground

Start with the most ordinary thing in programming. A domain offers operations, and a program uses them:

```
A();  
B();  
C();
```

Nothing here is new, and nothing is claimed about it. Ask the equally ordinary question:

What did that program just do?

There are good answers. One can enumerate the operations — it did A, then B, then C. One can put a method around them, and then the program did `CompletePurchase()`:

```
CompletePurchase() {  
    A();  
    B();  
    C();  
}
```

One can connect them with events instead of calls, or put a saga or a process manager in charge of the order, or spread them over services. All of these achieve the same thing, which deserves a plain name: **movement**. The operations run, in an order, with values flowing between them.

This paper concedes movement in full, and the concession is not decorative. Two open, independently authored reference systems were taken as a corpus — a widely used commerce reference implementation (.NET Foundation 2026) and a widely cited modular monolith (Grzybek 2024), both selected before any measurement for being well-regarded implementations of the design approach they follow (domain-driven design; Evans 2003). They achieve movement in several of the ways just listed, and they achieve it well. Nothing in this paper improves on how they move, and no engineering incapacity is being solved.

1.1 Two arrangements, side by side

Grant the movement. Now place a second arrangement beside the first:

arrangement one	arrangement two
-----	-----
A();	V := A(); B(); C();
B();	V(x);
C();	V(y);

and ask the question this paper is about:

What exists in the second arrangement that did not exist in the first?

Work through the candidate answers, because each elimination is doing real work.

Not *the sequence* — both arrangements have it. Not *the name* — `CompletePurchase()` can have a better one. Not *the coordination* — both can be coordinated, by any of the corpus’s mechanisms. Not *the result* — both can produce exactly the same effects. Not even *the composed act* — both can execute exactly the same operations, in the same order, with the same values.

What the second arrangement offers — and the paper must now earn rather than assert — is this: **a reusable capability whose constitution is A, B and C**. In the first arrangement the composition is consumed by the run that performs it. In the second it exists *between* its uses — $V(x)$ and $V(y)$ refer to it — and it remains afterwards. (A method can claim the same on the program-text side, and the claim is not dismissed in passing: §2.5 raises it and §3 decides it.)

before	A	B	C	
assembly	$V := A; B; C$			
after	A	B	C	+ V

The program of arrangement one *used* three capabilities. Assembly *produced* one. That is not a defect of arrangement one — nothing about it is supposed to produce a capability, and no system in the corpus is failing at a task it undertook. It is a difference in kind, and once it is on the table a second question becomes unavoidable, one level down from the first: **whose capability is V?** A capability, unlike an execution, has to live somewhere. §4 answers that; §2 and §3 first establish that the thing needing an owner exists.

The difference in kind also carries a quiet corollary, stated here and measured in §5: if V can be produced at assembly, then A, B and C did not need to anticipate it. **A repertoire need not know the verbs in which its operations will later participate.**

None of this is hypothetical, and §6 measures it in shipped software. Where the whole has nowhere to live as a unit, the movement is still needed — so it is authored into the pieces: events that carry an aggregate so a handler can fetch it again, correlation carried by a durable queue and a poller, coordinators holding the relation as state of their own. All of it real, all of it buying guarantees its system needs — and all of it work that exists because the composition has no place of its own. So the question underneath every section that follows is not whether the pieces can be tied together — today’s software is the standing proof that they can — but **where the continuity lives**, and what would have to exist for it to live somewhere other than inside the pieces. Stated as a promise, to be cashed at the end: the assembled verb will matter less because it adds one capability than because it gives continuity somewhere to live other than in

the capabilities it joins.

The paper’s claim, in one sentence, to be earned over the sections that follow:

Assembly does not merely execute a composition. It makes the composition available as a verb of a subject.

1.2 What is measured, what is conceded, and what is read

Seven bodies of evidence appear and answer different questions. Two are the corpus authors’ artifacts, four are constructed or measured by this paper, and one is not evidence of fact at all; a reader is owed every one of those differences.

	shows	kind
the corpus’s own domain tests	the operations compute on their own	observational
the corpus’s architecture	movement without a produced capability	observational
the substrate labs (§2–§5)	assembly producing one	author-constructed
the long scenario (§5)	the plane made visible across repertoires	author-constructed
the third-party rows and the prototype (§3.2)	the criterion’s cells, and the plane itself, on engines not this paper’s	author-driven, third-party substrate
the record’s prices (§8)	what the arrangement materially costs, against the direct baseline	author-measured
the calculus (Appendix C)	the separations as propositions with proofs	formal — construction, not evidence of fact

The constructed evidence proves realizability and nothing about how either corpus system was designed; no count from it is offered as evidence about their authors’ choices. The corpus was selected on stated criteria before it was read, and the census of §6 was run after the selection. No population-level generalisation is claimed anywhere.

The record’s cost is measured and nothing is ranked by it: §8 prices growth, replay and throughput against the direct baseline and claims no advantage anywhere.

Methodologically the paper is an analytic theory contribution in the sense of Gregor’s (2006) theory for analyzing (Type I): capability, exercise and constitution are constructs by which arrangements may be described and compared, and the paper offers no prescription — its instrument is the question of §3, askable

of a system already built. The distinctions themselves are stated operationally in Appendix C, over a core calculus small enough that each separation is a proposition with a constructive proof, and the labs double as its conformance checks. The labs of Appendix A are accordingly an existence proof of realizability — on two engines, one of them third-party — plus a measurement of the record’s material cost (§8); no advantage claim is made anywhere, and no design-science evaluation of benefit is attempted. That an artifact is built and measured invites the nearest alternative classification, so the boundary is stated checkably rather than trusted: a design-science contribution (Gregor’s Type V; Hevner et al. 2004; Peffers et al. 2007) states a problem, derives objectives of a solution, and evaluates the artifact’s utility against them, and this paper does none of the three, on purpose. What the labs evidence is therefore analytic, not evaluative: that the constructs denote — the category is instantiable, on more than one substrate — and what its instances measurably are. A reader who finds a utility claim anywhere in this paper has found an error.

The argument underneath runs in four steps, and the sections divide between *walking* it and *holding the ground it crosses*. §2 demonstrates the mechanism: V exists — defined once, exercised by reference, replayed from the record. §4 locates it: V belongs to none of the pieces, so a place other than the pieces exists. §5 cashes the deduction: two repertoires that never met participate in one verb of seventeen operations, unmodified — anticipation is not required. §6 shows the condition the corpus lives under where no such place exists: continuity authored into the pieces, measured. The rest holds ground: §3 verifies that the unit really resides where §2 put it; §2.4 establishes the modality property the conjunction of §7 needs; §4.2–§4.3 bound the subject’s position; §7 walks sixteen neighbour families. Necessary — the category would not survive scrutiny without them — and deliberately not the spine.

And one term is used before it is defined, deliberately. *Subject* carries, until §4, only its deflated sense — the thing an operation is attributable to. The strong sense is not in use until it is earned.

2. Assembly Produces a Capability

2.1 What the substrate provides

Two ordinary things. An **actor** (Hewitt, Bishop and Steiger 1973) holds state and is entered by performing acts on it. A **journal** records each act as it is performed. Neither idea is this paper’s, and no novelty is claimed for either.

One mechanical property of the record is the whole of what this section turns on, and it can be stated without any claim about what an entry *is*: **the record preserves enough of each performed operation that replay invokes it again**. Not enough to describe the operation, and not an outcome from which the operation could be inferred — enough to invoke it. The property is checkable by running the replay, and it is visible in the record’s two shapes:

a definition entry carries the statement text itself (`EventData.cs:49-52`), an invocation entry carries an integer identity and an argument string and nothing else (`EventData.cs:34-37`), and rehydration reads them back through one entry point (`ActorHandler.cs:3929`).

The domain assemblies used throughout are the corpus’s shipped binaries, unmodified, referenced as prebuilt binaries by a host their authors never anticipated. Nothing was edited to make any of this fit.

2.2 The entry

A composed verb is performed once — five domain operations and one register entry of the assembler’s own. This is what the journal contains afterwards, printed from it rather than transcribed:

```
define action 1 (payerId:string, country:string, period:string, amount:decimal,
                currency:string, street:string, city:string, state:string,
                zip:string, userId:string, userName:string, card:string,
                cvv:string, holder:string, sku:int, itemName:string,
                itemPrice:decimal, key:string) as
    payment    = payments.Buy(payerId, country, period, amount, currency);
    payment.MarkAsPaid();
    address    = Address(street, city, state, country, zip);
    welcomeKit = Order(userId, userName, address, 1, card, cvv, holder,
                        Now.AddYears(1), Null, Null);
    welcomeKit.AddOrderItem(sku, itemName, itemPrice, 0.0, '', 1);
    held.Keep(key, payment);
end;
```

Three facts about that entry, each about the record rather than an interpretation of it.

Its content is the composition. The entry does not merely correlate the operations, index them, or summarise their outcomes. It preserves the statements that constitute the definition, in their performed order, together with the values that flow among them.

The definition is not authored by the developer: the runtime emits it when a composed body is first performed (`DefineActionStatement.cs:25`).

It has no business name. The entry is `action 1` — an integer, because actions in this substrate carry no name at all: the invocation record has fields for an identity and an argument string, and no field a name could occupy (`EventData.cs:34-37`). It could as easily be `action 7` and everything in this section would hold. Whatever makes this entry one unit, naming is not it:

what makes the composition one unit	one definition
	an ordered set of operations
	the values flowing among them

what its status in the history rests on	the execution is attributed to the actor
	replay re-performs that definition

Replaying it re-performs the operations. Reading the entry back invokes the same statements, on the same repertoires, with the values the entry carries.

2.3 Reuse — the measurement that separates a capability from a receipt

Everything above is compatible with a weaker reading: the record holds a *receipt* of something that happened once. A capability is more than a receipt. **The test of a capability is reuse: it can be drawn on again without being restated.**

So the measurement (Appendix A: `TheCompositionBecomesACapabilityLab.cs`): the same composed verb is performed **twice**, with different arguments — different payer, amount, item, key. Then the journal is read, and then a fresh actor is built over it.

times the composed verb was performed	2
definitions of it in the journal	1
state after the live run	both purchases held
state after replaying the journal	both purchases held

The second use did not restate the composition. It **invoked** it: what travelled was the verb's identity and the new arguments. And a fresh actor over the same journal replays both invocations from the one recorded definition.

The seam where this happens is small enough to read (`ActorHandler.cs:2241-2255`): a body whose authored form is already known is journaled as an *invocation* of the existing action, identity plus fresh arguments; a body not yet known is journaled once as its *definition* and assigned the next integer identity. The branch is the mechanical difference between a capability and a copy.

An inline sequence can of course be executed twice — by writing it twice. What it cannot do is exist *between* the two executions as a thing the second one refers to. That existence-between-uses is what the figure of §1.1 asserts with its + V, and it is now exhibited in an inspectable record rather than asserted as notation:

Before assembly there were N capabilities. After assembly there are
N + 1, and the journal can be inspected for the +1.


The status of that sentence has to be stated exactly. *Capability* is operationalized here by the reuse test, and the arrangement measured implements that operationalization — so the N + 1 is not the confirmation of a hypothesis. It is two smaller things, each worth having. It is a **demonstration that the definition is realizable**: a record can hold the definition/exercise separation as data, which none of the ancestors examined in §7.1 establishes

for a *subject's own record*. And it is a **conformance check that could have failed**: had the second use re-emitted a second definition, the seam of `ActorHandler.cs:2241-2255` would have been broken and the lab would have counted two. What keeps the construct from being private to this paper is §7.1: *defined once, exercised many times* is not coined here — macro-operators, options and workflow cases hold it independently — and what this section adds to that inherited construct is only *where it lives*.

The measurement licenses a deduction stronger than the count. Its premises are measurements of this arrangement; only its *form* names no substrate:

1. A, B and C are available capabilities.
2. Assembly defines $V := A; B; C$.
3. V is invoked later, with values that did not exist when V was defined.
4. Two distinct invocations refer to the same V.
5. Replay preserves both the definition and the distinction between its exercises.


Therefore **V cannot be identified with any of its executions** — they differ and it does not. And V cannot be identified with A, B and C executing either: that is its *constitution*, not V itself. What remains is a structure with three levels, none reducible to the others:




Defining a capability is not exercising it. The journal of this section holds that separation as data: one definition, two exercises, and replay reconstructing both from the one. The deduction establishes what holds of this arrangement, measured — not what must hold of arrangements generally.

2.4 The capability is above its modalities

One narrowing has to be refused before it settles in, because the measurement above invites it: **the assembled verb is not defined by the journal**. What §2.2 and §2.3 exhibit is the *journalized command* modality of its exercises — the one where performances change state and persist. It is the modality where the definition/exercise separation can be shown as data, which is why it was measured first. It is not the definition of the term.





A verb exercised only as a query — operations of several repertoires read, a value computed that none of them holds, and the result pronounced — is, on the definition this section arrives at, as much an assembled verb as §2.2's. Its constitution may be *richer*: it can relate results no repertoire can relate. What differs is the modality of its exercises: a query executes under a read lock and does not persist to the journal (`ActorHandler.cs:2954-2969`), so the reuse measurement of §2.3 cannot be run on them. That bounds the *evidence*, not the *category*.

And the modality is not read off the operation — nor fixed on the verb: it is **attributed by the assembler, per performance**, and the record follows the attribution rather than the operation's effect. Measured (Appendix A: `WhoDecidesWhatCountsAsHistoryLab.cs`), with everything but the attribution held fixed — one body, one verb under §2.3's seam, three performances: exercised as a query it leaves the journal head unmoved; performed as a command the act is journaled and a fresh subject replays it from the record alone; queried again after the command, the head stays. Had modality been the verb's, the second and third performances could not have differed. The domain computes the same value each time and never changes.

journal head after setup	2	
one body exercised as a QUERY	2	unchanged
the same body performed as a COMMAND	3	advanced
the same body queried again	3	unchanged
after replaying the journal from scratch	3	

Nothing in the operation distinguishes the two runs. What distinguishes them is a decision that belongs to the party constituting the verb: **whether this performance counts as the subject's history**. Disclosing a value and looking at a value can be the same computation and different acts — a subject for whom the disclosure itself must stand as fact constitutes it as a command, and the effect on the domain is identical: none. Two subjects assembled over the same repertoire may therefore have different histories without the repertoire changing.

Practice already knows that a disclosure can be a fact. Access audit logging records reads precisely because the reading itself must stand — the health and payment-card regimes require it of systems handling protected data (HIPAA's audit controls, 45 C.F.R. § 164.312(b); PCI DSS Requirement 10) — so *jour-*

naling a read is not the observation here, and the section would be conceding its oldest neighbour by implying otherwise. Two separations hold, and both are checkable. **Where the decision lives:** an audit regime fixes the recording decision once, for a category of operations — all access to the protected data is logged — and implements it in the domain or its infrastructure: an interceptor, a middleware, a logging call inside the method. Here the attribution is per performance, taken at the plane of assembly, and the repertoire is untouched — the same operation, under the same verb, enters one act as a query and the next as a command, which no category-wide rule expresses. **Where the fact lands:** an audit log is a second record, kept beside the state the system runs on and consulted apart from it. Here the journaled disclosure enters the subject's one history — the record the subject is reconstructed from — and replays with it, as the measurement above shows. The kinship strengthens the observation rather than threatening it: what a regime must bolt on beside the domain, category-wide, the plane attributes verb by verb, editing no one.

What command-query separation (Meyer 1997) separates is therefore real, but it lives one level up from where it is usually placed: it is a property of the subject's acts, attributed at assembly per performance, not an intrinsic property of the domain's operations. An operation that changes nothing can constitute a command, because what the attribution governs is historicity rather than effect.

This is a geometry of authority the series has met before, twice. The output finding was that the producer of a value is not the authority on its destination; here, the author of an operation is not thereby the authority on the historicity of the acts that will incorporate it. And the movement of the audit-logging contrast — a property fixed once, system-wide, becoming a decision attributed at a named plane — is the movement the series made for eventual consistency, where a system-wide contract became a per-verb list the assembler wrote by hand; here the grain is finer still, one attribution per exercise. Consistency there, historicity here: the same register, a decision habitually fused into the system, found separately attributable — there verb by verb, here act by act. And where persistence lives inside a domain, its author necessarily takes both decisions at once — what the domain means, and what survives as history. Neither decision is wrong there. What the measurement shows is only that they are separable, and where the second one can live.

So the definition, stated at the level it actually lives:

An assembled verb is a capability of a subject, constituted from operations available to assembly. Command and query are modalities the assembler attributes to its exercises. History is what follows from that attribution — which of the subject's performances count as its history — not what makes the verb a verb.

Four things, then, that are not synonyms:

domain	what can be computed
assembly	what the subject can do

performance	what the subject does
history	which of its performances count as its history

If the term had been reduced to the journaled command, *assembled act* would have been the honest name. That the query enters naturally is what confirms *verb* was the right abstraction: a repertoire holds what a subject can do, and what a subject can do includes asking and saying, not only changing.

And the four lines are not left as a table. Appendix C states them as four projections of one small operational object — a core calculus in which assembly, exercise, modality and replay are rules — and shows by witness pairs that none of the projections determines the others. The table is the claim; the calculus is where the claim is a property.

2.5 The objection this does not yet answer

A reader should still be unconvinced, on one specific ground: `CompletePurchase()` also exists between two executions — in the program text. Call the method twice and nothing is restated either. Why is `action 1` not simply a method that happens to live in a log?

That objection deserves an instrument rather than a paragraph, and it gets one.

3. The Criterion

Call the composed work a verb performs *the whole*, for this section's purposes. Two other arrangements produce a record that also refers to the whole, and one of them gives it a better name than this one has:

COORDINATION	in two measured shapes. An event-sourced process manager, with its own identity, state and stream — it may emit <code>PurchaseCompleted</code> , and its name, lifecycle and durability are conceded in full. And a durable-execution orchestration, run on a third-party engine (§7.2), whose record is whatever that engine writes.
DESCRIPTION	a trace, in the distributed-tracing sense (Sigelman et al. 2010). Its root span is called <code>CompletePurchase</code> and every statement appears beneath it. It records the grouping faithfully.
CONSTITUTION	the entry of §2.2. It is called <code>`action 1`</code> .

The criterion owed here must be askable of any system — including one that has never heard of this substrate. A criterion of the form *is it a journal entry of the kind §2 describes* would answer nothing.

3.1 The question

Take the record that claims to be the whole. Ask two things of it, and nothing else:

- (a) **Content.** How many of the statements that constitute the whole does it contain?
- (b) **Replay.** Hold the arrangement's engine and the repertoires fixed, and supply nothing else: does **the record alone** determine a replay after which the whole's outcome exists again?

Both are answerable of any record by inspection and by running it, and both are **total**. (b) quantifies over the arrangement's own replay semantics, so an arrangement that supplies no replay operation answers *no* — vacuously, because nothing exists that could bring the outcome about — rather than falling outside the question. And *alone* is load-bearing: it constrains the replay's **input**, not its mechanics. A replay that reads anything beyond the record itself — beyond the engine that runs it and the repertoires it invokes — is not a replay of the record alone; in particular, a replay that reads a definition kept outside the record, as compiled code or an ambient program, leaves the outcome determined by the pair rather than by the record, and the same record means something else when the program changes (Appendix C states the distinction as a theorem: Proposition 5). The fixing of engine and repertoires is itself a relativization, carried as a limit in §8: what the record closes over is the composition, not the operations. Without the clause, (b) would credit to the record what an arrangement's own recovery guarantees — durable execution does bring its orchestration's result back, in its own operation — and the question would stop separating anything; the clause is what makes (b) a question about the *record* rather than about the arrangement's uptime. Neither question asks what the record is called, which tool produced it, or what its author intended.

And what the two questions classify is a *record*. The criterion is therefore a criterion for **recorded constitution** — not for the assembled verb as such — and the narrowing is load-bearing rather than a concession. §2.4 already met its reason: a verb whose exercises are queries deposits nothing, so no record exists for this section to ask anything of, and the verb is no less a verb. Three levels nest, each strictly wider than the next:

assembled verb	§2.4's definition –
queries included	
⊃ journaled assembled verb	the exercises that persist
⊃ record satisfying this criterion	recorded constitution –
what §3 decides	

§3 decides the innermost level, and only there is it decisive. Its verdicts are *evidence* for the outer levels — §2.3's reuse and §2.4's attribution read the same record — but the category is defined in §2.4 and owned in §4, and neither burden is carried here. A criterion that tried to decide verbhood directly would have

to see performances that never enter a record, and would stop being askable of arbitrary records — which is the property that makes it an instrument.

What counts as the outcome existing again. Outcome identity is taken *up to fresh values*: two outcomes are the same when they differ at most by a consistent renaming of identifiers minted during the run rather than supplied to it — the standard move on fresh names. The tolerance is principled, not convenient — and Appendix C proves it *necessary*: under the host axiom, minting is fresh by construction, so exact identity is not available and a bijection on minted names is (Proposition 4). Any criterion that rewards *re-performing* must state one, because re-performance re-mints whatever was minted; a criterion demanding bit-identity would hand the victory to arrangements that never re-perform and declare every re-performing system a failure. The tolerance also rewards nothing that does not re-perform — it only relates outcomes that exist again. And it is not free for this paper: §4.3 measures the value that forces it, in this paper’s own flagship entry, and §8 carries it as a limit. Without this clause the instrument would fail the paper’s own constitution, which is why it is stated here, before the criterion is used.

3.2 The measurement

The same work — six statements — performed under all arrangements, over the same two prebuilt domain assemblies (Appendix A: `TheCriterionForConstitutionLab.cs` with `ProcessManagerHarness.cs`; `TheEventSourcedCoordinatorOnAThirdPartyEngineLab.cs`; `TheCoordinationArmOnAThirdPartyEngineLab.cs`; `TheConstitutionPrototypedOnAThirdPartyEngineLab.cs`).

The provenance of each coordination row is declared, because a baseline the author configures is not a class. Both coordination rows are measured on third-party engines. The event-sourced shape runs as an Orleans journaled grain — the grain journals its own transitions, reconstructs from that journal, and the operations it coordinates are other grains’ acts — with the record read back through the framework’s public surface; the author-built harness is kept beside it as the corpus-mirror it always was, since the corpus’s own sagas have exactly this shape (§6). The durable-execution shape — the class at its strongest — is measured on the Durable Task Framework, the engine underneath Azure Durable Functions, with the history read through that framework’s public dispatch middleware. In both, the record counted is the engine’s own.

And the constitution row is not left as the only row realized solely on this paper’s engine, because a single implementation cannot separate two readings — *the assembly plane is a realizable arrangement* and *one framework has this feature*. The arrangement is therefore prototyped, in a few hundred lines, on the coordination row’s own engine: a journaled grain whose log holds `def` and `act` entries and whose state transition re-executes the definition’s statements against the same repertoires. The substrate is held constant and only the arrangement

varies, and one third-party engine realizes both cells — what decides the cell is what the journal holds and what the fold does with it. The prototype reproduces the row’s every property on a three-statement verb over the same repertoires: one definition under two exercises, the second stating nothing of the constitution; a query performance that executes and leaves the journal unmoved; reconstruction that re-executes every journaled statement from the record alone; and the domain-minted identity differing per replay — §4.3’s boundary, met unchanged on the second engine.

arrangement	names the whole	(a) statements in the record	(b) the record alone replays into the whole
coordination — event-sourced process manager (third-party engine; corpus-mirror harness agrees)	yes	0 — its own journal holds 6 transitions, none an act’s statement	no
coordination — durable execution (third-party engine)	yes	6	no
description	yes	6	no — nothing to replay
constitution	no	6	yes — up to fresh values (§3.1)

Naming decides nothing, and it runs the wrong way. The arrangements the last row must be distinguished from are the ones with the business names; the constituted verb is an integer. If naming carried the property, the other rows would have it.

(a) separates only the first shape. The event-sourced coordinator’s own journal holds its transitions and no operation: which step it reached, what it awaits — measured on Orleans, six progress marks and zero operations, the operations living in other subjects’ acts. The durable-execution record, by contrast, holds all six invocations — names, inputs, results. Content cannot tell it from the description row or from constitution.

(b) separates the rest, and it separates them differently. The trace answers *no* vacuously: it supplies no replay operation, so no replay of it exists after which the outcome could obtain. The durable-execution record has a replay, and in the model’s own operation the orchestration recovers and completes

— the row’s *no* does not deny what the model guarantees. What the row scores is (b)’s input clause: that replay re-executes **the orchestrator’s code** against the recorded history — a definition living outside the record — so its outcome is determined by the pair, not by the record alone. Measured: the orchestrator body ran in seven episodes, six of them replays, while every activity executed exactly once; the record alone re-performed nothing, and that is the model’s own determinism contract. The constitution row’s entry is read back, and reading it back performs the six statements again — nothing outside the record but the engine and the repertoires.

The input clause is not a footnote the strongest row forced; it is the half of the criterion that decides that row: **where the definition lives**. In durable execution the composition is compiled code, and the history holds exercises and outcomes that only that code can re-walk (§7.2). In the constitution row the definition is itself an entry in the record, which is why — with the repertoires held fixed, as (b) fixes them — the record alone suffices to re-perform. The difference from the strongest coordination is in what stays ambient: there, the composition and the operations; here, the operations only.

Only the conjunction classifies:

- (a) alone description, constitution and durable coordination are indistinguishable
- (b) alone the other three fail it, each for a different reason
- (a) and (b) exactly one arrangement contains the statements and re-performs them

Coordination preserves relations among acts — or, in its durable form, the invocations and their outcomes. Description preserves an account of a grouping. Constitution preserves the grouping as something that can be performed again.

3.3 The method objection, answered

Now §2.5’s objection can be put to the instrument. A method with exactly the same body — even the same six statements — leaves what behind after it runs?

method extraction	source: CompletePurchase = [a, b, c]
	history: a, b, c

the arrangement here	source: action 1 = [a, b, c]
	history: action 1 := [a, b, c]

Execution of a method does not by itself preserve its lexical grouping; unless something separately records that grouping, the history contains whatever its constituent operations leave behind. Something separate *can* be added — that is the description row, and it was measured rather than dismissed. What the third row has is different in kind:

**Lexical composition disappears when the program runs.
Constitutive composition survives into the record.**

And note the control hiding in the identifiers: the method is the one with the business name, and it is the arrangement that loses the composition. What differs is not what the composition is called but whether it is still there after the run — still there in the sense the criterion makes precise: containing its statements, and re-performing them when replayed.

3.4 What the criterion does not decide

It does not say a description could not be made replayable. It could — a trace enriched until its record suffices to re-invoke the operations would be a journal, and the criterion would classify it as constitution, which is the right answer. The criterion is indifferent to tools and names; it asks what replaying a record does.

It does not rank the three arrangements. They answer different questions: what follows this act, what happened during it, what constitutes it.

Nor are its two questions without ancestry of their own. The provenance literature has long distinguished prospective from retrospective provenance and asked when a captured record is re-executable (Freire et al. 2008; Moreau and Missier 2013). The kinship is acknowledged rather than annexed: what §3 adds is not either question but their use as a *joint* criterion over one record, whose conjunction — and only the conjunction — separates the three arrangements above.

And it classifies *records*, so its reach is the modalities that leave one. A query exercise of an assembled verb may leave no record at all, and the criterion is silent about it — which bounds the criterion, not the verb (§2.4).

3.5 What kind of instrument this is

One question has to be answered explicitly: is the classification empirical at all? The concession of §3.4 — a description enriched until re-invocable would classify as constitution — shows that the classification is **by construction**. *Constitution*, in this paper, names the conjunction; no observation falsifies a definition, and the paper says so plainly rather than dressing a definition as a discovery. This is the posture §1.2 declared, applied where it bites: the criterion is a Gregor Type I construct, a defined, observable property by which arrangements are described and compared.

What is empirical — and falsifiable — is everything the definition is then put to:

- **each cell of the table.** Every row re-runs from Appendix A, both third-party rows run on their engines' own semantics, and the constitution arrangement additionally runs as a prototype on one of those engines

(§3.2), so no cell rests solely on this paper’s engine. A Durable Task orchestration whose record, with the orchestrator’s code changed, still replayed to the same outcome would overturn its row — the record alone would be deciding; an Orleans coordinator whose journal carried the operations would overturn its.

- **the separations.** That the shipped arrangements occupy distinct cells was not derivable from the definition. The durable row’s 6 could not have been read off the definition; it could have come out otherwise, and for one arrangement family it did.
- **the constitution row itself, which is not satisfied by fiat.** Strictly — without §3.1’s fresh-value clause — the flagship entry of §2.2 fails (b): a domain-minted identity does not come about again on replay (§4.3). The equivalence is therefore not a convenience of this paper’s arm; it is *forced* by a measurement against that arm, declared before the criterion is used, and carried as a limit in §8. And the row stays falsifiable under it: a replay that missed the outcome even up to fresh renaming — a lost entry, a wrong count — would overturn it.

So §3 defines a property and measures where arrangements fall under it. The definition is the instrument; the measurements, the separations, and the one partial failure are the findings. And the two halves are kept formally apart: Appendix C restates the definitional half as propositions over a core calculus — so what is by-construction is exactly the part that admits proof — while the labs of Appendix A carry the empirical half, one conformance check per proposition.

And it does not yet say whose verb V is. That is the next section, and it is a separate question: a record can hold a composition as one replayable unit without anyone having asked to whom the unit belongs.

4. The Subject

4.1 To whom the verb belongs

The entry of §2.2 is not floating in a store. It is a record **of one actor’s history**: the execution is attributed to the actor that performed it, replay reconstructs that actor from its own record, and the reuse of §2.3 is that same actor drawing on its own verb a second time.

So the capability produced by assembly has an owner, and the owner is none of the ingredients. V does not belong to the repertoire that declares A, nor to B’s, nor to C’s — those repertoires know nothing of V, and §5 measures just how little. V belongs to the repertoire of the subject that performed the assembly.

Which gives the term its two faces, and both are already in the data of §2:

repertoire	history
-----	-----
V	V(p1)

V(p2)

As repertoire, V is a **capability** — what the subject *can do*. As performance, V(p) is an **exercise** of it — what the subject *did*. The two columns answer different questions, and for the journaled modality the record preserves precisely the relation between them: each entry on the right names the definition on the left. For other modalities the right-hand column may be empty and the left-hand one no less real — a query verb is exercised without depositing anything (§2.4). That is what the paper’s title means:

An assembled verb is a capability of a subject, constituted from operations available to assembly (§2.4) — a whole that is the verb of that subject, rather than a coordination among the acts of its parts. The term claims no mechanism. §2 showed one realization; §3 gave the external test; this section supplied the owner — and with the owner, §1.1’s argument gains its location: a place other than the pieces exists. §5 cashes what the place buys (the pieces need not anticipate), and §6 measures what the corpus pays where no such place is.

A sentence held back until here, because before §2 and §3 it would have been a metaphor and would have smuggled in subject, voice and attribution before any of them were shown:

There was no speaker among the pieces for whom the sentence was first-person speech.

The pieces were never missing speech of their own — each names every transition it is entitled to name. What was missing was a speaker for what spanned them, and the actor of §2 is that speaker in the only sense this paper uses: the composition is in *its* record, attributed to *it*, replayable as *its* act.

And “the only sense this paper uses” is a boundary, carried as a limit in §8. What the record establishes is **ownership**: the definition and its exercises are in one actor’s record, attributed to it, and replay reconstructs that actor from that record. **Subjecthood is not derived from constitution.** The same data admits a flatter description — *an actor instance owning a dynamically extended executable vocabulary* — and every measurement in this paper survives it unchanged; what the richer vocabulary adds is names for the relations measured — owner, exercise, history — not further claims. That the owner is a *subject* is the actor model’s frame and the series’ prior finding that identity precedes staging: assumed here, and argued elsewhere.

4.2 What the subject’s position is, and is not

Two misreadings have to be refused.

The subject does not know more about any repertoire than the repertoire does. Its additional knowledge is about the relation among the pieces,

not about the interior of any piece. It can know that one value participates in two operations, because it occupies the only position where both are present; it cannot decide what that value means inside either. (Both halves are exhibited concretely in §5: the co-binding scope of §5.1, and §5.2’s unstatable proposition.) Its reach has **width, and not depth** — reach rather than authority, because a runtime capability and a permission are different things and must not be fused:

The assembler can say more than either repertoire, because it sees both. It cannot say more about either repertoire than that repertoire permits.

Composition gives it no operation a repertoire did not already provide: no added method, no relaxed invariant, no authorship over decisions made inside an invoked operation. **It can say a sentence no part can say, and cannot correct a single word of the sentences the parts can say.**

And the subject’s position was not granted by the pieces’ authors. The repertoires here are consumed as compiled assemblies whose types the host language cannot even name — the compiler refuses the program that tries, and the refusal was obtained by invoking the compiler on one written to fail (Appendix A: `RepertoireBoundaryLab.cs`). The registration gate admits a repertoire’s internal types (`DomainLibraries.cs:151`) and the actor reaches their members by reflection (`ActorHandler.cs:4457-4516`). The composition happens anyway, which separates two properties usually fused: *composition does not require host-language accessibility*. But reaching a type the host cannot name is a capability of the runtime, not a permission granted by the type’s author, and whether standing on the far side of that boundary amounts to an entitlement is a normative reading the compiler is silent on. This paper declares it as a reading. A reader may decline it and keep every measurement.

4.3 What the arrangement cannot supply

The honest boundary of the subject’s position, measured rather than conceded in the abstract.

Values that enter an act from outside the computation cross a **parameter plane**: captured once at the moment of the act, journaled with it, re-injected on replay rather than taken again. A clock reading is handled that way — on replay the recorded occurrence time is injected back as the act’s **Now** rather than the clock being read again (`ActorHandler.cs:1456`) — and so is an expression parameter, resolved once and frozen into the entry’s arguments (`Parameter.cs:439-455`).

The boundary itself is not this paper’s discovery, and saying so costs nothing. Capturing nondeterminism at a recording frontier is the classic problem of record-and-replay (LeBlanc and Mellor-Crummey 1987; O’Callahan et al. 2017), and durable-execution runtimes enforce exactly this discipline for workflows: an

orchestration may not read the clock or mint a fresh identifier, replay-stable substitutes are supplied in their place, and the equivalence of the programming model with its record-replay implementation has been proven (Burckhardt et al. 2021). And beneath both, the discipline is the actor premise on the value axis: an actor’s next state is a function of the message it received and the state it held, so a value drawn from the machine behind the plane is an input with no sender — which is why Appendix C states the boundary as a condition on the host interpretation rather than as a rule of authorship. What this paper adds at the plane is only the measured failing case at its edge, met while composing repertoires the subject does not own — a value minted inside a foreign factory, which no substitute could have intercepted because the assembler is never asked for it. The plane has an edge, and one value in the composed verb falls behind it: one repertoire mints its own identity *inside* a factory, so the value never crosses the plane, because the assembler is never asked for it (the call site is marked, by number rather than by intent, in `SubscriptionPaymentBridge.cs:38-72`; the measurement is `RetrievalIsALookupLab.cs`, which measures this and only this: the supplied key stable, the minted identity fresh, and retrieval a lookup in the subject’s own state; it says nothing about capture). The repair the plane itself would offer is out of reach by the grammar of the case, and saying so exactly prevents a misreading: an expression parameter can draw a value once at the plane and freeze it into the entry, but only the repertoire could accept it, and this factory’s signature asks for none — the identity is born inside the operation, as part of its own domain event. Making the factory ask would repair replay by editing the repertoire, and an edited repertoire exits the paper’s category: the composition under measurement is of repertoires the subject does not own and did not touch (§5.2 counts the edits: zero), and the corpus entered under transformation rules fixed before measurement (Appendix A) — so the value was left as the repertoire produces it, and reported. Measured, live and replayed:

the assembler's key	stable across replay
the domain's identity	a different value on every replay

Replay re-runs the act, and the minted value returns fresh — an outcome the criterion’s equivalence absorbs (§3.1), and which Appendix C proves is the identity replay can promise, exact equality being unavailable under its host axiom under its host axiom. It is the measured value that forces §3.1’s clause to be declared before the criterion is used: without it, §3 would fail this paper’s own entry. In the calculus of Appendix C the same discipline is the axiom on the host interpretation, and this value is its licensed minting measured live: fresh names differ per replay by construction, which is the bijection Proposition 4 proves is the strongest identity replay can achieve, and the clause §3.1 states. Two things would take down replay equivalence and nothing above it — a true violation of that axiom (uncaptured I/O, an escaped clock, foreign mutable state), and the *capture* of a minted identity into a later act’s recorded arguments — and the definition and its exercises stand in the record under either: the verb is no less defined, exercised or owned. Nothing here is a defect of either party — the repertoire was written to be called once, by a caller that would store what

came back. And what the measurement locates is not a replay failure, since the equivalence holds; it is an **authority boundary**. Durable-execution runtimes stabilize such values by supplying replay-stable substitutes from the runtime's own facilities; the domain's author could stabilize this one by asking for it. The assembler alone can do neither, and the cost lands on the subject: **it inherits the decision without acquiring any authority over it**. One line separates properties usually spoken of together:

A repertoire can be complete, executable and testable — and still
mint identities that no assembler has the authority to stabilize.

5. The Case That Makes the Plane Visible

5.1 The constitution may draw from anywhere the subject reaches

Nothing in §2–§4 required two domains. An assembled verb may be constituted from one operation, from several operations of one repertoire, from operations of several repertoires, or from a mix that includes statements only the subject's own plane can contribute — its computation over what the repertoires said, and its speech:

Figure 1 — the forms a constitution may take.

<code>V := A</code>	
<code>V := A; B; C</code>	one repertoire
<code>V := A; X; Q</code>	several repertoires
<code>V := A; print x</code>	the subject's speech, outward
<code>V := A; tell t x</code>	the subject's speech, to another subject

What makes V assembled is not the plurality of its operations, nor of the repertoires they come from. It is that the resulting capability is authored at the assembly plane.

The constituents determine what is *available* to assembly. They do not determine the verbs assembly may author from it. And the caption says *authored* rather than *selected* because of the last two forms: assembly does not only compose repertoires — **the plane has vocabulary of its own**. *Print* and *tell* are the subject's speech, outward and to another subject, and neither is writable inside a domain: an operation has no voice, no say over destinations, and no acquaintance with the subjects that future verbs will address. Both are earlier papers' subjects and are not re-argued here. Read downward, the figure is the paper's ladder in miniature:

domain operation	<code>A()</code>	what a repertoire can compute
assembled computation	<code>A(); B()</code>	what this subject can do
assembled speech	<code>print x</code>	what this subject can say
assembled address	<code>tell t x</code>	whom this subject can act toward

And the deduction must be kept honest here, because a repertoire could not write the *tell* even if its author wanted to: **tell** is a statement of the subject's plane, not of the host language a domain is written in. What a repertoire's author *can* write is the equivalent the corpus actually contains — publish an event, call a bus — and that equivalent has a measured shape (§6): it authors a transport, a contract and an implied destination into the piece, and it fixes in advance one of the future acts the piece may participate in, because its author must anticipate the subject it will someday serve:

inside a repertoire	at the assembly plane
A() { ...; bus.Publish(ADone(x)); }	A remains A
a transport, a contract, a destination	V1 := A; print x
authored into the piece	V2 := A; tell t x

A did not change; what changed is who is using A to say what. This is the separation of authorities the series keeps meeting, on a third face now: producing a value grants no authority over its destination; providing an operation grants no authorship over the verbs that will incorporate it — and imposes no duty to know whom those verbs will someday address. A repertoire need not know the sentence in which one of its operations will someday become a word.

One more form deserves its own exhibit, because it may be the purest of all:

W := balance = a.Balance();	a query verb: reads across repertoires,
points = x.Points();	relates what no one of them holds,
print balance + points	and says it

A query verb such as *the customer's position* — a balance from one repertoire, points from another, pending orders from a third, related and pronounced — mutates nothing, may persist nothing, and is still undeniably a capability of the subject: it says something no piece can say, because it occupies the position that sees them jointly. The width of §4.2 is demonstrated there in the computation itself, with no boundary machinery in sight.

One discipline keeps that exhibit from licensing what the paper has already ruled out, and the sketch is written to obey it. The computation admitted at the assembly plane is **the relation itself**: arithmetic over what the repertoires said, of the kind no repertoire could host — a formula whose subject matter spans two repertoires cannot live inside one without lodging exactly the cross-repertoire knowledge §5.2 measures as absent. A formula whose subject matter lies **within** one repertoire is that repertoire's operation, and assembly *invokes* it rather than authoring it. This is not a style rule laid on top of §4.2; it is §4.2 applied to expressions: authoring single-repertoire logic at the assembly plane would be depth, and the assembler's reach has none.

The same honesty is owed to the sketch's humblest statements. Assignment is vocabulary the plane shares with every language, and the paper claims none of it. What no repertoire could write is the **scope** those assignments fill: **balance** and **points** bound side by side in one place, a pair neither piece could hold —

not because binding is special, but because those names cannot co-occur inside either (§5.2). The same scope holds values whose types the host language cannot even declare a variable of (§4.2). The subject’s typed scope is the position that sees the repertoires jointly, made operational; assignment is merely how a constitution moves values among its operations — the third of §2.2’s constituents of unity, at work.

The plane — verbs of the subject, above operations of the repertoires — is the same in all of them. But one form makes the plane *visible* in a way the others cannot, because it removes the reading that V is shorthand for something one repertoire could have said: when the operations come from repertoires that share nothing, no piece’s vocabulary can even in principle be the verb’s home.

5.2 One story over two repertoires that share no type

An author-constructed scenario — written for this study, and evidence of realizability only (Appendix A: `ALongTrajectoryOverTwoUntouchedRepertoiresLab.cs`). A customer buys a subscription, which is paid and granted; then places an order, adds items, has a payment method verified; the order is paid and shipped.

domain operations invoked	17
aggregates taking part	4
shipped assemblies, different authors, unmodified	2
types the two assemblies share	0
test doubles	0

It runs, end to end, with no handler, no bus, no repository, no mapper — the claim being narrower than their dispensability to the systems that use them: none is required to express *this trajectory* in the repertoires’ own operations.

And the run settles §1.1’s corollary by construction rather than by argument. The two assemblies were authored years apart, by parties who never met, and nothing in either was written toward this verb — yet seventeen of their operations constitute it. **A repertoire need not anticipate the verbs in which its operations will later participate** is not a hope here; it is the condition the scenario ran under — §1.1’s deduction, run as an experiment.

And there are things true of the trajectory that no participant in it can state. The customer appears as a Guid-bearing identity in one repertoire and a string-bearing one in the other; neither assembly names the other; so neither repertoire is in a position to state a proposition whose subject matter spans both — *this customer bought a subscription and later bought these items* has no possible position inside the run from which it can be said. For either repertoire to assert

it, the relation between the two representations would have to become knowledge held inside one of them, and the measured zeros show neither published repertoire occupies that position.

The whole is knowable, but not by any of its parts.

The run raises nine domain events — the corpus domains’ own, emitted by their aggregates, in a scenario this paper composed, so the count carries the constructed kind of §1.2’s table and no weight about either corpus system’s design — every one naming a transition of a single aggregate, none naming the trajectory:

The repertoires name every transition they are entitled to name.
What they do not name is the trajectory that exists only in their composition.

Under the arrangement of §2 that same trajectory is one entry — the entry printed there — and the subject whose verb it is can be asked for it again.

6. Returning to the Corpus: Where Is V, and Where Did the Continuity Go?

The order of the argument matters here, and it is the reverse of the archaeological one. This paper did not infer a missing category from the corpus and then build something to fill it. The category was demonstrated positively first — §2 constituted V, exhibited its one definition, exercised it twice, and replayed both exercises — and the corpus now serves as a **control**: does it contain a capability at V’s level?

That inversion changes what the corpus evidence has to carry. It no longer has to prove an absence by exhaustion; it has to answer a yes-or-no question about a category whose existence is already established. The question is one word: **where is V?** And behind it, §1.1’s question in the past tense: **where did the continuity go?** The trajectory runs — the corpus works — so if no unit holds its continuity, the continuity went somewhere, and this section finds it.

What is found is A, B and C — and their events, their handlers, their correlations, their coordinators’ state. What is not found, anywhere, is V as a capability of a subject.

The census. Every named artifact of both subsystems’ vocabularies was enumerated and coded, and both steps run from published rules rather than from a reading (Appendix A: `EveryNameStaysWithinOneSubjectsVocabularyLab.cs`, whose header is the codebook and whose body is a coder with no judgment to exercise). The population rule reproduces the denominator mechanically from the pinned corpus: 27 domain events, 16 integration events — the subsystem’s declared events plus the foreign events it handles — and 29 commands, with

handlers, generic infrastructure wrappers and queries out of population: **72 artifacts** across three layers and two grammatical moods. The construct — *names a trajectory spanning subjects* — is coded by two rules, because one alone is blind in exactly the form that matters. Rule one is mechanical, and generous within its form: an artifact is coded as naming a trajectory if its identifier references two or more distinct aggregates of its own system — the lexicon derived mechanically from the corpus — with no verb, no ordering, and no judgment required. But a trajectory can be named without naming any constituent: *CompletePurchase*, this paper’s own opening example, carries no aggregate noun, and a rule that counts aggregate references cannot see it. So rule two hands to **manual coding**, row by row and each with its reason published, exactly the identifiers rule one cannot see: every identifier referencing no aggregate at all, and every identifier bearing a word from a published process lexicon (*checkout*, *onboarding*, *fulfillment*, *purchase* and their kin) — and a residual identifier without a published coding fails the run, since silence is not a coding. Coded under both rules the count is **zero across 72**: mechanically zero under rule one, and a three-artifact residual under rule two, each naming a single subject’s transition, each reported with its reason. What this is not has to be said as plainly as what it is: **no independent coder was employed**. The two codings that agree — the author’s reading, and rules the author wrote after that reading — are two instruments of one authorship, so their 72/72 is internal consistency made auditable, not inter-coder reliability; and with every artifact in one category there would be no variance for a chance-corrected coefficient to correct anyway (Cohen 1960). Re-running the lab verifies that the published rules produce these codings from the pinned corpus — verification, not a further coding, since a deterministic rule is one coder however many times it runs. A further coder is a reader who codes the 72 against the construct and compares, which the codebook and the coded table exist to make cheap. Every artifact stays within the vocabulary of a single subject. The absence cannot be explained by a naming convention: it persists in the two layers where naming was unconstrained. The clearest single case is a pair of commands, one asking for a subscription’s payment, the other for the subscription that payment entitles — two imperatives for the two halves of one act, and no imperative for the pair. And after §2.2 the census means something sharper than it could have meant as an opening exhibit: **the missing name was evidence, not the thing missing**. The constituted verb of §2 has no name either. What the corpus lacks at the level of the trajectory is not a word — it is the unit the word would have named.

Defeat condition, cheap to run: exhibit one artifact in either system — event, command, type or method — whose name refers to a trajectory spanning operations of more than one subject. The count is zero across 72; one takes it to one, and either form is accepted: two subjects’ nouns in one name, which rule one takes mechanically, or a single word naming the whole — a CompletePurchase, a Checkout — which rule two’s manual lane exists to see, and whose lexicon any reader may extend.

The dissection. Six rules, pre-registered and committed before any removal,

were applied to both systems’ write surfaces in audited two-commit steps, with the hole count read from compiler diagnostics; the full procedure, its amendments, its stopping rule and its unequal-objectivity declaration are in the artifact. Two results matter here. The authors’ own domain suites — 49 tests — pass unedited against the dissected assemblies, so **the operations already computed on their own; what the dissection had to free was the build.** And what remains of a write step after persistence, ambient context, transport, wiring and observability are removed is: invocations of domain operations, plus a residue that articulates them. The pieces of V are everywhere. V is nowhere.

One reading of the persistence rule follows from §2.4 and is offered as a reading: removing persistence from a domain does not merely make it portable — it returns a decision to a plane that has context for it. A domain with persistence inside it has one author deciding both what the domain means and what survives as history; a domain without it leaves the second decision free to be taken per performance, by the party constituting the verb. Nothing in that is a judgement of either corpus system, whose authors took both decisions deliberately and well.

The compensations — where the continuity went. Where the plane has no author, the software compensates, and the compensations are measurable. An arrow between two halves of one act is carried by a durable enqueue into the module’s own schema, read by a poller in the same process — transport machinery supplying only *and then*, since the same act composed directly, through an in-memory queue, or through a keyed register delivers identical values, differing only in the identity a repertoire mints for itself (Appendix A: `WhatTheCarrierCarriesLab.cs`). Events that carry a whole aggregate are consumed by handlers that read an identity off them and fetch the aggregate again. Each mechanism buys real guarantees its system needs, and they are what movement looks like when the composition has nowhere to live as a unit — **continuity authored into the pieces**, because no other place could hold it. What §2 changes is not that continuity exists but where it can live: once the whole is a verb of a subject, the pieces no longer have to carry it.

7. The Neighbours

A claim this visible owes an answer to an uncomfortable question: if assembly producing a capability were real, how could five decades of languages, patterns, planners and runtimes not have produced it already? The honest answer is that most of V’s ingredients were produced, some long ago, and this section says so before saying what is left. The search behind it was run to find the paper that would make this one unnecessary, and the neighbours below are what survived it. The search is itself reported as method in Appendix B — its query strings, its sources, and the adversarial-review channel whose seven contributions measure what the strings missed; the claim it supports is scoped there: absence over the examined set, with §3’s questions as a finite test for any candidate beyond it.

One objection organises them all:

“You have merely made composition first-class.”

Every neighbour is put the same single question — **what exactly is V in this model?** — unfolded into the battery the paper has already answered for itself:

does V exist separately from A, B and C?

can V be reused, with parameters?

are V and V(p) distinct?

whose capability is V?

who assigns each exercise's command/query modality?

who decides whether V(p) belongs to the subject's history?

The first three are old, and the neighbours below hold them in various combinations. The last two are where this paper's difference proves hardest to reduce, and no examined neighbour holds either (Appendix B).

7.1 The composition ancestors — conceded first

Procedural abstraction is the oldest and deepest. A function is a definition exercised many times; that structure is the founding gift of every programming language, and §2.3's three levels reproduce it deliberately. What a function definition lacks is not reuse — it is everything below the third question: it lives in program text, belongs to a namespace rather than a subject, and §3.3 measured what remains of it after a run.

The Command pattern and its composites (Gamma et al. 1994) make the *request* a first-class object; a MacroCommand executes a sequence of Commands, and macros are built from existing commands. The pattern even keeps a history — the invoker's undo stack is a record of invocations. Each piece of the battery's first half is here. What is not: the macro is program structure, not an entry in anyone's history; the undo stack is the invoker's ephemeral memory, not the subject's record; and the modality is intrinsic — a Command *is* an execute, not a verb whose historicity the assembler attributes per constitution.

Transaction Script and the application service (Fowler 2002) organise domain operations to carry out a use case, and the corpus of §6 uses them well. The body can be identical to V's, which is why §2.5 refuses to answer this neighbour from the body. §3.3 answers it from what survives execution.

Workflow systems (van der Aalst and van Hee 2002) already hold the definition/instance separation at the process level: a workflow definition is instantiated as cases, and subworkflows make pieces of a definition reusable across it. So *defined once, exercised many times* cannot be claimed even at the granularity of business processes. What a workflow case's audit trail is, though, is the engine's record of an enactment — everything is logged because the engine logs, not because a party attributed historicity to this verb and not that one; and

the definition belongs to the process model, not to a subject with a history of its own.

Algebraic effects and handlers (Bauer and Pretnar 2015; their language Eff) make computational effects first-class, definable and composable — one more form of *composing operations yields a new abstraction*, from the programming-languages side, and conceded with the rest.

Multi-stage programming (Taha and Sheard 1997; MetaML and its descendants) builds, at run time, a program value that is then run with different arguments — the battery’s first three questions answered from the same side, and more literally than by any neighbour above: the composed definition exists at run time, is distinct from its executions, and parametrizes. What the staged program is not is anyone’s capability. It is a value in the host program’s hands; its persistence is the host’s own concern; no record holds its definition or its exercises; and no modality attaches to either. Stage separation answers *when* code exists — the plane here answers *whose record it enters*.

Macro-operators are the sharpest ancestor, and the concession here has a date: Fikes, Hart and Nilsson (1972) generalise a STRIPS plan so that problem-specific constants become parameters, store it as a triangle table, and reuse it — in whole or in part — as a single macro action in later problems. *One definition, many exercises, with different bindings* is fifty years old, and this paper claims no part of it. **Options** (Sutton, Precup and Singh 1999) do the modern form: temporally extended actions added to an agent’s set, used interchangeably with primitives. Both put a composed capability into a subject’s repertoire.

What the battery still separates, and it is the same two questions each time: the MACROP lives in a plan library and the option in a policy set — neither the definition nor its exercises enters the subject’s *history*, and nothing attributes modality or historicity per exercise, because in a planner there is nothing for that distinction to govern. The constituents, moreover, are the agent’s own operators; none of these ancestors composes repertoires their subject does not own and cannot name (§5.2).

Definitions emitted at first execution have an ancestry of their own, and §2.2’s runtime-emitted definition belongs to it. Programming by demonstration records a reusable macro from an enacted sequence (Cypher 1993); a trace-based JIT emits a reusable compiled trace the first time a path runs hot (Bala, Duesterwald and Banerjia 2000); and procedural reflection made the running program an object within its own system (Smith 1984) — a tradition made operative by the metaobject protocol (Maes 1987; Kiczales, des Rivières and Bobrow 1991), in which a running system extends its own executable vocabulary. All conceded — with the by-now familiar difference of address: what those systems emit or extend lands in a macro library, a code cache, or the image’s own machinery — an artifact of the program, not an entry in the emitting subject’s record.

7.2 Durable execution

The nearest living neighbour is not a pattern but a shipped runtime family, and the first round of discovery strings missed it; the adversarial-review channel of Appendix B supplied it, and the index pass then reached it on its own under revised strings — the miss and its bound are both recorded there. In **Durable Functions** (Burckhardt et al. 2021), an orchestration is a parametrized definition invocable many times with different arguments; the activities it composes come from libraries the orchestration does not own; *entities* supply actors with state and a journaled history; the runtime records invocations and re-executes by replay — and the paper proves the equivalence of the high-level semantics with the record-replay implementation. Orleans is the actor ancestry (Bykov et al. 2011). Temporal and Cadence carry the same model as industrial practice without a refereed venue; that Temporal’s replay has the same shape measured below — the workflow’s code re-executed, its activities run once with results cached, determinism enforced by rule — is documented from the actor community’s own side (Bernhardt 2021), which extends the engine measurement to that sibling documentarily.

Put the battery to it without flinching. The first three questions are satisfied outright. So is much of §5: the composed pieces need not be the orchestration’s own. And the determinism frontier this paper met at §4.3 is there *formalized and enforced* — the prohibition on reading the clock or minting an identifier inside an orchestration, with replay-stable substitutes supplied, is the parameter plane as a language rule. Nothing about that frontier is this paper’s, and §4.3 now says so.

What the criterion of §3 still separates is stated here from the published semantics rather than measured, and is flagged as such. In durable execution **the definition is code**: the orchestration function lives in the program, and the history stores the events and outcomes of its activities; replay re-executes the code against the recorded history to reconstruct position. Ask §3’s question (a) of that history and the answer is invocations and results — closer to constitution than any neighbour above — but the constitution itself never enters the record: change the code, and the same history replays against a different definition, which is precisely the versioning hazard that model’s own documentation manages. In the arrangement of §2 the definition *is* an entry; the constitution is data in the subject’s history, and §2.3’s second use referenced it there. The hazard is thereby narrowed, not escaped: the repertoires’ implementations stay ambient in both models, and a changed operation replays the same entry differently — what moves into the record is the composition, and §8 carries the remainder as a limit. The second separation stands whole: durable execution attributes no command/query modality to any exercise — recording is how execution works, not a decision anyone takes per performance, so *which performances count as the subject’s history* is not a question the model offers to anyone. And the first separation is not documentary either: Appendix C’s Proposition 5 derives it from the shape of the arrangements — its negative half

needs neither the host axiom nor anything about this engine or that one, only that an ambient program is free to vary — so what changes with a runtime’s documentation is the illustration, not the result. Both separations are now more than documentary. §3’s coordination arm was run on this family’s own engine — the Durable Task Framework under its in-memory emulator, with the history read through its public dispatch middleware — and the measurements landed where the semantics say: the record holds all six invocations, and replay re-executed the orchestrator’s code in six replay episodes while re-performing no operation (Appendix A: `TheCoordinationArmOnAThirdPartyEngineLab.cs`). Its own authors describe the model in exactly these terms — “an abstraction of a persistent call stack”, every step “transparently persisted by the server”, and “a workflow is a durable actor with all its state fully preserved” (Bykov and Fateev 2021) — a stack persisted, a definition that stays in the program. The versioning hazard remains cited from their semantics rather than measured.

7.3 Behaviour acquired by an object

One neighbour is closer than any pattern, and it was surfaced by the adversarial-review channel of Appendix B rather than by the paper’s own genealogy. **Talents** (Ressia, Gîrba, Nierstrasz, Perin and Renggli 2014) are dynamically composable units of reuse applied to *an individual object*: a talent specifies methods that can be added to — and removed from — the behaviour of one object, without touching the other instances of its class. That is dangerously near *a subject acquires a repertoire of its own*, and the paper cannot brush past it.

The answer is in what each thing adds and where. A talent adds **methods to the object** — the object’s own behaviour grows, and the new methods live where its old ones do. Assembly adds **nothing to any object**. The repertoires’ types are unchanged — §5.2’s are shipped binaries — and V is constituted on another plane, out of operations that remain exactly where and what they were. What grows is not an object’s method set but a subject’s repertoire, and the subject is not one of the objects. From there the last two questions of the battery finish it: a talent’s methods have whatever modality methods have, and no notion of which of their invocations count as anyone’s history.

Prototype languages put the same acquisition in the object system itself — in Self (Ungar and Smith 1987) an object holds behaviour without a class, and can acquire it — and the disposition is Talents’: the repertoire that grows is the object’s own, the growth lives in the running image, and no record of the definition or its exercises survives into any subject’s history.

7.4 Command sourcing

Recording invocations and replaying them through their handlers is an established alternative to recording outcomes, and it is the neighbour closest to §2.1’s mechanical property. Two things separate it. Each recorded command is self-contained — there is no derived definition that later entries reference, so the

record holds exercises of nothing. And the known objection to replaying commands — re-execution meets nondeterminism — is precisely the boundary the parameter plane draws (§4.3): what crossed the plane replays; what was minted behind it does not, and this paper reports its own failing case there.

7.5 The positive control: the materials without the category

No new neighbour is needed for this one; two already present supply it jointly. Command sourcing (§7.4) and the process manager (§7.6) live in the ordinary event-sourced stack: commands received, events appended to a journal, state rebuilt by replay, reactions and coordinators reading the log. Every material of §2 is present in that stack, at industrial maturity — and the unit its literature is about remains the aggregate, the handler, the process. The question its records answer stays *given this state and this event, what follows?*, not *which verb did this subject acquire?*

That is not a shortcoming to be argued with; it is a control this paper needs. Having the materials does not by itself produce the category — a system can hold commands, logs, events and reactions and never constitute a verb of a subject. Which is exactly why the category had to be measured (§2–§3) rather than asserted, and why this paper never claims that the pieces were unseen.

7.6 The process manager

One arrangement deserves its own subsection because it is the one this paper is most easily mistaken for competing with, and it is conceded entirely.

A process manager (Hohpe and Woolf 2003; its transactional ancestor is the saga, Garcia-Molina and Salem 1987) relates already constituted acts into a trajectory: correlation, lifecycle, resumption, compensation. It is mature and well understood, and nothing here improves on it, competes with it, or is offered as an alternative to it. The two relations are orthogonal, and the difference is not temporal — an assembled verb may contain operations separated by any amount of time; verticality is not simultaneity:

The horizontal relation presupposes its acts. The vertical relation constitutes one.

The same community also runs the arrow in reverse, and the reverse deserves naming. **Process mining** (van der Aalst 2016) discovers a process definition from event logs — from exercises to a definition, the inverse of §2.2, and precisely the archaeological direction §6 declined for itself. What discovery yields is an analyst’s model *about* the subjects whose logs were mined: it lands outside every subject’s record, no subject can exercise it by reference, running it executes a model rather than re-performing anyone’s acts, and it attributes nothing — its input is whatever the systems happened to record, which is the attribution

this paper locates at assembly. The kinship is exact and inverted: mining reconstructs, with effort and approximation, a definition the record never contained; constitution writes the definition into the record so that nothing ever needs to reconstruct it.

Because the axes are independent, the two appear together naturally: five operations may constitute *Purchase* vertically while *Purchase*, *Fulfilment* and *Delivery* form a coordinated trajectory horizontally. **An assembled verb may itself be one of the acts a process manager later relates; constituting an act does not remove the need to coordinate acts.**

A process manager answers what follows this act. An assembled verb answers what belongs to this act.

The observable separation between the two is §3's, and it is about what one record contains — not about identity, streams, state or names, all of which a process manager legitimately has.

7.7 The older neighbourhood

That some acts are *constituted* rather than merely caused — that a movement of pieces can **count as** an act of another subject — is not a question software raised first. It is formalised in deontic logic and in the study of normative multi-agent systems, where *counts-as* is a term of art (Jones and Sergot 1996), with the constitutive-rules lineage running through Austin (1962) and Searle (1969, 1995); none of that is claimed as new here. From the process-modelling side, business-artifact approaches centre a process on an entity with a lifecycle and a history of its own (Nigam and Caswell 2003) — subject-shaped modelling, without the constitution or attribution this paper measures. Enterprise ontology reaches the adjacent ground from organisational modelling: in DEMO (Dietz 2006), an organisation is a network of actors and transactions, actor roles are the units of authority and responsibility, and authority is assigned on the basis of competence — a normative account of who may perform what, stated for organisations of people. This paper's §4.2 is the mechanical shadow of that question. What this paper adds is not the question but a mechanical residue of it in an ordinary artifact: the boundary of a composing party's reach is measurable in a build, with the normative half explicitly left open (§4.2), and whether a composition was constituted or merely coordinated is decidable from a record (§3). The literature supplies the concepts; the arrangement supplies something the concepts can be checked against.

7.8 The neighbours in one table

Seventeen rows: the sixteen neighbour families of Appendix B, plus the event-sourced stack, which is §7.5's positive control rather than a family the search found.

neighbour	already does	the question this paper adds
procedural abstraction	definition exercised many times, in program text	what remains of the composition after a run?
Command / MacroCommand	the request as object; macros from commands; an undo history	whose capability, and who attributes its modality?
workflow / subworkflow	a definition instantiated as cases	whose history is the audit trail, and who chose what enters it?
algebraic effects	composed operations as a new first-class abstraction	a subject, a record, a modality — none is in the construction
multi-stage programming	a program builds a program at run time, reused with arguments	a value in the host's hands: whose capability, whose record?
Talents	methods added dynamically to one object	V adds no method to any object; the repertoire that grows is another subject's
reflection / MOP / prototypes	a running system extends its own executable vocabulary	the growth lands in the image's machinery, not in any subject's record
transaction script / application service	A; B; C as a use case	did a capability of a subject appear?
macro-operators, options	a parametrized composition enters the agent's repertoire	do the definition and its exercises enter the subject's <i>history</i> ?
macro-by-demonstration, trace JIT	a definition emitted at first execution	emitted into whose record?
durable execution	parametrized orchestrations, replayed from a journaled history, determinism enforced	is the <i>constitution</i> in the record, and who attributes modality?
command sourcing	invocations recorded and replayed	exercises of <i>what?</i> and what replays deterministically?
access audit logging	reads journaled because the disclosure must stand	attributed per performance, at assembly, into the subject's own history? (§2.4)

neighbour	already does	the question this paper adds
the event-sourced stack	commands, events, logs, reactions — every material of §2	the control: materials alone do not produce the category
process manager	relates acts into a trajectory	which operations constitute one act?
process mining	recovers a process definition from event logs	the definition lands in an analyst’s model — the inverse of writing it into the record
counts-as, DEMO	constitution and authority as concepts	a build and a record the concepts can be checked against

The neighbours attack different rungs, and none needs to be defeated — each is conceded at its own task. Composition has been made first-class, definitions reusable, behaviour dynamically acquirable, coordination explicit and constitution formal — separately, and in combinations that come close. What this search did not find said together is the conjunction — stated as a finding of the search reported in Appendix B, its adversarial-review channel included, scoped to its examined set, rather than as a proof of novelty:

Assembly is a plane at which operations available from repertoires become verbs of another subject; V is distinct from its exercises; and each exercise’s command/query character — and therefore which performances belong to that subject’s history — is attributed at that plane rather than inherited from the constituent operations.

The pieces of this paper are not new. The claim is about the planes they usually collapse: that *what can be computed*, *what the subject can do*, *what the subject does* and *what counts as its history* are not the same decision — which is the register this series has used from the start, discovering that things habitually fused are separately decidable.

8. Limits, and What Is Not Claimed

The cost measured is the record’s; no advantage is claimed. Three prices are measured against the direct baseline the criterion already uses (Appendix A: `WhatTheRecordCostsLab.cs`, the run published as `data/paper0A-assembled-verb/price-release.log`), every timed quantity over ten runs and reported as its median with the min–max spread, on a declared environment: Release build of engine and harness, .NET 9 on Windows

11, one 32-thread desktop processor, workstation GC, wall clock, the corpus assembles prebuilt binaries in both arms — and the direct baseline performs the same three statements and allocates the same domain objects, so what it lacks is only the record. *Growth* (exact counts, not samples): the act carrying the definition and first exercise costs 554 bytes of journal; each additional exercise costs ~112 — reuse is visibly cheaper than restating, which is the capability claim’s material side. *Replay*: reconstructing the subject from a record of 100 exercises takes 4.6 ms [4.4–5.6]; from 1,000, 15.0 ms [13.9–20.2]. *Throughput*: the plane performs ~2,490 exercises per second live [2,310–2,562], where the same statements as direct in-memory calls run at ~529,000 loops per second [476,000–542,000] — a ~**210x** price, the ratio of the two medians, reported to two significant figures because that is what the instrument supports: repeated invocations of the lab, each with its own warm-up and ten samples per arm, move the multiplier by a few percent. Two sampling decisions are part of the measurement rather than incidental to it. A warm-up run is discarded in every arm, so that JIT compilation is not sampled as steady-state work. And a direct *sample* is ten loops of the same thousand, each over a fresh register, timed together and divided back: one loop of a thousand lasts about two milliseconds — the scheduler’s noise floor rather than a measurement — while scaling a single loop to ten thousand would have changed the work itself, since the register would grow ten times larger. The comparison is asymmetric by construction and the asymmetry is the finding: the direct loop buys execution only; the plane’s loop buys execution plus a definition exercised by reference, an attributed modality, and a record that reconstructs the subject. The decomposition the two medians support: replay re-performs the same thousand exercises in ~15 ms — about eight times the direct loop’s median, against the plane’s ~210 — which locates the dominant term outside the plane’s interpretation, in the durable write — an inference from two measured quantities rather than a third measurement. What remains unmeasured: concurrency, scale beyond one machine, and every competing arrangement’s own price — whether this arrangement is cheaper or faster *than any other* is unexamined.

Two implementations, one author. The plane’s realizability no longer rests on a single engine: a prototype of a few hundred lines implements it on a third-party engine, the same engine that carries a coordination row, so the substrate is controlled (§3.2). What remains true is that both implementations are this paper’s. An implementation by an independent party, put to §3’s criterion, is the measurement this paper cannot supply itself.

The corpus is two systems, selected for inspectability and reputation, not representativeness. The census was run after the selection, and its zero is a statement about these two systems: *the plane did not appear of its own accord in either*. Whether it fails to appear elsewhere is not asserted, in the abstract or anywhere — no population-level claim is made. And the census’s coding employed no independent coder: the codings that agree are the author’s reading and rules of the author’s writing (§6). The published codebook makes an independent recoding cheap, and none has yet been performed.

Ownership is measured; subjecthood is assumed. The paper establishes ownership relative to the actor whose record contains the definition and its exercises: V belongs to no constituent repertoire (§5), its exercises are attributed to one actor, and replay reconstructs that actor from its own record. It does not derive subjecthood from constitution — and Appendix C says the same of itself: its Op/Verb partition takes the subject as given, so the calculus carries the separations *given* a subject and derives no subject from them. That the owner is a subject — the frame in which its record is a history and its outputs are speech — is inherited from the actor model (Hewitt, Bishop and Steiger 1973) and from this series’ prior finding that identity precedes staging. A reader who declines the frame keeps every measurement and loses only the names.

The entitlement reading remains a reading. §4.2 concedes that reaching past host-language accessibility is a runtime capability, not an author’s permission. A reader who declines to call the subject’s position a *right* loses nothing measured in this paper.

One value in the composed verb returns fresh on every replay, and no one at the plane has the authority to stabilize it. An identity minted inside a repertoire’s factory changes on every replay; the criterion’s equivalence absorbs the freshness, and Appendix C proves that exact identity is unavailable under its host axiom while that equivalence is — so what stands is not a replay failure but an **authority boundary**, reported against the paper’s own flagship entry: the assembler can neither supply the value nor stabilize it, where durable-execution runtimes stabilize theirs by design and a domain’s author could by asking (§4.3). The calculus also states the condition under which the guarantee does hold — every reference to a minted value travelling between acts as a lookup under a key that crossed the plane — which is a characterization of when replay is safe, not advice about how to write verbs (Appendix C). What the arrangement supplies is otherwise detection rather than immunity: the corpus’s own suites never meet this value — its 49 tests pass over it — and replay is the instrument that makes the undeclared inbound visible, against the actor premise, a standard the repertoire never adopted, which is why §4.3 charges neither party. The value is also the measured instance that forces §3’s outcome equivalence to be declared before use (§3.1). The subject inherits a decision it cannot author around. Completeness, executability and testability are the repertoire’s properties; a replay-stable minted identity is not among them, and no assembler can add it from outside — this corpus separates the four.

A captured mint breaks replay equivalence outright, and this one the calculus found first. Past the licensed freshness sits a sharper hazard: a minted identity read by the assembler and *captured* into a later act’s recorded arguments. The recorded literal names a value the live run held and the replayed run never minted, so no renaming relates the two (Appendix C, Proposition 4’s boundary). Measured, the composition’s own invariant — the register’s key is the identity of the payment kept under it — holds live and fails after replay; the same composition referring to the value by a supplied key holds through both,

its mint just as fresh (Appendix A: `WhereReplayEquivalenceBreaksLab.cs`). This is the paper’s one finding obtained by reasoning about the arrangement rather than by measuring what it was built to do: the derivation is committed before the lab that confirms it exists, and both commits are cited in Appendix A, so the order is checkable rather than asserted. Durable-execution practice prevents the case by design, supplying replay-stable identifiers from the runtime; this arrangement leaves it to what the assembler records, and the calculus states the sufficient condition under which the guarantee holds.

The record closes over the composition, not the operations. Replay re-performs by reading the repertoires’ code, which lives outside the record: change an operation’s implementation, and the same record replays to a different state. The criterion holds the engine and the repertoires fixed in (b) for exactly this reason, and Appendix C proves the relativized form: replay is a function of the record *and the interpretation* in every arrangement measured, and what distinguishes constitution from the strongest coordination is the remainder — an ambient program carrying the composition there, nothing but the operations here (Proposition 5). The versioning hazard §7.2 cites against durable execution is therefore inherited at the repertoire level, not eliminated. A subject that needed its operations’ semantics to survive their authors’ edits would need more than this arrangement supplies.

The query modality is clean about the record, not about state. A query leaves no entry, and that is what §2.4 measures. What follows for state is narrower than it may read: the engine prevents a query from persisting, and in an arrangement where durable state *is* reconstruction from the record, what a query touched does not survive rehydration. Repertoire state held outside that record — statics, caches, external stores — is reached by a query like any other act, and returns from nowhere. Appendix C carries the same boundary as its second idealization.

Reuse is measured for one modality. §2.3 shows the verb drawn on again, and shows it for the journaled command, where exercises persist and can be counted. Query exercises may leave no record, so the same measurement cannot be run on them; their standing as exercises of the same capability is argued from the definition plane (§2.4), not measured on a record. The paper keeps the two apart.

Specialization is not measured at all. Whether assembled verbs can be *inherited* or *specialized* — whether a subject’s repertoire admits the structure that types give operations — is future work, and nothing in this paper asserts it. What is asserted is only the precondition such work would need: that the verb exists as a unit at all.

Recognition is a corollary, not a contribution. A reaction can bind values from two constituent operations of the composed verb precisely because both are in one recorded entry — the correlation is carried by the act itself. The mechanism runs on both planes of §2.3 at

once: the pattern executes against the verb's cached definition while the captures read the exercise's journaled arguments (`Reaction.cs:1770-1795`), and entries whose composition was consumed at execution are explicitly skipped, having nothing stable to hold (`Reaction.cs:1825-1843`; Appendix A: `TheTrajectoryBecomesObservableLab.cs`). That is reported as a consequence: once something exists as a unit in the history, the rest of the system can refer to that same unit without reconstituting it. Reading those units — what recognition is, and what it entitles — is the next paper's subject, not this one's.

9. Conclusion

The paper opened with two arrangements side by side and one question: what exists in the second that did not exist in the first. The answer was earned in an order, and each step of it declares its kind: a measurement where the claim is empirical, a construction or a proof where the distinction is definitional — the division §3.5 states and Appendix C carries — and a declared reading where it is neither.

A capability was **produced**, in the operational sense §2.3 fixes and flags: one definition, two exercises with different arguments, both replayed from the one record. Its recorded constitution is **distinguishable** from the arrangements it resembles, by a criterion any record can be asked — an event-sourced coordinator's record contains none of the constituting statements, durable execution's contains all of them yet re-performs none on replay, a trace contains all with no replay; only the constituted verb both contains and re-performs — an arrangement realized twice, once on this paper's engine and once, in a few hundred lines, on a third-party one (§3.2). It is **anonymous**, an integer, while the method and the trace carry the business names — so naming was never the property. Its **modality is attributed, per performance**: one body, one verb, left no history exercised as a query and was journaled, replayable, performed as a command, the domain unchanged either way. Its constitution **drew on repertoires that share no type** and were never edited. And the corpus, asked *where is V?*, answered with 72 named artifacts none of which names a trajectory spanning subjects — coded twice, by the author and by a published mechanical rule, in full agreement — and 49 of its own domain tests passing over the freed pieces — the ingredients everywhere, the capability nowhere.

What survives the limits of §8 is therefore not a proposal but a conjunction, found standing among the examined neighbours of a search built to knock it down (Appendix B): a capability of a subject, constituted from repertoires the subject does not own, whose definition and exercises survive into the subject's record, and whose modality — including which performances count as the subject's history — is attributed at the plane of assembly.

As this series has done before, the paper closes by handing the reader its instrument rather than its enthusiasm. Of any system, take the record that claims to

be the whole and ask: **how many of the statements that constitute the whole does it contain, and does replaying it alone bring the outcome about again?** And then the two questions no neighbour of §7 was found to hold: **who assigns each exercise's command/query modality, and who decides which performances enter the subject's history?** Systems will answer differently, and the answers are informative wherever they land; nothing obliges a system to have V, and §8 said so.

And every item just listed survives the flatter description §4.1 admits, for a reader who declines this paper's vocabulary: the criterion, the ambient hierarchy, the attribution of historicity and the census are findings about records and engines, not about the words for them. The declarations of §8 fence that field; they do not empty it.

Two things are deliberately left standing. Whether occupying the assembly plane is an entitlement rather than a capability remains the open reading of §4.2. And once wholes exist as units in a subject's history, what *reading* them amounts to — what recognition is, and what it entitles — is not begun here. A subject that can acquire verbs has a history worth reading, and that is the next question, not this paper's.

The last consequence is the one the title was carrying all along, in the word *authorship*. The repertoire need not know the verbs in which it will participate. The subject need not have authored the operations from which its verbs are made. Between those two freedoms sits one operation:

Assembly turns the use of capabilities into a capability of a subject.

That is the finding. The consequence the series takes forward is the line after it, and it is the one §1.1 promised and §6 measured from the other side: the assembled verb matters less because it adds one capability than because it gives continuity somewhere to live other than in the capabilities it joins. The dissection showed the pieces already computing on their own; the carrier measurement showed the wires between them delivering the same values three ways — supplying *and then*, not meaning; and the long scenario showed two repertoires that never met joined in one verb, unmodified. The pieces kept their domain meaning without the machinery that carried their continuity. Give that continuity a place of its own, and:

Once the composition has a place of its own, continuity no longer has to be authored into its constituents — and the constituents need not know the future wholes in which they will participate.

Disclosure

Large language models were used in the preparation of this work: as the adversarial-review channel reported in Appendix B, and as drafting and verification assistance under the author’s direction. Nothing entered the paper on a model’s assertion alone — every contributed source was verified against the primary record, every count resolves to the labs or the dissection, and the model identities and versions are recorded at deposit. Venue-specific disclosure requirements are checked at each deposit.

Appendix A — The Labs

Every count in the body resolves to one of two published artifacts, and the split is stated so it can be checked. The measured counts — the criterion’s cells, the reuse and modality counts, the census’s 72 and its 27 + 16 + 29, the prototype’s counts, the record’s prices — resolve to assertions or printed output of one MSTest suite, published in the papers repository as `labs/paper0A-assembled-verb/`. The dissection’s counts — the 49 author-written tests passing over the freed pieces, and the per-rule hole counts — resolve to the dissection’s audited history: two self-contained git bundles in `paper0A-assets/dissection/`, replayable as ordinary clones.

Pins. The engine is the public repository (`github.com/alvaronCubo/puppeteer`), cloned as the sibling `../.../puppeteer` every lab in that repository uses, at commit `3160e39` — re-pinned, with every file:line anchor re-resolved, at deposit. The corpus is `dotnet/eShop` at `9b4f943` and `kgrzybek/modular-monolith-with-ddd` at `91c8ef2`, both cited in References — the same commits the dissection bundles are based on, so the binaries the paper composes and the sources it dissects are one corpus, pinned once. Neither corpus repository is redistributed: `fetch-corpus.ps1` clones both at the pinned commits and builds the three assemblies the suite references — the two domain repertoires the body composes, plus the seedwork assembly one of them depends on.

Run. `./fetch-corpus.ps1` once, then `dotnet test`. Twenty-one labs, all green. The suite compiles against the engine’s public surface only — no signing and no internals grant.

The labs the body cites:

lab	§	measures
TheCompositionBecome sACapabilityLab	§2.2–2.3	one definition, two invocations with different arguments, replay re-performs both from the one record

lab	§	measures
WhoDecidesWhatCounts AsHistoryLab	§2.4	one body, one verb, three performances: query 2→2, command 2→3 (journalized, replays), query again 3→3 — modality attributed per performance
TheCriterionForConst itutionLab (+ Process ManagerHarness)	§3	the event-sourced coordination arm (the corpus's shape, author-built and labelled): 0/no · description 6/no (nothing to replay) · constitution 6/yes
TheEventSourcedCoord inatorOnAThirdPartyE ngineLab	§3	the event-sourced coordinator on Orleans (in-process TestingHost, log-storage provider): 6 transitions and 0 operations in its own journal; reconstruction restores position and re-performs nothing
TheCoordinationArmOn AThirdPartyEngineLab	§3, §7.2	the coordination arm on the Durable Task Framework (in-memory emulator; history via its public dispatch middleware): 6 invocations in the engine's own record; orchestrator code replayed in 6 episodes; no operation re-performed
RepertoireBoundaryLa b (+ RepertoireBounda ryFixture/)	§4.2	the compiler refuses the host (CS0122, obtained by compiling a program written to fail); the assembler reaches the internal piece

lab	§	measures
RetrievalIsALookupLab (+ SubscriptionPaymentBridge)	§4.3	the licensed mint, and nothing about capture: the assembler's key stable across replay, the domain-minted identity fresh, retrieval a lookup in the subject's own state; the minting site marked R_99
ALongTrajectoryOverTwoUntouchedRepertoiresLab	§5.2	17 operations, 4 aggregates, 2 shipped assemblies sharing no type, 0 doubles; 9 events (corpus-emitted, author-composed scenario), none naming the trajectory
WhatTheCarrierCarriesLab	§6	one act carried three ways delivers identical values; only the repertoire-minted identity differs
EveryNameStaysWithinOneSubjectsVocabularyLab	§6	the census under two coding rules: population 27 + 16 + 29 = 72 reproduced mechanically; rule one (2 aggregates in one name) codes 0 T; rule two hands its blind spot — zero-aggregate names, process words — to manual coding: 3 residuals, each published with its reason, all S; agreement 72/72 (§6 on what that agreement is not)
TheTrajectoryBecomesObservableLab	§8	two reactions over unrelated patterns answer one question identically, correlated by the recorded act

lab	§	measures
TheConstitutionProto typedOnAThirdPartyEn gineLab	§3.2, §8	the constitution arrangement prototyped on Orleans, the coordination row's engine: 1 definition and 2 exercises in the engine's own journal; a query performance leaves it unmoved; reconstruction re-executes every journalled statement; the domain-minted identity differs per replay
WhereReplayEquivalen ceBreaksLab	§4.3, §8, App. C	the calculus's prediction, put to the engine: a minted identity captured into a later act's recorded arguments breaks the composition's invariant on replay (true→ false), while the same composition referring to it by a supplied key holds (true→true) — the mint fresh in both arms. Order: derivation 0001735 (15:08), lab 08 13fec (17:06), same day

lab	§	measures
WhatTheRecordCostsLab	§8	the record’s price, Release, n=10 after a discarded warm-up, median [min..max]: 554 bytes for definition + first exercise, ~ 112 per additional (exact); replay 4.6 ms at 100 and 15.0 ms at 1,000 exercises; ~ 2,490/s through the plane vs ~ 529,000/s direct — ~ 210x ; published as <code>data/paper0A-assembled-verb/price-release.log</code>

The dissection. The bundles replay as ordinary clones (`git clone eshop-dissection.bundle`); `git log --reverse` lists the pre-registered rules in the order applied with each step’s counts in its commit message, and `git blame` on any surviving line names the rule that spared it. The rule file — with its amendments, its stopping rule, its unequal-objectivity declaration and the withdrawn R7 — is also copied out beside the bundles for direct reading.

Additional instruments retained in the suite but not cited by the body are listed in the lab’s README; they guard the same corpus the cited labs run against.

Appendix B — The Search, Reported as Method

The claim this appendix underwrites does not rest on any search having been exhaustive. It is scoped as *absence over the examined set*, and §3’s two questions with §7’s battery are a finite decision procedure for any candidate beyond it — a form that survives the discovery of a further neighbour, which matters below, because one of this search’s channels is of a kind whose coverage cannot be characterized. The search is reported anyway, in the light form of Kitchenham and Charters (2007), because a claim of the form *not found said together* is a claim about a search, and an instrument must be reported — including what it is not: **this was not a systematic literature review**.

Objective. Find a construction holding the conjunction of §7.7: a capability of a subject, constituted from repertoires the subject does not own, whose definition and exercises survive into the subject’s record, with modality attributed at assembly.

Sources, with dates.

- *General web search.* Discovery strings, verbatim, run on 15 and 17 August 2026 — the claim being scoped to what was examined, not to when the examining stopped: `composite command pattern macro command first-class reusable composition object-oriented design`; Dietz DEMO enterprise ontology actor role competence authority responsibility transaction pattern production act; event sourcing "command definition" recorded journal replay "invocation" reuse parametrized action definition once invoked many. Verification strings — used to confirm candidates' bibliographic facts, not to discover — covered the 1972 macro-operators paper, the options framework, Talents, and durable functions. One further string was seeded with its intended answer, is recorded here as biased by construction, and its subject was excluded from the paper on provenance grounds.
- *Open bibliographic APIs* (17 August 2026). **DBLP**, title-level, eight strings: `durable execution workflow` — 2 hits, both 2026 agent-workflow papers, screened out (neither concerns a subject-owned record of a definition and its exercises); `first-class workflow` — 1 hit, workflow-centric research objects (2012), screened out (the first-class citizen is a scholarly artifact, not a capability of a recorded subject); `workflow as code` — 9 hits, screened out (policy-as-code, LLM-assisted workflows, one adoption study; none concerns such a record); `stateful serverless replay`, `composable capability actor`, `macro operators planning reuse`, `action definition journal replay`, `reified action record subject` — 0 hits each. Three of those five are this paper's own coinages, and a zero at title level is evidence about the string, not about the literature: the zeros are recorded as weak instruments, not as absences. **OpenAlex** — which indexes the metadata of the paywalled venues, narrowing the limitation declared below — the same eight strings, relevance-ranked, top six screened per string: the pass surfaced Burckhardt et al. (2021) at or near the top of two strings, independently reaching the family the adversarial-review channel had contributed, which bounds the earlier miss; the remaining top results belong to families already dispositioned in §7 (the workflow-patterns literature, the Actors model, macro-operators) plus one family screened out here: deterministic record-and-replay for debugging, which records one execution's nondeterminism so that run can be reproduced — an account of a single run, with no definition, no reuse, and no subject owning the record. **arXiv**, three phrase strings — 0 entries each. No candidate holding the conjunction surfaced in any row.
- *Paywalled indexes.* ACM DL, IEEE Xplore and Scopus were not directly searchable from the execution environment; their metadata is covered by the OpenAlex pass above. The residual limitation is full-text search within those venues.

- *Adversarial review by large language models.* Successive drafts were submitted to large language models instructed to act as program-committee reviewers for a software-engineering venue, across review rounds between 15 and 18 August 2026, each round given the current draft and, where applicable, its prior reviews; the model identities and versions are recorded at deposit alongside the archival snapshots. This channel contributed **seven neighbour families** the string-based rounds above had not surfaced: Talents, the workflow definition/case separation, and algebraic effects; and, in later rounds, durable execution (the Durable Functions semantics with its Temporal/Cadence practice), process mining, multi-stage programming, and computational reflection with the metaobject protocol. The channel’s standing has to be stated exactly, because it is easy to overstate: **it is neither corroboration nor expertise.** Sessions of systems of this kind share training corpora and therefore share blind spots, so agreement between them evidences nothing; their coverage of any literature is unknown and unmeasurable; no output is reproducible across runs; and none can be cited as authority for any claim. Every contributed family was verified by the author against primary sources before entering §7, and the bibliographic facts of each citation were confirmed against the publishers’ registers — nothing in this paper rests on a model’s assertion. What the channel does evidence is the **incompleteness of the string-based rounds**, and that is why it is reported rather than absorbed: seven families the strings did not reach are seven measurements of what the strings’ vocabulary did not cover. A verification pass then located the deficiency: under each field’s own vocabulary — one title-level string per family — the index instrument alone reaches **all seven** (the Talents article itself, the Durable Functions paper itself, and hit counts from 23 to 199 for the rest). The misses were the first round’s vocabulary, not the instrument’s reach, and both facts are retained here as the measurements they are.

Inclusion / exclusion. A candidate entered examination if it plausibly held *any* component of the conjunction; it was then put to the battery of §7. Name-only matches — systems merely called “capability”, “verb” or “assembly” — were excluded.

Examined set and disposition. Sixteen neighbour families, each dispositioned in §7 or where it arises: procedural abstraction; Command and Macro-Command; transaction script and the application service; workflow definition/case; algebraic effects; multi-stage programming; macro-operators; options; macro-by-demonstration and trace JITs, with procedural reflection, the metaobject protocol and prototype languages; Talents; command sourcing; durable execution, with its Orleans ancestry and its Temporal/Cadence practice; the process manager and saga; process mining; counts-as, DEMO and business artifacts; and access audit logging, dispositioned where it arises (§2.4).

Snowballing. Backward from the anchor citations (Gamma et al.; Fowler; Fikes et al.; Sutton et al.). Forward from Orleans (Bykov et al. 2011) to the

Temporal practitioner discussion and Bernhardt (2021). DEMO entered from this series’ prior related-work notes.

Threats. Not systematic; a single screener; web ranking is not reproducible; relevance-ranked sources were screened to their top six results per string; full text behind the paywalled venues was not searched; grey literature was reached only through the practitioner sources cited; a zero-hit string measures its own distance from the field’s vocabulary, not any absence in the field; and seven neighbour families entered through an adversarial-review channel whose coverage cannot be characterized and whose outputs are not reproducible — their absence from the string-based rounds is reported as a measured deficiency of those rounds, located by the verification pass in the first round’s vocabulary, and no agreement between sessions of that channel is treated as corroboration. Each of these argues for the scoped form of the claim, and none is argued away.

Appendix C — The Separation, Stated Over a Small Object

§2.4’s four questions — what is computed, what the subject can do, what it does, what counts as its history — are distinctions, and §2 and §3 measure them on running arrangements. A distinction, however, is *stated over an object*, and the smallest object that can carry these four is worth writing down. This appendix gives it: a core calculus, in the structural style of Plotkin (2004), just large enough that assembly, exercise, modality attribution and replay have operational rules — and small enough that each load-bearing separation is a short proposition, with what is merely book-keeping labelled an observation. The labs of Appendix A are its conformance checks, one per claim: `TheCompositionBecomesACapabilityLab` for Observations 1 and 2, `WhoDecidesWhatCountsAsHistoryLab` for Proposition 3, `RetrievalIsALookupLab` with `WhereReplayEquivalenceBreaksLab` for Proposition 4 and its boundary, and the two coordination arms with `TheCriterionForConstitutionLab` for Proposition 5.

Syntax. A sequence $x_1..x_n$ is written x^* . Two disjoint name spaces are assumed: operation names f belong to repertoires the subject does not own (the set Op); verb names V are the subject’s own. The partition takes the subject as given, deliberately: the calculus carries the separations *given* a subject, and derives no subject from them (§8). That a verb name is, in the engine, an anonymous integer rather than a word (§2.2) changes nothing here — the calculus needs only that the name denotes.

$v ::= \text{literal} \mid n$	values; n a name
$e ::= v \mid x$	expressions
$s ::= x := f(e^*) \mid f(e^*)$	statements over the repertoires
$b ::= s ; \dots ; s$	a body (nonempty)
$m ::= \text{cmd} \mid \text{qry}$	modalities

```

a ::= define V(x*) = b                                assembly
    | perform V(v*) : m                                exercise; m is attributed here

```

Two scope decisions are made at the grammar and declared as such. Bodies are first-order: no statement invokes a verb, and composition of verbs out of verbs is out of scope here. And the calculus carries no speech of the plane's own — §5.1's **print** and **tell** are outside the object, which covers the composition of operations and not the subject's vocalization; the four questions need no more, and §5.1's claim exceeds this appendix on purpose.

Configurations, and what is given. $\langle D, \sigma, J \rangle$ — definitions $D : \text{Verb} \rightarrow (x^*, b)$; domain state σ ; the journal J , a sequence of entries $\text{def}(V, x^*, b)$ and $\text{act}(V, v^*)$. One constant comes with the calculus: an initial state σ_0 . A configuration is **reachable** when the rules below reach it from $\langle \emptyset, \sigma_0, \varepsilon \rangle$; every claim quantifies over these.

The host as a parameter — axiom (HOST). The repertoires enter the calculus as an interpretation governed by one axiom, named once and cited by name:

(HOST) $\llbracket f \rrbracket$ relates (σ, v^*) to (σ', v, N) , totally on (σ, v^*) , where N is the finite set of fresh names the application mints:
 $N \cap \text{names}(\sigma, v^*) = \emptyset$; $|N|$ is determined by (σ, v^*) and is invariant under bijective renaming of the names in (σ, v^*) ;
 (σ', v) is determined by (σ, v^*) up to bijection on N ; and
the relation is equivariant — renaming the names of (σ, v^*) renames the outcome accordingly.

Freshness is the *only* mechanism of the calculus for new names — there is no name supply in the configuration, so nothing re-determinizes the mint. Two applications of the same f to the same arguments agree except in *which* fresh names they drew: which is what the prose of §4.3 says, what the measured factory does, and what makes Proposition 4's bijection necessary rather than decorative. (HOST) is the closure condition §4.3 speaks of, written as an axiom: whatever an execution depends on either crossed the plane (it is in v^*) or is fresh (it is in N). It is a **closure**, and it is worth saying plainly what it closes out: external I/O, randomness beyond the minting the axiom licenses, foreign mutable state reached behind the plane, callbacks, and clock readings that escape it — §4.3's undeclared inbounds. One further axiom gives the calculus something to distinguish, and Proposition 5 cites it:

(NONTRIV) there exist f, g in $0p$ and (σ, v^*) whose applications differ in σ' .

Execution. A statement executes under a local environment $\rho : \text{Var} \rightarrow \text{Val}$ — the values flowing between statements, the third constituent §2.2 names — and emits its operation's result as output:

$\langle \sigma, \rho \rangle \vdash s \Downarrow \langle \sigma', \rho' \rangle ! v$

For $x := f(e^*)$: the arguments e^* are read under ρ (a variable ρ does not

bind is stuck), $\llbracket \mathbf{f} \rrbracket$ is applied, minting per (HOST), and $\rho' = \rho[x \mapsto v]$; for bare $\mathbf{f}(\mathbf{e}^*)$, $\rho' = \rho$. A body folds its statements left to right, $\langle \sigma, \rho \rangle \vdash \mathbf{b} \Downarrow \langle \sigma', \rho' \rangle ! v$, emitting the last statement's output. The fold always terminates: bodies are finite, invoke no verbs, and $\llbracket \mathbf{f} \rrbracket$ is total.

Rules.

$$\begin{array}{c}
\text{ops}(\mathbf{b}) \subseteq \text{Op} \quad V \notin \text{dom}(\mathbf{D}) \\
\hline
\langle \mathbf{D}, \sigma, \mathbf{J} \rangle \text{--define } V(x^*)=\mathbf{b} \rightarrow \langle \mathbf{D}[V \mapsto (x^*, \mathbf{b})], \sigma, \mathbf{J} \cdot \text{def}(V, x^*, \mathbf{b}) \rangle \quad (\text{DEFINE}) \\
\\
\mathbf{D}(V) = (x^*, \mathbf{b}) \quad |v^*| = |x^*| \quad \langle \sigma, [x^* \mapsto v^*] \rangle \vdash \mathbf{b} \Downarrow \langle \sigma', \rho' \rangle ! v \\
\hline
\langle \mathbf{D}, \sigma, \mathbf{J} \rangle \text{--perform } V(v^*):\text{cmd} \rightarrow \langle \mathbf{D}, \sigma', \mathbf{J} \cdot \text{act}(V, v^*) \rangle ! v \quad (\text{PERFORM-CMD}) \\
\\
\mathbf{D}(V) = (x^*, \mathbf{b}) \quad |v^*| = |x^*| \quad \langle \sigma, [x^* \mapsto v^*] \rangle \vdash \mathbf{b} \Downarrow \langle \sigma', \rho' \rangle ! v \\
\hline
\langle \mathbf{D}, \sigma, \mathbf{J} \rangle \text{--perform } V(v^*):\text{qry} \rightarrow \langle \mathbf{D}, \sigma, \mathbf{J} \rangle ! v \quad (\text{PERFORM-QRY})
\end{array}$$

Two readings of PERFORM-QRY are deliberate. The execution happens — the premise is the same judgment as the command's, and the output v leaves, which is what keeps a query performance observable at all — and the configuration keeps nothing: not the effect, not the trace. What a query computes is visible at the boundary as its output and durable nowhere. In the engine measured this is literal, because state *is* reconstruction from $\mathbf{J}$ (§2.4): what a query changed does not exist after rehydration.

Four idealizations, declared. (1) The define/perform split is primitive; in the engine the two arrive fused in one act and the handler splits them at the reuse seam, and §2.3 measures that it does. (2) PERFORM-QRY is a rollback rule: the premise executes the body and the conclusion restores σ whole, silently undoing whatever a state-changing operation did. That is a substantive modeling decision, and it is **laxer than the engine it models**, which prevents rather than undoes: a query runs under a read lock with the symbol table in read-only mode, and its parser rejects the statements that would persist (§2.4). The calculus admits a query whose body would change state and unwinds it; the engine refuses to run one. The rules' conclusions agree — the configuration retains nothing either way — and every claim below depends on that conclusion rather than on the path to it, so a proposition proved under the laxer rule holds a fortiori under the stricter engine. What neither reaches is repertoire state held outside σ altogether (statics, caches, external stores), which rolls back nowhere. (3) $\llbracket \mathbf{f} \rrbracket$ is a relation over one store: no aliasing across repertoires, no shared mutable state outside σ , no reentrancy. The shape assumes by construction the tractability of exactly the situation §5.2 exhibits as hard — two repertoires whose only meeting point is the assembler. (HOST) bounds what $\llbracket \mathbf{f} \rrbracket$ may *depend on*; this shape bounds what it may *be*. (4) $\llbracket \mathbf{f} \rrbracket$ is total: no operation fails, so the calculus expresses no failed exercise — misfire and its kin are outside the object.

Replay. $\text{replay}(J)$ is a configuration reached from $\langle \emptyset, \sigma_0, \varepsilon \rangle$ by taking each **def** entry as its DEFINE step and each **act** entry as its PERFORM-CMD step, in J 's order — *a* configuration, not *the*: under (HOST) each run of the replay draws its own fresh names, and the propositions quantify accordingly. For the coordination variant of Proposition 5, $\text{replay}_P(J)$ starts instead from $\langle D_P, \sigma_0, \varepsilon \rangle$, with D_P read off the ambient program: the variant's journal holds no **def** entries, so a replay started from $\emptyset$ is stuck at its first **act** — which is itself the point, stated operationally: the record alone does not carry what its replay needs.

The five numbered claims are not of one kind, and the difference is declared rather than left to be discovered. Observations 1 and 2 are **bookkeeping**: they record what the definitions make true — Observation 2 follows from the side condition $V \notin \text{dom}(D)$, a fiat of the calculus — and their value is exactness, not discovery. Propositions 3, 4 and 5 carry the **content**: each could have been false of rules shaped otherwise, and each has an engine measurement that could have disagreed.

Observation 1 (defining computes nothing). DEFINE leaves σ fixed; hence there are reachable configurations equal in σ and different in D . *Why.* Immediate from the rule's conclusion; the witness pair is any configuration and its DEFINE-successor. On the engine: §2.2 exhibits the definition as an entry — content, not effect.

Observation 2 (a capability is not its exercises). After **define** $V(x^*) = b$ and n command performances with arguments $v^*_1 \dots v^*_n$, J holds exactly one **def** entry for V and the n entries $\text{act}(V, v^*_i)$, none containing b and all referencing the same V . *Why.* DEFINE's side condition $V \notin \text{dom}(D)$ forbids a second definition of V , and PERFORM-CMD appends entries that carry a name and arguments only. The definition/exercise separation is therefore *data in J*, not an interpretation of it. The idealization at work is (1) of the declared four: in the engine, assembly and exercise arrive fused in a single act, and the handler splits them at the reuse seam — a known body is an invocation, a new one a definition (§2.3). The calculus takes the split as primitive; §2.3 measures that the engine realizes it.

Proposition 3 (what it does is not what counts as its history). Fix D, σ, V, v^* with $D(V)$ defined and $|v^*|$ matching its parameters. The command and query performances execute under the same premise — one judgment, the same body from the same state, the same output emitted — and differ in journal membership, which is therefore a function of m alone. And m is a component of the *performance*, supplied at the exercise rule: it is not computed from the body, the state, or the outcome. The calculus establishes the **location** of the choice, not its author — there is no attribution judgment here, and none is needed for the separation. *Proof.* The rules share their premise and differ only in conclusion; the emitted output is what keeps the query half of the witness observable, since the query's configuration is unchanged. Who occupies that position in the measured arrangement is §2.4's question, and its answer there

is the assembler. On the engine: §2.4 — the journal head unmoved under the query, advanced under the command, the domain computing the same value both times.

Definition ($\cong$, determination, and what it is for). Two configurations are **equivalent**, $C \cong C'$, when some bijection π relates them componentwise, mapping the names minted during C' 's trace to those minted during C 's and fixing every other name — the two traces are distinct, and π is exactly the correspondence between what each of them minted. A pair $(J, \llbracket \cdot \rrbracket)$ **determines** replay's outcome up to $\cong$ when every configuration any run of $\text{replay}(J)$ reaches falls in one $\cong$ -class. And $\cong$ is what §3.1's *outcome exists again* means, made precise: **the whole's outcome exists again exactly when the replayed configuration is $\cong$ to the original** — which is the same clause §3.1 states in its own words as identity up to fresh values, and the reason it must be stated before the criterion is used. Propositions 4 and 5 use exactly these notions. (The local environment of the execution judgment is written ρ ; the bijection here and in what follows is always π .)

Proposition 4 (the record constitutes, up to fresh names — and no closer). Assume (HOST); let $\langle D, \sigma, J \rangle$ be reachable, and let the trace be **journal-closed**: no entry of J — neither a **def** body nor an **act**'s arguments — contains a name minted during that trace. Names are values in this grammar, so a body can carry one and the hypothesis has to range over the whole journal, not only its arguments; what it says is that everything the record holds crossed the plane from outside, which is exactly what makes π fix J . Then every run of $\text{replay}(J)$ reproduces D exactly — definitions mint nothing, so that much is a strict equality — reproduces J (by journal-closedness no entry contains a trace-minted name, so π fixes every one of them: $\pi(J) = J$, hence $J_r = J$) — and reaches a state σ_r with a bijection π on minted names such that $\pi(\sigma_r) = \sigma$. The three conclusions together are $\text{replay}(J) \cong \langle D, \sigma, J \rangle$, which is what Proposition 5 cites. The renaming cannot be dropped: **exact identity is not available for every interpretation satisfying (HOST)**, and equivalence up to bijective renaming is. The witness is immediate — an interpretation that draws a different fresh name on each application satisfies the axiom and defeats name equality across runs — so §3.1's clause is proved necessary rather than conceded. No maximality beyond that is claimed, and none is needed. *Proof sketch.* Induction over the define/command sub-trace with hypothesis $\pi_k(\sigma_r^k) = \sigma^k$; query steps change nothing in the configuration, so they drop out. At the next command step both sides execute the same journaled v^* , which by journal-closedness contains no trace-minted name, so π_k fixes it; (HOST)'s equivariance applied to π_k transports the replay's application at (σ_r^k, v^*) to an application at (σ^k, v^*) — the live one — whose outcomes therefore correspond under an extension $\pi_{k+1} \supseteq \pi_k$ over the step's fresh names, equal in number by (HOST)'s count invariance. The extensions compose because fresh names are new on both sides. On the engine: the constitution row of §3.2, whose *yes* is stated under exactly this equivalence (§3.1), and §4.3's measured identity — fresh on every replay while the supplied key holds — which

is π made visible.

Where the hypothesis fails. Let a trace-minted name be *captured* — read by the assembler and passed into a later act’s journaled arguments. The live application’s argument then aliases a name of the live state; the replayed application’s argument is the recorded literal, which aliases nothing in the replayed state, since replay minted fresh; and no bijection can both align the states and fix the literal. Replay equivalence breaks — not at the mint, but at the **capture** of the mint.

The converse is the calculus’s usable half: if every reference to a minted value travels between acts as a lookup under a key that crossed the plane, rather than as the value itself, the trace is journal-closed and Proposition 4 applies. That is a sufficient condition — a characterization of when the guarantee holds, not advice about how to write verbs.

Both halves were then measured, in a lab written after this derivation (§8 for the finding and its standing; Appendix A for the arms, the numbers and the commit order). One clause of that measurement belongs here rather than there: what the lab observes is a composition’s invariant holding live and failing after replay, and that this *is* a violation of $\cong$ is the calculus’s step — a configuration whose invariant fails cannot be related by any renaming to one whose invariant holds, since π maps values to values and cannot make an equality false.

What the hypothesis costs the measured arrangements. Nothing, in all three: §2.2’s flagship is a single act whose arguments are all supplied, so its trace is journal-closed outright; §5.2’s scenario is journaled nowhere, exercising the two repertoires directly to exhibit what neither can say; and §5.1’s co-binding travels inside one body, through the local environment. The hypothesis bites only where the measurement puts it — a value minted inside an operation, read out, and recorded as a later act’s argument — and what would violate (HOST) itself is different again: uncaptured I/O, escaped clocks, foreign mutable state.

The proposition’s scope has to be stated as exactly as its proof: **Proposition 4 characterizes *replayable* constitution under (HOST); it does not characterize constitution itself.** And it is doubly relative: to (HOST), and to the interpretation — replay reads $\llbracket f \rrbracket$ from outside the record, and Proposition 5 makes that second relativization explicit. Where (HOST) fails, what fails is replay equivalence — the (b)-half of §3’s innermost level — and nothing above it: DEFINE still extended $\mathbf{D}$, the definition and its exercises still stand in $\mathbf{J}$, and Observations 1 and 2 and Proposition 3 hold without the axiom, since none of their arguments touches $\llbracket f \rrbracket$. A verb whose repertoire reaches behind the plane in the ways (HOST) closes out is still defined, still exercised, still owned; its record has lost a guarantee, not the verb its existence.

Proposition 5 (what the criterion separates is how much stays ambient). Under Proposition 4’s hypotheses, $(\mathbf{J}, \llbracket \cdot \rrbracket)$ determines constitution’s replay up to $\cong$; coordination’s replay additionally depends on the ambient program $\mathbf{P}$.

Proof. The positive half is Proposition 4, which concludes $\text{replay}(\mathbf{J}) \cong \langle \mathbf{D}, \sigma, \mathbf{J} \rangle$ at the level of whole configurations — the level $\cong$ is defined at — and inherits its hypotheses, journal-closedness included. The negative half needs neither (HOST)’s freshness clauses nor anything about a particular engine: by (NONTRIV) there are operations whose outcomes differ at some state, reachable in the variant since the two calculi share their execution rules and differ only in where a definition lives; so ambient bodies $\mathbf{b} \neq \mathbf{b}'$ differing in such an operation make the same $\mathbf{J}$ replay to different states under $\mathbf{P}[\mathbf{V} \mapsto \mathbf{b}]$ and $\mathbf{P}[\mathbf{V} \mapsto \mathbf{b}']$, the interpretation held fixed.

Interpretation. Neither calculus lets the record settle its replay absolutely — re-performing reads $\llbracket \mathbf{f} \rrbracket$, code outside $\mathbf{J}$, in both — which is why §3.1’s (b) holds the interpretation and σ_0 fixed. What separates them is the remainder beyond $(\mathbf{J}, \llbracket \cdot \rrbracket)$: what moves into the record is the *composition*, never the operations, a difference in how much stays ambient rather than a categorical “record alone”, and §8 carries that remainder as a limit. Because the negative half rests only on a $\mathbf{P}$ free to vary, the separation survives any further narrowing of the positive half. On the engine: the durable-execution row of §3.2, where the history holds the exercises and replay re-walks compiled code, so the same history under changed code means something else; and the constitution row, where with the repertoires held fixed, changing nothing but the record is enough. §3.1’s (b) asks this proposition’s question of a record, and its two halves are answered here and at Proposition 4 with the measurement of §3.2.

The four questions, and the one arrow among them. Over a reachable configuration, three of §2.4’s four lines are projections: *what the subject can do* is $\text{dom}(\mathbf{D})$; *what counts as its history* is $\mathbf{J}$; *what is computed* is σ . The fourth — *what the subject does* — is deliberately not in the configuration: performances are the trace’s property, and Proposition 3 witnesses their separation at the rules themselves, one premise under two conclusions. What is claimed is exactly the separations §2.4 needs, each with its witness: equal σ with different $\mathbf{D}$ (Observation 1), so what is computed does not fix what the subject can do; equal $\mathbf{J}$ with different executions (Proposition 3’s query half), so the history does not reveal everything the subject did; equal $\mathbf{D}$ with different $\mathbf{J}$ (runs of different exercise counts), so the repertoire does not fix the history. And one positive dependency, which the paper uses: under Proposition 4’s hypotheses $\mathbf{J}$ reconstructs $\mathbf{D}$ exactly and σ up to $\cong$. That arrow is the finding rather than the table’s defect — recorded constitution *is* the arrangement in which the history suffices for the rest — and no wider independence is asserted here.

What is not claimed. No metatheory beyond these claims is developed. In particular, no equivalence between a programming model and its record-replay implementation is proven — the durable-execution neighbour proves exactly such a theorem for its model (Burckhardt et al., 2021), and this appendix’s object is smaller and its aim narrower: to carry the distinctions. The calculus is a characterization of the measured arrangement. Conformance of the engine to the rules is checked by the labs, one per claim, and checked is all it is.

References

Entries without a refereed venue — practitioner forum posts — are given with their access date and are marked as such; an archival snapshot is pinned at deposit.

The letter suffixes on this series’ own entries denote the paper’s number in the series, not the order of appearance here — 2026c is Paper 3, 2026d Paper 4, 2026h Paper 8, 2026i Paper 9 — so that a given paper carries the same label wherever in the series it is cited. In the blind review copy the four series entries and the *Dependencies* note are replaced by anonymised placeholders.

Austin, J. L. (1962). *How to do things with words*. Oxford University Press.

Bala, V., Duesterwald, E., & Banerjia, S. (2000). Dynamo: A transparent dynamic optimization system. *Proceedings of the ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI 2000)*, 1–12.

Bauer, A., & Pretnar, M. (2015). Programming with algebraic effects and handlers. *Journal of Logical and Algebraic Methods in Programming*, 84(1), 108–123.

Bernhardt, M. (2021). *Tour of Temporal: Welcome to the workflow* [Blog post; practitioner source, no refereed venue]. <https://manuel.bernhardt.io/2021/04/12/tour-of-temporal-welcome-to-the-workflow/> (accessed 17 August 2026; archival snapshot to be pinned at deposit).

Burckhardt, S., Gillum, C., Justo, D., Kallas, K., McMahon, C., & Meiklejohn, C. S. (2021). Durable functions: Semantics for stateful serverless. *Proceedings of the ACM on Programming Languages*, 5(OOPSLA), Article 133. <https://doi.org/10.1145/3485510>

Bykov, S., & Fateev, M. (2021). *Temporal vs Akka and Lagom* [Forum discussion; practitioner source, no refereed venue]. Temporal Community Forum. <https://community.temporal.io/t/temporal-vs-akka-and-lagom/2589> (accessed 17 August 2026; archival snapshot to be pinned at deposit).

Bykov, S., Geller, A., Kliot, G., Larus, J. R., Pandya, R., & Thelin, J. (2011). Orleans: Cloud computing for everyone. *Proceedings of the 2nd ACM Symposium on Cloud Computing (SoCC 2011)*, Article 16.

Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1), 37–46.

Cypher, A. (Ed.). (1993). *Watch what I do: Programming by demonstration*. MIT Press.

Dietz, J. L. G. (2006). *Enterprise ontology: Theory and methodology*. Springer.

Evans, E. (2003). *Domain-driven design: Tackling complexity in the heart of software*. Addison-Wesley.

- Fikes, R. E., Hart, P. E., & Nilsson, N. J. (1972). Learning and executing generalized robot plans. *Artificial Intelligence*, 3(4), 251–288.
- Fowler, M. (2002). *Patterns of enterprise application architecture*. Addison-Wesley.
- Freire, J., Koop, D., Santos, E., & Silva, C. T. (2008). Provenance for computational tasks: A survey. *Computing in Science & Engineering*, 10(3), 11–21.
- Gamma, E., Helm, R., Johnson, R., & Vlissides, J. (1994). *Design patterns: Elements of reusable object-oriented software*. Addison-Wesley.
- Garcia-Molina, H., & Salem, K. (1987). Sagas. *Proceedings of the ACM SIGMOD International Conference on Management of Data*, 249–259.
- Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3), 611–642.
- Grzybek, K. (2024). *Modular monolith with DDD* [Computer software]. GitHub. <https://github.com/kgrzybek/modular-monolith-with-ddd> (commit 91c8ef2, accessed 17 August 2026).
- Hevner, A. R., March, S. T., Park, J., & Ram, S. (2004). Design science in information systems research. *MIS Quarterly*, 28(1), 75–105.
- Hewitt, C., Bishop, P., & Steiger, R. (1973). A universal modular ACTOR formalism for artificial intelligence. *Proceedings of the 3rd International Joint Conference on Artificial Intelligence (IJCAI 1973)*, 235–245.
- Hohpe, G., & Woolf, B. (2003). *Enterprise integration patterns: Designing, building, and deploying messaging solutions*. Addison-Wesley.
- Jones, A. J. I., & Sergot, M. (1996). A formal characterisation of institutionalised power. *Logic Journal of the IGPL*, 4(3), 427–443.
- Kiczales, G., des Rivières, J., & Bobrow, D. G. (1991). *The art of the metaobject protocol*. MIT Press.
- Kitchenham, B., & Charters, S. (2007). *Guidelines for performing systematic literature reviews in software engineering* (Technical Report EBSE-2007-01). Keele University and Durham University.
- LeBlanc, T. J., & Mellor-Crummey, J. M. (1987). Debugging parallel programs with Instant Replay. *IEEE Transactions on Computers*, C-36(4), 471–482.
- Maes, P. (1987). Concepts and experiments in computational reflection. *Proceedings of the ACM Conference on Object-Oriented Programming Systems, Languages and Applications (OOPSLA 1987)*, 147–155.
- Meyer, B. (1997). *Object-oriented software construction* (2nd ed.). Prentice Hall.
- Moreau, L., & Missier, P. (Eds.). (2013). *PROV-DM: The PROV data model* (W3C Recommendation). World Wide Web Consortium.

- .NET Foundation. (2026). *eShop: A reference .NET application* [Computer software]. GitHub. <https://github.com/dotnet/eShop> (commit **9b4f943**, accessed 17 August 2026).
- Nigam, A., & Caswell, N. S. (2003). Business artifacts: An approach to operational specification. *IBM Systems Journal*, 42(3), 428–445.
- O’Callahan, R., Jones, C., Froyd, N., Huey, K., Noll, A., & Partush, N. (2017). Engineering record and replay for deployability. *Proceedings of the USENIX Annual Technical Conference (ATC 2017)*, 377–389.
- PCI Security Standards Council. (2022). *Payment Card Industry Data Security Standard: Requirements and testing procedures* (v4.0), Requirement 10.
- Peffer, K., Tuunanen, T., Rothenberger, M. A., & Chatterjee, S. (2007). A design science research methodology for information systems research. *Journal of Management Information Systems*, 24(3), 45–77.
- Plotkin, G. D. (2004). A structural approach to operational semantics. *Journal of Logic and Algebraic Programming*, 60–61, 17–139.
- Ressia, J., Gırba, T., Nierstrasz, O., Perin, F., & Renggli, L. (2014). Talents: An environment for dynamically composing units of reuse. *Software: Practice and Experience*, 44(4), 413–432. <https://doi.org/10.1002/spe.2160>
- Rivera, A. (2026c). Reactions and the partition: opt-in eventual consistency in actor-native systems. *Puppeteer Papers Series*, Paper 3 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.20792156>
- Rivera, A. (2026d). Preserving semantic continuity across actors: a tell-based approach without orchestration. *Puppeteer Papers Series*, Paper 4 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21207062>
- Rivera, A. (2026h). Inference without Authority: the three authorities that govern an output. *Puppeteer Papers Series*, Paper 8 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21499637>
- Rivera, A. (2026i). Identity precedes staging: one play, many stages. *Puppeteer Papers Series*, Paper 9 [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21894206>
- Searle, J. R. (1969). *Speech acts: An essay in the philosophy of language*. Cambridge University Press.
- Searle, J. R. (1995). *The construction of social reality*. Free Press.
- Shaw, M. (2003). Writing good software engineering research papers. *Proceedings of the 25th International Conference on Software Engineering (ICSE 2003)*, 726–736.
- Sigelman, B. H., Barroso, L. A., Burrows, M., Stephenson, P., Plakal, M., Beaver, D., Jaspán, S., & Shanbhag, C. (2010). *Dapper, a large-scale distributed systems tracing infrastructure* (Technical Report dapper-2010-1). Google.

- Smith, B. C. (1984). Reflection and semantics in Lisp. *Proceedings of the 11th ACM Symposium on Principles of Programming Languages (POPL 1984)*, 23–35.
- Stol, K.-J., & Fitzgerald, B. (2018). The ABC of software engineering research. *ACM Transactions on Software Engineering and Methodology*, 27(3), 11:1–11:51.
- Sutton, R. S., Precup, D., & Singh, S. (1999). Between MDPs and semi-MDPs: A framework for temporal abstraction in reinforcement learning. *Artificial Intelligence*, 112(1–2), 181–211.
- Taha, W., & Sheard, T. (1997). Multi-stage programming with explicit annotations. *Proceedings of the ACM SIGPLAN Symposium on Partial Evaluation and Semantics-Based Program Manipulation (PEPM 1997)*, 203–217.
- Ungar, D., & Smith, R. B. (1987). Self: The power of simplicity. *Proceedings of the ACM Conference on Object-Oriented Programming Systems, Languages and Applications (OOPSLA 1987)*, 227–242.
- U.S. Department of Health and Human Services. (2003). Security standards for the protection of electronic protected health information: Audit controls. 45 C.F.R. § 164.312(b).
- van der Aalst, W. M. P. (2016). *Process mining: Data science in action* (2nd ed.). Springer.
- van der Aalst, W. M. P., & van Hee, K. (2002). *Workflow management: Models, methods, and systems*. MIT Press.